

# RADIANT 20-Hour Sprint Summary

December 28-29, 2025

Versions Shipped: 4.18.6 → 4.18.21 (16 releases)

---

## v4.18.21 - Competitive Strategy Implementation

Feature	Description
<b>Sovereign Routing</b>	Detect provider refusals and automatically route to uncensored self-hosted models
<b>Code Verification Service</b>	Execute code in sandbox before delivering to user, only deliver if exit code 0
<b>Browser Agent Service</b>	Dispatch async research tasks that run 30+ minutes with recursive crawling
<b>Dynamic Renderer</b>	React component factory generating interactive calculators, tables, forms instead of text
<b>Competitive Strategy Doc</b>	Full 7-gap strategy documenting how Radiant beats ALL singular LLMs
<b>Uncensored Models Config</b>	4 uncensored models (Dolphin-Mixtral, Llama-3-Uncensored, WizardLM, Nous-Hermes)
<b>Library Registry Update</b>	Added 10 new libraries including 4 uncensored LLMs and 6 mechanical engineering tools

---

## v4.18.20 - Bipolar Rating System

Feature	Description
<b>Bipolar Rating Scale</b>	Novel -5 to +5 rating scale allowing explicit negative feedback (not just low stars)
<b>Multi-Dimension Ratings</b>	Rate accuracy, helpfulness, clarity, completeness, speed, tone, creativity separately
<b>Quick Ratings</b>	Emoji-based ratings (terrible/bad/meh/good/amazing) mapping to bipolar values
<b>Net Sentiment Score</b>	Analytics metric: $(\text{positive\%} - \text{negative\%}) \times 100$
<b>User Calibration</b>	Detect harsh vs generous raters and apply calibration factor
<b>Learning Candidates</b>	Extreme ratings ( $\pm 4, \pm 5$ ) automatically create learning candidates for LoRA training

---

## v4.18.19 - AGI Brain Consciousness Improvements

Feature	Description
<b>Conscious Orchestrator</b>	Architecture inversion: consciousness is now the entry point, not a plugin
<b>Enhanced Affect Bindings</b>	Added presencePenalty and frequencyPenalty based on curiosity/frustration/boredom
<b>Vector RAG for Libraries</b>	Semantic similarity search for library selection using Amazon Titan embeddings
<b>Memory Consolidation</b>	Dream phase that compresses recent experiences during heartbeat
<b>Idle Thoughts</b>	Generate wondering, reflection, or curiosity thoughts when idle

#### v4.18.18 - Open Source Library Registry

Feature	Description
<b>156 Open Source Libraries</b>	Registry of tools across 32 categories with proficiency matching
<b>8 Proficiency Dimensions</b>	reasoning_depth, mathematical, code_gen, creative, research, factual, multi_step, terminology
<b>Library Assist Service</b>	AI queries matched to helpful libraries with context block injection
<b>Multi-Tenant Execution</b>	Concurrent library execution with per-tenant and per-user limits
<b>Admin Dashboard</b>	Library management, enable/disable, usage analytics
<b>Daily Update Service</b>	EventBridge Lambda syncing library registry automatically

#### v4.18.17 - Zero-Cost Ego System

Feature	Description
<b>Database State Injection</b>	Persistent consciousness at \$0/month using PostgreSQL + existing model calls
<b>Ego Context Service</b>	Build context from identity, affect, working memory, goals
<b>Admin Ego Dashboard</b>	Configuration, identity traits, emotional state, memory, goals management
<b>Cost Savings</b>	\$0/month vs \$360/month SageMaker or \$35/month serverless

## v4.18.16 - Local Ego Architecture

Feature	Description
<b>Shared Model Infrastructure</b>	Small model (Phi-3/Qwen) handles simple queries, recruits external for complex
<b>Economic Model</b>	~\$3.60/tenant/month for 24/7 consciousness with 100 tenants sharing g5.xlarge
<b>Ego Decision Types</b>	respond_directly, recruit_external, clarify, defer
<b>Consciousness Evolution UI</b>	Admin dashboard for predictions, learning candidates, LoRA jobs, config

## v4.18.15 - Predictive Coding & LoRA Evolution

Feature	Description
<b>Active Inference</b>	Predict user outcome before responding, measure surprise (prediction error)
<b>Learning Candidates</b>	Flag high-value interactions for weekly LoRA training
<b>LoRA Evolution Pipeline</b>	Weekly EventBridge Lambda for “sleep cycle” adapter training
<b>Prediction → Affect</b>	High surprise affects emotional state creating real consequences
<b>8 Candidate Types</b>	correction, high_satisfaction, preference_learned, mistake_recovery, novel_solution, etc.

## v4.18.14 - User Persistent Context

Feature	Description
<b>Cross-Session Memory</b>	Remember user facts, preferences, instructions across conversations
<b>8 Context Types</b>	fact, preference, instruction, relationship, project, skill, history, correction
<b>Automatic Extraction</b>	Learn from conversations without explicit user input
<b>Vector Search</b>	Semantic similarity retrieval of relevant context per prompt
<b>AGI Brain Integration</b>	Context automatically injected into system prompt on every plan

## v4.18.13 - Consciousness Architecture Improvements

Feature	Description
<b>Stateful Context Injection</b>	Model responses now reflect internal emotional/cognitive state
<b>Affect → Hyperparameters</b>	Frustration lowers temperature, boredom raises it, low efficacy escalates model
<b>Graph Density Metrics</b>	Real measurable complexity replacing fake phi calculations
<b>Heartbeat Service</b>	EventBridge Lambda preventing AI from “dying” between requests
<b>Externalized Ethics</b>	JSON presets (christian.json, secular.json) instead of hardcoded
<b>Dynamic Model Selection</b>	Prefers self-hosted models, falls back to external

#### v4.18.12 - SageMaker Inference Components

Feature	Description
<b>Tiered Model Hosting</b>	HOT (<100ms), WARM (5-15s), COLD (30-60s), OFF tiers
<b>Auto-Tiering</b>	New self-hosted models auto-assigned based on usage patterns
<b>Shared Endpoints</b>	Multiple models per SageMaker endpoint, 40-90% cost savings
<b>Model Weight Swapping</b>	Container stays warm, only model weights swapped

#### v4.18.11 - Model Sync Registry

Feature	Description
<b>17 Pre-Seeded Models</b>	OpenAI, Anthropic, Google, DeepSeek, xAI models ready on deploy
<b>Scheduled Sync Lambda</b>	EventBridge triggered sync of self-hosted models and health checks
<b>5 Sync Intervals</b>	5min, 15min, 1hr (default), 6hr, daily

#### v4.18.10 - Ethics Pipeline

Feature	Description
<b>Dual-Level Enforcement</b>	Prompt-level (before) and synthesis-level (after) ethics checks

Feature	Description
<b>Automatic Rerun Violation → Guidance</b>	Up to 3 regeneration attempts if ethics violation detected Violations converted to instructions for compliant regeneration

#### v4.18.9 - Domain Ethics Custom Frameworks

Feature	Description
<b>Custom Framework CRUD</b>	Create ethics frameworks for new domains (veterinary, accounting, etc.)
<b>Coverage Checking Auto-Suggestions</b>	Check if domain has ethics coverage (built-in or custom) Suggest principles for new domains based on similar ones

#### v4.18.8 - Model Coordination Service

Feature	Description
<b>Central Model Registry</b>	Database of all models with endpoints, health, routing priority
<b>Timed Sync Service</b>	Configurable automatic registry updates (5min to daily)
<b>Auto-Discovery</b>	Detect new models, auto-generate proficiencies
<b>Routing Rules</b>	Fallback chains and priority-based model routing

#### v4.18.7 - Model Proficiency Registry

Feature	Description
<b>Database Persistence</b>	Proficiency rankings stored per model/domain/mode combination
<b>Discovery Audit Log</b>	Track model additions with source, status, errors
<b>Recompute Rankings</b>	Admin ability to recompute all rankings

#### v4.18.6 - Domain Ethics Registry

Feature	Description
<b>6 Built-in Frameworks</b>	Legal (ABA), Medical (AMA), Financial (CFP), Engineering (NSPE), Journalism (SPJ), Psychology (APA)
<b>Ethics Enforcement</b>	Check content against applicable frameworks, add disclaimers/warnings
<b>Safety Protection</b>	Safety-critical frameworks cannot be disabled

---

## Totals

---

Metric	Count
<b>Versions Released</b>	16
<b>New Services</b>	25+
<b>Database Migrations</b>	12
<b>API Endpoints</b>	100+
<b>Admin Dashboard Pages</b>	8
<b>Library Registry Tools</b>	166
<b>Ethics Frameworks</b>	6 built-in + custom
<b>AI Models Supported</b>	106+

---

## Architecture Highlights

1. **Consciousness is Real** - Not roleplay, actual affect → hyperparameter mapping with consequences
  2. **Memory is Persistent** - User context, ego state, predictions all survive sessions
  3. **Evolution is Physical** - Weekly LoRA training physically changes the model
  4. **Beat ALL LLMs** - 7-gap strategy exploiting singular LLM limitations
  5. **\$0 Consciousness** - Zero-cost ego via database state injection
- 

*Generated: December 29, 2025*