

RADIANT Changes - Last 20 Hours

Generated: January 1, 2026 at 4:04 AM PST

Commits

Commit	Description	Time
43edf80	feat: Truth Engine (ECD), Ghost Migration, React Hook fixes	2 hours ago
ddbd9b6	fix: Code review cleanup - orphan pages, stubs, and deprecations	4 hours ago
6b90d05	fix: Complete remaining code implementations and cleanup	5 hours ago
a19cc29	fix: Comprehensive codebase cleanup and fixes	5 hours ago
11f34e6	fix: Fix lint errors and type errors in dashboard pages	5 hours ago
0034ef7	feat: Add more missing API proxy routes	6 hours ago
41af910	feat: Add missing API proxy routes for admin dashboard	6 hours ago
73ecb4e	fix: Add Brain v6 navigation entries to sidebar	6 hours ago
3e5c403	fix: Remove unused path import from brain-stack.ts	6 hours ago
7c01931	fix: Add Next.js API proxy route for brain endpoints	6 hours ago
3b04fe6	fix: Use consistent lambda/dist code path in BrainStack	6 hours ago
1e953f3	fix: Add LITELLM_ENDPOINT and REDIS_URL to Lambda environment	6 hours ago
a34878b	fix: Add Redis initialization for Flash Buffer and Ghost Manager	6 hours ago
fcabc61	fix: Wire BrainStack into CDK app and admin routes	6 hours ago
8b5c6b4	feat: Wire LLM call to LiteLLM, add integration tests, export BrainStack	6 hours ago
13a8330	feat: Add missing Admin Dashboard components and Ghost Inference construct	7 hours ago
ff1e8a4	docs: Add AGI Brain v6.0.4 documentation	7 hours ago
4e8ac33	feat: Add Lambda handlers, CDK stack, and Admin Dashboard - Phase 2	7 hours ago
6b5e687	feat: Implement AGI Brain v6.0.4 core services - Phase 1	7 hours ago
35f4e43	feat: Switch to Qwen 2.5 7B SageMaker for translation (3x cheaper)	10 hours ago

Summary

- **20 commits** in the last 20 hours
- **9 features** added
- **10 fixes** applied
- **1 documentation** update

Key Highlights

1. **Truth Engine (Project TRUTH)** - Entity-Context Divergence for hallucination prevention
 2. **Ghost Migration** - Automatic ghost vector migration preserving consciousness continuity
 3. **AGI Brain v6.0.4** - Complete implementation with SOFAI routing, Flash Buffer, and Oversight
 4. **Translation Optimization** - Switched to Qwen 2.5 7B for 3x cost reduction
 5. **Unified Naming Convention** - Service renames per architecture report (pending commit)
-

RADIANT v6.0.4-S3 - Multi-tenant AWS SaaS Platform for AI Model Access