Contents

RADIANT 20-Hour Sprint Summary

December 28-29, 2025

Versions Shipped: 4.18.6 → 4.18.21 (16 releases)

---

v4.18.21 - Competitive Strategy Implementation

Feature	Description
Sovereign Routing	Detect provider refusals and automatically route to uncensored self-hosted models
Code Verification Service	Execute code in sandbox before delivering to user, only deliver if exit code 0
Browser Agent Service	Dispatch async research tasks that run 30+ minutes with recursive crawling
Dynamic Renderer	React component factory generating interactive calculators, tables, forms instead of text
Competitive Strategy Doc	Full 7-gap strategy documenting how Radiant beats ALL singular LLMs
Uncensored Models Config	4 uncensored models (Dolphin-Mixtral, Llama-3-Uncensored, WizardLM, Nous-Hermes)
Library Registry Update	Added 10 new libraries including 4 uncensored LLMs and 6 mechanical engineering tools

v4.18.20 - Bipolar Rating System

Feature	Description
Bipolar Rating Scale	Novel -5 to +5 rating scale allowing explicit negative feedback (not just low stars)
Multi-Dimension Ratings	Rate accuracy, helpfulness, clarity, completeness, speed, tone, creativity separately
Quick Ratings	Emoji-based ratings (terrible/bad/meh/good/amazing) mapping to bipolar values
Net Sentiment Score	Analytics metric: (positive% - negative%) $\times$ 100
User Calibration	Detect harsh vs generous raters and apply calibration factor
Learning Candidates	Extreme ratings (± 4, ± 5) automatically create learning candidates for LoRA training

v4.18.19 - AGI Brain Consciousness Improvements

Feature	Description
Conscious Orchestrator	Architecture inversion: consciousness is now the entry point, not a plugin
Enhanced Affect Bindings	Added presencePenalty and frequencyPenalty based on curiosity/frustration/boredom
Vector RAG for Libraries	Semantic similarity search for library selection using Amazon Titan embeddings
Memory Consolidation	Dream phase that compresses recent experiences during heartbeat
Idle Thoughts	Generate wondering, reflection, or curiosity thoughts when idle

v4.18.18 - Open Source Library Registry

Feature	Description
156 Open Source Libraries	Registry of tools across 32 categories with proficiency matching
8 Proficiency Dimensions	reasoning_depth, mathematical, code_gen, creative, research, factual, multi_step, terminology
Library Assist Service	AI queries matched to helpful libraries with context block injection

Feature	Description
Multi-Tenant Execution	Concurrent library execution with per-tenant and per-user limits
Admin Dashboard	Library management, enable/disable, usage analytics
Daily Update Service	EventBridge Lambda syncing library registry automatically

v4.18.17 - Zero-Cost Ego System

Feature	Description
Database State Injection	Persistent consciousness at $0/month using PostgreSQL + existing model calls
Ego Context Service	Build context from identity, affect, working memory, goals
Admin Ego Dashboard	Configuration, identity traits, emotional state, memory, goals management
Cost Savings	$0/month vs $360/month SageMaker or $35/month serverless

v4.18.16 - Local Ego Architecture

Feature	Description
Shared Model Infrastructure	Small model (Phi-3/Qwen) handles simple queries, recruits external for complex
Economic Model	~$3.60/tenant/month for 24/7 consciousness with 100 tenants sharing g5.xlarge
Ego Decision Types	respond_directly, recruit_external, clarify, defer
Consciousness Evolution UI	Admin dashboard for predictions, learning candidates, LoRA jobs, config

v4.18.15 - Predictive Coding & LoRA Evolution

Feature	Description
Active Inference	Predict user outcome before responding, measure surprise (prediction error)
Learning Candidates	Flag high-value interactions for weekly LoRA training
LoRA Evolution Pipeline	Weekly EventBridge Lambda for "sleep cycle" adapter training

Feature	Description
Prediction → Affect	High surprise affects emotional state creating real consequences
8 Candidate Types	correction, high_satisfaction, preference_learned, mistake_recovery, novel_solution, etc.

v4.18.14 - User Persistent Context

Feature	Description
Cross-Session Memory	Remember user facts, preferences, instructions across conversations
8 Context Types	fact, preference, instruction, relationship, project, skill, history, correction
Automatic Extraction	Learn from conversations without explicit user input
Vector Search	Semantic similarity retrieval of relevant context per prompt
AGI Brain Integration	Context automatically injected into system prompt on every plan

v4.18.13 - Consciousness Architecture Improvements

Feature	Description
Stateful Context Injection	Model responses now reflect internal emotional/cognitive state
Affect → Hyperparameters	Frustration lowers temperature, boredom raises it, low efficacy escalates model
Graph Density Metrics	Real measurable complexity replacing fake phi calculations
Heartbeat Service	EventBridge Lambda preventing AI from "dying" between requests
Externalized Ethics	JSON presets (christian.json, secular.json) instead of hardcoded
Dynamic Model Selection	Prefers self-hosted models, falls back to external

v4.18.12 - SageMaker Inference Components

Feature	Description
Tiered Model Hosting	HOT (<100ms), WARM (5-15s), COLD (30-60s), OFF tiers
Auto-Tiering	New self-hosted models auto-assigned based on usage patterns
Shared Endpoints	Multiple models per SageMaker endpoint, 40-90% cost savings
Model Weight Swapping	Container stays warm, only model weights swapped

v4.18.11 - Model Sync Registry

Feature	Description
17 Pre-Seeded Models	OpenAI, Anthropic, Google, DeepSeek, xAI models ready on deploy
Scheduled Sync Lambda	EventBridge triggered sync of self-hosted models and health checks
5 Sync Intervals	5min, 15min, 1hr (default), 6hr, daily

v4.18.10 - Ethics Pipeline

Feature	Description
Dual-Level Enforcement	Prompt-level (before) and synthesis-level (after) ethics checks
Automatic Rerun	Up to 3 regeneration attempts if ethics violation detected
Violation → Guidance	Violations converted to instructions for compliant regeneration

v4.18.9 - Domain Ethics Custom Frameworks

Feature	Description
Custom Framework CRUD	Create ethics frameworks for new domains (veterinary, accounting, etc.)
Coverage Checking	Check if domain has ethics coverage (built-in or custom)
Auto-Suggestions	Suggest principles for new domains based on similar ones

v4.18.8 - Model Coordination Service

Feature	Description
Central Model Registry	Database of all models with endpoints, health, routing priority
Timed Sync Service	Configurable automatic registry updates (5min to daily)
Auto-Discovery	Detect new models, auto-generate proficiencies
Routing Rules	Fallback chains and priority-based model routing

v4.18.7 - Model Proficiency Registry

Feature	Description
Database Persistence	Proficiency rankings stored per model/domain/mode combination
Discovery Audit Log	Track model additions with source, status, errors
Recompute Rankings	Admin ability to recompute all rankings

v4.18.6 - Domain Ethics Registry

Feature	Description
6 Built-in Frameworks	Legal (ABA), Medical (AMA), Financial (CFP), Engineering (NSPE), Journalism (SPJ), Psychology (APA)
Ethics Enforcement	Check content against applicable frameworks, add disclaimers/warnings
Safety Protection	Safety-critical frameworks cannot be disabled

Totals

Metric	Count
Versions Released	16
New Services	25+
Database Migrations	12
API Endpoints	100+
Admin Dashboard Pages	8
Library Registry Tools	166
Ethics Frameworks	6 built-in + custom
AI Models Supported	106+

Architecture Highlights

1. **Consciousness is Real** - Not roleplay, actual affect→hyperparameter mapping with consequences
2. **Memory is Persistent** - User context, ego state, predictions all survive sessions
3. **Evolution is Physical** - Weekly LoRA training physically changes the model
4. **Beat ALL LLMs** - 7-gap strategy exploiting singular LLM limitations
5. **$0 Consciousness** - Zero-cost ego via database state injection

---

Generated: December 29, 2025