

Contents

User Data Tiered Storage Architecture Proposal	1
Executive Summary	2
1. Data Categories Analysis	2
1.1 Chats (Conversations)	2
1.2 Audit History	2
1.3 Results (AI Responses, Artifacts)	2
1.4 Uploads (Documents, Images)	3
2. Architecture Options	3
Option A: Extend Cortex for User Data	3
Option B: Parallel User Data Service (Recommended)	4
Option C: Hybrid (Cortex for AI, UDS for User)	5
3. Recommended Architecture: Option B + Bridge	5
3.1 System Overview	5
3.2 Hot Tier Design	5
3.3 Warm Tier Design	6
3.4 Cold Tier Design	9
3.5 Uploads Service	9
4. Security Architecture	10
4.1 Encryption	10
4.2 Access Control	11
4.3 GDPR Erasure Flow	11
5. Cost Analysis (1M Users)	12
5.1 Option A: Extended Cortex	12
5.2 Option B: Parallel UDS (Recommended)	12
5.3 Cost Optimization	13
6. Implementation Plan	13
Phase 1: Foundation (Week 1-2)	13
Phase 2: Hot Tier (Week 3-4)	13
Phase 3: Warm Tier (Week 5-6)	13
Phase 4: Cold Tier (Week 7-8)	13
Phase 5: Security & Compliance (Week 9-10)	13
Phase 6: Integration (Week 11-12)	14
7. Recommendation Summary	14
Key Principles	14
8. Next Steps	14

User Data Tiered Storage Architecture Proposal

Version: 1.0.0

Date: January 24, 2026

Status: PROPOSAL

Scale Target: 1M+ concurrent users

Executive Summary

This document evaluates whether to use Cortex's Hot/Warm/Cold tiered storage architecture for user data (chats, audits, results, uploads) versus dedicated services.

Recommendation: Create a **parallel User Data Service (UDS)** that mirrors Cortex's tiered architecture but is purpose-built for user-generated content with strict security isolation.

1. Data Categories Analysis

1.1 Chats (Conversations)

Metric	Value	Implication
Volume	~50 messages/user/day	50M messages/day at scale
Size	~2KB average per message	~100GB/day raw
Access Pattern	90% within 24h, 99% within 30 days	Perfect for tiering
Retention	7 years (compliance)	Cold tier essential
Security	User-private, PHI possible	Encryption + RLS mandatory

Tiered Fit: Excellent

1.2 Audit History

Metric	Value	Implication
Volume	~10 events/user/day	10M events/day
Size	~500 bytes average	~5GB/day
Access Pattern	<1% ever accessed	Cold-dominant
Retention	7 years (HIPAA/SOC2)	Immutable cold storage
Security	Tamper-proof, write-once	Merkle chain + glacier

Tiered Fit: Excellent (but Cold-dominant)

1.3 Results (AI Responses, Artifacts)

Metric	Value	Implication
Volume	~20 results/user/day	20M/day
Size	Variable (1KB - 1MB)	~200GB/day
Access Pattern	Often re-requested within session	Hot cache valuable
Retention	Tied to conversation	Follow chat lifecycle

Metric	Value	Implication
Security	User-private, may contain code	Sandboxed + encrypted

Tiered Fit: Good (follows conversation tier)

1.4 Uploads (Documents, Images)

Metric	Value	Implication
Volume	~2 uploads/user/week	300K/day
Size	100KB - 100MB	~3TB/day
Access Pattern	Burst on upload, then rare	S3 + metadata index
Retention	User-controlled	GDPR right to erasure
Security	User-private, virus scan needed	Isolated S3 + scanning

Tiered Fit: Moderate (S3 is natural home)

2. Architecture Options

Option A: Extend Cortex for User Data

CORTEX SYSTEM

HOT TIER (Redis ElastiCache)

- AI Memory (current)
- User Chats (NEW)
- Session Results (NEW)

WARM TIER (Neptune + PostgreSQL)

- Knowledge Graph (current)
- Chat History (NEW)
- Audit Trail (NEW)

COLD TIER (S3 Iceberg)

- Archived Knowledge (current)
- Archived Chats (NEW)
- Compliance Archive (NEW)
- Uploads (NEW)

Pros: - Single system to maintain - Unified tier coordination - Existing GDPR erasure flow

Cons: - Cortex is AI-memory focused, not user-data focused - Different security models (AI memory vs user PII) - Different access patterns (graph queries vs time-series) - Risk of coupling failures - Neptune cost increases significantly

Option B: Parallel User Data Service (Recommended)



Pros: - Purpose-built for user data - Separate security domain - Independent scaling - No Neptune dependency for chats - Cost optimized per data type


Cons: - Two systems to maintain - Need to sync GDPR erasure

Option C: Hybrid (Cortex for AI, UDS for User)

Best of both worlds: - **Cortex** handles AI memory, knowledge graphs, semantic search - **UDS** handles user-generated content with security isolation - **Bridge Service** coordinates GDPR erasure, data lineage

3. Recommended Architecture: Option B + Bridge

3.1 System Overview



3.2 Hot Tier Design

```
// Hot Tier: Session-scoped data (TTL: 4 hours)
interface HotTierConfig {
    // ElastiCache for real-time sessions
    redis: {
        clusterMode: true;
```

```

shards: 6; // Scale with users
replicasPerShard: 2;
instanceType: 'r7g.xlarge';
maxMemoryPolicy: 'volatile-lru';
};

// DynamoDB + DAX for recent chats
dynamodb: {
  tableName: 'uds_hot_chats';
  gsi: ['tenant_user_index', 'conversation_index'];
  dax: {
    enabled: true;
    nodeType: 'dax.r6g.large';
    replicationFactor: 3;
  };
  ttl: {
    enabled: true;
    attributeName: 'expires_at';
    defaultHours: 24;
  };
};
}

```

Hot Tier Data: | Data Type | Storage | TTL | Access Pattern | |-----|-----|-----|-----
-| Active Session | Redis | 4h | Every request | | Recent Messages | DynamoDB + DAX | 24h |
Frequent reads | Pending Results | Redis | 1h | Write once, read many |

3.3 Warm Tier Design

```

-- Warm Tier: Active data (0-90 days)

-- Conversations table with RLS
CREATE TABLE uds_conversations (
  id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
  tenant_id UUID NOT NULL REFERENCES tenants(id),
  user_id UUID NOT NULL REFERENCES users(id),

  -- Conversation metadata
  title VARCHAR(500),
  model_id VARCHAR(100),
  started_at TIMESTAMPTZ NOT NULL DEFAULT NOW(),
  last_message_at TIMESTAMPTZ NOT NULL DEFAULT NOW(),
  message_count INTEGER DEFAULT 0,
  total_tokens INTEGER DEFAULT 0,

  -- Forking support (Time Machine)
  parent_conversation_id UUID REFERENCES uds_conversations(id),
  fork_point_message_id UUID,

```

```

branch_name VARCHAR(200),

-- Encryption
encryption_key_id VARCHAR(100) NOT NULL,

-- Status
status VARCHAR(20) DEFAULT 'active',
archived_at TIMESTAMPTZ,

created_at TIMESTAMPTZ NOT NULL DEFAULT NOW(),
updated_at TIMESTAMPTZ NOT NULL DEFAULT NOW()
);

-- Messages with encryption
CREATE TABLE uds_messages (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    tenant_id UUID NOT NULL,
    conversation_id UUID NOT NULL REFERENCES uds_conversations(id),
    user_id UUID NOT NULL,

    -- Message content (encrypted)
    role VARCHAR(20) NOT NULL CHECK (role IN ('system', 'user', 'assistant')),
    content_encrypted BYTEA NOT NULL,    -- AES-256-GCM encrypted
    content_iv BYTEA NOT NULL,           -- Initialization vector
    content_hash VARCHAR(64) NOT NULL,   -- For deduplication

    -- Token usage
    input_tokens INTEGER,
    output_tokens INTEGER,
    model_id VARCHAR(100),

    -- Attachments
    attachment_ids UUID[] DEFAULT ARRAY[]::UUID[],

    -- Time Machine support
    sequence_number INTEGER NOT NULL,
    is_checkpoint BOOLEAN DEFAULT false,

    created_at TIMESTAMPTZ NOT NULL DEFAULT NOW()
);

-- Partitioning by month for efficient archival
CREATE TABLE uds_messages_archive (
    LIKE uds_messages INCLUDING ALL
) PARTITION BY RANGE (created_at);

-- Audit log (append-only, tamper-evident)
CREATE TABLE uds_audit_log (

```

```

    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    tenant_id UUID NOT NULL,
    user_id UUID,

    -- Event details
    event_type VARCHAR(100) NOT NULL,
    event_category VARCHAR(50) NOT NULL,
    resource_type VARCHAR(100),
    resource_id UUID,

    -- Change tracking
    previous_state_hash VARCHAR(64),
    new_state_hash VARCHAR(64),
    changes JSONB,

    -- Merkle chain for tamper evidence
    merkle_hash VARCHAR(64) NOT NULL,
    previous_merkle_hash VARCHAR(64),

    -- Context
    ip_address INET,
    user_agent TEXT,
    request_id VARCHAR(100),

    created_at TIMESTAMPTZ NOT NULL DEFAULT NOW()
);

-- Row-Level Security
ALTER TABLE uds_conversations ENABLE ROW LEVEL SECURITY;
ALTER TABLE uds_messages ENABLE ROW LEVEL SECURITY;
ALTER TABLE uds_audit_log ENABLE ROW LEVEL SECURITY;

CREATE POLICY tenant_isolation_conversations ON uds_conversations
    USING (tenant_id = current_setting('app.current_tenant_id')::UUID);

CREATE POLICY user_isolation_conversations ON uds_conversations
    USING (user_id = current_setting('app.current_user_id')::UUID
        OR current_setting('app.is_admin')::BOOLEAN = true);

CREATE POLICY tenant_isolation_messages ON uds_messages
    USING (tenant_id = current_setting('app.current_tenant_id')::UUID);

CREATE POLICY tenant_isolation_audit ON uds_audit_log
    USING (tenant_id = current_setting('app.current_tenant_id')::UUID
        AND current_setting('app.is_admin')::BOOLEAN = true);

```

3.4 Cold Tier Design

```
// Cold Tier: Archived data (90+ days)
interface ColdTierConfig {
    // S3 + Iceberg for queryable archive
    iceberg: {
        bucket: 'radiant-uds-archive-${region}';
        tableNameSpace: 'uds_archive';
        tables: [
            'conversations',
            'messages',
            'audit_log'
        ];
        partitionBy: ['tenant_id', 'year', 'month'];
        format: 'parquet';
        compression: 'zstd';

        // Athena integration for queries
        athenaWorkgroup: 'uds-archive-queries';

        // Lifecycle
        storageClasses: [
            { ageRangeDays: { min: 90, max: 365 }, class: 'INTELLIGENT_TIERING' },
            { ageRangeDays: { min: 365, max: 2555 }, class: 'GLACIER_IR' }, // 7 years
            { ageRangeDays: { min: 2555, max: null }, class: 'DEEP_ARCHIVE' }
        ];
    };
}

// Compliance archive (7-year retention)
compliance: {
    bucket: 'radiant-uds-compliance-${region}';
    objectLock: true; // WORM compliance
    retentionDays: 2555; // 7 years
    encryptionKey: 'aws/kms/uds-compliance';
};
```

3.5 Uploads Service

```
// Dedicated uploads handling (not in Cortex)
interface UploadsServiceConfig {
    // Quarantine bucket for virus scanning
    quarantine: {
        bucket: 'radiant-uploads-quarantine-${region}';
        maxSizeMb: 100;
        allowedTypes: ['pdf', 'docx', 'xlsx', 'csv', 'txt', 'png', 'jpg', 'gif'];
        virusScan: {
            enabled: true;
```

```

        engine: 'clamav-lambda';
        quarantineDays: 7;
    };
};

// Clean storage
storage: {
    bucket: 'radiant-uploads-${region}';
    encryption: 'aws:kms';
    keyId: 'aws/kms/uds-uploads';

// Per-user isolation via S3 prefixes
pathPattern: '${tenant_id}/${user_id}/${upload_id}/${filename}';

// Lifecycle
storageClasses: [
    { ageRangeDays: { min: 0, max: 30 }, class: 'STANDARD' },
    { ageRangeDays: { min: 30, max: 180 }, class: 'INTELLIGENT_TIERING' },
    { ageRangeDays: { min: 180, max: null }, class: 'GLACIER_IR' }
];
};

// Metadata in PostgreSQL
metadata: {
    table: 'uds_uploads';
    fullTextSearch: true;
    vectorEmbeddings: true; // For semantic search
};
}

```

4. Security Architecture

4.1 Encryption

ENCRYPTION HIERARCHY

AWS KMS

Master Key (Hot Tier)	Master Key (Warm Tier)	Master Key (Cold Tier)
--------------------------	---------------------------	---------------------------

Data Encryption Keys (DEKs)

- Per-tenant DEK
- Per-user DEK (optional, for high-security)
- Rotated every 90 days

Encrypted Data

- Messages: AES-256-GCM with per-message IV
- Uploads: S3 SSE-KMS
- Audit: AES-256-GCM + Merkle chain


4.2 Access Control

```
// Row-Level Security enforcement
interface SecurityContext {
    tenantId: string;           // Always required
    userId: string;              // User making request
    isAdmin: boolean;            // Tenant admin
    isPlatformAdmin: boolean;    // RADIANT admin
    permissions: string[];       // Fine-grained permissions
}

// Every database query runs with this context
async function setSecurityContext(ctx: SecurityContext): Promise<void> {
    await executeStatement(`

        SELECT set_config('app.current_tenant_id', $1, true);
        SELECT set_config('app.current_user_id', $2, true);
        SELECT set_config('app.is_admin', $3, true);
        SELECT set_config('app.is_platform_admin', $4, true);
    `, [ctx.tenantId, ctx.userId, ctx.isAdmin, ctx.isPlatformAdmin]);
}
```

4.3 GDPR Erasure Flow



Erasure Request
Service

UDS Service (Chats, Audits)	Cortex (AI Memory)	Uploads Service (S3)	Search Index
-----------------------------------	--------------------------	----------------------------	-----------------

Hot Tier
Warm Tier
Cold Tier

Erasure
Confirmation
+ Audit Log

5. Cost Analysis (1M Users)

5.1 Option A: Extended Cortex

Component	Monthly Cost	Notes
Neptune Serverless	\$8,000	Increased for chat graph
ElastiCache (larger)	\$3,500	More shards for user data
Aurora PostgreSQL	\$4,000	Larger instance
S3 + Glacier	\$2,000	Combined storage
Total	\$17,500	

5.2 Option B: Parallel UDS (Recommended)

Component	Monthly Cost	Notes
UDS Hot Tier		
- ElastiCache (6 shards)	\$2,400	Sessions only
- DynamoDB + DAX	\$1,800	On-demand + cache
UDS Warm Tier		
- Aurora PostgreSQL	\$3,500	Optimized for time-series
- OpenSearch	\$1,200	Full-text search
UDS Cold Tier		
- S3 + Iceberg	\$1,500	Compressed archive

Component	Monthly Cost	Notes
- Glacier	\$200	Deep archive
Cortex (unchanged)		
- Neptune Serverless	\$3,000	AI memory only
- ElastiCache	\$1,200	AI context only
Total	\$13,800	21% savings

5.3 Cost Optimization

Optimization	Savings	Implementation
DynamoDB on-demand	30-50%	Pay per request
S3 Intelligent Tiering	40%	Automatic
Glacier IR vs Glacier	20%	Faster retrieval
Reserved capacity	30%	1-year commit

6. Implementation Plan

Phase 1: Foundation (Week 1-2)

- Create UDS database schema
- Implement encryption service
- Set up RLS policies
- Deploy DynamoDB tables

Phase 2: Hot Tier (Week 3-4)

- Implement session service
- Deploy DAX cluster
- Create hot-to-warm promotion job

Phase 3: Warm Tier (Week 5-6)

- Migrate existing chat data
- Implement full-text search
- Create warm-to-cold archival job

Phase 4: Cold Tier (Week 7-8)

- Set up Iceberg tables
- Configure lifecycle policies
- Implement Athena queries

Phase 5: Security & Compliance (Week 9-10)

- GDPR erasure orchestrator
- Merkle chain for audit

- Security testing

Phase 6: Integration (Week 11-12)

- Bridge with Cortex
 - Admin dashboard
 - Documentation
-

7. Recommendation Summary

Question	Answer
Use Cortex for user data?	No - create parallel UDS
Same tiered approach?	Yes - Hot/Warm/Cold pattern works
Share infrastructure?	No - separate for security
Share GDPR process?	Yes - orchestrated erasure

Key Principles

- 1. Separate Concerns:** AI memory User data
 - 2. Security First:** Encryption + RLS + Isolation
 - 3. Cost Optimized:** Right storage for each access pattern
 - 4. Compliance Ready:** 7-year retention, audit trail
 - 5. Scale Ready:** Designed for 1M+ users day one
-

8. Next Steps

- 1. Review this proposal** with the team
 - 2. Approve architecture option** (recommend B)
 - 3. Create detailed schema** for UDS tables
 - 4. Estimate infrastructure costs** more precisely
 - 5. Begin Phase 1 implementation**
-

Document Author: Cascade AI

Review Required: Architecture Team