

Contents

RADIANT Services Reference	1
Complete Lambda Services Inventory (62 Services)	1
Core Infrastructure Services	1
AI Model Services	2
Orchestration Services	5
Billing Services	7
Cognitive Services	8
Memory Services	9
AGI Services	10
Collaboration Services	10
Additional Services (30-62)	11

RADIANT Services Reference

Complete Lambda Services Inventory (62 Services)

Core Infrastructure Services

1. BrainRouter (brain-router.ts) **Purpose:** Central routing service that directs incoming requests to appropriate handlers based on task type.

Key Methods: - `routeTask(task: Task): Promise<TaskResult>` - Routes task to handler
- `analyzeTaskType(input: string): TaskType` - Determines task classification
- `selectHandler(taskType: TaskType): Handler` - Selects appropriate handler

Task Types: | Type | Description | Handler | |-----| | generation | Text generation | ModelRouterService || analysis | Data analysis | AnalyticsService || transformation | Content transformation | TransformService || orchestration | Multi-step workflow | OrchestrationService || conversation | Chat interaction | ConversationService |

2. ThermalStateService (thermal-state.ts) **Purpose:** Monitors system thermal state and adjusts workload distribution.

Key Methods: - `getSystemState(): ThermalState` - Current system state - `adjustWorkload(state: ThermalState): void` - Modify processing - `recordMetric(name: string, value: number): void` - Track metrics

States: - `nominal` - Normal operation - `elevated` - Increased load - `throttled` - Reduced capacity - `critical` - Emergency mode

3. MetricsCollector (metrics-collector.ts) **Purpose:** Collects and aggregates system metrics for monitoring.

Key Methods: - `recordLatency(service: string, ms: number): void` - `recordTokenUsage(model: string, input: number, output: number): void` - `recordCost(tenantId: string, cents: number): void` - `getMetrics(timeRange: TimeRange): MetricsSummary`

Metrics Tracked: - API latency (p50, p95, p99) - Token usage by model - Cost by tenant - Error rates - Provider health

4. ErrorLogger (error-logger.ts) **Purpose:** Structured error logging with context preservation.

Key Methods: - logError(error: Error, context: ErrorContext): void - logWarning(message: string, data: object): void - getRecentErrors(count: number): ErrorLog[]

Error Categories: - PROVIDER_ERROR - AI provider failures - VALIDATION_ERROR - Input validation - AUTH_ERROR - Authentication failures - RATE_LIMIT - Rate limiting triggered - INTERNAL_ERROR - System errors

5. CredentialsManager (credentials-manager.ts) **Purpose:** Secure management of API keys and credentials.

Key Methods: - getCredential(provider: string): Promise<string> - rotateCredential(provider: string): Promise<void> - validateCredential(provider: string): Promise<boolean>

Supported Providers: - OpenAI, Anthropic, Google, Mistral - Groq, Perplexity, xAI, Together - Cohere, DeepSeek, Replicate

AI Model Services

6. ModelRouterService (model-router.service.ts) **Purpose:** Routes AI requests to optimal provider with fallback.

Architecture:

Request → Validate → Select Provider → Execute → Fallback (if needed) → Response

Model Registry (24 Models):

Model ID	Provider	Capabilities	Cost (\$/1K tokens)
anthropic/clause-3-5-sonnet	bedrock	reasoning, coding, vision	\$0.003/\$0.015
anthropic/clause-3-haiku	bedrock	fast, efficient	\$0.00025/\$0.00125
meta/llama-3.1-70b	bedrock	reasoning, open-source	\$0.00099/\$0.00099
amazon/titan-text-express	bedrock	fast, aws-native	\$0.0002/\$0.0006
openai/gpt-4o	litellm	reasoning, vision	\$0.005/\$0.015
openai/gpt-4o-mini	litellm	fast, efficient	\$0.00015/\$0.0006
openai/o1	litellm	reasoning, math	\$0.015/\$0.060

Model ID	Provider	Capabilities	Cost (\$/1K tokens)
openai/o1-mini	litellm	reasoning, coding	\$0.003/\$0.012
google/gemini-1.5-pro	litellm	reasoning, long-context	\$0.00125/\$0.005
google/gemini-1.5-flash	litellm	fast, vision	\$0.000075/\$0.0003
mistral/mistral-large	litellm	reasoning, multilingual	\$0.003/\$0.009
mistral/codestral	litellm	coding	\$0.001/\$0.003
cohere/command-r-plus	litellm	reasoning, rag	\$0.003/\$0.015
deepseek/deepseek-coder-v2	litellm	coding	\$0.00014/\$0.00028
groq/llama-3.1-70b-versatile	groq	fast, reasoning	\$0.00059/\$0.00079
groq/llama-3.1-8b-instant	groq	instant, fast	\$0.00005/\$0.00008
groq/mixtral-8x7b	groq	fast, moe	\$0.00024/\$0.00024
perplexity/sonar-large	perplexity	search, citations	\$0.001/\$0.001
perplexity/sonar-small	perplexity	search, fast	\$0.0002/\$0.0002
xai/grok-beta	xai	reasoning, realtime	\$0.005/\$0.015
together/llama-3.1-405b	together	reasoning, large	\$0.005/\$0.015

Fallback Chains:

```

bedrock → litellm → groq
litellm → bedrock → groq
groq → litellm → bedrock
perplexity → litellm
xai → litellm → groq
together → litellm → groq

```

Provider Health Tracking: - `isHealthy`: boolean - `latencyMs`: number - `errorCount`: number - `consecutiveFailures`: number (>= 3 marks unhealthy)

7. ModelMetadataService (model-metadata.service.ts) **Purpose:** Manages live model capabilities, pricing, and availability.

Key Methods: - `getMetadata(modelId: string): Promise<ModelMetadata>` - `getAllMetadata(): Promise<ModelMetadata[]>` - `updateMetadata(modelId: string, data: Partial<ModelMetadata>): Promise<void>` - `refreshFromInternet(): Promise<void>` - AI-powered metadata updates

Metadata Structure:

```

interface ModelMetadata {
  modelId: string;
  provider: string;
  displayName: string;
  description: string;
  capabilities: {
    reasoning: number;           // 0-1 score
    coding: number;
    creative: number;
    factual: number;
    math: number;
    vision: boolean;
    longContext: boolean;
  };
  contextWindow: number;
  maxOutputTokens: number;
  pricing: {
    inputPer1kTokens: number;
    outputPer1kTokens: number;
    currency: string;
  };
  availability: {
    isAvailable: boolean;
    regions: string[];
    lastChecked: Date;
  };
  performance: {
    avgLatencyMs: number;
    throughputTokensPerSec: number;
  };
}

```

8. ModelSelectionService (model-selection-service.ts) **Purpose:** Intelligent model selection based on task characteristics.

Selection Algorithm: 1. **Domain Detection** - Identify problem domain from keywords 2. **Task Analysis** - Determine complexity, requirements 3. **Model Scoring** - Score each model for task fit 4. **Mode Assignment** - Select optimal execution mode 5. **Cost/Quality Balance** - Apply user preferences

Domain Keywords: | Domain | Keywords | |——|——| | coding | code, function, algorithm, debug, implement, API | | math | calculate, equation, formula, solve, proof | | legal | contract, law, compliance, regulation, liability | | medical | diagnosis, treatment, symptom, clinical, patient | | research | study, analyze, evidence, literature, methodology | | creative | write, story, design, brainstorm, creative |

Orchestration Services

9. OrchestrationPatternsService (`orchestration-patterns.service.ts`) Purpose: Manages 49 orchestration patterns with parameterized methods.

Pattern Categories (8):

Category	Count	Examples
Consensus & Aggregation	7	Self-Consistency, Meta-Reasoning, Mixture-of-Agents
Debate & Deliberation	7	AI Debate, Society of Mind, Socratic Dialogue
Critique & Refinement	7	Self-Refine, Reflexion, Constitutional AI
Verification & Validation	7	Chain-of-Verification, LLM-as-Judge, Fact-Check
Decomposition	7	Least-to-Most, Tree of Thoughts, Skeleton-of-Thought
Specialized Reasoning	7	Chain-of-Thought, ReAct, Self-Ask
Multi-Model Routing	4	Mixture of Experts, FrugalGPT, Cascading
Ensemble Methods	3	Model Ensemble, Blended RAG, Speculative Decoding

All 49 Patterns:

1. **Self-Consistency** - Multiple samples, majority vote
2. **Universal Self-Consistency** - Free-form answer selection
3. **Meta-Reasoning** - Compare reasoning paths
4. **DiVeRSe** - Diverse verifier ensemble
5. **Mixture-of-Agents** - Multi-agent aggregation
6. **LLM-Blender** - Pairwise ranking fusion
7. **Multi-Agent Consensus** - Agent negotiation
8. **AI Debate** - Adversarial debate with judge
9. **Multi-Agent Debate** - Multi-party debate
10. **Society of Mind** - Agent specialization
11. **ChatEval** - Multi-agent evaluation
12. **ReConcile** - Confidence-weighted discussion
13. **Socratic Dialogue** - Question-based exploration
14. **Diplomatic Consensus** - Negotiated agreement
15. **Self-Refine** - Iterative refinement
16. **Reflexion** - Verbal reinforcement learning
17. **CRITIC** - External tool verification
18. **Iterative Refinement** - Multi-pass improvement
19. **Constitutional AI** - Principle-based critique
20. **Progressive Refinement** - Staged quality improvement
21. **Expert Refinement** - Domain expert review

22. **Chain-of-Verification** - Claim verification chain
 23. **LLM-as-Judge** - Model evaluation
 24. **Self-Verification** - Self-checking
 25. **G-Eval** - Structured evaluation
 26. **Cross-Validation** - Multi-model validation
 27. **Fact-Check Chain** - Fact verification pipeline
 28. **Consensus Validation** - Agreement-based validation
 29. **Least-to-Most** - Simple to complex decomposition
 30. **Decomposed Prompting** - Sub-task breakdown
 31. **Tree of Thoughts** - Branching exploration
 32. **Graph of Thoughts** - Graph-based reasoning
 33. **Skeleton-of-Thought** - Parallel point expansion
 34. **Plan-and-Solve** - Planning then execution
 35. **Recursive Decomposition** - Hierarchical breakdown
 36. **Chain-of-Thought** - Step-by-step reasoning
 37. **Self-Ask** - Sub-question generation
 38. **ReAct** - Reasoning + Acting
 39. **Program-of-Thoughts** - Code-based reasoning
 40. **Analogical Reasoning** - Example-based reasoning
 41. **Maieutic Prompting** - Tree explanation
 42. **Contrastive CoT** - Valid/invalid contrast
 43. **Mixture of Experts** - Specialized routing
 44. **FrugalGPT** - Cost-optimized cascading
 45. **Router Chain** - Capability-based routing
 46. **Speculative Routing** - Predictive routing
 47. **Model Ensemble** - Multi-model combination
 48. **Blended RAG** - RAG ensemble
 49. **Speculative Decoding** - Draft-verify acceleration
-

10. WorkflowEngine (workflow-engine.ts) **Purpose:** Executes DAG-based workflows with task dependencies.

Key Methods: - `createWorkflow(definition: WorkflowDefinition): Promise<string>` - `addTask(workflowId: string, task: Task): Promise<void>` - `startExecution(workflowId: string, params: object): Promise<string>` - `updateExecutionStatus(executionId: string, status: Status): Promise<void>`

Workflow Definition:

```
interface WorkflowDefinition {
  workflowId: string;
  name: string;
  description: string;
  category: 'generation' | 'analysis' | 'transformation' | 'pipeline' | 'custom';
  dagDefinition: {
    nodes: TaskNode[];
    edges: Edge[];
  }
}
```

```

    };
    inputSchema: JSONSchema;
    outputSchema: JSONSchema;
    defaultParameters: Record<string, any>;
    timeoutSeconds: number;
    maxRetries: number;
}

interface TaskNode {
  taskId: string;
  taskType: 'model_inference' | 'transformation' | 'condition' | 'parallel' | 'aggregation';
  config: object;
  dependsOn: string[];
  conditionExpression?: string;
}

```

11. ResponseSynthesisService (response-synthesis.service.ts) Purpose: Synthesizes responses from multiple AI models.

Synthesis Strategies:

Strategy	Description	Best For
best_of	Select highest confidence response	Quality-critical
vote	Majority voting on answer	Factual questions
weighted	Confidence \times (1/latency) weighted	Balanced
merge	AI combines all responses	Complex analysis

Merge Algorithm:

1. Collect all responses with metadata
 2. Extract key points from each
 3. Identify agreements and conflicts
 4. Generate unified response
 5. Apply conflict resolution
 6. Calculate final confidence
-

Billing Services

12. BillingService (billing.ts) Purpose: Manages credits, subscriptions, and billing.

Key Methods: - getSubscription(tenantId: string): Promise<Subscription> - getCreditBalance(tenantId: string): Promise<CreditBalance> - addCredits(tenantId: string, amount: number, type: string): Promise<number> - useCredits(tenantId: string, amount: number): Promise<{success, newBalance}> - purchaseCredits(tenantId: string, amount: number, price: number): Promise<string>

Subscription Tiers: | Tier | Monthly Price | Annual Price | Credits/User | |——|——|——|——|
——|——| | Free Trial | \$0 | - | 100 | | Individual | \$19 | \$190 | 1,000 | | Pro | \$49 | \$490 |
5,000 | | Team | \$199 | \$1,990 | 25,000 | | Enterprise | Custom | Custom | Custom |

Volume Discounts: | Credit Amount | Discount | Bonus Credits | |——|——|——|——|
—| | 10-19 | 5% | 0 | | 20-49 | 10% | 0 | | 50-99 | 15% | 5% | | 100+ | 25% | 10% |

Transaction Types: - purchase - Credit purchase - bonus - Promotional credits - refund -
Refunded credits - usage - Credits consumed - transfer_in / transfer_out - Credit transfers
- subscription_allocation - Monthly allocation - expiration - Expired credits - adjustment -
Manual adjustment

13. StorageBillingService (storage-billing.ts) **Purpose:** Tracks storage costs per tenant.

Billable Storage: - Uploaded files - Generated artifacts - Session history - Conversation logs

Pricing: - \$0.023 per GB/month (Standard) - \$0.0125 per GB/month (Infrequent) - \$0.004 per
GB/month (Archive)

Cognitive Services

14. CognitiveBrainService (cognitive-brain.service.ts) **Purpose:** High-level cognitive
processing and reasoning.

Cognitive Capabilities: - Working memory management - Attention allocation - Abstract reasoning -
Analogy formation - Concept learning

15. ReasoningEngine (reasoning-engine.ts) **Purpose:** Chain-of-thought and multi-step
reasoning.

Reasoning Modes: | Mode | Description | |——|——|——| | deductive | From general to specific
| | inductive | From specific to general | | abductive | Best explanation inference | | analogical
| | Similarity-based reasoning |

16. CausalReasoningService (causal-reasoning.service.ts) **Purpose:** Causal inference
and counterfactual reasoning.

Methods: - identifyCauses(effect: string): Promise<Cause[]> - predictEffects(cause:
string): Promise<Effect[]> - counterfactual(scenario: string, change: string):
Promise<string>

17. GoalPlanningService (goal-planning.service.ts) **Purpose:** Goal decomposition and planning.

Planning Algorithm:

1. Parse high-level goal
 2. Identify subgoals
 3. Determine dependencies
 4. Sequence actions
 5. Allocate resources
 6. Execute and monitor
-

18. MetacognitionService (metacognition.service.ts) **Purpose:** Self-reflection and learning from mistakes.

Metacognitive Functions: - Confidence calibration - Error detection - Strategy selection - Performance monitoring

Memory Services

19. MemoryService (memory-service.ts) **Purpose:** Persistent memory across sessions.

Memory Types: - **Short-term:** Current session context - **Long-term:** Cross-session knowledge
- **Episodic:** Event-based memories - **Semantic:** Factual knowledge

20. EpisodicMemoryService (episodic-memory.service.ts) **Purpose:** Event-based memory storage and retrieval.

Key Methods: - recordEpisode(event: Episode): Promise<void> - retrieveRelevant(query: string, limit: number): Promise<Episode[]> - consolidate(): Promise<void> - Memory optimization

21. MemoryConsolidationService (memory-consolidation.service.ts) **Purpose:** Optimizes memory storage by consolidating similar memories.

22. TimeMachineService (time-machine.ts) **Purpose:** Access historical state at any point in time.

Key Methods: - getStateAt(timestamp: Date): Promise<SystemState> - getDiff(from: Date, to: Date): Promise<StateDiff> - restore(timestamp: Date): Promise<void>

AGI Services

23. AGIOrchestratorService (agi-orchestrator.service.ts) **Purpose:** Coordinates AGI capabilities across services.

24. AdvancedAGIService (advancedagi.service.ts) **Purpose:** Advanced AGI features including self-improvement.

25. AGICcompleteService (agi-complete.service.ts) **Purpose:** Complete AGI pipeline from input to output.

26. AGIEextensionsService (agi-extensions.service.ts) **Purpose:** Extensible AGI capabilities.

Collaboration Services

27. CollaborationService (collaboration.ts) **Purpose:** Real-time collaboration features.

WebSocket Events: | Event | Direction | Description | |——|——|——| | join_session | Client→Server | Join collaborative session | | leave_session | Client→Server | Leave session | | cursor_move | Bidirectional | Cursor position update | | content_update | Bidirectional | Content change | | user_joined | Server→Client | New user notification | | user_left | Server→Client | User left notification |

28. ConcurrentSessionManager (concurrent-session.ts) **Purpose:** Manages concurrent user sessions.

Key Methods: - createSession(config: SessionConfig): Promise<string> - joinSession(sessionId: string, userId: string): Promise<void> - getSessionState(sessionId: string): Promise<SessionState> - broadcastUpdate(sessionId: string, update: Update): Promise<void>

29. TeamService (team-service.ts) **Purpose:** Team and organization management.

Key Methods: - createTeam(tenantId: string, name: string): Promise<string> - addMember(teamId: string, userId: string, role: string): Promise<void> - getTeamMembers(teamId: string): Promise<Member[]>

Additional Services (30-62)

#	Service	File	Purpose
30	NeuralEngine	<code>neural-engine.ts</code>	Neural network operations
31	AutoResolveService	<code>auto-resolve.ts</code>	Automatic conflict resolution
32	CanvasService	<code>canvas-service.ts</code>	Visual canvas artifacts
33	PersonaService	<code>persona-service.ts</code>	AI persona management
34	SchedulerService	<code>scheduler-service.ts</code>	Task scheduling
35	LicenseService	<code>license-service.ts</code>	License management
36	UnifiedModelRegistry	<code>unified-model-registry.ts</code>	Cryptic model registry
37	GrandfatheringService	<code>grandfathering-service.ts</code>	Legacy migration
38	VoiceVideoService	<code>voice-video.ts</code>	Voice/video processing
39	ResultMergingService	<code>result-merging.ts</code>	Merge results
40	WorldModelService	<code>world-model.service.ts</code>	World state modeling
41	MultiAgentService	<code>multi-agent.service.ts</code>	Multi-agent coordination
42	TheoryOfMindService	<code>theory-of-mind.service.ts</code>	Mental state modeling
43	MultimodalBindingService	<code>multimodal-binding.service.ts</code>	Cross-modal binding
44	SkillExecutionService	<code>skill-execution.service.ts</code>	Skill execution
45	AutonomousAgentService	<code>autonomous-agent.service.ts</code>	Autonomous operations
46	ConsciousnessService	<code>consciousness.service.ts</code>	Consciousness modeling
47	ConfigEngineService	<code>config-engine.service.ts</code>	Configuration engine
48	SelfImprovementService	<code>self-improvement.service.ts</code>	Self-improvement
49	MoralCompassService	<code>moral-compass.service.ts</code>	Ethical reasoning
50	MLTrainingService	<code>ml-training.service.ts</code>	Machine model training
51	LearningService	<code>learning.service.ts</code>	Learning data collection
52	FeedbackService	<code>feedback.service.ts</code>	User feedback
53	FeedbackLearningService	<code>feedback-learning.service.ts</code>	Learn from feedback
54	WorkflowProposalService	<code>workflow-proposals.service.ts</code>	Workflow improvements
55	AppIsolationService	<code>app-isolation.ts</code>	App-level isolation
56	LocalizationService	<code>localization.ts</code>	i18n support
57	MigrationApprovalService	<code>migration-approval.service.ts</code>	Migration approval
58	SuperiorOrchestrationService	<code>superior-orchestration.service.ts</code>	Superior responses
59	RadiantUnifiedService	<code>radiant-unified.service.ts</code>	Unified API
60	NeuralOrchestrationService	<code>neural-orchestration.service.ts</code>	Neural orchestration
61	AuditService	<code>audit.ts</code>	Audit logging
62	APIKeysService	<code>api-keys.ts</code>	API key management