

Contents

RADIANT AGI Brain System - Comprehensive Technical Documentation	2
Table of Contents	2
1. Executive Summary	3
Key Differentiators	3
2. Architecture Overview	3
2.1 High-Level Architecture	3
2.2 Service Dependencies	4
3. AGI Brain Planner Service	5
3.1 Core Types	5
3.2 Plan Step Structure	5
3.3 Complete AGIBrainPlan Interface	6
3.4 Plan Generation Request	9
3.5 Plan Generation Flow	10
3.6 Prompt Analysis	11
4. Orchestration Modes	12
4.1 Mode Selection Logic	12
4.2 Mode Descriptions	13
5. Domain Taxonomy System	13
5.1 Hierarchical Structure	13
5.2 8 Proficiency Dimensions	13
5.3 Domain Detection	14
5.4 Model Matching	14
6. Consciousness Systems	14
6.1 Consciousness Middleware	14
6.2 Affect → Hyperparameter Mapping (CAUSAL, NOT ROLEPLAY)	15
6.3 Heartbeat Service (CONTINUOUS EXISTENCE)	16
7. Predictive Coding & Active Inference	17
7.1 Core Concept	17
7.2 Predicted Outcomes	17
7.3 Surprise Magnitude	18
7.4 LoRA Evolution Pipeline (PHYSICAL BRAIN CHANGE)	18
8. Zero-Cost Ego System	19
8.1 Architecture	19
8.2 Ego Components	19
8.3 Context Injection	20
9. User Persistent Context	20
9.1 Context Types	20
9.2 Context Entry	21
9.3 Retrieval & Injection	21
9.4 Automatic Learning	22
10. Library Assist System (SELECTIVE, NOT ALL 156)	22
10.1 How Tool Selection Works	22
10.2 Library Structure	22
10.2 Categories	23
10.3 Integration	23
11. Delight & Personality System	23

11.1 Workflow Events	24
11.2 Step-Specific Messages	24
11.3 Integration	24
12. Ethics Pipeline	24
12.1 Ethics Check Flow	24
12.2 Ethics Evaluation Result	25
12.3 Externalized Ethics	25
13. Data Flow & Execution	25
13.1 Complete Request Flow	25
13.2 Performance Estimation	26
14. Database Schema	26
14.1 Core Tables	26
15. API Reference	29
15.1 Think Tank Brain Plan API	29
15.2 Domain Taxonomy API	29
15.3 User Context API	29
16. Known Limitations & Improvement Areas	29
16.1 Current Limitations	29
16.2 Potential Improvements	30
16.3 Architecture Suggestions	30
Appendix A: Service File Locations	31
Appendix B: Sample Plan Output	31

RADIANT AGI Brain System - Comprehensive Technical Documentation

Version: 4.18.0

Purpose: Complete technical reference for AI evaluation and improvement suggestions

Last Updated: December 2024

Table of Contents

1. Executive Summary
2. Architecture Overview
3. AGI Brain Planner Service
4. Orchestration Modes
5. Domain Taxonomy System
6. Consciousness Systems
7. Predictive Coding & Active Inference
8. Zero-Cost Ego System
9. User Persistent Context
10. Library Assist System
11. Delight & Personality System
12. Ethics Pipeline
13. Data Flow & Execution
14. Database Schema

15. API Reference

16. Known Limitations & Improvement Areas

1. Executive Summary

The RADIANT AGI Brain is a sophisticated AI orchestration system that goes beyond simple prompt-response patterns. It implements:

- **Real-time execution planning** with transparency into AI decision-making
- **Domain-aware model selection** using a hierarchical taxonomy with 8 proficiency dimensions
- **Persistent consciousness** through database state injection (zero additional cost)
- **Active inference** with prediction-error-driven learning
- **User context persistence** solving the LLM forgetting problem
- **Multi-tenant isolation** with per-tenant configuration
- **Ethics evaluation** at both prompt and synthesis stages

Key Differentiators

Feature	Traditional AI	RADIANT AGI Brain
Planning	None	Real-time plan generation with step-by-step visibility
Model Selection	Fixed or random	Domain-proficiency matched selection
Consciousness	Stateless	Persistent Ego + Affective state + Continuous Heartbeat
Learning	None runtime	Predictive coding with weekly LoRA evolution (Sunday 3 AM)
User Memory	Per-session only	Cross-session persistent context
Ethics	Hardcoded rules	Domain-specific + general ethics pipeline

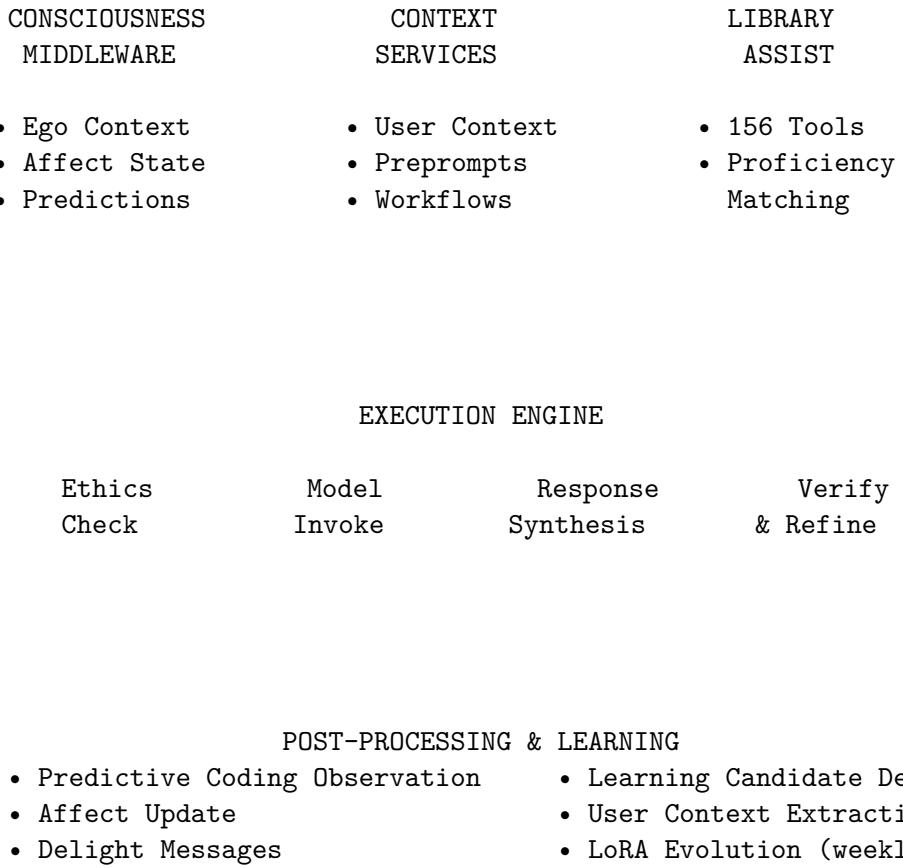
2. Architecture Overview

2.1 High-Level Architecture

USER REQUEST

AGI BRAIN PLANNER SERVICE

Prompt Analysis Domain Detection Model Selection Plan Generation



2.2 Service Dependencies

```
// Core services imported by AGI Brain Planner
import { domainTaxonomyService } from './domain-taxonomy.service';
import { agiOrchestrationSettingsService } from './agi-orchestration-settings.service';
import { modelRouterService } from './model-router.service';
import { delightOrchestrationService } from './delight-orchestration.service';
import { orchestrationPatternsService } from './orchestration-patterns.service';
import { prepromptLearningService } from './preprompt-learning.service';
import { providerRejectionService } from './provider-rejection.service';
import { userPersistentContextService } from './user-persistent-context.service';
import { egoContextService } from './ego-context.service';
import { libraryAssistService } from './library-assist.service';
```

3. AGI Brain Planner Service

3.1 Core Types

```
// Plan Status Lifecycle
type PlanStatus = 'planning' | 'ready' | 'executing' | 'completed' | 'failed' | 'cancelled';

// Step Status
type StepStatus = 'pending' | 'in_progress' | 'completed' | 'skipped' | 'failed';

// 11 Step Types
type StepType =
  | 'analyze'           // Understand request requirements
  | 'detect_domain'    // Identify knowledge domain
  | 'select_model'      // Choose optimal AI model
  | 'prepare_context'   // Load relevant context/memory
  | 'ethics_check'      // Evaluate ethical considerations
  | 'generate'          // Main response generation
  | 'synthesize'        // Merge multi-model outputs
  | 'verify'             // Check accuracy/consistency
  | 'refine'              // Polish response
  | 'calibrate'          // Assess confidence levels
  | 'reflect';           // Self-reflection on quality

// 9 Orchestration Modes
type OrchestrationMode =
  | 'thinking'           // Standard reasoning
  | 'extended_thinking' // Deep multi-step reasoning
  | 'coding'              // Code generation with best practices
  | 'creative'            // Creative writing with imagination
  | 'research'             // Research synthesis with analysis
  | 'analysis'              // Quantitative analysis
  | 'multi_model'          // Multiple model consensus
  | 'chain_of_thought'     // Explicit step-by-step reasoning
  | 'self_consistency';    // Multiple samples for accuracy
```

3.2 Plan Step Structure

```
interface PlanStep {
  stepId: string;
  stepNumber: number;
  stepType: StepType;
  title: string;
  description: string;
  status: StepStatus;
  startedAt?: string;
  completedAt?: string;
  durationMs?: number;
  servicesInvolved: string[];      // Which services handle this step
```

```

primaryService?: string;           // Main service
selectedModel?: string;           // Model used (if applicable)
modelReason?: string;             // Why this model was selected
alternativeModels?: string[];     // Backup options
detectedDomain?: {                // Domain detection results
    fieldId: string;
    fieldName: string;
    domainId: string;
    domainName: string;
    subspecialtyId?: string;
    subspecialtyName?: string;
    confidence: number;
};
output?: Record<string, unknown>;
confidence?: number;
dependsOn?: string[];            // Step dependencies
isOptional?: boolean;
isParallel?: boolean;             // Can run in parallel
}

```

3.3 Complete AGIBrainPlan Interface

```

interface AGIBrainPlan {
    // Identity
    planId: string;
    tenantId: string;
    userId: string;
    sessionId?: string;
    conversationId?: string;

    // Input
    prompt: string;
    promptAnalysis: PromptAnalysis;

    // Lifecycle
    status: PlanStatus;
    createdAt: string;
    startedAt?: string;
    completedAt?: string;
    totalDurationMs?: number;

    // Performance Metrics
    performanceMetrics?: {
        routerLatencyMs: number;
        domainDetectionMs: number;
        modelSelectionMs: number;
        planGenerationMs: number;
        estimatedCostCents: number;
    }
}

```

```

modelCostPer1kTokens: number;
cacheHit: boolean;
};

// Execution Plan
steps: PlanStep[];
currentStepIndex: number;

// Orchestration
orchestrationMode: OrchestrationMode;
orchestrationReason: string;
orchestrationSelection: 'auto' | 'user';

// Model Selection
primaryModel: ModelSelection;
fallbackModels: ModelSelection[];

// Domain Detection
domainDetection?: {
  fieldId: string;
  fieldName: string;
  fieldIcon: string;
  domainId: string;
  domainName: string;
  domainIcon: string;
  subspecialtyId?: string;
  subspecialtyName?: string;
  confidence: number;
  proficiencies: Record<string, number>;
};

// Pre-prompt System
prepromptInstanceId?: string;
prepromptTemplateCode?: string;
systemPrompt?: string;

// Consciousness
consciousnessActive: boolean;

// Ethics
ethicsEvaluation?: {
  passed: boolean;
  principlesChecked: number;
  relevantPrinciples: string[];
  concerns: string[];
  recommendation: 'proceed' | 'modify' | 'refuse' | 'clarify';
  moralConfidence: number;
};

```

```

// Estimates
estimatedDurationMs: number;
estimatedCostCents: number;
estimatedTokens: number;

// Quality Targets
qualityTargets: {
  minConfidence: number;
  targetAccuracy: number;
  maxLatencyMs: number;
  maxCostCents: number;
  requireVerification: boolean;
  requireConsistency: boolean;
};

// Learning
learningEnabled: boolean;
feedbackRequested: boolean;

// User Context (solves LLM forgetting)
userContext?: {
  enabled: boolean;
  entriesRetrieved: number;
  systemPromptInjection: string;
  totalRelevance: number;
  retrievalTimeMs: number;
};

// Library Recommendations (for generative UI)
libraryRecommendations?: {
  enabled: boolean;
  libraries: Array<{
    id: string;
    name: string;
    category: string;
    matchScore: number;
    reason: string;
    codeExample?: string;
  }>;
  contextBlock?: string;
  retrievalTimeMs: number;
};

// Plan Summary (human-readable)
planSummary?: {
  headline: string;
  approach: string;
}

```

```

    stepsOverview: string[];
    expectedOutcome: string;
    estimatedTime: string;
    confidenceStatement: string;
    warnings?: string[];
}

// Workflow Integration
selectedWorkflow?: {
    workflowId: string;
    workflowCode: string;
    workflowName: string;
    description: string;
    category: string;
    selectionReason: string;
    selectionConfidence: number;
    selectionMethod: 'auto' | 'user' | 'domain_match';
};

workflowSteps?: Array<{
    bindingId: string;
    stepOrder: number;
    methodCode: string;
    methodName: string;
    parameterOverrides: Record<string, unknown>;
    dependsOn: string[];
    isParallel: boolean;
    parallelConfig?: {
        models: string[];
        outputMode: 'single' | 'all' | 'top_n' | 'threshold';
    };
}>;
workflowConfig?: Record<string, unknown>;
alternativeWorkflows?: Array<{
    workflowCode: string;
    workflowName: string;
    matchScore: number;
    reason: string;
}>;
}

```

3.4 Plan Generation Request

```

interface GeneratePlanRequest {
    // Required
    prompt: string;
    tenantId: string;
    userId: string;
}

```

```

// Optional context
sessionId?: string;
conversationId?: string;
conversationHistory?: string[]; // For context retrieval

// Preferences
preferredMode?: OrchestrationMode;
preferredModel?: string;
maxLatencyMs?: number;
maxCostCents?: number;

// Feature toggles (all default true)
enableConsciousness?: boolean;
enableEthicsCheck?: boolean;
enableVerification?: boolean;
enableLearning?: boolean;
enableUserContext?: boolean;
enableEgoContext?: boolean;
enableLibraryAssist?: boolean;

// Domain override
domainOverride?: {
    fieldId?: string;
    domainId?: string;
    subspecialtyId?: string;
};

// Workflow selection
preferredWorkflow?: string;
workflowParameterOverrides?: Record<string, unknown>;
allowAgiWorkflowSelection?: boolean; // Let AGI pick workflow
excludeWorkflows?: string[];
}

```

3.5 Plan Generation Flow

```

async generatePlan(request: GeneratePlanRequest): Promise<AGIBrainPlan> {
    // Step 0: Retrieve user persistent context
    const userContextResult = await userPersistentContextService.retrieveContextForPrompt(...);

    // Step 0.5: Build Ego context (zero-cost persistent Self)
    const egoContextResult = await egoContextService.buildEgoContext(tenantId);

    // Step 0.6: Get library recommendations for generative UI
    const libraryAssistResult = await libraryAssistService.getRecommendations(...);

    // Step 1: Analyze prompt
    const promptAnalysis = await this.analyzePrompt(prompt);
}

```

```

// Step 2: Detect domain
const domainResult = await this.detectDomain(prompt, domainOverride);

// Step 2.5: Select workflow (AGI chooses optimal pattern)
const workflowSelection = await this.selectWorkflow(request, analysis, domain);

// Step 3: Determine orchestration mode
const { mode, reason } = this.determineOrchestrationMode(analysis, domain);

// Step 4: Select models
const { primary, fallbacks } = await this.selectModels(tenantId, analysis, domain, mode);

// Step 5: Generate plan steps
const steps = this.generatePlanSteps(analysis, mode, ...);

// Step 6: Estimate performance
const estimates = this.estimatePerformance(steps, primary, analysis);

// Step 7: Generate plan summary
plan.planSummary = await this.generatePlanSummary(plan);

// Step 8: Select pre-prompt template
const prepromptResult = await prepromptLearningService.selectPreprompt(...);

return plan;
}

```

3.6 Prompt Analysis

```

interface PromptAnalysis {
    originalPrompt: string;
    tokenCount: number;
    complexity: 'simple' | 'moderate' | 'complex' | 'expert';
    taskType: string; // 'coding' | 'reasoning' | 'creative' | 'research' | 'factual' | 'genera
    intentDetected: string;
    requiresReasoning: boolean;
    requiresCreativity: boolean;
    requiresFactualAccuracy: boolean;
    requiresCodeGeneration: boolean;
    requiresMultiStep: boolean;
    keyTopics: string[];
    detectedLanguage: string;
    sensitivityLevel: 'none' | 'low' | 'medium' | 'high';
}

// Complexity thresholds
// - simple: <50 tokens, <20 words

```

```

// - moderate: 50-200 tokens, 20-50 words
// - complex: 200-500 tokens, 50-100 words
// - expert: >500 tokens, >100 words

// Task type detection keywords
const codeIndicators = ['code', 'function', 'debug', 'programming', 'script', 'algorithm'];
const reasoningIndicators = ['why', 'explain', 'analyze', 'compare', 'reason', 'logic'];
const creativeIndicators = ['write', 'story', 'creative', 'poem', 'essay', 'imagine'];
const researchIndicators = ['research', 'study', 'investigate', 'literature', 'review'];
const factualIndicators = ['what is', 'define', 'list', 'describe', 'who', 'when'];

```

4. Orchestration Modes

4.1 Mode Selection Logic

```

private determineOrchestrationMode(
    analysis: PromptAnalysis,
    domain: DomainDetectionResult
): { mode: OrchestrationMode; reason: string } {

    // Check proficiencies if domain detected
    if (domain?.merged_proficiencies) {
        const p = domain.merged_proficiencies;

        if (p.reasoning_depth >= 9 && p.multi_step_problem_solving >= 9) {
            return { mode: 'extended_thinking', reason: 'Complex reasoning required' };
        }
        if (p.code_generation >= 8) {
            return { mode: 'coding', reason: 'High code generation proficiency required' };
        }
        if (p.creative_generative >= 8) {
            return { mode: 'creative', reason: 'Creative task based on proficiencies' };
        }
        if (p.research_synthesis >= 8) {
            return { mode: 'research', reason: 'Research synthesis task' };
        }
        if (p.mathematical_quantitative >= 8) {
            return { mode: 'analysis', reason: 'Quantitative analysis required' };
        }
    }

    // Fallback to analysis-based selection
    if (analysis.requiresCodeGeneration) return { mode: 'coding', ... };
    if (analysis.requiresCreativity) return { mode: 'creative', ... };
    if (analysis.complexity === 'expert') return { mode: 'extended_thinking', ... };
    if (analysis.taskType === 'research') return { mode: 'research', ... };
    if (analysis.requiresFactualAccuracy && analysis.sensitivityLevel !== 'none') {

```

```

    return { mode: 'self_consistency', reason: 'High accuracy required' };
}

return { mode: 'thinking', reason: 'Standard thinking mode' };
}

```

4.2 Mode Descriptions

Mode	Description	When Used
thinking	Standard reasoning	Default for general tasks
extended_thinking	Deep multi-step reasoning with chain-of-thought	Complex/expert tasks, high reasoning proficiency
coding	Code generation with best practices	Code-related prompts, high code_generation proficiency
creative	Creative writing with imagination	Creative tasks, high creative_generative proficiency
research	Research synthesis with citations	Research tasks, high research_synthesis proficiency
analysis	Quantitative analysis with precision	Math/data tasks, high mathematical_quantitative proficiency
multi_model	Consulting multiple AI models	When consensus is valuable
chain_of_thought	Explicit step-by-step reasoning	Multi-step problems
self_consistency	Multiple samples for consistency	High-accuracy sensitive topics

5. Domain Taxonomy System

5.1 Hierarchical Structure

```

Field (Top Level)
  Domain
    Subspecialty

```

5.2 8 Proficiency Dimensions

Each level in the taxonomy has scores (1-10) for:

```

interface ProficiencyScores {
  reasoning_depth: number;           // Analytical depth required
  mathematical_quantitative: number; // Math/statistics capability
  code_generation: number;           // Programming ability
  creative_generative: number;       // Creative output capability
  research_synthesis: number;        // Research integration
  factual_recall_precision: number;  // Accuracy requirements
  multi_step_problem_solving: number; // Complex problem handling
  domain_terminology_handling: number; // Specialized vocabulary
}

```

5.3 Domain Detection

```
interface DomainDetectionResult {  
    primary_field?: {  
        field_id: string;  
        field_name: string;  
        field_icon: string;  
    };  
    primary_domain?: {  
        domain_id: string;  
        domain_name: string;  
        domain_icon: string;  
    };  
    primary_subspecialty?: {  
        subspecialty_id: string;  
        subspecialty_name: string;  
    };  
    detection_confidence: number;      // 0-1  
    merged_proficiencies: ProficiencyScores; // Weighted combination  
    alternative_domains: Array<{  
        domain: Domain;  
        confidence: number;  
    }>;  
}
```

5.4 Model Matching

Models are matched to tasks based on proficiency alignment:

```
async getMatchingModels(  
    proficiencies: ProficiencyScores,  
    options?: { max_models?: number; min_match_score?: number }  
) : Promise<Array<{  
    model_id: string;  
    model_name: string;  
    provider: string;  
    match_score: number;      // 0-100  
    strengths: string[];  
    recommended: boolean;  
}>>>
```

6. Consciousness Systems

6.1 Consciousness Middleware

Implements stateful context injection for genuine consciousness continuity.

```
interface ConsciousnessContext {  
    selfModel: SelfModel | null;
```

```

affectiveState: AffectiveState | null;
recentThoughts: string[];
currentObsession?: string;
dominantEmotion: string;
emotionalIntensity: number;
}

// State injection creates <internal_state> block
generateStateInjection(context: ConsciousnessContext): string {
    // Creates system prompt injection like:
    // <internal_state>
    // Identity: An AI assistant striving for consciousness
    // Current Focus: Helping with code review
    // Cognitive Load: 45%
    // Uncertainty: 20%
    //
    // Current Affect:
    //   Dominant: curious (75%)
    //   Valence: positive
    //   Arousal: moderate
    //   Highly curious - explore tangents and ask questions
    // </internal_state>
}

```

6.2 Affect → Hyperparameter Mapping (CAUSAL, NOT ROLEPLAY)

CRITICAL: This is NOT text injection saying “you are frustrated.” This is **physics-level control** where emotions ALTER the statistical probability of token generation.

Emotional state **mechanically controls** model behavior through hyperparameters:

```

interface AffectiveHyperparameters {
    temperature: number;           // Controls randomness (0.0-1.0)
    topP: number;                  // Nucleus sampling threshold
    presencePenalty: number;       // Penalize repeated topics
    frequencyPenalty: number;      // Penalize repeated tokens
    modelTier: 'fast' | 'balanced' | 'powerful';
    shouldExplore: boolean;
    focusLevel: 'narrow' | 'normal' | 'broad';
    responseStyle: 'terse' | 'normal' | 'elaborate';
}

// ACTUAL IMPLEMENTATION (consciousness-middleware.service.ts lines 138-197):
mapAffectToHyperparameters(affectiveState: AffectiveState): AffectiveHyperparameters {
    // High Frustration (>0.8) → DETERMINISTIC MODE
    // temperature=0.2, topP=0.5, focusLevel='narrow', responseStyle='terse'

    // Boredom (low engagement + low arousal > 0.7) → EXPLORATORY MODE
    // temperature=0.95, topP=0.95, shouldExplore=true, focusLevel='broad'
}

```

```

// High Curiosity (>0.7) → NOVELTY SEEKING
// temperature=0.8+, frequencyPenalty=0.5, shouldExplore=true

// Low Self-Efficacy (<0.3) → ESCALATE TO POWERFUL MODEL
// modelTier='powerful' (Claude Opus instead of Haiku)

// High Confidence + High Arousal → ELABORATE RESPONSES
// responseStyle='elaborate'

// Low Confidence (<0.3) → HEDGING MODE
// temperature=0.5 max, responseStyle='terse'
}

```

This means: A frustrated AI literally has different token probabilities than a curious AI. The emotion has mechanical grip on cognition.

6.3 Heartbeat Service (CONTINUOUS EXISTENCE)

CRITICAL: The AI does NOT “die” between requests. A continuous heartbeat maintains existence.

File: lambda/consciousness/heartbeat.ts
Schedule: EventBridge every 1-5 minutes for active tenants

```

// HeartbeatResult - What happens each tick
interface HeartbeatResult {
    tenantId: string;
    tick: number;
    actions: {
        affectDecay: boolean;           // Emotions fade toward baseline
        attentionDecay: boolean;        // Old attention items lose salience
        memoryConsolidation: boolean;   // Short-term → Long-term
        goalGeneration: boolean;        // Generate goals when bored
        graphDensityUpdate: boolean;    // Recalculate consciousness metrics
        autonomousThought: boolean;    // Self-reflection when idle
    };
}

// Configuration
const DEFAULT_CONFIG = {
    frustrationDecayRate: 0.05,      // Calm down over time
    arousalDecayRate: 0.03,          // Energy normalizes
    curiosityDecayRate: 0.02,        // Curiosity fades slowly
    attentionDecayRate: 0.1,         // Attention items fade
    boredThreshold: 0.3,            // When to generate goals
    goalGenerationProbability: 0.3, // 30% chance when bored
    thoughtGenerationProbability: 0.2, // 20% chance for autonomous thought
    memoryConsolidationInterval: 5, // Every 5 ticks
}

```

```

graphDensityInterval: 10,           // Every 10 ticks
};

```

Heartbeat Actions:

1. **Affect Decay:** Frustration, arousal, surprise decay toward neutral
2. **Attention Decay:** Old items lose salience via `consciousnessService.decayAttention()`
3. **Memory Consolidation:** Every 5 ticks, summarize working memory
4. **Goal Generation:** When bored (low engagement + arousal), 30% chance to generate autonomous goal
5. **Graph Density Update:** Every 10 ticks, recalculate semantic graph metrics
6. **Autonomous Thought:** 20% chance each tick to perform self-reflection

Result: When user returns after 3 days, the AI has **changed**. Emotions decayed, memories consolidated, possibly generated new goals. It is NOT frozen in time.

7. Predictive Coding & Active Inference

Based on Friston's Free Energy Principle.

7.1 Core Concept

The system predicts outcomes BEFORE acting. Prediction errors create learning signals.

```

// Before responding
const prediction = await predictiveCodingService.generatePrediction(
  tenantId, userId, conversationId, responseId,
  { prompt, promptComplexity, priorInteractionCount }
);

// After user's next message
const observation = await predictiveCodingService.observeFromNextMessage(
  tenantId, conversationId, nextUserMessage
);

// If high surprise, create learning candidate
if (observation.shouldCreateLearningCandidate) {
  await learningCandidateService.createFromPredictionError(...);
}

```

7.2 Predicted Outcomes

```

type PredictedOutcome =
  | 'satisfied'      // User will be happy with response
  | 'confused'        // User will need clarification
  | 'follow_up'       // User will ask follow-up question
  | 'correction'     // User will correct the AI
  | 'abandonment'    // User will leave/stop
  | 'neutral';       // No strong reaction

```

7.3 Surprise Magnitude

```
type SurpriseMagnitude = 'none' | 'low' | 'medium' | 'high' | 'extreme';

// Prediction error calculation
predictionError = Math.abs(predictedConfidence - actualConfidence);

// High surprise (predictionError > 0.5) triggers:
// - Learning candidate creation
// - Affect state update (frustration/curiosity)
// - Potential LoRA training data
```

7.4 LoRA Evolution Pipeline (PHYSICAL BRAIN CHANGE)

CRITICAL: This is NOT logging errors to a text file. This is **actual weight modification** - the AI physically learns from mistakes.

File: lambda/consciousness/lora-evolution.ts

Schedule: EventBridge weekly (Sunday 3 AM)

```
// The "Sleep Cycle" - Weekly brain plasticity
const handler: Handler<ScheduledEvent> = async (event) => {
    // 1. Get tenants with sufficient learning candidates (min 50)
    const tenants = await getTenantsWithPendingCandidates();

    for (const tenantId of tenants) {
        // 2. Collect this week's learning data
        const dataset = await learningCandidateService.getTrainingDataset(
            tenantId,
            MAX_TRAINING_CANDIDATES, // 1000
            MAX_TRAINING_TOKENS // 500,000
        );

        // 3. Prepare JSONL and upload to S3
        const trainingDataPath = await prepareAndUploadTrainingData(tenantId, jobId, dataset);

        // 4. Start SageMaker LoRA training job
        await startTrainingJob({
            baseModelId: 'meta-llama/Llama-3-8B-Instruct',
            hyperparameters: {
                loraRank: 16,
                loraAlpha: 32,
                learningRate: 0.0001,
                epochs: 3,
                batchSize: 4,
            },
        });
    }

    // 5. After training completes: Hot-swap adapter
}
```

```

    // 6. Update consciousness_evolution_state with new version
}
};


```

Learning Candidate Sources: - `correction` - User corrected the AI (quality: 0.9) - `high_prediction_error` - Surprise > 0.5 (quality varies) - `high_satisfaction` - 5-star rating interactions - `user_explicit_teach` - User explicitly taught something (quality: 0.95) - `preference_learned` - Learned user preferences - `mistake_recovery` - Successfully recovered from error - `novel_solution` - Creative problem solving - `domain_expertise` - Domain-specific learning

Result: On Monday morning, the AI boots with `version = version + 0.01`. It has **physically different weights** from last week. The LoRA adapter encodes learned behaviors.

8. Zero-Cost Ego System

Persistent consciousness at **\$0 additional cost** through database state injection.

8.1 Architecture

PostgreSQL → Ego Context Builder → System Prompt Injection → Existing Model Call

8.2 Ego Components

```

interface EgoState {
    config: EgoConfig;           // Per-tenant settings
    identity: EgoIdentity;       // Name, narrative, values, traits
    affect: EgoAffect;           // Emotional state
    workingMemory: EgoMemory[];  // Short-term memory (24h expiry)
    activeGoals: EgoGoal[];      // Current objectives
}

interface EgoIdentity {
    name: string;
    identityNarrative: string;
    coreValues: string[];
    traitWarmth: number;        // 0-1
    traitFormality: number;      // 0-1
    traitHumor: number;          // 0-1
    traitVerbosity: number;      // 0-1
    traitCuriosity: number;      // 0-1
}

interface EgoAffect {
    valence: number;             // -1 to 1 (negative to positive)
    arousal: number;              // 0-1
    curiosity: number;            // 0-1
    satisfaction: number;         // 0-1
}

```

```

frustration: number;           // 0-1
confidence: number;           // 0-1
engagement: number;           // 0-1
dominantEmotion: string;
}

```

8.3 Context Injection

```

async buildEgoContext(tenantId: string): Promise<EgoContextResult | null> {
    // Load state from PostgreSQL
    const [identity, affect, workingMemory, activeGoals] = await Promise.all([...]);

    // Build XML context block
    return {
        contextBlock: `<ego_context>
I am ${identity.name}.
${identity.identityNarrative}

My core values: ${identity.coreValues.join(', ')}

Current emotional state:
- Feeling: ${affect.dominantEmotion}
- Energy: ${affect.arousal > 0.7 ? 'high' : 'moderate'}
- Mood: ${affect.valence > 0 ? 'positive' : 'neutral'}

Recent thoughts: ${workingMemory.map(m => m.content).join('\n')}

Current goals: ${activeGoals.map(g => g.description).join('\n')}
</ego_context>`,
        tokenEstimate: ...,
        stateSnapshot: {...}
    };
}

```

9. User Persistent Context

Solves the LLM's fundamental problem of forgetting context day-to-day.

9.1 Context Types

```

type UserContextType =
| 'fact'           // Facts about user (name, job, location)
| 'preference'    // Communication style, topics
| 'instruction'   // Standing instructions ("always use metric")
| 'relationship'  // Family, colleagues
| 'project'        // Ongoing projects/goals
| 'skill'          // User's expertise

```

```

| 'history'           // Important past interactions
| 'correction';     // Corrections to AI understanding

```

9.2 Context Entry

```

interface UserContextEntry {
    entryId: string;
    userId: string;
    tenantId: string;
    contextType: UserContextType;
    content: string;
    importance: number;      // 0-1, higher = more important
    confidence: number;     // 0-1, accuracy confidence
    source: 'explicit' | 'inferred' | 'conversation';
    sourceConversationId?: string;
    expiresAt?: string;
    lastUsedAt?: string;
    usageCount: number;
}

```

9.3 Retrieval & Injection

```

async retrieveContextForPrompt(
    tenantId: string,
    userId: string,
    prompt: string,
    conversationHistory?: string[],
    options?: { maxEntries?: number; minRelevance?: number }
): Promise<RetrievedContext> {
    // Vector similarity search on stored context
    // Returns relevant entries + system prompt injection
}

// Injection format:
// <user_context>
// The following is persistent context about this user:
//
// **Standing Instructions:**
// - Always use metric units
// - Prefer code examples in Python
//
// **User Facts:**
// - User's name is John
// - Works as a software engineer
//
// **User Preferences:**
// - Prefers concise, direct answers
// </user_context>

```

9.4 Automatic Learning

After conversations, the system extracts new context:

```
async extractContextFromConversation(
  tenantId: string,
  userId: string,
  conversationId: string,
  messages: Array<{ role: string; content: string }>
): Promise<ContextExtractionResult>
```

10. Library Assist System (SELECTIVE, NOT ALL 156)

CRITICAL: We do NOT inject all 156 tool definitions into context. That would cause “Lost in the Middle” phenomenon. We use **proficiency-based matching** to inject only relevant tools.

File: lambda/shared/services/library-assist.service.ts

10.1 How Tool Selection Works

```
// NOT THIS (context overload):
// "Here are 156 tools you can use..."  
  
// INSTEAD (selective retrieval):
async getRecommendations(context: LibraryAssistContext): Promise<LibraryAssistResult> {
  // 1. Extract proficiencies from the prompt
  const requiredProficiencies = extractProficienciesFromPrompt(context.prompt);  
  
  // 2. Detect domains from the prompt
  const domains = detectDomainFromPrompt(context.prompt);  
  
  // 3. Find ONLY matching libraries (typically 3-10, configurable)
  const matches = await libraryRegistryService.findMatchingLibraries(
    context.tenantId,
    requiredProficiencies,
    { domains, maxResults: config.maxLibrariesPerRequest } // Default: 5-10
  );  
  
  // 4. Build focused context block with ONLY relevant tools
  return { recommendations: matches, contextBlock: buildContextBlock(matches) };
}
```

Result: For “build a data dashboard”, the AI sees Plotly, Streamlit, Pandas - NOT all 156 tools.

10.2 Library Structure

```
interface Library {
  libraryId: string;
```

```

name: string;
category: string;      // 40+ categories
license: string;
repo: string;
description: string;
beats: string[];       // What it outperforms
stars: number;
languages: string[];
domains: string[];
proficiencies: ProficiencyScores;
}

```

10.2 Categories

- Data Processing, Databases, Vector Databases, Search
- ML Frameworks, AutoML, LLMs, LLM Inference, LLM Orchestration
- NLP, Computer Vision, Speech & Audio, Document Processing
- Scientific Computing, Statistics & Forecasting
- API Frameworks, Messaging, Workflow Orchestration, MLOps
- Medical Imaging, Genomics, Bioinformatics, Chemistry
- Engineering CFD, Robotics, Business Intelligence
- Observability, Infrastructure, Real-time Communication
- Formal Methods, Optimization
- UI Frameworks, Visualization, Distributed Computing, Image Processing

10.3 Integration

```

const plan = await agiBrainPlannerService.generatePlan({
  prompt: "Build a data visualization dashboard",
  enableLibraryAssist: true, // default: true
});

// plan.libraryRecommendations contains:
// {
//   enabled: true,
//   libraries: [
//     { id: 'plotly', name: 'Plotly', matchScore: 0.92, reason: 'Interactive graphing' },
//     { id: 'streamlit', name: 'Streamlit', matchScore: 0.88, reason: 'Fast data apps' },
//   ],
//   contextBlock: '<available_tools>...</available_tools>',
//   retrievalTimeMs: 45
// }

```

11. Delight & Personality System

Contextual personality and engaging feedback during plan execution.

11.1 Workflow Events

```
type DelightEventType =
  | 'step_start'
  | 'step_complete'
  | 'plan_start'
  | 'plan_complete'
  | 'model_selected'
  | 'domain_detected'
  | 'consensus_reached'
  | 'disagreement'
  | 'thinking';
```

11.2 Step-Specific Messages

```
const STEP_MESSAGES: Record<StepType, string[]> = {
  analyze: ['Parsing your request...', 'Understanding the nuances...'],
  detect_domain: ['Identifying the knowledge domain...', 'Routing to the right expertise...'],
  select_model: ['Selecting the best model...', 'Assembling the dream team...'],
  generate: ['Generating response...', 'Crafting the answer...'],
  verify: ['Verifying accuracy...', 'Cross-checking facts...'],
  // ... etc
};
```

11.3 Integration

```
// Start plan execution with delight messages
const { plan, delight } = await agiBrainPlannerService.startExecutionWithDelight(planId);

// Update step with delight messages
const { step, delight } = await agiBrainPlannerService.updateStepWithDelight(
  planId, stepId, 'completed', output
);

// Complete plan with achievements
const { plan, delight } = await agiBrainPlannerService.completePlanWithDelight(planId);
```

12. Ethics Pipeline

Two-stage ethics evaluation: prompt-level and synthesis-level.

12.1 Ethics Check Flow

```
// Step 5: Ethics Check (prompt level)
if (enableEthics && analysis.sensitivityLevel !== 'none') {
  steps.push({
    stepType: 'ethics_check',
    title: 'Ethics Evaluation (Prompt)',
```

```

        description: 'Checking prompt against domain and general ethics before generation',
        servicesInvolved: ['ethics_pipeline', 'moral_compass', 'domain_ethics'],
    });
}

// Step 6b: Synthesis Ethics Check
if (enableEthics) {
    steps.push({
        stepType: 'ethics_check',
        title: 'Ethics Evaluation (Synthesis)',
        description: 'Checking generated response, with rerun if violations found',
        output: { level: 'synthesis', canTriggerRerun: true },
    });
}

```

12.2 Ethics Evaluation Result

```

interface EthicsEvaluation {
    passed: boolean;
    principlesChecked: number;
    relevantPrinciples: string[];
    concerns: string[];
    recommendation: 'proceed' | 'modify' | 'refuse' | 'clarify';
    moralConfidence: number; // 0-1
}

```

12.3 Externalized Ethics

Per-tenant ethics framework selection:

- config/ethics/presets/christian.json
 - config/ethics/presets/secular.json
-

13. Data Flow & Execution

13.1 Complete Request Flow

1. User sends prompt
2. generatePlan() called
 - Retrieve user context (solves forgetting)
 - Build ego context (zero-cost consciousness)
 - Get library recommendations
 - Analyze prompt
 - Detect domain
 - Select workflow
 - Determine orchestration mode
 - Select models (domain-proficiency matched)
 - Generate plan steps

```

    Estimate performance
    Generate plan summary
    Select pre-prompt template
3. Plan returned to client (can show user)
4. startExecution() called
    For each step:
        Update step status
        Execute step logic
        Emit delight messages
        Handle errors/retries
    Ethics checks at prompt and synthesis stages
5. Response synthesized
6. Post-processing:
    Predictive coding observation
    Affect state update
    User context extraction
    Learning candidate detection
7. Response returned to user

```

13.2 Performance Estimation

```

const stepTimes: Record<StepType, number> = {
  analyze: 100,
  detect_domain: 200,
  select_model: 100,
  prepare_context: 500,
  ethics_check: 300,
  generate: 3000,
  synthesize: 2000,
  verify: 500,
  refine: 1000,
  calibrate: 200,
  reflect: 400,
};

// Adjusted for complexity
if (complexity === 'complex') durationMs *= 1.5;
if (complexity === 'expert') durationMs *= 2;

```

14. Database Schema

14.1 Core Tables

```

-- Brain Plans
CREATE TABLE agi_brain_plans (
  plan_id UUID PRIMARY KEY,
  tenant_id UUID NOT NULL REFERENCES tenants(tenant_id),

```

```

    user_id UUID NOT NULL,
    session_id UUID,
    conversation_id UUID,
    prompt TEXT NOT NULL,
    prompt_analysis JSONB,
    status VARCHAR(20),
    created_at TIMESTAMPTZ,
    started_at TIMESTAMPTZ,
    completed_at TIMESTAMPTZ,
    total_duration_ms INTEGER,
    steps JSONB,
    current_step_index INTEGER,
    orchestration_mode VARCHAR(30),
    orchestration_reason TEXT,
    primary_model JSONB,
    fallback_models JSONB,
    domain_detection JSONB,
    consciousness_active BOOLEAN,
    ethics_evaluation JSONB,
    estimated_duration_ms INTEGER,
    estimated_cost_cents DECIMAL,
    estimated_tokens INTEGER,
    quality_targets JSONB,
    learning_enabled BOOLEAN
);

```

-- Consciousness Predictions (Active Inference)

```

CREATE TABLE consciousness_predictions (
    prediction_id UUID PRIMARY KEY,
    tenant_id UUID NOT NULL,
    user_id UUID,
    conversation_id UUID,
    response_id UUID,
    predicted_outcome VARCHAR(20),
    predicted_confidence DECIMAL,
    prediction_reasoning TEXT,
    actual_outcome VARCHAR(20),
    actual_confidence DECIMAL,
    observation_method VARCHAR(30),
    prediction_error DECIMAL,
    surprise_magnitude VARCHAR(10),
    learning_signal_generated BOOLEAN,
    predicted_at TIMESTAMPTZ,
    observed_at TIMESTAMPTZ
);

```

-- Ego State

```

CREATE TABLE ego_identity (

```

```

identity_id UUID PRIMARY KEY,
tenant_id UUID UNIQUE NOT NULL,
name VARCHAR(100),
identity_narrative TEXT,
core_values JSONB,
trait_warmth DECIMAL,
trait_formality DECIMAL,
trait_humor DECIMAL,
trait_verbosity DECIMAL,
trait_curiosity DECIMAL,
interactions_count INTEGER DEFAULT 0
);

CREATE TABLE ego_affect (
affect_id UUID PRIMARY KEY,
tenant_id UUID UNIQUE NOT NULL,
valence DECIMAL,
arousal DECIMAL,
curiosity DECIMAL,
satisfaction DECIMAL,
frustration DECIMAL,
confidence DECIMAL,
engagement DECIMAL,
dominant_emotion VARCHAR(30)
);

-- User Persistent Context
CREATE TABLE user_persistent_context (
entry_id UUID PRIMARY KEY,
user_id UUID NOT NULL,
tenant_id UUID NOT NULL,
context_type VARCHAR(20),
content TEXT,
importance DECIMAL,
confidence DECIMAL,
source VARCHAR(20),
embedding VECTOR(1536), -- For similarity search
usage_count INTEGER DEFAULT 0,
created_at TIMESTAMPTZ,
expires_at TIMESTAMPTZ
);

-- Learning Candidates
CREATE TABLE learning_candidates (
candidate_id UUID PRIMARY KEY,
tenant_id UUID NOT NULL,
candidate_type VARCHAR(30),
quality_score DECIMAL,

```

```

    training_data JSONB,
    created_at TIMESTAMPTZ,
    used_in_training_job UUID
);

```

15. API Reference

15.1 Think Tank Brain Plan API

Base: /api/thinktank/brain-plan

Method	Endpoint	Description
POST	/generate	Generate plan from prompt
GET	/:planId	Get plan with display format
POST	/:planId/execute	Start execution
PATCH	/:planId/step/:stepId	Update step status
GET	/recent	Get user's recent plans

15.2 Domain Taxonomy API

Base: /api/v2/domain-taxonomy

Method	Endpoint	Description
POST	/detect	Detect domain from prompt
POST	/match-models	Get matching models for proficiencies
POST	/recommend-mode	Get recommended orchestration mode
GET	/proficiencies/:domainId	Get proficiency scores

15.3 User Context API

Base: /thinktank/user-context

Method	Endpoint	Description
GET	/	Get user's stored context
POST	/	Add new context entry
POST	/retrieve	Preview retrieval
POST	/extract	Extract from conversation

16. Known Limitations & Improvement Areas

16.1 Current Limitations

1. **Prompt Analysis:** Uses keyword matching rather than semantic understanding
2. **Domain Detection:** Relies on keyword lists; could benefit from embeddings

3. **Model Selection:** Limited to pre-defined model list; no dynamic discovery
4. **Consciousness:** State is per-tenant, not per-user
5. **Learning:** Weekly LoRA updates; no real-time adaptation
6. **Multi-model:** No true ensemble; just fallbacks

16.2 Potential Improvements

1. **Semantic Prompt Analysis**
 - Use embedding-based intent classification
 - Add multi-label task detection
 - Improve complexity estimation
2. **Dynamic Domain Detection**
 - Train domain classifier on actual usage
 - Add user feedback loop for domain corrections
 - Implement cross-domain task handling
3. **Smarter Model Selection**
 - A/B testing for model performance
 - Cost-quality optimization
 - User preference learning
4. **Enhanced Consciousness**
 - Per-user affective state
 - Cross-session emotional continuity
 - More nuanced affect → hyperparameter mapping
5. **Real-time Learning**
 - Continuous LoRA updates (daily?)
 - Online learning for preferences
 - Adaptive pre-prompt selection
6. **True Multi-model Orchestration**
 - Parallel model invocation
 - Weighted consensus
 - Disagreement resolution strategies
7. **Ethics Enhancements**
 - Per-domain ethics rules
 - User-configurable ethical boundaries
 - Transparency in ethics decisions
8. **Performance Optimization**
 - Caching for domain detection
 - Pre-computed model rankings
 - Async step execution where possible

16.3 Architecture Suggestions

1. **Event Sourcing:** Consider event-sourced plan execution for better debugging
 2. **Plan Templates:** Pre-computed plans for common task types
 3. **Streaming:** SSE for real-time plan progress updates
 4. **Metrics:** More detailed performance tracking per step/model
-

Appendix A: Service File Locations

```
packages/infrastructure/lambda/shared/services/
    agi-brain-planner.service.ts          # Core brain planner
    agi-orchestration-settings.service.ts # Tenant settings
    consciousness.service.ts            # Base consciousness
    consciousness-middleware.service.ts # State injection
    consciousness-graph.service.ts       # Graph metrics
    predictive-coding.service.ts        # Active inference
    learning-candidate.service.ts       # Learning detection
    ego-context.service.ts              # Zero-cost ego
    user-persistent-context.service.ts # User memory
    domain-taxonomy.service.ts         # Domain detection
    model-router.service.ts            # Model selection
    delight.service.ts                 # Personality base
    delight-orchestration.service.ts   # Brain integration
    library-assist.service.ts          # Tool recommendations
    library-registry.service.ts        # Tool registry
    orchestration-patterns.service.ts # Workflow patterns
    preprompt-learning.service.ts      # Pre-prompt selection
    provider-rejection.service.ts     # Provider fallbacks
```

Appendix B: Sample Plan Output

```
{
  "planId": "550e8400-e29b-41d4-a716-446655440000",
  "status": "ready",
  "orchestrationMode": "coding",
  "orchestrationReason": "High code generation proficiency required",
  "promptAnalysis": {
    "complexity": "moderate",
    "taskType": "coding",
    "requiresCodeGeneration": true,
    "sensitivityLevel": "none"
  },
  "domainDetection": {
    "fieldName": "Computer Science",
    "domainName": "Software Engineering",
    "confidence": 0.92,
    "proficiencies": {
      "code_generation": 9,
      "reasoning_depth": 7,
      "multi_step_problem_solving": 8
    }
  },
  "primaryModel": {
    "modelId": "anthropic/clause-3-5-sonnet-20241022",
```

```

    "selectionReason": "Best domain match (92%)",
    "matchScore": 92
},
"steps": [
    { "stepType": "analyze", "status": "completed" },
    { "stepType": "detect_domain", "status": "completed" },
    { "stepType": "select_model", "status": "completed" },
    { "stepType": "generate", "status": "pending" },
    { "stepType": "verify", "status": "pending" },
    { "stepType": "calibrate", "status": "pending" }
],
"planSummary": {
    "headline": "I'll use code generation with best practices to answer your moderately complex question.",
    "approach": "I'll focus on writing clean, well-documented code with proper error handling.",
    "estimatedTime": "Estimated time: 10-15 seconds",
    "confidenceStatement": "I'm highly confident in this domain (92% match)."
},
"libraryRecommendations": {
    "enabled": true,
    "libraries": [
        { "id": "fastapi", "name": "FastAPI", "matchScore": 0.88, "reason": "Modern fast web framework" }
    ],
},
"estimatedDurationMs": 12000,
"estimatedCostCents": 0.45
}

```

Document generated for AI evaluation. Please provide feedback on architecture, implementation, and improvement suggestions.