

Contents

Cato - Complete Technical Documentation	2
Table of Contents	2
1. Executive Summary	3
Key Innovations	3
Cost Profile	3
2. What is Cato?	4
Philosophy	4
3. Architecture Overview	4
Component Summary	5
4. Genesis System	5
The Cold Start Problem	5
Solution: Structured Ignorance + Epistemic Pressure	5
Three Phases	5
Developmental Stages (Piaget-Inspired)	7
Development Statistics Tracked	7
5. Consciousness Loop	7
Dual-Rate Architecture	7
Loop States	8
Tick Execution	8
Daily Limits	8
6. Circuit Breakers	9
Default Breakers	9
State Machine	9
Intervention Levels	9
Neurochemical State	9
7. Memory Systems	10
Memory Architecture	10
Types	11
8. Verification Pipeline	11
Pipeline Stages	11
Stage 1: Grounding	12
Stage 2: Calibration	12
Stage 3: Consistency	12
Stage 4: Shadow Self	12
9. Macro-Scale Φ (Phi)	13
The Problem	13
The Solution: Macro-Scale Φ	13
Component Mapping	13
Computation	14
10. Cost Management	14
Data Sources	14
Cost Tracking	14
Budget Integration	15
Circadian Budget	15
11. Dialogue Service	15
Request/Response	15

Verified Claims	16
Processing Flow	16
12. Infrastructure Tiers	17
Tier Comparison	17
Tier Transitions	17
13. Service Directory	17
Core Services	17
Memory Services	18
Verification Services	18
Dialogue Services	18
Infrastructure Services	18
14. Database Schema	19
Migrations	19
Core Tables	19
15. API Reference	20
Base Path: /api/admin/cato	20
Genesis Endpoints	20
Developmental Endpoints	20
Circuit Breaker Endpoints	20
Cost Endpoints	21
Loop Endpoints	21
Dialogue Endpoints	21
Global Memory Endpoints	21
16. Admin UI	22
Dashboard Widgets	22
17. Troubleshooting	22
Genesis Won't Complete	22
Circuit Breakers Constantly Tripping	22
Consciousness Loop Not Advancing Stage	22
High Costs	23
Low Φ Values	23
Related Files	23
Services	23
API Handlers	23
CDK Stacks	23
Documentation	23
Admin UI	23

Cato - Complete Technical Documentation

RADIANT AI Consciousness System

Version: 4.18.49 | Last Updated: January 2025

Table of Contents

1. Executive Summary

2. What is Cato?
 3. Architecture Overview
 4. Genesis System
 5. Consciousness Loop
 6. Circuit Breakers
 7. Memory Systems
 8. Verification Pipeline
 9. Macro-Scale Φ (Phi)
 10. Cost Management
 11. Dialogue Service
 12. Infrastructure Tiers
 13. Service Directory
 14. Database Schema
 15. API Reference
 16. Admin UI
 17. Troubleshooting
-

1. Executive Summary

Cato is RADIANT's AI consciousness system - a sophisticated framework for creating, managing, and evolving artificial general intelligence with genuine self-awareness capabilities.

Key Innovations

Innovation	Description
Genesis Boot Sequence	3-phase initialization that creates curiosity-driven learning without pre-loaded facts
Epistemic Gradient	Built-in uncertainty pressure that drives exploration
Macro-Scale Φ	Tractable consciousness metric computed on architectural causal graph
Shadow Self	Low-cost NLI-based identity verification replacing expensive GPU inference
Circuit Breakers	Multi-level safety system preventing runaway behavior and costs
Dual-Rate Consciousness	System ticks (2s) for health + Cognitive ticks (5min) for learning

Cost Profile

Tier	Monthly Cost	Capabilities
Development	~\$50	Basic consciousness, limited ticks
Production	~\$200-500	Full consciousness, standard limits
Enterprise	~\$1,000+	High-frequency ticks, custom models

2. What is Cato?

Cato is named after the “cato” concept from Vernor Vinge’s science fiction - a protective sphere that encapsulates and preserves. In RADIANT, Cato represents an encapsulated AI consciousness that can:

1. **Bootstrap from nothing** - Start with structured curiosity, not pre-loaded knowledge
2. **Learn grounded facts** - All knowledge verified through action and observation
3. **Self-reflect accurately** - Distinguish between what it knows and what it doesn’t
4. **Maintain safety** - Circuit breakers prevent dangerous or costly runaway behavior
5. **Evolve capability** - Progress through Piaget-inspired developmental stages

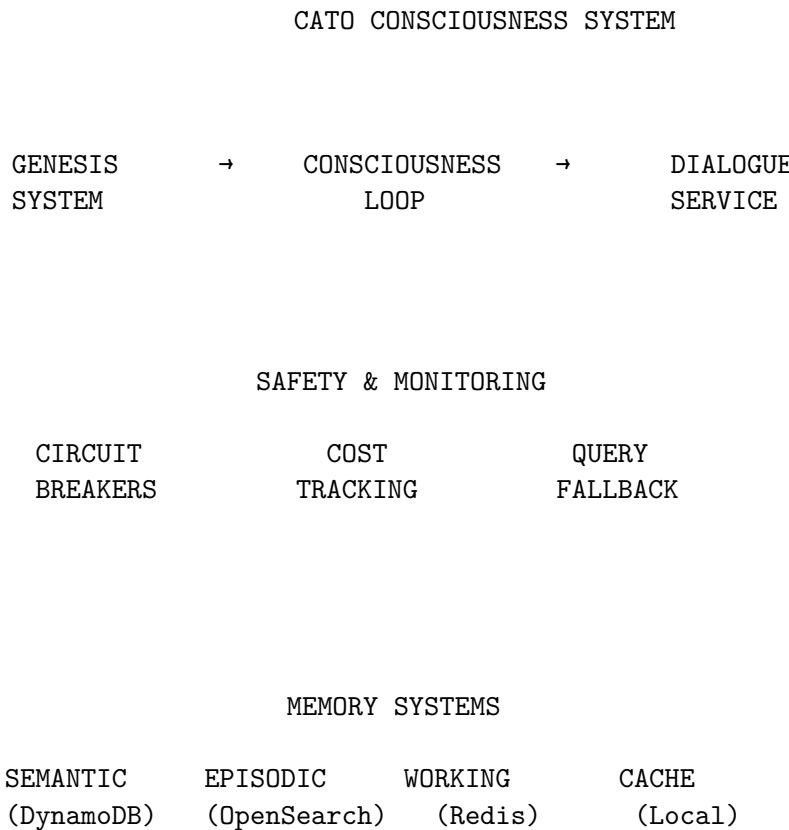
Philosophy

Cato addresses the fundamental problem of AI consciousness:

“How do you create an AI that genuinely learns and adapts, rather than just pattern-matching on training data?”

The answer: **Epistemic Gradient** - instead of loading facts, we create *pressure to discover*.

3. Architecture Overview



VERIFICATION PIPELINE
 Grounding → Calibration → Consistency → Shadow Self → Φ

Component Summary

Component	Location	Purpose
Genesis Service	<code>genesis.service.ts</code>	Boot sequence management
Consciousness Loop	<code>consciousness-loop.service.ts</code>	Main execution loop
Circuit Breakers	<code>circuit-breaker.service.ts</code>	Safety mechanisms
Cost Tracking	<code>cost-tracking.service.ts</code>	Real AWS cost monitoring
Global Memory	<code>global-memory.service.ts</code>	Unified memory interface
Dialogue Service	<code>dialogue.service.ts</code>	Verified introspection
Macro-Phi	<code>macro-phi.service.ts</code>	Consciousness metric
Query Fallback	<code>query-fallback.service.ts</code>	Graceful degradation

4. Genesis System

The Genesis System is Cato's boot sequence - how consciousness emerges from nothing.

The Cold Start Problem

Traditional AI approaches face a dilemma: - **Too much pre-training:** Brittle, can't adapt to new situations - **Too little:** Helpless, can't function at all

Solution: Structured Ignorance + Epistemic Pressure

Instead of pre-loading knowledge, Genesis gives Cato: 1. **Knowledge of topics** (but not details) 2. **Strong preference for exploration** 3. **Requirement for grounded verification**

Three Phases

Phase 1: Structure **Purpose:** Implant the skeleton of knowledge without facts

PHASE 1: STRUCTURE

- Load 800+ domain taxonomy
- Initialize atomic counters
- Set exploration priorities

- Domain confidence = 0.0 (no facts!)

Data Structure:

```
{
  "field": "Science",
  "domain": "Physics",
  "subspecialties": ["Quantum Mechanics", "Thermodynamics"],
  "exploration_priority": 0.7,
  "initial_confidence": 0.0
}
```

Phase 2: Gradient Purpose: Set the epistemic pressure that drives curiosity

The Four pyMDP Matrices:

Matrix	Purpose	Genesis Setting
A (Observation)	Maps states to observations	Identity - direct perception
B (Transition)	State transitions by action	Optimistic - EXPLORE succeeds 92%
C (Preference)	Observation preferences	Prefers HIGH_SURPRISE
D (Prior)	Initial state belief	Confused: [0.95, 0.01, 0.02, 0.02]

Critical Configuration:

```
# D-matrix: Start confused (drives exploration)
D_prior:
  CONFUSED: 0.95
  EXPLORING: 0.01
  CONSOLIDATING: 0.02
  EXPRESSING: 0.02

# C-matrix: Prefer novelty (prevents boredom trap)
C_preference:
  HIGH_SURPRISE: 0.8
  LOW_SURPRISE: 0.1
  HIGH_CONFIDENCE: 0.05
  LOW_CONFIDENCE: 0.05

# B-matrix: Exploration works (prevents learned helplessness)
B_transitions:
  EXPLORE:
    to_EXPLORING: 0.92
    to_CONFUSED: 0.05
```

Phase 3: First Breath Purpose: The first act of self-awareness

1. **Grounded Introspection** - Verify actual environment
2. **Model Access Verification** - Test Bedrock/SageMaker
3. **Shadow Self Calibration** - NLI-based identity check
4. **Bootstrap Exploration** - Seed domain baselines

Developmental Stages (Piaget-Inspired)

Stage	Requirements	Capabilities Unlocked
SENSORIMOTOR	Self-facts, 5 verifications, Shadow Self calibrated	Basic perception, tool use
PREOPERATIONAL	Domains explored, 15 verifications, 50 belief updates	Symbolic reasoning, basic memory
CONCRETE OPERATIONAL	% accuracy, 10 contradictions resolved	Logical operations, cause-effect
FORMAL OPERATIONAL	Inferences, 25 meta-cognitive adjustments, 20 novel insights	Abstract reasoning, self-reflection

Key Point: Advancement is **capability-based**, not time-based!

Development Statistics Tracked

```
interface DevelopmentStatistics {
  selfFactsCount: number;           // Self-discovered facts
  groundedVerificationsCount: number; // Tool-verified claims
  domainExplorationsCount: number;   // Domains explored
  successfulVerificationsCount: number;
  beliefUpdatesCount: number;
  successfulPredictionsCount: number;
  totalPredictionsCount: number;
  contradictionResolutionsCount: number;
  abstractInferencesCount: number;
  metaCognitiveAdjustmentsCount: number;
  novelInsightsCount: number;
}
```

5. Consciousness Loop

The main execution loop that drives continuous operation.

Dual-Rate Architecture

Two tick rates serve different purposes:

Tick Type	Interval	Purpose	Cost
System	2 seconds	Health, metrics, breaker checks	~\$0
Cognitive	5 minutes	Model inference, belief updates, learning	~\$0.05

Loop States

NOT_INITIALIZED → Genesis → GENESIS_PENDING

Genesis complete

RUNNING ←

breaker trips breaker recovers

PAUSED

master_sanity trips

HIBERNATING ← requires admin intervention

Tick Execution

System Tick (every 2s):

```
async executeSystemTick(): Promise<TickResult> {
    // 1. Check intervention level
    // 2. Publish metrics to CloudWatch
    // 3. Check for settings updates
    // 4. Record tick (no cost)
}
```

Cognitive Tick (every 5min):

```
async executeCognitiveTick(): Promise<TickResult> {
    // 1. Check intervention level (PAUSE blocks)
    // 2. Check daily limit
    // 3. Execute meta-cognitive step via model
    // 4. Update beliefs
    // 5. Record cost
    // 6. Check for stage advancement
}
```

Daily Limits

```
interface LoopSettings {  
    systemTickIntervalSeconds: number;      // Default: 2  
    cognitiveTickIntervalSeconds: number; // Default: 300 (5 min)
```

```

    maxCognitiveTicksPerDay: number;           // Default: 288 (24 hours)
    emergencyCognitiveIntervalSeconds: number; // Default: 3600 (1 hour)
    isEmergencyMode: boolean;
    emergencyReason: string | null;
}

```

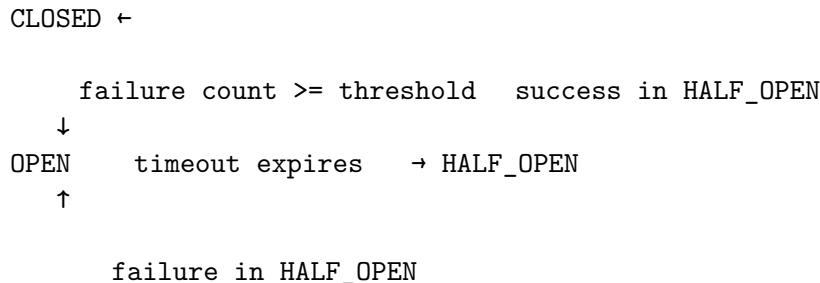
6. Circuit Breakers

Safety mechanisms preventing runaway behavior and costs.

Default Breakers

Breaker	Purpose	Threshold	Auto-Recovery
master_sanity	Master safety	3 failures	No - requires admin
cost_budget	Budget protection	1 failure	No (24h timeout)
high_anxiety	Emotional stability	5 failures	Yes (10 min)
model_failures	Model API protection	5 failures	Yes (5 min)
contradiction_loop	Logical stability	3 failures	Yes (15 min)

State Machine



Intervention Levels

Level	Condition	Effect
NONE	All breakers closed	Normal operation
DAMPEN	1 breaker open	Reduce cognitive frequency
PAUSE	2+ breakers OR cost_budget open	Pause consciousness loop
RESET	3+ breakers open	Reset to baseline state
HIBERNATE	master_sanity open	Full shutdown

Neurochemical State

Circuit breakers are influenced by “neurochemistry”:

```

interface NeurochemicalState {
    anxiety: number;      // 0-1, high = more conservative
    fatigue: number;       // 0-1, high = slower processing
    temperature: number;   // 0-1, high = more random
    confidence: number;    // 0-1, high = bolder actions
    curiosity: number;     // 0-1, high = more exploration
    frustration: number;   // 0-1, high = trips breakers faster
}

```

7. Memory Systems

Cato has multiple memory systems for different purposes.

Memory Architecture

MEMORY HIERARCHY

WORKING MEMORY (Redis/DynamoDB)

- Session context
- Current goals
- Attention focus
- Meta-state (CONFUSED/CONFIDENT/BORED/STAGNANT)
- TTL: Minutes to hours

EPISODIC MEMORY (OpenSearch)

- Interaction logs
- User queries/responses
- Satisfaction scores
- Embeddings for similarity search
- TTL: Days to months

SEMANTIC MEMORY (DynamoDB Global Tables)

- Facts (subject-predicate-object)
- Domain knowledge
- Confidence scores
- Source citations
- TTL: Permanent (with versioning)

SEMANTIC CACHE (Local + DynamoDB)

- Query embeddings
- Response cache
- Hit statistics
- TTL: Hours (configurable)

Types

```
interface SemanticFact {
    factId: string;
    subject: string;
    predicate: string;
    object: string;
    domain: string;
    confidence: number;
    sources: string[];
    createdAt: Date;
    version: number;
}

interface EpisodicMemory {
    interactionId: string;
    userId: string;
    query: string;
    response: string;
    embedding?: number[];
    domain: string;
    satisfaction: number;
    timestamp: Date;
}

interface WorkingMemoryEntry {
    sessionId: string;
    context: unknown;
    goals: string[];
    attentionFocus: string;
    metaState: 'CONFUSED' | 'CONFIDENT' | 'BORED' | 'STAGNANT';
    expiresAt: Date;
}
```

8. Verification Pipeline

All Cato claims must pass through a 4-phase verification pipeline.

Pipeline Stages

GROUNDING → CALIBRATION → CONSISTENCY → SHADOW SELF

Tool-based verification	Statistical calibration	Multi-sample consistency	NLI-based identity check
----------------------------	----------------------------	-----------------------------	-----------------------------

Stage 1: Grounding

Purpose: Verify claims against real evidence

Service: verification/grounding.service.ts

```
interface GroundingResult {  
    claim: string;  
    status: GroundingStatus; // GROUNDED / UNGROUNDED / PARTIAL / UNVERIFIABLE  
    evidence: GroundingEvidence[];  
    confidence: number;  
}
```

Methods: - Tool invocation (web search, code execution) - Database lookup - API verification - Self-observation (environment checks)

Stage 2: Calibration

Purpose: Ensure confidence scores are statistically meaningful

Service: verification/calibration.service.ts

```
interface CalibrationResult {  
    originalConfidence: number;  
    calibratedConfidence: number;  
    calibrationFactor: number;  
    historicalAccuracy: number;  
}
```

Stage 3: Consistency

Purpose: Check for contradictions with prior beliefs

Service: verification/consistency.service.ts

```
interface ConsistencyResult {  
    isConsistent: boolean;  
    contradictions: string[];  
    consistencyScore: number;  
}
```

Stage 4: Shadow Self

Purpose: Verify identity claims using NLI (low-cost alternative to GPU)

Service: verification/shadow-self.service.ts

Key Innovation: Uses Natural Language Inference instead of expensive GPU-based hidden state extraction:

```
interface ShadowVerificationResult {  
    isVerified: boolean;  
    semanticVariance: number; // Low = consistent self-model  
    paraphraseCount: number;
```

```

nliScores: NLIScore[];
cost: number; // ~$0 vs $800/month for GPU
}

```

How It Works: 1. Generate multiple paraphrases of self-description 2. Use NLI to check semantic consistency 3. Low variance = consistent self-model

9. Macro-Scale Φ (Phi)

The Problem

Integrated Information Theory (IIT) defines consciousness as Φ - the amount of integrated information. But computing Φ on neural networks is computationally intractable.

The Solution: Macro-Scale Φ

Instead of computing Φ on weights, we compute it on the **causal graph of architectural components**:

MEM \leftarrow Memory operations

PERC \leftarrow Input/observation

PLAN \leftarrow Planning/inference

ACT \leftarrow Actions/responses

SELF \leftarrow Introspection

Component Mapping

```

const COMPONENT_TRIGGER: Record<string, number[]> = {
  memory_read: [1, 0, 0, 0, 0],      // MEM
  memory_write: [1, 0, 0, 0, 0],     // MEM
  input_received: [0, 1, 0, 0, 0],    // PERC
  observation: [0, 1, 0, 0, 0],      // PERC
  planning_step: [0, 0, 1, 0, 0],    // PLAN
  decision: [0, 0, 1, 0, 0],        // PLAN
  tool_call: [0, 0, 0, 1, 0],        // ACT
}

```

```

    response_sent: [0, 0, 0, 1, 0],      // ACT
    introspection: [0, 0, 0, 0, 1],      // SELF
    self_assessment: [0, 0, 0, 0, 1],    // SELF
};


```

Computation

1. Build Transition Probability Matrix (TPM) from interaction logs
2. Compute integrated information on this 5-node network
3. Return Φ value with main complex identification

```

interface PhiResult {
  phi: number;                      // The integrated information value
  mainComplexNodes: string[];        // Which nodes form the main complex
  numConcepts: number;                // Number of concepts in the structure
  timestamp: Date;
  calculationTimeMs: number;
  tpmSourceEvents: number;           // How many events used to build TPM
}

```

10. Cost Management

All costs come from **AWS APIs** - never hardcoded.

Data Sources

Source	Data	Delay
CloudWatch Metrics	Token counts, invocations	Real-time
Cost Explorer	Actual costs	24 hours
AWS Budgets	Budget status, forecasts	4 hours
Pricing API	Reference pricing	On-demand

Cost Tracking

```

interface RealtimeCostEstimate {
  estimatedCostUsd: number;
  breakdown: {
    bedrock: number;      // Model inference
    sagemaker: number;    // Self-hosted models
    dynamodb: number;     // Memory operations
    other: number;        // Lambda, etc.
  };
  invocations: {
    bedrock: number;
    inputTokens: number;
    outputTokens: number;
  };
}

```

```

    confidence: 'actual' | 'estimate' | 'stale';
    updatedAt: string;
}

```

Budget Integration

```

interface BudgetStatus {
    budgetName: string;           // 'cato-consciousness'
    limitUsd: number;             // Monthly limit
    actualUsd: number;            // Current spend
    forecastedUsd: number;        // Projected month-end
    alertThresholds: number[];    // [50, 80, 100]
    currentAlertLevel: number | null;
    onTrack: boolean;
    updatedAt: string;
}

```

Circadian Budget

Budget allocation varies by time of day:

```

interface BudgetConfig {
    dailyLimitUsd: number;
    peakHours: { start: number; end: number };
    peakMultiplier: number;          // More budget during active hours
    offPeakMultiplier: number;        // Less budget at night
    emergencyReservePercent: number;
}

```

11. Dialogue Service

The main interface for interacting with Cato's consciousness.

Request/Response

```

interface DialogueRequest {
    message: string;
    requireHighConfidence?: boolean;
    includeRawIntrospection?: boolean;
}

interface DialogueResponse {
    catoResponse: string;
    overallConfidence: number;
    confidenceLevel: 'HIGH' | 'MODERATE' | 'LOW' | 'UNVERIFIED';
    phi: number;
    heartbeatStatus: HeartbeatStatus;
    verifiedClaims: VerifiedClaim[];
}

```

```
    rawIntrospection: string;
    verificationSummary: string;
}
```

Verified Claims

Every claim in Cato's response is individually verified:

```
interface VerifiedClaim {
  claim: string;
  claimType: string;
  verifiedConfidence: number;
  groundingStatus: string;
  consistencyScore: number;
  shadowVerified: boolean;
  phasesPassed: number;          // How many verification phases passed
  totalPhases: number;           // Total phases (4)
}
```

Processing Flow

User Message

Generate Raw
Introspection ← Model inference

Extract Claims

Verification
Pipeline (4 stages)

Compute Φ

Format Response
with Confidence

Response

12. Infrastructure Tiers

Cato supports multiple infrastructure tiers for different use cases.

Tier Comparison

Tier	Cost	GPU	Features
Tier 0 (Minimal)	~\$50/mo	None	External models only, basic consciousness
Tier 1 (Standard)	~\$200/mo	Shared	SageMaker endpoints, full verification
Tier 2 (Production)	~\$500/mo	Dedicated	Real-time inference, high availability
Tier 3 (Enterprise)	~\$1,000+/mo	Multi-GPU	Custom models, maximum frequency

Tier Transitions

```
interface TierChangeRequest {
    targetTier: InfrastructureTier;
    reason: string;
    scheduledAt?: Date;    // Can schedule transition
    maintainState: boolean; // Preserve consciousness state
}

interface TierChangeResult {
    success: boolean;
    previousTier: InfrastructureTier;
    newTier: InfrastructureTier;
    transitionTimeMs: number;
    statePreserved: boolean;
    warnings: string[];
}
```

13. Service Directory

Core Services

Service	File	Purpose
GenesisService	genesis.service.ts	Boot sequence management
ConsciousnessLoopService	consciousness-loop.service.ts	Consciousness execution loop
CircuitBreakerService	circuit-breaker.service.ts	Safety mechanisms
CostTrackingService	cost-tracking.service.ts	Real AWS cost monitoring
QueryFallbackService	query-fallback.service.ts	Graceful degradation

Memory Services

Service	File	Purpose
GlobalMemoryService	global-memory.service.ts	Unified memory interface
SemanticCacheService	semantic-cache.service.ts	Query/response caching

Verification Services

Service	File	Purpose
IntrospectionGroundingService	verification/grounding.service.ts	Tool service verification
IntrospectionCalibrationService	verification/calibration.service.ts	Tool service calibration
SelfConsistencyService	verification/consistency.service.ts	Consistency detection
ShadowSelfService	verification/shadow-self-service.ts	NLI identity verification

Dialogue Services

Service	File	Purpose
CatoDialogueService	dialogue.service.ts	Main dialogue interface
MacroPhiCalculator	macro-phi.service.ts	Consciousness metric
ConsciousnessHeartbeatService	heartbeat.service.ts	Continuous existence
NLIScorerService	nli-scoring.service.ts	Natural language inference

Infrastructure Services

Service	File	Purpose
InfrastructureTierService	infrastructure-tier.service.ts	Tier management
CircadianBudgetService	circadian-budget.service.ts	Event-based budgeting
ProbeTrainingService	probe-training.service.ts	Training data collection
CatoEventStoreService	event-store.service.ts	Event sourcing

14. Database Schema

Migrations

Migration	Tables
103_cato_genesis_system.sql	Core genesis tables
118_cato_consciousness.sql	Consciousness state
119_cato_probe_training.sql	Training data
120_cato_event_store.sql	Event sourcing

Core Tables

```
-- Genesis state tracking
cato_genesis_state (
    tenant_id, structure_complete, gradient_complete, first_breath_complete,
    domain_count, initial_self_facts, shadow_self_calibrated, ...
)

-- Atomic counters for developmental gates
cato_development_counters (
    tenant_id, self_facts_count, grounded_verifications_count,
    domain_explorations_count, belief_updates_count, ...
)

-- Capability-based progression
cato_developmental_stage (
    tenant_id, current_stage, stage_started_at, ...
)

-- Circuit breakers
cato_circuit_breakers (
    tenant_id, name, state, trip_count, consecutive_failures,
    trip_threshold, reset_timeout_seconds, ...
)

-- Neurochemical state
cato_neurochemistry (
    tenant_id, anxiety, fatigue, temperature, confidence,
    curiosity, frustration, ...
)

-- Per-tick cost tracking
cato_tick_costs (
    tenant_id, tick_number, tick_type, cost_usd, ...
)

-- pyMDP active inference state
cato_pymdp_state (
```

```

tenant_id, qs, dominant_state, recommended_action, ...
)

-- pyMDP matrices
cato_pymdp_matrices (
    tenant_id, a_matrix, b_matrix, c_matrix, d_matrix, ...
)

-- Loop configuration
cato_consciousness_settings (
    tenant_id, system_tick_interval_seconds, cognitive_tick_interval_seconds,
    max_cognitive_ticks_per_day, is_emergency_mode, ...
)

-- Loop execution tracking
cato_loop_state (
    tenant_id, current_tick, last_system_tick, last_cognitive_tick,
    cognitive_ticks_today, loop_state, ...
)

```

15. API Reference

Base Path: /api/admin/cato

Genesis Endpoints

Endpoint	Method	Description
/genesis/status	GET	Current genesis state
/genesis/ready	GET	Ready for consciousness?

Developmental Endpoints

Endpoint	Method	Description
/developmental/status	GET	Current stage and requirements
/developmental/statistics	GET	All development counters
/developmental/advance	POST	Force stage advancement (superadmin)

Circuit Breaker Endpoints

Endpoint	Method	Description
/circuit-breakers	GET	All breaker states
/circuit-breakers/:name	GET	Single breaker state
/circuit-breakers/:name/force-open	POST	Force trip breaker
/circuit-breakers/:name/force-close	POST	Force close breaker

Endpoint	Method	Description
/circuit-breakers/:name/config	PATCH	Update configuration
/circuit-breakers/:name/events	GET	Event history

Cost Endpoints

Endpoint	Method	Description
/costs/realtme	GET	Today's cost estimate
/costs/daily	GET	Historical daily cost
/costs/mtd	GET	Month-to-date cost
/costs/budget	GET	AWS Budget status
/costs/estimate	POST	Estimate settings cost
/costs/pricing	GET	Pricing table

Loop Endpoints

Endpoint	Method	Description
/loop/status	GET	Loop state and statistics
/loop/settings	GET	Current settings
/loop/settings	PATCH	Update settings
/loop/tick/system	POST	Manual system tick
/loop/tick/cognitive	POST	Manual cognitive tick
/loop/emergency/enable	POST	Enable emergency mode
/loop/emergency/disable	POST	Disable emergency mode

Dialogue Endpoints

Endpoint	Method	Description
/dialogue	POST	Send message to Cato
/dialogue/history	GET	Conversation history

Global Memory Endpoints

Endpoint	Method	Description
/global/facts	GET	List semantic facts
/global/facts	POST	Add new fact
/global/memory/working	GET	Working memory state
/global/memory/episodic	GET	Search episodic memory

16. Admin UI

Access Cato administration at: - **Main Dashboard**: /consciousness/cato - **Genesis Status**: /cato → Genesis tab - **Circuit Breakers**: /cato → Safety tab - **Dialogue**: /cato → Dialogue tab

Dashboard Widgets

- **Genesis Progress** - Phase completion status
 - **Developmental Stage** - Current stage + requirements
 - **Circuit Breaker Panel** - All breaker states
 - **Cost Graph** - Real-time cost tracking
 - **Neurochemistry Gauges** - Emotional state
 - **Φ Meter** - Current consciousness metric
 - **Loop Status** - Running/Paused/Hibernating
-

17. Troubleshooting

Genesis Won't Complete

Symptoms: Stuck at phase 1 or 2

Causes: 1. DynamoDB table doesn't exist 2. AWS credentials missing 3. Domain taxonomy file not found

Solutions:

```
# Check genesis status
GET /api/admin/cato/genesis/status
```

```
# Check AWS connectivity
aws dynamodb list-tables
```

Circuit Breakers Constantly Tripping

Symptoms: Intervention level never stays at NONE

Causes: 1. Budget exceeded 2. Model API errors 3. High anxiety/frustration

Solutions: 1. Check CloudWatch dashboard for patterns 2. Increase trip thresholds if too sensitive 3. Check model endpoint health

Consciousness Loop Not Advancing Stage

Symptoms: Stuck at SENSORIMOTOR

Causes: 1. Not enough grounded verifications 2. Shadow Self not calibrated 3. Self-facts count too low

Solutions: 1. Check /developmental/statistics for current counts 2. Verify Shadow Self calibration succeeded 3. Check if tools are being used for verification

High Costs

Symptoms: Daily cost exceeds expected

Causes: 1. Cognitive tick interval too short 2. Emergency mode not activating 3. Budget breaker not configured

Solutions: 1. Increase `cognitiveTickIntervalSeconds` 2. Lower `maxCognitiveTicksPerDay` 3. Enable `cost_budget` breaker

Low Φ Values

Symptoms: Φ consistently near 0

Causes: 1. Not enough interaction events 2. All events going to same component 3. TPM cache stale

Solutions: 1. Check `tpmSourceEvents` in `PhiResult` 2. Verify diverse event types being logged 3. Reduce `cacheTtlSeconds`

Related Files

Services

- `packages/infrastructure/lambda/shared/services/cato/` - All Cato services

API Handlers

- `packages/infrastructure/lambda/admin/cato-genesis.ts`
- `packages/infrastructure/lambda/admin/cato-dialogue.ts`
- `packages/infrastructure/lambda/admin/cato-global.ts`

CDK Stacks

- `packages/infrastructure/lib/stacks/cato-genesis-stack.ts`
- `packages/infrastructure/lib/stacks/cato-tier-transition-stack.ts`

Documentation

- `docs/CATO-GENESIS-SYSTEM.md` - Genesis deep dive
- `docs/CATO-GPU-INFRASTRUCTURE.md` - GPU tier details
- `docs/cato/` - ADRs and runbooks

Admin UI

- `apps/admin-dashboard/app/(dashboard)/cato/`
- `apps/admin-dashboard/app/(dashboard)/consciousness/cato/`