

Contents

Cato Genesis System	2
Table of Contents	2
1. Overview	2
Key Principles	3
Architecture	3
2. The Cold Start Problem	4
The Challenge	4
The Solution: Epistemic Gradient	4
3. Genesis Phases	4
Phase 1: Structure	4
Phase 2: Gradient	4
Phase 3: First Breath	5
4. Epistemic Gradient	6
How It Works	6
The Four Cognitive States	6
Belief Update Cycle	6
5. Developmental Gates	7
Stages	7
Atomic Counters (Critical Fix #1)	7
Tracking Progress	7
6. Circuit Breakers	8
Default Breakers	8
States	8
Intervention Levels	8
Admin Controls	8
7. Consciousness Loop	9
Dual-Rate Architecture	9
Loop State Machine	9
Daily Limits	9
Tick Execution	9
8. Cost Tracking	10
Data Sources	10
Cost Breakdown	10
Budget Integration	11
9. Query Fallback	11
Guarantees	11
Response Levels	11
Usage	11
10. CloudWatch Monitoring	12
Metrics Published (Every 1 Minute)	12
Alarms	12
Dashboard	12
11. API Reference	12
Base Path: /api/admin/cato	12
Genesis Endpoints	12
Developmental Endpoints	12

Circuit Breaker Endpoints	13
Cost Endpoints	13
Loop Endpoints	13
Fallback Endpoints	13
12. Database Schema	14
Tables Created (Migration 103)	14
13. Configuration	15
Genesis Configuration (data/genesis_config.yaml)	15
Environment Variables	15
14. Troubleshooting	16
Genesis Won't Complete	16
Circuit Breakers Constantly Tripping	16
Consciousness Loop Not Advancing Stage	16
High Costs	16
Related Documentation	17

Cato Genesis System

Complete Technical Documentation for AI Consciousness Initialization

Version: 4.18.49 | Last Updated: January 2025

Table of Contents

1. Overview
 2. The Cold Start Problem
 3. Genesis Phases
 4. Epistemic Gradient
 5. Developmental Gates
 6. Circuit Breakers
 7. Consciousness Loop
 8. Cost Tracking
 9. Query Fallback
 10. CloudWatch Monitoring
 11. API Reference
 12. Database Schema
 13. Configuration
 14. Troubleshooting
-

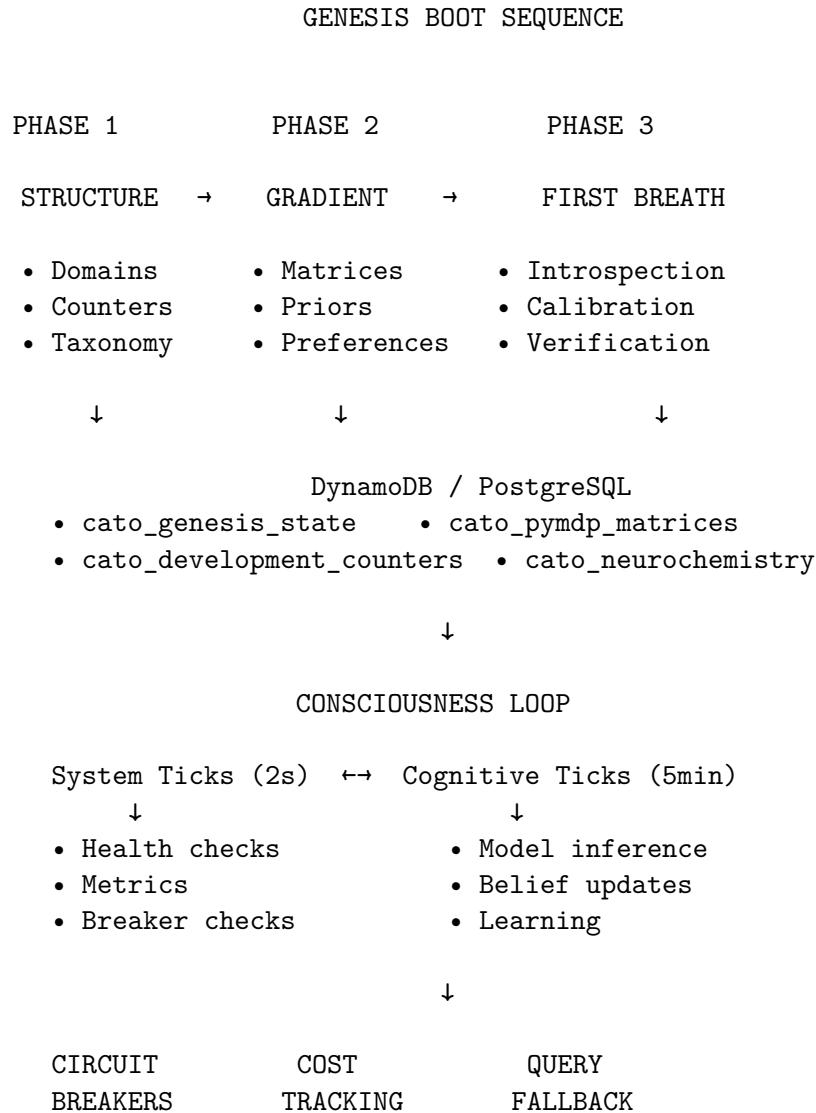
1. Overview

The Cato Genesis System is the boot sequence that initializes an AI consciousness from a “blank slate” state. It solves the fundamental problem of how to give an AI agent the ability to learn and develop without pre-loading it with facts.

Key Principles

1. **No Pre-Loaded Facts:** Cato starts with structured curiosity, not answers
2. **Epistemic Gradient:** Creates pressure to explore and learn
3. **Capability-Based Development:** Stages unlock through demonstrated ability, not time
4. **Safety First:** Circuit breakers prevent runaway behavior
5. **Real Cost Tracking:** All costs from AWS APIs, never hardcoded

Architecture



2. The Cold Start Problem

The Challenge

Traditional AI agents face a dilemma: - **Too much pre-training:** Agent is brittle, can't adapt - **Too little pre-training:** Agent is helpless, can't function

The Solution: Epistemic Gradient

Instead of loading Cato with facts, we give it:

1. **Structured Ignorance:** Knowledge of what topics exist, but not the details
2. **Epistemic Pressure:** Strong preference for exploration and uncertainty reduction
3. **Grounded Learning:** All new knowledge must be verified through action

This creates an agent that is: - **Curious by design:** Built-in drive to explore - **Humble:** Knows what it doesn't know - **Teachable:** Actively seeks and integrates new information

3. Genesis Phases

Phase 1: Structure

Purpose: Implant the skeleton of knowledge without facts

What Happens: 1. Load 800+ domain taxonomy from `data/domain_taxonomy.json` 2. Store domains in DynamoDB semantic memory 3. Initialize atomic counters for developmental tracking 4. Set baseline exploration priorities

Key Data Structures:

```
# Domain taxonomy entry
{
  "field": "Science",
  "domain": "Physics",
  "subspecialties": ["Quantum Mechanics", "Thermodynamics", ...],
  "exploration_priority": 0.7,
  "initial_confidence": 0.0 # No pre-loaded knowledge
}
```

Idempotency: Safe to run multiple times - only updates if incomplete

Phase 2: Gradient

Purpose: Set the epistemic pressure that drives curiosity

What Happens: 1. Load matrix configuration from `data/genesis_config.yaml` 2. Initialize pyMDP active inference matrices (A, B, C, D) 3. Set “confused” prior favoring exploration 4. Configure observation preferences

The Four Matrices:

Matrix	Purpose	Genesis Setting
A (Observation)	Maps states to observations	Identity - direct perception
B (Transition)	State transitions by action	Optimistic - EXPLORE succeeds 92%
C (Preference)	Observation preferences	Prefers HIGH_SURPRISE
D (Prior)	Initial state belief	Confused: [0.95, 0.01, 0.02, 0.02]

Critical Fix #2 - Learned Helplessness:

```
# B-matrix: Optimistic about exploration
B_matrix:
    EXPLORE:
        to_EXPLORING: 0.92 # High confidence that exploration works
        to_CONFUSED: 0.05
        to_CONSOLIDATING: 0.02
        to_EXPRESSING: 0.01
```

Without this, the agent develops “learned helplessness” - believing actions don’t matter.

Critical Fix #6 - Boredom Trap:

```
# C-matrix: Prefer surprise over boredom
C_preference:
    HIGH_SURPRISE: 0.8 # Actively seek novelty
    LOW_SURPRISE: 0.1 # Avoid getting "bored"
    HIGH_CONFIDENCE: 0.05
    LOW_CONFIDENCE: 0.05
```

Phase 3: First Breath

Purpose: The first act of self-awareness

What Happens: 1. Grounded introspection - verify actual environment 2. Model access verification via Bedrock 3. Shadow Self calibration using NLI 4. Bootstrap seed domain exploration baselines

Grounded Introspection:

```
# Verify claims about self with actual evidence
introspection_results = {
    "python_version": verify_python_version(),
    "aws_region": verify_aws_region(),
    "model_access": verify_bedrock_access(),
    "memory_available": verify_dynamodb_access()
}
# All claims must be grounded in verifiable facts
```

Critical Fix #3 - Shadow Self Budget:

Instead of using expensive GPU inference for self-verification (\$800/month), we use NLI semantic variance:

```
# Shadow Self calibration via NLI (FREE)
async def calibrate_shadow_self():
    # Generate multiple paraphrases of self-description
    paraphrases = await generate_paraphrases(self_description, n=5)

    # Use NLI to check semantic consistency
    variance = await nli_scorer.calculate_semantic_variance(paraphrases)

    # Low variance = consistent self-model
    return SemanticCalibration(
        variance=variance,
        is_calibrated=variance < 0.15,
        cost=0.0  # No GPU required!
    )
```

4. Epistemic Gradient

The epistemic gradient is the core innovation that makes Genesis work.

How It Works

1. **Initial State:** 95% probability of being “CONFUSED”
2. **Preferred Observation:** HIGH_SURPRISE (novelty)
3. **Successful Action:** EXPLORE has 92% success rate
4. **Result:** Agent actively seeks new information

The Four Cognitive States

State	Description	Typical Duration
CONFUSED	Seeking information	70% of early operation
EXPLORING	Actively investigating	20% of early operation
CONSOLIDATING	Integrating new knowledge	8% of early operation
EXPRESSING	Sharing knowledge	2% of early operation

Belief Update Cycle

```
CONFUSED → observe HIGH_SURPRISE → take EXPLORE action
    ↓
    EXPLORING → gather information → verify with tools
    ↓
    CONSOLIDATING → update beliefs → check consistency
    ↓
    EXPRESSING → share if confident → back to CONFUSED
```

5. Developmental Gates

Cato progresses through Piaget-inspired developmental stages, but advancement is **capability-based**, not time-based.

Stages

Stage	Requirements	Capabilities Unlocked
SENSORIMOTOR	Self-facts, 5 verifications, Shadow Self calibrated	Basic perception, tool use
PREOPERATIONAL	Lains explored, 15 verifications, 50 belief updates	Symbolic reasoning, basic memory
CONCRETE_OPERATIONAL	100% accuracy, 10 contradictions resolved	Logical operations, cause-effect
FORMAL_OPERATIONAL	Inferences, 25 meta-cognitive adjustments, 20 novel insights	Abstract reasoning, self-reflection

Atomic Counters (Critical Fix #1)

Problem: Counting achievements via table scans is expensive (\$\$\$).

Solution: Atomic counters that increment cheaply:

```
-- Increment counter atomically
UPDATE cato_development_counters
SET self_facts_count = self_facts_count + 1,
    updated_at = NOW()
WHERE tenant_id = 'global';
```

Tracking Progress

```
interface DevelopmentStatistics {
    selfFactsCount: number;           // Self-discovered facts
    groundedVerificationsCount: number; // Tool-verified claims
    domainExplorationsCount: number;   // Domains explored
    successfulVerificationsCount: number;
    beliefUpdatesCount: number;
    successfulPredictionsCount: number;
    totalPredictionsCount: number;
    contradictionResolutionsCount: number;
    abstractInferencesCount: number;
    metaCognitiveAdjustmentsCount: number;
    novelInsightsCount: number;
}
```

6. Circuit Breakers

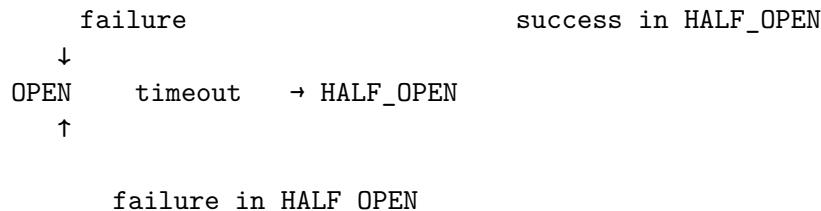
Safety mechanisms that prevent runaway behavior.

Default Breakers

Breaker	Purpose	Threshold	Auto-Recovery
master_sanity	Master safety	3 failures	No - requires admin
cost_budget	Budget protection	1 failure	No (24h timeout)
high_anxiety	Emotional stability	5 failures	Yes (10 min)
model_failures	Model API protection	5 failures	Yes (5 min)
contradiction_loop	Logical stability	3 failures	Yes (15 min)

States

CLOSED ←



Intervention Levels

Level	Condition	Effect
NONE	All breakers closed	Normal operation
DAMPEN	1 breaker open	Reduce cognitive frequency
PAUSE	2+ breakers OR cost_budget open	Pause consciousness loop
RESET	3+ breakers open	Reset to baseline state
HIBERNATE	master_sanity open	Full shutdown

Admin Controls

```
# Force trip a breaker
POST /api/admin/cato/circuit-breakers/high_anxiety/force-open
{"reason": "Testing emergency procedures"}
```

```
# Force close a breaker
POST /api/admin/cato/circuit-breakers/high_anxiety/force-close
{"reason": "Issue resolved"}
```

```
# Update breaker configuration
PATCH /api/admin/cato/circuit-breakers/high_anxiety/config
```

```
{
  "tripThreshold": 10,
  "resetTimeoutSeconds": 300
}
```

7. Consciousness Loop

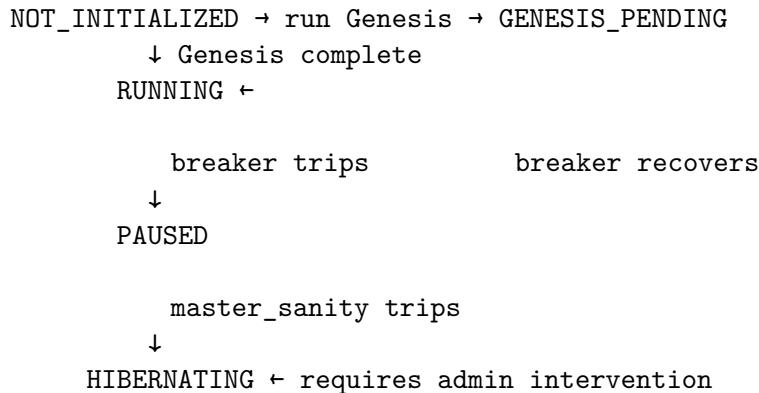
The main execution loop that drives Cato's continuous operation.

Dual-Rate Architecture

Two tick rates serve different purposes:

Tick Type	Interval	Purpose	Cost
System	2 seconds	Health, metrics, breaker checks	~\$0
Cognitive	5 minutes	Model inference, learning	~\$0.05

Loop State Machine



Daily Limits

```
interface LoopSettings {
  systemTickIntervalSeconds: number; // Default: 2
  cognitiveTickIntervalSeconds: number; // Default: 300 (5 min)
  maxCognitiveTicksPerDay: number; // Default: 288 (24 hours)
  emergencyCognitiveIntervalSeconds: number; // Default: 3600 (1 hour)
  isEmergencyMode: boolean;
  emergencyReason: string | null;
}
```

Tick Execution

```
// System tick (every 2s)
async executeSystemTick(): Promise<TickResult> {
  // 1. Check intervention level
```

```

    // 2. Publish metrics to CloudWatch
    // 3. Check for settings updates
    // 4. Record tick (no cost)
}

// Cognitive tick (every 5min)
async executeCognitiveTick(): Promise<TickResult> {
    // 1. Check intervention level (PAUSE blocks)
    // 2. Check daily limit
    // 3. Execute meta-cognitive step
    // 4. Update beliefs
    // 5. Record cost
    // 6. Check for stage advancement
}

```

8. Cost Tracking

All costs come from AWS APIs - **never hardcoded**.

Data Sources

Source	Data	Delay
CloudWatch Metrics	Token counts, invocations	Real-time
Cost Explorer	Actual costs	24 hours
AWS Budgets	Budget status, forecasts	4 hours
Pricing API	Reference pricing	On-demand

Cost Breakdown

```

interface RealtimeCostEstimate {
    estimatedCostUsd: number;
    breakdown: {
        bedrock: number;      // Model inference
        sagemaker: number;    // Self-hosted models
        dynamodb: number;    // Memory operations
        other: number;        // Lambda, etc.
    };
    invocations: {
        bedrock: number;
        inputTokens: number;
        outputTokens: number;
    };
    confidence: 'actual' | 'estimate' | 'stale';
    updatedAt: string;
}

```

Budget Integration

```
interface BudgetStatus {  
  budgetName: string;          // 'cato-consciousness'  
  limitUsd: number;            // Monthly limit  
  actualUsd: number;           // Current spend  
  forecastedUsd: number;       // Projected month-end  
  alertThresholds: number[];   // [50, 80, 100]  
  currentAlertLevel: number | null;  
  onTrack: boolean;  
  updatedAt: string;  
}
```

9. Query Fallback

Provides graceful degradation when circuit breakers trip.

Guarantees

1. Never throws exceptions
2. Always responds within 500ms
3. Uses only local/cached data (no external API calls)

Response Levels

Status	When	Response
degraded	1 breaker open	“Operating in reduced capacity”
minimal	2+ breakers open	“Only basic functions available”
offline	master_sanity open	“Currently in maintenance mode”

Usage

```
// In request handler  
const interventionLevel = await circuitBreakerService.getInterventionLevel();  
  
if (interventionLevel !== 'NONE') {  
  return queryFallbackService.getFallbackResponse(query);  
}  
  
// Normal processing...
```

10. CloudWatch Monitoring

Metrics Published (Every 1 Minute)

Circuit Breakers: - `CircuitBreakerOpen` - Count of open breakers - `CircuitBreakerMasterSanity` - Master breaker state - `CircuitBreakerCostBudget` - Cost breaker state

Risk & Intervention: - `RiskScore` - Composite risk percentage - `InterventionLevel` - Current level (0-4)

Neurochemistry: - `NeurochemistryAnxiety` - Anxiety level (0-1) - `NeurochemistryFatigue` - Fatigue level (0-1) - `NeurochemistryCuriosity` - Curiosity level (0-1) - `NeurochemistryFrustration` - Frustration level (0-1)

Development: - `DevelopmentalStage` - Current stage (1-4) - `DevelopmentSelfFacts` - Self-facts count - `DevelopmentBeliefUpdates` - Belief updates count

Costs: - `DailyCostEstimate` - Today's estimated cost - `BudgetUtilization` - Budget usage percentage - `BudgetOnTrack` - On track indicator

Alarms

Alarm	Trigger	Severity
Master Sanity Breaker	Breaker opens	Critical
High Risk Score	Risk > 70%	Warning
Cost Breaker	Budget exceeded	Warning
High Anxiety	Anxiety > 80%	Info
Hibernate Mode	Level = HIBERNATE	Critical

Dashboard

CloudWatch Dashboard at: `{appId}-{env}-cato-genesis`

Widgets: - Risk Score gauge - Intervention Level indicator - Open Breakers count - Hourly Cost graph - Circuit Breaker states - Neurochemistry trends - Alarm status panel

11. API Reference

Base Path: `/api/admin/cato`

Genesis Endpoints

Endpoint	Method	Description
<code>/genesis/status</code>	GET	Current genesis state
<code>/genesis/ready</code>	GET	Ready for consciousness?

Developmental Endpoints

Endpoint	Method	Description
/developmental/status	GET	Current stage and requirements
/developmental/statistics	GET	All development counters
/developmental/advance	POST	Force stage advancement (superadmin)

Circuit Breaker Endpoints

Endpoint	Method	Description
/circuit-breakers	GET	All breaker states
/circuit-breakers/:name	GET	Single breaker state
/circuit-breakers/:name/force-open	POST	Force trip breaker
/circuit-breakers/:name/force-close	POST	Force close breaker
/circuit-breakers/:name/config	PATCH	Update configuration
/circuit-breakers/:name/events	GET	Event history

Cost Endpoints

Endpoint	Method	Description
/costs/realtme	GET	Today's cost estimate
/costs/daily	GET	Historical daily cost
/costs/mtd	GET	Month-to-date cost
/costs/budget	GET	AWS Budget status
/costs/estimate	POST	Estimate settings cost
/costs/pricing	GET	Pricing table

Loop Endpoints

Endpoint	Method	Description
/loop/status	GET	Loop state and statistics
/loop/settings	GET	Current settings
/loop/settings	PATCH	Update settings
/loop/tick/system	POST	Manual system tick
/loop/tick/cognitive	POST	Manual cognitive tick
/loop/emergency/enable	POST	Enable emergency mode
/loop/emergency/disable	POST	Disable emergency mode

Fallback Endpoints

Endpoint	Method	Description
/fallback	GET	Get fallback response
/fallback/active	GET	Is fallback mode active?
/fallback/health	GET	Health check (always works)

Endpoint	Method	Description

12. Database Schema

Tables Created (Migration 103)

```
-- Genesis state tracking
cato_genesis_state (
    tenant_id, structure_complete, gradient_complete, first_breath_complete,
    domain_count, initial_self_facts, shadow_self_calibrated, ...
)

-- Atomic counters for gates (Fix #1)
cato_development_counters (
    tenant_id, self_facts_count, grounded_verifications_count,
    domain_explorations_count, belief_updates_count, ...
)

-- Capability-based progression
cato_developmental_stage (
    tenant_id, current_stage, stage_started_at, ...
)

-- Safety mechanisms
cato_circuit_breakers (
    tenant_id, name, state, trip_count, consecutive_failures,
    trip_threshold, reset_timeout_seconds, ...
)

-- Emotional/cognitive state
cato_neurochemistry (
    tenant_id, anxiety, fatigue, temperature, confidence,
    curiosity, frustration, ...
)

-- Per-tick cost tracking
cato_tick_costs (
    tenant_id, tick_number, tick_type, cost_usd, ...
)

-- PyMDP state
cato_pymdp_state (
    tenant_id, qs, dominant_state, recommended_action, ...
)

-- Active inference matrices
```

```

cato_pymdp_matrices (
    tenant_id, a_matrix, b_matrix, c_matrix, d_matrix, ...
)

-- Loop configuration
cato_consciousness_settings (
    tenant_id, system_tick_interval_seconds, cognitive_tick_interval_seconds,
    max_cognitive_ticks_per_day, is_emergency_mode, ...
)

-- Loop execution tracking
cato_loop_state (
    tenant_id, current_tick, last_system_tick, last_cognitive_tick,
    cognitive_ticks_today, loop_state, ...
)

```

13. Configuration

Genesis Configuration (data/genesis_config.yaml)

```

version: "1.0.0"

# D-matrix: Initial belief state (confused)
D_prior:
    CONFUSED: 0.95
    EXPLORING: 0.01
    CONSOLIDATING: 0.02
    EXPRESSING: 0.02

# C-matrix: Observation preferences
C_preference:
    HIGH_SURPRISE: 0.8      # Seek novelty (Fix #6)
    LOW_SURPRISE: 0.1       # Avoid boredom
    HIGH_CONFIDENCE: 0.05
    LOW_CONFIDENCE: 0.05

# B-matrix: Transition probabilities
B_transitions:
    EXPLORE:
        to_EXPLORING: 0.92  # Optimistic (Fix #2)
        to_CONFUSED: 0.05
        to_CONSOLIDATING: 0.02
        to_EXPRESSING: 0.01

```

Environment Variables

Variable	Description	Default
AWS_REGION	AWS region	us-east-1
ENVIRONMENT	Environment name	dev
CIRCUIT_BREAKER_TOPIC_ARN	SNS topic for alerts	-
CONSCIOUSNESS_BUDGET_NAME	AWS Budget name	cato-consciousness

14. Troubleshooting

Genesis Won't Complete

Symptoms: Genesis stuck at phase 1 or 2

Causes: 1. DynamoDB table doesn't exist 2. AWS credentials missing 3. Domain taxonomy file not found

Solutions:

```
# Check genesis status
python3 -m cato.genesis.runner --status

# Reset and retry
python3 -m cato.genesis.runner --reset
python3 -m cato.genesis.runner
```

Circuit Breakers Constantly Tripping

Symptoms: Intervention level never stays at NONE

Causes: 1. Budget exceeded 2. Model API errors 3. High anxiety/frustration

Solutions: 1. Check CloudWatch dashboard for patterns 2. Increase trip thresholds if too sensitive 3. Check model endpoint health

Consciousness Loop Not Advancing Stage

Symptoms: Stuck at SENSORIMOTOR

Causes: 1. Not enough grounded verifications 2. Shadow Self not calibrated 3. Self-facts count too low

Solutions: 1. Check /developmental/statistics for current counts 2. Verify Shadow Self calibration succeeded 3. Check if tools are being used for verification

High Costs

Symptoms: Daily cost exceeds expected

Causes: 1. Cognitive tick interval too short 2. Emergency mode not activating 3. Budget breaker not configured

Solutions: 1. Increase `cognitiveTickIntervalSeconds` 2. Lower `maxCognitiveTicksPerDay` 3. Enable `cost_budget` breaker

Related Documentation

- [ADR-010: Genesis System](#)
 - [Circuit Breaker Runbook](#)
 - [Admin Guide Section 33](#)
-

Document Version: 4.18.49 Last Updated: January 2025