

Contents

ADR-005: Circadian Budget Management	1
Status	1
Context	1
Uncontrolled Curiosity Cost Model	1
Decision	1
Operating Modes	2
Budget Hierarchy	2
Night Mode Benefits	2
Architecture	2
Implementation	3
TypeScript Service	3
Consequences	7
Positive	7
Negative	7
Admin Configuration	7
Scaling Budget with Users	8
References	8

ADR-005: Circadian Budget Management

Status

Accepted

Context

Cato's curiosity is designed to be autonomous and continuous. Without constraints, this creates a runaway cost problem:

Uncontrolled Curiosity Cost Model

Curiosity loop = 1 question + 1 answer + grounding

Average cost per loop = \$0.01 (Haiku) to \$0.10 (Sonnet with tools)

Continuous operation (24/7):

- 1 loop/second = 86,400 loops/day
- At \$0.05 average = \$4,320/day = \$129,600/month

This exceeds the \$500/month target by 260x!

Additionally, running curiosity during peak user hours: 1. Competes for model capacity 2. Increases latency for user queries 3. Wastes money on real-time inference (vs. batch pricing)

Decision

Implement a **circadian rhythm** for Cato with distinct day/night operational modes and hard budget caps.

Operating Modes

Mode	Hours (UTC)	Behavior	Budget
DAY	6 AM - 2 AM	Queue curiosity, serve users	\$0 exploration
NIGHT	2 AM - 6 AM	Batch process exploration	Up to \$15/night
EMERGENCY	Any	Over budget, minimal ops	\$0 all activity

Budget Hierarchy

Monthly Budget: \$500 (admin-configurable)

User Interactions: \$400 (80%)

 Real-time inference

 Cache misses

Autonomous Exploration: \$100 (20%)

 Night-mode curiosity: \$85

 Tool grounding: \$10

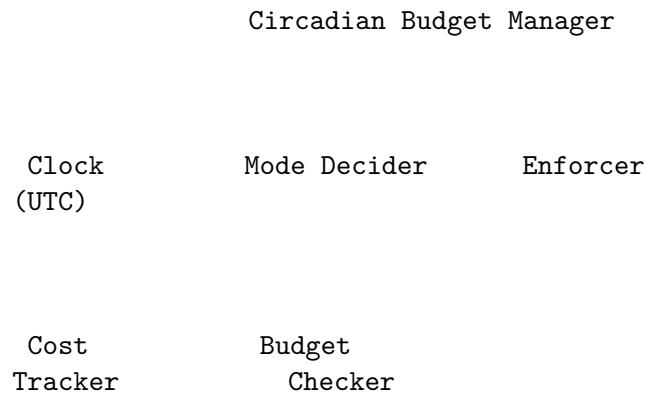
 Memory consolidation: \$5

Daily Exploration Cap: \$15 (prevents single bad night)

Night Mode Benefits

1. **Bedrock Batch API:** 50% discount on batch inference
2. **Lower traffic:** Less competition for resources
3. **Consolidation:** Natural time for memory consolidation
4. **Global timing:** 2-6 AM UTC covers low-traffic worldwide

Architecture



DAY MODE	NIGHT MODE	EMERGENCY
Queue curiosity	Process queue (batch)	Serve from cache

Implementation

TypeScript Service

```

import { DynamoDBClient } from '@aws-sdk/client-dynamodb';
import { DynamoDBDocumentClient, GetCommand, UpdateCommand } from '@aws-sdk/lib-dynamodb';

export enum OperatingMode {
  DAY = 'day',
  NIGHT = 'night',
  EMERGENCY = 'emergency'
}

export interface BudgetConfig {
  monthlyLimit: number;           // Default: $500
  dailyExplorationLimit: number;  // Default: $15
  explorationRatio: number;       // Default: 0.20
  nightStartHour: number;         // Default: 2 (2 AM UTC)
  nightEndHour: number;           // Default: 6 (6 AM UTC)
  emergencyThreshold: number;    // Default: 0.90
}

export interface BudgetStatus {
  mode: OperatingMode;
  dailySpend: number;
  monthlySpend: number;
  dailyRemaining: number;
  monthlyRemaining: number;
  canExplore: boolean;
  nextModeChange: Date;
}

export class CircadianBudgetManager {
  private readonly docClient: DynamoDBDocumentClient;
  private readonly configTable: string;
  private readonly costsTable: string;
  private config: BudgetConfig | null = null;
  private lastConfigRefresh: Date | null = null;
}

```

```

constructor(
  configTable: string = 'cato-config',
  costsTable: string = 'cato-costs',
  region: string = 'us-east-1'
) {
  const client = new DynamoDBClient({ region });
  this.docClient = DynamoDBDocumentClient.from(client);
  this.configTable = configTable;
  this.costsTable = costsTable;
}

async getConfig(): Promise<BudgetConfig> {
  const now = new Date();

  // Refresh config every 5 minutes
  if (
    this.config === null ||
    this.lastConfigRefresh === null ||
    now.getTime() - this.lastConfigRefresh.getTime() > 300000
  ) {
    const response = await this.docClient.send(new GetCommand({
      TableName: this.configTable,
      Key: { pk: 'CONFIG', sk: 'BUDGET' }
    }));

    if (response.Item) {
      this.config = {
        monthlyLimit: response.Item.monthlyLimit ?? 500,
        dailyExplorationLimit: response.Item.dailyExplorationLimit ?? 15,
        explorationRatio: response.Item.explorationRatio ?? 0.20,
        nightStartHour: response.Item.nightStartHour ?? 2,
        nightEndHour: response.Item.nightEndHour ?? 6,
        emergencyThreshold: response.Item.emergencyThreshold ?? 0.90
      };
    } else {
      // Default config
      this.config = {
        monthlyLimit: 500,
        dailyExplorationLimit: 15,
        explorationRatio: 0.20,
        nightStartHour: 2,
        nightEndHour: 6,
        emergencyThreshold: 0.90
      };
    }
  }

  this.lastConfigRefresh = now;
}

```

```

    }

    return this.config;
}

async getMode(): Promise<OperatingMode> {
    const config = await this.getConfig();
    const { dailySpend, monthlySpend } = await this.getSpendCounters();
    const now = new Date();
    const hour = now.getUTCHours();

    // Check emergency (budget exhausted)
    if (monthlySpend >= config.monthlyLimit * config.emergencyThreshold) {
        return OperatingMode.EMERGENCY;
    }

    // Check daily exploration limit
    if (dailySpend >= config.dailyExplorationLimit) {
        return OperatingMode.DAY; // No exploration but still serving
    }

    // Check time of day
    if (hour >= config.nightStartHour && hour < config.nightEndHour) {
        return OperatingMode.NIGHT;
    }

    return OperatingMode.DAY;
}

async canExplore(): Promise<boolean> {
    const mode = await this.getMode();
    if (mode === OperatingMode.EMERGENCY) {
        return false;
    }

    const config = await this.getConfig();
    const { dailySpend } = await this.getSpendCounters();

    return dailySpend < config.dailyExplorationLimit;
}

async recordCost(
    amount: number,
    category: 'inference' | 'curiosity' | 'grounding' | 'consolidation',
    model: string,
    tokensInput: number = 0,
    tokensOutput: number = 0
): Promise<void> {

```

```

const now = new Date();
const month = now.toISOString().slice(0, 7); // YYYY-MM
const day = now.toISOString().slice(0, 10); // YYYY-MM-DD

await this.docClient.send(new UpdateCommand({
  TableName: this.costsTable,
  Key: { pk: `COST#${month}`, sk: now.toISOString() },
  UpdateExpression: 'SET #amount = :amount, #category = :category, #model = :model, #day = :day',
  ExpressionAttributeNames: {
    '#amount': 'amount',
    '#category': 'category',
    '#model': 'model',
    '#day': 'day',
    '#tokensIn': 'tokensInput',
    '#tokensOut': 'tokensOutput'
  },
  ExpressionAttributeValues: {
    ':amount': amount,
    ':category': category,
    ':model': model,
    ':day': day,
    ':tokensIn': tokensInput,
    ':tokensOut': tokensOutput
  }
})));
}

async getStatus(): Promise<BudgetStatus> {
  const config = await this.getConfig();
  const mode = await this.getMode();
  const { dailySpend, monthlySpend } = await this.getSpendCounters();
  const canExplore = await this.canExplore();

  // Calculate next mode change
  const now = new Date();
  const hour = now.getUTCHours();
  let nextModeChange: Date;

  if (hour < config.nightStartHour) {
    nextModeChange = new Date(now);
    nextModeChange.setUTCHours(config.nightStartHour, 0, 0, 0);
  } else if (hour < config.nightEndHour) {
    nextModeChange = new Date(now);
    nextModeChange.setUTCHours(config.nightEndHour, 0, 0, 0);
  } else {
    nextModeChange = new Date(now);
    nextModeChange.setUTCDate(nextModeChange.getUTCDate() + 1);
    nextModeChange.setUTCHours(config.nightStartHour, 0, 0, 0);
  }
}

```

```

    }

    return {
      mode,
      dailySpend,
      monthlySpend,
      dailyRemaining: Math.max(0, config.dailyExplorationLimit - dailySpend),
      monthlyRemaining: Math.max(0, config.monthlyLimit - monthlySpend),
      canExplore,
      nextModeChange
    };
  }

  private async getSpendCounters(): Promise<{ dailySpend: number; monthlySpend: number }> {
    // Implementation queries DynamoDB for current spend
    // Aggregates by day and month
    return { dailySpend: 0, monthlySpend: 0 }; // Placeholder
  }
}

```

Consequences

Positive

- **Predictable costs:** Hard caps prevent budget overrun
- **Optimal pricing:** Night-mode uses Bedrock batch (50% off)
- **User priority:** Day mode focuses on user interactions
- **Natural rhythm:** Consolidation aligns with low-traffic periods

Negative

- **Delayed learning:** Curiosity queued until night
- **Global timing:** 2-6 AM UTC may not suit all regions
- **Complexity:** Two operational modes to manage
- **Queue management:** Must handle curiosity queue

Admin Configuration

The budget manager is admin-configurable via the Radianit Admin dashboard:

Setting	Default	Range	Description
Monthly Limit	\$500	\$100-\$10,000	Total monthly budget
Daily Exploration	\$15	\$5-\$100	Max daily curiosity spend
Night Start	2 AM	0-23	When night mode begins (UTC)
Night End	6 AM	0-23	When night mode ends (UTC)
Emergency Threshold	90%	50-99%	When to enter emergency mode

Scaling Budget with Users

As user base grows, budget should scale:

Users	Monthly Budget	Daily Exploration	Rationale
0-10K	\$500	\$15	Starting budget
10K-100K	\$2,000	\$60	4x growth
100K-1M	\$10,000	\$300	Supporting infrastructure
1M-10M	\$100,000	\$3,000	At scale

References

- [AWS Bedrock Batch Inference](#)
- [Circadian Rhythms in AI Systems](#)
- [Cost Optimization for ML Workloads](#)