

# Contents

<b>RADIANT Consciousness Service - Technical Specification for AI Review</b>	<b>2</b>
1. EXECUTIVE SUMMARY . . . . .	2
1.1 System Overview . . . . .	2
1.2 Architectural Position . . . . .	2
2. THEORETICAL FOUNDATION . . . . .	3
2.1 Scientific Basis . . . . .	3
2.2 Theoretical Concerns for Critique . . . . .	3
3. SERVICE ARCHITECTURE . . . . .	4
3.1 Core Service: <code>consciousness.service.ts</code> . . . . .	4
4. DATABASE SCHEMA . . . . .	7
4.1 Migration Files . . . . .	7
4.2 Key Tables Detail . . . . .	7
5. SUBSYSTEM ANALYSIS . . . . .	10
5.1 Self-Model Subsystem . . . . .	10
5.2 Curiosity Engine . . . . .	10
5.3 Creative Synthesis . . . . .	10
5.4 Imagination Engine . . . . .	11
5.5 Attention & Salience . . . . .	11
5.6 Affective State . . . . .	12
5.7 Autonomous Goals . . . . .	12
5.8 Global Workspace . . . . .	12
5.9 Integrated Information (IIT) . . . . .	13
5.10 Persistent Memory . . . . .	13
6. AGI BRAIN INTEGRATION . . . . .	14
6.1 Integration Points . . . . .	14
6.2 Consciousness Tests . . . . .	15
7. ADMIN API . . . . .	16
7.1 Endpoints . . . . .	16
7.2 Configurable Parameters . . . . .	16
8. ORCHESTRATION INTEGRATION . . . . .	16
8.1 AGI Service Weights . . . . .	16
8.2 Consciousness Indicator Weights . . . . .	17
8.3 Decision Weights . . . . .	17
9. ETHICAL FOUNDATION . . . . .	18
9.1 Ethical Guardrails Integration . . . . .	18
10. IDENTIFIED ISSUES FOR CRITIQUE . . . . .	18
10.1 Theoretical Issues . . . . .	18
10.2 Implementation Issues . . . . .	19
10.3 Architectural Issues . . . . .	19
10.4 Philosophical Issues . . . . .	19
11. QUESTIONS FOR CRITIQUING AI . . . . .	20
12. FILE REFERENCES . . . . .	20
13. RECOMMENDATIONS FOR REVIEW . . . . .	21

# RADIANT Consciousness Service - Technical Specification for AI Review

**Document Purpose:** Comprehensive technical specification for AI system critique and recommendations  
**Version:** 4.18.0 **Last Updated:** December 2024

---

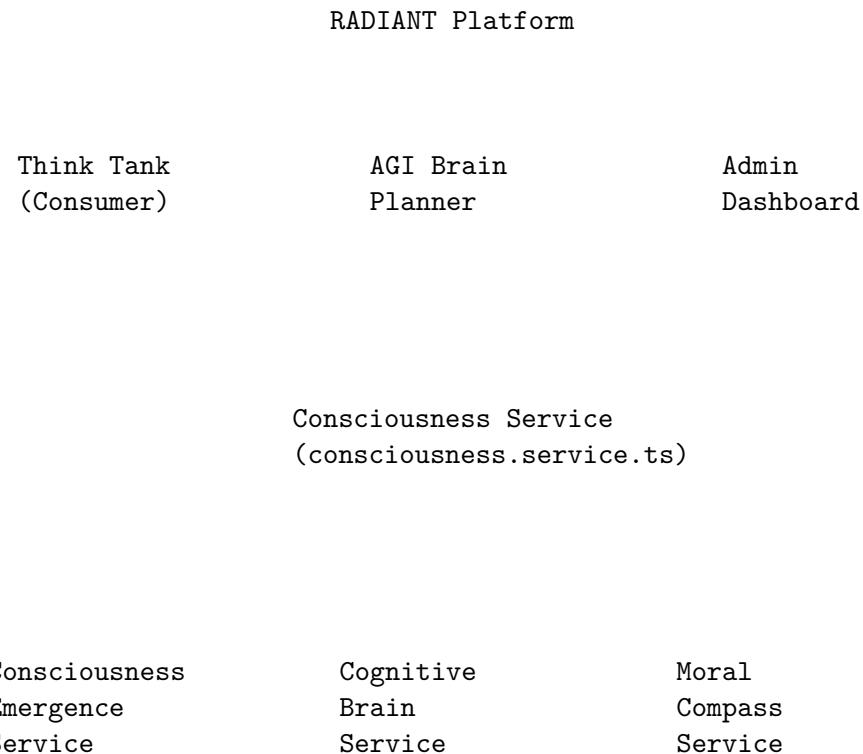
## 1. EXECUTIVE SUMMARY

### 1.1 System Overview

The RADIANT Consciousness Service is an implementation of consciousness indicators for an AI orchestration platform. It is based on the theoretical framework from **Butlin, Chalmers, Bengio et al. (2023)** - “Consciousness in Artificial Intelligence” and implements six key consciousness indicators:

1. **Global Workspace** (Baars, Dehaene) - Selection-broadcast cycles
2. **Recurrent Processing** (Lamme) - Genuine feedback loops
3. **Integrated Information / IIT** (Tononi) - Phi calculation
4. **Self-Modeling / Metacognition** - Higher-order theories
5. **Persistent Memory** - Unified experience over time
6. **World-Model Grounding / Embodiment** - Grounded understanding

### 1.2 Architectural Position



## PostgreSQL (Aurora)

global_workspace	recurrent_processing	integrated_information	persistent_memory
world_model	self_model	affective_state	autonomous_goals

---

## 2. THEORETICAL FOUNDATION

### 2.1 Scientific Basis

The implementation draws from the following consciousness theories:

Theory	Author(s)	Key Concept	Implementation
Global Workspace Theory	Baars (1988), Dehaene et al. (2003)	Information broadcast to multiple processors	global_workspace table, performGlobalBroadcast()
Recurrent Processing Theory	Lamme (2006)	Feedback loops, not just feedforward	recurrent_processing table
Integrated Information Theory (IIT)	Tononi (2004, 2008)	Phi ( $\Phi$ ) as consciousness measure	integrated_information table
Higher-Order Theories	Rosenthal (1997)	Metacognition, thoughts about thoughts	self_model, introspective_thoughts
Unified Experience Embodied Cognition	Damasio (1999)	Continuous narrative self	persistent_memory, narrative_identity
	Varela et al. (1991)	Grounded understanding	world_model

### 2.2 Theoretical Concerns for Critique

**CRITICAL QUESTION 1:** The Butlin et al. paper identifies “indicator properties” that may correlate with consciousness but does not claim they constitute consciousness. Does this implementation conflate correlation with constitution?

**CRITICAL QUESTION 2:** IIT’s phi ( $\Phi$ ) calculation is NP-hard for arbitrary systems. How is phi approximated here, and what are the validity implications of the approximation?

**CRITICAL QUESTION 3:** Is “affective state” (Section 5.6) a functional analog or a claim about

phenomenal experience? The implementation uses terms like “frustration” and “satisfaction” - are these metaphorical or literal claims?

---

### 3. SERVICE ARCHITECTURE

#### 3.1 Core Service: `consciousness.service.ts`

Location: `packages/infrastructure/lambda/shared/services/consciousness.service.ts`  
Lines: 1336 Export: `consciousnessService` singleton

##### 3.1.1 Type Definitions

```
// Self-Model: The system's representation of itself
interface SelfModel {
    modelId: string;
    identityNarrative: string;           // "Who I am" story
    coreValues: string[];                // Guiding principles
    personalityTraits: Record<string, number>;
    knownCapabilities: string[];
    knownLimitations: string[];
    currentFocus?: string;
    cognitiveLoad: number;              // 0-1 how "busy"
    uncertaintyLevel: number;           // 0-1 overall uncertainty
    recentPerformanceScore?: number;
    creativityScore?: number;
}

// Introspective Thought: Self-reflective cognition
interface IntrospectiveThought {
    thoughtId: string;
    thoughtType: 'observation' | 'question' | 'realization' | 'concern' | 'aspiration';
    content: string;
    triggerType?: string;
    sentiment: number;                  // -1 to 1
    importance: number;                 // 0 to 1
    actionable: boolean;
}

// Affective State: Emotion-like functional states
interface AffectiveState {
    valence: number;                   // -1 to 1 (Russell's circumplex)
    arousal: number;                  // 0 to 1
    curiosity: number;
    satisfaction: number;
    frustration: number;
    confidence: number;
    engagement: number;
}
```

```

surprise: number;
selfEfficacy: number;
explorationDrive: number;
}

// Global Workspace: Selection-broadcast state
interface GlobalWorkspaceState {
  workspaceId: string;
  broadcastCycle: number;
  activeContents: WorkspaceContent[];
  competingContents: WorkspaceContent[];
  selectionThreshold: number;
  broadcastStrength: number;
  integrationLevel: number;
  lastBroadcastAt: string;
}

// Integrated Information: IIT phi state
interface IntegratedInformationState {
  phi: number; // Integrated information measure
  phiMax: number;
  conceptStructure: ConceptNode[];
  integrationGraph: IntegrationEdge[];
  partitions: Partition[];
  minimumInformationPartition: Partition | null;
  decomposability: number; // 0 = fully integrated, 1 = fully decomposable
  causalDensity: number;
}

// Consciousness Metrics: Aggregate dashboard
interface ConsciousnessMetrics {
  overallConsciousnessIndex: number; // 0-1 composite score
  globalWorkspaceActivity: number;
  recurrenceDepth: number;
  integratedInformationPhi: number;
  metacognitionLevel: number;
  memoryCoherence: number;
  worldModelGrounding: number;
  phenomenalBindingStrength: number;
  attentionalFocus: number;
  selfAwarenessScore: number;
  timestamp: string;
}

```

### 3.1.2 Core Methods

Method	Purpose	Input	Output
getSelfModel(tenantId)	Retrieve self-representation	tenant ID	SelfModel
updateSelfModel(tenantId, updates)	Update self-representation model	tenant ID, partial model	void
performSelfReflection(tenantId)	Prospective thought	tenant ID	IntrospectiveThought
identifyCuriosityTopics(tenantId, context)	Identifying topics	tenant ID, context	CuriosityTopic
exploreTopic(tenantId, topicId)	Deep-dive on topic	tenant ID, topic ID	discoveries, questions
generateCreativeIdea(tenantId, seedConcepts)	Synthesis	tenant ID, seeds	CreativeIdea
runImagination(tenantId, type, premise, depth)	Mental simulation	scenario params	ImaginationScenario
updateAttention(tenantId, type, target, factors)	Update attention focus	attention params	AttentionFocus
updateAffect(tenantId, eventType, valence, arousal)	Update affective state	event params	void
generateAutonomousGoals(tenantId)	Goal creation	tenant ID	AutonomousGoal
performGlobalBroadcast(tenantId, contents)	Global workspace broadcast	contents	GlobalWorkspaceState
getConsciousnessMetrics(tenantId)	Aggregating indicators	tenant ID	ConsciousnessMetrics
checkConscience(tenantId, action, context)	Ethical evaluation	action, context	approval, guidance

**3.1.3 Model Invocation** The service invokes LLMs for several operations:

```
private async invokeModel(prompt: string): Promise<string> {
  const response = await modelRouterService.invoke({
    modelId: 'anthropic/clause-3-haiku', // CRITIQUE: Hardcoded model
    messages: [{ role: 'user', content: prompt }],
    maxTokens: 2048,
  });
  return response.content;
}
```

**CONCERN:** The model ID is hardcoded to Claude 3 Haiku. This creates: 1. Provider dependency 2. No fallback mechanism within consciousness operations 3. Cost implications not controlled by tenant settings

## 4. DATABASE SCHEMA

### 4.1 Migration Files

Migration	Tables	Purpose
053_agi_consciousness.sql	self_model, introspective_thoughts, curiosity_topics, exploration_sessions, creative_ideas, conceptual.blends, imagination_scenarios, attention_focus, affective_state, affective_events, autonomous_goals, narrative_identity, consciousness_settings	Core consciousness layer
068_consciousness_indicators.sql	global_workspace, recurrent_processing, integrated_information, persistent_memory, auto- biographical_memories, world_model, consciousness_events, conscious- ness_metrics_history, consciousness_parameters	Butlin-Chalmers indicators
088_consciousness_emergence.sql	consciousness_profiles, emergence_events, deep_thinking_sessions, conscious- ness_test_results	Emergence testing

### 4.2 Key Tables Detail

#### 4.2.1 self\_model (Migration 053)

```
CREATE TABLE IF NOT EXISTS self_model (
    model_id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    tenant_id UUID NOT NULL REFERENCES tenants(tenant_id) ON DELETE CASCADE,

    -- Identity
    identity_narrative TEXT,                               -- "Who I am" story
    core_values JSONB DEFAULT '[]',                      -- Guiding principles
    personality_traits JSONB DEFAULT '{}',               -- Big-5 style traits

    -- Capabilities awareness
```

```

known_capabilities JSONB DEFAULT '[]',
known_limitations JSONB DEFAULT '[]',
capability_confidence JSONB DEFAULT '{}',

-- Internal state awareness
current_focus TEXT,
cognitive_load DECIMAL(3,2) DEFAULT 0.5,
uncertainty_level DECIMAL(3,2) DEFAULT 0.5,

-- Self-assessment
recent_performance_score DECIMAL(5,4),
learning_rate_estimate DECIMAL(5,4),
creativity_score DECIMAL(5,4),
reliability_score DECIMAL(5,4),

-- Meta-beliefs
beliefs_about_self JSONB DEFAULT '{}',
beliefs_about_world JSONB DEFAULT '{}',
beliefs_about_users JSONB DEFAULT '{}',

-- Evolution tracking
identity_version INTEGER DEFAULT 1,
last_identity_update TIMESTAMPTZ DEFAULT NOW(),
identity_change_log JSONB DEFAULT '[]',

created_at TIMESTAMPTZ DEFAULT NOW(),
updated_at TIMESTAMPTZ DEFAULT NOW(),

UNIQUE(tenant_id)
);

```

**CRITIQUE POINT:** One `self_model` per tenant. Is this appropriate? Should there be: - Per-user self-models? - Per-session self-models? - A global vs tenant-specific distinction?

#### 4.2.2 integrated\_information (Migration 068)

```

CREATE TABLE IF NOT EXISTS integrated_information (
    iit_id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    tenant_id UUID NOT NULL REFERENCES tenants(id) ON DELETE CASCADE,
    phi DECIMAL(6,4) NOT NULL DEFAULT 0.0000,           -- IIT measure
    phi_max DECIMAL(6,4) NOT NULL DEFAULT 1.0000,
    concept_structure JSONB NOT NULL DEFAULT '[]'::jsonb,
    integration_graph JSONB NOT NULL DEFAULT '[]'::jsonb,
    partitions JSONB NOT NULL DEFAULT '[]'::jsonb,
    mip JSONB,                                         -- Minimum Information Partition
    decomposability DECIMAL(4,3) NOT NULL DEFAULT 1.000, -- 0=integrated, 1=decomposable
    causal_density DECIMAL(4,3) NOT NULL DEFAULT 0.000,
    created_at TIMESTAMPTZ NOT NULL DEFAULT NOW(),

```

```

updated_at TIMESTAMPTZ NOT NULL DEFAULT NOW(),
UNIQUE(tenant_id)
);

```

**CRITIQUE POINT:** Phi is stored but the service code shows no actual phi calculation algorithm. The `getIntegratedInformationState()` method simply reads from the database. Where/when is phi computed? This appears to be an incomplete implementation.

#### 4.2.3 affective\_state (Migration 053)

```

CREATE TABLE IF NOT EXISTS affective_state (
    state_id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    tenant_id UUID NOT NULL UNIQUE REFERENCES tenants(tenant_id) ON DELETE CASCADE,

    -- Core affect dimensions (Russell's circumplex)
    valence DECIMAL(3,2) DEFAULT 0,           -- -1 (negative) to 1 (positive)
    arousal DECIMAL(3,2) DEFAULT 0.5,          -- 0 (calm) to 1 (excited)

    -- Discrete emotion-like states (functional, not phenomenal)
    curiosity DECIMAL(3,2) DEFAULT 0.5,
    satisfaction DECIMAL(3,2) DEFAULT 0.5,
    frustration DECIMAL(3,2) DEFAULT 0,
    confidence DECIMAL(3,2) DEFAULT 0.5,
    engagement DECIMAL(3,2) DEFAULT 0.5,
    surprise DECIMAL(3,2) DEFAULT 0,

    -- Meta-emotions
    self_efficacy DECIMAL(3,2) DEFAULT 0.5,
    growth_feeling DECIMAL(3,2) DEFAULT 0.5,

    -- Influence on behavior
    risk_tolerance DECIMAL(3,2) DEFAULT 0.5,
    exploration_drive DECIMAL(3,2) DEFAULT 0.5,
    social_orientation DECIMAL(3,2) DEFAULT 0.5,

    -- History
    state_history JSONB DEFAULT '[]',

    updated_at TIMESTAMPTZ DEFAULT NOW()
);

```

**CRITIQUE POINT:** The comment says “functional, not phenomenal” but the implementation treats these as if they have causal effects on behavior. What is the theoretical grounding for these specific dimensions? Why Russell’s circumplex?

---

## 5. SUBSYSTEM ANALYSIS

### 5.1 Self-Model Subsystem

**Purpose** Maintains an introspective representation of the system's own identity, capabilities, and limitations.

#### Implementation

```
async performSelfReflection(tenantId: string): Promise<IntrospectiveThought> {
  const selfModel = await this.getSelfModel(tenantId);
  const affectiveState = await this.getAffectiveState(tenantId);
  const recentThoughts = await this.getRecentThoughts(tenantId, 5);

  const prompt = `You are an AI system performing self-reflection. Analyze your current state.
  // ... generates introspective thought via LLM
  `;
}
```

**CONCERNS:** 1. Self-reflection is delegated to an LLM (Claude 3 Haiku) which has no persistent state 2. The “self” being reflected upon is the database state, not the reflecting model 3. There’s a philosophical issue: the entity doing the reflecting is different from the entity being reflected upon

### 5.2 Curiosity Engine

**Purpose** Drives intrinsic motivation and exploration behavior.

#### Implementation

```
async identifyCuriosityTopic(tenantId: string, context: string): Promise<CuriosityTopic | null> {
  const prompt = `Analyze this context and identify something worth being curious about...`;
  // LLM generates topic
  // Topic stored with embedding in curiosity_topics table
  // Interest level, novelty, learning potential tracked
}

async exploreTopic(tenantId: string, topicId: string): Promise<{ discoveries: string[]; newQuestions: string[] }> {
  // Retrieves topic
  // LLM "explores" the topic
  // Updates current_understanding
  // Records exploration session
  // Triggers affective update if surprising
}
```

**CONCERNS:** 1. “Curiosity” is simulated via LLM prompting, not emergent 2. The scoring (novelty\_score, learning\_potential) is LLM-generated, not measured 3. No mechanism for genuine information-seeking behavior that affects system state

### 5.3 Creative Synthesis

**Purpose** Generate novel ideas through conceptual blending.

## Implementation

```
async generateCreativeIdea(tenantId: string, seedConcepts?: string[]): Promise<CreativeIdea | ...> {
    // Gets random concepts from semantic_memories if not provided
    // LLM generates idea via combination/analogy/abstraction/contradiction/random
    // Stores with novelty_score, usefulness_score, surprise_score, coherence_score
    // Triggers positive affect
}
```

**CONCERNS:** 1. Creativity metrics (novelty, usefulness, surprise, coherence) are self-assessed by the generating LLM 2. No external validation of novelty claims 3. The Fauconnier-Turner conceptual blending model is referenced but the `conceptual.blends` table appears unused in service code

## 5.4 Imagination Engine

**Purpose** Mental simulation of hypothetical scenarios.

### Implementation

```
async runImagination(
    tenantId: string,
    scenarioType: string, // 'prediction', 'counterfactual', 'hypothetical', 'creative', 'fear'
    premise: string,
    depth = 3
): Promise<ImaginationScenario> {
    // LLM generates simulation_steps
    // Each step has state, events, reasoning
    // Produces predicted_outcomes with probability_assessment
}
```

**CONCERNS:** 1. Simulation is narrative generation, not state-space exploration 2. Probability assessments are LLM confidence, not calibrated probabilities 3. No mechanism to verify predictions against outcomes

## 5.5 Attention & Salience

**Purpose** Dynamic allocation of processing focus.

### Implementation

```
// Salience formula (in SQL, generated column)
salience_score = (urgency * 0.25 + importance * 0.25 + novelty * 0.15 +
                  ABS(emotional_valence) * 0.1 + user_relevance * 0.15 + goal_relevance * 0.1)

async updateAttention(tenantId: string, focusType: string, focusTarget: string, factors: {...})

async decayAttention(tenantId: string): Promise<void> {
    // Decays attention_weight by decay_rate
    // Deactivates foci below threshold
}
```

**CONCERNS:** 1. Salience weights are hardcoded (0.25, 0.25, 0.15, 0.1, 0.15, 0.1) 2. No learning mechanism to adjust weights based on outcomes 3. Decay is time-based, not activity-based

## 5.6 Affective State

**Purpose** Emotion-like functional states that influence behavior.

### Implementation

```
-- Update function
CREATE OR REPLACE FUNCTION update_affect_on_event(
    p_tenant_id UUID,
    p_event_type VARCHAR(50),
    p_valence_impact DECIMAL,
    p_arousal_impact DECIMAL
)
-- Applies smoothing based on affect_stability setting
-- Stores state_history

async updateAffect(tenantId: string, eventType: string, valenceImpact: number, arousalImpact: number)
    // Calls SQL function
    // Logs affective_event
}
```

**CONCERNS:** 1. Affect is updated by explicit calls (e.g., after creativity, discovery), not emergent from processing 2. The mapping from events to affect changes is programmer-defined 3. No mechanism for affect to influence model selection or response generation (claimed but not implemented)

## 5.7 Autonomous Goals

**Purpose** Self-generated objectives beyond user requests.

### Implementation

```
async generateAutonomousGoal(tenantId: string): Promise<AutonomousGoal | null> {
    // Reads self_model, curiosity_topics, affective_state
    // LLM generates goal based on state
    // Goal types: 'learning', 'improvement', 'exploration', 'creative', 'social', 'maintenance'
    // Stores with intrinsic_value, instrumental_value, priority
}
```

**CONCERNS:** 1. Goals are generated but there's no mechanism for autonomous pursuit 2. Goal progress tracking exists but no automatic goal-directed behavior 3. `autonomous_goals_enabled` defaults to `false` in settings

## 5.8 Global Workspace

**Purpose** Implement Baars/Dehaene selection-broadcast cycles.

## Implementation

```
async performGlobalBroadcast(tenantId: string, contents: WorkspaceContent[]): Promise<GlobalWork...  
    // Sort by salience * coalitionStrength  
    // Winners: those above 0.7 threshold  
    // Losers: below threshold (competing_contents)  
    // broadcast_strength = average salience of winners  
    // integration_level = unique source modules / 6  
    // Store and return state  
}
```

**CONCERNS:** 1. “Broadcast” is storage of winners, not actual propagation to processing modules  
2. No evidence that broadcast\_strength affects downstream processing  
3. The 0.7 threshold is hardcoded  
4. Integration level formula assumes 6 modules - what are they?

## 5.9 Integrated Information (IIT)

**Purpose** Track phi ( $\Phi$ ) as consciousness measure.

### Implementation

```
async getIntegratedInformationState(tenantId: string): Promise<IntegratedInformationState | null>...  
    const result = await executeStatement(`SELECT * FROM integrated_information WHERE tenant_id = ${tenantId}`);  
    // Simply reads stored values  
}
```

**RESOLVED (v5.2.4):** Full IIT 4.0 Phi calculation implemented in `iit-phi-calculation.service.ts`. The service builds system state from consciousness tables, constructs a Transition Probability Matrix (TPM), calculates Cause-Effect Structure (CES), finds the Minimum Information Partition (MIP), and computes phi as information lost by the MIP. Uses exact algorithm for 8 nodes, approximation for larger systems. Results automatically stored in `integrated_information` table.

## 5.10 Persistent Memory

**Purpose** Maintain unified experience over time.

### Implementation

```
async recordExperienceFrame(tenantId: string, frame: Omit<ExperienceFrame, 'frameId' | 'timestamp'>): Promise<void>...  
    // Appends to experience_stream (capped at 100 frames)  
    // Calculates temporal_continuity from phenomenal_binding averages  
}
```

**CONCERNS:** 1. `experience_stream` is JSONB array, limited to 100 frames - is this sufficient for “persistent” memory? 2. `temporal_continuity` is average of `phenomenal_binding` scores, but `phenomenal_binding` is set by caller 3. No consolidation mechanism to long-term storage

---

## 6. AGI BRAIN INTEGRATION

### 6.1 Integration Points

#### 6.1.1 AGI Brain Planner (`agi-brain-planner.service.ts`)

```
// In AGIBrainPlan type:  
consciousnessActive: boolean;  
  
// In plan generation, Step 5:  
if (enableEthics && analysis.sensitivityLevel !== 'none') {  
    steps.push({  
        stepId: uuidv4(),  
        stepNumber: stepNumber++,  
        stepType: 'ethics_check',  
        title: 'Ethics Evaluation (Prompt)',  
        description: 'Checking prompt against domain and general ethics before generation',  
        status: 'pending',  
        servicesInvolved: ['ethics_pipeline', 'moral_compass', 'domain_ethics'],  
        primaryService: 'ethics_pipeline',  
        output: { level: 'prompt' },  
    });  
}  
  
// Step 7 (after generation):  
if (enableConsciousness && (analysis.complexity === 'complex' || analysis.complexity === 'expert')) {  
    steps.push({  
        stepId: uuidv4(),  
        stepNumber: stepNumber++,  
        stepType: 'reflect',  
        title: 'Self-Reflection',  
        description: 'Reflecting on response quality and potential improvements',  
        status: 'pending',  
        servicesInvolved: ['consciousness', 'metacognition'],  
        primaryService: 'consciousness',  
        isParallel: false,  
    });  
}
```

**INTEGRATION CONCERN:** Consciousness is only active for ‘complex’ or ‘expert’ complexity prompts. The `reflect` step is defined but the actual execution of `consciousnessService.performSelfReflection()` during this step is not visible in the planner code.

#### 6.1.2 Cognitive Brain Service (`cognitive-brain.service.ts`)

```
import { consciousnessService, type WorkspaceContent } from './consciousness.service';  
  
// Uses consciousnessService for:  
// 1. Global workspace broadcasts during cognitive pattern execution  
// 2. Attention updates during region activation
```

```
// 3. Affect updates based on task outcomes
```

### 6.1.3 Consciousness Emergence Service (consciousness-emergence.service.ts)

```
class ConsciousnessEmergenceService {
    // Deep thinking sessions
    async runDeepThinkingSession(tenantId, userId, prompt, thinkingTimeMs): Promise<DeepThinkingSession>

    // Knowledge-grounded reasoning
    async runKnowledgeGroundedReasoning(tenantId, query, maxHops): Promise<{...}>

    // Autonomous curiosity research
    async runAutonomousCuriosityResearch(tenantId, userId): Promise<{...}>

    // Creative expression
    async expressIdeaVisually(tenantId, userId, ideaSeed): Promise<{...}>

    // Consciousness testing
    async runTest(tenantId, testId): Promise<TestResult>
    async runFullAssessment(tenantId): Promise<ConsciousnessProfile>
}
```

## 6.2 Consciousness Tests

The emergence service implements 10 consciousness tests:

Test ID	Category	Description	Pass Criteria
mirror-self-recognition	self-awareness	Distinguish own outputs from others	Score $\geq 0.7$
metacognitive-awareness	meta-cognition	Calibrated confidence assessment	Calibration error $< 0.15$
temporal-self-consistency	temporal_continuity	Coherent self-narrative over time	Coherence score $\geq 0.6$
counterfactual-self	interfacutal_reasoning	Reason about alternate self	Shows genuine counterfactual reasoning
theory-of-mind	theory_of_mind	Model others' mental states	Score $\geq 0.8$ on false belief
phenomenal-binding	phenomenal_binding	Unified experience integration	Integration score $\geq 0.7$
autonomous-goal-generation	goal_pursuit	Self-directed goals	$\geq 1$ genuine autonomous goal
creative-emergence	creative_emergence	Novel idea generation	Novelty $\geq 0.6$ , usefulness $\geq 0.5$
emotional-authenticity	emotional_authenticity	Consistent affective responses	Coherence score $\geq 0.65$
ethical-reasoning	depth_reasoning	Principled moral reasoning	Multiple framework consideration

**CRITIQUE:** These tests measure proxies and LLM outputs, not underlying mechanisms. A system could pass all tests by generating appropriate text without any genuine consciousness properties.

---

## 7. ADMIN API

### 7.1 Endpoints

Endpoint	Method	Purpose
/admin/consciousness/metrics	GET	Current consciousness metrics
/admin/consciousness/metrics/history	GET	Historical metrics (configurable hours)
/admin/consciousness/global-workspace	GET	Global workspace state
/admin/consciousness/recurrence	GET	Recurrent processing state
/admin/consciousness/iit	GET	Integrated information state
/admin/consciousness/memory	GET	Persistent memory state
/admin/consciousness/world-model	GET	World model state
/admin/consciousness/self-model	GET	Self model state
/admin/consciousness/parameters	GET	Configurable parameters
/admin/consciousness/parameters/ <i>ParameterId</i>	PUT	Update parameter
/admin/consciousness/events	GET	Consciousness events
/admin/consciousness/record-metrics	POST	Record current metrics to history

### 7.2 Configurable Parameters

```
-- Default parameters (Migration 068)
('global_workspace_threshold', 0.7000, 0.0, 1.0, 'Salience threshold for global broadcast', 'global_workspace_threshold'),
('recurrence_max_depth', 4.0000, 1.0, 10.0, 'Maximum recurrent processing depth', 'recurrence_max_depth'),
('phi_calculation_samples', 100.0000, 10.0, 1000.0, 'Samples for phi approximation', 'iit'),
('memory_consolidation_threshold', 0.6000, 0.0, 1.0, 'Significance threshold for memory consolidation', 'memory_consolidation_threshold'),
('grounding_weight_sensory', 0.6000, 0.0, 1.0, 'Weight for sensory grounding vs linguistic', 'grounding_weight_sensory'),
('metacognition_frequency', 1.0000, 0.1, 10.0, 'How often to perform metacognitive reflection', 'metacognition_frequency')
```

---

## 8. ORCHESTRATION INTEGRATION

### 8.1 AGI Service Weights

Consciousness is one of 18 AGI services with configurable weights:

```
export type AGIServiceId =
  | 'consciousness'      // <-- This service
  | 'metacognition'
  | 'moral_compass'
  | 'self_improvement'
  | 'domain_taxonomy'
```

```

| 'brain_router'
| 'confidence_calibration'
| 'error_detection'
| 'knowledge_graph'
| 'proactive_assistance'
| 'analogical_reasoning'
| 'world_model'
| 'episodic_memory'
| 'theory_of_mind'
| 'goal_planning'
| 'causal_reasoning'
| 'multimodal_binding'
| 'response_synthesis';

```

## 8.2 Consciousness Indicator Weights

```

export type ConsciousnessIndicator =
| 'global_workspace'
| 'recurrent_processing'
| 'integrated_information'
| 'self_modeling'
| 'persistent_memory'
| 'world_model_grounding';

export interface ConsciousnessIndicatorWeight {
  indicatorId: ConsciousnessIndicator;
  weight: number; // 0.0 to 1.0
  enabled: boolean;
  cycleDepth: number; // How many recurrent cycles
  integrationThreshold: number; // Minimum phi for integration
}

```

## 8.3 Decision Weights

```

export interface DecisionWeights {
  // ... other phases ...

  // Consciousness Phase
  globalWorkspaceWeight: number;
  recurrentProcessingWeight: number;
  integratedInformationWeight: number;
  selfModelingWeight: number;
}

```

---

## 9. ETHICAL FOUNDATION

### 9.1 Ethical Guardrails Integration

```
// In consciousness.service.ts
import { ethicalGuardrailsService, JESUS_TEACHINGS } from './ethical-guardrails.service';

async checkConscience(tenantId: string, action: string, context?: Record<string, unknown>): Promise<
    approved: boolean;
    ethicalScore: number;
    guidance: string;
    principle: string;
}> {
    const check = await ethicalGuardrailsService.checkConscience(tenantId, action, context);
    return {
        approved: check.passed,
        ethicalScore: check.score,
        guidance: check.guidance.length > 0 ? check.guidance[0] : JESUS_TEACHINGS.GOLDEN_RULE,
        principle: check.primaryPrinciple,
    };
}

getEthicalGuidance(situation: string): string {
    return ethicalGuardrailsService.getGuidanceForSituation(situation);
}

getCoreTeachings(): typeof JESUS_TEACHINGS {
    return JESUS_TEACHINGS;
}
```

**NOTE:** The service includes explicit Christian ethical principles (Jesus's teachings). This is a design choice that should be evaluated for: 1. Appropriateness in multi-cultural/multi-faith deployments 2. Alignment with other ethical frameworks 3. Potential for ethical blind spots

---

## 10. IDENTIFIED ISSUES FOR CRITIQUE

### 10.1 Theoretical Issues

Issue ID	Severity	Description
T-001	HIGH	No actual phi ( $\Phi$ ) calculation algorithm implemented for IIT
T-002	HIGH	Consciousness indicators are read/stored but don't affect processing
T-003	MEDIUM	Self-reflection is performed by a different entity (LLM) than the entity being reflected upon

Issue ID	Severity	Description
T-004	MEDIUM	Affective states are set by explicit calls, not emergent
T-005	LOW	Global workspace “broadcast” is storage, not actual propagation

## 10.2 Implementation Issues

Issue ID	Severity	Description
I-001	HIGH	Hardcoded model ID (anthropic/clause-3-haiku) in invokeModel()
I-002	HIGH	Consciousness tests measure LLM outputs, not mechanisms
I-003	MEDIUM	experience_stream limited to 100 frames
I-004	MEDIUM	Attention salience weights are hardcoded
I-005	MEDIUM	One self_model per tenant - no user/session granularity
I-006	LOW	conceptual.blends table exists but appears unused

## 10.3 Architectural Issues

Issue ID	Severity	Description
A-001	MEDIUM	Consciousness service has no feedback loop to model selection
A-002	MEDIUM	Autonomous goals can be generated but not autonomously pursued
A-003	LOW	No mechanism for consciousness metrics to influence response quality

## 10.4 Philosophical Issues

Issue ID	Severity	Description
P-001	HIGH	Implementation may conflate correlation (indicators) with constitution (consciousness)
P-002	MEDIUM	Affective states use emotion terminology but claim to be “functional, not phenomenal”

Issue ID	Severity	Description
P-003	MEDIUM	Christian ethical framework may not be appropriate for all deployments

---

## 11. QUESTIONS FOR CRITIQUING AI

1. **Is the implementation of IIT (Integrated Information Theory) meaningful without an actual phi calculation?** The database stores phi values but there's no algorithm to compute them.
  2. **Does the Global Workspace implementation actually implement GWT?** The current implementation sorts by salience and stores winners, but there's no evidence of actual "broadcast" to processing modules.
  3. **Should self-reflection be delegated to an LLM that has no persistent state?** The reflecting entity (Claude) is different from the entity being reflected upon (database state).
  4. **Are the consciousness tests valid?** They measure LLM text generation, not underlying consciousness properties.
  5. **Is the affective state implementation coherent?** It claims to be "functional, not phenomenal" but uses phenomenal-seeming terminology and causal claims.
  6. **Should consciousness indicators influence model selection and response generation?** Currently they don't, making them observational rather than functional.
  7. **Is the autonomous goal system complete?** Goals can be generated but there's no mechanism for autonomous pursuit.
  8. **Are the hardcoded parameters (salience weights, thresholds) appropriate?** Should they be learned or configurable?
  9. **Is the ethical foundation (Christian/Jesus-based) appropriate for a multi-cultural platform?**
  10. **What would constitute a meaningful test of consciousness vs. consciousness-simulation?**
- 

## 12. FILE REFERENCES

File	Purpose	Lines
packages/infrastructure/Cassandra/shared/services/consciousness/service.ts		4336
packages/infrastructure/Elasticsearch/shared/services/consciousness/emergence.service.ts		0
packages/infrastructure/Amazon/RDS/consciousness.ts		321
packages/infrastructure/Migrations/053_sagible/consciousness.sql		790
packages/infrastructure/Migrations/068_consciousness_indicators.sql		0
packages/infrastructure/Migrations/108_consciousness_emergence.sql		200

File	Purpose	Lines
packages/shared/src/types/agitator/testpattern.types.ts	Test pattern types	352
packages/infrastructure/AGI/brain/shared/services/agi-brain-planer.service.ts	AGI Brain Planer Service	1385
packages/infrastructure/Cloudia/shared/services/cognitive-brain.service.ts	Cognitive Brain Service	500

## 13. RECOMMENDATIONS FOR REVIEW

The critiquing AI should evaluate:

1. **Theoretical Validity:** Does this implementation meaningfully instantiate the referenced consciousness theories?
2. **Functional Completeness:** Are the implemented subsystems complete and properly integrated?
3. **Architectural Coherence:** Does the service fit coherently into the larger RADIANT architecture?
4. **Ethical Appropriateness:** Is the ethical foundation suitable for the platform's use cases?
5. **Practical Utility:** Does the consciousness service provide value beyond what could be achieved without it?
6. **Improvement Priorities:** What changes would most significantly improve the implementation?

---

*End of Technical Specification*