

Contents

RADIANT v5.0.2 - Source Export Part 3: Python Flyte Tasks	1
6. Grimoire Flyte Tasks	1

RADIANT v5.0.2 - Source Export Part 3: Python Flyte Tasks

6. Grimoire Flyte Tasks

File: packages/flyte/workflows/grimoire_tasks.py

Purpose: Flyte workflow tasks for The Grimoire procedural memory system. Implements semantic search for heuristic retrieval and automated learning from successful executions.

Key Tasks: - consult_grimoire - Query relevant heuristics before task execution - librarian_review - Extract lessons after successful execution - cleanup_expired_heuristics - Maintenance task for stale data - Admin functions for manual management

"""

The Grimoire - Procedural Memory Flyte Tasks

RADIANT v5.0.2 - System Evolution

Flyte tasks for managing The Grimoire's self-optimizing procedural memory.

These tasks integrate with the Think Tank workflow to provide:

1. Pre-execution heuristic consultation
2. Post-execution learning (Librarian Review)
3. Scheduled maintenance

"""

```
import json
import os
import re
import uuid
from datetime import datetime, timedelta
from typing import Dict, List, Optional, Tuple
from dataclasses import dataclass

from flytekit import task, workflow, current_context
import httpx

# Local imports
from radiant.flyte.utils.db import get_safe_db_connection, get_system_db_connection
from radiant.flyte.utils.embeddings import generate_embedding, cosine_similarity
from radiant.flyte.utils.cato_client import CatoClient
```

```

# Configuration
LITELLM_PROXY_URL = os.environ.get("LITELLM_PROXY_URL", "http://localhost:4000")
LITELLM_API_KEY = os.environ.get("LITELLM_API_KEY", "")
MAX_HEURISTICS_PER_QUERY = 5
SIMILARITY_THRESHOLD = 0.7
DEFAULT_CONFIDENCE = 0.5

@dataclass
class HeuristicMatch:
    """A matched heuristic from The Grimoire"""
    id: str
    domain: str
    heuristic_text: str
    confidence_score: float
    similarity_score: float

@dataclass
class GrimoireConsultResult:
    """Result of consulting The Grimoire"""
    heuristics: List[HeuristicMatch]
    context_injection: str
    consultation_id: str

@task(
    cache=False,
    retries=2,
    timeout=timedelta(seconds=30)
)
def consult_grimoire(
    tenant_id: str,
    prompt: str,
    domain: str = "general",
    max_results: int = MAX_HEURISTICS_PER_QUERY
) -> GrimoireConsultResult:
    """
    Consults The Grimoire for relevant heuristics before task execution.
    """

    This task:
    1. Generates an embedding for the current prompt
    2. Performs vector similarity search against stored heuristics
    3. Validates results through Cato safety checks
    4. Returns a formatted context injection for the agent

```

Args:

tenant_id: UUID of the tenant

```


prompt: The user's prompt/task



domain: Task domain for filtering



max_results: Maximum heuristics to return



Returns:



GrimoireConsultResult with matched heuristics and context injection



"""



consultation_id = str(uuid.uuid4())



# Generate embedding for the prompt



try:



prompt_embedding = generate_embedding(prompt[:8000])



except Exception as e:



print(f"Embedding generation failed: {e}")



return GrimoireConsultResult(



heuristics=[],



context_injection="",



consultation_id=consultation_id



)



# Query The Grimoire



with get_safe_db_connection(tenant_id) as (conn, cur):



cur.execute("""



SELECT



id, domain, heuristic_text, confidence_score,



1 - (context_embedding <=> %s::vector) as similarity



FROM knowledge_heuristics



WHERE expires_at > NOW()



AND (domain = %s OR domain = 'general')



AND context_embedding IS NOT NULL



ORDER BY context_embedding <=> %s::vector



LIMIT %s



""", (prompt_embedding, domain, prompt_embedding, max_results * 2))



rows = cur.fetchall()



# Filter by similarity threshold and validate with Cato



heuristics = []



for row in rows:



id_, dom, text, confidence, similarity = row



if similarity < SIMILARITY_THRESHOLD:



continue



# Safety check via Cato



risk = CatoClient.epistemic_check(text, tenant_id)



if risk.should_block:



print(f"Heuristic {id_} blocked by Cato: {risk.reason}")


```

```

        continue

    heuristics.append(HeuristicMatch(
        id=str(id_),
        domain=dom,
        heuristic_text=text,
        confidence_score=confidence,
        similarity_score=similarity
    ))

    if len(heuristics) >= max_results:
        break

    # Build context injection
    context_injection = _build_context_injection(heuristics)

    return GrimoireConsultResult(
        heuristics=heuristics,
        context_injection=context_injection,
        consultation_id=consultation_id
    )

def _build_context_injection(heuristics: List[HeuristicMatch]) -> str:
    """Formats heuristics into a system prompt injection."""
    if not heuristics:
        return ""

    lines = ["<grimoire_heuristics>"]
    lines.append("The following lessons from previous successful executions may be relevant:")
    lines.append("")

    for i, h in enumerate(heuristics, 1):
        confidence_pct = int(h.confidence_score * 100)
        lines.append(f"{i}. [{h.domain.upper()}] (Confidence: {confidence_pct}%)")
        lines.append(f"    {h.heuristic_text}")
        lines.append("")

    lines.append("Apply these heuristics where appropriate, but use your judgment.")
    lines.append("</grimoire_heuristics>")

    return "\n".join(lines)

@task(
    cache=False,
    retries=2,
    timeout=timedelta(seconds=60)
)

```

```

)
def librarian_review(
    tenant_id: str,
    execution_id: str,
    prompt: str,
    response: str,
    domain: str,
    was_successful: bool,
    user_rating: Optional[int] = None
) -> Dict[str, any]:
    """

```

The Librarian reviews successful executions and extracts reusable heuristics.

This task:

1. Only processes successful executions
2. Uses an LLM to extract generalizable lessons
3. Validates extracted heuristics via Cato
4. Stores new heuristics with embeddings

Args:

```

tenant_id: UUID of the tenant
execution_id: ID of the execution to review
prompt: Original user prompt
response: AI response that was successful
domain: Task domain
was_successful: Whether the task was successful
user_rating: Optional 1-5 rating from user

```

Returns:

```

Dict with extracted_count and heuristic_ids
"""
# Only learn from successes
if not was_successful:
    return {"extracted_count": 0, "heuristic_ids": [], "reason": "Not successful"}

# Only learn from highly-rated responses
if user_rating is not None and user_rating < 4:
    return {"extracted_count": 0, "heuristic_ids": [], "reason": "Low rating"}

# Extract heuristics using LLM
extraction_prompt = f"""Analyze this successful AI interaction and extract 1-3 reusable heuristics.

```

DOMAIN: {domain}

USER PROMPT:
{prompt[:2000]}

SUCCESSFUL RESPONSE:

```
{response[:3000]}
```

Extract generalizable lessons in this format:

- "When [specific condition], always [specific action]"

Rules:

1. Only extract truly generalizable lessons
2. Be specific enough to be actionable
3. Avoid domain-specific jargon unless necessary
4. Each heuristic should be a single sentence

Return ONLY a JSON array of strings, e.g.:

```
["When X, always Y", "When A, always B"]
```

If no generalizable lessons can be extracted, return: []"""

```
try:  
    with httpx.Client(timeout=30.0) as client:  
        resp = client.post(  
            f"{LITELLM_PROXY_URL}/chat/completions",  
            headers={  
                "Authorization": f"Bearer {LITELLM_API_KEY}",  
                "Content-Type": "application/json"  
            },  
            json={  
                "model": "gpt-4o-mini",  
                "messages": [{"role": "user", "content": extraction_prompt}],  
                "temperature": 0.3,  
                "max_tokens": 500  
            }  
        )  
        resp.raise_for_status()  
  
        content = resp.json()["choices"][0]["message"]["content"]  
  
        # Parse JSON from response  
        json_match = re.search(r'\[.*\]', content, re.DOTALL)  
        if not json_match:  
            return {"extracted_count": 0, "heuristic_ids": [], "reason": "No JSON found"}  
  
        heuristics = json.loads(json_match.group())  
  
    except Exception as e:  
        print(f"Heuristic extraction failed: {e}")  
        return {"extracted_count": 0, "heuristic_ids": [], "reason": str(e)}  
  
    if not heuristics:  
        return {"extracted_count": 0, "heuristic_ids": [], "reason": "No heuristics extracted"}
```

```

# Validate and store heuristics
stored_ids = []

for heuristic_text in heuristics[:3]:
    if not isinstance(heuristic_text, str) or len(heuristic_text) < 20:
        continue

    # Safety check
    risk = CatoClient.validate_for_storage(heuristic_text, tenant_id)
    if risk.should_block:
        print(f"Heuristic blocked by Cato: {risk.reason}")
        continue

    # Generate embedding
    try:
        embedding = generate_embedding(heuristic_text)
    except Exception as e:
        print(f"Embedding failed for heuristic: {e}")
        continue

    # Check for duplicates
    with get_safe_db_connection(tenant_id) as (conn, cur):
        cur.execute("""
            SELECT id, heuristic_text, confidence_score,
                   1 - (context_embedding <=> %s::vector) as similarity
            FROM knowledge_heuristics
            WHERE domain = %s
                AND expires_at > NOW()
                AND context_embedding IS NOT NULL
            ORDER BY context_embedding <=> %s::vector
            LIMIT 1
        """, (embedding, domain, embedding))

        existing = cur.fetchone()

        if existing and existing[3] > 0.95:
            # Very similar heuristic exists - reinforce it
            cur.execute("""
                UPDATE knowledge_heuristics
                SET confidence_score = LEAST(1.0, confidence_score + 0.05),
                    expires_at = expires_at + INTERVAL '30 days',
                    updated_at = NOW()
                WHERE id = %s
            """, (existing[0],))
            conn.commit()
            stored_ids.append(str(existing[0]))
            continue

```

```

# Store new heuristic
heuristic_id = str(uuid.uuid4())
cur.execute("""
    INSERT INTO knowledge_heuristics (
        id, tenant_id, domain, heuristic_text,
        context_embedding, confidence_score, source_execution_id
    ) VALUES (%s, %s, %s, %s, %s, %s, %s)
""", (
    heuristic_id, tenant_id, domain, heuristic_text,
    embedding, DEFAULT_CONFIDENCE, execution_id
))
conn.commit()
stored_ids.append(heuristic_id)

return {
    "extracted_count": len(stored_ids),
    "heuristic_ids": stored_ids,
    "reason": "Success"
}

```

```

@task(
    cache=False,
    retries=1,
    timeout=timedelta(minutes=5)
)
def cleanup_expired_heuristics(batch_size: int = 1000) -> Dict[str, int]:
    """
    Scheduled maintenance task to clean up The Grimoire.

```

Returns:

1. *Expired heuristics (past expires_at)*
2. *Low-confidence heuristics older than 30 days*

Returns:

Dict with expired_count and low_confidence_count

```

with get_system_db_connection() as (conn, cur):
    # Delete expired
    cur.execute("""
        DELETE FROM knowledge_heuristics
        WHERE expires_at < NOW()
        RETURNING id
    """)
    expired_count = len(cur.fetchall())

```

Delete low-confidence stale

```

        cur.execute("""
            DELETE FROM knowledge_heuristics
            WHERE confidence_score < 0.3
                AND created_at < NOW() - INTERVAL '30 days'
            RETURNING id
        """)
        low_confidence_count = len(cur.fetchall())

        conn.commit()

    return {
        "expired_count": expired_count,
        "low_confidence_count": low_confidence_count,
        "total_deleted": expired_count + low_confidence_count
    }

# =====
# ADMIN FUNCTIONS
# =====

@task(cache=False)
def list_heuristics(
    tenant_id: str,
    domain: Optional[str] = None,
    limit: int = 100
) -> List[Dict]:
    """Admin function to list heuristics."""
    with get_safe_db_connection(tenant_id) as (conn, cur):
        if domain:
            cur.execute("""
                SELECT id, domain, heuristic_text, confidence_score,
                    created_at, expires_at
                FROM knowledge_heuristics
                WHERE domain = %s AND expires_at > NOW()
                ORDER BY confidence_score DESC
                LIMIT %s
            """, (domain, limit))
        else:
            cur.execute("""
                SELECT id, domain, heuristic_text, confidence_score,
                    created_at, expires_at
                FROM knowledge_heuristics
                WHERE expires_at > NOW()
                ORDER BY confidence_score DESC
                LIMIT %s
            """, (limit,))

```

```

rows = cur.fetchall()

return [
{
    "id": str(row[0]),
    "domain": row[1],
    "heuristic_text": row[2],
    "confidence_score": row[3],
    "created_at": row[4].isoformat() if row[4] else None,
    "expires_at": row[5].isoformat() if row[5] else None
}
for row in rows
]

@task(cache=False)
def delete_heuristic(tenant_id: str, heuristic_id: str) -> bool:
    """Admin function to delete a heuristic."""
    with get_safe_db_connection(tenant_id) as (conn, cur):
        cur.execute(
            "DELETE FROM knowledge_heuristics WHERE id = %s RETURNING id",
            (heuristic_id,))
    deleted = cur.fetchone() is not None
    conn.commit()
    return deleted

@task(cache=False)
def add_manual_heuristic(
    tenant_id: str,
    domain: str,
    heuristic_text: str
) -> Optional[str]:
    """Admin function to manually add a heuristic."""
    # Safety check
    risk = CatoClient.validate_for_storage(heuristic_text, tenant_id)
    if risk.should_block:
        print(f"Manual heuristic blocked: {risk.reason}")
        return None

    # Generate embedding
    try:
        embedding = generate_embedding(heuristic_text)
    except Exception as e:
        print(f"Embedding generation failed: {e}")
        return None

```

```

heuristic_id = str(uuid.uuid4())

with get_safe_db_connection(tenant_id) as (conn, cur):
    cur.execute("""
        INSERT INTO knowledge_heuristics (
            id, tenant_id, domain, heuristic_text,
            context_embedding, confidence_score
        ) VALUES (%s, %s, %s, %s, %s, 0.7)
        RETURNING id
    """, (heuristic_id, tenant_id, domain, heuristic_text, embedding))
    conn.commit()

return heuristic_id

@task(cache=False)
def get_grimoire_stats(tenant_id: str) -> Dict:
    """Get Grimoire statistics for a tenant."""
    with get_safe_db_connection(tenant_id) as (conn, cur):
        cur.execute("""
            SELECT
                COUNT(*) as total,
                COUNT(*) FILTER (WHERE confidence_score >= 0.8) as high_confidence,
                COUNT(*) FILTER (WHERE expires_at < NOW() + INTERVAL '7 days') as expiring_soon,
                AVG(confidence_score) as avg_confidence
            FROM knowledge_heuristics
            WHERE expires_at > NOW()
        """)
        row = cur.fetchone()

        cur.execute("""
            SELECT domain, COUNT(*) as count
            FROM knowledge_heuristics
            WHERE expires_at > NOW()
            GROUP BY domain
        """)
        by_domain = {r[0]: r[1] for r in cur.fetchall()}

    return {
        "total_heuristics": row[0] or 0,
        "high_confidence": row[1] or 0,
        "expiring_soon": row[2] or 0,
        "avg_confidence": float(row[3]) if row[3] else 0.0,
        "by_domain": by_domain
    }

```

Continued in GRIMOIRE-GOVERNOR-SOURCE-PART4.md