

# Contents

<b>RADIANT Performance Optimization Guide</b>	<b>1</b>
Version: 5.42.0 . . . . .	1
1. Lambda Cold Start Optimization . . . . .	1
Current Configuration . . . . .	1
Recommendations . . . . .	2
2. Database Query Optimization . . . . .	2
2.1 Query Patterns . . . . .	2
2.2 Connection Pooling . . . . .	3
3. Caching Strategy . . . . .	3
3.1 In-Memory Caching (Lambda) . . . . .	3
3.2 Redis/ElastiCache (Production) . . . . .	3
3.3 API Gateway Caching . . . . .	4
4. Response Optimization . . . . .	4
4.1 Compression . . . . .	4
4.2 Pagination . . . . .	4
5. AI Model Call Optimization . . . . .	5
5.1 Streaming Responses . . . . .	5
5.2 Request Batching . . . . .	5
5.3 Prompt Caching (Claude) . . . . .	6
6. Frontend Performance . . . . .	6
6.1 Code Splitting . . . . .	6
6.2 Data Fetching . . . . .	6
6.3 Image Optimization . . . . .	6
7. Monitoring & Alerts . . . . .	7
7.1 Key Metrics . . . . .	7
7.2 CloudWatch Alarms . . . . .	7
8. Cost Optimization . . . . .	7
8.1 Right-Sizing . . . . .	7
8.2 Reserved Concurrency . . . . .	7
8.3 Spot Instances (ECS/Fargate) . . . . .	7
9. Quick Wins Checklist . . . . .	8

## RADIANT Performance Optimization Guide

**Version: 5.42.0**

This document outlines performance optimizations implemented and recommended for the RADIANT platform.

---

### 1. Lambda Cold Start Optimization

#### Current Configuration

Lambda	Memory	Timeout	Provisioned Concurrency
Admin API	1024 MB	30s	0 (on-demand)
Think Tank API	2048 MB	60s	0 (on-demand)
Agent Worker	2048 MB	300s	0 (configurable)
Transparency Worker	512 MB	30s	0

## Recommendations

### 1.1 Bundle Optimization

```
// Use esbuild for smaller bundles
// Current: ~5MB bundled
// Target: <2MB bundled

// In CDK stack:
bundling: {
  minify: true,
  sourceMap: false,
  treeshaking: true,
  externalModules: ['@aws-sdk/*'], // Use Lambda's built-in SDK
}
```

### 1.2 Lazy Loading

```
// Defer heavy imports until needed
let bedrockClient: BedrockRuntimeClient | null = null;

function getBedrockClient() {
  if (!bedrockClient) {
    bedrockClient = new BedrockRuntimeClient({});
  }
  return bedrockClient;
}
```

### 1.3 Connection Reuse

```
// Reuse HTTP connections
const httpAgent = new https.Agent({
  keepAlive: true,
  maxSockets: 50,
});
```

---

## 2. Database Query Optimization

### 2.1 Query Patterns

Use Indexes Effectively:

```
-- Ensure indexes exist for common queries
CREATE INDEX CONCURRENTLY IF NOT EXISTS idx_ai_reports_tenant_updated
ON ai_reports(tenant_id, updated_at DESC);

CREATE INDEX CONCURRENTLY IF NOT EXISTS idx_ai_report_insights_tenant_created
ON ai_report_insights(tenant_id, created_at DESC);
```

### Batch Operations:

```
// Instead of N individual inserts
for (const item of items) {
    await executeStatement('INSERT INTO...', [item]);
}

// Use batch insert
const values = items.map((_, i) => `(${i*3+1}, ${i*3+2}, ${i*3+3})`).join(',');
await executeStatement(`INSERT INTO table (a, b, c) VALUES ${values}`, flatParams);
```

## 2.2 Connection Pooling

Aurora Data API handles connection pooling automatically. Ensure: - `maxConnections` is appropriate for tier - Idle connections are released

---

## 3. Caching Strategy

### 3.1 In-Memory Caching (Lambda)

```
// Cache expensive computations during Lambda lifetime
const cache = new Map<string, { data: unknown; expiry: number }>();

function getCached<T>(key: string, ttlMs: number, compute: () => T): T {
    const now = Date.now();
    const cached = cache.get(key);

    if (cached && cached.expiry > now) {
        return cached.data as T;
    }

    const data = compute();
    cache.set(key, { data, expiry: now + ttlMs });
    return data;
}
```

### 3.2 Redis/ElastiCache (Production)

For high-traffic endpoints: - Model configurations: 5 min TTL - Tenant settings: 1 min TTL - User sessions: 15 min TTL

### 3.3 API Gateway Caching

```
// CDK configuration
const api = new apigateway.RestApi(this, 'Api', {
  deployOptions: {
    cachingEnabled: true,
    cacheClusterEnabled: true,
    cacheClusterSize: '0.5',
    cacheTtl: cdk.Duration.minutes(1),
  },
});

// Per-method caching
method.addMethodResponse({
  statusCode: '200',
  responseParameters: {
    'method.response.header.Cache-Control': true,
  },
});
```

---

## 4. Response Optimization

### 4.1 Compression

```
// Enable gzip for large responses
if (body.length > 1024) {
  const compressed = zlib.gzipSync(body);
  return {
    statusCode: 200,
    headers: {
      'Content-Encoding': 'gzip',
      'Content-Type': 'application/json',
    },
    body: compressed.toString('base64'),
    isBase64Encoded: true,
  };
}
```

### 4.2 Pagination

```
// Always paginate large result sets
const DEFAULT_PAGE_SIZE = 50;
const MAX_PAGE_SIZE = 200;

function paginate<T>(items: T[], page: number, pageSize: number): PaginatedResult<T> {
  const size = Math.min(pageSize || DEFAULT_PAGE_SIZE, MAX_PAGE_SIZE);
  const offset = (page - 1) * size;
```

```

    return {
      items: items.slice(offset, offset + size),
      total: items.length,
      page,
      pageSize: size,
      totalPages: Math.ceil(items.length / size),
    };
}

```

---

## 5. AI Model Call Optimization

### 5.1 Streaming Responses

```

// Use streaming for long generations
const stream = await bedrockClient.send(new InvokeModelWithResponseStreamCommand({
  modelId,
  body: JSON.stringify(request),
}));

for await (const event of stream.body) {
  if (event.chunk) {
    yield JSON.parse(new TextDecoder().decode(event.chunk.bytes));
  }
}

```

### 5.2 Request Batching

```

// Batch similar requests
const batchWindow = 50; // ms
const pendingRequests: Map<string, Promise<Response>> = new Map();

async function batchedInvoke(prompt: string): Promise<Response> {
  const key = hashPrompt(prompt);

  if (pendingRequests.has(key)) {
    return pendingRequests.get(key)!;
  }

  const promise = invoke(prompt);
  pendingRequests.set(key, promise);

  setTimeout(() => pendingRequests.delete(key), batchWindow);

  return promise;
}

```

### 5.3 Prompt Caching (Claude)

```
// Use prompt caching for repeated system prompts
const request = {
  anthropic_version: 'bedrock-2023-05-31',
  system: [
    {
      type: 'text',
      text: systemPrompt,
      cache_control: { type: 'ephemeral' }, // Enable caching
    },
  ],
  messages,
};
```

---

## 6. Frontend Performance

### 6.1 Code Splitting

```
// Dynamic imports for large components
const AIReportsPage = dynamic(() => import('./ai-reports'), {
  loading: () => <Skeleton />,
  ssr: false,
});
```

### 6.2 Data Fetching

```
// Use SWR for caching and revalidation
const { data, error, isLoading } = useSWR(
  '/api/admin/ai-reports',
  fetcher,
  {
    revalidateOnFocus: false,
    dedupingInterval: 30000,
  }
);
```

### 6.3 Image Optimization

```
// Use Next.js Image component
import Image from 'next/image';

<Image
  src={logoUrl}
  width={200}
  height={50}
  priority={false}
```

```
loading="lazy"  
/>>
```

---

## 7. Monitoring & Alerts

### 7.1 Key Metrics

Metric	Warning	Critical
P95 Latency	> 1s	> 3s
Error Rate	> 1%	> 5%
Cold Start Rate	> 10%	> 25%
Cache Hit Rate	< 80%	< 60%

---

### 7.2 CloudWatch Alarms

```
new cloudwatch.Alarm(this, 'HighLatencyAlarm', {  
  metric: api.metricLatency({ statistic: 'p95' }),  
  threshold: 1000,  
  evaluationPeriods: 3,  
  comparisonOperator: cloudwatch.ComparisonOperator.GREATER_THAN_THRESHOLD,  
});
```

---

## 8. Cost Optimization

### 8.1 Right-Sizing

- Start with minimum viable memory
- Use Lambda Power Tuning to find optimal
- Monitor actual memory usage

### 8.2 Reserved Concurrency

```
// For predictable workloads  
const fn = new lambda.Function(this, 'Function', {  
  // ...  
  reservedConcurrentExecutions: 100, // Limit max concurrency  
});
```

### 8.3 Spot Instances (ECS/Fargate)

For self-hosted models:

```
fargateService.taskDefinition.addContainer('model', {  
  // ...  
  capacityProviderStrategies: [  
    { capacityProvider: 'FARGATE_SPOT', weight: 2 },
```

```
{ capacityProvider: 'FARGATE', weight: 1 },  
],  
});
```

---

## 9. Quick Wins Checklist

- Enable API Gateway caching for read endpoints
- Add database indexes for frequent queries
- Implement response compression for large payloads
- Use streaming for AI model responses
- Enable HTTP keep-alive for external calls
- Lazy-load heavy SDK clients
- Paginate all list endpoints
- Add CloudWatch alarms for latency
- Review Lambda memory settings monthly
- Enable prompt caching for Claude models