

Contents

RADIANT Performance Guide	1
Overview	1
Architecture Performance	2
Request Flow	2
Latency Targets	2
Caching Strategy	2
Multi-Layer Caching	2
Cache Keys	2
Cache Invalidation	3
Database Optimization	3
Connection Pooling	3
Query Optimization	3
RLS Performance	4
Lambda Optimization	4
Cold Start Reduction	4
Memory Optimization	4
Bundling	4
Rate Limiting	5
Tier Limits	5
Implementation	5
Load Testing	5
Running Tests	5
Performance Baselines	5
Scaling	6
Horizontal Scaling	6
Vertical Scaling	6
Monitoring	6
Key Metrics	6
Alerting Thresholds	7
Best Practices	7
Do's	7
Don'ts	7
Troubleshooting	7
High Latency	7
High Error Rate	7
Scaling Issues	8

RADIANT Performance Guide

Overview

This guide covers performance optimization, caching strategies, and scalability considerations for the RADIANT platform.

Architecture Performance

Request Flow

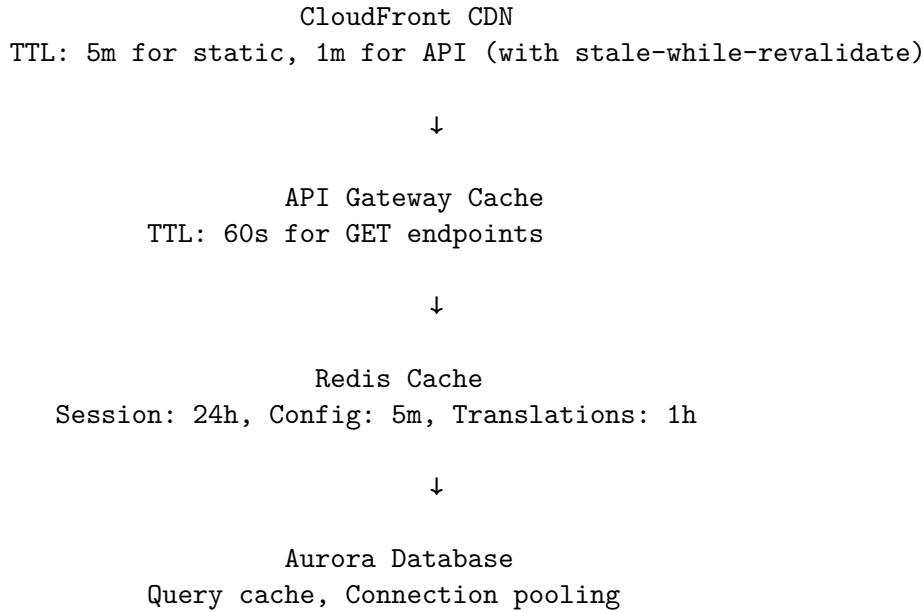
```
Client → CloudFront → WAF → API Gateway → Lambda → Aurora  
                                ↓  
                                Redis Cache
```

Latency Targets

Component	Target	Max Acceptable
CloudFront edge	< 50ms	100ms
WAF processing	< 5ms	20ms
API Gateway	< 20ms	50ms
Lambda cold start	< 500ms	1000ms
Lambda execution	< 200ms	500ms
Database query	< 50ms	200ms
Total P95	< 500ms	2000ms

Caching Strategy

Multi-Layer Caching



Cache Keys

```
// Session cache  
`session:${tenantId}:${userId}` → TTL: 24h  
  
// Configuration cache
```

```

`config:${tenantId}:${key}` → TTL: 5m
`config:global:${key}` → TTL: 5m

// Translation cache
`i18n:${language}:bundle` → TTL: 1h
`i18n:${language}:${key}` → TTL: 1h

// Model cache
`models:${tenantId}:list` → TTL: 5m
`models:${tenantId}:${modelId}` → TTL: 5m

// Rate limit cache
`ratelimit:${tenantId}:${endpoint}` → TTL: 1m
`ratelimit:ip:${ip}` → TTL: 5m

```

Cache Invalidation

```

// Pattern-based invalidation
await redis.del(`config:${tenantId}:*`);

// Event-driven invalidation
eventBridge.putEvents({
  Entries: [
    {
      Source: 'radianc.config',
      DetailType: 'ConfigUpdated',
      Detail: JSON.stringify({ tenantId, key }),
    },
  ],
});

```

Database Optimization

Connection Pooling

```

// RDS Proxy configuration
const pool = {
  min: 2,
  max: 10,
  idleTimeoutMillis: 30000,
  connectionTimeoutMillis: 5000,
};

```

Query Optimization

```

-- Always use indexes
CREATE INDEX idx_models_tenant_status ON ai_models(tenant_id, status);
CREATE INDEX idx_transactions_tenant_date ON credit_transactions(tenant_id, created_at DESC);

-- Use covering indexes for common queries
CREATE INDEX idx_models_list ON ai_models(tenant_id, status, is_enabled)

```

```

INCLUDE (display_name, category, input_cost_per_1k);

-- Partition large tables by date
CREATE TABLE audit_logs_2024_01 PARTITION OF audit_logs
FOR VALUES FROM ('2024-01-01') TO ('2024-02-01');

```

RLS Performance

```

-- Set tenant context once per request
SET app.current_tenant_id = 'tenant-123';

-- All subsequent queries automatically filtered
SELECT * FROM models; -- Implicitly filtered by RLS

```

Lambda Optimization

Cold Start Reduction

```

// 1. Minimize dependencies
// 2. Use Lambda layers for shared code
// 3. Enable provisioned concurrency for critical functions

// Provisioned concurrency config
new lambda.Function(this, 'Router', {
  // ... config
  provisionedConcurrentExecutions: 10, // Keep 10 warm
});

```

Memory Optimization

```

// Memory vs CPU tradeoff
// More memory = more CPU = faster execution

// Recommended settings by function type:
const memoryConfig = {
  router: 1024,      // Main API - balanced
  billing: 512,      // Light compute
  aiProxy: 2048,      // Heavy compute for AI
  migration: 256,     // Infrequent, light
};

```

Bundling

```

// esbuild configuration for minimal bundle size
{
  bundle: true,
  minify: true,
  treeShaking: true,
  external: ['aws-sdk'], // Use Lambda runtime SDK
}

```

```
    target: 'node20',
}
```

Rate Limiting

Tier Limits

Tier	RPS	Burst	Daily
Free	10	20	1,000
Starter	50	100	10,000
Professional	100	200	50,000
Business	500	1,000	250,000
Enterprise	2,000	5,000	Unlimited

Implementation

```
// Token bucket algorithm in Redis
const rateLimiter = {
  async checkLimit(tenantId: string, limit: number): Promise<boolean> {
    const key = `ratelimit:${tenantId}`;
    const current = await redis.incr(key);

    if (current === 1) {
      await redis.expire(key, 1); // 1 second window
    }

    return current <= limit;
  }
};
```

Load Testing

Running Tests

```
# Install k6
brew install k6

# Run smoke test
k6 run --env BASE_URL=https://api-dev.radiant.example.com tests/load/k6-config.js

# Run with specific scenario
k6 run --env BASE_URL=https://api-dev.radiant.example.com \
-e SCENARIO=load tests/load/k6-config.js
```

Performance Baselines

Metric	Baseline	Target
Throughput	500 RPS	2000 RPS
P50 Latency	100ms	50ms
P95 Latency	500ms	200ms
P99 Latency	1000ms	500ms
Error Rate	< 1%	< 0.1%

Scaling

Horizontal Scaling

Component	Scaling Method
Lambda	Automatic (up to account limit)
Aurora	Read replicas, Serverless v2
Redis	ElastiCache cluster mode
API Gateway	Automatic

Vertical Scaling

```
// Aurora Serverless v2 scaling
const database = new rds.DatabaseCluster(this, 'Database', {
  serverlessV2MinCapacity: 0.5, // Minimum ACUs
  serverlessV2MaxCapacity: 16, // Maximum ACUs
});

// Lambda memory scaling
const lambda = new lambda.Function(this, 'Function', {
  memorySize: 2048, // More memory = more CPU
});
```

Monitoring

Key Metrics

```
// CloudWatch metrics to monitor
const metrics = {
  // Latency
  'AWS/ApiGateway/Latency': 'p95 < 500ms',
  'AWS/Lambda/Duration': 'p95 < 200ms',
  'AWS/RDS/ReadLatency': 'avg < 50ms',

  // Throughput
  'AWS/ApiGateway/Count': 'track trends',
  'AWS/Lambda/Invocations': 'track trends',

  // Errors
  'AWS/ApiGateway/5XXError': 'rate < 1%',
```

```

'AWS/Lambda/Errors': 'rate < 1%',

// Resources
'AWS/Lambda/ConcurrentExecutions': '< 80% of limit',
'AWS/RDS/CPUUtilization': '< 80%',
'AWS/RDS/DatabaseConnections': '< 80% of max',
};


```

Alerting Thresholds

Metric	Warning	Critical
API P95 Latency	> 1s	> 3s
Error Rate	> 1%	> 5%
Lambda Concurrent	> 500	> 800
DB CPU	> 70%	> 85%
DB Connections	> 60%	> 80%

Best Practices

Do's

- Cache aggressively with proper invalidation
- Use connection pooling
- Minimize cold starts with provisioned concurrency
- Use read replicas for read-heavy workloads
- Implement circuit breakers for external services
- Use async processing for non-critical paths

Don'ts

- Don't make synchronous calls to external APIs in hot paths
- Don't use Lambda for long-running tasks (> 15 min)
- Don't store large objects in Redis
- Don't rely on API Gateway caching for dynamic data
- Don't skip database indexes

Troubleshooting

High Latency

1. Check Lambda cold starts (enable provisioned concurrency)
2. Check database query times (add indexes)
3. Check external API latency (add caching/circuit breaker)
4. Check connection pool exhaustion

High Error Rate

1. Check Lambda errors in CloudWatch Logs
2. Check database connection errors

3. Check rate limiting (429 errors)
4. Check WAF blocked requests

Scaling Issues

1. Check Lambda concurrent execution limit
2. Check database connection limit
3. Check API Gateway throttling
4. Check Redis memory usage