

Contents

RADIANT Platform Architecture	1
Table of Contents	1
1. Platform Overview	2
1.1 What is RADIANT?	2
1.2 Core Value Proposition	2
1.3 Platform Statistics	2
2. System Architecture	2
2.1 High-Level Architecture	2
2.2 Three-Component Structure	4
3. Component Deep Dive	5
3.1 Model Router Service	5
3.2 Service Architecture	6
4. Data Architecture	7
4.1 Database Schema Overview	7
4.2 Multi-Tenant Data Isolation	9
5. Security Architecture	9
5.1 Security Layers	9
6. Deployment Architecture	11
6.1 AWS Infrastructure	11
6.2 CDK Stack Dependencies	11
7. Integration Points	12
7.1 External API Integrations	12

RADIANT Platform Architecture

Enterprise Multi-Tenant AI Platform

Version 4.18.0 | December 2024

A comprehensive technical architecture document for the RADIANT AI orchestration platform

Table of Contents

- 1. Platform Overview
 - 2. System Architecture
 - 3. Component Deep Dive
 - 4. Data Architecture
 - 5. Security Architecture
 - 6. Deployment Architecture
 - 7. Integration Points
-

1. Platform Overview

1.1 What is RADIANT?

RADIANT (Real-time AI Distribution, Integration, and Automation Network for Tenants) is an enterprise-grade, multi-tenant SaaS platform that provides unified access to 106+ AI models across multiple providers, with intelligent orchestration, cost management, and comprehensive analytics.

1.2 Core Value Proposition

RADIANT VALUE PROPOSITION

UNIFIED ACCESS	INTELLIGENT ORCHESTRATION	COST MANAGEMENT	ENTERPRISE SECURITY
106+ AI Models One API	49 Multi- AI Patterns AGI Router	Credits, Budgets, Analytics	SOC2/HIPAA Compliant Multi-Tenant

1.3 Platform Statistics

Metric	Value
AI Models Supported	106+ (50 external + 56 self-hosted)
AI Providers Integrated	15+ (OpenAI, Anthropic, Google, Meta, etc.)
Orchestration Patterns	49 documented patterns
Model Execution Modes	9 (thinking, research, fast, creative, etc.)
Database Migrations	66+ schema migrations
CDK Stacks	14 infrastructure stacks

2. System Architecture

2.1 High-Level Architecture

RADIANT PLATFORM ARCHITECTURE

CLIENT LAYER

Swift Deployer (macOS)	Admin Dashboard (Next.js)	Think Tank Consumer App	SDK & API Clients
---------------------------	------------------------------	-------------------------	-------------------

API GATEWAY LAYER

- Amazon API Gateway
- REST APIs
 - JWT Auth
 - WebSocket APIs
 - API Keys
 - Rate Limiting
 - Usage Plans

COMPUTE LAYER (Lambda)

Model Router	Orchestration Engine	Billing Service	Admin Service
Think Tank	Analytics Service	Learning Service	Webhook Service

DATA LAYER

Aurora PostgreSQL (RLS)	DynamoDB (Sessions)	S3 (Storage)	Redis (Cache)
----------------------------	------------------------	-----------------	------------------

EXTERNAL AI PROVIDERS

OpenAI GPT-4o o1	Anthropic Claude 3.5	Google Gemini 2.0	Meta Llama 3.1	Mistral	+10 more
------------------------	----------------------------	-------------------------	----------------------	---------	----------

2.2 Three-Component Structure

RADIANT consists of three primary deployment components:

THREE COMPONENTS OF RADIANT

`packages/infrastructure/lambda/shared/services/`

CORE SERVICES

<code>model-router.service.ts</code>	Route requests to AI providers
<code>model-metadata.service.ts</code>	Live model data & capabilities
<code>orchestration-patterns.service</code>	49 multi-AI workflow patterns
<code>superior-orchestration.service</code>	Guaranteed superior responses
<code>learning.service.ts</code>	ML feedback & continuous learning

BILLING SERVICES

<code>billing.service.ts</code>	Credit & subscription management
<code>cost-management.service.ts</code>	Budget alerts & cost tracking
<code>usage-analytics.service.ts</code>	Usage metrics & reporting

PLATFORM SERVICES

<code>tenant.service.ts</code>	Multi-tenant management
<code>auth.service.ts</code>	Authentication & authorization
<code>api-key.service.ts</code>	API key lifecycle
<code>webhook.service.ts</code>	Event notifications
<code>storage.service.ts</code>	File & artifact storage

THINK TANK SERVICES

thinktank-engine.ts	Multi-step problem solving
thinktank-sessions.ts	Conversation management
collaboration.service.ts	Real-time collaboration

3. Component Deep Dive

3.1 Model Router Service

The intelligent core that routes AI requests to optimal providers:

MODEL ROUTER SERVICE

Incoming
Request

Request Validation
• API Key Check
• Rate Limiting
• Tenant Verification

Model Selection	Budget Check	Fallback Logic
• Metadata • Preferences • Capabilities	• Credits • Limits • Cost Est.	• Primary • Secondary • Tertiary

Provider Adapter Layer

OpenAI Anthropic ...
Adapter Adapter

Response Processing

- Token Counting
- Cost Calculation
- Usage Recording
- Analytics Event

3.2 Service Architecture

LAMBDA SERVICES ARCHITECTURE

packages/infrastructure/lambda/shared/services/

CORE SERVICES

model-router.service.ts	Route requests to AI providers
model-metadata.service.ts	Live model data & capabilities
orchestration-patterns.service	49 multi-AI workflow patterns
superior-orchestration.service	Guaranteed superior responses
learning.service.ts	ML feedback & continuous learning

BILLING SERVICES

billing.service.ts	Credit & subscription management
cost-management.service.ts	Budget alerts & cost tracking

usage-analytics.service.ts Usage metrics & reporting

PLATFORM SERVICES

tenant.service.ts	Multi-tenant management
auth.service.ts	Authentication & authorization
api-key.service.ts	API key lifecycle
webhook.service.ts	Event notifications
storage.service.ts	File & artifact storage

THINK TANK SERVICES

thinktank-engine.ts	Multi-step problem solving
thinktank-sessions.ts	Conversation management
collaboration.service.ts	Real-time collaboration

4. Data Architecture

4.1 Database Schema Overview

AURORA POSTGRESQL SCHEMA

66+ Migrations in packages/infrastructure/migrations/

CORE ENTITIES

tenants	Multi-tenant organizations
users	User accounts with roles
api_keys	API authentication keys
model_configurations	Per-tenant model settings
model_metadata	AI model capabilities & pricing

BILLING & CREDITS

credit_accounts	Tenant credit balances
credit_transactions	Credit usage history
subscriptions	Plan subscriptions
invoices	Billing invoices
budgets	Spending limits & alerts

ORCHESTRATION

orchestration_methods	Reusable AI method definitions
orchestration_workflows	49 workflow patterns
workflow_method_bindings	Steps linking workflows to methods
orchestration_executions	Execution history & results

THINK TANK

thinktank_sessions	Problem-solving sessions
thinktank_conversations	Conversation threads
thinktank_messages	Individual messages
thinktank_steps	Reasoning steps
thinktank_artifacts	Generated outputs

ANALYTICS & LEARNING

usage_events	API usage events
analytics_aggregates	Pre-computed metrics
learning_interactions	ML training data
model_performance	Model quality tracking

SECURITY

Row-Level Security (RLS) on all tenant tables
SET app.current_tenant_id for automatic filtering
Audit logging on sensitive operations

4.2 Multi-Tenant Data Isolation

ROW-LEVEL SECURITY (RLS) MODEL

Request from Tenant A

JWT Token
tenant_id=A

Request from Tenant B

JWT Token
tenant_id=B

Database Connection

```
SET app.current_tenant_id = 'tenant_id_from_jwt';
```

RLS Policy Applied

```
CREATE POLICY tenant_isolation ON table_name
USING (tenant_id = current_setting('app.current_tenant_id'));
```

Result: Each tenant ONLY sees their own data

Tenant A sees:

Only Tenant A's
- Users
- API Keys
- Usage Data
- Conversations

Tenant B sees:

Only Tenant B's
- Users
- API Keys
- Usage Data
- Conversations

5. Security Architecture

5.1 Security Layers

SECURITY ARCHITECTURE

LAYER 1: NETWORK SECURITY

- VPC with private subnets for database
- WAF rules for API Gateway
- CloudFront for DDoS protection
- TLS 1.3 for all connections

LAYER 2: AUTHENTICATION

- Cognito User Pools for user authentication
- JWT tokens with tenant claims
- API Keys with scoped permissions
- MFA support for admin users

LAYER 3: AUTHORIZATION

- Role-based access control (RBAC)
- Permission sets per tenant
- Resource-level policies
- API endpoint authorization

LAYER 4: DATA SECURITY

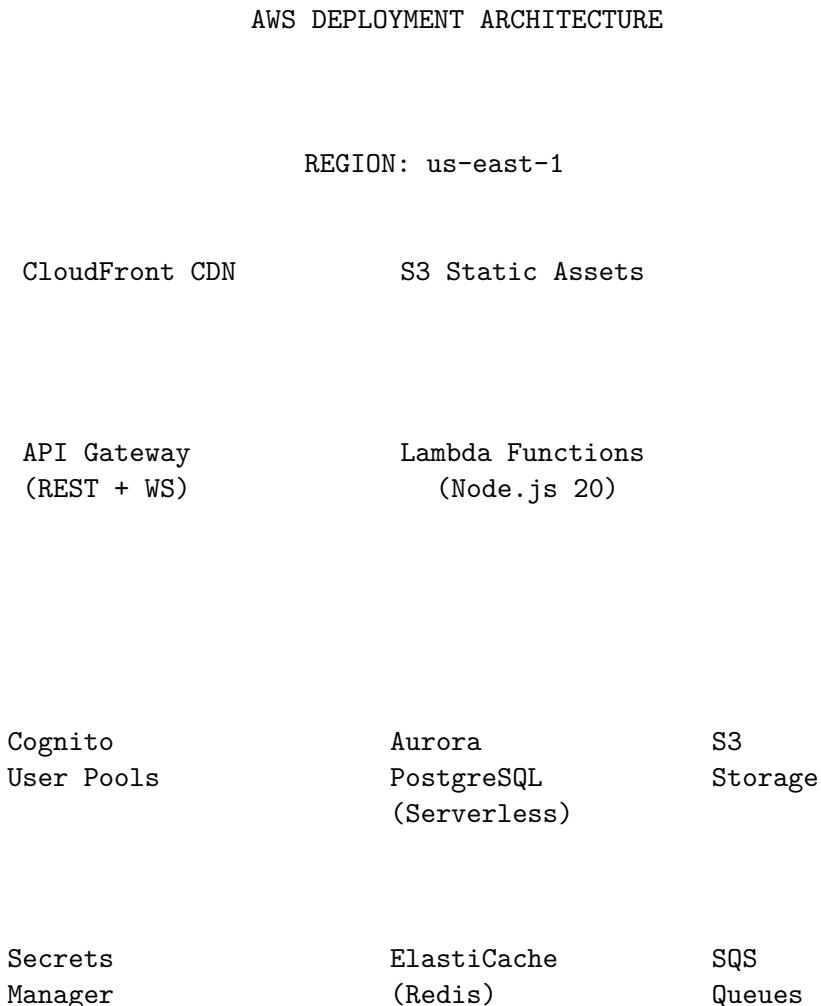
- Row-Level Security (RLS) in PostgreSQL
- Encryption at rest (AES-256)
- Encryption in transit (TLS)
- KMS for key management
- PHI sanitization for HIPAA compliance

LAYER 5: AUDIT & COMPLIANCE

- CloudTrail for API logging
- Audit tables for data changes
- Compliance reporting dashboard
- SOC2 Type II controls
- HIPAA compliance mode

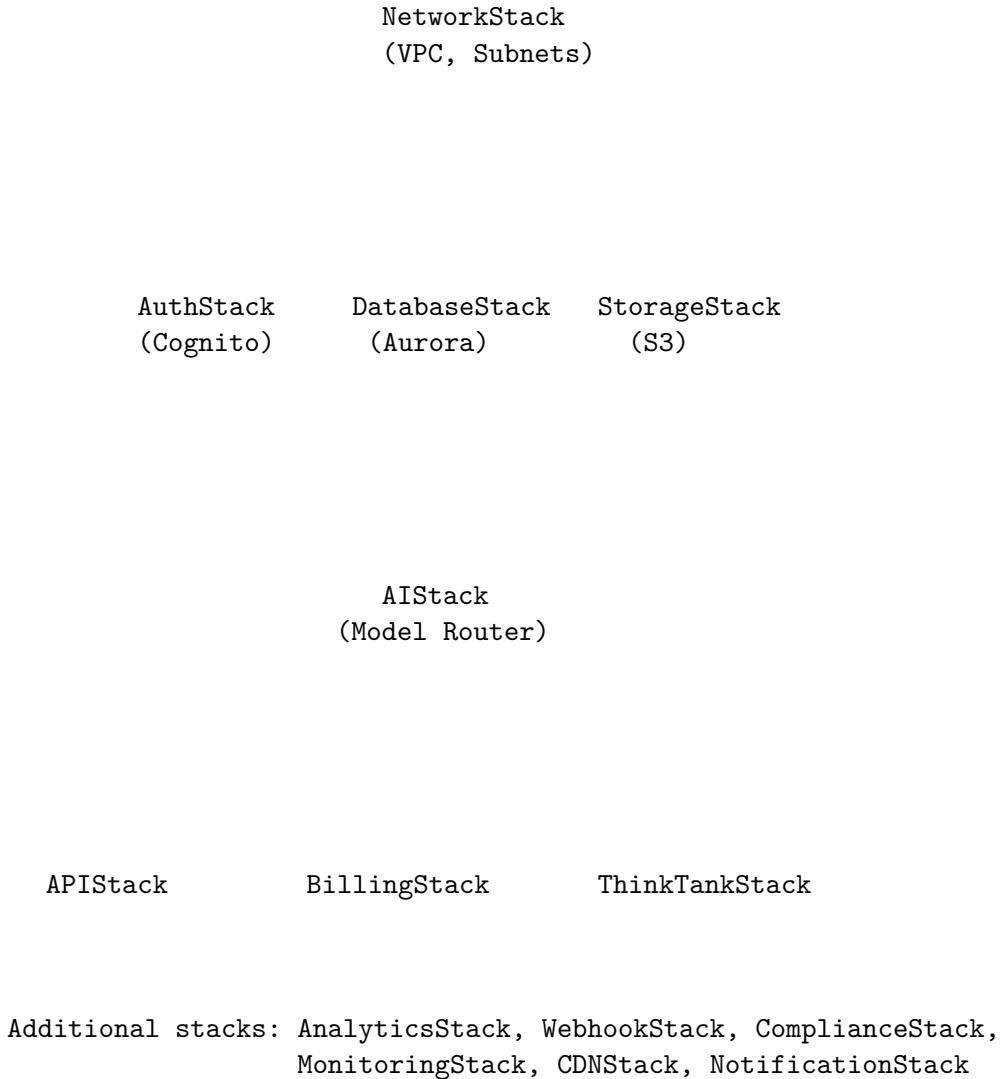
6. Deployment Architecture

6.1 AWS Infrastructure



6.2 CDK Stack Dependencies

CDK STACK DEPENDENCIES



7. Integration Points

7.1 External API Integrations

EXTERNAL INTEGRATIONS

AI PROVIDERS (15+)

OpenAI Anthropic Google Meta Mistral

GPT-4o, o1 Claude 3.5 Gemini 2.0 Llama 3.1 Large

Cohere AI21 Perplexity DeepSeek xAI
Sonar R1, Chat Grok

PAYMENT PROVIDERS

Stripe Credit card processing, subscriptions, invoicing

MONITORING & OBSERVABILITY

CloudWatch X-Ray Sentry Custom Analytics
(Logs) (Traces) (Errors) Dashboard

NOTIFICATIONS

SES SNS Webhooks Slack/Teams Integrations
(Email) (Push) (Custom)

RADIANT Platform Architecture v4.18.0

Building the future of enterprise AI

© 2024 RADIANT. All Rights Reserved.