Contents

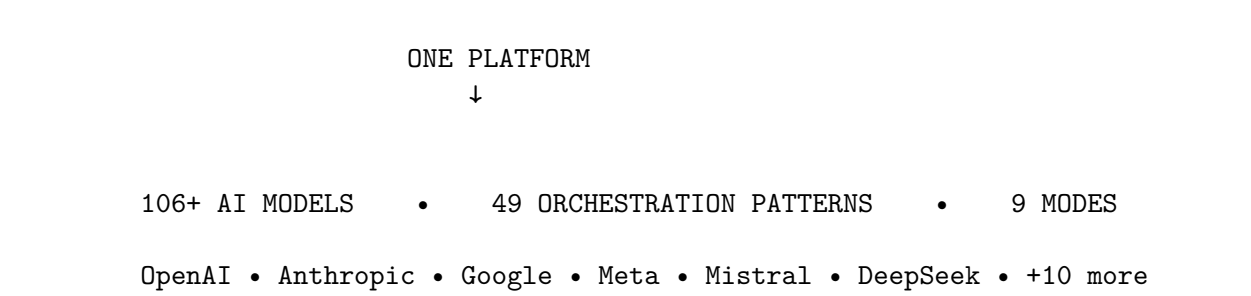# RADIANT & Think Tank Executive Summary

Enterprise AI Platform Overview

Version 4.18.0 | December 2024

For executives, investors, and decision-makers

What is RADIANT?

RADIANT is an enterprise-grade, multi-tenant AI platform that provides organizations with unified access to 106+ AI models through a single API, with intelligent orchestration that coordinates multiple AI systems to deliver superior results.

```
                    ONE PLATFORM
                         ↓


    106+ AI MODELS    •    49 ORCHESTRATION PATTERNS    •    9 MODES

    OpenAI • Anthropic • Google • Meta • Mistral • DeepSeek • +10 more
```

```
                                    ↓
            INTELLIGENT ROUTING   •   COST MANAGEMENT   •   COMPLIANCE
```

What is Think Tank?

Think Tank is RADIANT's advanced problem-solving platform that decomposes complex problems into manageable steps, applies multi-AI reasoning, and synthesizes comprehensive solutions with confidence scoring.

```
   COMPLEX PROBLEM   →   DECOMPOSE   →   REASON   →   EXECUTE   →   SYNTHESIZE

   "Design a scalable        Break into        Step-by-step    Multiple
    microservices            5 sub-            chain-of-       AI models
    architecture"            problems          thought         in parallel

                                     ↓

            COMPREHENSIVE SOLUTION + CONFIDENCE SCORE
```

Key Differentiators

1. AGI-Driven Model Selection

Unlike platforms that use a single AI model, RADIANT's AGI layer **automatically selects the optimal combination of models** based on task analysis:

What We Analyze	What We Select
Problem domain (coding, legal, medical…)	Best models for that domain
Task complexity	Number of models (2-5)
Reasoning requirements	Execution mode (thinking, fast, precise…)
Quality vs speed priority	Parallel execution strategy

Result: 50-300% better outcomes than single-model approaches.
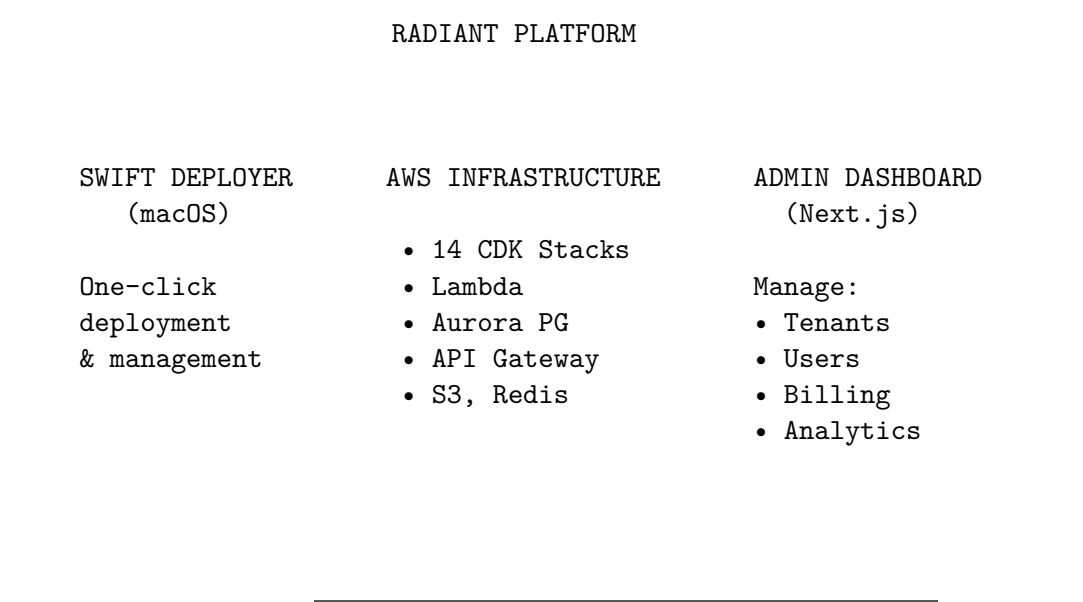
2. 49 Proven Orchestration Patterns

Research-backed workflows including: - **AI Debate** - Two AIs argue, judge decides - **Self-Refine** - Generate → Critique → Improve - **Chain-of-Verification** - Fact-check every claim - **Tree of**

Thoughts - Explore multiple solution paths

3. Enterprise-Grade Security

Capability	Description
Multi-Tenant Isolation	PostgreSQL Row-Level Security
Compliance	SOC2, HIPAA-ready
Encryption	At-rest (AES-256) and in-transit (TLS 1.3)
Audit Logging	Complete activity trail

Platform Components

```
                  RADIANT PLATFORM



SWIFT DEPLOYER      AWS INFRASTRUCTURE      ADMIN DASHBOARD
   (macOS)                                     (Next.js)
                  • 14 CDK Stacks
One-click         • Lambda               Manage:
deployment        • Aurora PG            • Tenants
& management      • API Gateway          • Users
                  • S3, Redis            • Billing
                                         • Analytics
```

By the Numbers

Metric	Value
AI Models	106+ (50 external + 56 self-hosted)
AI Providers	15+ integrated
Orchestration Patterns	49 documented workflows
Execution Modes	9 specialized modes
Database Migrations	66+ schema versions
CDK Stacks	14 infrastructure components

Use Cases

Enterprise AI Gateway

- Unified access to all major AI providers
- Centralized cost management and budgeting
- Consistent API regardless of backend model
- Automatic failover for reliability

Complex Problem Solving (Think Tank)

- Multi-step technical analysis
- Research synthesis with citations
- Architecture design with artifacts
- Decision support with confidence scores

Quality-Critical Applications

- Legal document analysis (precise mode)
- Medical information processing (HIPAA compliant)
- Financial analysis (multi-model verification)
- Code generation (AI debate + critique)

Cost Optimization

- Intelligent model routing (use cheaper models when appropriate)
- Budget alerts and limits
- Usage analytics by team/project
- Model performance vs cost analysis

Competitive Advantages

vs. Single-Model APIs	vs. Other Platforms
Multi-model orchestration	49 research-backed patterns
Built-in verification	AGI-driven model selection
Higher accuracy	9 execution modes
Reduced bias	Visual workflow editor
Confidence scoring	Think Tank problem solving

Technology Stack

Layer	Technology
Frontend	Next.js 14, TypeScript, Tailwind CSS, shadcn/ui
Backend	AWS Lambda (Node.js 20), API Gateway

Layer	Technology
Database	Aurora PostgreSQL (Serverless), DynamoDB, Redis
Infrastructure	AWS CDK (TypeScript), 14 stacks
Desktop	SwiftUI (macOS 13.0+, Swift 5.9+)
Security	Cognito, KMS, WAF, Row-Level Security

Deployment Model

RADIANT deploys to **your AWS account**:

```
YOUR AWS ACCOUNT


            RADIANT INFRASTRUCTURE

   • Your data stays in your account
   • Your compliance requirements met
   • Your region/residency requirements
   • Full control over infrastructure


   Deployed via Swift Deployer (macOS app) or CLI
```

Pricing Model

Tier	Target	Includes
Free	Developers	10K tokens/month, 3 models
Pro	Teams	1M tokens/month, all models, orchestration
Enterprise	Organizations	Unlimited, SLA, custom patterns, HIPAA

All tiers include: - Full API access - Admin dashboard - Basic analytics - Email support

Roadmap Highlights

Timeframe	Features
Q1 2026	Mobile SDK, more self-hosted models
Q2 2026	Fine-tuning pipeline, custom model hosting
Q3 2026	Multi-region deployment, advanced compliance
Q4 2026	Marketplace for custom patterns

Summary

RADIANT + Think Tank delivers:

1. **Unified AI Access** - One API for 106+ models across 15+ providers
2. **Intelligent Orchestration** - AGI selects optimal models and modes
3. **Superior Results** - 49 patterns achieve 50-300% better outcomes
4. **Enterprise Security** - Multi-tenant, SOC2, HIPAA-ready
5. **Cost Control** - Budgets, analytics, intelligent routing
6. **Problem Solving** - Think Tank for complex multi-step reasoning

---

RADIANT v4.18.0 + Think Tank v3.2.0

The enterprise platform for intelligent AI orchestration

Contact: info@radiant.ai | **Documentation:** docs.radiant.ai