

Contents

SECTION 4: LAMBDA FUNCTIONS - CORE (v2.1.0)	2
Type Imports	2
RADIANT v2.2.0 - Prompt 4: Lambda Functions - Core	2
OVERVIEW	2
LAMBDA DIRECTORY STRUCTURE	2
PART 1: LAMBDA CONFIGURATION	3
packages/infrastructure/lambda/package.json	3
packages/infrastructure/lambda/tsconfig.json	4
PART 2: SHARED UTILITIES	4
packages/infrastructure/lambda/shared/index.ts	4
packages/infrastructure/lambda/shared/config.ts	5
packages/infrastructure/lambda/shared/logger.ts	6
packages/infrastructure/lambda/shared/errors.ts	9
packages/infrastructure/lambda/shared/response.ts	12
packages/infrastructure/lambda/shared/auth.ts	16
PART 3: DATABASE CLIENT	20
packages/infrastructure/lambda/shared/db/index.ts	20
packages/infrastructure/lambda/shared/db/types.ts	20
packages/infrastructure/lambda/shared/db/client.ts	23
packages/infrastructure/lambda/shared/db/queries.ts	30
PART 4: LITELLM CLIENT	37
packages/infrastructure/lambda/shared/litellm/index.ts	37
packages/infrastructure/lambda/shared/litellm/types.ts	37
packages/infrastructure/lambda/shared/litellm/client.ts	41
PART 5: PHI SANITIZATION	46
packages/infrastructure/lambda/shared/phi/index.ts	46
packages/infrastructure/lambda/shared/phi/types.ts	46
packages/infrastructure/lambda/shared/phi/patterns.ts	48
packages/infrastructure/lambda/shared/phi/sanitizer.ts	51
PART 6: ROUTER LAMBDA	56
packages/infrastructure/lambda/api/router.ts	56
PART 7: CHAT LAMBDA	60
packages/infrastructure/lambda/api/chat.ts	60
PART 8: MODELS LAMBDA	69
packages/infrastructure/lambda/api/models.ts	69
PART 9: PROVIDERS LAMBDA	75
packages/infrastructure/lambda/api/providers.ts	75
DEPLOYMENT VERIFICATION	81
NEXT PROMPTS	82
END OF SECTION 4	82

SECTION 4: LAMBDA FUNCTIONS - CORE (v2.1.0)

Dependencies: Sections 0-3 **Creates:** Router, Chat, Models, Providers, PHI handlers

Type Imports

```
import { AIModel, AIProvider, ThermalState, PHICConfig, UsageEvent } from '@radiant/shared';
```

RADIANT v2.2.0 - Prompt 4: Lambda Functions - Core

Prompt 4 of 9 | Target Size: ~55KB | Version: 3.7.0 | December 2024

OVERVIEW

This prompt creates the core Lambda functions for the RADIANT API:

1. **Router Lambda** - Main API entry point, request routing, health checks
 2. **Chat Lambda** - AI completions via LiteLLM, streaming support
 3. **Models Lambda** - Dynamic model registry CRUD operations
 4. **Providers Lambda** - Provider management and configuration
 5. **PHI Lambda** - HIPAA-compliant PHI sanitization and re-identification

Plus shared utilities: - Database client (Aurora PostgreSQL via Data API) - LiteLLM client (HTTP client for AI routing) - Authentication helpers (JWT validation, tenant extraction) - Response helpers (standardized API responses) - Logging and tracing utilities

LAMBDA DIRECTORY STRUCTURE

```
packages/infrastructure/lambda/
  Ä€ "Ä€ " - tsconfig.json
  Ä€ "Ä€ " - package.json
  Ä€ "Ä€ " - shared/
    Ä€ "Ä€ " - index.ts
    Ä€ "Ä€ " - config.ts
    Ä€ "Ä€ " - logger.ts
    Ä€ "Ä€ " - errors.ts
    Ä€ "Ä€ " - response.ts
    Ä€ "Ä€ " - auth.ts
```

Æ¢â€ ªæš Æ¢â€ ª"Æ¢â€ ª,¬Æ¢â€ ª,¬ db/
 Æ¢â€ ªæš Æ¢â€ ªæš Æ¢â€ ª"Æ¢â€ ª,¬Æ¢â€ ª,¬ index.ts
 Æ¢â€ ªæš Æ¢â€ ªæš Æ¢â€ ª"Æ¢â€ ª,¬Æ¢â€ ª,¬ client.ts
 Æ¢â€ ªæš Æ¢â€ ªæš Æ¢â€ ª"Æ¢â€ ª,¬Æ¢â€ ª,¬ queries.ts
 Æ¢â€ ªæš Æ¢â€ ªæš Æ¢â€ ªæ€ ª,¬Æ¢â€ ª,¬ types.ts
 Æ¢â€ ªæš Æ¢â€ ª"Æ¢â€ ª,¬Æ¢â€ ª,¬ litellm/
 Æ¢â€ ªæš Æ¢â€ ªæš Æ¢â€ ª"Æ¢â€ ª,¬Æ¢â€ ª,¬ index.ts
 Æ¢â€ ªæš Æ¢â€ ªæš Æ¢â€ ª"Æ¢â€ ª,¬Æ¢â€ ª,¬ client.ts
 Æ¢â€ ªæš Æ¢â€ ªæš Æ¢â€ ªæ€ ª,¬Æ¢â€ ª,¬ types.ts
 Æ¢â€ ªæš Æ¢â€ ªæ€ ª,¬Æ¢â€ ª,¬ phi/
 Æ¢â€ ªæš Æ¢â€ ª"Æ¢â€ ª,¬Æ¢â€ ª,¬ index.ts
 Æ¢â€ ªæš Æ¢â€ ª"Æ¢â€ ª,¬Æ¢â€ ª,¬ sanitizer.ts
 Æ¢â€ ªæš Æ¢â€ ª"Æ¢â€ ª,¬Æ¢â€ ª,¬ patterns.ts
 Æ¢â€ ªæš Æ¢â€ ªæ€ ª,¬Æ¢â€ ª,¬ types.ts
 Æ¢â€ ª"Æ¢â€ ª,¬Æ¢â€ ª,¬ api/
 Æ¢â€ ªæš Æ¢â€ ª"Æ¢â€ ª,¬Æ¢â€ ª,¬ router.ts
 Æ¢â€ ªæš Æ¢â€ ª"Æ¢â€ ª,¬Æ¢â€ ª,¬ chat.ts
 Æ¢â€ ªæš Æ¢â€ ª"Æ¢â€ ª,¬Æ¢â€ ª,¬ models.ts
 Æ¢â€ ªæš Æ¢â€ ªæ€ ª,¬Æ¢â€ ª,¬ providers.ts
 Æ¢â€ ªæ€ ª,¬Æ¢â€ ª,¬ admin/
 Æ¢â€ ªæ€ ª,¬Æ¢â€ ª,¬ [See Prompt 5]

PART 1: LAMBDA CONFIGURATION

packages/infrastructure/lambda/package.json

```
{
  "name": "@radiant/lambda",
  "version": "2.2.0",
  "private": true,
  "scripts": {
    "build": "tsc",
    "clean": "rm -rf dist",
    "lint": "eslint . --ext .ts",
    "test": "jest"
  },
  "dependencies": {
    "@aws-sdk/client-dynamodb": "^3.470.0",
    "@aws-sdk/client-rds-data": "^3.470.0",
    "@aws-sdk/client-secrets-manager": "^3.470.0",
    "@aws-sdk/client-ssm": "^3.470.0",
    "@aws-sdk/client-s3": "^3.470.0",
    "@aws-sdk/client-sagemaker-runtime": "^3.470.0",
    "@aws-sdk/lib-dynamodb": "^3.470.0",
    "@aws-sdk/s3-request-presigner": "^3.470.0",
    "uuid": "^9.0.0",
    "zod": "^3.22.0"
  }
}
```

```

    },
    "devDependencies": {
        "@types/aws-lambda": "^8.10.130",
        "@types/node": "^20.10.0",
        "@types/uuid": "^9.0.0",
        "typescript": "^5.3.0",
        "jest": "^29.7.0",
        "@types/jest": "^29.5.11",
        "ts-jest": "^29.1.1"
    }
}

```

`packages/infrastructure/lambda/tsconfig.json`

```

{
    "compilerOptions": {
        "target": "ES2022",
        "module": "NodeNext",
        "moduleResolution": "NodeNext",
        "lib": ["ES2022"],
        "outDir": "./dist",
        "rootDir": ".",
        "strict": true,
        "esModuleInterop": true,
        "skipLibCheck": true,
        "forceConsistentCasingInFileNames": true,
        "declaration": true,
        "declarationMap": true,
        "sourceMap": true,
        "resolveJsonModule": true,
        "noUnusedLocals": true,
        "noUnusedParameters": true,
        "noImplicitReturns": true,
        "noFallthroughCasesInSwitch": true
    },
    "include": ["**/*.ts"],
    "exclude": ["node_modules", "dist"]
}

```

PART 2: SHARED UTILITIES

`packages/infrastructure/lambda/shared/index.ts`

```

// Re-export all shared utilities
export * from './config';
export * from './logger';
export * from './errors';

```

```

export * from './response';
export * from './auth';
export * from './db';
export * from './litellm';
export * from './phi';

packages/infrastructure/lambda/shared/config.ts

/**
 * Environment configuration with validation
 */

import { z } from 'zod';

const envSchema = z.object({
  APP_ID: z.string().min(1),
  ENVIRONMENT: z.enum(['dev', 'staging', 'prod']),
  TIER: z.string().transform(Number).pipe(z.number().min(1).max(5)),
  LITELLM_URL: z.string().url(),
  AURORA_SECRET_ARN: z.string().startsWith('arn:aws:secretsmanager:'),
  AURORA_CLUSTER_ARN: z.string().startsWith('arn:aws:rds:'),
  USAGE_TABLE: z.string().min(1),
  SESSIONS_TABLE: z.string().min(1),
  CACHE_TABLE: z.string().min(1),
  MEDIA_BUCKET: z.string().min(1),
  USER_POOL_ID: z.string().min(1),
  LOG_LEVEL: z.enum(['debug', 'info', 'warn', 'error']).default('info'),
  AWS_REGION: z.string().default('us-east-1'),
});

export type Config = z.infer<typeof envSchema>;

let cachedConfig: Config | null = null;

export function getConfig(): Config {
  if (cachedConfig) return cachedConfig;

  const result = envSchema.safeParse(process.env);

  if (!result.success) {
    console.error('Configuration validation failed:', result.error.flatten());
    throw new Error(`Invalid configuration: ${JSON.stringify(result.error.flatten())}`);
  }

  cachedConfig = result.data;
  return cachedConfig;
}

```

```

/**
 * Feature flags based on tier
 */
export interface FeatureFlags {
  multiRegion: boolean;
  waf: boolean;
  guardDuty: boolean;
  sagemaker: boolean;
  elasticache: boolean;
  xray: boolean;
  phiSanitization: boolean;
  advancedMetrics: boolean;
}

export function getFeatureFlags(tier: number): FeatureFlags {
  return {
    multiRegion: tier >= 4,
    waf: tier >= 2,
    guardDuty: tier >= 2,
    sagemaker: tier >= 3,
    elasticache: tier >= 2,
    xray: tier >= 2,
    phiSanitization: true, // Always available
    advancedMetrics: tier >= 3,
  };
}

```

packages/infrastructure/lambda/shared/logger.ts

```

/**
 * Structured logging with correlation IDs
 */

export type LogLevel = 'debug' | 'info' | 'warn' | 'error';

interface LogContext {
  requestId?: string;
  tenantId?: string;
  userId?: string;
  appId?: string;
  environment?: string;
  [key: string]: unknown;
}

interface LogEntry {
  timestamp: string;
  level: LogLevel;
  message: string;
}

```

```

context?: LogContext;
error?: {
  name: string;
  message: string;
  stack?: string;
};
duration?: number;
[key: string]: unknown;
}

const LOG_LEVELS: Record<LogLevel, number> = {
  debug: 0,
  info: 1,
  warn: 2,
  error: 3,
};

export class Logger {
  private context: LogContext;
  private minLevel: LogLevel;
  private startTime: number;

  constructor(context: LogContext = {}, minLevel?: LogLevel) {
    this.context = context;
    this.minLevel = minLevel || (process.env.LOG_LEVEL as LogLevel) || 'info';
    this.startTime = Date.now();
  }

  private shouldLog(level: LogLevel): boolean {
    return LOG_LEVELS[level] >= LOG_LEVELS[this.minLevel];
  }

  private formatEntry(level: LogLevel, message: string, extra?: Record<string, unknown>): LogEntry {
    return {
      timestamp: new Date().toISOString(),
      level,
      message,
      context: this.context,
      duration: Date.now() - this.startTime,
      ...extra,
    };
  }

  private log(level: LogLevel, message: string, extra?: Record<string, unknown>): void {
    if (!this.shouldLog(level)) return;

    const entry = this.formatEntry(level, message, extra);
    const output = JSON.stringify(entry);
  }
}

```

```

switch (level) {
  case 'error':
    console.error(output);
    break;
  case 'warn':
    console.warn(output);
    break;
  default:
    console.log(output);
}
}

debug(message: string, extra?: Record<string, unknown>): void {
  this.log('debug', message, extra);
}

info(message: string, extra?: Record<string, unknown>): void {
  this.log('info', message, extra);
}

warn(message: string, extra?: Record<string, unknown>): void {
  this.log('warn', message, extra);
}

error(message: string, error?: Error, extra?: Record<string, unknown>): void {
  this.log('error', message, {
    ...extra,
    error: error ? {
      name: error.name,
      message: error.message,
      stack: error.stack,
    } : undefined,
  });
}

child(additionalContext: LogContext): Logger {
  return new Logger(
    { ...this.context, ...additionalContext },
    this.minLevel
  );
}

setRequestId(requestId: string): void {
  this.context.requestId = requestId;
}

setTenantId(tenantId: string): void {

```

```

    this.context.tenantId = tenantId;
}

setUserId(userId: string): void {
    this.context.userId = userId;
}
}

// Default logger instance
export const logger = new Logger({
    appId: process.env.APP_ID,
    environment: process.env.ENVIRONMENT,
});

packages/infrastructure/lambda/shared/errors.ts

/**
 * Custom error types for API responses
 */

export abstract class AppError extends Error {
    abstract readonly statusCode: number;
    abstract readonly code: string;
    readonly isOperational = true;

    constructor(message: string) {
        super(message);
        this.name = this.constructor.name;
        Error.captureStackTrace(this, this.constructor);
    }

    toJSON() {
        return {
            code: this.code,
            message: this.message,
            statusCode: this.statusCode,
        };
    }
}

// 400 Bad Request
export class ValidationError extends AppError {
    readonly statusCode = 400;
    readonly code = 'VALIDATION_ERROR';
    readonly details?: Record<string, string[]>;

    constructor(message: string, details?: Record<string, string[]>) {
        super(message);
    }
}

```

```

    this.details = details;
}

toJSON() {
  return {
    ...super.toJSON(),
    details: this.details,
  };
}
}

// 401 Unauthorized
export class UnauthorizedError extends AppError {
  readonly statusCode = 401;
  readonly code = 'UNAUTHORIZED';

  constructor(message = 'Authentication required') {
    super(message);
  }
}

// 403 Forbidden
export class ForbiddenError extends AppError {
  readonly statusCode = 403;
  readonly code = 'FORBIDDEN';

  constructor(message = 'Access denied') {
    super(message);
  }
}

// 404 Not Found
export class NotFoundError extends AppError {
  readonly statusCode = 404;
  readonly code = 'NOT_FOUND';
  readonly resource?: string;

  constructor(resource?: string) {
    super(resource ? `${resource} not found` : 'Resource not found');
    this.resource = resource;
  }
}

// 409 Conflict
export class ConflictError extends AppError {
  readonly statusCode = 409;
  readonly code = 'CONFLICT';
}

```

```

constructor(message: string) {
    super(message);
}
}

// 422 Unprocessable Entity
export class UnprocessableError extends AppError {
    readonly statusCode = 422;
    readonly code = 'UNPROCESSABLE_ENTITY';

    constructor(message: string) {
        super(message);
    }
}

// 429 Too Many Requests
export class RateLimitError extends AppError {
    readonly statusCode = 429;
    readonly code = 'RATE_LIMITED';
    readonly retryAfter?: number;

    constructor(retryAfter?: number) {
        super('Rate limit exceeded');
        this.retryAfter = retryAfter;
    }
}

// 500 Internal Server Error
export class InternalError extends AppError {
    readonly statusCode = 500;
    readonly code = 'INTERNAL_ERROR';

    constructor(message = 'An unexpected error occurred') {
        super(message);
    }
}

// 502 Bad Gateway (AI provider errors)
export class ProviderError extends AppError {
    readonly statusCode = 502;
    readonly code = 'PROVIDER_ERROR';
    readonly provider?: string;

    constructor(message: string, provider?: string) {
        super(message);
        this.provider = provider;
    }
}

```

```

// 503 Service Unavailable
export class ServiceUnavailableError extends AppError {
  readonly statusCode = 503;
  readonly code = 'SERVICE_UNAVAILABLE';

  constructor(message = 'Service temporarily unavailable') {
    super(message);
  }
}

/**
 * Check if error is an operational error (expected)
 */
export function isOperationalError(error: unknown): error is AppError {
  return error instanceof AppError && error.isOperational;
}

/**
 * Convert unknown error to AppError
 */
export function toAppError(error: unknown): AppError {
  if (error instanceof AppError) {
    return error;
  }

  if (error instanceof Error) {
    return new InternalError(error.message);
  }

  return new InternalError('Unknown error occurred');
}

```

packages/infrastructure/lambda/shared/response.ts

```

/**
 * Standardized API response helpers
 */

import type { APIGatewayProxyResult } from 'aws-lambda';
import { AppError, toAppError } from './errors';
import { Logger } from './logger';

interface SuccessResponse<T> {
  success: true;
  data: T;
  meta?: ResponseMeta;
}

```

```

interface ErrorResponse {
  success: false;
  error: {
    code: string;
    message: string;
    details?: unknown;
  };
}

interface ResponseMeta {
  requestId?: string;
  pagination?: {
    page: number;
    limit: number;
    total: number;
    hasMore: boolean;
  };
  timing?: {
    duration: number;
  };
}

type ApiResponse<T> = SuccessResponse<T> | ErrorResponse;

const DEFAULT_HEADERS = {
  'Content-Type': 'application/json',
  'Access-Control-Allow-Origin': '*',
  'Access-Control-Allow-Headers': 'Content-Type,Authorization,X-API-Key,X-Tenant-Id',
  'Access-Control-Allow-Methods': 'GET,POST,PUT,DELETE,OPTIONS',
  'X-Content-Type-Options': 'nosniff',
  'X-Frame-Options': 'DENY',
  'Strict-Transport-Security': 'max-age=31536000; includeSubDomains',
};

/**
 * Create success response
 */
export function success<T>(
  data: T,
  statusCode = 200,
  meta?: ResponseMeta
): APIGatewayProxyResult {
  const response: SuccessResponse<T> = {
    success: true,
    data,
    meta,
  };
}

```

```

    return {
      statusCode,
      headers: DEFAULT_HEADERS,
      body: JSON.stringify(response),
    };
}

/**
 * Create created response (201)
 */
export function created<T>(data: T, meta?: ResponseMeta): APIGatewayProxyResult {
  return success(data, 201, meta);
}

/**
 * Create no content response (204)
 */
export function noContent(): APIGatewayProxyResult {
  return {
    statusCode: 204,
    headers: DEFAULT_HEADERS,
    body: '',
  };
}

/**
 * Create error response
 */
export function error(
  err: AppError,
  logger?: Logger
): APIGatewayProxyResult {
  if (logger) {
    if (err.statusCode >= 500) {
      logger.error('Internal error', err);
    } else {
      logger.warn('Client error', { error: err.toJSON() });
    }
  }

  const response: ErrorResponse = {
    success: false,
    error: {
      code: err.code,
      message: err.message,
      details: (err as any).details,
    },
  };
}

```

```

};

const headers = { ...DEFAULT_HEADERS };

if ('retryAfter' in err && err.retryAfter) {
  headers['Retry-After'] = String(err.retryAfter);
}

return {
  statusCode: err.statusCode,
  headers,
  body: JSON.stringify(response),
};
}

/**
 * Handle errors uniformly
 */
export function handleError(
  err: unknown,
  logger?: Logger
): APIGatewayProxyResult {
  const appError = toAppError(err);
  return error(appError, logger);
}

/**
 * Create streaming response headers
 */
export function streamingHeaders(): Record<string, string> {
  return {
    ...DEFAULT_HEADERS,
    'Content-Type': 'text/event-stream',
    'Cache-Control': 'no-cache',
    Connection: 'keep-alive',
  };
}

/**
 * Format Server-Sent Event
 */
export function formatSSE(data: unknown, event?: string): string {
  let output = '';
  if (event) {
    output += `event: ${event}\n`;
  }
  output += `data: ${JSON.stringify(data)}\n\n`;
  return output;
}

```

```
}
```

```
packages/infrastructure/lambda/shared/auth.ts
```

```
/**  
 * Authentication and authorization utilities  
 */  
  
import type { APIGatewayProxyEvent } from 'aws-lambda';  
import { UnauthorizedError, ForbiddenError } from './errors';  
import { Logger } from './logger';  
  
/**  
 * Enhanced AuthContext with app isolation (v4.6.0)  
 */  
export interface AuthContext {  
    // Identity  
    userId: string;           // Cognito sub  
    appUserId: string;        // App-scoped user ID from app_users table  
    tenantId: string;  
    appId: string;            // Application identifier (thinktank, launchboard, etc.)  
    email: string;  
  
    // Roles & permissions  
    roles: string[];  
    groups: string[];  
    isAdmin: boolean;  
    isSuperAdmin: boolean;  
  
    // Session  
    sessionId?: string;  
    tokenExpiry: number;  
}  
  
export interface TokenClaims {  
    sub: string;  
    email?: string;  
    'cognito:username'?: string;  
    'cognito:groups'?: string[];  
    'custom:tenantId'?: string;  
    'custom:tenant_id'?: string;  
    'custom:appId'?: string;  
    'custom:app_id'?: string;  
    'custom:appUserId'?: string;  
    'custom:app_user_id'?: string;  
    'custom:role'?: string;  
    iss: string;  
    aud: string;
```

```

    exp: number;
    iat: number;
}

/**
 * Extract and validate authentication context from API Gateway event
 * Includes app isolation validation (v4.6.0)
 */
export function extractAuthContext(event: APIGatewayProxyEvent): AuthContext {
    const claims = event.requestContext.authorizer?.claims as TokenClaims | undefined;

    if (!claims) {
        throw new UnauthorizedError('No authentication claims found');
    }

    // Check token expiration
    if (claims.exp && claims.exp < Date.now() / 1000) {
        throw new UnauthorizedError('Token has expired');
    }

    // Extract core identifiers
    const userId = claims.sub;
    const tenantId = claims['custom:tenantId'] || claims['custom:tenant_id'];
    const appId = claims['custom:appId'] || claims['custom:app_id'];
    const appUserId = claims['custom:appUserId'] || claims['custom:app_user_id'];
    const email = claims.email || claims['cognito:username'] || '';

    // Validate required claims
    if (!userId) throw new UnauthorizedError('Missing user ID');
    if (!tenantId) throw new UnauthorizedError('Missing tenant ID');

    // App isolation - appId and appUserId required for v4.6.0+
    // Fallback for backward compatibility with pre-v4.6.0 tokens
    const resolvedAppId = appId || extractAppIdFromRoute(event) || 'default';
    const resolvedAppUserId = appUserId || userId; // Fallback to userId for legacy tokens

    // Validate app_id matches route (defense in depth)
    const routeAppId = extractAppIdFromRoute(event);
    if (routeAppId && appId && routeAppId !== appId) {
        throw new ForbiddenError(`Token app_id (${appId}) does not match route (${routeAppId})`);
    }

    // Extract roles and groups
    const groups = claims['cognito:groups'] || [];
    const roles = claims['custom:role'] ? [claims['custom:role']] : [];

    const isAdmin = groups.some(g =>
        ['super_admin', 'admin', 'operator', 'auditor'].includes(g)
    )
}

```

```

);
const isSuperAdmin = groups.includes('super_admin');

return {
  userId,
  appUserId: resolvedAppUserId,
  tenantId,
  appId: resolvedAppId,
  email,
  roles,
  groups,
  isAdmin,
  isSuperAdmin,
  tokenExpiry: claims.exp || 0,
};
}

/**
 * Extract app_id from route for validation (v4.6.0)
 */
function extractAppIdFromRoute(event: APIGatewayProxyEvent): string | null {
  // Extract from subdomain: thinktank.domain.com -> thinktank
  const host = event.headers.Host || event.headers['host'];
  if (host) {
    const subdomain = host.split('.')[0];
    if(['thinktank', 'launchboard', 'alwaysme', 'mechanicalmaker'].includes(subdomain)) {
      return subdomain;
    }
  }
}

// Extract from path: /api/thinktank/... -> thinktank
const pathMatch = event.path.match(/^\/api\/(thinktank|launchboard|alwaysme|mechanicalmaker)/);
if (pathMatch) {
  return pathMatch[1];
}

return null;
}

/**
 * Require specific roles
 */
export function requireRoles(auth: AuthContext, requiredRoles: string[]): void {
  const hasRole = requiredRoles.some(role =>
    auth.roles.includes(role) || auth.groups.includes(role)
  );

  if (!hasRole) {

```

```

        throw new ForbiddenError(
            `Required roles: ${requiredRoles.join(', ')}`)
    );
}
}

/** 
 * Require admin access
 */
export function requireAdmin(auth: AuthContext): void {
    if (!auth.isAdmin) {
        throw new ForbiddenError('Admin access required');
    }
}

/** 
 * Require super admin access
 */
export function requireSuperAdmin(auth: AuthContext): void {
    if (!auth.groups.includes('super_admin')) {
        throw new ForbiddenError('Super admin access required');
    }
}

/** 
 * Check if user can access tenant
 */
export function canAccessTenant(auth: AuthContext, tenantId: string): boolean {
    // Super admins can access any tenant
    if (auth.groups.includes('super_admin')) {
        return true;
    }

    // Users can only access their own tenant
    return auth.tenantId === tenantId;
}

/** 
 * Require tenant access
 */
export function requireTenantAccess(auth: AuthContext, tenantId: string): void {
    if (!canAccessTenant(auth, tenantId)) {
        throw new ForbiddenError('Access to tenant denied');
    }
}

/** 
 * Extract API key from header (for API key auth)

```

```

/*
export function extractApiKey(event: APIGatewayProxyEvent): string | undefined {
  return event.headers['X-Api-Key'] || event.headers['x-api-key'];
}

/**
 * Log authentication context (sanitized)
 */
export function logAuthContext(auth: AuthContext, logger: Logger): void {
  logger.info('Authenticated request', {
    userId: auth.userId,
    tenantId: auth.tenantId,
    isAdmin: auth.isAdmin,
    groupCount: auth.groups.length,
  });
}


```

PART 3: DATABASE CLIENT

packages/infrastructure/lambda/shared/db/index.ts

```

export * from './client';
export * from './queries';
export * from './types';


```

packages/infrastructure/lambda/shared/db/types.ts

```

/**
 * Database types matching Aurora PostgreSQL schema
 */

// =====
// PROVIDERS
// =====

export interface DBProvider {
  id: string;
  name: string;
  type: 'external' | 'self-hosted' | 'mid-tier';
  status: 'active' | 'inactive' | 'deprecated';
  hipaa_compliant: boolean;
  baa_available: boolean;
  base_url: string | null;
  auth_type: 'api_key' | 'oauth' | 'iam' | 'none';
  capabilities: string[];
  config: Record<string, unknown>;
  created_at: string;
}


```

```

    updated_at: string;
}

// =====
// MODELS
// =====

export interface DBModel {
  id: string;
  provider_id: string;
  name: string;
  display_name: string;
  description: string | null;
  type: 'external' | 'self-hosted';
  specialty: string;
  capabilities: string[];
  context_window: number | null;
  max_output_tokens: number | null;
  supports_functions: boolean;
  supports_vision: boolean;
  supports_streaming: boolean;
  has_thinking_mode: boolean;
  thinking_budget_tokens: number | null;
  pricing: DBModelPricing;
  thermal_state: string | null;
  thermal_config: DBThermalConfig | null;
  status: 'active' | 'inactive' | 'deprecated' | 'coming_soon';
  release_date: string | null;
  deprecation_date: string | null;
  created_at: string;
  updated_at: string;
}

export interface DBModelPricing {
  input_tokens?: number;
  output_tokens?: number;
  per_image?: number;
  per_minute_audio?: number;
  per_minute_video?: number;
  per_3d_model?: number;
  billed_markup: number;
}

export interface DBThermalConfig {
  state: 'OFF' | 'COLD' | 'WARM' | 'HOT' | 'AUTOMATIC';
  min_instances: number;
  max_instances: number;
  scale_to_zero_after_minutes?: number;
}

```

```

    warmup_time_seconds?: number;
}

// =====
// TENANTS
// =====

export interface DBTenant {
    id: string;
    app_id: string;
    name: string;
    domain: string | null;
    status: 'active' | 'suspended' | 'deleted';
    settings: DBTenantSettings;
    phi_config: DBPhiConfig | null;
    created_at: string;
    updated_at: string;
}

export interface DBTenantSettings {
    default_model?: string;
    allowed_providers?: string[];
    max_tokens_per_request?: number;
    rate_limit?: {
        requests_per_minute: number;
        tokens_per_day: number;
    };
}

export interface DBPhiConfig {
    mode: 'auto' | 'manual' | 'disabled';
    categories: Record<string, boolean>;
    reidentification: {
        allowed: boolean;
        requires_approval: boolean;
        mapping_ttl_hours: number;
    };
}

// =====
// USAGE
// =====

export interface DBUsageRecord {
    id: string;
    tenant_id: string;
    user_id: string | null;
    session_id: string | null;
}

```

```

model_id: string;
provider_id: string;
input_tokens: number;
output_tokens: number;
total_tokens: number;
cost: number;
billed_amount: number;
request_type: string;
latency_ms: number;
created_at: string;
}

// =====
// AUDIT LOG
// =====

export interface DBAuditLog {
  id: string;
  tenant_id: string;
  user_id: string | null;
  admin_id: string | null;
  action: string;
  resource_type: string;
  resource_id: string | null;
  details: Record<string, unknown>;
  ip_address: string | null;
  user_agent: string | null;
  created_at: string;
}

```

packages/infrastructure/lambda/shared/db/client.ts

```

/**
 * Aurora PostgreSQL client using Data API
 */

import {
  RDSDataClient,
  ExecuteStatementCommand,
  BatchExecuteStatementCommand,
  BeginTransactionCommand,
  CommitTransactionCommand,
  RollbackTransactionCommand,
  Field,
  SqlParameter,
} from '@aws-sdk/client-rds-data';
import {
  SecretsManagerClient,

```

```

GetSecretValueCommand,
} from '@aws-sdk/client-secrets-manager';
import { getConfig } from '../config';
import { Logger } from '../logger';
import { InternalError } from '../errors';

// Initialize clients
const rdsClient = new RDSDataClient({});
const secretsClient = new SecretsManagerClient({});

// Cache database credentials
let dbCredentials: { username: string; password: string } | null = null;

export interface QueryResult<T = Record<string, unknown>> {
  rows: T[];
  rowCount: number;
  columnMetadata?: { name: string; type: string }[];
}

export interface TransactionContext {
  transactionId: string;
}

/**
 * Get database credentials from Secrets Manager
 */
async function getDbCredentials(): Promise<{ username: string; password: string }> {
  if (dbCredentials) return dbCredentials;

  const config = getConfig();

  try {
    const command = new GetSecretValueCommand({
      SecretId: config.AURORA_SECRET_ARN,
    });

    const response = await secretsClient.send(command);

    if (!response.SecretString) {
      throw new Error('Secret value is empty');
    }

    dbCredentials = JSON.parse(response.SecretString);
    return dbCredentials!;
  } catch (error) {
    throw new InternalError(`Failed to retrieve database credentials: ${error}`);
  }
}

```

```

/**
 * Convert JavaScript value to SQL parameter
 */
function toSqlParameter(name: string, value: unknown): SqlParameter {
    if (value === null || value === undefined) {
        return { name, value: {isNull: true} };
    }

    if (typeof value === 'string') {
        return { name, value: {stringValue: value} };
    }

    if (typeof value === 'number') {
        if (Number.isInteger(value)) {
            return { name, value: {longValue: value} };
        }
        return { name, value: {doubleValue: value} };
    }

    if (typeof value === 'boolean') {
        return { name, value: {booleanValue: value} };
    }

    if (Array.isArray(value)) {
        return { name, value: {stringValue: JSON.stringify(value)}, typeHint: 'JSON' };
    }

    if (typeof value === 'object') {
        return { name, value: {stringValue: JSON.stringify(value)}, typeHint: 'JSON' };
    }

    return { name, value: {stringValue: String(value)} };
}

/**
 * Convert SQL field to JavaScript value
 */
function fromSqlField(field: Field): unknown {
    if (field.isNull) return null;
    if (field.stringValue !== undefined) return field.stringValue;
    if (field.longValue !== undefined) return field.longValue;
    if (field.doubleValue !== undefined) return field.doubleValue;
    if (field.booleanValue !== undefined) return field.booleanValue;
    if (field.blobValue !== undefined) return field.blobValue;
    if (field.arrayValue !== undefined) {
        return field.arrayValue.stringValues ||
               field.arrayValue.longValues ||

```

```

        field.arrayValue.doubleValues ||  

        field.arrayValue.booleanValues ||  

        [];  

    }  

    return null;  

}  
  

/**  

 * Execute a SQL query  

 */  

export async function query<T = Record<string, unknown>>(  

    sql: string,  

    params: Record<string, unknown> = {},  

    logger?: Logger,  

    transactionId?: string  

): Promise<QueryResult<T>> {  

    const config = getConfig();  

    const startTime = Date.now();  
  

    try {  

        const command = new ExecuteStatementCommand({  

            resourceArn: config.AURORA_CLUSTER_ARN,  

            secretArn: config.AURORA_SECRET_ARN,  

            database: 'radiant',  

            sql,  

            parameters: Object.entries(params).map(([name, value]) =>  

                toSqlParameter(name, value)
            ),  

            includeResultMetadata: true,
            transactionId,
        });
    }  
  

    const response = await rdsClient.send(command);
    const duration = Date.now() - startTime;  
  

    if (logger) {
        logger.debug('Database query executed', {
            sql: sql.substring(0, 100),
            duration,
            rowCount: response.numberOfRecordsUpdated,
        });
    }
}  
  

// Convert response to rows
const rows: T[] = [];  
  

if (response.records && response.columnMetadata) {
    for (const record of response.records) {

```

```

        const row: Record<string, unknown> = {};

        for (let i = 0; i < response.columnMetadata.length; i++) {
            const columnName = response.columnMetadata[i].name || `col${i}`;
            row[columnName] = fromSqlField(record[i]);
        }

        rows.push(row as T);
    }
}

return {
    rows,
    rowCount: response.numberOfRecordsUpdated || rows.length,
    columnMetadata: response.columnMetadata?.map(col => ({
        name: col.name || '',
        type: col.typeName || '',
    })),
};

} catch (error) {
    const duration = Date.now() - startTime;

    if (logger) {
        logger.error('Database query failed', error as Error, {
            sql: sql.substring(0, 100),
            duration,
        });
    }

    throw new InternalError(`Database query failed: ${error}`);
}
}

/**
 * Execute a single row query
 */
export async function queryOne<T = Record<string, unknown>>(
    sql: string,
    params: Record<string, unknown> = {},
    logger?: Logger
): Promise<T | null> {
    const result = await query<T>(sql, params, logger);
    return result.rows[0] || null;
}

/**
 * Begin a transaction
 */

```

```

export async function beginTransaction(logger?: Logger): Promise<TransactionContext> {
  const config = getConfig();

  try {
    const command = new BeginTransactionCommand({
      resourceArn: config.AURORA_CLUSTER_ARN,
      secretArn: config.AURORA_SECRET_ARN,
      database: 'radianit',
    });

    const response = await rdsClient.send(command);

    if (!response.transactionId) {
      throw new Error('No transaction ID returned');
    }

    if (logger) {
      logger.debug('Transaction started', { transactionId: response.transactionId });
    }

    return { transactionId: response.transactionId };
  } catch (error) {
    throw new InternalError(`Failed to begin transaction: ${error}`);
  }
}

/**
 * Commit a transaction
 */
export async function commitTransaction(
  ctx: TransactionContext,
  logger?: Logger
): Promise<void> {
  const config = getConfig();

  try {
    const command = new CommitTransactionCommand({
      resourceArn: config.AURORA_CLUSTER_ARN,
      secretArn: config.AURORA_SECRET_ARN,
      transactionId: ctx.transactionId,
    });

    await rdsClient.send(command);

    if (logger) {
      logger.debug('Transaction committed', { transactionId: ctx.transactionId });
    }
  } catch (error) {

```

```

        throw new InternalError(`Failed to commit transaction: ${error}`);
    }
}

/** 
 * Rollback a transaction
 */
export async function rollbackTransaction(
    ctx: TransactionContext,
    logger?: Logger
): Promise<void> {
    const config = getConfig();

    try {
        const command = new RollbackTransactionCommand({
            resourceArn: config.AURORA_CLUSTER_ARN,
            secretArn: config.AURORA_SECRET_ARN,
            transactionId: ctx.transactionId,
        });

        await rdsClient.send(command);

        if (logger) {
            logger.debug('Transaction rolled back', { transactionId: ctx.transactionId });
        }
    } catch (error) {
        // Log but don't throw - rollback errors are usually not critical
        if (logger) {
            logger.warn('Transaction rollback failed', {
                transactionId: ctx.transactionId,
                error: String(error),
            });
        }
    }
}

/** 
 * Execute a callback within a transaction
 */
export async function withTransaction<T>(
    callback: (transactionId: string) => Promise<T>,
    logger?: Logger
): Promise<T> {
    const ctx = await beginTransaction(logger);

    try {
        const result = await callback(ctx.transactionId);
        await commitTransaction(ctx, logger);
    }
}

```

```

        return result;
    } catch (error) {
        await rollbackTransaction(ctx, logger);
        throw error;
    }
}

```

packages/infrastructure/lambda/shared/db/queries.ts

```

/*
 * Pre-built database queries
 */

import { query, queryOne, QueryResult } from './client';
import { DBProvider, DBModel, DBTenant, DBUsageRecord, DBAuditLog } from './types';
import { Logger } from '../logger';
import { NotFoundError } from '../errors';

// =====
// PROVIDERS
// =====

export async function getProviders(
    filters: {
        type?: string;
        status?: string;
        hipaaCompliant?: boolean;
        limit?: number;
        offset?: number;
    } = {},
    logger?: Logger
): Promise<QueryResult<DBProvider>> {
    let sql = `
        SELECT * FROM providers
        WHERE 1=1
    `;
    const params: Record<string, unknown> = {};

    if (filters.type) {
        sql += ` AND type = :type`;
        params.type = filters.type;
    }

    if (filters.status) {
        sql += ` AND status = :status`;
        params.status = filters.status;
    }
}

```

```

if (filters.hipaaCompliant !== undefined) {
  sql += ` AND hipaa_compliant = :hipaaCompliant`;
  params.hipaaCompliant = filters.hipaaCompliant;
}

sql += ` ORDER BY name ASC`;

if (filters.limit) {
  sql += ` LIMIT :limit`;
  params.limit = filters.limit;
}

if (filters.offset) {
  sql += ` OFFSET :offset`;
  params.offset = filters.offset;
}

return query<DBProvider>(sql, params, logger);
}

export async function getProviderById(
  id: string,
  logger?: Logger
): Promise<DBProvider> {
  const result = await queryOne<DBProvider>(
    `SELECT * FROM providers WHERE id = :id`,
    { id },
    logger
  );

  if (!result) {
    throw new NotFoundError(`Provider ${id}`);
  }

  return result;
}

export async function updateProvider(
  id: string,
  updates: Partial<Pick<DBProvider, 'status' | 'hipaa_compliant' | 'config'>>,
  logger?: Logger
): Promise<DBProvider> {
  const setClauses: string[] = ['updated_at = NOW()'];
  const params: Record<string, unknown> = { id };

  if (updates.status !== undefined) {
    setClauses.push('status = :status');
    params.status = updates.status;
  }
}

```

```

}

if (updates.hipaa_compliant !== undefined) {
  setClauses.push('hipaa_compliant = :hipaaCompliant');
  params.hipaaCompliant = updates.hipaa_compliant;
}

if (updates.config !== undefined) {
  setClauses.push('config = :config');
  params.config = updates.config;
}

const sql = `
  UPDATE providers
  SET ${setClauses.join(', ')}
  WHERE id = :id
  RETURNING *
`;

const result = await queryOne<DBProvider>(sql, params, logger);

if (!result) {
  throw new NotFoundError(`Provider ${id}`);
}

return result;
}

// =====
// MODELS
// =====

export async function getModels(
  filters: {
    providerId?: string;
    specialty?: string;
    status?: string;
    type?: string;
    supportsVision?: boolean;
    supportsStreaming?: boolean;
    limit?: number;
    offset?: number;
  } = {},
  logger?: Logger
): Promise<QueryResult<DBModel>> {
  let sql = `
    SELECT * FROM models
    WHERE 1=1

```

```

`;

const params: Record<string, unknown> = {};

if (filters.providerId) {
  sql += ` AND provider_id = :providerId`;
  params.providerId = filters.providerId;
}

if (filters.specialty) {
  sql += ` AND specialty = :specialty`;
  params.specialty = filters.specialty;
}

if (filters.status) {
  sql += ` AND status = :status`;
  params.status = filters.status;
}

if (filters.type) {
  sql += ` AND type = :type`;
  params.type = filters.type;
}

if (filters.supportsVision !== undefined) {
  sql += ` AND supports_vision = :supportsVision`;
  params.supportsVision = filters.supportsVision;
}

if (filters.supportsStreaming !== undefined) {
  sql += ` AND supports_streaming = :supportsStreaming`;
  params.supportsStreaming = filters.supportsStreaming;
}

sql += ` ORDER BY display_name ASC`;

if (filters.limit) {
  sql += ` LIMIT :limit`;
  params.limit = filters.limit;
}

if (filters.offset) {
  sql += ` OFFSET :offset`;
  params.offset = filters.offset;
}

return query<DBModel>(sql, params, logger);
}

```

```

export async function getModelById(
  id: string,
  logger?: Logger
): Promise<DBModel> {
  const result = await queryOne<DBModel>(
    `SELECT * FROM models WHERE id = :id`,
    { id },
    logger
  );

  if (!result) {
    throw new NotFoundError(`Model ${id}`);
  }

  return result;
}

export async function getModelByName(
  name: string,
  logger?: Logger
): Promise<DBModel | null> {
  return queryOne<DBModel>(
    `SELECT * FROM models WHERE name = :name AND status = 'active'`,
    { name },
    logger
  );
}

export async function updateModel(
  id: string,
  updates: Partial<Pick<DBModel, 'status' | 'thermal_state' | 'thermal_config' | 'display_name'>>,
  logger?: Logger
): Promise<DBModel> {
  const setClauses: string[] = ['updated_at = NOW()'];
  const params: Record<string, unknown> = { id };

  if (updates.status !== undefined) {
    setClauses.push('status = :status');
    params.status = updates.status;
  }

  if (updates.thermal_state !== undefined) {
    setClauses.push('thermal_state = :thermalState');
    params.thermalState = updates.thermal_state;
  }

  if (updates.thermal_config !== undefined) {
    setClauses.push('thermal_config = :thermalConfig');
    params.thermalConfig = updates.thermal_config;
  }
}

```

```

    params.thermalConfig = updates.thermal_config;
}

if (updates.display_name !== undefined) {
  setClauses.push('display_name = :displayName');
  params.displayName = updates.display_name;
}

if (updates.description !== undefined) {
  setClauses.push('description = :description');
  params.description = updates.description;
}

const sql = `
UPDATE models
SET ${setClauses.join(', ')}
WHERE id = :id
RETURNING *
`;

const result = await queryOne<DBModel>(sql, params, logger);

if (!result) {
  throw new NotFoundError(`Model ${id}`);
}

return result;
}

// =====
// TENANTS
// =====

export async function getTenantById(
  id: string,
  logger?: Logger
): Promise<DBTenant> {
  const result = await queryOne<DBTenant>(
    `SELECT * FROM tenants WHERE id = :id AND status = 'active'`,
    { id },
    logger
  );

  if (!result) {
    throw new NotFoundError(`Tenant ${id}`);
  }

  return result;
}

```

```

}

export async function getTenantPhiConfig(
  tenantId: string,
  logger?: Logger
): Promise<DBTenant['phi_config']> {
  const tenant = await getTenantById(tenantId, logger);
  return tenant.phi_config;
}

// =====
// USAGE
// =====

export async function recordUsage(
  usage: Omit<DBUsageRecord, 'id' | 'created_at'>,
  logger?: Logger
): Promise<DBUsageRecord> {
  const sql = `
    INSERT INTO usage_records (
      id, tenant_id, user_id, session_id, model_id, provider_id,
      input_tokens, output_tokens, total_tokens, cost, billed_amount,
      request_type, latency_ms
    ) VALUES (
      gen_random_uuid(), :tenantId, :userId, :sessionId, :modelId, :providerId,
      :inputTokens, :outputTokens, :totalTokens, :cost, :billedAmount,
      :requestType, :latencyMs
    )
    RETURNING *
  `;
  const result = await queryOne<DBUsageRecord>(sql, {
    tenantId: usage.tenant_id,
    userId: usage.user_id,
    sessionId: usage.session_id,
    modelId: usage.model_id,
    providerId: usage.provider_id,
    inputTokens: usage.input_tokens,
    outputTokens: usage.output_tokens,
    totalTokens: usage.total_tokens,
    cost: usage.cost,
    billedAmount: usage.billed_amount,
    requestType: usage.request_type,
    latencyMs: usage.latency_ms,
  }, logger);

  return result!;
}

```

```

// =====
// AUDIT LOG
// =====

export async function createAuditLog(
  log: Omit<DBAuditLog, 'id' | 'created_at'>,
  logger?: Logger
): Promise<DBAuditLog> {
  const sql = `
    INSERT INTO audit_logs (
      id, tenant_id, user_id, admin_id, action, resource_type,
      resource_id, details, ip_address, user_agent
    ) VALUES (
      gen_random_uuid(), :tenantId, :userId, :adminId, :action, :resourceType,
      :resourceId, :details, :ipAddress, :userAgent
    )
    RETURNING *
  `;
  const result = await queryOne<DBAuditLog>(sql, {
    tenantId: log.tenant_id,
    userId: log.user_id,
    adminId: log.admin_id,
    action: log.action,
    resourceType: log.resource_type,
    resourceId: log.resource_id,
    details: log.details,
    ipAddress: log.ip_address,
    userAgent: log.user_agent,
  }, logger);

  return result!;
}

```

PART 4: LITELLM CLIENT

packages/infrastructure/lambda/shared/litellm/index.ts

```

export * from './client';
export * from './types';

```

packages/infrastructure/lambda/shared/litellm/types.ts

```

/**
 * LitellM API types
 */

```

```

// =====
// CHAT COMPLETION
// =====

export interface ChatCompletionRequest {
    model: string;
    messages: ChatMessage[];
    max_tokens?: number;
    temperature?: number;
    top_p?: number;
    stream?: boolean;
    stop?: string | string[];
    presence_penalty?: number;
    frequency_penalty?: number;
    user?: string;
    metadata?: Record<string, unknown>;
}

export interface ChatMessage {
    role: 'system' | 'user' | 'assistant' | 'function' | 'tool';
    content: string | ContentPart[];
    name?: string;
    function_call?: FunctionCall;
    tool_calls?: ToolCall[];
}

export interface ContentPart {
    type: 'text' | 'image_url';
    text?: string;
    image_url?: {
        url: string;
        detail?: 'low' | 'high' | 'auto';
    };
}

export interface FunctionCall {
    name: string;
    arguments: string;
}

export interface ToolCall {
    id: string;
    type: 'function';
    function: FunctionCall;
}

export interface ChatCompletionResponse {
}

```

```

id: string;
object: 'chat.completion';
created: number;
model: string;
choices: ChatChoice[];
usage: TokenUsage;
system_fingerprint?: string;
}

export interface ChatChoice {
index: number;
message: ChatMessage;
finish_reason: 'stop' | 'length' | 'function_call' | 'tool_calls' | 'content_filter' | null;
logprobs?: unknown;
}

export interface TokenUsage {
prompt_tokens: number;
completion_tokens: number;
total_tokens: number;
}

// =====
// STREAMING
// =====

export interface ChatCompletionChunk {
id: string;
object: 'chat.completion.chunk';
created: number;
model: string;
choices: StreamChoice[];
usage?: TokenUsage;
}

export interface StreamChoice {
index: number;
delta: {
role?: string;
content?: string;
function_call?: Partial<FunctionCall>;
tool_calls?: Partial<ToolCall>[];
};
finish_reason: string | null;
}

// =====
// EMBEDDINGS

```

```

// =====

export interface EmbeddingRequest {
  model: string;
  input: string | string[];
  encoding_format?: 'float' | 'base64';
  dimensions?: number;
  user?: string;
}

export interface EmbeddingResponse {
  object: 'list';
  data: EmbeddingData[];
  model: string;
  usage: {
    prompt_tokens: number;
    total_tokens: number;
  };
}

export interface EmbeddingData {
  object: 'embedding';
  index: number;
  embedding: number[];
}

// =====
// MODELS
// =====

export interface LiteLLMModel {
  id: string;
  object: 'model';
  created: number;
  owned_by: string;
}

export interface LiteLLMModelList {
  object: 'list';
  data: LiteLLMModel[];
}

// =====
// HEALTH
// =====

export interface HealthResponse {
  status: 'healthy' | 'unhealthy';
}

```

```

    version?: string;
    models?: string[] ;
}

// =====
// ERRORS
// =====

export interface LiteLLMError {
  error: {
    message: string;
    type: string;
    param?: string;
    code?: string;
  };
}

```

packages/infrastructure/lambda/shared/litellm/client.ts

```

/*
 * LiteLLM HTTP client
 */

import { getConfig } from '../config';
import { Logger } from '../logger';
import { ProviderError, ServiceUnavailableError, RateLimitError } from '../errors';
import {
  ChatCompletionRequest,
  ChatCompletionResponse,
  ChatCompletionChunk,
  EmbeddingRequest,
  EmbeddingResponse,
  LiteLLMModelList,
  HealthResponse,
  LiteLLMError,
} from './types';

const DEFAULT_TIMEOUT = 120000; // 2 minutes for AI requests

interface FetchOptions {
  method?: string;
  body?: unknown;
  timeout?: number;
  stream?: boolean;
}

/*
 * Make HTTP request to LiteLLM

```

```

*/
async function fetchLiteLLM<T>(
  path: string,
  options: FetchOptions = {},
  logger?: Logger
): Promise<T> {
  const config = getConfig();
  const url = `${config.LITELLM_URL}${path}`;
  const timeout = options.timeout || DEFAULT_TIMEOUT;

  const controller = new AbortController();
  const timeoutId = setTimeout(() => controller.abort(), timeout);

  try {
    const startTime = Date.now();

    const response = await fetch(url, {
      method: options.method || 'GET',
      headers: {
        'Content-Type': 'application/json',
        Accept: options.stream ? 'text/event-stream' : 'application/json',
      },
      body: options.body ? JSON.stringify(options.body) : undefined,
      signal: controller.signal,
    });

    const duration = Date.now() - startTime;

    if (logger) {
      logger.debug('LiteLLM request completed', {
        path,
        status: response.status,
        duration,
      });
    }
  }

  if (!response.ok) {
    const errorBody = await response.text();
    let errorMessage = `LiteLLM error: ${response.statusText}`;

    try {
      const errorJson = JSON.parse(errorBody) as LiteLLMError;
      errorMessage = errorJson.error?.message || errorMessage;
    } catch {
      errorMessage = errorBody || errorMessage;
    }

    if (response.status === 429) {

```

```

        const retryAfter = parseInt(response.headers.get('Retry-After') || '60');
        throw new RateLimitError(retryAfter);
    }

    if (response.status >= 500) {
        throw new ServiceUnavailableError(errorMessage);
    }

    throw new ProviderError(errorMessage, 'litellm');
}

if (options.stream) {
    return response as unknown as T;
}

return await response.json() as T;
} catch (error) {
    if (error instanceof Error && error.name === 'AbortError') {
        throw new ServiceUnavailableError('LiteLLM request timed out');
    }

    if (error instanceof ProviderError ||
        error instanceof ServiceUnavailableError ||
        error instanceof RateLimitError) {
        throw error;
    }

    throw new ServiceUnavailableError(`LiteLLM connection failed: ${error}`);
} finally {
    clearTimeout(timeoutId);
}
}

/**
 * Create chat completion
 */
export async function createChatCompletion(
    request: ChatCompletionRequest,
    logger?: Logger
): Promise<ChatCompletionResponse> {
    return fetchLiteLLM<ChatCompletionResponse>(
        '/v1/chat/completions',
        {
            method: 'POST',
            body: request,
        },
        logger
    );
}

```

```

}

/**
 * Create streaming chat completion
 */
export async function* streamChatCompletion(
  request: ChatCompletionRequest,
  logger?: Logger
): AsyncGenerator<ChatCompletionChunk> {
  const streamRequest = { ...request, stream: true };

  const response = await fetchLiteLLM<Response>(
    '/v1/chat/completions',
    {
      method: 'POST',
      body: streamRequest,
      stream: true,
    },
    logger
  );

  const reader = response.body?.getReader();
  if (!reader) {
    throw new ServiceUnavailableError('No response body for streaming');
  }

  const decoder = new TextDecoder();
  let buffer = '';

  try {
    while (true) {
      const { done, value } = await reader.read();

      if (done) break;

      buffer += decoder.decode(value, { stream: true });
      const lines = buffer.split('\n');
      buffer = lines.pop() || '';

      for (const line of lines) {
        const trimmed = line.trim();

        if (!trimmed || !trimmed.startsWith('data: ')) continue;

        const data = trimmed.slice(6);

        if (data === '[DONE]') {
          return;
        }
      }
    }
  } catch (err) {
    logger.error(`Error reading response: ${err}`);
  }
}

```

```

        }

        try {
            const chunk = JSON.parse(data) as ChatCompletionChunk;
            yield chunk;
        } catch {
            if (logger) {
                logger.warn('Failed to parse SSE chunk', { data });
            }
        }
    }
}

} finally {
    reader.releaseLock();
}
}

/** 
 * Create embeddings
 */
export async function createEmbedding(
    request: EmbeddingRequest,
    logger?: Logger
): Promise<EmbeddingResponse> {
    return fetchLiteLLM<EmbeddingResponse>(
        '/v1/embeddings',
        {
            method: 'POST',
            body: request,
        },
        logger
    );
}

/** 
 * List available models
 */
export async function listModels(logger?: Logger): Promise<LiteLLMModelList> {
    return fetchLiteLLM<LiteLLMModelList>('/v1/models', {}, logger);
}

/** 
 * Health check
 */
export async function checkHealth(logger?: Logger): Promise<HealthResponse> {
    try {
        const response = await fetchLiteLLM<HealthResponse>(
            '/health',

```

```

        { timeout: 5000 },
        logger
    );
    return response;
} catch {
    return { status: 'unhealthy' };
}
}

/**
 * Calculate cost for a completion
 */
export function calculateCost(
    usage: { prompt_tokens: number; completion_tokens: number },
    pricing: { input_tokens?: number; output_tokens?: number; billed_markup: number }
): { cost: number; billed: number } {
    const inputCost = (usage.prompt_tokens / 1_000_000) * (pricing.input_tokens || 0);
    const outputCost = (usage.completion_tokens / 1_000_000) * (pricing.output_tokens || 0);
    const cost = inputCost + outputCost;
    const billed = cost * (1 + pricing.billed_markup);

    return { cost, billed };
}

```

PART 5: PHI SANITIZATION

packages/infrastructure/lambda/shared/phi/index.ts

```

export * from './sanitizer';
export * from './patterns';
export * from './types';

```

packages/infrastructure/lambda/shared/phi/types.ts

```

/**
 * PHI (Protected Health Information) types
 */

```

export type PHICategory =

```

| 'NAME'
| 'SSN'
| 'DOB'
| 'ADDRESS'
| 'PHONE'
| 'EMAIL'
| 'DIAGNOSIS'
| 'TREATMENT'

```

```

| 'MEDICAL_RECORD'
| 'INSURANCE_ID';

export interface PHIConfig {
  mode: 'auto' | 'manual' | 'disabled';
  categories: Record<PHICategory, boolean>;
  reidentification: {
    allowed: boolean;
    requires_approval: boolean;
    mapping_ttl_hours: number;
  };
}

export interface PHIMatch {
  category: PHICategory;
  original: string;
  placeholder: string;
  startIndex: number;
  endIndex: number;
  confidence: number;
}

export interface SanitizationResult {
  sanitizedText: string;
  matches: PHIMatch[];
  mappingId: string;
}

export interface ReidentificationResult {
  originalText: string;
  matches: PHIMatch[];
}

export interface PHIMapping {
  id: string;
  tenant_id: string;
  session_id: string | null;
  mappings: Record<string, string>; // placeholder -> original
  created_at: string;
  expires_at: string;
}

export const DEFAULT_PHI_CONFIG: PHIConfig = {
  mode: 'auto',
  categories: {
    NAME: true,
    SSN: true,
    DOB: true,
  }
}

```

```

ADDRESS: true,
PHONE: true,
EMAIL: true,
DIAGNOSIS: false, // Often needed for AI analysis
TREATMENT: false, // Often needed for AI analysis
MEDICAL_RECORD: true,
INSURANCE_ID: true,
},
reidentification: {
    allowed: true,
    requires_approval: true,
    mapping_ttl_hours: 24,
},
};


```

packages/infrastructure/lambda/shared/phi/patterns.ts

```

/**
 * PHI detection patterns
 */

import { PHICategory } from './types';

export interface PHIPattern {
    category: PHICategory;
    patterns: RegExp[];
    validator?: (match: string) => boolean;
    confidence: number;
}

/**
 * PHI detection patterns by category
 */
export const PHI_PATTERNS: PHIPattern[] = [
    // Social Security Numbers
    {
        category: 'SSN',
        patterns: [
            /\b\d{3}-\d{2}-\d{4}\b/g,
            /\b\d{3}\s\d{2}\s\d{4}\b/g,
            /\bSSN[:\s]*\d{3}[-\s]?\d{2}[-\s]?\d{4}\b/gi,
        ],
        validator: (match) => {
            const digits = match.replace(/\D/g, '');
            if (digits.length !== 9) return false;
            // Invalid SSN patterns
            if (digits.startsWith('000') || digits.startsWith('666')) return false;
            if (digits.substring(0, 3) === '900' && parseInt(digits.substring(0, 3)) <= 999) return false;
        }
    }
];

```

```

        return true;
    },
    confidence: 0.95,
},

// Dates of Birth
{
category: 'DOB',
patterns: [
    /\b(?:DOB|Date of Birth|Birthday|Born) [:s]*(\d{1,2}[-/] \d{1,2}[-/] \d{2,4})\b/gi,
    /\b(?:DOB|Date of Birth|Birthday|Born) [:s]*([A-Z] [a-z]+ \s+ \d{1,2}, ?\s+ \d{4})\b/gi,
],
confidence: 0.90,
},

// Phone Numbers
{
category: 'PHONE',
patterns: [
    /\b\(?[2-9]\d{2}\)\?[-.\s]? \d{3}[-.\s]? \d{4}\b/g,
    /\b(?:Phone|Tel|Mobile|Cell) [:s]*\(?[2-9]\d{2}\)\?[-.\s]? \d{3}[-.\s]? \d{4}\b/gi,
    /\b\+1[-.\s]\?(\? [2-9]\d{2}\)\?[-.\s]\? \d{3}[-.\s]\? \d{4}\b/g,
],
confidence: 0.85,
},

// Email Addresses
{
category: 'EMAIL',
patterns: [
    /\b[A-Za-z0-9._%+-]+\@[A-Za-z0-9.-]+\.[A-Za-z]{2,}\b/g,
],
validator: (match) => {
    // Exclude common system emails
    const systemDomains = ['YOUR_DOMAIN.com', 'test.com', 'localhost'];
    const domain = match.split('@')[1].toLowerCase();
    return !systemDomains.includes(domain);
},
confidence: 0.90,
},

// Medical Record Numbers
{
category: 'MEDICAL_RECORD',
patterns: [
    /\b(?:MRN|Medical Record|Patient ID) [:s#]*([A-Z0-9]{6,12})\b/gi,
    /\bMRN [:s#]*\d{6,12}\b/gi,
],
}

```

```

confidence: 0.85,
},

// Insurance IDs
{
category: 'INSURANCE_ID',
patterns: [
  /\b(?:Insurance ID|Policy Number|Member ID) [:\s#]*([A-Z0-9]{8,15})\b/gi,
  /\b(?:Group|Plan)\s*(?:Number|ID|#) [:\s]*([A-Z0-9]{6,12})\b/gi,
],
confidence: 0.80,
},

// Street Addresses
{
category: 'ADDRESS',
patterns: [
  /\b\d{1,5}\s+[A-Za-z]+(?:\s+[A-Za-z]+)*\s+(?:Street|St|Avenue|Ave|Road|Rd|Boulevard|Blvd
  /\b(?:P\.?O\.?\s*Box|PO Box)\s+\d+\b/gi,
],
confidence: 0.75,
},

// Names (context-dependent)
{
category: 'NAME',
patterns: [
  /\b(?:Patient|Client|Name) [:\s]*([A-Z] [a-z]+(?:\s+[A-Z] [a-z]+)+)\b/gi,
  /\b(?:Dr\.|Doctor|Mr\.|Mrs\.|Ms\. )\s+([A-Z] [a-z]+(?:\s+[A-Z] [a-z]+)+*)\b/g,
],
confidence: 0.70,
},

// Diagnoses (ICD codes and common conditions)
{
category: 'DIAGNOSIS',
patterns: [
  /\b(?:ICD-?10) [:\s]*([A-Z] \d{2}(?:\.\d{1,4}))?\b/gi,
  /\b(?:Diagnosis|Dx) [:\s]*([A-Za-z\s]+(?:syndrome|disease|disorder|condition))\b/gi,
],
confidence: 0.80,
},

// Treatments
{
category: 'TREATMENT',
patterns: [
  /\b(?:Rx|Prescription|Medication) [:\s]*([A-Za-z]+(?:\s+\d+\s*mg))?\b/gi,
]
}

```

```

    /\b(?:Treatment|Procedure) [:\s]*([A-Za-z\s]+(?:ectomy|plasty|scopy|therapy))\b/gi,
],
confidence: 0.75,
},
];

```

```

/**
 * Get patterns for enabled categories
 */
export function getEnabledPatterns(
  categories: Record<PHICategory, boolean>
): PHI_PATTERN[] {
  return PHI_PATTERNS.filter(pattern => categories[pattern.category]);
}

```

packages/infrastructure/lambda/shared/phi/sanitizer.ts

```

/**
 * PHI sanitization engine
 */

import { v4 as uuid } from 'uuid';
import {
  PHIConfig,
  PHICategory,
  PHIMatch,
  SanitizationResult,
  ReidentificationResult,
  PHIMapping,
  DEFAULT_PHI_CONFIG,
} from './types';
import { getEnabledPatterns, PHI_PATTERN } from './patterns';
import { Logger } from '../logger';

// Placeholder format: [PHI_CATEGORY_INDEX]
const PLACEHOLDER_PREFIX = '[PHI_';
const PLACEHOLDER_SUFFIX = ']';

/**
 * Generate a placeholder for a PHI match
 */
function generatePlaceholder(category: PHICategory, index: number): string {
  return `${PLACEHOLDER_PREFIX}${category}_${index}${PLACEHOLDER_SUFFIX}`;
}

/**
 * Parse a placeholder back to its parts
*/

```

```

function parsePlaceholder(placeholder: string): { category: PHICategory; index: number } | null {
  const match = placeholder.match(/\b[PHI_([A-Z_]+)_(\d+)\]\b/);
  if (!match) return null;
  return {
    category: match[1] as PHICategory,
    index: parseInt(match[2]),
  };
}

/**
 * Detect PHI in text
 */
export function detectPHI(
  text: string,
  config: PHICConfig = DEFAULT_PHI_CONFIG,
  logger?: Logger
): PHIMatch[] {
  if (config.mode === 'disabled') {
    return [];
  }

  const matches: PHIMatch[] = [];
  const patterns = getEnabledPatterns(config.categories);
  const categoryCounters: Record<string, number> = {};

  for (const patternDef of patterns) {
    for (const regex of patternDef.patterns) {
      // Reset regex lastIndex
      regex.lastIndex = 0;
      let match: RegExpExecArray | null;

      while ((match = regex.exec(text)) !== null) {
        const original = match[1] || match[0];

        // Skip if validator fails
        if (patternDef.validator && !patternDef.validator(original)) {
          continue;
        }

        // Check for overlapping matches
        const startIndex = text.indexOf(original, match.index);
        const endIndex = startIndex + original.length;

        const isOverlapping = matches.some(
          m => (startIndex >= m.startIndex && startIndex < m.endIndex) ||
            (endIndex > m.startIndex && endIndex <= m.endIndex)
        );
      }
    }
  }
}

```

```

    if (isOverlapping) continue;

    // Generate unique placeholder
    categoryCounters[patternDef.category] = (categoryCounters[patternDef.category] || 0) +
      const placeholder = generatePlaceholder(
        patternDef.category,
        categoryCounters[patternDef.category]
      );

    matches.push({
      category: patternDef.category,
      original,
      placeholder,
      startIndex,
      endIndex,
      confidence: patternDef.confidence,
    });
  }
}

// Sort by start index (descending) for safe replacement
matches.sort((a, b) => b.startIndex - a.startIndex);

if (logger && matches.length > 0) {
  logger.info('PHI detected', {
    matchCount: matches.length,
    categories: [...new Set(matches.map(m => m.category))],
  });
}

return matches;
}

/**
 * Sanitize PHI in text
 */
export function sanitizePHI(
  text: string,
  config: PHICConfig = DEFAULT_PHI_CONFIG,
  logger?: Logger
): SanitizationResult {
  const matches = detectPHI(text, config, logger);

  if (matches.length === 0) {
    return {
      sanitizedText: text,
      matches: [],
    };
  }

  // Sanitize PHI
  const sanitizedText = matches.map(m => {
    const start = text.substring(0, m.startIndex);
    const end = text.substring(m.endIndex);
    const placeholder = generatePlaceholder(m.category);
    return start + placeholder + end;
  }).join('');

  return {
    sanitizedText,
    matches,
  };
}

```

```

        mappingId: '',
    };
}

let sanitizedText = text;

// Replace in reverse order to preserve indices
for (const match of matches) {
    sanitizedText =
        sanitizedText.substring(0, match.startIndex) +
        match.placeholder +
        sanitizedText.substring(match.endIndex);
}

// Generate mapping ID for re-identification
const mappingId = uuid();

return {
    sanitizedText,
    matches: matches.reverse(), // Return in original order
    mappingId,
};
}

/**
* Re-identify PHI in text
*/
export function reidentifyPHI(
    sanitizedText: string,
    mapping: PHIMapping,
    logger?: Logger
): ReidentificationResult {
    let originalText = sanitizedText;
    const matches: PHIMatch[] = [];

    // Find all placeholders in the text
    const placeholderRegex = /\b[PHI_][A-Z_]+\d+\b/g;
    let match: RegExpExecArray | null;

    while ((match = placeholderRegex.exec(sanitizedText)) !== null) {
        const placeholder = match[0];
        const original = mapping.mappings[placeholder];

        if (original) {
            const parsed = parsePlaceholder(placeholder);

            if (parsed) {
                matches.push({

```

```

        category: parsed.category,
        original,
        placeholder,
        startIndex: match.index,
        endIndex: match.index + placeholder.length,
        confidence: 1.0,
    });
}
}
}

// Replace in reverse order
matches.sort((a, b) => b.startIndex - a.startIndex);

for (const m of matches) {
    originalText =
        originalText.substring(0, m.startIndex) +
        m.original +
        originalText.substring(m.endIndex);
}

if (logger) {
    logger.info('PHI re-identified', {
        matchCount: matches.length,
    });
}

return {
    originalText,
    matches: matches.reverse(),
};
}

/**
 * Create a mapping record for storage
 */
export function createMappingRecord(
    tenantId: string,
    sessionId: string | null,
    result: SanitizationResult,
    ttlHours: number
): PHI_mapping {
    const mappings: Record<string, string> = {};

    for (const match of result.matches) {
        mappings[match.placeholder] = match.original;
    }
}

```

```

const now = new Date();
const expiresAt = new Date(now.getTime() + ttlHours * 60 * 60 * 1000);

return {
  id: result.mappingId,
  tenant_id: tenantId,
  session_id: sessionId,
  mappings,
  created_at: now.toISOString(),
  expires_at: expiresAt.toISOString(),
};
}

/**
 * Check if text contains any PHI
 */
export function containsPHI(
  text: string,
  config: PHICConfig = DEFAULT_PHI_CONFIG
): boolean {
  const matches = detectPHI(text, config);
  return matches.length > 0;
}

/**
 * Get PHI summary for audit logging
 */
export function getPHISummary(
  matches: PHIMatch[]
): Record<PHICCategory, number> {
  const summary: Record<string, number> = {};

  for (const match of matches) {
    summary[match.category] = (summary[match.category] || 0) + 1;
  }

  return summary as Record<PHICCategory, number>;
}

```

PART 6: ROUTER LAMBDA

packages/infrastructure/lambda/api/router.ts

```

/**
 * Router Lambda - Main API entry point
 *
 * Handles:

```

```

* - Health checks
* - Request routing
* - CORS preflight
* - Error handling
*/
```

```

import type {
  APIGatewayProxyEvent,
  APIGatewayProxyResult,
  Context,
} from 'aws-lambda';
import { Logger } from '../shared/logger';
import { getConfig } from '../shared/config';
import { success, handleError } from '../shared/response';
import { NotFoundError } from '../shared/errors';
import { checkHealth as checkLiteLLMHealth } from '../shared/litellm';
import { query } from '../shared/db';

// Initialize logger
const logger = new Logger({ handler: 'router' });

/**
 * Main handler
 */
export async function handler(
  event: APIGatewayProxyEvent,
  context: Context
): Promise<APIGatewayProxyResult> {
  // Set request context
  const requestLogger = logger.child({
    requestId: context.awsRequestId,
    path: event.path,
    method: event.httpMethod,
  });

  try {
    const config = getConfig();

    requestLogger.info('Request received', {
      queryParams: event.queryStringParameters,
      hasBody: !!event.body,
    });
  }

  // Route based on path
  const path = event.path.replace(/^\api\/v2/, '');

  switch (true) {
    // Health check

```

```

    case path === '/health' || path === '/':
      return await handleHealthCheck(requestLogger);

    // Ready check (deep health)
    case path === '/ready':
      return await handleReadyCheck(requestLogger);

    // Version info
    case path === '/version':
      return handleVersionInfo();

    // Metrics (admin only)
    case path === '/metrics':
      return await handleMetrics(event, requestLogger);

    default:
      throw new NotFoundError(`Route ${event.httpMethod} ${event.path}`);
  }
} catch (error) {
  return handleError(error, requestLogger);
}
}

/**
 * Basic health check
 */
async function handleHealthCheck(logger: Logger): Promise<APIGatewayProxyResult> {
  const config = getConfig();

  return success({
    status: 'healthy',
    timestamp: new Date().toISOString(),
    environment: config.ENVIRONMENT,
    tier: config.TIER,
  });
}

/**
 * Deep health check (ready probe)
 */
async function handleReadyCheck(logger: Logger): Promise<APIGatewayProxyResult> {
  const config = getConfig();
  const checks: Record<string, { status: string; latency?: number }> = {};
  const startTime = Date.now();

  // Check database
  try {
    const dbStart = Date.now();

```

```

    await query('SELECT 1', {}, logger);
    checks.database = {
      status: 'healthy',
      latency: Date.now() - dbStart,
    };
  } catch (error) {
    checks.database = { status: 'unhealthy' };
    logger.error('Database health check failed', error as Error);
  }

  // Check LiteLLM
  try {
    const llmStart = Date.now();
    const health = await checkLiteLLMHealth(logger);
    checks.litellm = {
      status: health.status,
      latency: Date.now() - llmStart,
    };
  } catch (error) {
    checks.litellm = { status: 'unhealthy' };
    logger.error('LiteLLM health check failed', error as Error);
  }

  // Determine overall status
  const allHealthy = Object.values(checks).every(c => c.status === 'healthy');
  const overallStatus = allHealthy ? 'ready' : 'degraded';

  logger.info('Ready check completed', {
    status: overallStatus,
    checks,
    totalLatency: Date.now() - startTime,
  });

  return success({
    status: overallStatus,
    timestamp: new Date().toISOString(),
    environment: config.ENVIRONMENT,
    tier: config.TIER,
    checks,
  }, allHealthy ? 200 : 503);
}

/**
 * Version information
 */
function handleVersionInfo(): APIGatewayProxyResult {
  const config = getConfig();

```

```

    return success({
      version: '2.2.0',
      appId: config.APP_ID,
      environment: config.ENVIRONMENT,
      tier: config.TIER,
      apiVersion: 'v2',
      buildDate: '2024-12',
    });
}

/**
 * Metrics endpoint (admin only)
 */
async function handleMetrics(
  event: APIGatewayProxyEvent,
  logger: Logger
): Promise<APIGatewayProxyResult> {
  // Metrics collection - use Section 12 MetricsCollector service
  // See Section 12.2 for implementation details
  // await metricsCollector.recordUsage({ ... });
  // This would aggregate data from CloudWatch, DynamoDB usage table, etc.

  return success({
    message: 'Metrics endpoint - coming in Prompt 5',
    timestamp: new Date().toISOString(),
  });
}

```

PART 7: CHAT LAMBDA

packages/infrastructure/lambda/api/chat.ts

```

/**
 * Chat Lambda - AI completions handler
 *
 * Handles:
 * - Chat completions via LiteLLM
 * - Streaming responses
 * - PHI sanitization
 * - Usage tracking
 */

import type {
  APIGatewayProxyEvent,
  APIGatewayProxyResult,
  Context,
} from 'aws-lambda';

```

```

import { z } from 'zod';
import { v4 as uuid } from 'uuid';
import { Logger } from '../shared/logger';
import { getConfig, getFeatureFlags } from '../shared/config';
import { success, handleError, formatSSE, streamingHeaders } from '../shared/response';
import { extractAuthContext, logAuthContext } from '../shared/auth';
import { ValidationError, NotFoundError } from '../shared/errors';
import {
  createChatCompletion,
  streamChatCompletion,
  calculateCost,
  ChatCompletionRequest,
  ChatMessage,
} from '../shared/litellm';
import { getModelByName, getTenantPhiConfig, recordUsage } from '../shared/db';
import { sanitizePHI, createMappingRecord, DEFAULT_PHI_CONFIG } from '../shared/phi';
import { DynamoDBClient } from '@aws-sdk/client-dynamodb';
import { DynamoDBDocumentClient, PutCommand } from '@aws-sdk/lib-dynamodb';

// Initialize clients
const ddbClient = DynamoDBDocumentClient.from(new DynamoDBClient({}));
const logger = new Logger({ handler: 'chat' });

// Request validation schema
const chatRequestSchema = z.object({
  model: z.string().min(1),
  messages: z.array(z.object({
    role: z.enum(['system', 'user', 'assistant', 'function', 'tool']),
    content: z.union([
      z.string(),
      z.array(z.object({
        type: z.enum(['text', 'image_url']),
        text: z.string().optional(),
        image_url: z.object({
          url: z.string(),
          detail: z.enum(['low', 'high', 'auto']).optional(),
        }).optional(),
      })),
    ]),
    name: z.string().optional(),
  })).min(1),
  max_tokens: z.number().int().positive().max(128000).optional(),
  temperature: z.number().min(0).max(2).optional(),
  stream: z.boolean().optional().default(false),
  session_id: z.string().uuid().optional(),
  enable_phi: z.boolean().optional(),
  phi_categories: z.array(z.string()).optional(),
});

```

```

type ChatRequest = z.infer<typeof chatRequestSchema>;

/**
 * Main handler
 */
export async function handler(
  event: APIGatewayProxyEvent,
  context: Context
): Promise<APIGatewayProxyResult> {
  const requestLogger = logger.child({
    requestId: context.awsRequestId,
    path: event.path,
  });

  try {
    // Extract authentication
    const auth = extractAuthContext(event);
    logAuthContext(auth, requestLogger);
    requestLogger.setTenantId(auth.tenantId);
    requestLogger.setUserId(auth.userId);

    // Parse and validate request
    const body = event.body ? JSON.parse(event.body) : {};
    const parseResult = chatRequestSchema.safeParse(body);

    if (!parseResult.success) {
      throw new ValidationError(
        'Invalid request body',
        parseResult.error.flatten().fieldErrors as Record<string, string[]>
      );
    }
  }

  const request = parseResult.data;
  const config = getConfig();
  const features = getFeatureFlags(config.TIER);

  // Get model information
  const model = await getModelByName(request.model, requestLogger);
  if (!model) {
    throw new NotFoundError(`Model ${request.model}`);
  }

  // Get tenant PHI config
  let phiConfig = DEFAULT_PHI_CONFIG;
  if (features.phiSanitization && request.enable_phi !== false) {
    try {
      const tenantConfig = await getTenantPhiConfig(auth.tenantId, requestLogger);

```

```

        if (tenantConfig) {
            phiConfig = tenantConfig;
        }
    } catch {
        // Use default if tenant config not found
    }
}

// Sanitize PHI in messages
const sanitizedMessages = await sanitizeMessages(
    request.messages,
    phiConfig,
    auth.tenantId,
    request.session_id || null,
    requestLogger
);

// Build LiteLLM request
const litellmRequest: ChatCompletionRequest = {
    model: model.name,
    messages: sanitizedMessages.messages,
    max_tokens: request.max_tokens,
    temperature: request.temperature,
    stream: request.stream,
    user: auth.userId,
    metadata: {
        tenant_id: auth.tenantId,
        session_id: request.session_id,
        app_id: config.APP_ID,
    },
};

// Handle streaming vs non-streaming
if (request.stream) {
    return await handleStreamingRequest(
        litellmRequest,
        model,
        auth,
        request,
        sanitizedMessages.mappingId,
        requestLogger
    );
} else {
    return await handleNonStreamingRequest(
        litellmRequest,
        model,
        auth,
        request,

```

```

        sanitizedMessages.mappingId,
        requestLogger
    );
}
} catch (error) {
    return handleError(error, requestLogger);
}
}

/**
 * Sanitize PHI in messages
 */
async function sanitizeMessages(
    messages: ChatRequest['messages'],
    phiConfig: typeof DEFAULT_PHI_CONFIG,
    tenantId: string,
    sessionId: string | null,
    logger: Logger
): Promise<{ messages: ChatMessage[]; mappingId: string }> {
    if (phiConfig.mode === 'disabled') {
        return {
            messages: messages as ChatMessage[],
            mappingId: '',
        };
    }

    const config = getConfig();
    const sanitizedMessages: ChatMessage[] = [];
    let combinedMappingId = '';
    const allMappings: Record<string, string> = {};

    for (const message of messages) {
        if (typeof message.content === 'string') {
            const result = sanitizePHI(message.content, phiConfig, logger);

            if (result.matches.length > 0) {
                combinedMappingId = combinedMappingId || result.mappingId;

                // Collect mappings
                for (const match of result.matches) {
                    allMappings[match.placeholder] = match.original;
                }
            }
        }

        sanitizedMessages.push({
            ...message,
            content: result.sanitizedText,
        } as ChatMessage);
    }
}

```

```

} else {
  // Handle content parts (images, etc.)
  const sanitizedParts = [];

  for (const part of message.content) {
    if (part.type === 'text' && part.text) {
      const result = sanitizePHI(part.text, phiConfig, logger);

      if (result.matches.length > 0) {
        combinedMappingId = combinedMappingId || result.mappingId;
        for (const match of result.matches) {
          allMappings[match.placeholder] = match.original;
        }
      }

      sanitizedParts.push({
        ...part,
        text: result.sanitizedText,
      });
    } else {
      sanitizedParts.push(part);
    }
  }

  sanitizedMessages.push({
    ...message,
    content: sanitizedParts,
  } as ChatMessage);
}

}

// Store PHI mapping in DynamoDB for re-identification
if (combinedMappingId && Object.keys(allMappings).length > 0) {
  const ttlHours = phiConfig.reidentification.mapping_ttl_hours;
  const expiresAt = Math.floor(Date.now() / 1000) + (ttlHours * 60 * 60);

  await ddbClient.send(new PutCommand({
    TableName: config.CACHE_TABLE,
    Item: {
      pk: `phi#${combinedMappingId}`,
      tenant_id: tenantId,
      session_id: sessionId,
      mappings: allMappings,
      created_at: new Date().toISOString(),
      ttl: expiresAt,
    },
  }));
}

```

```

        logger.info('PHI mappings stored', {
            mappingId: combinedMappingId,
            mappingCount: Object.keys(allMappings).length,
        });
    }

    return {
        messages: sanitizedMessages,
        mappingId: combinedMappingId,
    };
}

/**
 * Handle non-streaming request
 */
async function handleNonStreamingRequest(
    request: ChatCompletionRequest,
    model: Awaited<ReturnType<typeof getModelByName>>,
    auth: ReturnType<typeof extractAuthContext>,
    originalRequest: ChatRequest,
    phiMappingId: string,
    logger: Logger
): Promise<APIGatewayProxyResult> {
    const startTime = Date.now();

    // Call LiteLLM
    const response = await createChatCompletion(request, logger);
    const latency = Date.now() - startTime;

    // Calculate costs
    const pricing = model!.pricing;
    const { cost, billed } = calculateCost(response.usage, {
        input_tokens: pricing.input_tokens || 0,
        output_tokens: pricing.output_tokens || 0,
        billed_markup: pricing.billed_markup,
    });

    // Record usage
    await recordUsage({
        tenant_id: auth.tenantId,
        user_id: auth.userId,
        session_id: originalRequest.session_id || null,
        model_id: model!.id,
        provider_id: model!.provider_id,
        input_tokens: response.usage.prompt_tokens,
        output_tokens: response.usage.completion_tokens,
        total_tokens: response.usage.total_tokens,
        cost,
    });
}

```

```

    billed_amount: billed,
    request_type: 'chat.completion',
    latency_ms: latency,
}, logger);

// Store in sessions table if session_id provided
if (originalRequest.session_id) {
  const config = getConfig();
  await ddbClient.send(new PutCommand({
    TableName: config.SESSIONS_TABLE,
    Item: {
      pk: originalRequest.session_id,
      gsi1pk: auth.userId,
      gsi1sk: new Date().toISOString(),
      messages: [
        ...originalRequest.messages,
        response.choices[0]?.message,
      ],
      model: request.model,
      tenant_id: auth.tenantId,
      phi_mapping_id: phiMappingId || undefined,
      created_at: new Date().toISOString(),
      updated_at: new Date().toISOString(),
      ttl: Math.floor(Date.now() / 1000) + (30 * 24 * 60 * 60), // 30 days
    },
  }));
}

logger.info('Chat completion successful', {
  model: request.model,
  inputTokens: response.usage.prompt_tokens,
  outputTokens: response.usage.completion_tokens,
  latency,
  cost,
});

return success({
  id: response.id,
  object: response.object,
  created: response.created,
  model: response.model,
  choices: response.choices,
  usage: response.usage,
  session_id: originalRequest.session_id,
  phi_mapping_id: phiMappingId || undefined,
});
}

```

```

/**
 * Handle streaming request
 * Note: API Gateway doesn't support true streaming, so we buffer and return
 * For true streaming, use WebSocket API or Lambda Function URLs
 */
async function handleStreamingRequest(
    request: ChatCompletionRequest,
    model: Awaited<ReturnType<typeof getModelByName>>,
    auth: ReturnType<typeof extractAuthContext>,
    originalRequest: ChatRequest,
    phiMappingId: string,
    logger: Logger
): Promise<APIGatewayProxyResult> {
    const startTime = Date.now();
    const chunks: string[] = [];
    let totalContent = '';
    let finishReason: string | null = null;
    let usage = { prompt_tokens: 0, completion_tokens: 0, total_tokens: 0 };

    try {
        for await (const chunk of streamChatCompletion(request, logger)) {
            chunks.push(formatSSE(chunk));

            if (chunk.choices[0]?.delta?.content) {
                totalContent += chunk.choices[0].delta.content;
            }

            if (chunk.choices[0]?.finish_reason) {
                finishReason = chunk.choices[0].finish_reason;
            }

            if (chunk.usage) {
                usage = chunk.usage;
            }
        }

        // Add done marker
        chunks.push(formatSSE('[DONE]'));
    } catch (error) {
        logger.error('Streaming error', error as Error);
        throw error;
    }

    const latency = Date.now() - startTime;

    // Record usage if we have it
    if (usage.total_tokens > 0) {
        const pricing = model!.pricing;
    }
}

```

```

const { cost, billed } = calculateCost(usage, {
  input_tokens: pricing.input_tokens || 0,
  output_tokens: pricing.output_tokens || 0,
  billed_markup: pricing.billed_markup,
});

await recordUsage({
  tenant_id: auth.tenantId,
  user_id: auth.userId,
  session_id: originalRequest.session_id || null,
  model_id: model!.id,
  provider_id: model!.provider_id,
  input_tokens: usage.prompt_tokens,
  output_tokens: usage.completion_tokens,
  total_tokens: usage.total_tokens,
  cost,
  billed_amount: billed,
  request_type: 'chat.completion.stream',
  latency_ms: latency,
}, logger);
}

logger.info('Streaming completion successful', {
  model: request.model,
  chunkCount: chunks.length,
  contentLength: totalContent.length,
  latency,
});
}

// Return SSE formatted response
// Note: For true streaming, implement Lambda Function URL or WebSocket
return {
  statusCode: 200,
  headers: streamingHeaders(),
  body: chunks.join(''),
};
}

```

PART 8: MODELS LAMBDA

packages/infrastructure/lambda/api/models.ts

```

/**
 * Models Lambda - Dynamic model registry
 *
 * Handles:
 * - List all models

```

```

* - Get model by ID
* - Filter by specialty, provider, status
* - Admin: Update model status, thermal state
*/

import type {
    APIGatewayProxyEvent,
    APIGatewayProxyResult,
    Context,
} from 'aws-lambda';
import { z } from 'zod';
import { Logger } from '../shared/logger';
import { getConfig } from '../shared/config';
import { success, handleError, noContent } from '../shared/response';
import { extractAuthContext, requireAdmin, logAuthContext } from '../shared/auth';
import { ValidationError, NotFoundError } from '../shared/errors';
import { getModels, getModelById, updateModel as dbUpdateModel } from '../shared/db';
import { createAuditLog } from '../shared/db';

// Initialize logger
const logger = new Logger({ handler: 'models' });

// Query parameters schema
const listQuerySchema = z.object({
    provider_id: z.string().optional(),
    specialty: z.string().optional(),
    status: z.enum(['active', 'inactive', 'deprecated', 'coming_soon']).optional(),
    type: z.enum(['external', 'self-hosted']).optional(),
    supports_vision: z.string().transform(v => v === 'true').optional(),
    supports_streaming: z.string().transform(v => v === 'true').optional(),
    limit: z.string().transform(Number).pipe(z.number().int().min(1).max(100)).optional(),
    offset: z.string().transform(Number).pipe(z.number().int().min(0)).optional(),
});

// Update request schema (admin only)
const updateModelSchema = z.object({
    status: z.enum(['active', 'inactive', 'deprecated', 'coming_soon']).optional(),
    thermal_state: z.enum(['OFF', 'COLD', 'WARM', 'HOT', 'AUTOMATIC']).optional(),
    display_name: z.string().min(1).max(200).optional(),
    description: z.string().max(2000).optional(),
});

/**
* Main handler
*/
export async function handler(
    event: APIGatewayProxyEvent,
    context: Context
)

```

```

): Promise<APIGatewayProxyResult> {
  const requestLogger = logger.child({
    requestId: context.awsRequestId,
    path: event.path,
    method: event.httpMethod,
  });

  try {
    // Extract authentication
    const auth = extractAuthContext(event);
    logAuthContext(auth, requestLogger);

    // Parse path parameters
    const modelId = event.pathParameters?.modelId;

    switch (event.httpMethod) {
      case 'GET':
        if (modelId) {
          return await handleGetModel(modelId, requestLogger);
        }
        return await handleListModel(event, requestLogger);

      case 'PUT':
      case 'PATCH':
        if (!modelId) {
          throw new ValidationError('Model ID required for update');
        }
        requireAdmin(auth);
        return await handleUpdateModel(modelId, event, auth, requestLogger);

      default:
        throw new ValidationError(`Method ${event.httpMethod} not allowed`);
    }
  } catch (error) {
    return handleError(error, requestLogger);
  }
}

/**
 * List models with filtering
 */
async function handleListModel(
  event: APIGatewayProxyEvent,
  logger: Logger
): Promise<APIGatewayProxyResult> {
  // Parse and validate query parameters
  const queryResult = listQuerySchema.safeParse(event.queryStringParameters || {});
}

```

```

    if (!queryResult.success) {
      throw new ValidationError(
        'Invalid query parameters',
        queryResult.error.flatten().fieldErrors as Record<string, string[]>
      );
    }

    const filters = queryResult.data;

    // Query database
    const result = await getModels({
      providerId: filters.provider_id,
      specialty: filters.specialty,
      status: filters.status,
      type: filters.type,
      supportsVision: filters.supports_vision,
      supportsStreaming: filters.supports_streaming,
      limit: filters.limit || 50,
      offset: filters.offset || 0,
    }, logger);

    // Transform to API response format
    const models = result.rows.map(transformModel);

    logger.info('Models listed', {
      count: models.length,
      filters,
    });

    return success({
      models,
      pagination: {
        limit: filters.limit || 50,
        offset: filters.offset || 0,
        total: result.rowCount,
        hasMore: result.rowCount > (filters.offset || 0) + models.length,
      },
    });
  }

  /**
   * Get single model by ID
   */
  async function handleGetModel(
    modelId: string,
    logger: Logger
  ): Promise<APIGatewayProxyResult> {
    const model = await getModelById(modelId, logger);
  }
}

```

```

logger.info('Model retrieved', { modelId });

return success({
  model: transformModel(model),
});
}

/**
 * Update model (admin only)
 */
async function handleUpdateModel(
  modelId: string,
  event: APIGatewayProxyEvent,
  auth: ReturnType<typeof extractAuthContext>,
  logger: Logger
): Promise<APIGatewayProxyResult> {
  // Parse and validate request body
  const body = event.body ? JSON.parse(event.body) : {};
  const parseResult = updateModelSchema.safeParse(body);

  if (!parseResult.success) {
    throw new ValidationError(
      'Invalid request body',
      parseResult.error.flatten().fieldErrors as Record<string, string[]>
    );
  }
}

const updates = parseResult.data;

if (Object.keys(updates).length === 0) {
  throw new ValidationError('No updates provided');
}

// Get current model
const currentModel = await getModelById(modelId, logger);

// Update model
const updatedModel = await dbUpdateModel(modelId, updates, logger);

// Create audit log
await createAuditLog({
  tenant_id: auth.tenantId,
  user_id: null,
  admin_id: auth.userId,
  action: 'model.update',
  resource_type: 'model',
  resource_id: modelId,
}

```

```

details: {
  before: {
    status: currentModel.status,
    thermal_state: currentModel.thermal_state,
    display_name: currentModel.display_name,
  },
  after: updates,
},
ip_address: event.requestContext.identity?.sourceIp || null,
user_agent: event.headers['User-Agent'] || null,
}, logger);

logger.info('Model updated', {
  modelId,
  updates,
  adminId: auth.userId,
});

return success({
  model: transformModel(updatedModel),
});
}

/**
 * Transform database model to API response format
 */
function transformModel(dbModel: any) {
  return {
    id: dbModel.id,
    providerId: dbModel.provider_id,
    name: dbModel.name,
    displayName: dbModel.display_name,
    description: dbModel.description,
    type: dbModel.type,
    specialty: dbModel.specialty,
    capabilities: dbModel.capabilities,
    contextWindow: dbModel.context_window,
    maxOutputTokens: dbModel.max_output_tokens,
    supportsFunctions: dbModel.supports_functions,
    supportsVision: dbModel.supports_vision,
    supportsStreaming: dbModel.supports_streaming,
    hasThinkingMode: dbModel.has_thinking_mode,
    thinkingBudgetTokens: dbModel.thinking_budget_tokens,
    pricing: {
      inputTokens: dbModel.pricing?.input_tokens,
      outputTokens: dbModel.pricing?.output_tokens,
      perImage: dbModel.pricing?.per_image,
      perMinuteAudio: dbModel.pricing?.per_minute_audio,
    }
  }
}

```

```

    perMinuteVideo: dbModel.pricing?.per_minute_video,
    per3DModel: dbModel.pricing?.per_3d_model,
    billedMarkup: dbModel.pricing?.billed_markup,
),
thermalState: dbModel.thermal_state,
thermalConfig: dbModel.thermal_config,
status: dbModel.status,
releaseDate: dbModel.release_date,
deprecationDate: dbModel.deprecation_date,
createdAt: dbModel.created_at,
updatedAt: dbModel.updated_at,
);
}

```

PART 9: PROVIDERS LAMBDA

packages/infrastructure/lambda/api/providers.ts

```

/*
 * Providers Lambda - AI provider management
 *
 * Handles:
 * - List all providers
 * - Get provider by ID
 * - Get provider models
 * - Admin: Update provider status
 */

import type {
  APIGatewayProxyEvent,
  APIGatewayProxyResult,
  Context,
} from 'aws-lambda';
import { z } from 'zod';
import { Logger } from '../shared/logger';
import { getConfig } from '../shared/config';
import { success, handleError } from '../shared/response';
import { extractAuthContext, requireAdmin, logAuthContext } from '../shared/auth';
import { ValidationError } from '../shared/errors';
import {
  getProviders,
  getProviderById,
  updateProvider as dbUpdateProvider,
  getModels,
  createAuditLog,
} from '../shared/db';

```

```

// Initialize logger
const logger = new Logger({ handler: 'providers' });

// Query parameters schema
const listQuerySchema = z.object({
  type: z.enum(['external', 'self-hosted', 'mid-tier']).optional(),
  status: z.enum(['active', 'inactive', 'deprecated']).optional(),
  hipaa_compliant: z.string().transform(v => v === 'true').optional(),
  limit: z.string().transform(Number).pipe(z.number().int().min(1).max(100)).optional(),
  offset: z.string().transform(Number).pipe(z.number().int().min(0)).optional(),
});

// Update request schema (admin only)
const updateProviderSchema = z.object({
  status: z.enum(['active', 'inactive', 'deprecated']).optional(),
  hipaa_compliant: z.boolean().optional(),
  config: z.record(z.unknown()).optional(),
});

/**
 * Main handler
 */
export async function handler(
  event: APIGatewayProxyEvent,
  context: Context
): Promise<APIGatewayProxyResult> {
  const requestLogger = logger.child({
    requestId: context.awsRequestId,
    path: event.path,
    method: event.httpMethod,
  });

  try {
    // Extract authentication
    const auth = extractAuthContext(event);
    logAuthContext(auth, requestLogger);

    // Parse path
    const providerId = event.pathParameters?.providerId;
    const subPath = event.path.split('/').pop();

    switch (event.httpMethod) {
      case 'GET':
        if (providerId) {
          if (subPath === 'models') {
            return await handleGetProviderModels(providerId, requestLogger);
          }
        }
        return await handleGetProvider(providerId, requestLogger);
    }
  } catch (error) {
    logError(error, requestLogger);
    return {
      statusCode: 500,
      body: JSON.stringify({ error: 'Internal Server Error' }),
    };
  }
}

```

```

        }

        return await handleListProviders(event, requestLogger);

    case 'PUT':
    case 'PATCH':
        if (!providerId) {
            throw new ValidationError('Provider ID required for update');
        }
        requireAdmin(auth);
        return await handleUpdateProvider(providerId, event, auth, requestLogger);

    default:
        throw new ValidationError(`Method ${event.httpMethod} not allowed`);
    }
} catch (error) {
    return handleError(error, requestLogger);
}
}

/**
 * List providers with filtering
 */
async function handleListProviders(
    event: APIGatewayProxyEvent,
    logger: Logger
): Promise<APIGatewayProxyResult> {
    // Parse and validate query parameters
    const queryResult = listQuerySchema.safeParse(event.queryStringParameters || {});

    if (!queryResult.success) {
        throw new ValidationError(
            'Invalid query parameters',
            queryResult.error.flatten().fieldErrors as Record<string, string[]>
        );
    }
}

const filters = queryResult.data;

// Query database
const result = await getProviders({
    type: filters.type,
    status: filters.status,
    hipaaCompliant: filters.hipaa_compliant,
    limit: filters.limit || 50,
    offset: filters.offset || 0,
}, logger);

// Transform to API response format

```

```

const providers = result.rows.map(transformProvider);

logger.info('Providers listed', {
  count: providers.length,
  filters,
});

return success({
  providers,
  pagination: {
    limit: filters.limit || 50,
    offset: filters.offset || 0,
    total: result.rowCount,
    hasMore: result.rowCount > (filters.offset || 0) + providers.length,
  },
});
}

/**
 * Get single provider by ID
 */
async function handleGetProvider(
  providerId: string,
  logger: Logger
): Promise<APIGatewayProxyResult> {
  const provider = await getProviderById(providerId, logger);

  logger.info('Provider retrieved', { providerId });

  return success({
    provider: transformProvider(provider),
  });
}

/**
 * Get models for a provider
 */
async function handleGetProviderModels(
  providerId: string,
  logger: Logger
): Promise<APIGatewayProxyResult> {
  // Verify provider exists
  await getProviderById(providerId, logger);

  // Get provider's models
  const result = await getModels({
    providerId,
    status: 'active',
  });
}

```

```

}, logger);

const models = result.rows.map(model => ({
  id: model.id,
  name: model.name,
  displayName: model.display_name,
  specialty: model.specialty,
  status: model.status,
  thermalState: model.thermal_state,
}));

logger.info('Provider models retrieved', {
  providerId,
  modelCount: models.length,
});

return success({
  providerId,
  models,
});
}

/**
 * Update provider (admin only)
 */
async function handleUpdateProvider(
  providerId: string,
  event: APIGatewayProxyEvent,
  auth: ReturnType<typeof extractAuthContext>,
  logger: Logger
): Promise<APIGatewayProxyResult> {
  // Parse and validate request body
  const body = event.body ? JSON.parse(event.body) : {};
  const parseResult = updateProviderSchema.safeParse(body);

  if (!parseResult.success) {
    throw new ValidationError(
      'Invalid request body',
      parseResult.error.flatten().fieldErrors as Record<string, string[]>
    );
  }

  const updates = parseResult.data;

  if (Object.keys(updates).length === 0) {
    throw new ValidationError('No updates provided');
  }
}

```

```

// Get current provider
const currentProvider = await getProviderById(providerId, logger);

// Update provider
const updatedProvider = await dbUpdateProvider(providerId, {
  status: updates.status,
  hipaa_compliant: updates.hipaa_compliant,
  config: updates.config,
}, logger);

// Create audit log
await createAuditLog({
  tenant_id: auth.tenantId,
  user_id: null,
  admin_id: auth.userId,
  action: 'provider.update',
  resource_type: 'provider',
  resource_id: providerId,
  details: {
    before: {
      status: currentProvider.status,
      hipaa_compliant: currentProvider.hipaa_compliant,
    },
    after: updates,
  },
  ip_address: event.requestContext.identity?.sourceIp || null,
  user_agent: event.headers['User-Agent'] || null,
}, logger);

logger.info('Provider updated', {
  providerId,
  updates,
  adminId: auth.userId,
});

return success({
  provider: transformProvider(updatedProvider),
});
}

/**
 * Transform database provider to API response format
 */
function transformProvider(dbProvider: any) {
  return {
    id: dbProvider.id,
    name: dbProvider.name,
    type: dbProvider.type,
}
}

```

```

    status: dbProvider.status,
    hipaaCompliant: dbProvider.hipaa_compliant,
    baaAvailable: dbProvider.baa_available,
    baseUrl: dbProvider.base_url,
    authType: dbProvider.auth_type,
    capabilities: dbProvider.capabilities,
    createdAt: dbProvider.created_at,
    updatedAt: dbProvider.updated_at,
};

}

```

DEPLOYMENT VERIFICATION

After deploying the Lambda functions, verify they work:

```

# 1. Health check
curl https://api.thinktank.YOUR_DOMAIN.com/api/v2/health

# Expected response:
# {
#   "success": true,
#   "data": {
#     "status": "healthy",
#     "timestamp": "2024-12-21T...",
#     "environment": "dev",
#     "tier": 1
#   }
# }

# 2. List models (requires auth token)
curl -H "Authorization: Bearer $TOKEN" \
  https://api.thinktank.YOUR_DOMAIN.com/api/v2/models

# 3. List providers
curl -H "Authorization: Bearer $TOKEN" \
  https://api.thinktank.YOUR_DOMAIN.com/api/v2/providers

# 4. Chat completion
curl -X POST \
  -H "Authorization: Bearer $TOKEN" \
  -H "Content-Type: application/json" \
  -d '{
    "model": "gpt-4",
    "messages": [{"role": "user", "content": "Hello!"}]
}' \
  https://api.thinktank.YOUR_DOMAIN.com/api/v2/chat/completions

```

NEXT PROMPTS

Continue with: - **Prompt 5:** Lambda Functions - Admin & Billing (Invitations, Approvals, Metering) - **Prompt 6:** Self-Hosted Models & Mid-Level Services Configuration - **Prompt 7:** External Providers & Database Schema/Migrations - **Prompt 8:** Admin Web Dashboard (Next.js) - **Prompt 9:** Assembly & Deployment Guide

End of Prompt 4: Lambda Functions - Core RADIANT v2.2.0 - December 2024

â• â•

END OF SECTION 4

â• â•

â• â•