

# Contents

<b>SECTION 5: LAMBDA FUNCTIONS - ADMIN &amp; BILLING (v2.1.0)</b>	<b>2</b>
Type Imports . . . . .	2
<b>RADIANT v2.2.0 - Prompt 5: Lambda Functions - Admin &amp; Billing</b>	<b>2</b>
OVERVIEW . . . . .	2
KEY FEATURES . . . . .	2
Administrator Management . . . . .	2
Billing & Metering . . . . .	2
LAMBDA DIRECTORY STRUCTURE . . . . .	3
PART 1: SHARED ADMIN UTILITIES . . . . .	3
packages/infrastructure/lambda/shared/admin/index.ts . . . . .	3
packages/infrastructure/lambda/shared/admin/types.ts . . . . .	3
packages/infrastructure/lambda/shared/admin/tokens.ts . . . . .	4
packages/infrastructure/lambda/shared/admin/email.ts . . . . .	5
packages/infrastructure/lambda/shared/admin/stripe.ts . . . . .	6
PART 2: INVITATIONS LAMBDA . . . . .	7
packages/infrastructure/lambda/admin/invitations.ts . . . . .	7
PART 3: TWO-PERSON APPROVALS LAMBDA . . . . .	12
packages/infrastructure/lambda/admin/approvals.ts . . . . .	12
PART 4: METERING LAMBDA . . . . .	17
packages/infrastructure/lambda/billing/metering.ts . . . . .	17
PART 5: DATABASE SCHEMA ADDITIONS . . . . .	20
packages/infrastructure/migrations/005_admin_billing.sql . . . . .	20
API ROUTES SUMMARY . . . . .	22
Admin Routes . . . . .	22
Billing Routes . . . . .	22
Metering Routes . . . . .	23
DEPLOYMENT VERIFICATION . . . . .	23
 2. List Pending Approvals (for me to approve) . . . . .	23
3. Get Current Billing Period . . . . .	23
4. Get Usage Rollups . . . . .	23
5. Generate Invoice . . . . .	23
ESTIMATED COSTS BY TIER . . . . .	24
NEXT PROMPTS . . . . .	24
 END OF SECTION 5 . . . . .	24

SECTION 5: LAMBDA FUNCTIONS - ADMIN & BILLING  
(v2.1.0)

**Dependencies:** Sections 0-4 **Creates:** Invitations, Approvals, Metering, Billing handlers

## Type Imports

```
import { Administrator, AdminRole, Invitation, ApprovalRequest, Invoice } from '@radiant/shared'
```

**IMPORTANT:** Use canonical table names: administrators, invitations, approval requests

RADIANT v2.2.0 - Prompt 5: Lambda Functions - Admin & Billing

**Prompt 5 of 9** | Target Size: ~60KB | Version: 3.7.0 | December 2024

## OVERVIEW

This prompt creates the Admin & Billing Lambda functions for the RADIANT platform:

1. **Invitations Lambda** - Email-based administrator invitations with secure tokens
  2. **Approvals Lambda** - Two-person approval workflow for production deployments
  3. **Admin Users Lambda** - Administrator CRUD, roles, and MFA management
  4. **Admin Profiles Lambda** - Preferences, notifications, and settings
  5. **Metering Lambda** - Real-time usage event collection and tracking
  6. **Billing Lambda** - Cost aggregation, invoicing, and payment processing
  7. **Audit Lambda** - Comprehensive audit logging and compliance reporting
  8. **Notifications Lambda** - Admin notification delivery and management

## KEY FEATURES

## **Administrator Management**

- **Email Invitations:** Secure token-based invitation system with expiry
  - **Role-Based Access:** super\_admin, admin, operator, auditor roles
  - **MFA Requirement:** Mandatory MFA for production environment access
  - **Two-Person Approval:** Production deployments require separate initiator and approver

## Billing & Metering

- **Real-Time Tracking:** Every API call metered with token counts and costs

- **Dynamic Pricing:** Costs pulled from Dynamic Model Registry
  - **Configurable Margin:** Default 20% markup on provider costs
  - **Stripe Integration:** Automatic billing and payment processing
  - **Usage Rollups:** Daily aggregation for efficient querying
- 

## LAMBDA DIRECTORY STRUCTURE

```
```
packages/infrastructure/lambda/
  -- Admin
    - invitations.ts # Email invitations
    - approvals.ts # Two-person approval workflow
    - users.ts # Admin user CRUD
    - profiles.ts # Admin preferences
    - notifications.ts # Notification management
    - audit.ts # Audit log queries
    - billing/
      - billing.ts # Cost aggregation & invoices
      - pricing.ts # Dynamic pricing sync
    - payments.ts # Stripe integration
    - shared/
      - admin/
        - index.ts # Admin-specific types
      - tokens.ts # Secure token generation
    - stripe.ts # Stripe client
```

```

---

## PART 1: SHARED ADMIN UTILITIES

`packages/infrastructure/lambda/shared/admin/index.ts`

```
```
typescript // Re-export admin utilities
export * from './types';
export * from './email';
export * from './tokens';
export * from './stripe';
```

```

`packages/infrastructure/lambda/shared/admin/types.ts`

```
```
typescript /**
 * Admin-specific types for RADIANT v2.2.0 */
import { z } from 'zod';

// =====
// ADMINISTRATOR ROLES // =====

export const AdminRole = {
  SUPER_ADMIN: 'super_admin',
  ADMIN: 'admin',
  OPERATOR: 'operator',
  AUDITOR: 'auditor',
} as const;

export type AdminRoleType = typeof AdminRole[keyof typeof AdminRole];

export const ROLE_HIERARCHY: Record<AdminRoleType, number> = {
  [AdminRole.SUPER_ADMIN]: 100,
  [AdminRole.ADMIN]: 75,
  [AdminRole.OPERATOR]: 50,
  [AdminRole.AUDITOR]: 25,
};

export const ROLE_PERMISSIONS: Record<AdminRoleType, string[]> = {
  [AdminRole.SUPER_ADMIN]: [
    'admin:',
    'billing:',
    'settings:',
    'deployments:',
    'approvals:',
    'audit:',
  ],
  [AdminRole.ADMIN]: [
    'admin:read',
    'admin:write',
    'billing:read',
    'settings:read',
    'settings:write',
    'deployments:read',
    'deployments:write',
    'approvals:read',
    'approvals:initiate',
  ],
};
```

```

```

[AdminRole.OPERATOR]: [ 'admin:read', 'billing:read', 'settings:read', 'deployments:read', ],
[AdminRole.AUDITOR]: [ 'admin:read', 'billing:read', 'audit:read', ], };

// =====
// INVITATION TYPES // =====

export interface Invitation { id: string; email: string; role: AdminRoleType; invitedBy: string; invitedByName: string; appId: string; tenantId: string; environment: 'dev' | 'staging' | 'prod'; token: string; tokenHash: string; expiresAt: string; status: 'pending' | 'accepted' | 'expired' | 'revoked'; createdAt: string; acceptedAt?: string; acceptedByIp?: string; }

export const createInvitationSchema = z.object({ email: z.string().email(), role: z.enum(['super_admin', 'admin', 'operator', 'auditor']), message: z.string().max(500).optional(), expiresInHours: z.number().int().min(1).max(168).default(48), });

export const acceptInvitationSchema = z.object({ token: z.string().min(32), firstName: z.string().min(1).max(50), lastName: z.string().min(1).max(50), password: z.string().min(12).max(128), mfaMethod: z.enum(['authenticator', 'sms', 'email']).default('authenticator'), phone: z.string().optional(), });

// =====
// APPROVAL TYPES // =====

export type ApprovalType = | 'deployment' | 'promotion' | 'model_activation' | 'provider_change' | 'user_role_change' | 'billing_change';

export type ApprovalStatus = | 'pending' | 'approved' | 'rejected' | 'expired' | 'cancelled';

export interface ApprovalRequest { id: string; type: ApprovalType; appId: string; tenantId: string; environment: 'dev' | 'staging' | 'prod'; requestedBy: string; requestedByName: string; requestedAt: string; expiresAt: string; status: ApprovalStatus; approvedBy?: string; approvedByName?: string; approvedAt?: string; rejectedReason?: string; action: string; resourceType: string; resourceId: string; details: Record<string, unknown>; priority: 'low' | 'medium' | 'high' | 'critical'; notes?: string; }

export const createApprovalSchema = z.object({ type: z.enum(['deployment', 'promotion', 'model_activation', 'provider_change', 'user_role_change', 'billing_change']), action: z.string(), resourceType: z.string(), resourceId: z.string(), details: z.record(z.unknown()), priority: z.enum(['low', 'medium', 'high', 'critical']).default('medium'), notes: z.string().max(1000).optional(), expiresInHours: z.number().int().min(1).max(168).default(24), });

export const processApprovalSchema = z.object({ action: z.enum(['approve', 'reject']), reason: z.string().max(500).optional(), });


```

## **packages/infrastructure/lambda/shared/admin/tokens.ts**

```

``typescript /**
 * Secure token generation and validation */
import { randomBytes, createHash, timingSafeEqual } from 'crypto';

export function generateToken(length: number = 32): string { return randomBytes(length).toString('base64url'); }
``
```

```

export function generateCode(length: number = 6): string { const chars = '0123456789ABCDE-  
FGHIJKLMNOPQRSTUVWXYZ'; const bytes = randomBytes(length); let code = ""; for (let i =  
0; i < length; i++) { code += chars[bytes[i] % chars.length]; } return code; }

export function hashToken(token: string): string { return createHash('sha256').update(token).digest('hex'); }

export function verifyToken(token: string, hash: string): boolean { const tokenHash = hashTo-  
ken(token); try { return timingSafeEqual( Buffer.from(tokenHash, 'hex'), Buffer.from(hash, 'hex') ) } catch { return false; } }

export function generateInvitationToken(): { token: string; tokenHash: string } { const token =  
generateToken(48); const tokenHash = hashToken(token); return { token, tokenHash }; }

export function calculateExpiry(hoursFromNow: number): string { const expiry = new Date();  
expiry.setHours(expiry.getHours() + hoursFromNow); return expiry.toISOString(); }

export function isExpired(expiresAt: string): boolean { return new Date(expiresAt) < new Date(); }
}

```

## packages/infrastructure/lambda/shared/admin/email.ts

```

``typescript /**
 * Email utilities using AWS SES
 */

import { SESClient, SendEmailCommand } from '@aws-sdk/client-ses';
import { Logger } from './logger';

const sesClient = new SESClient({});

export interface EmailOptions {
    to: string | string[];
    subject: string;
    html: string;
    text?: string;
    replyTo?: string;
}

export async function sendEmail(options: EmailOptions, logger: Logger): Promise<void> {
    const fromEmail = 'noreply@${process.env.DOMAIN} || "radiant.cloud"';
    const toAddresses = Array.isArray(options.to) ? options.to : [options.to];

    const command = new SendEmailCommand({
        Source: fromEmail,
        Destination: { ToAddresses },
        Message: {
            Subject: { Data: options.subject, Charset: 'UTF-8' },
            Body: {
                Html: { Data: options.html, Charset: 'UTF-8' },
                ...(
                    options.text &&
                    { Text: { Data: options.text, Charset: 'UTF-8' } }
                )
            },
            ReplyToAddresses: options.replyTo ? [options.replyTo] : undefined
        }
    });

    try {
        await sesClient.send(command);
        logger.info('Email sent successfully', { to: toAddresses, subject: options.subject });
    } catch (error) {
        logger.error('Failed to send email', error as Error, { to: toAddresses });
        throw error;
    }
}

export function generateInvitationEmail(params: { inviteeName: string; inviterName: string; role: string; appName: string; environment: string; acceptUrl: string; expiresAt: string; message?: string }): { html: string; text: string } {
    const html = `<!DOCTYPE html>
    

# You're Invited!


    <p><strong>${params.inviterName}</strong> has invited you to join <strong>${params.appName}</strong>
    ${params.message ? `<blockquote style="background: #f8f9fa; padding: 15px; border-left: 4px solid #1a1a1a; margin-left: 20px;">` : ''}
    ${params.message}
    ${params.message ? `</blockquote>` : ''}
    <div style="text-align: center; margin: 30px 0;">

```

```

<a href="\${params.acceptUrl}" style="display: inline-block; padding: 14px 32px; background:
</div>
<p style="color: #666; font-size: 14px;">This invitation expires on <strong>\${new Date(params
";
const text = 'You\'ve been invited to ${params.appName} by ${params.inviterName} as a
${params.role}. Accept: ${params.acceptUrl}'; return { html, text }; }

export function generateApprovalEmail(params: { approverName: string; requesterName: string;
appName: string; environment: string; action: string; resourceType: string; resourceId: string;
approveUrl: string; rejectUrl: string; expiresAt: string; }): { html: string; text: string } { const
html = `<!DOCTYPE html>

Approval Required - ${params.appName}

<div style="background: #ffc107; padding: 20px; text-align: center;">
  <h1 style="margin: 0; color: #1a1a1a;">Ã¢Â¡ÃÃ¢Â,Ã¢Â Approval Required</h1>
</div>
<div style="padding: 30px;">
  <p>Hi \${params.approverName},</p>
  <p><strong>\${params.requesterName}</strong> has requested approval for:</p>
  <div style="background: #f8f9fa; padding: 20px; border-radius: 6px; margin: 20px 0;">
    <p><strong>Environment:</strong> \${params.environment.toUpperCase()}</p>
    <p><strong>Action:</strong> \${params.action}</p>
    <p><strong>Resource:</strong> \${params.resourceType} (\${params.resourceId})</p>
  </div>
  <div style="text-align: center; margin: 30px 0;">
    <a href="\${params.approveUrl}" style="display: inline-block; padding: 14px 32px; background:
      <a href="\${params.rejectUrl}" style="display: inline-block; padding: 14px 32px; background:
    </div>
    <p style="color: #dc3545; font-size: 12px; text-align: center;">Ã¢Â¡ÃÃ¢Â,Ã¢Â Production deploy
  </div>
</div>

const text = 'APPROVAL REQUIRED: ${params.requesterName} requests approval for
${params.action} on ${params.resourceType}. Approve: ${params.approveUrl} Reject:
${params.rejectUrl}'; return { html, text }; } ```

```

## packages/infrastructure/lambda/shared/admin/stripe.ts

```

``typescript /**
 * Stripe payment integration */

import Stripe from 'stripe'; import { SecretsManagerClient, GetSecretValueCommand } from
'@aws-sdk/client-secrets-manager'; import { Logger } from '../logger';

const secretsClient = new SecretsManagerClient({}); let stripeClient: Stripe | null = null;

async function getStripeClient(): Promise<{ if (stripeClient) return stripeClient;

const secretArn = process.env.STRIPE_SECRET_ARN; if (!secretArn) throw new Error('STRIPE_SECRET_ARN not configured');
```

```

```

const command = new GetSecretValueCommand({ SecretId: secretArn });
const response = await secretsClient.send(command);
if (!response.SecretString) throw new Error('Stripe secret not found');

const secrets = JSON.parse(response.SecretString);
stripeClient = new Stripe(secrets.apiKey, {
  apiVersion: '2023-10-16'
});
return stripeClient;

export async function getOrCreateCustomer( tenantId: string, email: string, name: string, metadata: Record<string, string>, logger: Logger ): Promise<{ const stripe = await getStripeClient(); }>
{
  const existing = await stripe.customers.search({
    query: `metadata['tenantId']: ${tenantId}`,
    limit: 1,
  });

  if (existing.data.length > 0) {
    logger.info('Found existing Stripe customer', { tenantId, customerId: existing.data[0].id });
    return existing.data[0].id;
  }

  const customer = await stripe.customers.create({
    email,
    name,
    metadata: { tenantId, ...metadata },
  });

  logger.info('Created Stripe customer', { tenantId, customerId: customer.id });
  return customer.id;
}

export async function createInvoice(params: { customerId: string; tenantId: string; periodStart: Date; periodEnd: Date; lineItems: Array<{ description: string; amount: number; quantity: number; metadata?: Record<string, string> }>; metadata?: Record<string, string>; logger: Logger }): Promise<Stripe.Invoice>
{
  const stripe = await getStripeClient();

  for (const item of params.lineItems) {
    await stripe.invoiceItems.create({
      customer: params.customerId,
      amount: item.amount,
      currency: 'usd',
      description: item.description,
      quantity: item.quantity,
      metadata: item.metadata,
    });
  }

  const invoice = await stripe.invoices.create({
    customer: params.customerId,
    auto_advance: true,
    collection_method: 'charge Automatically',
    metadata: { tenantId: params.tenantId },
    period_start: params.periodStart.toISOString(),
    period_end: params.periodEnd.toISOString(),
    ...params.metadata,
  });
}

logger.info('Created Stripe invoice', { tenantId: params.tenantId, invoiceId: invoice.id });
return invoice;
}

export async function getInvoiceStatus(invoiceId: string, logger: Logger): Promise<{ status: string; paid: boolean; amountDue: number; amountPaid: number }>
{
  const stripe = await getStripeClient();
  const invoice = await stripe.invoices.retrieve(invoiceId);
  return {
    status: invoice.status || 'unknown',
    paid: invoice.paid,
    amountDue: invoice.amount_due,
    amountPaid: invoice.amount_paid,
  };
}

```

---

## PART 2: INVITATIONS LAMBDA

[packages/infrastructure/lambda/admin/invitations.ts](#)

```
``typescript /**
 * Administrator Invitation Lambda
 * Handles email-based administrator invitations with secure tokens
 */
```

```

import { APIGatewayProxyEvent, APIGatewayProxyResult, Context } from 'aws-lambda';
import { v4 as uuidv4 } from 'uuid'; import { Logger } from '../shared/logger'; import {
success, created, handleError } from '../shared/response'; import { extractAuthContext, requireAdmin, requirePermission } from '../shared/auth'; import { ValidationError, NotFoundError, ForbiddenError } from '../shared/errors'; import { executeQuery, createAuditLog } from '../shared/db'; import { createInvitationSchema, acceptInvitationSchema, ROLE_HIERARCHY } from '../shared/admin/types'; import { generateInvitationToken, hashToken, calculateExpiry, isExpired } from '../shared/admin/tokens'; import { sendEmail, generateInvitationEmail } from '../shared/admin/email'; import { CognitoIdentityProviderClient, AdminCreateUserCommand, AdminAddUserToGroupCommand, AdminSetUserMFAPreferenceCommand, } from '@aws-sdk/client-cognito-identity-provider';

const logger = new Logger({ handler: 'invitations' }); const cognitoClient = new CognitoIdentityProviderClient({});

export async function handler( event: APIGatewayProxyEvent, context: Context ): Promise<APIGatewayProxyResult> {
const requestLogger = logger.child({ requestId: context.awsRequestId, path: event.path });

try { const invitationId = event.pathParameters?.invitationId; const action = event.path.split('/').pop();

switch (event.httpMethod) {
  case 'GET':
    if (invitationId) return await handleGetInvitation(invitationId, event, requestLogger);
    return await handleListInvitations(event, requestLogger);

  case 'POST':
    if (action === 'accept') return await handleAcceptInvitation(event, requestLogger);
    if (action === 'resend' && invitationId) return await handleResendInvitation(invitationId, event, requestLogger);
    return await handleCreateInvitation(event, requestLogger);

  case 'DELETE':
    if (!invitationId) throw new ValidationError('Invitation ID required');
    return await handleRevokeInvitation(invitationId, event, requestLogger);

  default:
    throw new ValidationError(`Method ${event.httpMethod} not allowed`);
}

} catch (error) { return handleError(error, requestLogger); }

async function handleCreateInvitation(event: APIGatewayProxyEvent, logger: Logger): Promise<APIGatewayProxyResult> {
const auth = extractAuthContext(event); requireAdmin(auth); requirePermission(auth, 'admin:write');

const body = event.body ? JSON.parse(event.body) : {}; const parseResult = createInvitationSchema.safeParse(body); if (!parseResult.success) { throw new ValidationError('Invalid request body', parseResult.error.flatten().fieldErrors as Record<string, string[]>); }

const { email, role, message, expiresInHours } = parseResult.data;

// Validate role hierarchy if (ROLE_HIERARCHY[role] > ROLE_HIERARCHY[auth.role as keyof typeof ROLE_HIERARCHY]) { throw new ForbiddenError('Cannot invite administrator with');
}

```

```

higher role than your own'); }

// Check existing admin/invitation const existingAdmin = await executeQuery( 'SELECT id
FROM administrators WHERE email = $1 AND tenant_id = $2', [email, auth.tenantId], logger );
if (existingAdmin.rowCount > 0) { throw new ValidationError('An administrator with this
email already exists'); }

const pendingInvitation = await executeQuery( 'SELECT id FROM invitations WHERE email = $1
AND tenant_id = $2 AND status = \'pending\' AND expires_at > NOW()', [email, auth.tenantId], logger );
if (pendingInvitation.rowCount > 0) { throw new ValidationError('A pending invitation
already exists for this email'); }

// Generate token and create invitation const { token, tokenHash } = generateInvitationToken();
const expiresInAt = calculateExpiry(expiresInHours); const invitationId = uuidv4();

const inviterResult = await executeQuery('SELECT first_name, last_name FROM administrators
WHERE id = $1', [auth.userId], logger); const inviterName = inviterResult.rows[0] ? `${inviterRe-
sult.rows[0].first_name} ${inviterResult.rows[0].last_name}` : 'An administrator';

const appResult = await executeQuery('SELECT name FROM apps WHERE id = $1',
[auth.appId], logger); const appName = appResult.rows[0]?.name || auth.appId;

await executeQuery( 'INSERT INTO invitations (id, email, role, invited_by, app_id, tenant_id,
environment, token_hash, expires_at, status, message, created_at) VALUES ($1, $2, $3, $4,
$5, $6, $7, $8, $9, \'pending\', $10, NOW())', [invitationId, email, role, auth.userId, auth.appId,
auth.tenantId, auth.environment, tokenHash, expiresInAt, message || null], logger );

// Send email const acceptUrl = `${process.env.ADMIN_URL}/invite/accept?token=${token}`;
const emailContent = generateInvitationEmail({ inviteeName: '', inviterName, role, appName,
environment: auth.environment, acceptUrl, expiresInAt, message, }); await sendEmail({ to: email,
subject: `You've been invited to ${appName}`, html: emailContent.html, text: emailContent.text },
logger);

await createAuditLog({ tenant_id: auth.tenantId, user_id: null, admin_id: auth.userId,
action: 'invitation.create', resource_type: 'invitation', resource_id: invitationId, details: {
email, role, expiresInAt }, ip_address: event.requestContext.identity?.sourceIp || null, user_agent:
event.headers['User-Agent'] || null, }, logger);

logger.info('Invitation created and sent', { invitationId, email, role });
return created({ invitation:
{ id: invitationId, email, role, status: 'pending', expiresInAt, createdAt: new Date().toISOString() }
}); }

async function handleAcceptInvitation(event: APIGatewayProxyEvent, logger: Logger): Promise<any>
{ const body = event.body ? JSON.parse(event.body) : {};
const parseResult = acceptInvitationSchema.safeParse(body);
if (!parseResult.success) { throw new ValidationError('Invalid request
body', parseResult.error.flatten().fieldErrors as Record<string, string[]>); }

const { token, firstName, lastName, password, mfaMethod, phone } = parseResult.data;

const tokenHash = hashToken(token);
const inviteResult = await executeQuery( 'SELECT *
FROM invitations WHERE token_hash = $1 AND status = \'pending\'', [tokenHash], logger );
if (inviteResult.rowCount === 0) throw new NotFoundError('Invalid or expired invitation');
}

```

```

const invitation = inviteResult.rows[0]; if (isExpired(invitation.expires_at)) { await executeQuery(`UPDATE invitations SET status = 'expired' WHERE id = $1`, [invitation.id], logger); throw new ValidationError('This invitation has expired'); }

if (invitation.environment === 'prod' && mfaMethod === 'sms' && !phone) { throw new ValidationError('Phone number required for SMS MFA'); }

const userPoolId = process.env.ADMIN_USER_POOL_ID; const adminUserId = uuidv4();

// Create Cognito user await cognitoClient.send(new AdminCreateUserCommand({ UserPoolId: userPoolId, Username: invitation.email, TemporaryPassword: password, UserAttributes: [ { Name: 'email', Value: invitation.email }, { Name: 'email_verified', Value: 'true' }, { Name: 'given_name', Value: firstName }, { Name: 'family_name', Value: lastName }, { Name: 'custom:adminId', Value: adminUserId }, { Name: 'custom:tenantId', Value: invitation.tenant_id }, { Name: 'custom:role', Value: invitation.role }, ], MessageAction: 'SUPPRESS', }));;

await cognitoClient.send(new AdminAddUserToGroupCommand({ UserPoolId: userPoolId, Username: invitation.email, GroupName: invitation.role, }));;

if (invitation.environment === 'prod') { await cognitoClient.send(new AdminSetUserMFAPreferenceCommand({ UserPoolId: userPoolId, Username: invitation.email, SoftwareTokenMfaSettings: mfaMethod === 'authenticator' ? { Enabled: true, PreferredMfa: true } : undefined, SMSMfaSettings: mfaMethod === 'sms' ? { Enabled: true, PreferredMfa: true } : undefined, })); }

// Create admin user record await executeQuery( 'INSERT INTO administrators (id, cognito_user_id, email, first_name, last_name, display_name, role, app_id, tenant_id, mfa_enabled, mfa_method, status, created_at, updated_at, created_by, invitation_id) VALUES ($1, $2, $3, $4, $5, $6, $7, $8, $9, $10, $11, 'active', NOW(), NOW(), $12, $13)', [adminUserId, invitation.email, invitation.email, firstName, lastName, `${firstName} ${lastName}`, invitation.role, invitation.app_id, invitation.tenant_id, invitation.environment === 'prod', mfaMethod, invitation.invited_by, invitation.id], logger );;

// Create default profile await executeQuery( 'INSERT INTO admin_profiles (admin_id, notifications, timezone, language, date_format, time_format, currency, theme, default_environment, sidebar_collapsed, table_rows_per_page, updated_at) VALUES ($1, $2, 'America/New_York', 'en', 'MM/DD/YYYY', '12h', 'USD', 'system', $3, false, 25, NOW())', [adminUserId, JSON.stringify({ method: 'email', frequency: 'immediate', categories: { security: true, billing: true, deployments: true, approvals: true, system: true } }), invitation.environment], logger );;

await executeQuery( 'UPDATE invitations SET status = 'accepted', accepted_at = NOW(), accepted_by_ip = $2 WHERE id = $1', [invitation.id, event.requestContext.identity?.sourceIp || null], logger );

await createAuditLog({ tenant_id: invitation.tenant_id, user_id: null, admin_id: adminUserId, action: 'invitation.accept', resource_type: 'invitation', resource_id: invitation.id, details: { email: invitation.email, role: invitation.role, firstName, lastName }, ip_address: event.requestContext.identity?.sourceIp || null, user_agent: event.headers['User-Agent'] || null, }, logger);

logger.info('Invitation accepted', { invitationId: invitation.id, adminUserId, email: invitation.email }); return success({ message: 'Invitation accepted successfully', adminUser: { id:

```

```

adminUserId, email: invitation.email, firstName, lastName, role: invitation.role, mfaRequired: invitation.environment === 'prod' }, });
}

async function handleListInvitations(event: APIGatewayProxyEvent, logger: Logger): Promise<any> {
  const auth = extractAuthContext(event);
  requireAdmin(auth);
  requirePermission(auth, 'admin:read');

  const status = event.queryStringParameters?.status;
  const limit = parseInt(event.queryStringParameters?.limit || '50');
  const offset = parseInt(event.queryStringParameters?.offset || '0');

  let query = 'SELECT i.* , a.first_name AS inviter_first_name, a.last_name AS inviter_last_name FROM invitations i LEFT JOIN administrators a ON i.invited_by = a.id WHERE i.tenant_id = $1';
  const params: any[] = [auth.tenantId];

  if (status) {
    params.push(status);
    query += ' AND i.status = $$' + params.length;
  }
  query += ' ORDER BY i.created_at DESC LIMIT $$' + (params.length + 1) + ' OFFSET $$' + (params.length + 2);
  params.push(limit, offset);

  const result = await executeQuery(query, params, logger);
  const invitations = result.rows.map((row) => ({
    id: row.id, email: row.email, role: row.role, status: row.status, environment: row.environment, invitedBy: { id: row.invited_by, name: row.inviter_first_name ? `${row.inviter_first_name} ${row.inviter_last_name}` : 'Unknown' }, message: row.message, expiresAt: row.expires_at, createdAt: row.created_at, acceptedAt: row.accepted_at, }));
}

return success({ invitations, pagination: { limit, offset, hasMore: invitations.length === limit } });
}

async function handleGetInvitation(invitationId: string, event: APIGatewayProxyEvent, logger: Logger): Promise<any> {
  const auth = extractAuthContext(event);
  requireAdmin(auth);
  requirePermission(auth, 'admin:read');

  const result = await executeQuery(`SELECT i.* , a.first_name AS inviter_first_name, a.last_name AS inviter_last_name FROM invitations i LEFT JOIN administrators a ON i.invited_by = a.id WHERE i.id = $1 AND i.tenant_id = $2`, [invitationId, auth.tenantId], logger);
  if (result.rowCount === 0) throw newNotFoundError('Invitation not found');

  const row = result.rows[0];
  return success({ invitation: { id: row.id, email: row.email, role: row.role, status: row.status, environment: row.environment, invitedBy: { id: row.invited_by, name: `${row.inviter_first_name} ${row.inviter_last_name}` }, message: row.message, expiresAt: row.expires_at, createdAt: row.created_at, acceptedAt: row.accepted_at } });
}

async function handleResendInvitation(invitationId: string, event: APIGatewayProxyEvent, logger: Logger): Promise<any> {
  const auth = extractAuthContext(event);
  requireAdmin(auth);
  requirePermission(auth, 'admin:write');

  const result = await executeQuery(`SELECT * FROM invitations WHERE id = $1 AND tenant_id = $2`, [invitationId, auth.tenantId], logger);
  if (result.rowCount === 0) throw newNotFoundError('Invitation not found');

  const invitation = result.rows[0];
  if (invitation.status !== 'pending') throw new ValidationError(`Cannot resend ${invitation.status} invitation`);

  const { token, tokenHash } = generateInvitationToken();
  const expiresAt = calculateExpiry(48);
}

```

```

await executeQuery('UPDATE invitations SET token_hash = $2, expires_at = $3 WHERE id = $1', [invitationId, tokenHash, expiresAt], logger);

const inviterResult = await executeQuery('SELECT first_name, last_name FROM administrators WHERE id = $1', [auth.userId], logger); const inviterName = inviterResult.rows[0] ? `${inviterResult.rows[0].first_name} ${inviterResult.rows[0].last_name}` : 'An administrator';

const appResult = await executeQuery('SELECT name FROM apps WHERE id = $1', [auth.appId], logger); const appName = appResult.rows[0]?.name || auth.appId;

const acceptUrl = `${process.env.ADMIN_URL}/invite/accept?token=${token}`; const emailContent = generateInvitationEmail({ inviteeName: '', inviterName, role: invitation.role, appName, environment: invitation.environment, acceptUrl, expiresAt, message: invitation.message }); await sendEmail({ to: invitation.email, subject: 'Reminder: You\'ve been invited to ${appName}', html: emailContent.html, text: emailContent.text }, logger);

logger.info('Invitation resent', { invitationId, email: invitation.email }); return success({ message: 'Invitation resent successfully', expiresAt });

async function handleRevokeInvitation(invitationId: string, event: APIGatewayProxyEvent, logger: Logger): Promise<any> {
  const auth = extractAuthContext(event);
  requireAdmin(auth);
  requirePermission(auth, 'admin:write');

  const result = await executeQuery('SELECT * FROM invitations WHERE id = $1 AND tenant_id = $2', [invitationId, auth.tenantId], logger);
  if (result.rowCount === 0) throw new NotFoundError('Invitation not found');

  const invitation = result.rows[0];
  if (invitation.status !== 'pending') throw new ValidationException(`Cannot revoke ${invitation.status} invitation`);

  await executeQuery('UPDATE invitations SET status = \'revoked\' WHERE id = $1', [invitationId], logger);

  await createAuditLog({ tenant_id: auth.tenantId, user_id: null, admin_id: auth.userId, action: 'invitation.revoke', resource_type: 'invitation', resource_id: invitationId, details: { email: invitation.email }, ip_address: event.requestContext.identity?.sourceIp || null, user_agent: event.headers['User-Agent'] || null }, logger);

  logger.info('Invitation revoked', { invitationId, email: invitation.email });
  return success({ message: 'Invitation revoked successfully' });
}

```

---

### PART 3: TWO-PERSON APPROVALS LAMBDA

`packages/infrastructure/lambda/admin/approvals.ts`

```
``typescript /**
 * Two-Person Approval Workflow Lambda
 * Production deployments require separate initiator and approver
 */
```

```
import { APIGatewayProxyEvent, APIGatewayProxyResult, Context } from 'aws-lambda';
import { v4 as uuidv4 } from 'uuid';
import { Logger } from '../shared/logger';
import { success, created, handleError } from '../shared/response';
import { extractAuthContext, requireAdmin, requirePermission } from '../shared/auth';
import { ValidationException, NotFoundError, ForbiddenError } from
```

```

'./shared/errors'; import { executeQuery, createAuditLog } from './shared/db'; import { createApprovalSchema, processApprovalSchema, ApprovalStatus, AdminRole, ROLE_HIERARCHY } from './shared/admin/types'; import { calculateExpiry, isExpired } from './shared/admin/tokens'; import { sendEmail, generateApprovalEmail } from './shared/admin/email';

const logger = new Logger({ handler: 'approvals' });

export async function handler(event: APIGatewayProxyEvent, context: Context): Promise<any> {
    const requestLogger = logger.child({ requestId: context.awsRequestId, path: event.path });

    try {
        const auth = extractAuthContext(event);
        requireAdmin(auth);

        const approvalId = event.pathParameters?.approvalId;
        const action = event.path.split('/').pop();

        switch (event.httpMethod) {
            case 'GET':
                if (approvalId) return await handleGetApproval(approvalId, auth, requestLogger);
                return await handleListApprovals(event, auth, requestLogger);
            case 'POST':
                if (action === 'process' && approvalId) return await handleProcessApproval(approvalId, event);
                return await handleCreateApproval(event, auth, requestLogger);
            case 'DELETE':
                if (!approvalId) throw new ValidationError('Approval ID required');
                return await handleCancelApproval(approvalId, event, auth, requestLogger);
            default:
                throw new ValidationError(`Method ${event.httpMethod} not allowed`);
        }
    } catch (error) {
        return handleError(error, requestLogger);
    }

    async function handleCreateApproval(event: APIGatewayProxyEvent, auth: ReturnType, logger: Logger): Promise<any> {
        const body = event.body ? JSON.parse(event.body) : {};
        const parseResult = createApprovalSchema.safeParse(body);
        if (!parseResult.success) throw new ValidationError('Invalid request body', parseResult.error.flatten().fieldErrors as Record<string, string[]>);

        const { type, action, resourceType, resourceId, details, priority, notes, expiresInHours } = parseResult.data;
        const requiresTwoPersonApproval = auth.environment === 'prod';

        const requesterResult = await executeQuery('SELECT first_name, last_name FROM administrators WHERE id = $1', [auth.userId], logger);
        const requesterName = requesterResult.rows[0] ? `${requesterResult.rows[0].first_name} ${requesterResult.rows[0].last_name}` : 'Unknown';

        const approvalId = uuidv4();
        const expiresAt = calculateExpiry(expiresInHours);

        await executeQuery(`INSERT INTO approval_requests (id, type, app_id, tenant_id, environment, requested_by, requested_at, expires_at, status, action, resource_type, resource_id, details, priority, notes, requires_two_person, created_at) VALUES ($1, $2, $3, $4, $5, $6, NOW(), $7, 'pending', $8, $9, $10, $11, $12, $13, $14, NOW())`, [approvalId, type, auth.appId, auth.tenantId, auth.environment, auth.userId, expiresAt, action, resourceType, resourceId, JSON.stringify(details), priority, notes || null, requiresTwoPersonApproval], logger);
    }
}

```

```

if (requiresTwoPersonApproval) { await notifyApprovers(approvalId, auth, requesterName, logger);
}

await createAuditLog({ tenant_id: auth.tenantId, user_id: null, admin_id: auth.userId, action: 'approval.create', resource_type: 'approval_request', resource_id: approvalId, details: { type, resourceType, resourceId, requiresTwoPersonApproval }, ip_address: event.requestContext.identity?.sourceIp || null, user_agent: event.headers['User-Agent'] || null, }, logger);

logger.info('Approval request created', { approvalId, type, requestedBy: auth.userId, requiresTwoPersonApproval }); return created({ approval: { id: approvalId, type, status: 'pending', action, resourceType, resourceId, priority, expiresAt, requiresTwoPersonApproval, createdAt: new Date().toISOString() }, }, );
}

async function handleProcessApproval(approvalId: string, event: APIGatewayProxyEvent, auth: ReturnType, logger: Logger): Promise { requirePermission(auth, 'approvals:*');

const body = event.body ? JSON.parse(event.body) : {}; const parseResult = processApprovalSchema.safeParse(body); if (!parseResult.success) throw new ValidationError('Invalid request body', parseResult.error.flatten().fieldErrors as Record<string, string[]>);

const { action, reason } = parseResult.data;

const result = await executeQuery('SELECT * FROM approval_requests WHERE id = $1 AND tenant_id = $2', [approvalId, auth.tenantId], logger); if (result.rowCount === 0) throw new NotFoundError('Approval request not found');

const approval = result.rows[0]; if (approval.status !== 'pending') throw new ValidationError('Cannot process ${approval.status} approval request'); if (isExpired(approval.expires_at)) { await executeQuery('UPDATE approval_requests SET status = \'expired\' WHERE id = $1', [approvalId], logger); throw new ValidationError('This approval request has expired'); }

// Two-person approval: cannot approve own request if (approval.requires_two_person && approval.requested_by === auth.userId) { throw new ForbiddenError('You cannot approve your own request. Production deployments require approval from a different administrator.'); }

const newStatus: ApprovalStatus = action === 'approve' ? 'approved' : 'rejected'; await executeQuery('UPDATE approval_requests SET status = $2, approved_by = $3, approved_at = NOW(), rejected_reason = $4 WHERE id = $1', [approvalId, newStatus, auth.userId, action === 'reject' ? reason : null], logger );

if (action === 'approve') { await executeApprovedAction(approval, logger); }

await createAuditLog({ tenant_id: auth.tenantId, user_id: null, admin_id: auth.userId, action: `approval.${action}`, resource_type: 'approval_request', resource_id: approvalId, details: { type: approval.type, resourceType: approval.resource_type, resourceId: approval.resource_id, reason }, ip_address: event.requestContext.identity?.sourceIp || null, user_agent: event.headers['User-Agent'] || null, }, logger);

logger.info('Approval request processed', { approvalId, action, processedBy: auth.userId }); return success({ approval: { id: approvalId, status: newStatus, processedBy: auth.userId, processedAt: new Date().toISOString(), reason: action === 'reject' ? reason : undefined } });
}

async function executeApprovedAction(approval: any, logger: Logger): Promise { logger.info('Executing approved action', { type: approval.type, resourceType: approval.resource_type,

```

```

resourceId: approval.resource_id });

switch (approval.type) { case 'deployment': // Trigger deployment workflow via Step Functions or CodePipeline break; case 'promotion': await executeQuery('UPDATE deployments SET promoted_to = $2, promoted_at = NOW() WHERE id = $1', [approval.resource_id, approval.details?.targetEnv], logger); break; case 'model_activation': await executeQuery('UPDATE ai_models SET status = $2, thermal_state = $3, updated_at = NOW() WHERE id = $1', [approval.details?.modelId, approval.details?.newStatus, approval.details?.thermalState], logger); break; case 'provider_change': await executeQuery('UPDATE ai_providers SET config = $2, updated_at = NOW() WHERE id = $1', [approval.details?.providerId, JSON.stringify(approval.details?.config)], logger); break; case 'user_role_change': await executeQuery('UPDATE administrators SET role = $2, updated_at = NOW() WHERE id = $1', [approval.details?.userId, approval.details?.newRole], logger); break; case 'billing_change': await executeQuery('UPDATE billing_settings SET margin_percent = COALESCE($2, margin_percent), tax_percent = COALESCE($3, tax_percent), updated_at = NOW() WHERE tenant_id = $1', [approval.details?.tenantId, approval.details?.marginPercent, approval.details?.taxPercent], logger); break; default: logger.warn('Unknown approval type', { type: approval.type }) } }

async function notifyApprovers(approvalId: string, auth: ReturnType, requesterName: string, logger: Logger): Promise {
  const result = await executeQuery(`SELECT id, email, first_name, last_name FROM administrators WHERE tenant_id = $1 AND id != $2 AND status = 'active' AND role IN ('super_admin', 'admin')`, [auth.tenantId, auth.userId], logger);
  if (result.rowCount === 0) {
    logger.warn('No other admins available to approve');
    return;
  }
}

const appResult = await executeQuery('SELECT name FROM apps WHERE id = $1', [auth.appId], logger);
const appName = appResult.rows[0] ?.name || auth.appId;

const approvalResult = await executeQuery('SELECT * FROM approval_requests WHERE id = $1', [approvalId], logger);
const approval = approvalResult.rows[0];

for (const admin of result.rows) {
  const approveUrl = `${process.env.ADMIN_URL}/approvals/${approvalId}?action=approve`;
  const rejectUrl = `${process.env.ADMIN_URL}/approvals/${approvalId}?action=reject`;
  const emailContent = generateApprovalEmail({
    approverName: admin.first_name,
    requesterName,
    appName,
    environment: auth.environment,
    action: approval.action,
    resourceType: approval.resource_type,
    resourceId: approval.resource_id,
    approveUrl,
    rejectUrl,
    expiresAt: approval.expires_at
  });
  await sendEmail({
    to: admin.email,
    subject: `Approval Required: ${approval.action} - ${appName}`,
    html: emailContent.html,
    text: emailContent.text
  }, logger);
  logger.info('Approvers notified', { approvalId, notifiedCount: result.rowCount });
}

async function handleListApprovals(event: APIGatewayProxyEvent, auth: ReturnType, logger: Logger): Promise {
  requirePermission(auth, 'approvals:read');
}

const status = event.queryStringParameters?.status;
const pendingForMe = event.queryStringParameters?.pending === 'true';
const limit = parseInt(event.queryStringParameters?.limit || '50');
const offset = parseInt(event.queryStringParameters?.offset || '0');

let query = 'SELECT ar.*,
req.first_name AS requester_first_name,
req.last_name AS requester_last_name
FROM approval_requests ar
LEFT JOIN administrators req
ON ar.requested_by = req.id
WHERE ar.tenant_id = $1';
const params: any[] = [auth.tenantId];
if (status) {
  params.push(status);
  query += ` AND ar.status = $$${params.length}`;
}
if (pendingForMe) {
  query += ` AND req.id = ${auth.userId}`;
}
query += ` LIMIT ${limit} OFFSET ${offset}`;
```

```

ForMe) { params.push(auth.userId); query += ' AND ar.status = 'pending' AND ar.requested_by != $$\{params.length\}'; } query += ' ORDER BY ar.created_at DESC LIMIT $$\{params.length + 1\} OFFSET $$\{params.length + 2\}'; params.push(limit, offset);

const result = await executeQuery(query, params, logger); const approvals = result.rows.map(row => ({ id: row.id, type: row.type, status: row.status, environment: row.environment, action: row.action, resourceType: row.resource_type, resourceId: row.resource_id, priority: row.priority, requestedBy: { id: row.requested_by, name: '$\{row.requester_first_name\} $\{row.requester_last_name\}' }, requestedAt: row.requested_at, expiresAt: row.expires_at, requiresTwoPersonApproval: row.requires_two_person, canApprove: row.status === 'pending' && row.requested_by !== auth.userId, }));}

return success({ approvals, pagination: { limit, offset, hasMore: approvals.length === limit } });

async function handleGetApproval(approvalId: string, auth: ReturnType, logger: Logger): Promise { requirePermission(auth, 'approvals:read');

const result = await executeQuery( 'SELECT ar.* , req.first_name as requester_first_name, req.last_name as requester_last_name, req.email as requester_email FROM approval_requests ar LEFT JOIN administrators req ON ar.requested_by = req.id WHERE ar.id = $1 AND ar.tenant_id = $2', [approvalId, auth.tenantId], logger ); if (result.rowCount === 0) throw newNotFoundError('Approval request not found');

const row = result.rows[0]; return success({ approval: { id: row.id, type: row.type, status: row.status, environment: row.environment, action: row.action, resourceType: row.resource_type, resourceId: row.resource_id, details: row.details, priority: row.priority, notes: row.notes, requestedBy: { id: row.requested_by, name: '$\{row.requester_first_name\} $\{row.requester_last_name\}', email: row.requester_email }, requestedAt: row.requested_at, expiresAt: row.expires_at, requiresTwoPersonApproval: row.requires_two_person, canApprove: row.status === 'pending' && row.requested_by !== auth.userId, } });

async function handleCancelApproval(approvalId: string, event: APIGatewayProxyEvent, auth: ReturnType, logger: Logger): Promise { const result = await executeQuery('SELECT * FROM approval_requests WHERE id = $1 AND tenant_id = $2', [approvalId, auth.tenantId], logger); if (result.rowCount === 0) throw newNotFoundError('Approval request not found');

const approval = result.rows[0]; if (approval.requested_by !== auth.userId && auth.role !== AdminRole.SUPER_ADMIN) { throw new ForbiddenError('Only the requester or a super admin can cancel this request'); } if (approval.status !== 'pending') throw new ValidationError('Cannot cancel ${approval.status} approval request');

await executeQuery('UPDATE approval_requests SET status = 'cancelled' WHERE id = $1', [approvalId], logger); await createAuditLog({ tenant_id: auth.tenantId, user_id: null, admin_id: auth.userId, action: 'approval.cancel', resource_type: 'approval_request', resource_id: approvalId, details: { type: approval.type }, ip_address: event.requestContext.identity?.sourceIp || null, user_agent: event.headers['User-Agent'] || null, }, logger);

logger.info('Approval request cancelled', { approvalId }); return success({ message: 'Approval request cancelled' });
}

```

---

## PART 4: METERING LAMBDA

`packages/infrastructure/lambda/billing/metering.ts`

```
“typescript /**
 * Usage Metering Lambda * Collects and stores usage events for billing calculations
 */

import { APIGatewayProxyEvent, APIGatewayProxyResult, Context } from ‘aws-lambda’; import
{ DynamoDBClient } from ‘@aws-sdk/client-dynamodb’; import { DynamoDBDocumentClient,
PutCommand, QueryCommand, UpdateCommand } from ‘@aws-sdk/lib-dynamodb’; import { v4
as uuidv4 } from ‘uuid’; import { z } from ‘zod’; import { Logger } from ‘./shared/logger’; import
{ success, created, handleError } from ‘./shared/response’; import { extractAuthContext } from
‘./shared/auth’; import { ValidationError } from ‘./shared/errors’; import { executeQuery } from
‘./shared/db’;

const logger = new Logger({ handler: ‘metering’ }); const ddbClient = DynamoDBDocument-
Client.from(new DynamoDBClient({}), { marshallOptions: { removeUndefinedValues: true } });

const USAGE_TABLE = process.env.USAGE_TABLE || ‘radiant-usage-events’; const
ROLLUP_TABLE = process.env.ROLLUP_TABLE || ‘radiant-usage-rollups’;

const recordUsageSchema = z.object({ requestId: z.string(), providerId: z.string(), modelId:
z.string(), modelName: z.string(), requestType: z.enum(['chat', ‘embedding’, ‘image’, ‘audio’,
‘video’]), inputTokens: z.number().int().min(0), outputTokens: z.number().int().min(0), latencyMs:
z.number().int().min(0), cached: z.boolean().default(false), phiDetected: z.boolean().default(false),
phiSanitized: z.boolean().default(false), userId: z.string().optional(), });

export async function handler(event: APIGatewayProxyEvent, context: Context): Promise {
    const requestLogger = logger.child({ requestId: context.awsRequestId, path: event.path });

    try {
        const auth = extractAuthContext(event);
        const action = event.path.split('/').pop();

        switch (event.httpMethod) {
            case ‘POST’:
                if (action === ‘record’) return await handleRecordUsage(event, auth, requestLogger);
                if (action === ‘batch’) return await handleBatchRecord(event, auth, requestLogger);
                break;
            case ‘GET’:
                if (action === ‘summary’) return await handleGetSummary(event, auth, requestLogger);
                if (action === ‘rollups’) return await handleGetRollups(event, auth, requestLogger);
                break;
        }
        throw new ValidationError(`Unknown action: ${action}`);
    } catch (error) {
        return handleError(error, requestLogger);
    }
}

async function handleRecordUsage(event: APIGatewayProxyEvent, auth: ReturnType, logger: Logger):
Promise {
    const body = event.body ? JSON.parse(event.body) : {};
    const parseResult = recordUsageSchema.safeParse(body);
    if (!parseResult.success) throw new ValidationError(`Invalid usage data`,
parseResult.error.flatten().fieldErrors as Record<string, string[]>);

    const data = parseResult.data;
    // Get pricing from model registry
    const pricing = await getModelPricing(data.modelId, logger);
}
```

```

const inputCost = (data.inputTokens / 1000000) * pricing.inputPricePerMillion; const outputCost
= (data.outputTokens / 1000000) * pricing.outputPricePerMillion; const providerCost = inputCost
+ outputCost;

// Apply margin const marginPercent = await getTenantMargin(auth.tenantId, logger); const
billedCost = providerCost * (1 + marginPercent / 100);

const usageEvent = { id: uuidv4(), timestamp: new Date().toISOString(), tenantId: auth.tenantId,
userId: data.userId || auth.userId, adminId: auth.isAdmin ? auth.userId : undefined, appId:
auth.appId, environment: auth.environment, providerId: data.providerId, modelId: data.modelId,
modelName: data.modelName, requestType: data.requestType, inputTokens: data.inputTokens,
outputTokens: data.outputTokens, totalTokens: data.inputTokens + data.outputTokens, provider-
Cost, billedCost, currency: 'USD', requestId: data.requestId, latencyMs: data.latencyMs, cached:
data.cached, phiDetected: data.phiDetected, phiSanitized: data.phiSanitized, };

await ddbClient.send(new PutCommand({ TableName: USAGE_TABLE, Item: { pk: `TEN-
ANT#${auth.tenantId}`, sk: `EVENT#${usageEvent.timestamp}#${usageEvent.id}`, ...us-
ageEvent, ttl: Math.floor(Date.now() / 1000) + (90 * 24 * 60 * 60) }, }));
```

await updateDailyRollup(usageEvent, logger);

```

logger.info('Usage recorded', { eventId: usageEvent.id, modelId: data.modelId, tokens: us-
ageEvent.totalTokens, cost: billedCost }); return created({ event: { id: usageEvent.id, billedCost,
providerCost } });
}

async function handleBatchRecord(event: APIGatewayProxyEvent, auth: ReturnType, logger: Logger): Promise {
  const body = event.body ? JSON.parse(event.body) : {};
  if (!Array.isArray(body.events) || body.events.length === 0) throw new ValidationError('events array
is required');
  if (body.events.length > 100) throw new ValidationError('Maximum 100 events per
batch');

  const results: Array<{ id: string; success: boolean; error?: string }> = [];
  for (const eventData of body.events) {
    try {
      const parseResult = recordUsageSchema.safeParse(eventData);
      if (!parseResult.success) {
        results.push({ id: eventData.requestId || 'unknown', success: false, error:
`Invalid data` });
        continue;
      }

      const data = parseResult.data;
      const pricing = await getModelPricing(data.modelId, logger);
      const providerCost = ((data.inputTokens / 1000000) * pricing.inputPricePerMillion) +
((data.o
      const marginPercent = await getTenantMargin(auth.tenantId, logger);
      const billedCost = providerCost * (1 + marginPercent / 100);

      const usageEvent = {
        id: uuidv4(), timestamp: new Date().toISOString(), tenantId: auth.tenantId, userId: data.us
        appId: auth.appId, environment: auth.environment, providerId: data.providerId, modelId: dat
        requestType: data.requestType, inputTokens: data.inputTokens, outputTokens: data.outputTok
        providerCost, billedCost, currency: 'USD', requestId: data.requestId, latencyMs: data.late
      };

      await ddbClient.send(new PutCommand({ TableName: USAGE_TABLE, Item: { pk: `TENANT#${auth.t
      await updateDailyRollup(usageEvent, logger);
    }
  }
}

```

```

        results.push({ id: usageEvent.id, success: true });
    } catch (error: any) {
        results.push({ id: eventData.requestId || 'unknown', success: false, error: error.message });
    }
}

logger.info('Batch usage recorded', { total: body.events.length, successful: results.filter(r => r.success).length });
return success({ results, summary: { total: results.length, successful: results.filter(r => r.success).length, failed: results.filter(r => !r.success).length } });
}

async function handleGetSummary(event: APIGatewayProxyEvent, auth: ReturnType, logger: Logger): Promise<any> {
    const startDate = event.queryStringParameters?.startDate || getDefaultStartDate();
    const endDate = event.queryStringParameters?.endDate || getTodayDate();

    const response = await ddbClient.send(new QueryCommand({
        TableName: ROLLUP_TABLE,
        KeyConditionExpression: 'pk = :pk AND sk BETWEEN :start AND :end',
        ExpressionAttributeValues: { ':pk': 'TENANT#${auth.tenantId}', ':start': 'DATE#${startDate}', ':end': 'DATE#${endDate}#~' },
    }));

    const rollups = response.Items || [];
    const totals = rollups.reduce((acc, r) => ({ requests: acc.requests + (r.requestCount || 0), inputTokens: acc.inputTokens + (r.inputTokens || 0), outputTokens: acc.outputTokens + (r.outputTokens || 0), providerCost: acc.providerCost + (r.providerCost || 0), billedCost: acc.billedCost + (r.billedCost || 0), }), { requests: 0, inputTokens: 0, outputTokens: 0, providerCost: 0, billedCost: 0 });

    return success({ period: { startDate, endDate }, totals });
}

async function handleGetRollups(event: APIGatewayProxyEvent, auth: ReturnType, logger: Logger): Promise<any> {
    const startDate = event.queryStringParameters?.startDate || getDefaultStartDate();
    const endDate = event.queryStringParameters?.endDate || getTodayDate();

    const response = await ddbClient.send(new QueryCommand({
        TableName: ROLLUP_TABLE,
        KeyConditionExpression: 'pk = :pk AND sk BETWEEN :start AND :end',
        ExpressionAttributeValues: { ':pk': 'TENANT#${auth.tenantId}', ':start': 'DATE#${startDate}', ':end': 'DATE#${endDate}#~' },
    }));

    const rollups = (response.Items || []).map(item => ({ date: item.date, modelId: item.modelId, providerId: item.providerId, requestCount: item.requestCount, inputTokens: item.inputTokens, outputTokens: item.outputTokens, totalTokens: item.totalTokens, providerCost: item.providerCost, billedCost: item.billedCost, avgLatencyMs: item.avgLatencyMs }));
}

return success({ rollups, period: { startDate, endDate } });

async function updateDailyRollup(event: any, logger: Logger): Promise<any> {
    const date = event.timestamp.split('T')[0];
    try {
        await ddbClient.send(new UpdateCommand({
            TableName: ROLLUP_TABLE,
            Key: { pk: 'TENANT#${event.tenantId}', sk: 'DATE#${date}#MODEL#${event.modelId}' },
            UpdateExpression: 'SET #date = :date, modelId = :modelId, providerId = :providerId, requestCount = if_not_exists(requestCount, :zero) + :one, inputTokens = if_not_exists(inputTokens, :zero) + :inputTokens, outputTokens = if_not_exists(outputTokens, :zero) + :outputTokens, totalTokens = if_not_exists(totalTokens, :zero) + :totalTokens, providerCost = if_not_exists(providerCost, :zero) + :providerCost, billedCost = if_not_exists(billedCost, :zero) + :billedCost, cachedRequests = if_not_exists(cachedRequests, :zero) + :cached, phiRequests = if_not_exists(phiRequests, :zero) + :phiRequests',
            ExpressionAttributeNames: { '#date': 'date' },
            ExpressionAttributeValues: { ':date': date, ':modelId': event.modelId, ':providerId': event.providerId, ':one': 1, ':zero': 0, ':inputTokens': 0, ':outputTokens': 0, ':totalTokens': 0, ':providerCost': 0, ':billedCost': 0, ':cached': 0, ':phiRequests': 0 }
        }));
    } catch (err) {
        logger.error(`Error updating daily rollup for date ${date}: ${err}`);
    }
}

```

```

if_not_exists(phiRequests, :zero) + :phi, updatedAt = :updatedAt', ExpressionAttributeNames:
{ ':#date': 'date' }, ExpressionAttributeValues: { ':date': date, ':modelId': event.modelId,
':providerId': event.providerId, ':zero': 0, ':one': 1, ':inputTokens': event.inputTokens, ':outputTo-
kens': event.outputTokens, ':totalTokens': event.totalTokens, ':providerCost': event.providerCost,
':billedCost': event.billedCost, ':cached': event.cached ? 1 : 0, ':phi': event.phiDetected ? 1 : 0,
':updatedAt': new Date().toISOString(), }, })); } catch (error) { logger.error('Failed to update
rollup', error as Error, { tenantId: event.tenantId, date, modelId: event.modelId }); } }

async function getModelPricing(modelId: string, logger: Logger): Promise<{ inputPricePerMil-
lion: number; outputPricePerMillion: number }> { const result = await executeQuery(`SELECT
input_price_per_million, output_price_per_million FROM ai_models WHERE id = $1`, [mod-
elId], logger); if (result.rowCount === 0) return { inputPricePerMillion: 1.0, outputPricePerMil-
lion: 2.0 }; return { inputPricePerMillion: parseFloat(result.rows[0].input_price_per_million) ||
1.0, outputPricePerMillion: parseFloat(result.rows[0].output_price_per_million) || 2.0 }; }

async function getTenantMargin(tenantId: string, logger: Logger): Promise<{ margin_percent: number }> { const result = await executeQuery(`SELECT margin_percent FROM billing_settings WHERE tenant_id = $1`, [ten-
antId], logger); return result.rowCount > 0 ? parseFloat(result.rows[0].margin_percent) : 20; }

function getDefaultStartDate(): string { const d = new Date(); d.setDate(d.getDate() - 30); return d.toISOString().split('T')[0]; } function getTodayDate(): string { return new
Date().toISOString().split('T')[0]; } ""

```

---

## PART 5: DATABASE SCHEMA ADDITIONS

[packages/infrastructure/migrations/005\\_admin\\_billing.sql](#)

```

--sql =====
--RADIANT v2.2.0 - Admin & Billing Schema =====

-- Admin Invitations CREATE TABLE IF NOT EXISTS invitations ( id UUID PRIMARY KEY
DEFAULT gen_random_uuid(), email VARCHAR(255) NOT NULL, role VARCHAR(50) NOT
NULL, invited_by UUID NOT NULL REFERENCES administrators(id), app_id VARCHAR(100)
NOT NULL, tenant_id VARCHAR(100) NOT NULL, environment VARCHAR(20) NOT NULL,
token_hash VARCHAR(64) NOT NULL, expires_at TIMESTAMP WITH TIME ZONE NOT
NULL, status VARCHAR(20) NOT NULL DEFAULT 'pending', message TEXT, accepted_at
TIMESTAMP WITH TIME ZONE, accepted_by_ip VARCHAR(45), created_at TIMESTAMP
WITH TIME ZONE NOT NULL DEFAULT NOW(), CONSTRAINT valid_invitation_status
CHECK (status IN ('pending', 'accepted', 'expired', 'revoked')) );

CREATE INDEX idx_invitations_email ON invitations(email); CREATE INDEX idx_invitations_tenant
ON invitations(tenant_id); CREATE INDEX idx_invitations_token ON invitations(token_hash);
CREATE INDEX idx_invitations_status ON invitations(status);

-- Approval Requests CREATE TABLE IF NOT EXISTS approval_requests ( id UUID PRIMARY
KEY DEFAULT gen_random_uuid(), type VARCHAR(50) NOT NULL, app_id VARCHAR(100)
NOT NULL, tenant_id VARCHAR(100) NOT NULL, environment VARCHAR(20) NOT NULL,
requested_by UUID NOT NULL REFERENCES administrators(id), requested_at TIMESTAMP
WITH TIME ZONE NOT NULL DEFAULT NOW(), expires_at TIMESTAMP WITH TIME
ZONE NOT NULL, status VARCHAR(20) NOT NULL DEFAULT 'pending', approved_by

```

UUID REFERENCES administrators(id), approved\_at TIMESTAMP WITH TIME ZONE, rejected\_reason TEXT, action VARCHAR(100) NOT NULL, resource\_type VARCHAR(100) NOT NULL, resource\_id VARCHAR(255), details JSONB, priority VARCHAR(20) NOT NULL DEFAULT ‘medium’, notes TEXT, requires\_two\_person BOOLEAN NOT NULL DEFAULT false, created\_at TIMESTAMP WITH TIME ZONE NOT NULL DEFAULT NOW(), CONSTRAINT valid\_approval\_status CHECK (status IN (‘pending’, ‘approved’, ‘rejected’, ‘expired’, ‘cancelled’)), CONSTRAINT valid\_priority CHECK (priority IN (‘low’, ‘medium’, ‘high’, ‘critical’)) );

CREATE INDEX idx\_approvals\_tenant ON approval\_requests(tenant\_id); CREATE INDEX idx\_approvals\_status ON approval\_requests(status); CREATE INDEX idx\_approvals\_requested\_by ON approval\_requests(requested\_by);

– Admin Profiles CREATE TABLE IF NOT EXISTS admin\_profiles ( admin\_id UUID PRIMARY KEY REFERENCES administrators(id) ON DELETE CASCADE, notifications JSONB NOT NULL DEFAULT ‘{}’, timezone VARCHAR(50) NOT NULL DEFAULT ‘America/New\_York’, language VARCHAR(10) NOT NULL DEFAULT ‘en’, date\_format VARCHAR(20) NOT NULL DEFAULT ‘MM/DD/YYYY’, time\_format VARCHAR(10) NOT NULL DEFAULT ‘12h’, currency VARCHAR(3) NOT NULL DEFAULT ‘USD’, theme VARCHAR(20) NOT NULL DEFAULT ‘system’, default\_environment VARCHAR(20) NOT NULL DEFAULT ‘dev’, sidebar\_collapsed BOOLEAN NOT NULL DEFAULT false, table\_rows\_per\_page INTEGER NOT NULL DEFAULT 25, updated\_at TIMESTAMP WITH TIME ZONE NOT NULL DEFAULT NOW() );

– Billing Settings CREATE TABLE IF NOT EXISTS billing\_settings ( tenant\_id VARCHAR(100) PRIMARY KEY, margin\_percent DECIMAL(5,2) NOT NULL DEFAULT 20.00, margin\_type VARCHAR(20) NOT NULL DEFAULT ‘fixed’, tiers JSONB, tax\_enabled BOOLEAN NOT NULL DEFAULT false, tax\_percent DECIMAL(5,2) NOT NULL DEFAULT 0.00, tax\_id VARCHAR(50), stripe\_customer\_id VARCHAR(100), default\_payment\_method\_id VARCHAR(100), auto\_pay BOOLEAN NOT NULL DEFAULT false, billing\_cycle\_day INTEGER NOT NULL DEFAULT 1, currency VARCHAR(3) NOT NULL DEFAULT ‘USD’, budget\_limit DECIMAL(12,2), alert\_thresholds INTEGER[] NOT NULL DEFAULT ‘{50, 75, 90, 100}’, created\_at TIMESTAMP WITH TIME ZONE NOT NULL DEFAULT NOW(), updated\_at TIMESTAMP WITH TIME ZONE NOT NULL DEFAULT NOW(), CONSTRAINT valid\_margin\_type CHECK (margin\_type IN (‘fixed’, ‘tiered’)), CONSTRAINT valid\_billing\_day CHECK (billing\_cycle\_day BETWEEN 1 AND 28) );

– Invoices CREATE TABLE IF NOT EXISTS invoices ( id UUID PRIMARY KEY DEFAULT gen\_random\_uuid(), tenant\_id VARCHAR(100) NOT NULL, app\_id VARCHAR(100) NOT NULL, period\_start DATE NOT NULL, period\_end DATE NOT NULL, subtotal DECIMAL(12,2) NOT NULL, margin\_percent DECIMAL(5,2) NOT NULL, tax DECIMAL(12,2) NOT NULL DEFAULT 0, tax\_percent DECIMAL(5,2) NOT NULL DEFAULT 0, total DECIMAL(12,2) NOT NULL, currency VARCHAR(3) NOT NULL DEFAULT ‘USD’, status VARCHAR(20) NOT NULL DEFAULT ‘draft’, due\_date TIMESTAMP WITH TIME ZONE NOT NULL, paid\_at TIMESTAMP WITH TIME ZONE, line\_items JSONB NOT NULL DEFAULT ‘[]’, stripe\_invoice\_id VARCHAR(100), stripe\_payment\_intent\_id VARCHAR(100), created\_at TIMESTAMP WITH TIME ZONE NOT NULL DEFAULT NOW(), updated\_at TIMESTAMP WITH TIME ZONE NOT NULL DEFAULT NOW(), CONSTRAINT valid\_invoice\_status CHECK (status IN (‘draft’, ‘pending’, ‘paid’, ‘overdue’, ‘cancelled’)) );

CREATE INDEX idx\_invoices\_tenant ON invoices(tenant\_id); CREATE INDEX idx\_invoices\_status

```
ON invoices(status); CREATE INDEX idx_invoices_period ON invoices(period_start, period_end);
```

– Admin Notifications CREATE TABLE IF NOT EXISTS admin\_notifications ( id UUID PRIMARY KEY DEFAULT gen\_random\_uuid(), admin\_id UUID NOT NULL REFERENCES administrators(id) ON DELETE CASCADE, tenant\_id VARCHAR(100) NOT NULL, type VARCHAR(50) NOT NULL, priority VARCHAR(20) NOT NULL DEFAULT ‘medium’, title VARCHAR(255) NOT NULL, message TEXT NOT NULL, action\_url VARCHAR(500), action\_label VARCHAR(100), read BOOLEAN NOT NULL DEFAULT false, read\_at TIMESTAMP WITH TIME ZONE, dismissed BOOLEAN NOT NULL DEFAULT false, dismissed\_at TIMESTAMP WITH TIME ZONE, email\_sent BOOLEAN NOT NULL DEFAULT false, email\_sent\_at TIMESTAMP WITH TIME ZONE, created\_at TIMESTAMP WITH TIME ZONE NOT NULL DEFAULT NOW(), expires\_at TIMESTAMP WITH TIME ZONE, CONSTRAINT valid\_notification\_type CHECK (type IN (‘security’, ‘billing’, ‘deployment’, ‘approval’, ‘system’, ‘alert’)), CONSTRAINT valid\_notification\_priority CHECK (priority IN (‘low’, ‘medium’, ‘high’, ‘critical’)) );

```
CREATE INDEX idx_notifications_admin ON admin_notifications(admin_id); CREATE INDEX idx_notifications_read ON admin_notifications(admin_id, read); ““
```

---

## API ROUTES SUMMARY

### Admin Routes

| Method | Path                     | Description             | Permission         |
|--------|--------------------------|-------------------------|--------------------|
| POST   | /admin/invitation        | Create invitation       | admin:write        |
| GET    | /admin/invitation        | List invitations        | admin:read         |
| GET    | /admin/invitation/:id    | Get invitation          | admin:read         |
| POST   | /admin/invitation/send   | Send invitation         | admin:write        |
| DELETE | /admin/invitation/revoke | Revoke invitation       | admin:write        |
| POST   | /admin/invitation/accept | Accept invitation       | (public)           |
| POST   | /admin/approval          | Create approval request | approvals:initiate |
| GET    | /admin/approval          | List approvals          | approvals:read     |
| GET    | /admin/approval/:id      | Get approval            | approvals:read     |
| POST   | /admin/approval/approve  | Approve request         | approvals:*        |
| DELETE | /admin/approval/cancel   | Cancel approval         | approvals:initiate |
| GET    | /admin/users             | List admins             | admin:read         |
| GET    | /admin/users/:id         | Get admin               | admin:read         |
| PUT    | /admin/users/:id         | Update admin            | admin:write        |
| DELETE | /admin/users/:id         | Delete admin            | admin:*            |

### Billing Routes

| Method | Path              | Description     | Permission   |
|--------|-------------------|-----------------|--------------|
| GET    | /billing/settings | Get settings    | billing:read |
| PUT    | /billing/settings | Update settings | billing:*    |

| Method | Path                  | Description          | Permission   |
|--------|-----------------------|----------------------|--------------|
| GET    | /billing/current      | Current period usage | billing:read |
| GET    | /billing/projections  | Cost projections     | billing:read |
| GET    | /billing/invoices     | List invoices        | billing:read |
| GET    | /billing/invoices/:id | Get invoice          | billing:read |
| POST   | /billing/generate     | Generate invoice     | billing:*    |

## Metering Routes

| Method | Path              | Description         | Permission   |
|--------|-------------------|---------------------|--------------|
| POST   | /metering/record  | Record usage event  | (internal)   |
| POST   | /metering/batch   | Batch record events | (internal)   |
| GET    | /metering/summary | Usage summary       | billing:read |
| GET    | /metering/rollups | Daily rollups       | billing:read |

---

## DEPLOYMENT VERIFICATION

```
“bash # 1. Create Invitation curl -X POST -H “Authorization: Bearer $ADMIN_TOKEN” -H “Content-Type: application/json” \ -d ‘{“email”: “new.admin@company.com”, “role”: “operator”, “message”: “Welcome!”, “expiresInHours”: 48}’ \ https://admin-api.thinktank.YOUR_DOMAIN.com/api/v2/admin/invitations
```

### 2. List Pending Approvals (for me to approve)

```
curl -H “Authorization: Bearer $ADMIN_TOKEN” \ “https://admin-api.thinktank.YOUR_DOMAIN.com/api/v2/approvals/pending”
```

### 3. Get Current Billing Period

```
curl -H “Authorization: Bearer $ADMIN_TOKEN” \ https://admin-api.thinktank.YOUR_DOMAIN.com/api/v2/billing/period
```

### 4. Get Usage Rollups

```
curl -H “Authorization: Bearer $ADMIN_TOKEN” \ “https://admin-api.thinktank.YOUR_DOMAIN.com/api/v2/metering/rollups?start=2024-12-01&endDate=2024-12-21”
```

### 5. Generate Invoice

```
curl -X POST -H “Authorization: Bearer $ADMIN_TOKEN” -H “Content-Type: application/json” \ -d ‘{“periodStart”: “2024-12-01”, “periodEnd”: “2024-12-31”, “sendToStripe”: true}’ \ https://admin-api.thinktank.YOUR_DOMAIN.com/api/v2/billing/generate “
```

---

## ESTIMATED COSTS BY TIER

| Component             | Tier 1   | Tier 2   | Tier 3   | Tier 4   | Tier 5   |
|-----------------------|----------|----------|----------|----------|----------|
| Lambda (Admin)        | \$5      | \$15     | \$50     | \$150    | \$500    |
| Lambda (Billing)      | \$5      | \$20     | \$75     | \$200    | \$600    |
| DynamoDB (Usage)      | \$10     | \$50     | \$200    | \$500    | \$1,500  |
| SES (Emails)          | \$1      | \$5      | \$20     | \$50     | \$150    |
| Stripe Fees           | Variable | Variable | Variable | Variable | Variable |
| <b>Prompt 5 Total</b> | ~\$21    | ~\$90    | ~\$345   | ~\$900   | ~\$2,750 |

*Note: Stripe fees are 2.9% + \$0.30 per transaction, not included in estimates.*

## NEXT PROMPTS

Continue with: - **Prompt 6:** Self-Hosted Models & Mid-Level Services Configuration - **Prompt 7:** External Providers & Database Schema/Migrations - **Prompt 8:** Admin Web Dashboard (Next.js) - **Prompt 9:** Assembly & Deployment Guide

*End of Prompt 5: Lambda Functions - Admin & Billing RADIANT v2.2.0 - December 2024*

END OF SECTION 5