

Contents

SECTION 22: PERSISTENT MEMORY (v3.6.0)	1
	1
22.1 Memory System Overview	1
22.2 Memory Database Schema	1
22.3 Memory Service	2
	5

SECTION 22: PERSISTENT MEMORY (v3.6.0)

22.1 Memory System Overview

Long-term memory storage using pgvector embeddings for semantic search.

22.2 Memory Database Schema

-- *migrations/031_persistent_memory.sql*

```
CREATE TABLE memory_stores (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    tenant_id UUID NOT NULL REFERENCES tenants(id),
    user_id UUID NOT NULL REFERENCES users(id),
    store_name VARCHAR(100) NOT NULL DEFAULT 'default',
    embedding_model VARCHAR(100) DEFAULT 'text-embedding-3-small',
    total_memories INTEGER DEFAULT 0,
    last_accessed TIMESTAMPTZ,
    created_at TIMESTAMPTZ NOT NULL DEFAULT CURRENT_TIMESTAMP,
    UNIQUE(tenant_id, user_id, store_name)
);

CREATE TABLE memories (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    store_id UUID NOT NULL REFERENCES memory_stores(id) ON DELETE CASCADE,
    content TEXT NOT NULL,
    embedding vector(1536),
    memory_type VARCHAR(50) DEFAULT 'fact',
    source VARCHAR(100),
    importance DECIMAL(3, 2) DEFAULT 0.5,
    access_count INTEGER DEFAULT 0,
    last_accessed TIMESTAMPTZ,
    expires_at TIMESTAMPTZ,
    metadata JSONB DEFAULT '{}',
    created_at TIMESTAMPTZ NOT NULL DEFAULT CURRENT_TIMESTAMP
);
```

```

CREATE TABLE memory_relationships (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    source_memory_id UUID NOT NULL REFERENCES memories(id) ON DELETE CASCADE,
    target_memory_id UUID NOT NULL REFERENCES memories(id) ON DELETE CASCADE,
    relationship_type VARCHAR(50) NOT NULL,
    strength DECIMAL(3, 2) DEFAULT 0.5,
    created_at TIMESTAMPTZ NOT NULL DEFAULT CURRENT_TIMESTAMP,
    UNIQUE(source_memory_id, target_memory_id, relationship_type)
);

CREATE INDEX idx_memories_embedding ON memories USING ivfflat (embedding vector_cosine_ops) WITH (faiss_index_type = IVFFLAT);
CREATE INDEX idx_memories_store ON memories(store_id);
CREATE INDEX idx_memory_relationships ON memory_relationships(source_memory_id);

ALTER TABLE memory_stores ENABLE ROW LEVEL SECURITY;
ALTER TABLE memories ENABLE ROW LEVEL SECURITY;
ALTER TABLE memory_relationships ENABLE ROW LEVEL SECURITY;

CREATE POLICY memory_stores_isolation ON memory_stores USING (tenant_id = current_setting('app.current_tenant'));
CREATE POLICY memories_isolation ON memories USING (
    store_id IN (SELECT id FROM memory_stores WHERE tenant_id = current_setting('app.current_tenant'))
);
CREATE POLICY memory_relationships_isolation ON memory_relationships USING (
    source_memory_id IN (SELECT m.id FROM memories m JOIN memory_stores ms ON m.store_id = ms.id)
);

```

22.3 Memory Service

```

// packages/core/src/services/memory-service.ts

import { Pool } from 'pg';
import { BedrockRuntimeClient, InvokeModelCommand } from '@aws-sdk/client-bedrock-runtime';

export class MemoryService {
    private pool: Pool;
    private bedrock: BedrockRuntimeClient;

    constructor(pool: Pool) {
        this.pool = pool;
        this.bedrock = new BedrockRuntimeClient({});
    }

    async getOrCreateStore(tenantId: string, userId: string, storeName: string = 'default'): Promise<any> {
        const result = await this.pool.query(`

            INSERT INTO memory_stores (tenant_id, user_id, store_name)
            VALUES ($1, $2, $3)

            ON CONFLICT (tenant_id, user_id, store_name) DO UPDATE SET last_accessed = NOW()
        `);
    }
}

```

```

        RETURNING id
    `, [tenantId, userId, storeName]);

    return result.rows[0].id;
}

async addMemory(
    storeId: string,
    content: string,
    options?: {
        type?: string;
        source?: string;
        importance?: number;
        metadata?: Record<string, any>;
    }
): Promise<string> {
    const embedding = await this.generateEmbedding(content);

    const result = await this.pool.query(`

        INSERT INTO memories (store_id, content, embedding, memory_type, source, importance)
        VALUES ($1, $2, $3::vector, $4, $5, $6, $7)
        RETURNING id
    `,
    [
        storeId,
        content,
        `${embedding.join(',')}`,
        options?.type || 'fact',
        options?.source,
        options?.importance || 0.5,
        JSON.stringify(options?.metadata || {})
    ]);
}

// Update store count
await this.pool.query(`

    UPDATE memory_stores SET total_memories = total_memories + 1 WHERE id = $1
``,
    [storeId]);

return result.rows[0].id;
}

async searchMemories(
    storeId: string,
    query: string,
    limit: number = 5,
    minSimilarity: number = 0.7
): Promise<any[]> {
    const embedding = await this.generateEmbedding(query);
}

```

```

const result = await this.pool.query(`

    SELECT
        id, content, memory_type, source, importance, metadata,
        1 - (embedding <=> $2::vector) as similarity
    FROM memories
    WHERE store_id = $1
    AND (expires_at IS NULL OR expires_at > NOW())
    AND 1 - (embedding <=> $2::vector) >= $4
    ORDER BY embedding <=> $2::vector
    LIMIT $3
`, [storeId, `${embedding.join(',')}`], limit, minSimilarity);

// Update access counts
const memoryIds = result.rows.map(r => r.id);
if (memoryIds.length > 0) {
    await this.pool.query(`

        UPDATE memories SET access_count = access_count + 1, last_accessed = NOW()
        WHERE id = ANY($1)
    `, [memoryIds]);
}

return result.rows;
}

async addRelationship(
    sourceId: string,
    targetId: string,
    relationshipType: string,
    strength: number = 0.5
): Promise<void> {
    await this.pool.query(`

        INSERT INTO memory_relationships (source_memory_id, target_memory_id, relationship_type)
        VALUES ($1, $2, $3, $4)
        ON CONFLICT (source_memory_id, target_memory_id, relationship_type)
        DO UPDATE SET strength = EXCLUDED.strength
    `, [sourceId, targetId, relationshipType, strength]);
}

async getRelatedMemories(memoryId: string, limit: number = 5): Promise<any[]> {
    const result = await this.pool.query(`

        SELECT m.*, mr.relationship_type, mr.strength
        FROM memory_relationships mr
        JOIN memories m ON mr.target_memory_id = m.id
        WHERE mr.source_memory_id = $1
        ORDER BY mr.strength DESC
        LIMIT $2
    `, [memoryId, limit]);
}

```

```
        return result.rows;
    }

private async generateEmbedding(text: string): Promise<number[]> {
    const response = await this.bedrock.send(new InvokeModelCommand({
        modelId: 'amazon.titan-embed-text-v1',
        body: JSON.stringify({ inputText: text }),
        contentType: 'application/json'
    }));

    const result = JSON.parse(new TextDecoder().decode(response.body));
    return result.embedding;
}
```