

Contents

SECTION 34: DATABASE-DRIVEN ORCHESTRATION ENGINE (v4.1.0)	1
	1
34.1 DATABASE SCHEMA	1
packages/database/migrations/034_orchestration_engine.sql	1
34.2 ALPHAFOOLD 2 SEED DATA	8
packages/database/migrations/034a_seed_alphafoold2.sql	8
34.3 ORCHESTRATION ENGINE SERVICE	10
packages/services/orchestration/OrchestrationEngine.ts	10
	15

SECTION 34: DATABASE-DRIVEN ORCHESTRATION ENGINE (v4.1.0)

CRITICAL: This section replaces ALL hardcoded model configurations with database-driven management. All model configs, workflows, and orchestration parameters are stored in PostgreSQL. Administrators can add/edit/delete models entirely through the Admin Dashboard UI.

34.1 DATABASE SCHEMA

packages/database/migrations/034_orchestration_engine.sql

```
-- =====
-- RADIANT v4.1.0 - DATABASE-DRIVEN ORCHESTRATION ENGINE
-- =====
-- This migration creates the foundation for dynamic model management.
-- ALL model configurations are stored in PostgreSQL - NO HARDCODING.
-- =====

-- AI Model Registry (Replaces hardcoded TypeScript configs)
CREATE TABLE IF NOT EXISTS ai_models (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    model_id VARCHAR(100) UNIQUE NOT NULL,
    name VARCHAR(100) NOT NULL,
    display_name VARCHAR(200) NOT NULL,
    description TEXT,
    category VARCHAR(50) NOT NULL,
    specialty VARCHAR(50),
    provider_type VARCHAR(20) NOT NULL DEFAULT 'self_hosted',
    deployment_config JSONB NOT NULL DEFAULT '{}',
```

```

parameters BIGINT DEFAULT 0,
accuracy VARCHAR(200),
benchmark TEXT,
capabilities TEXT[] DEFAULT '{}',
input_formats TEXT[] DEFAULT '{}',
output_formats TEXT[] DEFAULT '{}',
architecture JSONB DEFAULT '{}',
performance_metrics JSONB DEFAULT '{}',

thermal_config JSONB NOT NULL DEFAULT '{"defaultState":"OFF","scaleToZeroAfterMinutes":15,"w',
current_thermal_state VARCHAR(20) DEFAULT 'OFF',
last_thermal_change TIMESTAMPTZ,

pricing_config JSONB NOT NULL DEFAULT '{"hourlyRate":0,"perRequest":0,"markup":0.75}',

min_tier INTEGER DEFAULT 1,
requires_gpu BOOLEAN DEFAULT false,
gpu_memory_gb INTEGER DEFAULT 0,
enabled BOOLEAN DEFAULT true,
status VARCHAR(20) DEFAULT 'active',
version VARCHAR(50),
repository VARCHAR(500),
release_date DATE,

created_at TIMESTAMPTZ DEFAULT NOW(),
updated_at TIMESTAMPTZ DEFAULT NOW(),
created_by UUID REFERENCES administrators(id),
updated_by UUID REFERENCES administrators(id)
);

-- License Tracking
CREATE TABLE IF NOT EXISTS model_licenses (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    model_id UUID NOT NULL REFERENCES ai_models(id) ON DELETE CASCADE,
    license_type VARCHAR(50) NOT NULL,
    license_spdx VARCHAR(50) NOT NULL,
    license_url VARCHAR(500),
    commercial_use BOOLEAN DEFAULT true,
    commercial_notes TEXT,
    attribution_required BOOLEAN DEFAULT false,
    attribution_text TEXT,
    share_alike BOOLEAN DEFAULT false,
    compliance_status VARCHAR(20) DEFAULT 'compliant',
    last_compliance_review TIMESTAMPTZ,
    reviewed_by UUID REFERENCES administrators(id),
    compliance_notes TEXT,
    expires_at TIMESTAMPTZ,
    created_at TIMESTAMPTZ DEFAULT NOW(),

```

```

updated_at TIMESTAMPTZ DEFAULT NOW()
);

-- Model Dependencies
CREATE TABLE IF NOT EXISTS model_dependencies (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    model_id UUID NOT NULL REFERENCES ai_models(id) ON DELETE CASCADE,
    dependency_type VARCHAR(50) NOT NULL,
    dependency_name VARCHAR(200) NOT NULL,
    dependency_size_gb DECIMAL(10,2),
    dependency_license VARCHAR(50),
    required BOOLEAN DEFAULT true,
    notes TEXT,
    created_at TIMESTAMPTZ DEFAULT NOW()
);

-- Workflow Definitions
CREATE TABLE IF NOT EXISTS workflow_definitions (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    workflow_id VARCHAR(100) UNIQUE NOT NULL,
    name VARCHAR(200) NOT NULL,
    description TEXT,
    category VARCHAR(50) NOT NULL,
    version VARCHAR(20) NOT NULL DEFAULT '1.0.0',
    dag_definition JSONB NOT NULL,
    input_schema JSONB NOT NULL DEFAULT '{}',
    output_schema JSONB NOT NULL DEFAULT '{}',
    default_parameters JSONB NOT NULL DEFAULT '{}',
    timeout_seconds INTEGER DEFAULT 3600,
    max_retries INTEGER DEFAULT 3,
    min_tier INTEGER DEFAULT 1,
    enabled BOOLEAN DEFAULT true,
    requires_audit_trail BOOLEAN DEFAULT false,
    hipaa_compliant BOOLEAN DEFAULT false,
    created_at TIMESTAMPTZ DEFAULT NOW(),
    updated_at TIMESTAMPTZ DEFAULT NOW(),
    created_by UUID REFERENCES administrators(id)
);

-- Workflow Tasks
CREATE TABLE IF NOT EXISTS workflow_tasks (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    workflow_id UUID NOT NULL REFERENCES workflow_definitions(id) ON DELETE CASCADE,
    task_id VARCHAR(100) NOT NULL,
    name VARCHAR(200) NOT NULL,
    description TEXT,
    task_type VARCHAR(50) NOT NULL,
    model_id UUID REFERENCES ai_models(id),

```

```

service_id VARCHAR(100),
config JSONB NOT NULL DEFAULT '{}',
input_mapping JSONB DEFAULT '{}',
output_mapping JSONB DEFAULT '{}',
sequence_order INTEGER DEFAULT 0,
depends_on TEXT[] DEFAULT '{}',
condition_expression TEXT,
timeout_seconds INTEGER DEFAULT 300,
created_at TIMESTAMPTZ DEFAULT NOW(),
updated_at TIMESTAMPTZ DEFAULT NOW(),
UNIQUE(workflow_id, task_id)
);

-- Workflow Parameters
CREATE TABLE IF NOT EXISTS workflow_parameters (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    workflow_id UUID NOT NULL REFERENCES workflow_definitions(id) ON DELETE CASCADE,
    parameter_name VARCHAR(100) NOT NULL,
    display_name VARCHAR(200),
    description TEXT,
    data_type VARCHAR(50) NOT NULL,
    default_value JSONB,
    validation_rules JSONB DEFAULT '{}',
    ui_component VARCHAR(50) DEFAULT 'text',
    ui_config JSONB DEFAULT '{}',
    user_configurable BOOLEAN DEFAULT true,
    admin_only BOOLEAN DEFAULT false,
    sequence_order INTEGER DEFAULT 0,
    created_at TIMESTAMPTZ DEFAULT NOW(),
    UNIQUE(workflow_id, parameter_name)
);

-- Workflow Executions
CREATE TABLE IF NOT EXISTS workflow_executions (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    workflow_id UUID NOT NULL REFERENCES workflow_definitions(id),
    tenant_id UUID NOT NULL REFERENCES tenants(id),
    user_id UUID NOT NULL REFERENCES users(id),
    status VARCHAR(20) NOT NULL DEFAULT 'pending',
    input_parameters JSONB NOT NULL DEFAULT '{}',
    resolved_parameters JSONB DEFAULT '{}',
    output_data JSONB,
    error_message TEXT,
    error_details JSONB,
    started_at TIMESTAMPTZ,
    completed_at TIMESTAMPTZ,
    duration_ms INTEGER,
    estimated_cost_usd DECIMAL(10,4),
);

```

```

actual_cost_usd DECIMAL(10,4),
checkpoint_data JSONB,
priority INTEGER DEFAULT 5,
created_at TIMESTAMPTZ DEFAULT NOW(),
updated_at TIMESTAMPTZ DEFAULT NOW()
);

-- Task Executions
CREATE TABLE IF NOT EXISTS task_executions (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    workflow_execution_id UUID NOT NULL REFERENCES workflow_executions(id) ON DELETE CASCADE,
    task_id VARCHAR(100) NOT NULL,
    status VARCHAR(20) NOT NULL DEFAULT 'pending',
    attempt_number INTEGER DEFAULT 1,
    input_data JSONB,
    output_data JSONB,
    error_message TEXT,
    error_code VARCHAR(50),
    started_at TIMESTAMPTZ,
    completed_at TIMESTAMPTZ,
    duration_ms INTEGER,
    resource_usage JSONB DEFAULT '{}',
    cost_usd DECIMAL(10,4),
    model_id UUID REFERENCES ai_models(id),
    model_endpoint VARCHAR(500),
    created_at TIMESTAMPTZ DEFAULT NOW(),
    updated_at TIMESTAMPTZ DEFAULT NOW()
);

-- Orchestration Audit Log
CREATE TABLE IF NOT EXISTS orchestration_audit_log (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    entity_type VARCHAR(50) NOT NULL,
    entity_id UUID NOT NULL,
    action VARCHAR(50) NOT NULL,
    actor_id UUID REFERENCES administrators(id),
    actor_type VARCHAR(20) DEFAULT 'admin',
    previous_state JSONB,
    new_state JSONB,
    change_summary TEXT,
    ip_address INET,
    user_agent TEXT,
    created_at TIMESTAMPTZ DEFAULT NOW()
);

-- Indexes
CREATE INDEX IF NOT EXISTS idx_ai_models_category ON ai_models(category);
CREATE INDEX IF NOT EXISTS idx_ai_models_status ON ai_models(status, enabled);

```

```

CREATE INDEX IF NOT EXISTS idx_model_licenses_model ON model_licenses(model_id);
CREATE INDEX IF NOT EXISTS idx_model_licenses_compliance ON model_licenses(compliance_status);
CREATE INDEX IF NOT EXISTS idx_workflow_executions_tenant ON workflow_executions(tenant_id);
CREATE INDEX IF NOT EXISTS idx_workflow_executions_status ON workflow_executions(status);
CREATE INDEX IF NOT EXISTS idx_orchestration_audit_time ON orchestration_audit_log(created_at);

-- Function: Get full model config (replaces hardcoded lookups)
CREATE OR REPLACE FUNCTION get_model_config(p_model_id VARCHAR)
RETURNS JSONB AS $$

DECLARE
    result JSONB;
BEGIN
    SELECT jsonb_build_object(
        'id', m.model_id,
        'name', m.name,
        'displayName', m.display_name,
        'description', m.description,
        'category', m.category,
        'providerType', m.provider_type,
        'deployment', m.deployment_config,
        'thermal', m.thermal_config,
        'pricing', m.pricing_config,
        'parameters', m.parameters,
        'accuracy', m.accuracy,
        'minTier', m.min_tier,
        'requiresGpu', m.requires_gpu,
        'status', m.status,
        'version', m.version,
        'currentThermalState', m.current_thermal_state,
        'licenses', (
            SELECT COALESCE(jsonb_agg(jsonb_build_object(
                'id', l.id, 'type', l.license_type, 'spdx', l.license_spdx,
                'commercialUse', l.commercial_use, 'complianceStatus', l.compliance_status
            )), '[]'::jsonb)
            FROM model_licenses l WHERE l.model_id = m.id
        ),
        'dependencies', (
            SELECT COALESCE(jsonb_agg(jsonb_build_object(
                'name', d.dependency_name, 'type', d.dependency_type,
                'sizeGb', d.dependency_size_gb, 'required', d.required
            )), '[]'::jsonb)
            FROM model_dependencies d WHERE d.model_id = m.id
        )
    ) INTO result
    FROM ai_models m
    WHERE m.model_id = p_model_id AND m.enabled = true;

    RETURN result;

```

```

END;
$$ LANGUAGE plpgsql STABLE;

-- Function: Get workflow config
CREATE OR REPLACE FUNCTION get_workflow_config(p_workflow_id VARCHAR)
RETURNS JSONB AS $$

DECLARE
    result JSONB;
BEGIN
    SELECT jsonb_build_object(
        'id', w.workflow_id,
        'name', w.name,
        'description', w.description,
        'category', w.category,
        'version', w.version,
        'dag', w.dag_definition,
        'inputSchema', w.input_schema,
        'defaultParameters', w.default_parameters,
        'timeout', w.timeout_seconds,
        'minTier', w.min_tier,
        'tasks', (
            SELECT COALESCE(jsonb_agg(jsonb_build_object(
                'taskId', t.task_id, 'name', t.name, 'type', t.task_type,
                'modelId', (SELECT model_id FROM ai_models WHERE id = t.model_id),
                'config', t.config, 'dependsOn', t.depends_on, 'timeout', t.timeout_seconds
            ) ORDER BY t.sequence_order), '[]'::jsonb)
            FROM workflow_tasks t WHERE t.workflow_id = w.id
        ),
        'parameters', (
            SELECT COALESCE(jsonb_agg(jsonb_build_object(
                'name', p.parameter_name, 'displayName', p.display_name,
                'type', p.data_type, 'default', p.default_value,
                'validation', p.validation_rules, 'uiComponent', p.ui_component
            ) ORDER BY p.sequence_order), '[]'::jsonb)
            FROM workflow_parameters p WHERE p.workflow_id = w.id
        )
    ) INTO result
    FROM workflow_definitions w
    WHERE w.workflow_id = p_workflow_id AND w.enabled = true;

    RETURN result;
END;
$$ LANGUAGE plpgsql STABLE;

-- Audit Trigger
CREATE OR REPLACE FUNCTION log_model_changes()
RETURNS TRIGGER AS $$
BEGIN

```

```

IF TG_OP = 'INSERT' THEN
    INSERT INTO orchestration_audit_log (entity_type, entity_id, action, new_state, change_summary)
    VALUES ('model', NEW.id, 'created', to_jsonb(NEW), 'Model ' || NEW.display_name || ' created');
ELSIF TG_OP = 'UPDATE' THEN
    INSERT INTO orchestration_audit_log (entity_type, entity_id, action, previous_state, new_state, change_summary)
    VALUES ('model', NEW.id, 'updated', to_jsonb(OLD), to_jsonb(NEW), 'Model ' || NEW.display_name || ' updated');
ELSIF TG_OP = 'DELETE' THEN
    INSERT INTO orchestration_audit_log (entity_type, entity_id, action, previous_state, new_state, change_summary)
    VALUES ('model', OLD.id, 'deleted', to_jsonb(OLD), 'Model ' || OLD.display_name || ' deleted');
END IF;
RETURN COALESCE(NEW, OLD);
END;
$$ LANGUAGE plpgsql;

DROP TRIGGER IF EXISTS ai_models_audit_trigger ON ai_models;
CREATE TRIGGER ai_models_audit_trigger
AFTER INSERT OR UPDATE OR DELETE ON ai_models
FOR EACH ROW EXECUTE FUNCTION log_model_changes();

-- Row Level Security
ALTER TABLE ai_models ENABLE ROW LEVEL SECURITY;
ALTER TABLE workflow_executions ENABLE ROW LEVEL SECURITY;

CREATE POLICY ai_models_read ON ai_models FOR SELECT TO authenticated
    USING (enabled = true AND status IN ('active', 'beta'));

CREATE POLICY workflow_executions_owner ON workflow_executions FOR ALL TO authenticated
    USING (user_id = auth.uid() OR EXISTS (SELECT 1 FROM administrators WHERE user_id = auth.uid));

```

34.2 ALPHAFOLD 2 SEED DATA

packages/database/migrations/034a_seed_alphaFold2.sql

```

-- =====
-- ALPHAFOOLD 2 - Nobel Prize-Winning Protein Folding Model
-- =====

INSERT INTO ai_models (
    model_id, name, display_name, description, category, specialty, provider_type,
    deployment_config, parameters, accuracy, benchmark, capabilities,
    input_formats, output_formats, architecture, performance_metrics,
    thermal_config, pricing_config, min_tier, requires_gpu, gpu_memory_gb,
    status, version, repository, release_date
) VALUES (
    'alphaFold2',
    'alphaFold2',
    'AlphaFold 2',

```

```

'Nobel Prize-winning protein structure prediction with near-experimental accuracy. Won CASP14
'scientific_protein',
'protein_folding',
'self_hosted',
'{
  "image": "pytorch-inference:2.1-gpu-py310-cu121-ubuntu22.04",
  "instanceType": "ml.g5.12xlarge",
  "environment": {"MODEL_NAME": "alphafold2_ptm", "JAX_PLATFORM_NAME": "gpu"},
  "modelDataUrl": "s3://rariant-models/alphafold2/params_2022-12-06.tar",
  "resourceRequirements": {"minGpuMemoryGB": 16, "recommendedGpuMemoryGB": 96, "databaseStorage": 93000000,
}':::JSONB,
'GDT > 90 on CASP14, 92.4 median GDT_TS, 0.96Å median backbone RMSD',
'CASP14: Won 88/97 targets, far exceeding all other methods.',
ARRAY['protein_folding', 'structure_prediction', 'confidence_estimation', 'multimer_prediction'],
ARRAY['text/fasta', 'text/plain', 'application/json'],
ARRAY['application/pdb', 'application/mmcif', 'application/json'],
'{ "evoformerBlocks": 48, "structureBlocks": 8, "recyclingIterations": 4 }':::JSONB,
'{ "inferenceTime100Residues": 4.9, "inferenceTime500Residues": 29, "maxResiduesG5_12xlarge": 100000 },
{ "defaultState": "OFF", "scaleToZeroAfterMinutes": 10, "warmupTimeSeconds": 180, "minInstances": 1 },
{ "hourlyRate": 14.28, "perRequest": 2.00, "perResidueOver500": 0.002, "markup": 0.75 }':::JSONB,
4, true, 96, 'active', '2.3.2', 'https://github.com/google-deepmind/alphafold', '2023-04-05'
) ON CONFLICT (model_id) DO UPDATE SET updated_at = NOW();

-- Insert licenses
DO $$
DECLARE alphafold2_id UUID;
BEGIN
  SELECT id INTO alphafold2_id FROM ai_models WHERE model_id = 'alphafold2';

  INSERT INTO model_licenses (model_id, license_type, license_spdx, license_url, commercial_use)
  VALUES
    (alphafold2_id, 'source_code', 'Apache-2.0', 'https://github.com/google-deepmind/alphafold'),
    (alphafold2_id, 'model_weights', 'CC-BY-4.0', 'https://creativecommons.org/licenses/by/4.0/'),
    (alphafold2_id, 'database', 'CC-BY-SA-4.0', 'https://creativecommons.org/licenses/by-sa/4.0/');

  ON CONFLICT DO NOTHING;

  INSERT INTO model_dependencies (model_id, dependency_type, dependency_name, dependency_size_gb)
  VALUES
    (alphafold2_id, 'database', 'BFD', 1800, 'CC-BY-SA-4.0', true),
    (alphafold2_id, 'database', 'UniRef90', 103, 'CC-BY-4.0', true),
    (alphafold2_id, 'database', 'MGNify', 120, 'CC0-1.0', true),
    (alphafold2_id, 'database', 'PDB70', 56, 'CC0-1.0', true),
    (alphafold2_id, 'database', 'PDB mmCIF', 238, 'CC0-1.0', true);

  ON CONFLICT DO NOTHING;
END $$;

-- Insert protein folding workflow

```

```

INSERT INTO workflow_definitions (
    workflow_id, name, description, category, version, dag_definition, input_schema, output_schema,
    default_parameters, timeout_seconds, max_retries, min_tier, requires_audit_trail
) VALUES (
    'protein_folding_alphaFold2',
    'AlphaFold 2 Protein Folding Pipeline',
    'Complete protein structure prediction using AlphaFold 2 with MSA generation, template search and scoring',
    'scientific',
    '1.0.0',
    '{"nodes": [{"id": "validate_input", "type": "validation"}, {"id": "msa_generation", "type": "task", "children": [{"id": "sequence", "type": "object", "required": true}], "properties": {"sequence": {"type": "string", "minLength": 1}}, "children": [{"id": "structures", "type": "object", "required": true}, {"id": "confidence", "type": "object", "required": true}], "properties": {"structures": {"type": "array"}, "confidence": {"type": "object"}}, "recyclingIterations": 4, "relaxStructure": true, "useTemplates": true} } :: JSONB,
    7200, 2, 4, true
) ON CONFLICT (workflow_id) DO UPDATE SET updated_at = NOW();

```

34.3 ORCHESTRATION ENGINE SERVICE

packages/services/orchestration/OrchestrationEngine.ts

```

/**
 * RADIANT v4.1.0 - Database-Driven Orchestration Engine
 *
 * KEY PRINCIPLE: ZERO HARDCODING
 * All model configs stored in ai_models table - retrieved via get_model_config()
 */

import { Pool } from 'pg';
import { EventEmitter } from 'events';

export interface ModelConfig {
    id: string;
    name: string;
    displayName: string;
    description?: string;
    category: string;
    providerType: 'self_hosted' | 'external';
    deployment: Record<string, any>;
    thermal: { defaultState: string; scaleToZeroAfterMinutes: number; warmupTimeSeconds: number };
    pricing: { hourlyRate: number; perRequest: number; markup: number };
    parameters: number;
    minTier: number;
    requiresGpu: boolean;
    status: string;
    currentThermalState?: string;
    licenses: Array<{ id: string; type: string; spdx: string; commercialUse: boolean; compliance: string }>;
    dependencies: Array<{ name: string; type: string; sizeGb?: number; required: boolean }>;
}

```

```
}
```

```
export class OrchestrationEngine extends EventEmitter {
    private pool: Pool;
    private modelCache: Map<string, { config: ModelConfig; timestamp: number }> = new Map();
    private cacheExpiry = 5 * 60 * 1000; // 5 minutes

    constructor(pool: Pool) {
        super();
        this.pool = pool;
    }

    /**
     * Get model config from database - REPLACES hardcoded TypeScript configs
     */
    async getModelConfig(modelId: string): Promise<ModelConfig | null> {
        const cached = this.modelCache.get(modelId);
        if (cached && Date.now() - cached.timestamp < this.cacheExpiry) {
            return cached.config;
        }

        const result = await this.pool.query('SELECT get_model_config($1) as config', [modelId]);
        if (result.rows[0]?.config) {
            const config = result.rows[0].config as ModelConfig;
            this.modelCache.set(modelId, { config, timestamp: Date.now() });
            return config;
        }
        return null;
    }

    /**
     * List models with filters - Used by Admin Dashboard
     */
    async listModels(filters?: { category?: string; status?: string; minTier?: number; providerType?: string }): Promise<ModelConfig[]> {
        let query = `SELECT get_model_config(model_id) as config FROM ai_models WHERE enabled = true`;
        const params: any[] = [];
        let idx = 1;

        if (filters?.category) { query += ` AND category = ${idx++}`; params.push(filters.category); }
        if (filters?.status) { query += ` AND status = ${idx++}`; params.push(filters.status); }
        if (filters?.minTier) { query += ` AND min_tier <= ${idx++}`; params.push(filters.minTier); }
        if (filters?.providerType) { query += ` AND provider_type = ${idx++}`; params.push(filters.providerType); }

        query += ` ORDER BY category, display_name`;
        const result = await this.pool.query(query, params);
        return result.rows.map(r => r.config).filter(Boolean);
    }
}
```

```

/**
 * Create model via Admin Dashboard - NO code deployment needed
 */
async createModel(input: any, adminId: string): Promise<string> {
  const client = await this.pool.connect();
  try {
    await client.query('BEGIN');

    const result = await client.query(`

      INSERT INTO ai_models (
        model_id, name, display_name, description, category, specialty,
        provider_type, deployment_config, parameters, accuracy, benchmark,
        capabilities, input_formats, output_formats, thermal_config, pricing_config,
        min_tier, requires_gpu, gpu_memory_gb, status, version, repository, created_by
      ) VALUES ($1,$2,$3,$4,$5,$6,$7,$8,$9,$10,$11,$12,$13,$14,$15,$16,$17,$18,$19,$20,$21,$22)
      RETURNING id
    `, [
      input.modelId, input.name, input.displayName, input.description,
      input.category, input.specialty, input.providerType || 'self_hosted',
      JSON.stringify(input.deployment || {}), input.parameters || 0,
      input.accuracy, input.benchmark, input.capabilities || [],
      input.inputFormats || [], input.outputFormats || [],
      JSON.stringify(input.thermal || {}), JSON.stringify(input.pricing || {}),
      input.minTier || 1, input.requiresGpu || false, input.gpuMemoryGb || 0,
      input.status || 'active', input.version, input.repository, adminId
    ]);
  }

  const modelUuid = result.rows[0].id;

  // Insert licenses
  for (const license of (input.licenses || [])) {
    await client.query(
      `INSERT INTO model_licenses (model_id, license_type, license_spdx, license_url, commercial_use)
      [modelUuid, license.type, license.spdx, license.url, license.commercialUse ?? true]
    `);
  }

  // Insert dependencies
  for (const dep of (input.dependencies || [])) {
    await client.query(
      `INSERT INTO model_dependencies (model_id, dependency_type, dependency_name, dependency_size_gb, license, required)
      [modelUuid, dep.type, dep.name, dep.sizeGb, dep.license, dep.required ?? true]
    `);
  }

  await client.query('COMMIT');
  this.modelCache.delete(input.modelId);
  this.emit('model:created', { modelId: input.modelId, adminId });
}

```

```

        return modelUuid;
    } catch (error) {
        await client.query('ROLLBACK');
        throw error;
    } finally {
        client.release();
    }
}

/**
 * Update model via Admin Dashboard
 */
async updateModel(modelId: string, updates: any, adminId: string): Promise<void> {
    const sets: string[] = [];
    const params: any[] = [];
    let idx = 1;

    const fields: Record<string, string> = {
        displayName: 'display_name', description: 'description', category: 'category',
        status: 'status', minTier: 'min_tier', version: 'version'
    };

    for (const [key, col] of Object.entries(fields)) {
        if (updates[key] !== undefined) { sets.push(`$ ${col} = ${idx++}`); params.push(updates[key]) }

        if (updates.deployment) { sets.push(`deployment_config = ${idx++}`); params.push(JSON.stringify(updates.deployment)) }
        if (updates.thermal) { sets.push(`thermal_config = ${idx++}`); params.push(JSON.stringify(updates.thermal)) }
        if (updates.pricing) { sets.push(`pricing_config = ${idx++}`); params.push(JSON.stringify(updates.pricing)) }

        if (sets.length === 0) return;
    }

    sets.push(`updated_at = NOW()`, `updated_by = ${idx++}`);
    params.push(adminId, modelId);

    await this.pool.query(`UPDATE ai_models SET ${sets.join(', ')} WHERE model_id = ${idx}`, [modelId, adminId]);
    this.modelCache.delete(modelId);
    this.emit('model:updated', { modelId, adminId });
}

/**
 * Delete model (soft delete)
 */
async deleteModel(modelId: string, adminId: string): Promise<void> {
    await this.pool.query(
        `UPDATE ai_models SET enabled = false, status = 'disabled', updated_at = NOW(), updated_by = ${adminId} WHERE model_id = ${modelId}`,
        [modelId, adminId]
    );
}

```

```

    this.modelCache.delete(modelId);
    this.emit('model:deleted', { modelId, adminId });
}

/**
 * Get workflow config from database
 */
async getWorkflowConfig(workflowId: string): Promise<any> {
    const result = await this.pool.query('SELECT get_workflow_config($1) as config', [workflowId]);
    return result.rows[0].config || null;
}

/**
 * Execute workflow with parameters
 */
async executeWorkflow(workflowId: string, tenantId: string, userId: string, parameters: Record<string, any>): Promise<any> {
    const workflow = await this.getWorkflowConfig(workflowId);
    if (!workflow) throw new Error(`Workflow not found: ${workflowId}`);
    const resolvedParams = { ...workflow.defaultParameters, ...parameters };

    const result = await this.pool.query(`

        INSERT INTO workflow_executions (workflow_id, tenant_id, user_id, status, input_parameters)
        VALUES ((SELECT id FROM workflow_definitions WHERE workflow_id = $1), $2, $3, 'running'),
        RETURNING id
    `, [workflowId, tenantId, userId, JSON.stringify(parameters), JSON.stringify(resolvedParams)]);

    const executionId = result.rows[0].id;
    this.emit('workflow:started', { executionId, workflowId, tenantId, userId });
    return executionId;
}

/**
 * Get license summary for compliance dashboard
 */
async getLicenseSummary(): Promise<{ totalModels: number; commercialOk: number; reviewNeeded: number }> {
    const result = await this.pool.query(`

        SELECT
            COUNT(DISTINCT m.id) as total_models,
            COUNT(DISTINCT CASE WHEN NOT EXISTS (SELECT 1 FROM model_licenses l2 WHERE l2.model_id = m.id) THEN m.id END) as commercial_ok,
            COUNT(DISTINCT CASE WHEN l.compliance_status = 'review_needed' THEN m.id END) as review_needed,
            COUNT(DISTINCT CASE WHEN l.compliance_status = 'non_compliant' THEN m.id END) as non_compliant,
            COUNT(DISTINCT CASE WHEN l.expires_at IS NOT NULL AND l.expires_at < NOW() + INTERVAL '1 month' THEN m.id END) as soon_to_expire
        FROM ai_models m LEFT JOIN model_licenses l ON m.id = l.model_id WHERE m.enabled = true
    `);
    return {
        totalModels: parseInt(result.rows[0].total_models),
        commercialOk: parseInt(result.rows[0].commercial_ok),
        reviewNeeded: parseInt(result.rows[0].review_needed),
        nonCompliant: parseInt(result.rows[0].non_compliant),
        soonToExpire: parseInt(result.rows[0].soon_to_expire)
    };
}

```

```
    reviewNeeded: parseInt(result.rows[0].review_needed),
    nonCompliant: parseInt(result.rows[0].non_compliant),
    expiringIn30Days: parseInt(result.rows[0].expiring_soon),
  };
}

refreshCache(): void {
  this.modelCache.clear();
}
}

let instance: OrchestrationEngine | null = null;
export function getOrchestrationEngine(pool: Pool): OrchestrationEngine {
  if (!instance) instance = new OrchestrationEngine(pool);
  return instance;
}
```
