

Contents

SECTION 37: FEEDBACK LEARNING SYSTEM & NEURAL ENGINE LOOP (v4.3.0)	1
	1
37.1 Feedback System Overview	1
Architecture Flow	1
Learning Scope Hierarchy	3
37.2 Feedback Database Schema	3
37.3 Shared Types for Feedback System	15
37.4 API Endpoints Summary	21
Feedback Endpoints	21
Neural Engine Endpoints	22
Voice Processing Endpoints	22
Service Layer (Client Apps)	22
Admin Endpoints	22
37.5 Implicit Signal Sentiment Mapping	23
37.6 Learning Weighting Formula	23
Summary	24
Section 37: Feedback Learning System (v4.3.0)	24
	24

SECTION 37: FEEDBACK LEARNING SYSTEM & NEURAL ENGINE LOOP (v4.3.0)

Dependencies: Sections 0-36, especially 11 (Brain) and 13 (Neural Engine) **Creates:** Complete feedback loop from user signals → Neural Engine → Brain decisions

37.1 Feedback System Overview

The Feedback Learning System creates a closed-loop where user feedback continuously improves AI routing decisions. The system captures both explicit feedback (thumbs up/down) and implicit signals (regenerate, copy, abandon), ties each to a complete execution manifest, and feeds everything into the Neural Engine which advises the Brain.

Architecture Flow

FEEDBACK LEARNING LOOP

User Request

RADIANT BRAIN

- Consults Neural Engine for recommendations
- Considers user/tenant/global learning
- Applies confidence thresholds
- Selects orchestration → services → models

EXECUTION MANIFEST

Records: output_id, models[], versions[], orchestration_id,
services[], thermal_states{}, provider_health{},
brain_reasoning, latency_ms, cost, timestamp

AI RESPONSE

Delivered to user with output_id for feedback reference

EXPLICIT	IMPLICIT	VOICE
+ cat	regenerate	any lang
+ text	copy/share	transcribe
	abandon	translate
	switch model	

NEURAL ENGINE

- Aggregates feedback by model/orchestration/service
- Updates model scores (individual → tenant → global)
- Applies confidence thresholds and decay
- Generates routing recommendations
- Tracks A/B experiment outcomes

BRAIN INTELLIGENCE

Real-time: Neural recommendations inform next Brain decision

Batch: Nightly aggregation updates routing weights

Learning Scope Hierarchy

TIERED LEARNING SCOPE

Priority 1: USER SCOPE (Most Personalized)

- Alice's feedback improves Alice's experience
- Fast adaptation to individual preferences
- Requires minimum ~20 feedback samples for confidence

Priority 2: TENANT SCOPE (Organization-Wide)

- Company X's employees collectively train Company X's Brain
- Isolated from Company Y (privacy boundary)
- Captures organizational preferences (e.g., "we prefer Claude")

Priority 3: GLOBAL SCOPE (Platform-Wide, Anonymized)

- Aggregate patterns: "Claude wins 73% for legal tasks"
- Cold start defaults for new users/tenants
- Configurable opt-in/opt-out per tenant

37.2 Feedback Database Schema

```
-- migrations/037_feedback_learning_system.sql

-- =====
-- EXECUTION MANIFESTS: Full provenance for every AI output
-- =====

CREATE TABLE execution_manifests (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    output_id VARCHAR(100) UNIQUE NOT NULL, -- Reference ID for feedback
```

```

-- Context
tenant_id UUID NOT NULL REFERENCES tenants(id),
user_id UUID NOT NULL REFERENCES users(id),
conversation_id UUID REFERENCES thinktank_conversations(id),
message_id UUID REFERENCES thinktank_messages(id),

-- What was requested
request_type VARCHAR(50) NOT NULL, -- 'chat', 'completion', 'orchestration', 'service'
task_type VARCHAR(50), -- 'code', 'creative', 'analysis', 'medical', etc.
domain_mode VARCHAR(50), -- Think Tank domain mode if applicable
input_prompt_hash VARCHAR(64), -- SHA-256 of input for deduplication
input_tokens INTEGER,
input_language VARCHAR(10), -- Detected input language (ISO 639-1)

-- What was used (THE MANIFEST)
models_used TEXT[] NOT NULL,
model_versions JSONB DEFAULT '{}', -- {"claude-sonnet-4": "20241022", ...}
orchestration_id UUID REFERENCES workflow_definitions(id),
orchestration_name VARCHAR(100),
services_used TEXT[], -- ['perception', 'medical', ...]
thermal_states_at_execution JSONB DEFAULT '{}', -- {"whisper-large-v3": "HOT", ...}
provider_health_at_execution JSONB DEFAULT '{}', -- {"anthropic": {"latency_ms": 150, "er...

-- Brain's decision
brain_reasoning TEXT, -- Why Brain chose this path
brain_confidence DECIMAL(3, 2), -- 0.00-1.00
was_user_override BOOLEAN DEFAULT false, -- User manually selected model

-- Outcome metrics
output_tokens INTEGER,
total_latency_ms INTEGER,
time_to_first_token_ms INTEGER,
total_cost DECIMAL(10, 6),
was_streamed BOOLEAN DEFAULT false,

-- For multi-step orchestrations
step_count INTEGER DEFAULT 1,
step_details JSONB DEFAULT '[]', -- [{model, latency, tokens, cost}, ...]

created_at TIMESTAMPTZ NOT NULL DEFAULT CURRENT_TIMESTAMP,

CONSTRAINT valid_confidence CHECK (brain_confidence >= 0 AND brain_confidence <= 1)
);

CREATE INDEX idx_exec_manifest_output ON execution_manifests(output_id);
CREATE INDEX idx_exec_manifest_tenant_user ON execution_manifests(tenant_id, user_id);
CREATE INDEX idx_exec_manifest_conversation ON execution_manifests(conversation_id);

```

```

CREATE INDEX idx_exec_manifest_models ON execution_manifests USING GIN(models_used);
CREATE INDEX idx_exec_manifest_task ON execution_manifests(task_type);
CREATE INDEX idx_exec_manifest_created ON execution_manifests(created_at DESC);

-- =====
-- EXPLICIT FEEDBACK: Thumbs up/down with optional categories
-- =====

CREATE TYPE feedback_rating AS ENUM ('positive', 'negative', 'neutral');
CREATE TYPE feedback_category AS ENUM (
    'accuracy',      -- Factually correct
    'relevance',     -- Answered the question
    'tone',          -- Appropriate style
    'format',         -- Good structure/formatting
    'speed',          -- Fast enough
    'safety',         -- Appropriate content
    'creativity',    -- Novel/interesting
    'completeness',   -- Fully addressed request
    'other'
);

CREATE TABLE feedback_explicit (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

    -- Link to execution
    output_id VARCHAR(100) NOT NULL REFERENCES execution_manifests(output_id),
    tenant_id UUID NOT NULL REFERENCES tenants(id),
    user_id UUID NOT NULL REFERENCES users(id),

    -- The feedback
    rating feedback_rating NOT NULL,
    categories feedback_category[] DEFAULT '{}', -- What specifically was good/bad
    comment_text TEXT, -- Optional text feedback
    comment_language VARCHAR(10), -- Detected language of comment

    -- Metadata
    feedback_source VARCHAR(50) DEFAULT 'thinktank', -- 'thinktank', 'api', 'admin'
    user_agent TEXT,
    client_timestamp TIMESTAMPTZ, -- When user clicked

    created_at TIMESTAMPTZ NOT NULL DEFAULT CURRENT_TIMESTAMP,

    -- One feedback per output per user
    UNIQUE(output_id, user_id)
);

CREATE INDEX idx_feedback_explicit_output ON feedback_explicit(output_id);
CREATE INDEX idx_feedback_explicit_tenant ON feedback_explicit(tenant_id);

```

```

CREATE INDEX idx_feedback_explicit_rating ON feedback_explicit(rating);
CREATE INDEX idx_feedback_explicit_created ON feedback_explicit(created_at DESC);

-- =====
-- IMPLICIT FEEDBACK: Behavioral signals (higher volume, weighted less)
-- =====

CREATE TYPE implicit_signal_type AS ENUM (
    'regenerate',           -- User clicked regenerate (negative)
    'copy_response',        -- User copied response (positive)
    'share_response',       -- User shared response (very positive)
    'export_response',      -- User exported (positive)
    'continue_conversation', -- User sent follow-up (neutral-positive)
    'abandon_conversation', -- User left without follow-up (weak negative)
    'abandon_mid_response', -- User stopped streaming (negative)
    'manual_model_switch',  -- User changed model (negative for current)
    'edit_and_resend',      -- User edited their prompt (neutral)
    'response_time_short',  -- Quick reply = engaged (positive)
    'response_time_long',   -- Slow reply = confused/distracted (neutral)
    'scroll_to_end',         -- Read full response (positive)
    'scroll_bounce',         -- Scrolled away quickly (negative)
    'favorite_added',        -- Added to favorites (very positive)
    'report_content'        -- Reported for review (very negative)
);
;

CREATE TABLE feedback_implicit (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

    -- Link to execution
    output_id VARCHAR(100) NOT NULL REFERENCES execution_manifests(output_id),
    tenant_id UUID NOT NULL REFERENCES tenants(id),
    user_id UUID NOT NULL REFERENCES users(id),

    -- The signal
    signal_type implicit_signal_type NOT NULL,
    signal_value JSONB DEFAULT '{}', -- Additional context (e.g., time_to_next_message_ms)

    -- Computed sentiment score (-1.0 to +1.0)
    sentiment_score DECIMAL(3, 2) NOT NULL,

    created_at TIMESTAMPTZ NOT NULL DEFAULT CURRENT_TIMESTAMP
);

CREATE INDEX idx_feedback_implicit_output ON feedback_implicit(output_id);
CREATE INDEX idx_feedback_implicit_tenant ON feedback_implicit(tenant_id);
CREATE INDEX idx_feedback_implicit_signal ON feedback_implicit(signal_type);
CREATE INDEX idx_feedback_implicit_created ON feedback_implicit(created_at DESC);

```

```

-- =====
-- VOICE FEEDBACK: Multi-language voice input with transcription
-- =====

CREATE TABLE feedback_voice (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

    -- Link to execution
    output_id VARCHAR(100) NOT NULL REFERENCES execution_manifests(output_id),
    tenant_id UUID NOT NULL REFERENCES tenants(id),
    user_id UUID NOT NULL REFERENCES users(id),

    -- Audio storage
    audio_s3_key VARCHAR(500) NOT NULL,
    audio_duration_seconds DECIMAL(8, 2),
    audio_format VARCHAR(20), -- 'webm', 'mp3', 'wav', etc.

    -- Transcription
    transcription_text TEXT,
    original_language VARCHAR(10), -- Detected language (ISO 639-1)
    translated_text TEXT, -- English translation if not English
    transcription_confidence DECIMAL(3, 2),
    transcription_model VARCHAR(50), -- 'whisper-large-v3', etc.

    -- Sentiment analysis of voice feedback
    sentiment_score DECIMAL(3, 2), -- -1.0 to +1.0
    inferred_rating feedback_rating,
    inferred_categories feedback_category[],

    -- Processing status
    processing_status VARCHAR(20) DEFAULT 'pending', -- 'pending', 'processing', 'completed',
    processing_error TEXT,
    processed_at TIMESTAMPTZ,

    created_at TIMESTAMPTZ NOT NULL DEFAULT CURRENT_TIMESTAMP
);

CREATE INDEX idx_feedback_voice_output ON feedback_voice(output_id);
CREATE INDEX idx_feedback_voice_status ON feedback_voice(processing_status);

-- =====
-- NEURAL MODEL SCORES: Learned effectiveness per model/task/scope
-- =====

CREATE TYPE learning_scope AS ENUM ('user', 'tenant', 'global');

CREATE TABLE neural_model_scores (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

```

```

-- Scope (null = global)
scope learning_scope NOT NULL,
tenant_id UUID REFERENCES tenants(id), -- null for global
user_id UUID REFERENCES users(id), -- null for tenant/global

-- What we're scoring
model_id VARCHAR(100) NOT NULL,
task_type VARCHAR(50), -- null = overall score for model
domain_mode VARCHAR(50), -- null = all domains

-- Aggregated scores (0.0 to 1.0)
effectiveness_score DECIMAL(4, 3) NOT NULL DEFAULT 0.500,
accuracy_score DECIMAL(4, 3),
relevance_score DECIMAL(4, 3),
speed_score DECIMAL(4, 3),

-- Statistics
positive_count INTEGER DEFAULT 0,
negative_count INTEGER DEFAULT 0,
neutral_count INTEGER DEFAULT 0,
implicit_positive_count INTEGER DEFAULT 0,
implicit_negative_count INTEGER DEFAULT 0,
total_feedback_count INTEGER DEFAULT 0,

-- Confidence in this score (based on sample size)
confidence DECIMAL(3, 2) DEFAULT 0.00,

-- Model version tracking
last_model_version VARCHAR(50),
score_decay_applied_at TIMESTAMPTZ, -- When we last decayed old scores

created_at TIMESTAMPTZ NOT NULL DEFAULT CURRENT_TIMESTAMP,
updated_at TIMESTAMPTZ NOT NULL DEFAULT CURRENT_TIMESTAMP,

-- Unique per scope/model/task combination
UNIQUE NULLS NOT DISTINCT (scope, tenant_id, user_id, model_id, task_type, domain_mode)
);

CREATE INDEX idx_neural_scores_model ON neural_model_scores(model_id);
CREATE INDEX idx_neural_scores_scope ON neural_model_scores(scope);
CREATE INDEX idx_neural_scores_tenant ON neural_model_scores(tenant_id);
CREATE INDEX idx_neural_scores_effectiveness ON neural_model_scores(effectiveness_score DESC);

-- =====
-- NEURAL ROUTING RECOMMENDATIONS: Brain reads these for decisions
-- =====

```

```

CREATE TABLE neural_routing_recommendations (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

    -- Scope
    scope learning_scope NOT NULL,
    tenant_id UUID REFERENCES tenants(id),
    user_id UUID REFERENCES users(id),

    -- Recommendation context
    task_type VARCHAR(50) NOT NULL,
    domain_mode VARCHAR(50),
    input_characteristics JSONB DEFAULT '{}', -- {requires_vision, requires_audio, token_estimation}

    -- The recommendation
    recommended_model VARCHAR(100) NOT NULL,
    recommended_orchestration_id UUID REFERENCES workflow_definitions(id),
    recommended_services TEXT[],

    -- Alternatives (ordered by score)
    alternative_models TEXT[],

    -- Confidence
    recommendation_confidence DECIMAL(3, 2) NOT NULL,
    sample_size INTEGER, -- How much data this is based on

    -- Reasoning (for debugging/transparency)
    reasoning TEXT,

    -- Validity
    valid_from TIMESTAMPTZ NOT NULL DEFAULT CURRENT_TIMESTAMP,
    valid_until TIMESTAMPTZ, -- null = no expiry
    is_active BOOLEAN DEFAULT true,

    created_at TIMESTAMPTZ NOT NULL DEFAULT CURRENT_TIMESTAMP,
    updated_at TIMESTAMPTZ NOT NULL DEFAULT CURRENT_TIMESTAMP
);

CREATE INDEX idx_neural_rec_active ON neural_routing_recommendations(is_active, scope);
CREATE INDEX idx_neural_rec_task ON neural_routing_recommendations(task_type, domain_mode);
CREATE INDEX idx_neural_rec_tenant ON neural_routing_recommendations(tenant_id);

-- =====
-- USER TRUST SCORES: Anti-gaming feedback weighting
-- =====

CREATE TABLE user_trust_scores (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    tenant_id UUID NOT NULL REFERENCES tenants(id),

```

```

user_id UUID NOT NULL REFERENCES users(id),


-- Trust factors
trust_score DECIMAL(3, 2) NOT NULL DEFAULT 0.50, -- 0.00-1.00
account_age_days INTEGER,
total_feedback_count INTEGER DEFAULT 0,
feedback_diversity_score DECIMAL(3, 2), -- How varied their feedback is
feedback_alignment_score DECIMAL(3, 2), -- How aligned with population


-- Flags
is_outlier BOOLEAN DEFAULT false,
outlier_reason TEXT,
is_rate_limited BOOLEAN DEFAULT false,
rate_limit_until TIMESTAMPTZ,


-- Last update
last_feedback_at TIMESTAMPTZ,
last_trust_calculation_at TIMESTAMPTZ,


created_at TIMESTAMPTZ NOT NULL DEFAULT CURRENT_TIMESTAMP,
updated_at TIMESTAMPTZ NOT NULL DEFAULT CURRENT_TIMESTAMP,


UNIQUE(tenant_id, user_id)
);

CREATE INDEX idx_trust_scores_tenant ON user_trust_scores(tenant_id);
CREATE INDEX idx_trust_scores_outlier ON user_trust_scores(is_outlier);

--- =====
-- A/B TESTING: Measure if routing changes improve outcomes
--- =====


CREATE TYPE experiment_status AS ENUM ('draft', 'running', 'paused', 'completed', 'cancelled')


CREATE TABLE ab_experiments (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

    -- Experiment definition
    name VARCHAR(200) NOT NULL,
    description TEXT,
    hypothesis TEXT,

    -- Scope
    tenant_id UUID REFERENCES tenants(id), -- null = all tenants

    -- What we're testing
    experiment_type VARCHAR(50) NOT NULL, -- 'model_routing', 'orchestration', 'service'
    control_config JSONB NOT NULL, -- The current/default behavior

```

```

treatment_config JSONB NOT NULL, -- The new behavior being tested

-- Traffic allocation
traffic_percentage DECIMAL(5, 2) DEFAULT 10.00, -- % of users in treatment

-- Status
status experiment_status NOT NULL DEFAULT 'draft',
started_at TIMESTAMPTZ,
ended_at TIMESTAMPTZ,

-- Results
control_sample_size INTEGER DEFAULT 0,
treatment_sample_size INTEGER DEFAULT 0,
control_positive_rate DECIMAL(5, 4),
treatment_positive_rate DECIMAL(5, 4),
statistical_significance DECIMAL(5, 4), -- p-value
effect_size DECIMAL(5, 4), -- Cohen's d
winner VARCHAR(20), -- 'control', 'treatment', 'inconclusive'

created_by UUID REFERENCES administrators(id),
created_at TIMESTAMPTZ NOT NULL DEFAULT CURRENT_TIMESTAMP,
updated_at TIMESTAMPTZ NOT NULL DEFAULT CURRENT_TIMESTAMP
);

CREATE INDEX idx_ab_experiments_status ON ab_experiments(status);
CREATE INDEX idx_ab_experiments_tenant ON ab_experiments(tenant_id);

CREATE TABLE ab_experiment_assignments (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    experiment_id UUID NOT NULL REFERENCES ab_experiments(id) ON DELETE CASCADE,
    tenant_id UUID NOT NULL REFERENCES tenants(id),
    user_id UUID NOT NULL REFERENCES users(id),

    -- Assignment
    variant VARCHAR(20) NOT NULL, -- 'control' or 'treatment'
    assigned_at TIMESTAMPTZ NOT NULL DEFAULT CURRENT_TIMESTAMP,

    UNIQUE(experiment_id, user_id)
);

CREATE INDEX idx_ab_assignments_experiment ON ab_experiment_assignments(experiment_id);
CREATE INDEX idx_ab_assignments_user ON ab_experiment_assignments(user_id);

=====

-- ROW LEVEL SECURITY
=====

ALTER TABLE execution_manifests ENABLE ROW LEVEL SECURITY;

```

```

ALTER TABLE feedback_explicit ENABLE ROW LEVEL SECURITY;
ALTER TABLE feedback_implicit ENABLE ROW LEVEL SECURITY;
ALTER TABLE feedback_voice ENABLE ROW LEVEL SECURITY;
ALTER TABLE neural_model_scores ENABLE ROW LEVEL SECURITY;
ALTER TABLE neural_routing_recommendations ENABLE ROW LEVEL SECURITY;
ALTER TABLE user_trust_scores ENABLE ROW LEVEL SECURITY;
ALTER TABLE ab_experiments ENABLE ROW LEVEL SECURITY;
ALTER TABLE ab_experiment_assignments ENABLE ROW LEVEL SECURITY;

CREATE POLICY exec_manifest_isolation ON execution_manifests
    USING (tenant_id = current_setting('app.current_tenant_id')::UUID);
CREATE POLICY feedback_explicit_isolation ON feedback_explicit
    USING (tenant_id = current_setting('app.current_tenant_id')::UUID);
CREATE POLICY feedback_implicit_isolation ON feedback_implicit
    USING (tenant_id = current_setting('app.current_tenant_id')::UUID);
CREATE POLICY feedback_voice_isolation ON feedback_voice
    USING (tenant_id = current_setting('app.current_tenant_id')::UUID);
CREATE POLICY neural_scores_isolation ON neural_model_scores
    USING (tenant_id IS NULL OR tenant_id = current_setting('app.current_tenant_id')::UUID);
CREATE POLICY neural_rec_isolation ON neural_routing_recommendations
    USING (tenant_id IS NULL OR tenant_id = current_setting('app.current_tenant_id')::UUID);
CREATE POLICY trust_scores_isolation ON user_trust_scores
    USING (tenant_id = current_setting('app.current_tenant_id')::UUID);
CREATE POLICY ab_experiments_isolation ON ab_experiments
    USING (tenant_id IS NULL OR tenant_id = current_setting('app.current_tenant_id')::UUID);
CREATE POLICY ab_assignments_isolation ON ab_experiment_assignments
    USING (tenant_id = current_setting('app.current_tenant_id')::UUID);

-- =====
-- HELPER FUNCTIONS
-- =====

-- Get aggregated feedback score for an output
CREATE OR REPLACE FUNCTION get_feedback_score(p_output_id VARCHAR)
RETURNS JSONB AS $$
DECLARE
    explicit_score DECIMAL;
    implicit_score DECIMAL;
    combined_score DECIMAL;
    result JSONB;
BEGIN
    -- Get explicit feedback
    SELECT
        CASE rating
            WHEN 'positive' THEN 1.0
            WHEN 'negative' THEN -1.0
            ELSE 0.0
        END INTO explicit_score

```

```

    FROM feedback_explicit
    WHERE output_id = p_output_id
    LIMIT 1;

    -- Get average implicit feedback
    SELECT AVG(sentiment_score) INTO implicit_score
    FROM feedback_implicit
    WHERE output_id = p_output_id;

    -- Combine (explicit weighted 3x)
    combined_score := COALESCE(
        COALESCE(explicit_score, 0) * 3 + COALESCE(implicit_score, 0)) /
        CASE
            WHEN explicit_score IS NOT NULL AND implicit_score IS NOT NULL THEN 4
            WHEN explicit_score IS NOT NULL THEN 3
            WHEN implicit_score IS NOT NULL THEN 1
            ELSE 1
        END,
        0
    );

    result := jsonb_build_object(
        'explicit_score', explicit_score,
        'implicit_score', implicit_score,
        'combined_score', combined_score,
        'has_explicit', explicit_score IS NOT NULL,
        'has_implicit', implicit_score IS NOT NULL
    );
}

    RETURN result;
END;
$$ LANGUAGE plpgsql;

-- Get best model recommendation for context
CREATE OR REPLACE FUNCTION get_neural_recommendation(
    p_tenant_id UUID,
    p_user_id UUID,
    p_task_type VARCHAR,
    p_domain_mode VARCHAR DEFAULT NULL
)
RETURNS TABLE (
    model_id VARCHAR,
    confidence DECIMAL,
    scope learning_scope,
    reasoning TEXT
) AS $$

BEGIN
    -- Try user-specific first, then tenant, then global

```

```

RETURN QUERY
SELECT
    r.recommended_model,
    r.recommendation_confidence,
    r.scope,
    r.reasoning
FROM neural_routing_recommendations r
WHERE r.is_active = true
AND (r.valid_until IS NULL OR r.valid_until > NOW())
AND r.task_type = p_task_type
AND (r.domain_mode IS NULL OR r.domain_mode = p_domain_mode)
AND (
    (r.scope = 'user' AND r.tenant_id = p_tenant_id AND r.user_id = p_user_id) OR
    (r.scope = 'tenant' AND r.tenant_id = p_tenant_id AND r.user_id IS NULL) OR
    (r.scope = 'global' AND r.tenant_id IS NULL AND r.user_id IS NULL)
)
ORDER BY
    CASE r.scope
        WHEN 'user' THEN 1
        WHEN 'tenant' THEN 2
        WHEN 'global' THEN 3
    END,
    r.recommendation_confidence DESC
LIMIT 1;
END;
$$ LANGUAGE plpgsql;

-- Calculate implicit signal sentiment
CREATE OR REPLACE FUNCTION get_implicit_sentiment(p_signal_type implicit_signal_type)
RETURNS DECIMAL AS $$

BEGIN
    RETURN CASE p_signal_type
        WHEN 'copy_response' THEN 0.7
        WHEN 'share_response' THEN 0.9
        WHEN 'export_response' THEN 0.6
        WHEN 'favorite_added' THEN 0.9
        WHEN 'continue_conversation' THEN 0.3
        WHEN 'scroll_to_end' THEN 0.2
        WHEN 'response_time_short' THEN 0.2
        WHEN 'regenerate' THEN -0.8
        WHEN 'manual_model_switch' THEN -0.6
        WHEN 'abandon_conversation' THEN -0.3
        WHEN 'abandon_mid_response' THEN -0.7
        WHEN 'scroll_bounce' THEN -0.4
        WHEN 'report_content' THEN -1.0
        WHEN 'edit_and_resend' THEN 0.0
        WHEN 'response_time_long' THEN 0.0
        ELSE 0.0
    END;
END;
$$ LANGUAGE plpgsql;

```

```

    END;
END;
$$ LANGUAGE plpgsql IMMUTABLE;

```

37.3 Shared Types for Feedback System

```

// packages/shared/src/types/feedback.types.ts

/**
 * RADIANT v4.3.0 - Feedback System Types
 */

// =====
// EXECUTION MANIFEST
// =====

export interface ExecutionManifest {
  id: string;
  outputId: string; // Unique reference for feedback

  // Context
  tenantId: string;
  userId: string;
  conversationId?: string;
  messageId?: string;

  // Request
  requestType: 'chat' | 'completion' | 'orchestration' | 'service';
  taskType?: TaskType;
  domainMode?: DomainMode;
  inputPromptHash?: string;
  inputTokens?: number;
  inputLanguage?: string;

  // THE MANIFEST - What was used
  modelsUsed: string[];
  modelVersions: Record<string, string>;
  orchestrationId?: string;
  orchestrationName?: string;
  servicesUsed: string[];
  thermalStatesAtExecution: Record<string, ThermalState>;
  providerHealthAtExecution: Record<string, ProviderHealthSnapshot>;

  // Brain decision
  brainReasoning?: string;
  brainConfidence: number;
}

```

```

wasUserOverride: boolean;

// Outcome
outputTokens?: number;
totalLatencyMs: number;
timeToFirstTokenMs?: number;
totalCost: number;
wasStreamed: boolean;

// Multi-step
stepCount: number;
stepDetails: ExecutionStep[];

createdAt: Date;
}

export interface ExecutionStep {
  stepIndex: number;
  modelId: string;
  latencyMs: number;
  inputTokens: number;
  outputTokens: number;
  cost: number;
}

export interface ProviderHealthSnapshot {
  latencyMs: number;
  errorRate: number;
  status: 'healthy' | 'degraded' | 'unhealthy';
}

export type TaskType =
  | 'chat'
  | 'code'
  | 'analysis'
  | 'creative'
  | 'vision'
  | 'audio'
  | 'medical'
  | 'legal'
  | 'research'
  | 'translation';

export type DomainMode =
  | 'general'
  | 'medical'
  | 'legal'
  | 'code'

```

```

| 'creative'
| 'research'
| 'business';

// =====
// EXPLICIT FEEDBACK
// =====

export type FeedbackRating = 'positive' | 'negative' | 'neutral';

export type FeedbackCategory =
    | 'accuracy'
    | 'relevance'
    | 'tone'
    | 'format'
    | 'speed'
    | 'safety'
    | 'creativity'
    | 'completeness'
    | 'other';

export interface ExplicitFeedback {
    id: string;
    outputId: string;
    tenantId: string;
    userId: string;

    rating: FeedbackRating;
    categories: FeedbackCategory[];
    commentText?: string;
    commentLanguage?: string;

    feedbackSource: 'thinktank' | 'api' | 'admin';
    createdAt: Date;
}

export interface FeedbackSubmission {
    outputId: string;
    rating: FeedbackRating;
    categories?: FeedbackCategory[];
    commentText?: string;
}

// =====
// IMPLICIT FEEDBACK
// =====

export type ImplicitSignalType =

```

```

| 'regenerate'
| 'copy_response'
| 'share_response'
| 'export_response'
| 'continue_conversation'
| 'abandon_conversation'
| 'abandon_mid_response'
| 'manual_model_switch'
| 'edit_and_resend'
| 'response_time_short'
| 'response_time_long'
| 'scroll_to_end'
| 'scroll_bounce'
| 'favorite_added'
| 'report_content';

export interface ImplicitFeedback {
  id: string;
  outputId: string;
  tenantId: string;
  userId: string;

  signalType: ImplicitSignalType;
  signalValue: Record<string, unknown>;
  sentimentScore: number; // -1.0 to +1.0

  createdAt: Date;
}

export interface ImplicitSignalSubmission {
  outputId: string;
  signalType: ImplicitSignalType;
  signalValue?: Record<string, unknown>;
}

// =====
// VOICE FEEDBACK
// =====

export interface VoiceFeedback {
  id: string;
  outputId: string;
  tenantId: string;
  userId: string;

  audioS3Key: string;
  audioDurationSeconds: number;
  audioFormat: string;
}

```

```

transcriptionText?: string;
originalLanguage?: string;
translatedText?: string;
transcriptionConfidence?: number;
transcriptionModel?: string;

sentimentScore?: number;
inferredRating?: FeedbackRating;
inferredCategories?: FeedbackCategory[];

processingStatus: 'pending' | 'processing' | 'completed' | 'failed';
processedAt?: Date;

createdAt: Date;
}

// =====
// NEURAL LEARNING
// =====

export type LearningScope = 'user' | 'tenant' | 'global';

export interface NeuralModelScore {
  id: string;
  scope: LearningScope;
  tenantId?: string;
  userId?: string;

  modelId: string;
  taskType?: TaskType;
  domainMode?: DomainMode;

  effectivenessScore: number; // 0.0-1.0
  accuracyScore?: number;
  relevanceScore?: number;
  speedScore?: number;

  positiveCount: number;
  negativeCount: number;
  neutralCount: number;
  implicitPositiveCount: number;
  implicitNegativeCount: number;
  totalFeedbackCount: number;

  confidence: number; // 0.0-1.0

  updatedAt: Date;
}

```

```

}

export interface NeuralRecommendation {
  id: string;
  scope: LearningScope;
  tenantId?: string;
  userId?: string;

  taskType: TaskType;
  domainMode?: DomainMode;
  inputCharacteristics: Record<string, unknown>;

  recommendedModel: string;
  recommendedOrchestrationId?: string;
  recommendedServices: string[];
  alternativeModels: string[];

  recommendationConfidence: number;
  sampleSize: number;
  reasoning: string;

  isActive: boolean;
  validFrom: Date;
  validUntil?: Date;
}

// =====
// TRUST & ANTI-GAMING
// =====

export interface UserTrustScore {
  id: string;
  tenantId: string;
  userId: string;

  trustScore: number; // 0.0-1.0
  accountAgeDays: number;
  totalFeedbackCount: number;
  feedbackDiversityScore: number;
  feedbackAlignmentScore: number;

  isOutlier: boolean;
  outlierReason?: string;
  isRateLimited: boolean;
  rateLimitUntil?: Date;

  updatedAt: Date;
}

```

```

// =====
// A/B TESTING
// =====

export type ExperimentStatus = 'draft' | 'running' | 'paused' | 'completed' | 'cancelled';

export interface ABExperiment {
    id: string;
    name: string;
    description?: string;
    hypothesis?: string;

    tenantId?: string;
    experimentType: 'model_routing' | 'orchestration' | 'service';
    controlConfig: Record<string, unknown>;
    treatmentConfig: Record<string, unknown>;

    trafficPercentage: number;
    status: ExperimentStatus;

    controlSampleSize: number;
    treatmentSampleSize: number;
    controlPositiveRate?: number;
    treatmentPositiveRate?: number;
    statisticalSignificance?: number;
    effectSize?: number;
    winner?: 'control' | 'treatment' | 'inconclusive';

    startedAt?: Date;
    endedAt?: Date;
}


```

// Import existing types
`import { ThermalState } from './ai.types';`

Update the shared types index:

// packages/shared/src/types/index.ts

// Add to existing exports
`export * from './feedback.types';`

37.4 API Endpoints Summary

Feedback Endpoints

Method	Endpoint	Description
POST	/api/v2/feedback/explicit	Submit thumbs up/down with optional categories
POST	/api/v2/feedback/implicit	Record implicit signal (regenerate, copy, etc.)
POST	/api/v2/feedback/voice	Upload voice feedback (multipart)
GET	/api/v2/feedback/stats/{output}	Get aggregated feedback for output

Neural Engine Endpoints

Method	Endpoint	Description
GET	/api/v2/neural/recommendations	Get Neural Engine recommendation
GET	/api/v2/neural/scores	Get model scores for context
POST	/api/v2/neural/learn	Trigger learning for output (internal)

Voice Processing Endpoints

Method	Endpoint	Description
POST	/api/v2/voice/transcribe	Transcribe audio to text (any language)
POST	/api/v2/voice/translate	Translate text to English

Service Layer (Client Apps)

Method	Endpoint	Description
GET	/api/v2/service/feedback/config	Get feedback widget configuration
POST	/api/v2/service/feedback/implants/batch	Batch implicit signals
GET	/api/v2/service/feedback/conversations/{fid}/summary	Get summary

Admin Endpoints

Method	Endpoint	Description
GET	/api/v2/admin/feedback/stats	Feedback analytics dashboard
GET	/api/v2/admin/experiments	List A/B experiments
POST	/api/v2/admin/experiments	Create A/B experiment
PUT	/api/v2/admin/experiments/{id}	Update experiment status
GET	/api/v2/admin/trust-scores	View user trust scores

37.5 Implicit Signal Sentiment Mapping

Signal Type	Sentiment Score	Interpretation
favorite_added	+0.9	Very positive - user values this response
share_response	+0.9	Very positive - worth sharing
copy_response	+0.7	Positive - useful enough to copy
export_response	+0.6	Positive - keeping for later
continue_conversation	+0.3	Neutral-positive - engaged
scroll_to_end	+0.2	Weak positive - read full response
response_time_short	+0.2	Weak positive - quick engagement
edit_and_resend	0.0	Neutral - refining question
response_time_long	0.0	Neutral - thinking/distracted
abandon_conversation	-0.3	Weak negative - may be done or frustrated
scroll_bounce	-0.4	Negative - didn't read response
manual_model_switch	-0.6	Negative - current model wasn't working
abandon_mid_response	-0.7	Negative - stopped streaming early
regenerate	-0.8	Negative - response wasn't good
report_content	-1.0	Very negative - safety/quality issue

37.6 Learning Weighting Formula

Combined feedback score for model scoring:

```
weighted_score = (
    (explicit_score * 3.0 * trust_weight) +
    (implicit_score * 1.0) +
    (voice_score * 2.0 * trust_weight)
) / total_weight
```

Where: - **explicit_score**: -1.0 to +1.0 from thumbs rating - **implicit_score**: Average of implicit signal sentiments - **voice_score**: Sentiment analysis of transcribed voice feedback - **trust_weight**: 0.1 to 1.0 based on user trust score

Model effectiveness update:

```
new_score = (current_score * current_count + weighted_score) / (current_count + 1)
confidence = min(total_count / min_sample_size, 1.0)
```

Summary

RADIANT v4.3.0 (PROMPT-17) adds **Feedback Learning System & Neural Engine Loop**:

Section 37: Feedback Learning System (v4.3.0)

1. **Database Schema** (037_feedback_learning_system.sql)
 - **execution_manifests** - Full provenance for every AI output
 - **feedback_explicit** - Thumbs up/down with categories
 - **feedback_implicit** - Behavioral signals (regenerate, copy, abandon, etc.)
 - **feedback_voice** - Multi-language voice feedback with transcription
 - **neural_model_scores** - Learned effectiveness per model/task/scope
 - **neural_routing_recommendations** - Brain advice from Neural Engine
 - **user_trust_scores** - Anti-gaming trust levels
 - **ab_experiments** - A/B testing experiments
 2. **Learning Scopes**
 - User-level: Personal preferences, fastest adaptation
 - Tenant-level: Organization-wide patterns, privacy isolated
 - Global-level: Platform defaults, cold start handling
 3. **Signal Types**
 - Explicit: Thumbs up/down + 8 feedback categories + text comments
 - Implicit: 15 behavioral signals (regenerate, copy, share, abandon, etc.)
 - Voice: Multi-language voice feedback with Whisper transcription
 4. **Neural Engine Integration**
 - Brain consults Neural Engine before every routing decision
 - Neural scores weighted at 35% in model selection
 - Real-time + nightly batch learning
 5. **Anti-Gaming Protection**
 - User trust scores (0.0-1.0)
 - Rate limiting (50 feedback/hour)
 - Outlier detection (deviation from population)
 - Account age weighting
 6. **A/B Testing Framework**
 - Create experiments comparing routing strategies
 - Automatic user assignment to control/treatment
 - Statistical significance tracking
-