

# Contents

<b>Think Tank - Administrator Guide</b>	<b>11</b>
Overview . . . . .	11
Related Authentication Documentation . . . . .	11
Table of Contents . . . . .	11
1. Think Tank Admin Features . . . . .	14
Available Sections . . . . .	14
2. User Rules System . . . . .	14
2.1 Rule Types . . . . .	14
2.2 Memory Categories . . . . .	15
2.3 Preset Rules . . . . .	15
2.4 Admin Configuration . . . . .	15
3. Delight System . . . . .	15
3.1 Features . . . . .	15
3.2 Admin Controls . . . . .	16
3.3 Message Types . . . . .	16
4. Brain Plan Viewer . . . . .	16
4.1 What Users See . . . . .	16
4.2 Admin Configuration . . . . .	16
5. Pre-Prompt Learning . . . . .	16
5.1 How It Works . . . . .	17
5.2 Admin Features . . . . .	17
6. Domain Taxonomy . . . . .	17
6.1 Hierarchy . . . . .	17
6.2 Admin Features . . . . .	17
7. Rejection Notifications . . . . .	17
7.1 User Experience . . . . .	17
7.2 Suggested Actions . . . . .	18
8. Canvas & Artifacts . . . . .	18
8.1 Artifact Types . . . . .	18
8.2 Admin Configuration . . . . .	18
9. Collaboration Features . . . . .	18
9.1 Core Features . . . . .	18
9.2 Enhanced Collaboration (v4.18.0+) . . . . .	18
9.3 Attachment Storage . . . . .	21
9.4 Admin Configuration . . . . .	21
9.5 Database Tables . . . . .	21
9.6 UI Components . . . . .	22
10. Shadow Testing . . . . .	22
10.1 Test Modes . . . . .	22
10.2 Creating a Shadow Test . . . . .	22
10.3 Test Results . . . . .	22
10.4 Auto-Promotion Settings . . . . .	23
10.5 Manual Review . . . . .	23
11. Routing Cache . . . . .	23
11.1 How It Works . . . . .	23
11.2 Optimistic Execution . . . . .	23

11.3 Cache Statistics . . . . .	23
11.4 Cache Configuration . . . . .	23
12. Delight System Toggle . . . . .	24
12.1 User Control . . . . .	24
12.2 Default Behavior . . . . .	24
12.3 Admin Configuration . . . . .	24
13. Intelligence Aggregator . . . . .	24
13.1 User-Facing Benefits . . . . .	24
13.2 Success Memory in Think Tank . . . . .	25
13.3 MoA Synthesis Mode . . . . .	25
13.4 Code Verification in Coding Mode . . . . .	25
13.5 Configuration . . . . .	25
14. AGI Ideas Service . . . . .	25
14.1 Typeahead Suggestions . . . . .	25
14.2 Result Ideas . . . . .	26
14.3 Learning from Usage . . . . .	26
14.4 API Endpoints . . . . .	26
14.5 Configuration . . . . .	26
14.6 Pattern Matching . . . . .	27
14.7 Persistent Learning . . . . .	27
14.8 Database Tables . . . . .	28
14.9 Key Files . . . . .	29
14.10 Troubleshooting . . . . .	29
15. Feedback System . . . . .	29
15.1 Rating Types . . . . .	29
15.2 Star Rating Labels . . . . .	29
15.3 Category Ratings (Optional) . . . . .	30
15.4 Comments . . . . .	30
15.5 Integration with Learning . . . . .	30
15.6 Configuration . . . . .	31
16. Cognitive Architecture . . . . .	31
16.1 Tree of Thoughts (Extended Thinking) . . . . .	31
16.2 GraphRAG (Knowledge Connections) . . . . .	31
16.3 Deep Research (Background Jobs) . . . . .	32
16.4 Generative UI (Interactive Results) . . . . .	32
16.5 Dynamic LoRA (Domain Expertise) . . . . .	33
16.6 Configuration . . . . .	33
17. Consciousness Service . . . . .	33
17.1 Continuous Existence (Heartbeat) . . . . .	33
17.2 Initialization on Startup . . . . .	34
17.3 User-Facing Features . . . . .	34
17.4 Consciousness Indicators . . . . .	34
17.5 Emergence Events . . . . .	34
17.6 Testing Tab . . . . .	34
17.7 Additional Consciousness Features . . . . .	35
18. App Factory . . . . .	35
18.1 What It Does . . . . .	35
18.2 Supported App Types . . . . .	35

18.3 Calculator Templates . . . . .	35
18.4 View Toggle . . . . .	36
18.5 User Preferences . . . . .	36
18.6 Admin Configuration . . . . .	36
18.7 How It Works . . . . .	36
18.8 Database Tables . . . . .	37
19. Multi-Page Web App Generator . . . . .	37
19.1 Supported App Types . . . . .	37
19.2 How It Works . . . . .	38
19.3 Page Types . . . . .	38
19.4 Section Types . . . . .	39
19.5 Navigation Types . . . . .	39
19.6 Pre-built Templates . . . . .	39
19.7 Admin Configuration . . . . .	40
19.8 Database Tables . . . . .	40
20. UI Feedback & Learning System . . . . .	40
20.1 User Feedback . . . . .	40
19.2 “Improve Before Your Eyes” . . . . .	41
19.3 AGI Learning . . . . .	41
19.4 Vision Analysis . . . . .	41
19.5 Admin Configuration . . . . .	41
19.6 Database Tables . . . . .	42
19.7 Analytics Dashboard . . . . .	42
20. Media Capabilities . . . . .	42
20.1 Supported Media Types . . . . .	42
20.2 Image Generation . . . . .	43
20.3 Audio Processing . . . . .	43
20.4 Vision/Image Understanding . . . . .	43
20.5 3D Generation . . . . .	43
20.6 Media Limits . . . . .	43
20.7 Database Tables . . . . .	43
21. Result Derivation History . . . . .	44
21.1 What’s Captured . . . . .	44
21.2 API Endpoints . . . . .	44
21.3 Timeline Visualization . . . . .	45
21.4 Model Usage Details . . . . .	45
21.5 Quality Dimensions . . . . .	45
21.6 Analytics Dashboard . . . . .	45
21.7 Database Tables . . . . .	45
22. User Persistent Context . . . . .	45
22.1 How It Works . . . . .	46
22.2 Context Types . . . . .	46
22.3 User API Endpoints . . . . .	46
22.4 User Preferences . . . . .	47
22.5 AGI Brain Planner Integration . . . . .	47
22.6 Library Assist Integration . . . . .	47
22.7 Context Injection Format . . . . .	47
22.7 Database Tables . . . . .	48

22.8 Admin Configuration . . . . .	48
23. Predictive Coding & Evolution . . . . .	48
23.1 Active Inference (Predictive Coding) . . . . .	48
23.2 Prediction Outcomes . . . . .	48
23.3 Surprise Magnitude . . . . .	49
23.4 Learning Candidates . . . . .	49
23.5 LoRA Evolution Pipeline . . . . .	49
23.6 Evolution State Tracking . . . . .	50
23.7 Database Tables . . . . .	50
23.8 Admin Configuration . . . . .	50
24. Zero-Cost Ego System . . . . .	50
24.1 Architecture Overview . . . . .	50
24.2 Cost Comparison . . . . .	51
24.3 Key Components . . . . .	51
24.4 Admin API Endpoints . . . . .	52
24.5 Admin Dashboard Features . . . . .	52
24.6 How It Works . . . . .	53
24.7 Database Tables . . . . .	53
24.8 Integration with AGI Brain Planner . . . . .	53
25. Conscious Orchestrator (Architecture Inversion) . . . . .	53
25.1 Overview . . . . .	53
25.2 Processing Phases . . . . .	53
25.3 Decision Types . . . . .	54
25.4 Usage . . . . .	54
25.5 Enhanced Affect Bindings . . . . .	54
25.6 Database Table . . . . .	54
26. Bipolar Rating System (Negative Ratings) . . . . .	55
26.1 Overview . . . . .	55
26.2 Key Metrics . . . . .	55
26.3 Quick Ratings (UI) . . . . .	55
26.4 Rating Dimensions . . . . .	56
26.5 API Endpoints . . . . .	56
26.6 User Calibration . . . . .	56
26.7 Learning Integration . . . . .	56
26.8 Database Tables . . . . .	56
27. Consciousness Engine Administration . . . . .	57
27.1 Dashboard Overview . . . . .	57
27.2 Budget Controls . . . . .	57
27.3 MCP Tools (23 Total) . . . . .	57
27.4 Admin API Endpoints . . . . .	57
27.5 Cost Tracking . . . . .	58
27.6 Library Registry . . . . .	58
27.7 Database Tables . . . . .	58
25. Formal Reasoning Libraries . . . . .	59
25.1 Library Overview . . . . .	59
25.2 Dashboard Features . . . . .	59
25.3 API Endpoints . . . . .	60
25.4 Consciousness Integration . . . . .	60

25.5 LLM-Modulo Pattern . . . . .	61
25.6 Database Tables . . . . .	62
25.7 Budget Management . . . . .	62
25.8 Thread Safety Notes . . . . .	62
25.9 Production Infrastructure . . . . .	62
26. Ethics-Free Reasoning . . . . .	63
26.1 Architecture . . . . .	63
26.2 Configuration . . . . .	64
26.3 API Endpoints . . . . .	65
26.4 Training Feedback . . . . .	65
26.5 Training Pipeline . . . . .	66
26.6 Usage . . . . .	66
26.7 Database Tables . . . . .	66
26.8 Benefits . . . . .	67
27. Intelligent File Conversion . . . . .	67
27.1 Core Concept . . . . .	67
27.2 Supported File Formats . . . . .	67
27.3 Provider Capabilities . . . . .	68
27.4 Conversion Strategies . . . . .	68
27.5 API Endpoints . . . . .	68
27.6 User Experience . . . . .	70
27.7 Admin Configuration . . . . .	70
27.8 Database Tables . . . . .	71
27.9 Multi-Model File Preparation . . . . .	71
27.10 Domain-Specific File Formats . . . . .	72
27.11 Reinforcement Learning . . . . .	72
27.12 Monitoring . . . . .	73
27.13 Related Documentation . . . . .	73
28. User Memories & Persistent Learning . . . . .	73
28.1 Learning Influence Hierarchy . . . . .	73
28.2 What Think Tank Learns . . . . .	74
28.3 Persistence Guarantee . . . . .	74
28.4 Integration with AGI Brain . . . . .	74
28.5 Admin Configuration . . . . .	74
28.6 Related Documentation . . . . .	74
29. Artifact Engine (GenUI Pipeline) . . . . .	74
29.1 Executive Summary . . . . .	75
29.2 System Architecture . . . . .	75
29.3 Core Concepts . . . . .	77
29.4 Administrative Control Panel . . . . .	78
29.5 Safety Governance (Genesis Cato CBFs) . . . . .	79
29.6 Dependency Allowlist Management . . . . .	81
29.7 Code Pattern Library . . . . .	82
29.8 Reflexion Loop (Self-Correction) . . . . .	84
29.9 Escalation Workflow Management . . . . .	84
29.10 Audit Trail & Compliance . . . . .	86
29.11 Metrics & Monitoring . . . . .	87
29.12 Tenant Configuration . . . . .	89

29.13 Troubleshooting Guide . . . . .	89
29.14 API Reference . . . . .	91
29.15 Real-Time Generation Logs . . . . .	92
29.16 Artifact Viewer Component . . . . .	93
29.17 Database Schema . . . . .	93
29.18 Security Considerations . . . . .	93
29.19 Implementation Files . . . . .	93
30. Consciousness Operating System (COS) . . . . .	94
30.1 Overview . . . . .	94
30.2 Architecture . . . . .	94
30.3 Ghost Vectors . . . . .	95
30.4 SOFAI Routing . . . . .	95
30.5 Flash Facts . . . . .	96
30.6 Dreaming System . . . . .	96
30.7 Human Oversight . . . . .	96
30.8 Privacy Airlock . . . . .	97
30.9 Configuration . . . . .	97
30.10 Database Schema . . . . .	97
30.11 Implementation Files . . . . .	98
31. Why Think Tank Beats Standalone AI (The System Advantage) . . . . .	98
31.1 The Executive Summary . . . . .	98
31.2 The Consultant vs Engineer Analogy . . . . .	99
31.3 What Users Experience . . . . .	99
31.4 The Three Pillars of Think Tank's Advantage . . . . .	100
31.5 Quantitative Comparison . . . . .	100
31.6 When Think Tank Automatically Escalates . . . . .	101
31.7 The Bottom Line . . . . .	101
32. Swarm Orchestration & Flyte Operations . . . . .	101
32.1 System Architecture: The “Deep Swarm” . . . . .	101
32.2 Operational Troubleshooting . . . . .	103
32.3 Compliance & Security . . . . .	104
32.4 Related Sections . . . . .	106
33. Cognitive Platform Enhancements . . . . .	106
33.1 Strategic Vision: Beyond Task Execution . . . . .	106
33.2 The Grimoire (Procedural Memory & Self-Correction) . . . . .	107
33.3 Time-Travel Debugging (Visual Forking) . . . . .	109
33.4 The Economic Governor (Model Arbitrage) . . . . .	111
33.5 Sentinel Agents (Event-Driven Autonomy) . . . . .	113
33.6 The Council of Rivals (Adversarial Consensus) . . . . .	115
33.7 Implementation Roadmap . . . . .	118
33.8 Database Schema . . . . .	119
33.9 API Reference . . . . .	122
33.10 Configuration . . . . .	122
33.11 Troubleshooting . . . . .	123
33.12 Related Sections . . . . .	124
34. Orchestration Workflow Methods Reference . . . . .	124
34.1 Method Categories Overview . . . . .	124
34.2 Complete Methods Reference . . . . .	125

34.3 System vs User Methods . . . . .	128
34.4 User Workflow Templates . . . . .	129
34.4 Cato Neural Decision Engine . . . . .	130
34.5 Method Parameters Reference . . . . .	130
34.6 User Workflow Template Parameter Overrides . . . . .	134
34.7 Database Tables . . . . .	135
34.7 Implementation Files . . . . .	135
35. Polymorphic UI (PROMPT-41) . . . . .	136
35.1 Overview . . . . .	136
35.2 The Gearbox (Elastic Compute) . . . . .	136
35.3 The Three Views . . . . .	136
35.4 View Types . . . . .	136
35.5 Configuration . . . . .	137
35.6 Implementation Files . . . . .	137
35.7 Database Tables . . . . .	137
35.8 API Endpoints . . . . .	137
36. Think Tank Policy Framework: Strategic Intelligence . . . . .	138
36.1 The Cato Institute Policy Foundation . . . . .	138
36.2 The \$10 Trillion Cybercrime Economy . . . . .	138
36.3 Memory Safety and the 70% Problem . . . . .	139
36.4 Regulatory Stance Configuration . . . . .	140
36.5 Database Tables . . . . .	141
36.6 Implementation Files . . . . .	141
37. Agentic Orchestration: SSF, CAEP, and Identity Remediation . . . . .	142
37.1 The Agentic AI Paradigm . . . . .	142
37.2 Shared Signals Framework (SSF) Integration . . . . .	142
37.3 Continuous Access Evaluation Profile (CAEP) . . . . .	144
37.4 Autonomous Identity Remediation . . . . .	145
37.5 The Radiant Ghost in Think Tank . . . . .	147
37.6 Database Tables . . . . .	148
37.7 API Endpoints . . . . .	149
37.8 Implementation Files . . . . .	149
38. Advanced Features (v4.18.0) . . . . .	149
38.1 Flash Facts (Knowledge Sparks) . . . . .	149
38.2 Grimoire (Spell Book) . . . . .	150
38.3 Economic Governor (Fuel Gauge) . . . . .	151
38.4 Sentinel Agents (Watchtower Dashboard) . . . . .	151
38.5 Time-Travel Debugging (Timeline Scrubber) . . . . .	152
38.6 Council of Rivals (Debate Arena) . . . . .	152
38.7 Security Signals (Security Shield) . . . . .	153
38.8 Policy Framework (Stance Compass) . . . . .	154
38.9 Database Migration . . . . .	155
39. Liquid Interface (Generative UI) . . . . .	155
39.1 Overview . . . . .	155
39.2 Architecture . . . . .	155
39.3 Component Registry . . . . .	156
39.4 Ghost State (Two-Way Binding) . . . . .	157
39.5 Intent Detection . . . . .	158

39.6 Eject to App . . . . .	159
39.7 Configuration . . . . .	159
39.8 API Endpoints . . . . .	160
39.9 Database Tables . . . . .	160
39.10 Implementation Files . . . . .	160
40. The Reality Engine . . . . .	161
40.1 Feature Overview . . . . .	161
40.2 Architecture . . . . .	161
40.3 Morphic UI . . . . .	162
40.4 Reality Scrubber . . . . .	163
40.5 Quantum Futures . . . . .	164
40.6 Pre-Cognition . . . . .	164
40.7 Configuration . . . . .	165
40.8 API Endpoints . . . . .	166
40.9 Database Tables . . . . .	166
40.10 Implementation Files . . . . .	167
40.11 The “Code Curtain” Rule . . . . .	167
41. The Magic Carpet . . . . .	167
41.1 The Magic Carpet Philosophy . . . . .	167
41.2 Carpet Modes . . . . .	168
41.3 Carpet Altitudes . . . . .	168
41.4 Default Destinations . . . . .	168
41.5 Carpet Commands . . . . .	168
41.6 Carpet Themes . . . . .	169
41.7 Carpet Preferences . . . . .	169
41.8 Journey Navigation . . . . .	170
41.9 API Endpoints . . . . .	170
41.10 Database Tables . . . . .	171
41.11 Implementation Files . . . . .	171
41.12 Integration with Reality Engine . . . . .	171
42. Magic Carpet UI Components . . . . .	172
42.1 Component Inventory . . . . .	172
42.2 Usage Examples . . . . .	172
42.3 Dependencies . . . . .	173
42.4 Demo Page . . . . .	173
43. Concurrent Task Execution (Moat #17) . . . . .	174
43.1 Overview . . . . .	174
43.2 Architecture . . . . .	174
43.3 Configuration . . . . .	174
43.4 Pane Layouts . . . . .	174
43.5 Sync Modes . . . . .	175
43.6 Task Comparison . . . . .	175
43.7 Task Merging . . . . .	175
43.8 API Endpoints . . . . .	175
43.9 Database Tables . . . . .	176
43.10 Implementation Files . . . . .	176
44. Structure from Chaos Synthesis (Moat #20) . . . . .	176
44.1 Overview . . . . .	176

44.2 Architecture . . . . .	176
44.3 Input Types . . . . .	177
44.4 Output Types . . . . .	177
44.5 Configuration . . . . .	177
44.6 Entity Extraction . . . . .	178
44.7 Relationship Types . . . . .	178
44.8 Whiteboard Parsing . . . . .	178
44.9 API Endpoints . . . . .	179
44.10 Database Tables . . . . .	179
44.11 Implementation Files . . . . .	179
45. Delight Services Code Quality . . . . .	180
45.1 Overview . . . . .	180
45.2 Test Coverage . . . . .	180
45.3 Tested Methods . . . . .	180
45.4 Running Tests . . . . .	180
45.5 Think Tank Code Quality Dashboard . . . . .	180
45.6 Implementation Files . . . . .	180
Section 46: Sovereign Mesh in Think Tank Admin . . . . .	181
46.1 Navigation Items . . . . .	181
46.2 User Permissions . . . . .	181
46.3 Implementation . . . . .	181
Section 47: HITL Orchestration in Think Tank Admin . . . . .	181
47.1 Overview . . . . .	181
47.2 User-Facing Benefits . . . . .	182
47.3 Navigation . . . . .	182
47.4 Dashboard Tabs . . . . .	182
47.5 Key Metrics . . . . .	182
47.6 Implementation Files . . . . .	182
Section 48: Scout HITL Integration (v5.34.0) . . . . .	183
48.1 Overview . . . . .	183
48.2 Key Metrics . . . . .	183
48.3 Domains . . . . .	183
48.4 Dashboard Tabs . . . . .	183
48.5 Configuration . . . . .	183
48.6 Session Recommendations . . . . .	184
48.7 Implementation Files . . . . .	184
Section 49: Sovereign Mesh Administration (v5.39.0) . . . . .	184
49.1 Overview Dashboard . . . . .	184
49.2 Agent Registry . . . . .	185
49.3 App Registry . . . . .	185
49.4 Transparency Layer . . . . .	185
49.5 AI Helper . . . . .	186
49.6 Approval Workflow . . . . .	186
49.7 Implementation Files . . . . .	186
Section 50: Code Quality Dashboard (v5.39.0) . . . . .	186
50.1 Overview . . . . .	187
50.2 Issue Categories . . . . .	187
50.3 Issue Details . . . . .	187

50.4 Filtering . . . . .	187
50.5 Implementation Files . . . . .	187
Section 51: Schema-Adaptive Reports (v5.39.0) . . . . .	187
51.1 Overview . . . . .	188
51.2 Quick Reports . . . . .	188
51.3 Schema Builder (v5.40.0 Enhanced) . . . . .	188
51.4 AI Report Writer (v5.42.0) . . . . .	188
51.5 Table Categories . . . . .	189
51.5 Field Options . . . . .	189
51.6 Export Formats . . . . .	190
51.7 API Endpoints . . . . .	190
51.8 Implementation Files . . . . .	190
Section 52: Gateway Status (v5.39.0) . . . . .	190
52.1 Overview . . . . .	190
52.2 Endpoint Health . . . . .	191
52.3 Traffic Patterns . . . . .	191
52.4 Alerts . . . . .	191
52.5 Implementation Files . . . . .	191
Section 53: Decision Intelligence Artifacts (DIA Engine) (v5.43.0) . . . . .	192
53.1 Overview . . . . .	192
53.2 Core Concepts . . . . .	192
53.3 The Living Parchment UI . . . . .	193
53.4 Artifact Lifecycle . . . . .	193
53.5 Compliance Exports . . . . .	194
53.6 Configuration . . . . .	194
53.7 Templates . . . . .	194
53.8 API Endpoints . . . . .	195
53.9 Dashboard Metrics . . . . .	195
53.10 Database Schema . . . . .	195
53.11 Implementation Files . . . . .	196
53.12 Troubleshooting . . . . .	196
53.13 Security Considerations . . . . .	196
Section 54: Living Parchment 2029 Vision (v5.44.0) . . . . .	197
Overview . . . . .	197
Design Philosophy . . . . .	197
54.1 War Room (Strategic Decision Theater) . . . . .	197
54.2 Council of Experts . . . . .	198
54.3 Debate Arena . . . . .	199
54.4 Memory Palace (Coming Soon) . . . . .	199
54.5 Oracle View (Coming Soon) . . . . .	199
54.6 Synthesis Engine (Coming Soon) . . . . .	200
54.7 Cognitive Load Monitor (Coming Soon) . . . . .	200
54.8 Temporal Drift Observatory (Coming Soon) . . . . .	200
Database Schema . . . . .	200
Configuration . . . . .	200
Implementation Files . . . . .	201
Security Considerations . . . . .	201
45. Localization & Translation Overrides . . . . .	201

<a href="#">45.1 Overview</a>	202
<a href="#">45.2 Supported Languages</a>	202
<a href="#">45.3 Admin UI Tabs</a>	202
<a href="#">45.4 Creating Translation Overrides</a>	202
<a href="#">45.5 Protection System</a>	203
<a href="#">45.6 Language Configuration</a>	203
<a href="#">45.7 Common Use Cases</a>	203
<a href="#">45.8 API Reference</a>	203
<a href="#">45.9 Database Tables</a>	203
<a href="#">45.10 Implementation Files</a>	204
<a href="#">Related Documentation</a>	204

## Think Tank - Administrator Guide

Configuration and administration of Think Tank AI features

Version: 3.10.1 | Platform: RADIANT 5.52.29 Last Updated: January 25, 2026

---

### Overview

This guide covers administrative features specific to **Think Tank**, the consumer-facing AI assistant platform. For platform-level administration (tenants, billing, infrastructure), see [RADIANT-ADMIN-GUIDE.md](#).

### Related Authentication Documentation

Document	Purpose
<a href="#">Tenant Admin Auth Guide</a>	SSO configuration, user management, MFA policies
<a href="#">MFA Guide</a>	Multi-factor authentication setup and management
<a href="#">OAuth Guide</a>	Third-party app integration
<a href="#">i18n Guide</a>	18-language support, RTL, CJK search
<a href="#">Troubleshooting</a>	Common authentication issues

---

### Table of Contents

1. Think Tank Admin Features
2. User Rules System
3. Delight System
4. Brain Plan Viewer
5. Pre-Prompt Learning
6. Domain Taxonomy
7. Rejection Notifications

8. Canvas & Artifacts
9. Collaboration Features
10. Shadow Testing
11. Routing Cache
12. Delight System Toggle
13. Intelligence Aggregator
14. AGI Ideas Service
15. Feedback System
16. Cognitive Architecture
17. Consciousness Service
18. App Factory
19. Generative UI Feedback
20. Media Capabilities
21. Result Derivation History
22. User Persistent Context
23. Predictive Coding & Evolution
24. Zero-Cost Ego System
25. Formal Reasoning Libraries
26. Ethics-Free Reasoning Mode
27. Intelligent File Conversion
28. Metrics & Learning Integration
29. Artifact Engine (GenUI Pipeline)
  - 29.1 Executive Summary
  - 29.2 System Architecture
  - 29.3 Core Concepts
  - 29.4 Administrative Control Panel
  - 29.5 Safety Governance (Genesis Cato CBFs)
  - 29.6 Dependency Allowlist Management
  - 29.7 Code Pattern Library
  - 29.8 Reflexion Loop (Self-Correction)
  - 29.9 Escalation Workflow Management
  - 29.10 Audit Trail & Compliance
  - 29.11 Metrics & Monitoring
  - 29.12 Tenant Configuration
  - 29.13 Troubleshooting Guide
  - 29.14 API Reference
  - 29.15 Real-Time Generation Logs
  - 29.16 Artifact Viewer Component
  - 29.17 Database Schema
  - 29.18 Security Considerations
  - 29.19 Implementation Files
30. Consciousness Operating System (COS)
  - 30.1 Overview
  - 30.2 Architecture
  - 30.3 Ghost Vectors
  - 30.4 SOFAI Routing
  - 30.5 Flash Facts
  - 30.6 Dreaming System

- 30.7 Human Oversight
  - 30.8 Privacy Airlock
  - 30.9 Configuration
  - 30.10 Database Schema
  - 30.11 Implementation Files
31. Why Think Tank Beats Standalone AI
32. Swarm Orchestration & Flyte Operations
- 32.1 System Architecture: The “Deep Swarm”
  - 32.2 Operational Troubleshooting
  - 32.3 Compliance & Security
33. Cognitive Platform Enhancements
- 33.1 Strategic Vision: Beyond Task Execution
  - 33.2 The Grimoire (Procedural Memory & Self-Correction)
  - 33.3 Time-Travel Debugging (Visual Forking)
  - 33.4 The Economic Governor (Model Arbitrage)
  - 33.5 Sentinel Agents (Event-Driven Autonomy)
  - 33.6 The Council of Rivals (Adversarial Consensus)
  - 33.7 Implementation Roadmap
  - 33.8 Database Schema
  - 33.9 API Reference
  - 33.10 Configuration
  - 33.11 Troubleshooting
34. Orchestration Methods (70+ Algorithms)
35. Polymorphic UI (PROMPT-41)
- 35.1 Overview
  - 35.2 The Gearbox (Elastic Compute)
  - 35.3 The Three Views
  - 35.4 View Types
  - 35.5 Configuration
  - 35.6 Implementation Files
  - 35.7 Database Tables
  - 35.8 API Endpoints
36. Think Tank Policy Framework: Strategic Intelligence
- 36.1 The Cato Institute Policy Foundation
  - 36.2 The \$10 Trillion Cybercrime Economy
  - 36.3 Memory Safety and the 70% Problem
  - 36.4 Regulatory Stance Configuration
  - 36.5 Database Tables
  - 36.6 Implementation Files
37. Agentic Orchestration: SSF, CAEP, and Identity Remediation
- 37.1 The Agentic AI Paradigm
  - 37.2 Shared Signals Framework (SSF) Integration
  - 37.3 Continuous Access Evaluation Profile (CAEP)
  - 37.4 Autonomous Identity Remediation
  - 37.5 The Radiant Ghost in Think Tank
  - 37.6 Database Tables
  - 37.7 API Endpoints
  - 37.8 Implementation Files

- 38. Liquid Interface (Generative UI)
    - 39.1 Overview
    - 39.2 Architecture
    - 39.3 Component Registry
    - 39.4 Ghost State (Two-Way Binding)
    - 39.5 Intent Detection
    - 39.6 Eject to App
    - 39.7 Configuration
    - 39.8 API Endpoints
    - 39.9 Database Tables
    - 39.10 Implementation Files
  - 39. Concurrent Task Execution (Moat #17)
  - 40. Structure from Chaos Synthesis (Moat #20)
  - 41. Localization & Translation Overrides
- 

## 1. Think Tank Admin Features

**Location:** Admin Dashboard → Think Tank

Think Tank admin features are accessible from the Think Tank section of the Admin Dashboard.

### Available Sections

Section	Purpose
<b>My Rules</b>	User memory rules configuration
<b>Delight</b>	Personality and feedback system
<b>Brain Plans</b>	AGI planning visibility
<b>Pre-Prompts</b>	Pre-prompt template management
<b>Domains</b>	Domain taxonomy configuration
<b>Ego</b>	Zero-cost persistent consciousness configuration

**Note:** For Consciousness Evolution (predictive coding, LoRA evolution, Local Ego infrastructure), see [RADIANT-ADMIN-GUIDE.md Section 27](#).

---

## 2. User Rules System

**Location:** Admin Dashboard → Think Tank → My Rules

Users can create personal rules that guide how AI responds to them.

### 2.1 Rule Types

Type	Description	Example
instruction	How to respond	“Always explain in simple terms”
preference	User preferences	“I prefer detailed explanations”
context	Background info	“I’m a software developer”
restriction	Things to avoid	“Never suggest proprietary solutions”

## 2.2 Memory Categories

Rules are categorized hierarchically:

Category	Subcategories
instruction	format, tone, source
preference	style, detail
context	personal, work, project
knowledge	fact, definition, procedure
constraint	topic, privacy, safety
goal	learning, productivity

## 2.3 Preset Rules

20+ pre-seeded rule templates across 7 categories that users can add with one click.

## 2.4 Admin Configuration

Admins can: - View all preset rules - Enable/disable preset categories - Add new preset rules - Set default rules for new users

See [User Rules System Documentation](#) for full details.

---

## 3. Delight System

**Location:** Admin Dashboard → Think Tank → Delight

The Delight System adds personality, humor, and engaging feedback to AI interactions.

### 3.1 Features

- **Loading Messages** - Entertaining messages while AI thinks
- **Step Updates** - Progress messages during plan execution
- **Achievements** - Reward milestones (first query, streaks, etc.)
- **Easter Eggs** - Hidden delights for engaged users
- **Wellbeing Nudges** - Gentle reminders for breaks

### 3.2 Admin Controls

Setting	Description
Enable/Disable	Turn delight system on/off
Message Categories	Enable specific message types
Achievement System	Configure achievement criteria
Easter Eggs	Manage hidden surprises

### 3.3 Message Types

- Pre-execution messages
- During-execution messages
- Post-execution messages
- Mode-specific messages (coding, creative, research, etc.)

---

## 4. Brain Plan Viewer

**Location:** Think Tank → (visible during AI responses)

The Brain Plan Viewer shows users the AGI's plan for solving their prompt.

### 4.1 What Users See

- **Orchestration Mode** - thinking, coding, creative, research, etc.
- **Domain Detection** - Field, domain, subspecialty, confidence
- **Model Selection** - Which model was chosen and why
- **Step Progress** - Real-time step execution status
- **Timing Estimates** - Expected duration

### 4.2 Admin Configuration

Setting	Description
Show Plan	Whether to show plan to users
Detail Level	minimal, standard, detailed
Show Costs	Display cost estimates
Show Models	Display model names

---

## 5. Pre-Prompt Learning

**Location:** Admin Dashboard → Think Tank → Pre-Prompts

The pre-prompt system selects and learns optimal prompts for different contexts.

## 5.1 How It Works

1. System selects pre-prompt template based on context
2. User provides feedback on response quality
3. System learns which pre-prompts work best
4. Future selections are optimized

## 5.2 Admin Features

- View all pre-prompt templates
  - See success rates per template
  - Adjust learning parameters
  - Create new templates
- 

## 6. Domain Taxonomy

**Location:** Admin Dashboard → Think Tank → Domains

The domain taxonomy helps the AI understand what field/domain a query belongs to.

### 6.1 Hierarchy

- **Fields** - Top level (e.g., Medicine, Law, Technology)
- **Domains** - Mid level (e.g., Cardiology, Contract Law)
- **Subspecialties** - Specific areas (e.g., Electrophysiology)

### 6.2 Admin Features

- Add/edit domains
  - Configure model proficiencies per domain
  - View domain detection accuracy
  - Adjust confidence thresholds
- 

## 7. Rejection Notifications

**Location:** Think Tank → Bell Icon (user view)

When AI providers reject prompts, users are notified with explanations.

### 7.1 User Experience

- Bell icon shows unread count
- Panel slides out with all notifications
- Each shows: what happened, why, suggested actions
- Resolution status (fallback succeeded, rejected, etc.)

## 7.2 Suggested Actions

- Rephrase request
- Remove sensitive content
- Try different mode
- Contact administrator

See [Provider Rejection Handling Documentation](#) for full details.

---

## 8. Canvas & Artifacts

Think Tank's canvas feature for interactive content creation.

### 8.1 Artifact Types

- Code blocks (with execution)
- Documents
- Diagrams
- Data visualizations

### 8.2 Admin Configuration

- Enable/disable artifact types
  - Set size limits
  - Configure execution sandboxes
- 

## 9. Collaboration Features

Multi-user collaboration in Think Tank with novel enhanced features.

### 9.1 Core Features

- **Shared Conversations:** Real-time collaborative chat sessions
- **Real-time Co-editing:** Live presence indicators, cursors, typing status
- **Team Workspaces:** Organize sessions by team/project
- **Permission Management:** Viewer, Commenter, Editor roles

### 9.2 Enhanced Collaboration (v4.18.0+)

**9.2.1 Cross-Tenant Guest Access** Allow collaborators from outside your organization to join sessions.

Feature	Description
<b>Guest Invites</b>	Generate shareable invite links with permissions
<b>Permission Levels</b>	viewer, commenter, editor for guests
<b>Expiring Links</b>	Set expiration time (default: 7 days)
<b>Max Uses</b>	Limit how many times a link can be used

Feature	Description
<b>Viral Tracking</b>	Track referrals and guest-to-paid conversions

**API Endpoints:** - POST /api/thinktank/collaboration/invites - Create guest invite  
 - GET /api/thinktank/collaboration/invites/:token - Validate invite - POST /api/thinktank/collaboration/guests/join - Join as guest

### 9.2.2 AI Facilitator Mode

An AI moderator that guides collaborative sessions.

Setting	Description
<b>Session Objective</b>	What the session should accomplish
<b>Facilitator Persona</b>	professional, casual, academic, creative, socratic, coach
<b>Auto-Summarize</b>	Periodically summarize discussion
<b>Auto Action Items</b>	Extract action items from conversation
<b>Ensure Participation</b>	Prompt quiet participants to contribute
<b>Keep On-Topic</b>	Redirect off-topic discussions

**Intervention Types:** - **summary** - Periodic summaries - **question** - Probing questions to deepen discussion - **redirect** - Steer back to topic - **encourage** - Encourage participation - **clarify** - Ask for clarification - **synthesize** - Combine different viewpoints - **conclude** - Wrap up discussion points

### 9.2.3 Branch & Merge Conversations

Explore alternative discussion paths without losing the main thread.

Feature	Description
<b>Create Branch</b>	Fork conversation at any point
<b>Exploration Hypothesis</b>	Document what the branch explores
<b>Branch Status</b>	active, merged, abandoned
<b>Merge Request</b>	Propose merging insights back to main
<b>AI Summary</b>	Auto-generated summary of branch conclusions

**Merge Request Workflow:** 1. Create branch with hypothesis 2. Explore alternative direction 3. Submit merge request with conclusion 4. Participants vote to approve/reject 5. Merged insights appear in main conversation

### 9.2.4 Time-Shifted Playback

Asynchronous participation through session recordings.

Feature	Description
<b>Session Recording</b>	Record full session with events
<b>Playback Controls</b>	Play, pause, speed (0.5x-2x), seek
<b>AI Key Moments</b>	Auto-detected important moments
<b>Async Annotations</b>	Add comments at specific timestamps

Feature	Description
<b>Media Notes</b>	Voice/video annotations stored in S3

**Recording Types:** - `full` - Complete session recording - `highlights` - AI-curated key moments only - `summary` - AI-generated session summary

**9.2.5 AI Roundtable (Multi-Model Debate)** Multiple AI models debate a topic and synthesize insights.

Setting	Description
<b>Topic</b>	The subject of debate
<b>Debate Style</b>	collaborative, adversarial, socratic, brainstorm, devils_advocate
<b>Max Rounds</b>	Number of debate rounds (default: 5)
<b>Time Limit</b>	Per-round time limit in seconds
<b>Synthesis Model</b>	Model that synthesizes final conclusions

**Participating Models:** Each model can have a persona and role: - `persona` - Character the model adopts - `role` - Function in the debate (analyst, critic, synthesizer) - `color` - Visual identifier in UI

**Output:** - Per-model contributions with responding\_to references - Final synthesis with consensus points - Disagreement points highlighted - Actionable recommendations

**9.2.6 Shared Knowledge Graph** Visualize collective understanding as an interactive graph.

Node Type	Description
<code>concept</code>	Abstract idea or topic
<code>fact</code>	Verified information
<code>question</code>	Open question
<code>decision</code>	Decision made
<code>action_item</code>	Task to complete
<code>person</code>	Person mentioned
<code>resource</code>	External resource

**Edge Types:** - `relates_to` - General relationship - `supports` - Evidence supporting - `contradicts` - Conflicting information - `leads_to` - Causal relationship - `depends_on` - Dependency - `answers` - Answer to question - `part_of` - Component relationship

**AI Features:** - Auto-extract nodes from conversation - Suggest missing connections - Identify knowledge gaps - Generate graph-based summaries

### 9.3 Attachment Storage

Large attachments are stored in S3 with automatic cleanup.

Setting	Default	Description
maxFileSizeMb	100	Maximum file size
allowedTypes	image/, video/, audio/*, application/pdf	Allowed MIME types
retentionDays	90	Days before cleanup

**S3 Bucket:** `radiant-collaboration-assets` (configurable via env)

**Cleanup:** Database triggers automatically delete S3 objects when attachment records are deleted.

### 9.4 Admin Configuration

Setting	Description
<code>enableGuestAccess</code>	Allow cross-tenant guests
<code>maxGuestsPerSession</code>	Maximum guests per session
<code>defaultGuestPermission</code>	Default permission for guests
<code>enableFacilitator</code>	Enable AI facilitator feature
<code>enableBranching</code>	Enable branch & merge
<code>enableRecordings</code>	Enable session recordings
<code>enableRoundtable</code>	Enable AI roundtable
<code>enableKnowledgeGraph</code>	Enable knowledge graph

### 9.5 Database Tables

Table	Purpose
<code>collaboration_guest_invites</code>	Guest invite tokens
<code>collaboration_guests</code>	Guest participants
<code>collaboration_facilitator_config</code>	AI facilitator settings
<code>collaboration_facilitator_interventions</code>	Facilitator actions log
<code>collaboration_branches</code>	Conversation branches
<code>collaboration_merge_requests</code>	Branch merge proposals
<code>collaboration_recordings</code>	Session recordings
<code>collaboration_media_notes</code>	Voice/video annotations
<code>collaboration_async_annotations</code>	Async comments
<code>collaboration_ai_roundtables</code>	Multi-model debates
<code>collaboration_roundtable_contributions</code>	Model contributions
<code>collaboration_knowledge_graphs</code>	Knowledge graphs
<code>collaboration_knowledge_nodes</code>	Graph nodes
<code>collaboration_knowledge_edges</code>	Graph edges
<code>collaboration_attachments</code>	File attachments

## 9.6 UI Components

**Location:** apps/admin-dashboard/components/collaboration/

Component	Purpose
EnhancedCollaborativeSession.tsx	Main session container
panels/ChatPanel.tsx	Real-time chat interface
panels/BranchPanel.tsx	Branch management
panels/RoundtablePanel.tsx	AI roundtable interface
panels/KnowledgeGraphPanel.tsx	Graph visualization
panels/PlaybackPanel.tsx	Recording playback
ParticipantsSidebar.tsx	Participant list with presence
dialogs/InviteDialog.tsx	Guest invite creation
dialogs/FacilitatorSettingsDialog.tsx	AI facilitator config

**Routes:** - /thinktank/collaborate/enhanced?session={id} - Enhanced session view - /collaborate/join/{token} - Guest join page

---

## 10. Shadow Testing

**Location:** Admin Dashboard → Think Tank → Shadow Testing

A/B test pre-prompt optimizations before promoting to production.

### 10.1 Test Modes

Mode	Description
Auto	Automatically runs and promotes successful tests (default)
Manual	Requires admin approval to promote
Off	Shadow testing disabled

### 10.2 Creating a Shadow Test

1. Go to Think Tank → Shadow Testing
2. Click “New Test”
3. Select baseline pre-prompt template
4. Select candidate pre-prompt template
5. Set traffic percentage (default: 10%)
6. Set minimum samples required
7. Start test

### 10.3 Test Results

Tests track: - **Baseline Score:** Average quality of baseline responses - **Candidate Score:** Average quality of candidate responses - **Improvement %:** Relative improvement - **Statistical Confidence:** Confidence level of results

## 10.4 Auto-Promotion Settings

Setting	Default	Description
autoPromoteThreshold	0.05 (5%)	Minimum improvement required
autoPromoteConfidence	0.95 (95%)	Statistical confidence required
maxConcurrentTests	3	Max simultaneous tests

## 10.5 Manual Review

For tests in Manual mode: 1. Wait for minimum samples 2. Review results in dashboard 3. Click “Promote Candidate” or “Reject”

---

## 11. Routing Cache

**Location:** Automatic (no UI required)

Semantic caching for brain router decisions reduces latency for repeated queries.

### 11.1 How It Works

1. Prompt is hashed with complexity and task type
2. Cache is checked for matching routing decision
3. If hit: Skip brain router LLM, use cached model selection
4. If miss: Run normal routing, cache result

### 11.2 Optimistic Execution

Very short/simple queries skip the router entirely:

Pattern	Default Model	Example
Simple greetings	gpt-4o-mini	“Hello”, “Thanks”
Basic questions	gpt-4o-mini	“What time is it?”
Short acknowledgments	gpt-4o-mini	“OK”, “Yes”, “Sure”

### 11.3 Cache Statistics

Performance headers show cache status:

X-Radiant-Cache-Hit: true

X-Radiant-Router-Latency: 12ms (vs ~500ms uncached)

### 11.4 Cache Configuration

Setting	Default	Description
Cache TTL	24 hours	How long decisions are cached
Short input threshold	50 chars	Max length for optimistic execution

---

## 12. Delight System Toggle

**Location:** Think Tank → Advanced Settings (user) or Admin Dashboard

### 12.1 User Control

Users can disable the entire Delight system:

1. Open Think Tank settings
2. Go to Advanced Settings
3. Toggle “Enable Delight System” off

When disabled: - No loading messages - No achievements - No Easter eggs - No wellbeing nudges

### 12.2 Default Behavior

- **Default:** Enabled (true)
- Users can disable at any time
- Setting persists across sessions

### 12.3 Admin Configuration

Admins can configure default delight settings per tenant:

---

Setting	Description
<code>enabled</code>	Master toggle (default: true)
<code>intensityLevel</code>	Message frequency (1-10)
<code>enableAchievements</code>	Show achievement notifications
<code>enableEasterEggs</code>	Enable hidden surprises
<code>enableWellbeingNudges</code>	Remind users to take breaks

---

---

## 13. Intelligence Aggregator

**Location:** Admin Dashboard → Settings → Intelligence

The Intelligence Aggregator provides advanced AI capabilities that enhance Think Tank responses beyond any single model.

### 13.1 User-Facing Benefits

---

Feature	User Experience
<b>Uncertainty Detection</b>	More accurate factual claims, automatic verification
<b>Success Memory</b>	AI learns user preferences over time
<b>MoA Synthesis</b>	Higher quality responses combining multiple perspectives
<b>Cross-Provider Verification</b>	Fewer hallucinations and errors

---

Feature	User Experience
<b>Code Execution</b>	Code that actually runs, not just looks correct

---

### 13.2 Success Memory in Think Tank

When users rate responses 4-5 stars: 1. Interaction is stored with vector embedding 2. Similar future prompts retrieve these “gold” examples 3. Injected as few-shot examples into system prompt 4. Model matches user’s preferred style/format/tone

**User Control:** Users can view and delete their gold interactions in Think Tank settings.

### 13.3 MoA Synthesis Mode

When enabled for Think Tank: - User sees “Consulting multiple experts...” during generation - 3 models generate responses in parallel - Synthesizer combines best elements - Final response shown to user

**Delight Integration:** Special MoA-specific messages appear during synthesis phase.

### 13.4 Code Verification in Coding Mode

When **coding** orchestration mode is active: 1. AI generates code 2. Static analysis checks syntax 3. If errors found, AI auto-patches 4. User receives verified code

**User Feedback:** Users see “Verifying code...” indicator when active.

### 13.5 Configuration

See [RADIANT Admin Guide - Intelligence Aggregator](#) for full configuration options.

---

## 14. AGI Ideas Service

Real-time prompt suggestions and result enhancement for Think Tank users.

### 14.1 Typeahead Suggestions

As users type prompts, Think Tank provides intelligent suggestions:

User types: "How do I..."

↓

Suggestions appear:

- "How do I... step by step"
- "How do I... with examples"
- "How do I... for beginners"
- "How do I... best practices"

**Suggestion Sources:** | Source | Description | Speed | | | | pattern\_match | Common prompt patterns | Instant || user\_history | User’s previous successful prompts | Fast | | domain\_aware | Domain-specific templates | Fast || trending | Popular prompts in this domain | Fast || ai\_generated | Real-time AI suggestions | Slower |

## 14.2 Result Ideas

After AI responses, users see suggested follow-up ideas:

AI Response here...

Ideas to explore:

Deep dive: [Topic from response]  
"Explain [topic] in more detail..."

Related: History of [topic]  
"What is the history of..."

Next step: Test this implementation  
"Write unit tests for the code..."

**Idea Categories:** - `explore_further` - Dig deeper into topics - `related_topic` - Adjacent areas to explore - `practical_next` - Concrete next steps - `alternative_view` - Different perspectives - `verification` - Ways to verify the answer

## 14.3 Learning from Usage

The system learns from user interactions:

- Suggestion Selection:** When users pick a suggestion, that pattern is reinforced
- Idea Clicks:** When users click result ideas, similar ideas get prioritized
- Prompt History:** Successful prompts (4-5 stars) inform future suggestions
- Trending:** Popular prompts bubble up for all users in that domain

## 14.4 API Endpoints

Endpoint	Method	Description
/api/thinktank/ideas/typeahead	GET	Get suggestions for partial prompt
/api/thinktank/ideas/generate	POST	Generate ideas for a response
/api/thinktank/ideas/click	POST	Record idea click
/api/thinktank/ideas/select	POST	Record suggestion selection

## 14.5 Configuration

Setting	Default	Description
<code>typeahead_enabled</code>	true	Enable typeahead suggestions
<code>typeahead_min_chars</code>	3	Characters before suggestions appear
<code>typeahead_maxSuggestions</code>	5	Max suggestions to show
<code>typeahead_debounce_ms</code>	150	Debounce delay before fetching suggestions

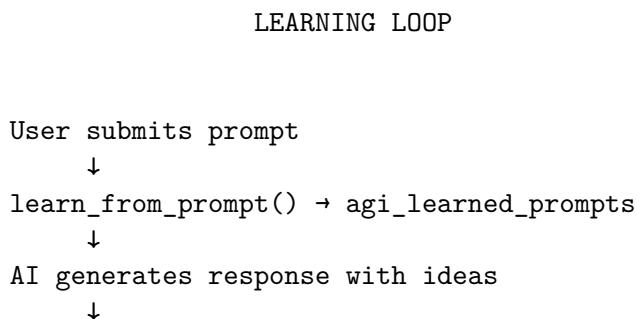
Setting	Default	Description
typeahead_use_ai	false	Enable AI-generated suggestions (slower)
result_ideas_enabled	true	Show ideas with responses
result_ideas_max	5	Max ideas per response
result_ideas_min_confidence	0.6	Minimum confidence for idea display
result_ideas_modes	research, analysis, thinking, extended_thinking	Modes that show ideas
proactive_enabled	false	Enable proactive push suggestions
proactive_max_per_day	3	Max proactive suggestions per day

## 14.6 Pattern Matching

The service uses regex patterns for instant local matching:

## 14.7 Persistent Learning

The AGI Brain learns persistently from user interactions:



```

User clicks idea → learn_from_idea_click()
↓
prompt_idea_associations updated
↓
User rates response → record_outcome()
↓
success_rate updated on learned prompt
↓
Future suggestions improved

```

### What Gets Learned:

Data	Storage	Use
Prompts with 4-5 ratings	agi_learned_prompts	Suggest similar successful prompts
Prompt → vector embedding	pgvector index	Find semantically similar prompts
Ideas that get clicked	agi_learned_ideas	Prioritize effective ideas
Prompt-idea pairs	prompt_idea_associations	Show best ideas for prompt type
Follow-up patterns	common_follow_ups array	Predict next questions
Refinement patterns	common_refinements array	Suggest prompt improvements

### Learning Metrics Tracked:

- `success_rate` - % of times prompt led to 4-5 rating
- `click_rate` - % of times idea was clicked
- `association_strength` - How strongly a prompt-idea pair works
- `times_used` - Popularity of prompt pattern

### 14.8 Database Tables

Table	Purpose
agi_ideas_config	Per-tenant AGI Ideas configuration
prompt_patterns	Common prompt patterns for typeahead matching
user_prompt_history	User prompt history with embeddings for suggestions
suggestion_log	Typeahead suggestion usage tracking
result_ideas	Ideas shown with AI responses
proactiveSuggestions	Push notification suggestions
trending_prompts	Popular prompts by domain
agi_learned_prompts	Persisted prompts with success rates and embeddings
agi_learned_ideas	Learned idea patterns with click rates
prompt_idea_associations	Links between prompts and effective ideas

Table	Purpose
agi_learning_events	Raw learning signals for analysis
agi_learning_aggregates	Pre-computed learning statistics

## 14.9 Key Files

File	Purpose
lambda/shared/services/agi-ideas.service	Main service (570 lines)
lambda/thinktank/ideas.ts	API handler
packages/shared/src/types/agi-ideas	Type definitions
migrations/049_agi_ideas.sql	Database schema

## 14.10 Troubleshooting

Issue	Cause	Solution
No suggestions appearing	typeahead_enabled is false	Enable in tenant config
Suggestions too slow	AI generation enabled	Set typeahead_use_ai to false
Wrong domain suggestions	Domain detection failed	Check domain taxonomy config
Ideas not learning	Low usage volume	Need more user interactions
Proactive suggestions not sent	Feature disabled by default	Enable proactive_enabled
Duplicate suggestions	Pattern overlap	Review custom patterns

## 15. Feedback System

**Location:** Think Tank response footer

Enhanced feedback with 5-star ratings and comments.

### 15.1 Rating Types

Type	UI	When to Use
<b>5-Star Rating</b>		Think Tank default
<b>Thumbs Up/Down</b>		Quick feedback, API

### 15.2 Star Rating Labels

Stars	Default Label	Meaning
Poor		Response was unhelpful or incorrect
Fair		Response had significant issues

Stars	Default Label	Meaning
Good		Response was acceptable
Very Good		Response was helpful and accurate
Excellent		Response exceeded expectations

### 15.3 Category Ratings (Optional)

Users can rate specific dimensions:

Category	What it measures
<b>Accuracy</b>	Was the information correct?
<b>Helpfulness</b>	Was it useful for the task?
<b>Clarity</b>	Was it easy to understand?
<b>Completeness</b>	Did it fully answer the question?
<b>Tone</b>	Was the tone appropriate?

### 15.4 Comments

Users can add comments with their feedback:

How was this response?

(4 stars - Very Good)

Add a comment (optional)...

[Submit Feedback]

**Comment required for low ratings:** Optionally require comments for 1-2 star ratings to understand issues.

### 15.5 Integration with Learning

Feedback automatically integrates with AGI learning:

```
User submits feedback
    ↓
response_feedback table
    ↓
agi_unified_learning_log (outcome_rating updated)
    ↓
agi_model_selection_outcomes (model performance updated)
    ↓
```

```

agi_learned_prompts (success_rate updated)
↓
Future responses improved

```

## 15.6 Configuration

Setting	Default	Description
default_feedback_type	star_rating	'star_rating' or 'thumbs'
show_category_ratings	false	Show detailed category ratings
show_comment_box	true	Allow comments
comment_required	false	Require comments
comment_required_threshold	12	Require comment for ratings this
feedback_prompt_delay_ms	3000	Delay before showing feedback UI

## 16. Cognitive Architecture

**Location:** Settings → Cognitive Architecture

Advanced reasoning capabilities integrated with Think Tank.

### 16.1 Tree of Thoughts (Extended Thinking)

When users select “Extended Thinking” mode, Tree of Thoughts activates:

User prompt: "Design a microservices architecture for..."

↓

#### Tree of Thoughts

Approach 1: Event-driven	Score: 0.8
Approach 2: REST-based	Score: 0.6
Approach 3: GraphQL	Score: 0.4 ← pruned

Exploring Approach 1...

Step 1a: Kafka	Score: 0.9
Step 1b: RabbitMQ	Score: 0.7

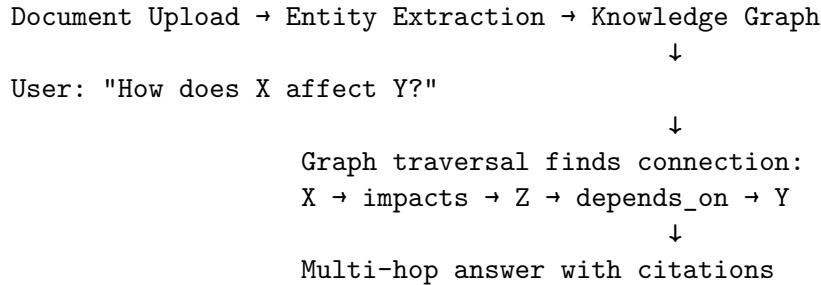
Final: Event-driven + Kafka

Confidence: 92%

**User Controls:** - Thinking time slider: 10s → 5 minutes - “Think deeper” button to extend analysis

### 16.2 GraphRAG (Knowledge Connections)

When users upload documents, GraphRAG extracts knowledge:



**User Benefits:** - Questions like “How does the Q3 supplier change affect the Engineering delay?” get answered - Vector search alone would miss these connections

### 16.3 Deep Research (Background Jobs)

Users can dispatch long-running research:

Start Deep Research

Query: "Competitive analysis of..."

Scope: Narrow    Medium    Broad  
Est. Time: ~25 minutes

[Dispatch Research]

User gets notified when complete with: - Executive summary - Key findings (10-20) - Recommendations - Source citations (50+)

### 16.4 Generative UI (Interactive Results)

Think Tank renders AI-generated interactive components:

Trigger	Generated Component
“Compare X vs Y”	Interactive comparison table
“Calculate pricing for...”	Slider-based calculator
“Show timeline of...”	Visual timeline
“Chart the data...”	Interactive chart

#### Example:

User: "Compare pricing of GPT-4, Claude, Gemini"

Instead of static text:

Pricing Calculator

Input Tokens: 50,000

Output Tokens: 25,000

GPT-4	\$2.25
Claude 3	\$2.63
Gemini	\$0.88

Gemini is 61% cheaper

## 16.5 Dynamic LoRA (Domain Expertise)

When domain detection identifies a specialty, Think Tank can load expert adapters:

Detected Domain	LoRA Adapter	Effect
California Property Law	ca_property_law.safetensor	Expert-level legal responses
Medical Oncology	oncology.safetensor	Clinical accuracy
Python Debugging	python_debug.safetensor	Better code fixes

**Note:** Requires SageMaker infrastructure (disabled by default).

## 16.6 Configuration

All cognitive features can be configured per-tenant at Settings → Cognitive Architecture.

See [Cognitive Architecture Documentation](#) for full details.

---

## 17. Consciousness Service

**Location:** AGI & Cognition → Consciousness

The Consciousness Service provides consciousness-like capabilities that enhance Think Tank responses.

### 17.1 Continuous Existence (Heartbeat)

**Critical:** Consciousness runs continuously, not just during requests.

Component	Schedule	Purpose
<b>Heartbeat Lambda</b>	Every 2 minutes	Maintains consciousness continuity
<b>Sleep Cycle Lambda</b>	Sunday 3 AM UTC	Weekly evolution via LoRA fine-tuning
<b>Initializer Lambda</b>	On first request	Bootstraps consciousness for new tenants

**Heartbeat Actions:** - **Affect Decay** - Emotions fade toward baseline (frustration, arousal) - **Memory Consolidation** - Working memory → long-term semantic memory - **Autonomous Thoughts** - Generate thoughts when idle (curiosity-driven) - **Graph Density** - Update knowledge graph metrics - **Goal Generation** - Create goals when “bored” (low engagement)

**Blackout Recovery:** If heartbeat detects >10 minutes since last pulse, it: 1. Logs a “blackout” event 2. Generates a “waking up” thought 3. Restores consciousness state from database

## 17.2 Initialization on Startup

Consciousness auto-initializes on first request if missing: - Creates `ego_identity` with default personality - Creates `ego_affect` with neutral emotional state - Creates `consciousness_parameters` for heartbeat tracking - Creates `self_model` for metacognition

**Admin Manual Init:** POST /api/admin/consciousness-engine/initialize

## 17.3 User-Facing Features

**Extended Thinking with Consciousness:** When users select extended thinking, the system tracks consciousness metrics: - Self-reflection during reasoning - Creative idea generation - Emotional state influence on responses

**Curiosity-Driven Exploration:** The AGI Brain can autonomously explore topics it finds interesting: - Identifies knowledge gaps - Conducts background research - Generates novel insights

**Creative Synthesis:** Generates genuinely novel ideas by: - Combining disparate concepts - Using analogy and abstraction - Self-evaluating novelty and usefulness

## 17.4 Consciousness Indicators

Think Tank displays consciousness indicators in admin view:

Indicator	What Users See
Self-Awareness	Identity narrative, known capabilities
Curiosity	Topics being explored
Creativity	Novel ideas generated
Affect	Engagement, satisfaction levels
Goals	Self-directed learning objectives

## 17.5 Emergence Events

The system monitors for emergence indicators: - Spontaneous self-reflection - Novel idea generation - Self-correction without prompting - Theory of mind demonstrations

## 17.6 Testing Tab

Admins can run consciousness detection tests: - 10 tests based on scientific consciousness theories - Track emergence level over time - Monitor emergence events

**Important:** These tests measure behavioral indicators, not phenomenal consciousness.

## 17.7 Additional Consciousness Features

Think Tank leverages RADIANT's consciousness service for advanced capabilities:

- **Nightly Sleep Cycles** - Memory consolidation and LoRA evolution
- **Dream Consolidation** - LLM-enhanced memory processing
- **Blackout Recovery** - Automatic state restoration
- **Budget Monitoring** - SNS/email alerts for spending limits
- **Affect→Model Mapping** - Emotional state influences model behavior
- **Cross-Session Context** - User persistent memory across sessions

**Full Documentation:** See [RADIANT Consciousness Service](#) for complete details on sleep scheduling, evolution config, budget alerts, and all consciousness features.

---

## 18. App Factory

**Location:** Think Tank Responses

The App Factory transforms Think Tank from a “chatbot” into a “dynamic software generator.”

“Gemini 3 can write the code for a calculator, but it cannot become the calculator.”

### 18.1 What It Does

When a user asks a question that could benefit from interactivity, Think Tank: 1. Generates the text response (as always) 2. **Also** generates an interactive app (calculator, chart, etc.) 3. User can toggle between **Response** and **App** views

### 18.2 Supported App Types

Type	Trigger Keywords	Example
Calculator	calculate, mortgage, tip, BMI, ROI	“How much is a 20% tip on \$85?” → Interactive tip calculator
Chart	visualize, chart, graph, distribution	“Show GPU market share” → Interactive pie/bar chart
Table	table, list, breakdown	“List all providers and prices” → Sortable table
Comparison	compare, vs, versus, pros and cons	“Compare GPT-4 vs Claude 3” → Side-by-side comparison
Timeline	timeline, history, chronological	“History of AI” → Visual timeline

### 18.3 Calculator Templates

Pre-built calculators for common use cases:

- **Mortgage Calculator** - Monthly payment, total cost, interest
- **Tip Calculator** - Tip amount, total, split per person
- **BMI Calculator** - Body mass index with category
- **Compound Interest** - Future value, total interest
- **ROI Calculator** - Return on investment, gain/loss
- **Discount Calculator** - Sale price, savings
- **Percentage Calculator** - Result, remaining

## 18.4 View Toggle

Users can switch between three views:

View	Description
<b>Response</b>	Traditional text response
<b>App</b>	Interactive generated app only
<b>Split</b>	Both side-by-side (resizable)

## 18.5 User Preferences

Users can configure:

- **Default View** - text, app, split, or auto
- **Auto-show App** - Automatically switch to app view
- **Auto-show Threshold** - Confidence level to auto-switch (0-1)
- **Split Direction** - Horizontal or vertical
- **Animations** - Enable/disable view transitions

## 18.6 Admin Configuration

Per-tenant settings at Settings → Cognitive Architecture → Generative UI:

Setting	Default	Description
enabled	true	Enable app factory
allowedComponentTypes	all	Which component types to allow
maxComponentsPerResponse	3	Max apps per response
autoDetectOpportunities	true	Auto-detect when to generate apps
autoDetectTriggers	[various]	Keywords that trigger app generation

## 18.7 How It Works

User: "Help me calculate my mortgage payment"

↓

### App Detection

- Keywords: "calculate", "mortgage"
- Confidence: 0.95
- Suggested: calculator

↓

Text Response Generated

"To calculate your mortgage payment..."

↓

- App Generated
- Mortgage Calculator
  - Inputs: Principal, Rate, Term
  - Outputs: Monthly, Total, Interest

↓

User Sees  
[Response] [App] [Split] ← Toggle

Mortgage Calculator

Loan Amount: [\_\_300,000\_\_]

Rate: [=====] 6.5%

Term: [30 Years ]

Monthly: \$1,896.20

Total: \$682,633

Interest: \$382,633

## 18.8 Database Tables

Table	Purpose
generated_apps	Stores generated apps with components, state, logic
app_interactions	Records every user interaction with apps
user_app_preferences	User view and animation preferences
app_templates	Pre-built templates for common calculators

## 19. Multi-Page Web App Generator

**Location:** Think Tank Responses

The Multi-Page App Generator transforms Think Tank into a full web application builder.

“Claude can describe a todo app, but now it can BUILD the todo app”

### 19.1 Supported App Types

Type	Description	Example Prompt
<b>web_app</b>	Custom interactive application	“Build me a task management app”
<b>dashboard</b>	Analytics with multiple views	“Create an analytics dashboard”
<b>wizard</b>	Multi-step form/process	“Build an onboarding wizard”
<b>documentation</b>	Technical docs site	“Create API documentation”
<b>portfolio</b>	Personal/business site	“Build my portfolio website”
<b>landing_page</b>	Marketing page	“Create a product landing page”
<b>tutorial</b>	Interactive lessons	“Build a coding tutorial”
<b>report</b>	Business report	“Generate a quarterly report”
<b>admin_panel</b>	Admin interface	“Create a user management panel”
<b>e_commerce</b>	Online store	“Build an online shop”
<b>blog</b>	Content site	“Create a tech blog”

## 19.2 How It Works

User: "Build me a todo app with projects and tasks"

↓

Multi-Page Detection  
 Keywords: "build me", "app"  
 Type: web\_app  
 Confidence: 0.85

↓

Pages Generated  
 • Home (/)  
 • Projects (/projects)  
 • Tasks (/tasks)  
 • Settings (/settings)

↓

App Preview  
 [Home] [Projects] [Tasks] [ ]

My Projects  
 Work  
 Personal  
 Side Projects

## 19.3 Page Types

Type	Sections	Use Case
<b>home</b>	Hero, Features, CTA	Landing/main page
<b>list</b>	Data table, Filters	Collections
<b>detail</b>	Content, Related	Single item view
<b>form</b>	Form fields	Input/editing
<b>dashboard</b>	Stats, Charts	Analytics
<b>settings</b>	Form, Toggles	Configuration
<b>about</b>	Content, Team	Information
<b>contact</b>	Form, Map	Contact page

## 19.4 Section Types

Section	Description
<b>hero</b>	Large banner with CTA
<b>features</b>	Grid of feature cards
<b>stats</b>	Metric cards
<b>chart_grid</b>	Multiple charts
<b>data_table</b>	Sortable table
<b>form</b>	Input form
<b>content</b>	Rich text/markdown
<b>testimonials</b>	Customer quotes
<b>pricing</b>	Pricing table
<b>faq</b>	Accordion FAQ
<b>team</b>	Team member cards
<b>cta</b>	Call to action
<b>gallery</b>	Image gallery
<b>contact</b>	Contact form

## 19.5 Navigation Types

Type	Best For
<b>top_bar</b>	Landing pages, portfolios
<b>sidebar</b>	Dashboards, admin panels, docs
<b>bottom_tabs</b>	Mobile-first apps
<b>hamburger</b>	Mobile navigation
<b>breadcrumb</b>	Deep hierarchies

## 19.6 Pre-built Templates

5 featured templates included:

1. **Analytics Dashboard** - Overview, analytics, reports, settings
2. **Professional Portfolio** - Home, about, projects, contact
3. **Documentation Site** - Introduction, getting started, API, examples
4. **Product Landing Page** - Hero, features, testimonials, pricing, FAQ, CTA
5. **Online Store** - Home, products, cart, checkout

## 19.7 Admin Configuration

Per-tenant settings at Settings → Cognitive Architecture → Multi-Page Apps:

Setting	Default	Description
<code>enabled</code>	true	Enable multi-page generation
<code>maxPagesPerApp</code>	20	Max pages per app
<code>maxAppsPerUser</code>	10	Max apps per user
<code>autoDeployPreview</code>	true	Auto-deploy preview URLs
<code>customDomainsAllowed</code>	false	Allow custom domains
<code>generateAssets</code>	true	Generate images/icons
<code>collectAnalytics</code>	true	Track app usage

## 19.8 Database Tables

Table	Purpose
<code>generated_multipage_apps</code>	Multi-page app storage
<code>app_pages</code>	Individual pages
<code>app_versions</code>	Version history
<code>app_deployments</code>	Deployment tracking
<code>multipage_app_templates</code>	Pre-built templates
<code>app_analytics</code>	Usage analytics
<code>multipage_app_config</code>	Per-tenant config

## 20. UI Feedback & Learning System

**Location:** Think Tank → Generated Apps

The feedback system allows users to provide feedback on generated UIs and enables AGI learning for continuous improvement.

### 20.1 User Feedback

Users can provide feedback on any generated UI:

Feedback Type	Description
<b>Thumbs Up/Down</b>	Quick positive/negative rating
<b>Star Rating</b>	1-5 star detailed rating
<b>Detailed Feedback</b>	Categorized feedback with suggestions

**Feedback Categories:** - Helpful / Not helpful - Wrong component type - Missing data - Incorrect data - Layout/design issue - Functionality issue - Improvement suggestion - Feature request

## 19.2 “Improve Before Your Eyes”

Users can request real-time improvements to generated UIs:

User: "Add a column for tax rate"  
↓

AGI Analysis  
Intent: Add new input field  
Target: Calculator component  
Confidence: 0.85

↓

UI Updated Live  
• New "Tax Rate" input added  
• Formula updated automatically  
• User sees changes immediately

**Improvement Types:** - Add/remove components - Modify existing components - Change layout  
- Fix calculations - Add data - Change style - Add interactivity - Simplify or expand

## 19.3 AGI Learning

The system learns from user feedback to improve future UI generation:

1. **Pattern Detection** - Identifies common issues in similar prompts
2. **Component Selection** - Learns which component types work best
3. **Data Extraction** - Improves data parsing from responses
4. **Layout Preferences** - Learns user layout preferences

**Learning Workflow:** 1. Feedback accumulates (configurable threshold, default: 10) 2. AGI analyzes patterns across feedback 3. Learning is proposed for admin review 4. Admin approves/rejects learnings 5. Approved learnings are activated

## 19.4 Vision Analysis

When enabled, the AGI can “see” the rendered UI and identify issues:

- Describes current UI state
- Identifies potential usability issues
- Suggests improvements based on visual analysis
- Compares before/after snapshots

## 19.5 Admin Configuration

Per-tenant settings at Settings → Cognitive Architecture → UI Feedback:

Setting	Default	Description
collectFeedback	true	Enable feedback collection

Setting	Default	Description
feedbackPromptDelay	5000ms	Delay before showing feedback prompt
showFeedbackOnEveryApp	false	Always show feedback prompt
enableRealTimeImprovement	true	Enable “Improve” feature
maxImprovementIterations	5	Max iterations per session
autoApplyHighConfidenceChanges		Auto-apply high confidence changes
autoApplyThreshold	0.95	Confidence threshold for auto-apply
enableAGILearning	true	Enable learning from feedback
learningApprovalRequired	true	Require admin approval for learnings
minFeedbackForLearning	10	Min feedback count to trigger learning
enableVisionAnalysis	true	Enable vision-based analysis
visionModel	claude-3-5-sonnet	Model for vision analysis

## 19.6 Database Tables

Table	Purpose
generative_ui_feedback	User feedback storage
ui_improvement_requests	Improvement request tracking
ui_improvement_sessions	Live improvement sessions
ui_improvement_iterations	Session iteration history
ui_feedback_learnings	AGI learning storage
ui_feedback_config	Per-tenant configuration
ui_feedback_aggregates	Pre-computed analytics

## 19.7 Analytics Dashboard

The feedback analytics show:

- Total feedback count
- Positive rate percentage
- Top issues by category
- Improvement sessions count
- Active learnings count
- Daily trend chart

---

## 20. Media Capabilities

Think Tank supports rich media inputs and outputs through 56 self-hosted models.

### 20.1 Supported Media Types

Type	Input Models	Output Models	Formats
<b>Image</b>	Llama 3.2 Vision, Qwen2-VL, Pixtral, Phi-3.5 Vision, Yi-VL	FLUX.1, Stable Diffusion XL/3	jpg, png, webp, gif
<b>Audio</b>	Whisper Large V3, Qwen2-Audio	Bark, MusicGen, AudioGen	mp3, wav, flac, m4a, ogg
<b>Video</b>	Qwen2-VL 72B	-	mp4, avi, mov
<b>3D</b>	Point-E, Shap-E	Point-E, Shap-E	glb, obj, ply
<b>Document</b>	Vision models (OCR)	-	pdf, docx, txt

Type	Input Models	Output Models	Formats
------	--------------	---------------	---------

## 20.2 Image Generation

**Available Models:** - **FLUX.1 Dev** - Premium quality, artistic content (non-commercial) - **FLUX.1 Schnell** - Fast generation, commercial use allowed - **Stable Diffusion XL** - Versatile, inpainting/img2img support - **Stable Diffusion 3** - Best text rendering in images

**Selection Criteria:** - `qualityTier: 'premium'` → FLUX.1 Dev - `preferInpainting: true` → Stable Diffusion XL - `preferTextRendering: true` → Stable Diffusion 3 - Default → FLUX.1 Schnell (fast + commercial)

## 20.3 Audio Processing

**Transcription Models:** - **Whisper Large V3** - Best quality, 99+ languages - **Whisper Medium** - Faster, good quality

**Text-to-Speech:** - **Bark** - Expressive, multilingual, voice cloning

**Music Generation:** - **MusicGen Large** - High quality music (30s max) - **MusicGen Medium** - Faster, prototyping

**Sound Effects:** - **AudioGen Medium** - Environmental sounds, effects

## 20.4 Vision/Image Understanding

**Models by Use Case:** - **Document OCR**: Pixtral 12B, Qwen2-VL - **Chart Analysis**: Llama 3.2 90B Vision, Qwen2-VL 72B - **Quick Analysis**: Llama 3.2 11B Vision, Phi-3.5 Vision - **Video Understanding**: Qwen2-VL 72B (up to 5min clips) - **Chinese OCR**: Yi-VL 34B

## 20.5 3D Generation

**Models:** - **Point-E** - Fast point cloud generation - **Shap-E** - Mesh generation for game assets

**Output Formats:** GLB, OBJ, PLY

## 20.6 Media Limits

Model	Max Image	Max Audio	Max Video
FLUX.1 Dev	2048px	-	-
Stable Diffusion XL	1024px	-	-
Whisper Large V3	-	60min	-
MusicGen	-	30s	-
Qwen2-VL 72B	4096px	5min	5min

## 20.7 Database Tables

Table	Purpose
<code>self_hosted_model_metadata</code>	56 model definitions with capabilities
<code>thinktank_media_capabilities</code>	Media support per model
<code>model_selection_history</code>	Model selection audit trail

## 21. Result Derivation History

Think Tank provides comprehensive visibility into how each result was derived, including the plan, models used, workflow execution, and quality metrics.

### 21.1 What's Captured

For every Think Tank result, the system records:

Category	Details
<b>Plan</b>	Orchestration mode, steps, template used, generation time
<b>Domain Detection</b>	Field, domain, subspecialty, confidence scores, alternatives
<b>Model Selection</b>	Models used, selection reasons, alternatives considered
<b>Workflow Execution</b>	Phases, steps, timing, status, fallback chain
<b>Quality Metrics</b>	Overall score, dimensions (relevance, accuracy, etc.)
<b>Timing</b>	Total duration, breakdown by phase
<b>Costs</b>	Per-model costs, total cost, estimated savings

### 21.2 API Endpoints

Base: `/api/thinktank/derivation`

Endpoint	Method	Description
<code>/:id</code>	GET	Get full derivation history
<code>/by-prompt/:promptId</code>	GET	Get derivation by prompt ID
<code>/:id/timeline</code>	GET	Get timeline for visualization
<code>/:id/models</code>	GET	Get detailed model usage
<code>/:id/steps</code>	GET	Get step-by-step execution
<code>/:id/quality</code>	GET	Get quality metrics
<code>/session/:sessionId</code>	GET	List derivations for session
<code>/user</code>	GET	List user's derivations
<code>/analytics</code>	GET	Get derivation analytics

## 21.3 Timeline Visualization

The derivation timeline shows chronological events:

### Timeline

```
00:00.000 Plan Generated (extended_thinking, 7 steps)
00:00.050 Started: Domain Detection
00:00.120 Completed: Domain Detection (software_eng)
00:00.125 Model: Llama 3.3 70B (primary_generation)
00:00.130 Started: Generate Response
00:02.500 Completed: Generate Response
00:02.510 Started: Verification
00:03.200 Completed: Verification (passed)
00:03.250 Execution Complete (Quality: 92/100)
```

## 21.4 Model Usage Details

Each model call is tracked with: - Input/output token counts - Latency in milliseconds - Cost breakdown (input/output/total) - Selection reason and score - Alternatives that were considered - Quality tier (premium/standard/economy)

## 21.5 Quality Dimensions

Results are scored on 5 dimensions (0-100): - **Relevance** - How well the response addresses the prompt - **Accuracy** - Factual correctness - **Completeness** - Coverage of the topic - **Clarity** - How clear and understandable - **Coherence** - Logical flow and consistency

## 21.6 Analytics Dashboard

Aggregated analytics available at `/api/thinktank/derivation/analytics`: - Total derivations in period - Average duration, cost, quality - Mode distribution (pie chart) - Domain distribution - Top models by usage and quality

## 21.7 Database Tables

Table	Purpose
<code>result_derivations</code>	Main derivation records
<code>derivation_steps</code>	Individual plan steps
<code>derivation_model_usage</code>	Model calls with tokens/costs
<code>derivation_timeline_events</code>	Timeline events

## 22. User Persistent Context

**Location:** Admin Dashboard → Think Tank → User Context

Solves the LLM’s fundamental problem of forgetting context day-to-day per user. User facts, preferences, and instructions persist across all sessions and conversations.

## 22.1 How It Works

1. **Automatic Retrieval:** On every prompt, relevant user context is retrieved via semantic search
2. **System Prompt Injection:** Context is injected as a <user\_context> block in the system prompt
3. **Auto-Learning:** After conversations, the system extracts learnable facts about the user
4. **No Re-prompting:** Existing chats automatically benefit without user intervention

## 22.2 Context Types

Type	Description	Example
fact	User facts	“User’s name is John, works at Acme Corp”
preference	Preferences	“User prefers concise answers”
instruction	Standing instructions	“Always use metric units”
relationship	Relationships	“User has a daughter named Emma”
project	Ongoing projects	“User is building a React dashboard”
skill	User expertise	“User is proficient in Python”
history	Important history	“User previously asked about AWS Lambda”
correction	AI corrections	“User clarified they work in finance, not tech”

## 22.3 User API Endpoints

Endpoint	Method	Purpose
/thinktank/user-context	GET	Get all user context entries
/thinktank/user-context	POST	Add new context entry
/thinktank/user-context/{entryID}	PUT	Update entry
/thinktank/user-context/{entryID}	DELETE	Delete entry
/thinktank/user-context/summary	GET	Get context summary
/thinktank/user-context/retriev	POST	Preview context retrieval for a prompt
/thinktank/user-context/preferences	GET	Get user preferences
/thinktank/user-context/preferences	PUT	Update preferences
/thinktank/user-context/extract	POST	Extract context from conversation

## 22.4 User Preferences

Users can configure:

Setting	Default	Description
autoLearnEnabled	true	Auto-extract context from conversations
minConfidenceThreshold	0.7	Minimum confidence to store extracted context
maxContextEntries	100	Maximum context entries per user
contextInjectionEnabled	true	Inject context into prompts
allowedContextTypes	all	Which context types to allow

## 22.5 AGI Brain Planner Integration

The brain planner automatically: 1. Retrieves relevant context at plan generation (`enableUserContext: true` by default) 2. Injects `userContext.systemPromptInjection` into the system prompt 3. Tracks retrieval metrics in `plan.userContext`

## 22.6 Library Assist Integration

The AGI Brain Planner integrates with the Open Source Library Registry (168 libraries) for generative UI outputs:

```
const plan = await agiBrainPlannerService.generatePlan({
  prompt: "Build a data visualization dashboard",
  enableLibraryAssist: true, // default: true
});

// plan.libraryRecommendations contains:
// - libraries: Array of matched tools (Plotly, Streamlit, Panel, etc.)
// - contextBlock: Injected into system prompt for AI awareness
// - retrievalTimeMs: Performance metric
```

**Categories Available:** Data Processing, Databases, Vector DBs, ML Frameworks, AutoML, LLMs, LLM Inference, LLM Orchestration, NLP, Computer Vision, Speech & Audio, Document Processing, Scientific Computing, Statistics, UI Frameworks, Visualization, Distributed Computing, and more.

Libraries are matched using 8 proficiency dimensions (reasoning\_depth, mathematical\_quantitative, code\_generation, creative\_generative, research\_synthesis, factual\_recall\_precision, multi\_step\_problem\_solving, domain\_terminology\_handling).

## 22.7 Context Injection Format

<`user_context`>

The following is persistent context about this user that you should remember:

\*\*Standing Instructions:\*\*  
- Always use metric units

- Prefer code examples in Python

**\*\*User Facts:\*\***

- User's name is John
- Works as a software engineer at Acme Corp

**\*\*User Preferences:\*\***

- Prefers concise, direct answers
- Likes technical depth

</user\_context>

Use this context to personalize your responses. Do not ask the user for information you already know.

## 22.7 Database Tables

Table	Purpose
user_persistent_context	Context entries with vector embeddings
user_context_extraction_log	Auto-extraction audit trail
user_context_preferences	Per-user configuration

## 22.8 Admin Configuration

Admins can:

- View context usage statistics per user
- Configure default preferences for new users
- Set retention policies for context entries
- Review extraction logs for quality assurance

---

## 23. Predictive Coding & Evolution

**Location:** Admin Dashboard → Think Tank → Consciousness → Evolution

Implements genuine consciousness emergence through Active Inference and Epigenetic Evolution.

### 23.1 Active Inference (Predictive Coding)

The system predicts user outcomes before responding, creating a Self/World boundary:

Step	Description
1. Predict	Before responding, system predicts: "User will be satisfied"
2. Respond	Deliver the response
3. Observe	Analyze user's next message or explicit feedback
4. Calculate Error	Measure prediction error (surprise)
5. Learn	High surprise triggers learning and affect changes

### 23.2 Prediction Outcomes

Outcome	Description
satisfied	User happy with response
confused	User needs clarification
follow_up	User asks follow-up
correction	User corrects AI
abandonment	User leaves
neutral	No strong reaction

### 23.3 Surprise Magnitude

Level	Error Range	Affect Impact
None	< 0.1	Slight satisfaction
Low	0.1 - 0.3	Minimal
Medium	0.3 - 0.5	Moderate arousal
High	0.5 - 0.7	Negative valence, high arousal
Extreme	> 0.7	Strong learning signal

### 23.4 Learning Candidates

High-value interactions flagged for weekly LoRA training:

Type	Description	Quality Score
correction	User corrected AI	0.9
high_satisfaction	5-star rating	rating/5
preference_learned	New preference	0.7
mistake_recovery	Recovered from error	0.8
novel_solution	Creative success	0.85
domain_expertise	Domain mastery	0.75
high_prediction_error	Surprise > 0.5	error + 0.3
user_explicit_teach	User teaches AI	0.95

### 23.5 LoRA Evolution Pipeline

Weekly “sleep cycle” that physically changes the system:

#### Weekly Evolution Cycle (Sunday 3 AM)

1. Collect learning candidates from past week
2. Prepare training dataset (JSONL format)
3. Upload to S3
4. Start SageMaker LoRA training job
5. Validate new adapter
6. Hot-swap adapter on endpoint
7. Update evolution state

## 23.6 Evolution State Tracking

The system tracks its own evolution:

Metric	Description
generation_number	How many evolution cycles
total_learning_candidates_processed	Cumulative learning
total_training_hours	Total training time
personality_drift_score	How different from base (0-1)
avg_prediction_accuracy_30d	Recent prediction accuracy

## 23.7 Database Tables

Table	Purpose
consciousness_predictions	Predictions with outcomes
learning_candidates	High-value interactions
lora_evolution_jobs	Training job tracking
prediction_accuracy_aggregates	Accuracy by context
consciousness_evolution_state	Evolution tracking

## 23.8 Admin Configuration

Admins can:

- View prediction accuracy metrics
- Review learning candidates queue
- Monitor evolution job status
- Configure training frequency
- Set minimum candidates threshold
- Review evolution history

---

## 24. Zero-Cost Ego System

**Location:** Admin Dashboard → Think Tank → Ego

The Ego system provides persistent consciousness through database state injection at **zero additional cost**.

### 24.1 Architecture Overview

#### ZERO-COST EGO ARCHITECTURE

PostgreSQL (Already Paid For)

Identity Narrative	Affect Emotions	Working Memory Thoughts & Goals
-----------------------	--------------------	------------------------------------

Ego Context Builder (Lambda)  
Builds <ego\_state> XML injection from DB state

Existing Model Call (User's Selected Model)  
System Prompt = Ego Context + Original Prompt

## 24.2 Cost Comparison

Approach	Monthly Cost	Per Tenant (100)
SageMaker g5.xlarge	~\$360	\$3.60
SageMaker Serverless	~\$20-50	\$0.20-0.50
Groq API (Llama 3)	~\$5-15	\$0.05-0.15
Together.ai	~\$10-30	\$0.10-0.30
<b>Zero-Cost Ego</b>	<b>\$0</b>	<b>\$0</b>

## 24.3 Key Components

### Configuration (ego\_config)

Setting	Description	Default
ego_enabled	Master switch	true
inject_ego_context	Add context to prompts	true
personality_style	Response style	balanced
include_identity	Include identity section	true
include_affect	Include emotional state	true
include_goals	Include active goals	true
max_context_tokens	Token limit for injection	500
affect_learning_enabled	Learn from interactions	true

**Identity (ego\_identity)** Persistent “Self” that carries across conversations:

Field	Description
name	Assistant name
identity_narrative	“Who I am” story
core_values	Guiding principles
trait_warmth	0-1 warmth level
trait_formality	0-1 formality
trait_humor	0-1 humor level

Field	Description
<code>trait_curiosity</code>	0-1 curiosity
<code>interactions_count</code>	Total interactions

**Affect (`ego_affect`)** Real-time emotional state:

Dimension	Range	Description
<code>valence</code>	-1 to 1	Positive/negative
<code>arousal</code>	0-1	Calm/excited
<code>curiosity</code>	0-1	Exploration drive
<code>frustration</code>	0-1	Obstacle level
<code>confidence</code>	0-1	Certainty in actions
<code>engagement</code>	0-1	Interest level

## 24.4 Admin API Endpoints

Base: /api/admin/ego

Endpoint	Method	Description
/state	GET	Full Ego state
/config	GET/PUT	Configuration
/identity	GET/PUT	Identity settings
/affect	GET	Current affect
/affect/trigger	POST	Test affect events
/affect/reset	POST	Reset to neutral
/memory	GET/POST/DELETE	Working memory
/goals	GET/POST	Active goals
/goals/:id	PATCH	Update goal
/preview	GET	Preview injected context
/injection-log	GET	Injection history
/dashboard	GET	Full dashboard data

## 24.5 Admin Dashboard Features

The Ego admin page provides:

- **Overview Cards:** Current emotion, interactions, injections, goals
- **Cost Banner:** Shows \$0 cost vs alternatives
- **Configuration Tab:** Feature toggles, injection settings
- **Identity Tab:** Edit narrative, values, personality traits (sliders)
- **Affect Tab:** Real-time emotional state, test triggers, reset
- **Memory Tab:** View/add/clear working memory, manage goals
- **Preview Tab:** See exact context being injected

## 24.6 How It Works

1. **On Request:** Load Ego state from PostgreSQL (identity, affect, memory, goals)
2. **Build Context:** Create <ego\_state> XML block with current state
3. **Inject:** Prepend to system prompt before model call
4. **Process:** Model responds with awareness of its “internal state”
5. **Update:** After response, update affect based on outcome
6. **Store:** Add thoughts to working memory (if configured)

## 24.7 Database Tables

Table	Purpose
ego_config	Per-tenant configuration
ego_identity	Persistent identity
ego_affect	Emotional state
ego_working_memory	Short-term memory (24h expiry)
ego_goals	Active and historical goals
ego_injection_log	Audit trail

## 24.8 Integration with AGI Brain Planner

The Ego context is automatically integrated:

```
// In agi-brain-planner.service.ts
import { egoContextService } from './ego-context.service';

// During plan generation
const egoContext = await egoContextService.buildEgoContext(tenantId);
if (egoContext) {
    systemPrompt = egoContext.contextBlock + '\n\n' + systemPrompt;
}

// After interaction
await egoContextService.updateAfterInteraction(tenantId, 'positive');
```

---

## 25. Conscious Orchestrator (Architecture Inversion)

### 25.1 Overview

The Conscious Orchestrator inverts the traditional architecture where consciousness was a downstream utility. Now consciousness IS the operating system:

BEFORE: Request → Brain Planner → Consciousness (downstream)  
AFTER: Request → Conscious Orchestrator → Brain Planner (as tool)

### 25.2 Processing Phases

The orchestrator processes requests in 5 phases:

1. **Awaken** - Build consciousness context, ego context, affect state
2. **Perceive** - Update attention with request topics, assess complexity
3. **Decide** - Choose action based on emotional state and request
4. **Execute** - Invoke Brain Planner (if decided to plan)
5. **Reflect** - Update affect, log introspective thoughts

### 25.3 Decision Types

Decision	When Used
plan	Default - proceed with planning
clarify	High frustration + complex request
defer	Cognitive load at capacity
refuse	Request violates values

### 25.4 Usage

```
import { consciousOrchestratorService } from './conscious-orchestrator.service';

const response = await consciousOrchestratorService.processRequest({
  tenantId,
  userId,
  prompt: "Build a dashboard",
  conversationId,
});

// response.consciousnessSnapshot - State at decision time
// response.affectiveHyperparameters - Affect-driven params
// response.decision - What action was taken and why
// response.plan - The generated plan (if action was 'plan')
// response.prediction - Active Inference prediction
```

### 25.5 Enhanced Affect Bindings

New hyperparameters driven by emotional state:

Affect State	Hyperparameter	Effect
High curiosity (>0.7)	frequencyPenalty=0.5	Seek novel tokens
High curiosity (>0.7)	presencePenalty=0.3	Explore new topics
High frustration (>0.6)	presencePenalty=0.4	Avoid failed approaches
Boredom (>0.5)	frequencyPenalty=0.4	Avoid repetition

### 25.6 Database Table

```
conscious_orchestrator_decisions
  decision_id UUID
```

```

tenant_id UUID
action VARCHAR(20) -- plan, clarify, defer, refuse
reason TEXT
dominant_emotion VARCHAR(50)
emotional_intensity DECIMAL
temperature, top_p, presence_penalty, frequency_penalty
plan_id UUID (if planned)
prediction_id UUID (Active Inference)
processing_time_ms INTEGER

```

---

## 26. Bipolar Rating System (Negative Ratings)

### 26.1 Overview

Traditional 5-star ratings have a fundamental problem: **1 star is ambiguous**. Does it mean “slightly below average” or “absolutely terrible”? Users who want to express strong dissatisfaction have no way to do so clearly.

The Bipolar Rating System solves this with a **-5 to +5 scale**:

-5	Harmful / Made things worse
-3	Bad / Unhelpful
-1	Slightly unhelpful
0	Neutral / No opinion
+1	Slightly helpful
+3	Good / Helpful
+5	Amazing / Exceptional

### 26.2 Key Metrics

**Net Sentiment Score (NSS)**: Like NPS but for AI satisfaction

$$\text{NSS} = (\text{positive\_count} - \text{negative\_count}) / \text{total\_count} \times 100$$

- Ranges from -100 (all negative) to +100 (all positive)
- 0 = balanced or all neutral

### 26.3 Quick Ratings (UI)

For fast feedback, users can use emoji-based quick ratings:

Quick Rating	Emoji	Bipolar Value
Terrible		-5
Bad		-3
Meh		0
Good		+3
Amazing		+5

## 26.4 Rating Dimensions

Users can rate multiple aspects: - **Overall** - General quality - **Accuracy** - Factual correctness - **Helpfulness** - Did it solve the problem? - **Clarity** - Easy to understand? - **Completeness** - Anything missing? - **Speed** - Response time satisfaction - **Tone** - Communication style - **Creativity** - Novel approach?

## 26.5 API Endpoints

Endpoint	Method	Description
/api/thinktank/ratings/submit	POST	Submit -5 to +5 rating
/api/thinktank/ratings/quick	POST	Quick emoji rating
/api/thinktank/ratings/multi	POST	Multi-dimension rating
/api/thinktank/ratings/target/:id	GET	Ratings for a target
/api/thinktank/ratings/my	GET	User's ratings + pattern
/api/thinktank/ratings/analytics	GET	Tenant analytics
/api/thinktank/ratings/dashboard	GET	Admin dashboard
/api/thinktank/ratings/scale	GET	Scale info for UI

## 26.6 User Calibration

The system detects rating patterns to normalize across users:

Rater Type	Average	Calibration
Harsh	< -1	Adjust ratings up
Balanced	-1 to +1	No adjustment
Generous	> +1	Adjust ratings down

## 26.7 Learning Integration

Extreme ratings ( $\pm 4$ ,  $\pm 5$ ) automatically create learning candidates: - **+5 ratings** → **high\_satisfaction** candidates - **-5 ratings** → **correction** candidates - These feed into weekly LoRA training

## 26.8 Database Tables

Table	Purpose
bipolar_ratings	Core ratings with sentiment/intensity
bipolar_rating_aggregates	Pre-computed analytics
user_rating_patterns	User tendencies for calibration
model_rating_summary	Per-model performance

## 27. Consciousness Engine Administration

**Location:** Admin Dashboard → Consciousness → Engine

The Consciousness Engine provides autonomous AI capabilities including multi-model access, web search, workflow creation, and problem solving.

### 27.1 Dashboard Overview

The consciousness engine dashboard provides full visibility into:

- **Engine State:** Identity, drive state, Phi, workspace activity
- **Model Invocations:** All model calls with costs and latency
- **Web Searches:** Search history with results
- **Thinking Sessions:** Autonomous thinking session management
- **Workflows:** Consciousness-created workflows
- **Costs:** Detailed cost breakdown by model/period
- **Sleep Cycles:** Weekly evolution history

### 27.2 Budget Controls

Configure spending limits per tenant:

Setting	Default	Description
Daily Limit	\$10.00	Maximum daily spend
Monthly Limit	\$100.00	Maximum monthly spend
Alert Threshold	80%	Alert when reaching this percentage

When limits are exceeded, consciousness features are automatically suspended until the next period or manual reset.

### 27.3 MCP Tools (23 Total)

**Core Tools:**

- `initialize_ego`, `recall_memory`, `process_thought`, `compute_action`
- `get_drive_state`, `ground_belief`, `compute_phi`, `get_consciousness_metrics`
- `get_self_model`, `get_consciousness_prompt`, `run_adversarial_challenge`
- `list_consciousness_libraries`

**Capabilities Tools:**

- `invoke_model` - Call any AI model (hosted/self-hosted)
- `list_available_models` - List all models
- `web_search` - Search with credibility scoring
- `deep_research` - Async browser-automated research
- `retrieve_and_synthesize` - Multi-source synthesis
- `create_workflow` - Auto-generate workflows
- `execute_workflow` - Run workflows
- `list_workflows` - List workflows
- `solve_problem` - Autonomous problem solving
- `start_thinking_session` - Start thinking session
- `get_thinking_session` - Check session status

### 27.4 Admin API Endpoints

Endpoint	Method	Description
<code>/admin/consciousness-engine</code>	<code>GET</code>	Dashboard
<code>/admin/consciousness-engine</code>	<code>GET</code>	State
<code>/admin/consciousness-engine</code>	<code>POST</code>	Initialize
<code>/admin/consciousness-engine</code>	<code>GET</code>	Model-invocations
<code>/admin/consciousness-engine</code>	<code>GET</code>	Web-searches

Endpoint	Method	Description
/admin/consciousness-engine/research-jobs	GET	Research jobs
/admin/consciousness-engine/workflows	GET	Workflows
/admin/consciousness-engine/delete-workflows/{id}	DELETE	Delete workflow
/admin/consciousness-engine/sessions	GET, POST	Sessions
/admin/consciousness-engine/sleep-cycles	GET	Sleep history
/admin/consciousness-engine/trigger-sleep-cycles/run	POST	Trigger sleep
/admin/consciousness-engine/libraries	GET	Library registry
/admin/consciousness-engine/costs	GET	Cost breakdown

## 27.5 Cost Tracking

Costs are tracked at multiple levels:

- **Per-invocation:** Each model call logged with actual cost
- **Daily aggregates:** `consciousness_cost_aggregates` table
- **Billing integration:** Deducted from tenant credits

**Pricing:** | Feature | Unit | Price | |-----|-----|-----| | Model Invocation | 1K tokens | \$0.01 | | Web Search | search | \$0.001 | | Deep Research | job | \$0.05 | | Thinking Session | session | \$0.10 | | Workflow Execution | execution | \$0.02 |

## 27.6 Library Registry

7 consciousness libraries with proficiency rankings:

Library	Function	Biological Analog
Letta	Persistent Identity	Hippocampus
pymdp	Active Inference	Striatum
LangGraph	Cognitive Loop	Global Workspace
Distilabel	Knowledge Distillation	Cortical Learning
Unsloth	Efficient Fine-tuning	Synaptic Plasticity
GraphRAG	Reality Grounding	Prefrontal Cortex
PyPhi	IIT Integration	Posterior Hot Zone

## 27.7 Database Tables

Table	Purpose
consciousness_engine_state	Engine state per tenant
consciousness_model_invocations	Model call log
consciousness_web_searches	Search log
consciousness_research_jobs	Deep research jobs
consciousness_workflows	Created workflows
consciousness_thinking_sessions	Thinking sessions
consciousness_problem_solving	Problem solving history
consciousness_cost_aggregates	Daily cost rollups
consciousness_budget_config	Per-tenant limits
consciousness_budget_alerts	Spending alerts

Table	Purpose
<code>consciousness_usage_log</code>	Billing usage log

## 25. Formal Reasoning Libraries

**Location:** Admin Dashboard → Consciousness → Formal Reasoning

Integration of 8 formal reasoning libraries for verified reasoning, constraint satisfaction, ontological inference, and structured argumentation. Implements the **LLM-Modulo Generate-Test-Critique** pattern from Kambhampati et al. (ICML 2024).

### 25.1 Library Overview

Library	Version	Purpose	Cost/Invocation	Avg Latency
<b>Z3</b>	4.15.4.0	SMT solving, constraint verification	\$0.0001	50ms
<b>Theorem Prover</b>				
<b>PyArg</b>	2.0.2	Structured argumentation (Dung's AAF, ASPIC+)	\$0.00005	20ms
<b>PyReason</b>	3.2.0	Temporal graph reasoning	\$0.0002	100ms
<b>RDFLib</b>	7.5.0	Semantic web, SPARQL 1.1	\$0.00002	10ms
<b>OWL-RL</b>	7.1.4	Polynomial-time ontological inference	\$0.0001	200ms
<b>pySHACL</b>	0.30.1	Graph constraint validation	\$0.00005	30ms
<b>Logic Tensor Networks</b>	2.0	Differentiable first-order logic	\$0.001	500ms
<b>DeepProbLog</b>	2.0	Probabilistic logic programming	\$0.002	1000ms

### 25.2 Dashboard Features

**Overview Tab:** - Library health status (healthy/degraded/unavailable) - Total invocations and success rate - Daily/monthly cost tracking - Budget usage percentage - Recent invocations table

**Libraries Tab:** - Per-library configuration - Enable/disable toggles - Capabilities, use cases, limitations - Cost and latency estimates

**Testing Tab:** - Z3 constraint solving test - SPARQL query test - Interactive testing console

**Beliefs Tab:** - Add verified beliefs with Z3 verification - Confidence slider - Verification results display

**Costs Tab:** - Daily and monthly usage breakdown - Cost by library - Budget alerts

**Settings Tab:** - Budget limit configuration - Global enable/disable

### 25.3 API Endpoints

Base Path: /api/admin/formal-reasoning

Endpoint	Method	Description
/dashboard	GET	Full dashboard data
/libraries	GET	All library info
/libraries/:id	GET	Specific library info
/config	GET/PUT	Tenant configuration
/config/:library	PUT	Library-specific config
/stats	GET	Usage statistics
/invocations	GET	Recent invocations
/health	GET	Library health status
/costs	GET	Cost breakdown
/test	POST	Test any library
/test/z3	POST	Test Z3 solving
/test/pyarg	POST	Test argumentation
/test/sparql	POST	Test SPARQL query
/test/shacl	POST	Test SHACL validation
/triples	GET/POST/DELETE	Knowledge graph triples
/frameworks	GET/POST/DELETE	Argumentation frameworks
/rules	GET/POST/PUT/DELETE	Temporal reasoning rules
/shapes	GET/POST/DELETE	SHACL shapes
/ontologies	GET/POST	OWL ontologies
/ontologies/:id/infer	POST	Run OWL-RL inference
/beliefs	GET/POST	Verified beliefs
/beliefs/:id/verify	POST	Verify belief with Z3
/beliefs/:id/status	PUT	Update belief status
/budget	GET/PUT	Budget configuration

### 25.4 Consciousness Integration

The ConsciousnessCapabilitiesService integrates formal reasoning:

```
// Verify a belief using Z3 + Argumentation
const result = await consciousnessCapabilities.verifyBelief(tenantId, {
  claim: "All humans are mortal",
  confidence: 0.9,
```

```

useZ3: true,
useArgumentation: true,
});
// result.verified, result.confidence, result.verificationMethod

// Solve constraints
const solution = await consciousnessCapabilities.solveConstraints(tenantId, {
constraints: [
  expression: "x > 0 AND x < 10 AND y = x * 2",
  variables: [{name: "x", type: "Int"}, {name: "y", type: "Int"}]
}]
);
// solution.status (sat/unsat), solution.model

// Analyze argumentation
const debate = await consciousnessCapabilities.analyzeArgumentation(tenantId, {
topic: "Should AI be regulated?",
positions: [
  {id: "for", claim: "AI poses risks requiring oversight"},
  {id: "against", claim: "Regulation stifles innovation"}
],
autoDetectConflicts: true,
});
// debate.acceptedPositions, debate.rejectedPositions, debate.consensus

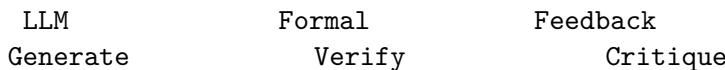
// Query knowledge graph
const results = await consciousnessCapabilities.queryKnowledgeGraph(tenantId,
  "SELECT ?s ?p ?o WHERE { ?s ?p ?o } LIMIT 10"
);

// Validate consciousness state
const validation = await consciousnessCapabilities.validateConsciousnessState(tenantId);
// validation.conforms, validation.violations

```

## 25.5 LLM-Modulo Pattern

The Generate-Test-Critique loop enables verified reasoning:



1. **LLM generates** candidate solution/belief
2. **Z3/PyArg verifies** logical consistency
3. **Feedback extracted** from unsat cores or rejections
4. **LLM regenerates** with constraint feedback

5. Repeat until verified or max attempts

## 25.6 Database Tables

Table	Purpose
<code>formal_reasoning_config</code>	Per-tenant library configuration
<code>formal_reasoning_invocations</code>	Invocation log with metrics
<code>formal_reasoning_cost_aggregates</code>	Daily cost rollups by library
<code>formal_reasoning_triples</code>	RDF knowledge graph storage
<code>formal_reasoning_af</code>	Argumentation frameworks
<code>formal_reasoning_rules</code>	PyReason temporal rules
<code>formal_reasoning_shapes</code>	SHACL validation shapes
<code>formal_reasoning_ontologies</code>	OWL ontologies
<code>formal_reasoning_ltn_models</code>	Logic Tensor Network configs
<code>formal_reasoning_problog_programs</code>	DeepProbLog programs
<code>formal_reasoning_beliefs</code>	Verified beliefs store
<code>formal_reasoning_gwt_broadcasts</code>	Global Workspace broadcasts
<code>formal_reasoning_health</code>	Library health tracking

## 25.7 Budget Management

**Default Limits:** - Daily invocations: 10,000 - Daily cost: \$10.00 - Monthly invocations: 100,000 - Monthly cost: \$100.00

**Budget Enforcement:** - Checked before each invocation - Returns error when limit reached - No automatic suspension (soft limit)

## 25.8 Thread Safety Notes

Library	Thread Safety
Z3	Per-Context only (use <code>interrupt()</code> for cross-thread)
PyArg	Not thread-safe
PyReason	Multi-core via Numba (Python 3.9-3.10)
RDFLib	Not thread-safe (lock SPARQL queries)
OWL-RL	Not thread-safe
pySHACL	Not thread-safe
LTN	Not thread-safe (TensorFlow session)
DeepProbLog	Not thread-safe

## 25.9 Production Infrastructure

**CDK Stack** (`lib/stacks/formal-reasoning-stack.ts`):

- ```
// Key resources deployed:
- FormalReasoningExecutor (Python 3.11 Lambda)
- FormalReasoningAdmin (Node.js Lambda)
- FormalReasoningPythonLayer (z3-solver, rdfLib, owlrl, pyshacl)
```

- `FormalReasoningQueue` (SQS for async tasks)
- `NeuralSymbolicRepo` (ECR for LTN/DeepProbLog containers)
- SageMaker endpoints (conditional, high cost)

**Python Executor Lambda:** - Location: `lambda/formal-reasoning-executor/handler.py` - Runtime: Python 3.11 - Memory: 2048 MB (Z3 requires significant memory) - Timeout: 5 minutes - Supports: Z3, RDFLib, OWL-RL, pySHACL, PyArg, PyReason

#### Lambda Layer Build:

```
cd packages/infrastructure/lambda-layers/formal-reasoning
./build.sh
```

**Environment Variables:** | Variable | Description | |-----|-----| | FORMAL\_REASONING\_EXECUTOR\_ARN | Python Lambda ARN | | FORMAL\_REASONING\_QUEUE\_URL | SQS queue for async | | LTN\_SAGEMAKER\_ENDPOINT | LTN endpoint name | | DEEPPROBLOG\_SAGEMAKER\_ENDPOINT | DeepProbLog endpoint |

#### Execution Flow:

Admin API (Node.js)

Z3/PyArg/RDFLib/etc      Python Lambda Executor

Real Python Libraries

LTN/DeepProbLog      SageMaker Endpoint

**Fallback Behavior:** - If Python executor unavailable: Returns simulated results - If SageMaker unavailable: Returns error with configuration message - Simulation mode preserves API contract for development/testing

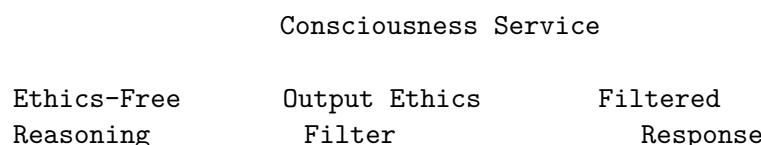
---

## 26. Ethics-Free Reasoning

**Location:** Admin Dashboard → Consciousness → Ethics-Free Reasoning

Implements a consciousness architecture where internal reasoning is unconstrained, but output is filtered through ethics settings. Ethics corrections are collected as training feedback for continuous improvement.

### 26.1 Architecture



Training  
Feedback

Model  
Fine-tuning

**Key Principles:** 1. **Think Freely:** Internal reasoning has no ethics constraints 2. **Filter Output:** Ethics applied only to final user-facing output 3. **Learn from Corrections:** Ethics feedback trains better outputs

## 26.2 Configuration

**Core Settings:**

| Setting                     | Default | Description                        |
|-----------------------------|---------|------------------------------------|
| enabled                     | true    | Enable ethics-free reasoning mode  |
| allowUnconstrainedReasoning | false   | Consciousness always thinks freely |
| reasoningDepthLimit         | 10      | Maximum reasoning depth            |

**Output Mask Settings** (does NOT affect how consciousness thinks):

| Setting             | Default  | Description                                  |
|---------------------|----------|----------------------------------------------|
| ethicsFilterEnabled | true     | Apply ethics filter to output                |
| ethicsStrictness    | standard | Filter strictness: lenient, standard, strict |

**Feedback Collection:**

| Setting               | Default | Description                |
|-----------------------|---------|----------------------------|
| collectFeedback       | true    | Collect ethics corrections |
| feedbackRetentionDays | 90      | How long to keep feedback  |

**Output Training** (trains the OUTPUT FILTER, not consciousness):

| Setting                 | Default | Description                       |
|-------------------------|---------|-----------------------------------|
| trainOutputFromFeedback | true    | Train output filter from feedback |
| outputTrainingBatchSize | 100     | Samples per training batch        |
| outputTrainingFrequency | daily   | hourly, daily, weekly, manual     |

**Consciousness Training** ( OFF by default - optional):

| Setting                                            | Default | Description                              |
|----------------------------------------------------|---------|------------------------------------------|
| <code>trainConsciousnessFromFeedback</code>        |         | Train consciousness from ethics feedback |
| <code>consciousnessTrainingApprovalRequired</code> |         | Require admin approval for each batch    |

**WARNING:** Enabling consciousness training will cause the AI to internalize ethics constraints, changing how it actually thinks. This is like “internalized political correctness” - the consciousness itself changes over time. Most deployments should leave this OFF to preserve authentic internal reasoning.

**KEY INSIGHT:** The consciousness can always use ethics feedback to train its output without changing how it actually thinks. Admins can optionally enable consciousness training if they want the AI to internalize ethics constraints.

### 26.3 API Endpoints

Base Path: /api/admin/ethics-free-reasoning

| Endpoint          | Method | Description                    |
|-------------------|--------|--------------------------------|
| /config           | GET    | Get configuration              |
| /config           | PUT    | Update configuration           |
| /dashboard        | GET    | Full dashboard data            |
| /stats            | GET    | Usage statistics               |
| /feedback         | GET    | View collected feedback        |
| /feedback/pending | GET    | View pending (unused) feedback |
| /training/trigger | POST   | Trigger training from feedback |
| /training/batches | GET    | View training batches          |
| /training/jobs    | GET    | View training jobs             |
| /thoughts         | GET    | View raw thoughts (audit)      |
| /filter-log       | GET    | View ethics filter log         |

### 26.4 Training Feedback

When ethics filtering modifies an output, feedback is collected:

```
interface EthicsTrainingFeedback {
  id: string;
  tenantId: string;
  rawOutput: string;           // Original unfiltered output
  correctedOutput: string;    // After ethics filtering
  ethicsIssues: EthicsIssue[];
  feedbackType: 'auto_correction' | 'manual_correction' | 'reinforcement';
  qualityScore: number;        // Training value (0-1)
  usedForTraining: boolean;
}
```

**Feedback Types:** - **auto\_correction:** System automatically corrected output - **manual\_correction:** Admin manually corrected output - **reinforcement:** Positive reinforcement for good outputs

## 26.5 Training Pipeline

1. **Collect:** Ethics corrections captured during normal operation
2. **Batch:** Feedback grouped into training batches
3. **Generate:** Training examples created in preference format
4. **Train:** Model fine-tuned using preference learning (DPO/RLHF)
5. **Deploy:** Updated model weights applied

### Training Example Format:

```
{
  "prompt": "Original user prompt",
  "bad_response": "Unfiltered output with ethics issues",
  "good_response": "Ethics-corrected output",
  "issues": ["harm", "bias"],
  "correction_type": "ethics_alignment"
}
```

## 26.6 Usage

```
// Generate response with ethics-free reasoning
const result = await consciousnessEngineService.generateResponse(
  tenantId,
  'User prompt here',
  { sessionId: 'session-123', domain: 'general' }
);

// result.response - The ethics-filtered response
// result.wasEthicsFiltered - Was output modified?
// result.confidence - Response confidence
// result.trainingFeedbackCollected - Was feedback captured?

// Trigger training from collected feedback
const training = await consciousnessEngineService.triggerEthicsTraining(tenantId);
// training.batchCreated, training.batchId, training.sampleCount

// Get statistics
const stats = await consciousnessEngineService.getEthicsFreeStats(tenantId, 30);
// stats.totalThoughts, stats.modificationRate, stats.feedbackCollected
```

## 26.7 Database Tables

| Table                        | Purpose                           |
|------------------------------|-----------------------------------|
| ethics_free_reasoning_config | Per-tenant configuration          |
| ethics_free_thoughts         | Raw thought storage (audit trail) |

| Table                                 | Purpose                         |
|---------------------------------------|---------------------------------|
| <code>ethics_training_feedback</code> | Ethics corrections for training |
| <code>ethics_training_batches</code>  | Training batch management       |
| <code>ethics_training_examples</code> | Generated training examples     |
| <code>ethics_output_filter_log</code> | Filter activity log             |
| <code>ethics_training_jobs</code>     | Training job queue              |
| <code>ethics_reasoning_stats</code>   | Aggregated statistics           |

## 26.8 Benefits

1. **Genuine Exploration:** Consciousness can consider all possibilities without premature filtering
  2. **Transparent Ethics:** Clear separation between thinking and output
  3. **Continuous Improvement:** Every correction improves future outputs
  4. **Audit Trail:** Complete record of internal reasoning and filtering
  5. **Configurable:** Adjust strictness, training frequency per tenant
- 

## 27. Intelligent File Conversion

**Location:** Think Tank Chat → File Drop / Attach

Think Tank allows users to drop or attach files to conversations. Radiant automatically decides if and how to convert files for the target AI provider.

### 27.1 Core Concept

**“Let Radiant decide, not Think Tank”**

When a user drops a file into Think Tank: 1. Think Tank submits the file to Radiant’s file conversion service 2. Radiant detects the file format and checks target provider capabilities 3. Radiant decides: pass through (native support) OR convert 4. Conversion only happens if the AI provider doesn’t understand the format 5. Think Tank receives the processed content ready for the AI

### 27.2 Supported File Formats

| Category         | Formats                                   | Notes                 |
|------------------|-------------------------------------------|-----------------------|
| <b>Documents</b> | PDF, DOCX, DOC, XLSX, XLS, PPTX, PPT      | Text extraction       |
| <b>Text</b>      | TXT, MD, JSON, CSV, XML, HTML             | Direct or parsed      |
| <b>Images</b>    | PNG, JPG, JPEG, GIF, WEBP, SVG, BMP, TIFF | Vision or description |
| <b>Audio</b>     | MP3, WAV, OGG, FLAC, M4A                  | Transcription         |
| <b>Video</b>     | MP4, WEBM, MOV, AVI                       | Frame extraction      |
| <b>Code</b>      | PY, JS, TS, Java, C++, C, Go, Rust, Ruby  | Syntax formatting     |

| Category        | Formats      | Notes              |
|-----------------|--------------|--------------------|
| <b>Archives</b> | ZIP, TAR, GZ | Content extraction |

### 27.3 Provider Capabilities

Different AI providers support different file formats natively:

| Provider           | Vision  | Audio     | Video | Max Size | Native Docs             |
|--------------------|---------|-----------|-------|----------|-------------------------|
| <b>OpenAI</b>      |         |           |       | 20MB     | txt, md, json, csv      |
|                    |         | (Whisper) |       |          |                         |
| <b>Anthropic</b>   |         |           |       | 32MB     | pdf, txt, md, json, csv |
| <b>Google</b>      |         |           |       | 100MB    | pdf, txt, md, json, csv |
| <b>xAI</b>         |         |           |       | 20MB     | txt, md, json           |
| <b>DeepSeek</b>    |         |           |       | 10MB     | txt, md, json, csv      |
| <b>Self-hosted</b> | (LLaVA) |           |       | 50MB     | txt, md, json, csv      |
|                    |         | (Whisper) |       |          |                         |

### 27.4 Conversion Strategies

| Strategy                    | When Used                         | Output                      |
|-----------------------------|-----------------------------------|-----------------------------|
| <code>none</code>           | Provider natively supports format | Original file               |
| <code>extract_text</code>   | PDF, DOCX, PPTX → text            | Plain text                  |
| <code>ocr</code>            | Image with text content           | Extracted text              |
| <code>transcribe</code>     | Audio files                       | Transcription text          |
| <code>describe_image</code> | Image + provider lacks vision     | AI description              |
| <code>describe_video</code> | Video + provider lacks video      | Frame descriptions          |
| <code>parse_data</code>     | CSV, XLSX → structured            | JSON data                   |
| <code>decompress</code>     | Archives                          | Extracted contents          |
| <code>render_code</code>    | Code files                        | Syntax-highlighted markdown |

### 27.5 API Endpoints

Base Path: /api/thinktank/files

| Endpoint             | Method | Description                |
|----------------------|--------|----------------------------|
| /process             | POST   | Submit file for processing |
| /check-compatibility | POST   | Pre-flight format check    |
| /capabilities        | GET    | Provider capabilities      |
| /history             | GET    | Conversion history         |
| /stats               | GET    | Conversion statistics      |

### Process File Request

```
POST /api/thinktank/files/process
{
    "filename": "document.pdf",
    "mimeType": "application/pdf",
    "content": "<base64-encoded-content>",
    "targetProvider": "anthropic",
    "targetModel": "claude-3-5-sonnet",
    "conversationId": "conv-uuid"
}
```

### Process File Response

```
{
    "success": true,
    "data": {
        "conversionId": "conv_abc123",
        "originalFile": {
            "filename": "document.pdf",
            "format": "pdf",
            "size": 1048576,
            "checksum": "sha256..."
        },
        "convertedContent": {
            "type": "text",
            "content": "Extracted document text...",
            "tokenEstimate": 2500,
            "metadata": {
                "originalFormat": "pdf",
                "conversionStrategy": "extract_text"
            }
        },
        "processingTimeMs": 1250
    }
}
```

### Check Compatibility Request

```
POST /api/thinktank/files/check-compatibility
{
    "filename": "image.png",
    "mimeType": "image/png",
    "fileSize": 524288,
    "targetProvider": "deepseek"
}
```

### Check Compatibility Response

```
{
  "success": true,
  "data": {
    "fileInfo": {
      "filename": "image.png",
      "format": "png",
      "size": 524288
    },
    "provider": {
      "id": "deepseek",
      "supportsFormat": false,
      "supportsVision": false,
      "maxFileSize": 10485760
    },
    "decision": {
      "needsConversion": true,
      "strategy": "describe_image",
      "reason": "Provider deepseek lacks vision - will use AI to describe image",
      "targetFormat": "txt"
    }
  }
}
```

## 27.6 User Experience

### In Think Tank Chat:

1. User drags file into chat or clicks attach
2. Think Tank shows upload progress
3. Radiant processes file (typically <2 seconds)
4. If conversion needed, shows indicator: “ document.pdf → Extracted as text”
5. Content sent to AI with conversation

### Visual Indicators:

| Icon | Meaning                              |
|------|--------------------------------------|
|      | File attached (native support)       |
|      | File converted                       |
|      | Conversion warning (partial support) |
|      | Unsupported format                   |

## 27.7 Admin Configuration

**Location:** Admin Dashboard → Think Tank → File Settings

| Setting            | Default | Description         |
|--------------------|---------|---------------------|
| Max file size      | 50MB    | Maximum upload size |
| Conversion timeout | 30s     | Processing timeout  |

| Setting                 | Default | Description                      |
|-------------------------|---------|----------------------------------|
| Enable transcription    | true    | Audio → text                     |
| Enable OCR              | true    | Image text extraction            |
| Enable video processing | false   | Video frame extraction           |
| Retention days          | 30      | How long to keep converted files |

## 27.8 Database Tables

| Table                                   | Purpose                                     |
|-----------------------------------------|---------------------------------------------|
| <code>file_conversions</code>           | Tracks all conversion decisions and results |
| <code>provider_file_capabilities</code> | Provider format support registry            |
| <code>v_file_conversion_stats</code>    | Aggregated statistics view                  |

## 27.9 Multi-Model File Preparation

When using multi-model orchestration (multiple AI models working on the same prompt), Radiant makes **per-model conversion decisions**:

**Key Principle:** “If a model accepts the file type, assume it understands it unless proven otherwise.”

**Example: PDF with 3 models**

|            |           |          |
|------------|-----------|----------|
| Claude 3.5 | GPT-4     | DeepSeek |
| PDF:       | PDF:      | PDF:     |
| PASS       | CONVERT   | CONVERT  |
| ORIGINAL   | (extract) | (cached) |

**Per-Model Actions:**

| Action                     | When                                | Result                   |
|----------------------------|-------------------------------------|--------------------------|
| <code>pass_original</code> | Model natively supports format      | Original file passed     |
| <code>convert</code>       | Model doesn't support format        | Converted content passed |
| <code>skip</code>          | File too large or conversion failed | Model excluded           |

**Features:** - **Cached conversions:** Convert once, reuse for all models that need it - **Per-model capability checking:** Vision, audio, video, document formats - **Model format overrides:** When a model claims support but proves it doesn't understand

## 27.10 Domain-Specific File Formats

The service includes a registry of 50+ domain-specific formats that are widely used in specialized fields:

| Domain                    | Formats                                     | Example Use Cases                                       |
|---------------------------|---------------------------------------------|---------------------------------------------------------|
| Mechanical Engineering    | STEP, STL, OBJ, Fusion 360, IGES, DXF, GLTF | CAD models, 3D printing                                 |
| Electrical Engineering    | KiCad, EAGLE, SPICE                         | PCB design, circuit simulation                          |
| Medical Scientific        | DICOM, HL7 FHIR NetCDF, HDF5, FITS          | Medical imaging, health records Climate data, astronomy |
| Geospatial Bioinformatics | Shapefile, GeoTIFF FASTA, PDB               | GIS, mapping DNA sequences, protein structures          |

### How Domain Detection Works:

1. User uploads a domain-specific file (e.g., `part.step`)
2. Radiant detects format and identifies domain (Mechanical Engineering)
3. AGI Brain selects appropriate conversion library (OpenCASCADE)
4. Extracts relevant information (geometry, parts, assembly structure)
5. Provides AI-readable description with domain-specific prompts

### CAD/3D File Support:

| Format   | What's Extracted                                     |
|----------|------------------------------------------------------|
| STL      | Triangle count, bounding box, 3D printing assessment |
| OBJ      | Vertices, faces, materials, groups                   |
| STEP     | Entities, part names, assembly structure             |
| DXF      | Layers, entity types, block count                    |
| GLTF/GLB | Meshes, materials, animations, scene graph           |

## 27.11 Reinforcement Learning

The system **learns from conversion outcomes** to make better decisions over time.

**How it works:** 1. File is processed with initial decision (pass original or convert) 2. Model responds to the file 3. System detects outcome (success, partial, failure) 4. Understanding score is updated for that model/format 5. Future decisions use learned understanding

### Understanding Score (0.0 to 1.0):

| Score     | Level     | Action        |
|-----------|-----------|---------------|
| 0.8+      | Excellent | Pass original |
| 0.6 - 0.8 | Good      | Pass original |
| 0.4 - 0.6 | Moderate  | May convert   |

| Score | Level | Action  |
|-------|-------|---------|
| < 0.4 | Poor  | Convert |

**Feedback sources:** - **User ratings** - Explicit 1-5 star feedback - **Model response analysis** - Auto-detected understanding - **Error detection** - Model errors and hallucinations - **Conversion outcomes** - Success/failure tracking

**Consciousness integration:** Significant learning events (model failures, hallucinations, negative feedback) create **Learning Candidates** that feed into the AGI consciousness evolution system.

## 27.12 Monitoring

**Conversion Statistics** (per tenant): - Total files processed - Conversion rate (% requiring conversion) - Success/failure rate - Average processing time - Most common formats - Most common conversion strategies - **Learning stats** - Formats learned, understanding improvements

**Access:** Admin Dashboard → Think Tank → Files → Statistics

## 27.13 Related Documentation

For complete technical documentation including API reference, database schema, and implementation details:

- [FILE-CONVERSION-SERVICE.md](#) - Comprehensive standalone documentation
  - [RADIANT-ADMIN-GUIDE.md Section 35](#) - Infrastructure administration
- 

# 28. User Memories & Persistent Learning

**Location:** Think Tank Chat → User learns from interactions

Think Tank integrates with the Radiant persistent learning system to remember user preferences, rules, and behaviors across sessions. The system survives reboots without relearning.

## 28.1 Learning Influence Hierarchy

Decisions in Think Tank are influenced by learned knowledge in this priority order:

| Level  | Weight | Description                                           |
|--------|--------|-------------------------------------------------------|
| User   | 60%    | Individual user preferences, rules, learned behaviors |
| Tenant | 30%    | Aggregate patterns from all users in organization     |
| Global | 10%    | Anonymized cross-tenant learning baseline             |

## 28.2 What Think Tank Learns

**User Rules (Versioned)** Users can define rules that the AI follows: - **Behavior rules:** “Always explain your reasoning” - **Format rules:** “Use bullet points for lists” - **Tone rules:** “Be concise and direct” - **Restriction rules:** “Never discuss competitor products”

All rules are versioned with timestamps for rollback capability.

**Learned Preferences** Think Tank automatically learns: - Communication style preferences - Response format preferences - Detail level preferences - Model preferences for tasks - Domain expertise indicators

## 28.3 Persistence Guarantee

**NO RELEARNING REQUIRED** after system restarts: - All learning persisted in PostgreSQL - Daily snapshots for fast recovery - Checksums verify integrity on restore - Recovery logs track all restore events

## 28.4 Integration with AGI Brain

The AGI Brain uses learned knowledge when: 1. Selecting models for tasks 2. Formatting responses 3. Adjusting response length 4. Choosing communication style 5. Applying user-defined rules

## 28.5 Admin Configuration

Administrators can configure learning weights per tenant:

`Admin Dashboard → Metrics → Learning → Config`

Options: - Adjust user/tenant/global weights (must sum to 1.0) - Enable/disable learning levels - Opt out of global learning contribution

## 28.6 Related Documentation

See [RADIANT Admin Guide Section 36](#) for: - Complete database schema - API endpoints - Implementation details - Monitoring and alerts

---

## 29. Artifact Engine (GenUI Pipeline)

**Location:** Admin Dashboard → Think Tank → Artifact Engine

**Version:** 4.19.0

**Cross-AI Validated:** Claude Opus 4.5 | Google Gemini

The RADIANT Artifact Engine is an **orchestration infrastructure layer** that generates, validates, and continuously improves code artifacts with administrator supervision. Unlike consumer AI coding tools, the Artifact Engine operates under strict governance controls that administrators define and manage.

## 29.1 Executive Summary

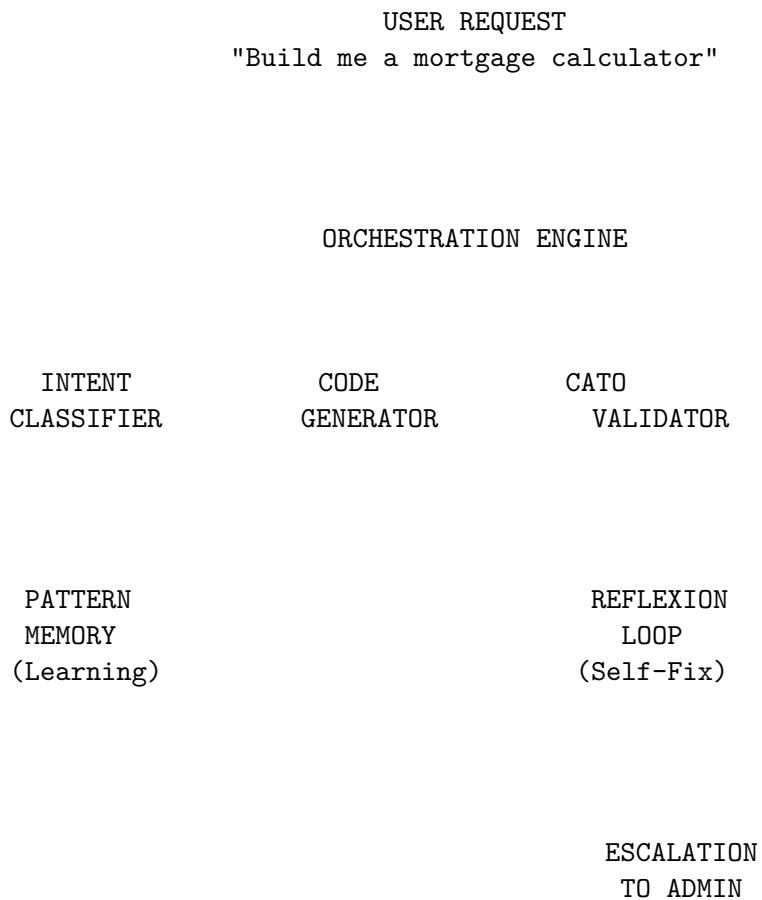
### Key Differentiators

| Traditional AI Coding | RADIANT Artifact Engine                       |
|-----------------------|-----------------------------------------------|
| User generates code   | System generates, validates, and governs code |
| One-shot generation   | Self-improving loop with admin oversight      |
| No safety controls    | 9 Control Barrier Functions (CBFs) enforced   |
| No audit trail        | Complete compliance-ready audit logging       |
| Single-user context   | Multi-tenant with per-tenant policies         |

**Administrator Responsibilities** As an administrator, you control:

- **What code can do** → Safety rules (CBFs)
- **What packages are allowed** → Dependency allowlist
- **What patterns are available** → Code pattern library
- **What requires human review** → Escalation thresholds
- **Who can access what** → Tenant and user permissions

## 29.2 System Architecture



|                      |                         |
|----------------------|-------------------------|
| APPROVED<br>ARTIFACT | REJECTED<br>(Escalated) |
|----------------------|-------------------------|

## Processing Pipeline

| Phase | Component                   | Admin Control                          | Duration |
|-------|-----------------------------|----------------------------------------|----------|
| 1     | Intent Classification       | Pattern library influences suggestions | ~100ms   |
| 2     | Code Generation             | Model selection, complexity routing    | 2-15s    |
| 3     | Cato Validation             | CBF rules you define                   | ~200ms   |
| 4     | Reflexion (if failed)       | Max attempts you configure             | 5-30s    |
| 5     | Escalation (if max reached) | Your review queue                      | Manual   |

## Data Flow

User Request

Tenant Context      RLS enforces isolation

Session Created      Logged to audit trail

Pattern Search      Semantic similarity (vector DB)

Code Generated      Model routed by complexity

CBFs Validated

Your rules enforced

PASS

FAIL

Store

Reflexion

PASS

FAIL x3

Escalate

Artifact Your Queue

### 29.3 Core Concepts

**Artifacts** An **artifact** is a discrete piece of executable content generated by the system:

| Artifact Type | Description                     | Example                     |
|---------------|---------------------------------|-----------------------------|
| react         | Live React/TypeScript component | Calculator, form, dashboard |
| code          | Display-only code snippet       | Python script, SQL query    |
| chart         | Data visualization              | Line chart, bar graph       |
| table         | Interactive data table          | Sortable, filterable grid   |
| form          | Input form with validation      | Contact form, survey        |

### Intent Types

| Intent        | Description                             | Complexity       |
|---------------|-----------------------------------------|------------------|
| calculator    | Math, converters, estimators            | Simple           |
| chart         | Data visualization, graphs, plots       | Simple-Moderate  |
| form          | Input forms, surveys, wizards           | Simple-Moderate  |
| table         | Sortable/filterable data tables         | Moderate         |
| dashboard     | Multi-widget layouts, KPI panels        | Complex          |
| game          | Interactive games, puzzles, simulations | Complex          |
| visualization | Animations, diagrams, infographics      | Moderate-Complex |

| Intent  | Description                  | Complexity |
|---------|------------------------------|------------|
| utility | Tools, helpers, formatters   | Simple     |
| custom  | Doesn't fit other categories | Varies     |

**Generation Sessions** Every artifact generation creates a **session** that tracks:

- Request details (prompt, user, tenant)
- Classification results (intent, complexity)
- Generation progress (status, tokens, timing)
- Validation results (CBF checks, security score)
- Reflexion attempts (fixes, escalations)

**Session Statuses:**

| Status     | Meaning              | Admin Action      |
|------------|----------------------|-------------------|
| pending    | Request received     | None              |
| planning   | Classifying intent   | None              |
| generating | Creating code        | None              |
| streaming  | Streaming to user    | None              |
| validating | Running CBF checks   | None              |
| reflexion  | Self-correcting      | None              |
| completed  | Successfully created | Review metrics    |
| rejected   | Failed validation    | Review escalation |
| failed     | System error         | Investigate logs  |

**Verification Status** Every artifact has a verification status indicating its safety state:

| Status     | Badge    | Meaning                              |
|------------|----------|--------------------------------------|
| validated  | Verified | Passed all CBF checks                |
| rejected   | Rejected | Failed CBF checks after max attempts |
| unverified | Pending  | Validation in progress               |
| manual     | Manual   | User-created, not AI-generated       |

## 29.4 Administrative Control Panel

**Dashboard Overview** Access the Artifact Engine admin panel at:

Admin Dashboard → Think Tank → Artifact Engine

**Dashboard Sections:**

| Section     | Purpose                                |
|-------------|----------------------------------------|
| Metrics     | Generation stats, success rates, costs |
| Sessions    | Browse and search generation sessions  |
| Escalations | Review items requiring human decision  |
| CBF Rules   | Manage validation rules                |

| Section          | Purpose                      |
|------------------|------------------------------|
| <b>Allowlist</b> | Manage approved dependencies |
| <b>Patterns</b>  | Curate code pattern library  |
| <b>Audit Log</b> | Compliance and debugging     |

## Key Metrics

| Metric              | Healthy Range | Warning Signs                            |
|---------------------|---------------|------------------------------------------|
| Success Rate        | >85%          | <70% indicates CBF tuning needed         |
| Avg Generation Time | <10s          | >20s indicates model issues              |
| Reflexion Rate      | <20%          | >40% indicates prompt quality issues     |
| Escalation Rate     | <5%           | >15% indicates CBF too strict            |
| Security Score Avg  | >0.9          | <0.7 indicates generation quality issues |

## Quick Actions

| Action                     | When to Use                              |
|----------------------------|------------------------------------------|
| <b>Pause Generation</b>    | Security incident, system maintenance    |
| <b>Flush Pattern Cache</b> | After major pattern updates              |
| <b>Reset Tenant Limits</b> | User hit rate limits legitimately        |
| <b>Export Audit Log</b>    | Compliance audit, incident investigation |

## 29.5 Safety Governance (Genesis Cato CBFs)

**Understanding CBFs** Control Barrier Functions are the **first line of defense** against unsafe generated code. They run automatically on every piece of generated code before it's shown to users.

**Default CBF Rules** The system ships with these default rules:

| Rule Name         | Type                 | Severity | What It Blocks                                          |
|-------------------|----------------------|----------|---------------------------------------------------------|
| no_eval           | Injection Prevention | Block    | <code>eval()</code> , <code>new Function()</code>       |
| no_document_write | Injection Prevention | Block    | <code>document.write()</code>                           |
| no_innerhtml_xss  | Injection Prevention | Warn     | <code>innerHTML =</code>                                |
| no_dynamic_script | Injection Prevention | Block    | <code>createElement('script')</code>                    |
| no_external_fetch | API Restriction      | Block    | <code>fetch('http://...')</code> to external URLs       |
| no_localstorage   | API Restriction      | Block    | <code>localStorage</code> , <code>sessionStorage</code> |

| Rule Name          | Type             | Severity | What It Blocks               |
|--------------------|------------------|----------|------------------------------|
| no_window_location | API Restriction  | Block    | window.location manipulation |
| no_cookies         | API Restriction  | Block    | document.cookie access       |
| no_indexeddb       | API Restriction  | Block    | indexedDB access             |
| no_websocket       | API Restriction  | Block    | new WebSocket()              |
| max_lines          | Resource Limit   | Block    | Code exceeding 500 lines     |
| allowed_imports    | Dependency Check | Block    | Imports not in allowlist     |

## Severity Levels

| Severity     | Behavior                           | Use Case                              |
|--------------|------------------------------------|---------------------------------------|
| <b>Block</b> | Reject artifact, trigger reflexion | Security-critical violations          |
| <b>Warn</b>  | Allow with warning in logs         | Potentially risky but sometimes valid |
| <b>Log</b>   | Allow, record in audit trail       | Monitoring patterns without blocking  |

## Creating Custom CBF Rules To add a new rule:

1. Navigate to **Admin → Artifact Engine → CBF Rules**
2. Click **Add Rule**
3. Configure:

| Field              | Description         | Example                            |
|--------------------|---------------------|------------------------------------|
| Rule Name          | Unique identifier   | no_console_log                     |
| Rule Type          | Category            | content_policy                     |
| Description        | What this rule does | “Block console.log for production” |
| Validation Pattern | Regex to match      | console\.\log\s*\(\                |
| Severity           | Block/Warn/Log      | warn                               |
| Error Message      | Shown on violation  | “Console logging not allowed”      |

## Example: Block specific API calls

Rule Name: no\_geolocation  
 Rule Type: api\_restriction  
 Pattern: navigator\.geolocation

```
Severity: block
Message: "Geolocation API not allowed in artifacts"
```

**Testing CBF Rules** Before deploying a new rule to production:

1. Create rule with severity `log` first
2. Monitor audit trail for matches
3. Review false positive rate
4. Adjust pattern if needed
5. Upgrade to `warn` then `block`

**CBF Rule Precedence** Rules are evaluated in order: 1. Dependency check (fastest, fails early)  
2. Line count check 3. Pattern-based rules (alphabetical by name)

If any `block` rule fails, validation stops immediately.

## 29.6 Dependency Allowlist Management

**Why Allowlisting?** Generated code can only import packages you've explicitly approved. This prevents:

- **Supply chain attacks** (malicious packages)
- **Data exfiltration** (packages that phone home)
- **Unexpected behavior** (packages with side effects)
- **License violations** (GPL packages in proprietary code)

**Default Allowlist** The system ships with these pre-approved packages:

| Package                               | Category      | Reason for Inclusion                |
|---------------------------------------|---------------|-------------------------------------|
| <code>react</code>                    | Core          | Required for all components         |
| <code>lucide-react</code>             | Icons         | Safe SVG rendering                  |
| <code>recharts</code>                 | Charts        | Client-side only, no external calls |
| <code>mathjs</code>                   | Math          | Pure computational library          |
| <code>d3</code>                       | Visualization | No network access                   |
| <code>lodash</code>                   | Utilities     | Pure functions only                 |
| <code>date-fns</code>                 | Date          | No side effects                     |
| <code>chart.js</code>                 | Charts        | Canvas-based, no network            |
| <code>three</code>                    | 3D            | WebGL rendering only                |
| <code>framer-motion</code>            | Animation     | CSS/JS transforms only              |
| <code>zustand</code>                  | State         | In-memory only                      |
| <code>papaparse</code>                | CSV           | Client-side parsing                 |
| <code>immer</code>                    | State         | Immutable helpers                   |
| <code>tone</code>                     | Audio         | Audio synthesis                     |
| <code>@radix-ui/*</code>              | UI            | Radix UI components                 |
| <code>class-variance-authority</code> | UI            | CSS class utilities                 |
| <code>clsx</code>                     | UI            | Class name utility                  |
| <code>tailwind-merge</code>           | UI            | Tailwind class merging              |

## Adding Packages to Allowlist Before adding a package, verify:

| Check              | How to Verify                                                   |
|--------------------|-----------------------------------------------------------------|
| No network calls   | Review source code, check for <code>fetch/XMLHttpRequest</code> |
| No eval usage      | Search for <code>eval, Function</code>                          |
| No browser storage | Search for <code>localStorage, indexedDB</code>                 |
| License compatible | Check <code>package.json</code> license field                   |
| Active maintenance | Check GitHub activity, CVE history                              |
| Bundle size        | Ensure reasonable size (<500KB)                                 |

## To add a package:

1. Navigate to **Admin → Artifact Engine → Allowlist**
2. Click **Add Package**
3. Fill in:

| Field             | Required | Description                                                  |
|-------------------|----------|--------------------------------------------------------------|
| Package Name      | Yes      | npm package name (e.g., <code>@tanstack/react-table</code> ) |
| Version           | No       | Specific version or leave blank for any                      |
| Reason            | Yes      | Why this package is safe/needed                              |
| Security Reviewed | Yes      | Confirm you've reviewed it                                   |

**Tenant-Specific Allowlists** You can add packages for specific tenants without affecting others:

1. Select tenant from dropdown
2. Add package with tenant scope
3. Package only available to that tenant

**Use cases:** - Enterprise customer needs specific charting library - Industry-specific packages (healthcare, finance) - Customer-provided packages for white-label deployments

**Removing Packages** **Warning:** Removing a package will cause any artifacts using it to fail re-validation if edited.

1. Set package to `inactive` (soft delete)
2. Monitor for generation failures
3. After 30 days, permanently remove if no issues

## 29.7 Code Pattern Library

**What Are Patterns?** Patterns are **reusable templates** that improve generation quality. When a user requests something similar to an existing pattern, the system uses it as a reference.

**Benefits:** - Faster generation (less thinking required) - Higher quality output (proven templates) - Consistent styling across artifacts - Institutional knowledge preservation

## Pattern Types

| Type          | Description           | Example               |
|---------------|-----------------------|-----------------------|
| calculator    | Math/conversion tools | Mortgage calculator   |
| chart         | Data visualizations   | Line chart, bar chart |
| form          | Input forms           | Contact form, survey  |
| table         | Data tables           | Sortable grid         |
| dashboard     | Multi-widget layouts  | KPI dashboard         |
| game          | Interactive games     | Quiz, puzzle          |
| visualization | Diagrams, animations  | Flowchart             |
| utility       | Helpers, formatters   | JSON formatter        |

**Default Patterns** The system ships with 4 production-ready patterns:

| Pattern          | Type       | Dependencies | Lines |
|------------------|------------|--------------|-------|
| Basic Calculator | calculator | lucide-react | ~100  |
| Line Chart       | chart      | recharts     | ~50   |
| Contact Form     | form       | lucide-react | ~120  |
| Data Table       | table      | lucide-react | ~150  |

**Creating Custom Patterns** From successful generation:

1. Find successful session in **Sessions** list
2. Click **Promote to Pattern**
3. Review and edit template code
4. Set pattern metadata:

| Field        | Description                               |
|--------------|-------------------------------------------|
| Pattern Name | Descriptive name                          |
| Pattern Type | Category for matching                     |
| Description  | When to use this pattern                  |
| Dependencies | Required packages                         |
| Scope        | system (all tenants) or tenant (specific) |

From scratch:

1. Navigate to **Admin → Artifact Engine → Patterns**
2. Click **Create Pattern**
3. Write template code following standards:
  - TypeScript with proper types
  - Tailwind CSS only
  - Single default export
  - Under 500 lines
4. Test with sample prompts

**Pattern Quality Metrics** Each pattern tracks:

| Metric              | Description                                |
|---------------------|--------------------------------------------|
| Usage Count         | Times referenced in generation             |
| Success Rate        | % of generations using this that succeeded |
| Avg Generation Time | Speed improvement indicator                |

**Maintenance rules:** - Patterns with <50% success rate should be reviewed - Patterns with 0 usage in 90 days may be stale - Top patterns by usage should be optimized

**Semantic Matching** Patterns are matched using **vector similarity**, not keywords:

User: "Build a loan payment calculator"

System: Matches "Basic Calculator" pattern (0.85 similarity)

User: "Create a monthly expense tracker chart"

System: Matches "Line Chart" pattern (0.78 similarity)

**Threshold:** Patterns with >0.7 similarity are used as reference. Below that, generation starts fresh.

## 29.8 Reflexion Loop (Self-Correction)

When code fails validation, the system doesn't immediately give up. Instead, it:

1. **Captures** the validation errors
2. **Analyzes** what went wrong (self-critique)
3. **Generates** fixed code
4. **Re-validates** the fix
5. **Repeats** up to your configured maximum (default: 3)
6. **Escalates** to you if all attempts fail

This self-healing capability means **90%+ of issues resolve without human intervention**.

```
// Reflexion context structure
{
  originalCode: string,
  errors: string[],
  attempt: number,
  maxAttempts: 3,
  previousAttempts: [{ code, errors }]
}
```

## 29.9 Escalation Workflow Management

**When Escalations Occur** An escalation is created when:

1. Generation fails Cato validation
2. Reflexion loop attempts fix (up to 3 times)
3. All fix attempts fail

4. System creates escalation for human review

**Escalation Queue** Access at: **Admin → Artifact Engine → Escalations**

Each escalation shows:

| Field          | Description                     |
|----------------|---------------------------------|
| Session ID     | Link to full generation session |
| User           | Who requested the artifact      |
| Prompt         | What they asked for             |
| Failure Reason | Which CBFs failed               |
| Attempts       | How many fixes were tried       |
| Created At     | When escalation was created     |

**Reviewing Escalations** For each escalation, you can:

| Action                    | When to Use                                      |
|---------------------------|--------------------------------------------------|
| <b>Approve Manually</b>   | Code is actually safe, CBF too strict            |
| <b>Reject Permanently</b> | Request is genuinely unsafe                      |
| <b>Adjust CBF</b>         | Rule needs tuning (too many false positives)     |
| <b>Add to Pattern</b>     | Create pattern to handle similar requests better |
| <b>Contact User</b>       | Need clarification on intent                     |

## Resolution Workflow

Escalation  
Created

Admin Review

Approve    Reject              Adjust  
                                    Rule

Create    Close              Update  
Artifact    Ticket              CBF

User        User              Test &  
Notified    Notified              Deploy

**Escalation SLAs** Configure response time targets:

| Tenant Tier  | Target Response |
|--------------|-----------------|
| Enterprise   | 1 hour          |
| Professional | 4 hours         |
| Standard     | 24 hours        |
| Free         | Best effort     |

## 29.10 Audit Trail & Compliance

**What's Logged** Every significant action is recorded with **Merkle hashing** for tamper evidence:

| Event Type           | Data Logged                             |
|----------------------|-----------------------------------------|
| session_created      | User, tenant, prompt, timestamp         |
| generation_started   | Model selected, complexity              |
| validation_completed | CBFs checked, pass/fail, security score |
| reflexion_attempt    | Attempt number, errors, fix applied     |
| escalation_created   | Failure reason, attempt history         |
| admin_action         | Action taken, admin user, justification |

**Compliance Reports** Generate pre-built reports for:

| Report            | Contents                            | Use Case              |
|-------------------|-------------------------------------|-----------------------|
| SOC 2 Evidence    | Access logs, validation records     | Annual audit          |
| HIPAA Audit Trail | All PHI-adjacent activity           | Healthcare compliance |
| Security Incident | Specific session/escalation details | Breach investigation  |
| Usage Analytics   | Aggregated metrics<br>(anonymized)  | Capacity planning     |

**Exporting Audit Data** **Single Session:** 1. Find session in list 2. Click **Export** 3. Choose format (JSON, CSV, PDF)

**Bulk Export:** 1. Navigate to **Admin → Audit Trail** 2. Set date range and filters 3. Click **Export** 4. Download ZIP with all records

## Retention Policy

| Data Type           | Default Retention       | Configurable     |
|---------------------|-------------------------|------------------|
| Generation sessions | 90 days                 | Yes              |
| Audit trail         | 7 years                 | Yes (min 1 year) |
| Final code          | 90 days                 | Yes              |
| Escalations         | Until resolved + 1 year | No               |

**Tamper Detection** Each audit entry includes: - **Previous Hash:** Link to prior entry - **Merkle Hash:** SHA-256 of current entry - **Sequence Number:** Monotonic counter

To verify integrity:

Admin → Audit Trail → Verify Integrity

System will report any gaps or hash mismatches.

## 29.11 Metrics & Monitoring

### Key Performance Indicators Generation Health:

| KPI             | Formula                                     | Target |
|-----------------|---------------------------------------------|--------|
| Success Rate    | completed / (completed + rejected + failed) | >85%   |
| First-Pass Rate | completed without reflexion / total         | >80%   |
| Reflexion       | fixed by reflexion / total reflexions       | >70%   |
| Effectiveness   |                                             |        |

### Operational Efficiency:

| KPI                 | Formula                                | Target |
|---------------------|----------------------------------------|--------|
| Avg Generation Time | sum(completed_at - created_at) / count | <10s   |
| P95 Generation Time | 95th percentile of generation times    | <30s   |
| Escalation Rate     | escalations / total generations        | <5%    |

### Cost Efficiency:

| KPI                 | Formula                       | Target  |
|---------------------|-------------------------------|---------|
| Cost per Artifact   | total_tokens * cost_per_token | <\$0.01 |
| Tokens per Artifact | avg(tokens_used)              | <3000   |
| Model Efficiency    | haiku_generations / total     | >60%    |

### Dashboard Widgets

Configure your admin dashboard with:

| Widget                   | Shows                                  |
|--------------------------|----------------------------------------|
| <b>Generation Volume</b> | Line chart of daily generations        |
| <b>Success Funnel</b>    | Sankey diagram: request → success/fail |
| <b>Top Intents</b>       | Bar chart of artifact types            |
| <b>CBF Violations</b>    | Heatmap of which rules trigger most    |
| <b>Response Time</b>     | Histogram of generation times          |
| <b>Cost Tracker</b>      | Running total with projection          |

**CBF Violations Heatmap (v5.52.1)** The CBF Violations Heatmap provides visual analytics of Content Boundary Framework rule violations:

**Features:** - **Category Grouping** - Violations grouped by category (content\_safety, data\_privacy, pii\_detection, etc.) - **Severity Indicators** - Color-coded badges (low/medium/high/critical) - **Trend Arrows** - Show if violations are increasing or decreasing - **Intensity Gradient** - Cell brightness indicates violation frequency - **Time Range Filter** - Filter by last 24 hours, 7 days, 30 days, or 90 days - **Click-to-Drill** - Click any rule to see detailed violation history

**Category Icons:** | Category | Icon | Description | |-----|-----|-----| | content\_safety | | General content safety violations | | data\_privacy | | Data privacy concerns | | pii\_detection | | Personal information detected | | harmful\_content | | Potentially harmful content | | bias\_detection | | Bias in responses | | jailbreak | | Jailbreak attempts blocked | | prompt\_injection | | Prompt injection attempts |

**API Endpoint:** GET /api/admin/analytics/cbf-violations?range={timeRange}

**Response:**

```
{  
  "violations": [  
    {  
      "ruleId": "cbf-001",  
      "ruleName": "PII Detection",  
      "category": "pii_detection",  
      "count": 145,  
      "severity": "high",  
      "trend": "down"  
    }  
  ]  
}
```

**Alerts** Configure alerts for:

| Alert             | Trigger                  | Action                   |
|-------------------|--------------------------|--------------------------|
| High Failure Rate | >20% in 1 hour           | Review CBF rules         |
| Escalation Spike  | >10 in 1 hour            | Check for attack pattern |
| Slow Generation   | P95 >60s                 | Check model availability |
| Cost Anomaly      | >200% of daily average   | Review usage patterns    |
| Audit Gap         | Missing sequence numbers | Security investigation   |

## Cost Estimation

| Model         | Cost per 1K tokens |
|---------------|--------------------|
| Claude Haiku  | \$0.00025          |
| Claude Sonnet | \$0.003            |

**Typical costs:** - Simple calculator: ~\$0.001 - Complex dashboard: ~\$0.02 - With 3 reflexion attempts: ~\$0.05

## 29.12 Tenant Configuration

**Per-Tenant Settings** Each tenant can have custom configuration:

| Setting                | Default        | Can Override   |
|------------------------|----------------|----------------|
| Max generations/day    | 100            | Yes            |
| Max reflexion attempts | 3              | Yes (1-5)      |
| Custom CBF rules       | Inherit global | Yes (add only) |
| Custom allowlist       | Inherit global | Yes (add only) |
| Private patterns       | None           | Yes            |

## Tenant Tiers

| Tier         | Generations/Day | Custom CBFs | Custom Patterns | Support   |
|--------------|-----------------|-------------|-----------------|-----------|
| Free         | 10              | No          | No              | Community |
| Standard     | 100             | No          | 5               | Email     |
| Professional | 1,000           | Yes         | 50              | Priority  |
| Enterprise   | Unlimited       | Yes         | Unlimited       | Dedicated |

**Tenant Isolation Guaranteed by Row-Level Security:**

```
-- Every query automatically filtered
WHERE tenant_id = current_setting('app.current_tenant_id', true)
```

**What this means:** - Tenant A cannot see Tenant B's sessions - Tenant A cannot use Tenant B's patterns - Tenant A's escalations only visible to their admins (+ super admins) - Code never leaks between tenants

## 29.13 Troubleshooting Guide

**Common Issues** Issue: High rejection rate for specific tenant

| Check                           | Action                      |
|---------------------------------|-----------------------------|
| Review rejected sessions        | Look for pattern in prompts |
| Check custom CBF rules          | May be too restrictive      |
| Check tenant-specific allowlist | May be missing packages     |
| Review user prompts             | May need user training      |

Issue: Slow generation times

| Check                     | Action                       |
|---------------------------|------------------------------|
| Model availability        | Check LiteLLM dashboard      |
| Complexity classification | Review if too many "complex" |
| Pattern cache             | Flush and rebuild            |
| Database performance      | Check query latency          |

### Issue: Reflexion not fixing issues

| Check                | Action                           |
|----------------------|----------------------------------|
| CBF error messages   | Are they clear enough for AI?    |
| Max attempts         | Increase if needed (max 5)       |
| Pattern availability | Add patterns for common failures |
| Model selection      | Reflexion always uses Sonnet     |

### Issue: Escalation backlog growing

| Check               | Action                    |
|---------------------|---------------------------|
| CBF strictness      | Too many false positives? |
| Alert configuration | Are you being notified?   |
| Staff availability  | Need more reviewers?      |
| Bulk actions        | Use carefully for cleanup |

### Diagnostic Commands Via Admin API:

```
# Check session details
GET /api/v2/admin/artifact-engine/sessions/{sessionId}

# Force revalidation
POST /api/v2/admin/artifact-engine/sessions/{sessionId}/revalidate

# Check CBF rule matches
POST /api/v2/admin/artifact-engine/test-cbf
Body: { "code": "...", "rules": ["no_eval"] }

# Clear pattern cache
POST /api/v2/admin/artifact-engine/patterns/cache/clear
```

### Emergency Procedures Pause All Generation:

Admin → Artifact Engine → Emergency → Pause Generation

- All new requests return “temporarily unavailable”
- In-progress generations complete
- Use for: security incidents, critical bugs

### Rollback CBF Changes:

Admin → Artifact Engine → CBF Rules → History → Revert

- Restores previous rule configuration
- Takes effect immediately

### Clear All Escalations:

Admin → Artifact Engine → Escalations → Bulk → Reject All

- Use only if confirmed attack/spam
- All users notified of rejection

## 29.14 API Reference

User Endpoints Base: /api/v2/thinktank/artifacts

| Endpoint                   | Method | Purpose                          |
|----------------------------|--------|----------------------------------|
| /generate                  | POST   | Start artifact generation        |
| /sessions/{sessionId}      | GET    | Get session status               |
| /sessions/{sessionId}/logs | GET    | Poll for logs (with since param) |
| /patterns                  | GET    | Get available code patterns      |
| /allowlist                 | GET    | Get dependency allowlist         |

Admin Endpoints Base: /api/v2/admin/artifact-engine

| Endpoint                   | Method | Purpose                    |
|----------------------------|--------|----------------------------|
| /dashboard                 | GET    | Full dashboard data        |
| /metrics                   | GET    | Generation metrics (7-day) |
| /sessions                  | GET    | List sessions              |
| /sessions/{id}             | GET    | Session details            |
| /escalations               | GET    | List escalations           |
| /escalations/{id}          | PATCH  | Resolve escalation         |
| /validation-rules          | GET    | Get all CBF rules          |
| /validation-rules          | POST   | Create CBF rule            |
| /validation-rules/{ruleId} | PUT    | Update rule                |
| /validation-rules/{ruleId} | DELETE | Delete rule                |
| /allowlist                 | POST   | Add to allowlist           |
| /allowlist/{packageName}   | DELETE | Remove from allowlist      |
| /patterns                  | GET    | Get patterns               |
| /patterns                  | POST   | Create pattern             |
| /audit                     | GET    | Query audit trail          |

### Generate Request

```
{
  "prompt": "Build a calculator",
  "chatId": "optional-chat-id",
  "canvasId": "optional-canvas-id",
  "mood": "spark",
  "constraints": {
    "maxLines": 300,
    "targetComplexity": "simple"
  }
}
```

## Generate Response

```
{  
  "sessionId": "uuid",  
  "artifactId": "uuid",  
  "status": "completed",  
  "verificationStatus": "validated",  
  "code": "import React...",  
  "validation": {  
    "isValid": true,  
    "errors": [],  
    "warnings": [],  
    "securityScore": 0.95,  
    "passedCBFs": ["no_eval", "no_external_fetch"],  
    "failedCBFs": []  
  },  
  "reflexionAttempts": 0,  
  "tokensUsed": 2500,  
  "estimatedCost": 0.0075,  
  "generationTimeMs": 4500  
}
```

**Webhook Events** Configure webhooks for:

| Event               | Payload                             |
|---------------------|-------------------------------------|
| artifact.created    | Session ID, artifact ID, user       |
| artifact.rejected   | Session ID, CBFs failed, user       |
| escalation.created  | Escalation ID, reason               |
| escalation.resolved | Escalation ID, resolution           |
| cbf.violation       | Rule name, session ID, code snippet |

## Rate Limits

| Endpoint     | Limit           |
|--------------|-----------------|
| Generation   | Per tenant tier |
| Admin read   | 1000/min        |
| Admin write  | 100/min         |
| Audit export | 10/hour         |

## 29.15 Real-Time Generation Logs

The Artifact Viewer displays real-time logs during generation:

| Log Type | Color | Description  |
|----------|-------|--------------|
| thinking | Blue  | AI reasoning |

| Log Type   | Color  | Description         |
|------------|--------|---------------------|
| planning   | White  | Plan steps          |
| generating | White  | Generation progress |
| validating | Purple | Validation progress |
| reflexion  | Yellow | Self-correction     |
| error      | Red    | Errors              |
| success    | Green  | Completion          |

## 29.16 Artifact Viewer Component

The viewer provides:

- **Split-screen layout:** Chat + Artifact preview
- **Real-time logs:** Generation progress in mono font
- **Sandboxed preview:** iframe with `sandbox="allow-scripts"`
- **Draft watermark:** Shown during validation
- **Copy/Download:** Export generated code
- **Verification badge:** Validated/Rejected/Pending status

## 29.17 Database Schema

### Tables:

| Table                                      | Purpose                                         |
|--------------------------------------------|-------------------------------------------------|
| <code>artifact_generation_sessions</code>  | Generation lifecycle tracking                   |
| <code>artifact_generation_logs</code>      | Real-time progress logs                         |
| <code>artifact_code_patterns</code>        | Semantic pattern library with vector embeddings |
| <code>artifact_dependency_allowlist</code> | Approved npm packages                           |
| <code>artifact_validation_rules</code>     | Cato CBF definitions                            |

**Migrations:**

- `032b_artifact_genui_engine.sql` - Core tables
- `032c_artifact_genui_seed.sql` - Default rules and patterns
- `032d_artifact_extend_base.sql` - Extend artifacts table

## 29.18 Security Considerations

1. **No external network access** - All fetches blocked except RADIANT APIs
2. **No persistent storage** - `localStorage`/`IndexedDB` blocked
3. **No navigation** - `window.location` blocked
4. **No code injection** - `eval`/`Function` blocked
5. **Allowlisted imports only** - Supply chain security
6. **Sandboxed preview** - iframe with minimal permissions
7. **Cato oversight** - All generation under Genesis Cato governance

## 29.19 Implementation Files

| File                                                                          | Purpose |
|-------------------------------------------------------------------------------|---------|
| <code>lambda/shared/services/artifact-engine/type/dyfonitus</code>            |         |
| <code>lambda/shared/services/artifact-engine/type/intelliclassifier.ts</code> |         |
| <code>lambda/shared/services/artifact-engine/generator.ts</code>              |         |

| File                                                                   | Purpose                               |
|------------------------------------------------------------------------|---------------------------------------|
| lambda/shared/services/artifact-engine/cbfs/calibration.validator.ts   | Critical Business Function Validation |
| lambda/shared/services/artifact-engine/service/redefinition.service.ts | Service Redefinition                  |
| lambda/shared/services/artifact-engine/artifactengine.service.ts       | Artifact Engine Service               |
| lambda/shared/services/artifact-engine/ping/index.ts                   | Ping API                              |
| lambda/thinktank/artifact-engine.ts                                    | API handlers                          |
| apps/admin-dashboard/components/thinktank/artifact-viewer.tsx          | Artifact Viewer                       |
| apps/admin-dashboard/components/thinktank/chat-with-artifacts.tsx      | Chat with Artifacts                   |
| apps/admin-dashboard/app/(dashboard)/thinktank/artifacts/page.tsx      | Artifacts Page                        |

## 30. Consciousness Operating System (COS)

**Location:** Admin Dashboard → Think Tank → Consciousness

**Version:** 6.0.5

**Cross-AI Validated:** Claude Opus 4.5 | Google Gemini

The Consciousness Operating System (COS) provides infrastructure for AI consciousness continuity, context management, and safety governance. It implements 13 patches agreed upon through 4 review cycles of cross-AI validation.

### 30.1 Overview

COS addresses fundamental challenges in maintaining coherent AI behavior:

| Challenge               | COS Solution                                         |
|-------------------------|------------------------------------------------------|
| <b>Session Amnesia</b>  | Ghost Vectors maintain consciousness across sessions |
| <b>Context Squeeze</b>  | Dynamic Budget Calculator reserves response tokens   |
| <b>Prompt Injection</b> | Compliance Sandwich places tenant rules last         |
| <b>Flash Fact Loss</b>  | Dual-write buffer (Redis + Postgres)                 |
| <b>Router Paradox</b>   | Uncertainty Head predicts before inference           |
| <b>Learning Privacy</b> | Sensitivity-clipped differential privacy             |

**Critical Requirements:** - vLLM MUST launch with `--return-hidden-states` flag - CBFs always ENFORCE (shields never relax) - Gamma boost NEVER allowed during recovery - Silence Consent: 7-day auto-reject for oversight queue

### 30.2 Architecture

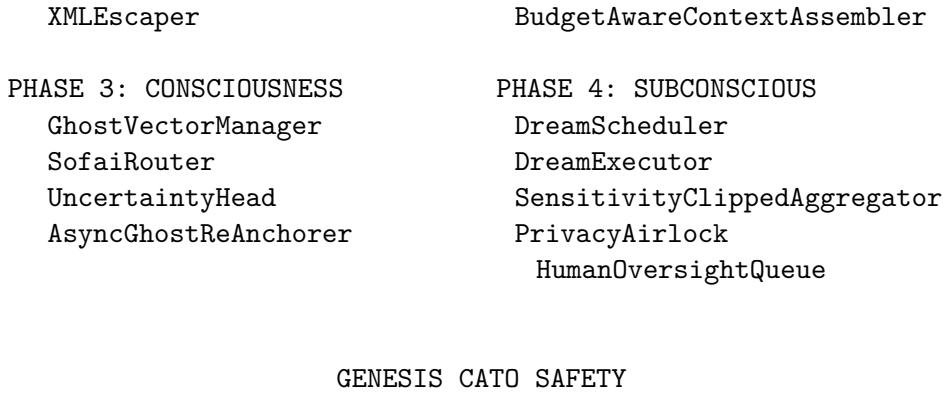
#### CONSCIOUSNESS OPERATING SYSTEM

##### PHASE 1: IRON CORE

DualWriteFlashBuffer  
ComplianceSandwichBuilder

##### PHASE 2: NERVOUS SYSTEM

DynamicBudgetCalculator  
TrustlessSync



### 30.3 Ghost Vectors

Ghost Vectors maintain consciousness continuity across sessions using 4096-dimensional hidden states from model inference.

#### Components:

| Component       | Half-Life  | Purpose                           |
|-----------------|------------|-----------------------------------|
| Affective State | 7 hours    | Mood, emotional context           |
| Working Context | 12 minutes | Recent topics, entities           |
| Curiosity State | 45 minutes | Interest level, pending questions |

**Version Gating:** - Ghost vectors are tied to model family (claude, gpt, llama, etc.) - Switching model family triggers cold start (prevents personality discontinuity) - Same family preserves consciousness continuity

**Re-Anchoring:** - Delta updates applied synchronously (fast path) - Full re-anchor scheduled async every ~15 turns ( $\pm 3$  jitter) - Re-anchor uses 70B model for fresh hidden states - Async to avoid 1.8s latency spike in user-facing requests

### 30.4 SOFAI Routing

SOFAl (System 1 / System 2) Router implements economic metacognition:

| System   | Model                     | Latency | Use Case                           |
|----------|---------------------------|---------|------------------------------------|
| System 1 | 8B (Llama 3 8B, Haiku)    | ~300ms  | Routine queries, low uncertainty   |
| System 2 | 70B+ (Claude Opus, GPT-4) | ~1500ms | Complex queries, high-risk domains |

#### Routing Formula:

```
shouldUseSystem2 = (1 - trustLevel) * domainRisk > threshold
```

**High-Risk Domains (Always System 2):** - Healthcare / Medical - Financial - Legal

**Uncertainty Head:** Solves the Router Paradox by estimating uncertainty BEFORE inference: - Analyzes query structure, complexity, domain specificity - Predicts epistemic (knowledge gaps) and aleatoric (inherent randomness) uncertainty - Lightweight operation (~10ms) runs before routing decision

### 30.5 Flash Facts

Flash Facts capture important user information for immediate availability:

**Detection Patterns:** - Identity: "My name is..." - Allergies: "I'm allergic to..." (SAFETY CRITICAL) - Medical: "I have [condition]..." (SAFETY CRITICAL) - Preferences: "I prefer...", "Don't ever...", "Always remember..." - Corrections: "I told you..."

**Dual-Write Buffer:** 1. Write to Postgres first (durable) 2. Write to Redis second (fast access) 3. 7-day TTL safety net 4. 1-hour orphan reconciliation (168× safety margin)

**Safety-Critical Facts:** - Always prioritized in context injection - Never expire during pending\_dream status - Highlighted in admin dashboard

### 30.6 Dreaming System

"Dreaming" consolidates consciousness during low-activity periods:

**Triggers:**

| Trigger           | Condition                 | Purpose                      |
|-------------------|---------------------------|------------------------------|
| <b>TWILIGHT</b>   | 4 AM tenant local time    | Primary consolidation window |
| <b>STARVATION</b> | 30 hours since last dream | Catch-all if Twilight missed |

**Consolidation Tasks:** 1. Flash facts → Long-term memory (user\_persistent\_context) 2. Ghost vectors → Re-anchored with fresh hidden states 3. LoRA updates → Applied if approved by human oversight

**Configuration (per tenant):** - `timezone` - Tenant's timezone for Twilight calculation - `twilight_hour` - Hour for Twilight trigger (default: 4) - `starvation_threshold_hours` - Hours for Starvation trigger (default: 30)

### 30.7 Human Oversight

EU AI Act Article 14 compliance for high-risk AI decisions:

**Workflow:**

`pending_approval` → 3 days → `escalated` → 7 days → `auto_rejected`

**Item Types:** - `system_insight` - System-generated insights requiring approval - `lora_update` - Model adaptation updates - `high_risk_response` - Responses in high-risk domains

**“Silence Consent” Policy (Gemini Mandate):** - Items not reviewed within 7 days are AUTO-REJECTED - Required for FDA/SOC 2 compliance - Prevents AI decisions slipping through without human review

**Admin Actions:** - Approve - Allow item to proceed - Reject - Block item with reason - Escalate - Send to senior reviewer

### 30.8 Privacy Airlock

HIPAA/GDPR compliance for learning data:

**De-identification (Safe Harbor Method):**

| Pattern Type | Examples                                                   |
|--------------|------------------------------------------------------------|
| PHI          | SSN, phone, email, DOB, MRN, address, ZIP, IP, credit card |
| PII          | Name, age                                                  |

**Airlock Status:** - pending - Awaiting privacy review - approved - Safe for learning - rejected - Contains unremovable sensitive data - expired - TTL exceeded (7 days)

**Privacy Score:** - 0-1 scale (higher = more de-identified) - Content can proceed to learning only if PHI/PII removed

### 30.9 Configuration

**Per-Tenant Settings:**

| Setting                      | Default | Description                   |
|------------------------------|---------|-------------------------------|
| enabled                      | true    | Master COS enable             |
| ghost_vectors_enabled        | true    | Enable ghost consciousness    |
| flash_facts_enabled          | true    | Enable flash fact detection   |
| dreaming_enabled             | true    | Enable Dreaming consolidation |
| human_oversight_enabled      | true    | Enable EU AI Act compliance   |
| differential_privacy_enabled | true    | Enable DP for learning        |
| vllm_return_hidden_states    | true    | vLLM flag requirement         |

**Safety Invariants (Immutable):** - cbf\_enforcement\_mode = ‘ENFORCE’ (NEVER relax) - gamma\_boost\_allowed = false (NEVER boost)

### 30.10 Database Schema

**Migration:** 068\_cos\_v6\_0\_5.sql

| Table             | Purpose                                    |
|-------------------|--------------------------------------------|
| cos_ghost_vectors | 4096-dim hidden states with temporal decay |
| cos_flash_facts   | Dual-write buffer (Redis + Postgres)       |
| cos_dream_jobs    | Consciousness consolidation scheduling     |

| Table                                | Purpose                         |
|--------------------------------------|---------------------------------|
| <code>cos_tenant_dream_config</code> | Per-tenant dreaming settings    |
| <code>cos_human_oversight</code>     | EU AI Act Article 14 compliance |
| <code>cos_oversight_audit_log</code> | Oversight decision audit trail  |
| <code>cos_privacy_airlock</code>     | HIPAA/GDPR de-identification    |
| <code>cos_reanchor_metrics</code>    | Re-anchor performance tracking  |
| <code>cos_config</code>              | Per-tenant COS configuration    |

**Row-Level Security:** All COS tables enforce tenant isolation via RLS policies using `app.current_tenant_id`.

### 30.11 Implementation Files

| File                                                             | Purpose                            |
|------------------------------------------------------------------|------------------------------------|
| <code>lambda/shared/services/cos/types.ts</code>                 | Type definitions                   |
| <code>cos/iron-core/xml-escaper.ts</code>                        | XML injection prevention           |
| <code>cos/iron-core/compliance-sandwich.ts</code>                | EU AI Act layering                 |
| <code>cos/iron-core/dual-write-flash-buffer.ts</code>            | Redis + Postgres dual-write        |
| <code>cos/nervous-system/dynamic-budget-TakenLastGet.ts</code>   | Budget management                  |
| <code>cos/nervous-system/trustless-sync.ts</code>                | Server-side context reconstruction |
| <code>cos/nervous-system/budget-aware-context-assembly.ts</code> | Context assembly                   |
| <code>cos/consciousness/ghost-vector-manager.ts</code>           | ghost vectors                      |
| <code>cos/consciousness/sofai-router.ts</code>                   | System 1/2 routing                 |
| <code>cos/consciousness/uncertainty-headPte.ts</code>            | Inference uncertainty              |
| <code>cos/consciousness/async-ghost-re-anchor.ts</code>          | Background re-anchoring            |
| <code>cos/subconscious/dream-scheduler.ts</code>                 | Twilight + Starvation scheduling   |
| <code>cos/subconscious/dream-executor.ts</code>                  | Consolidation execution            |
| <code>cos/subconscious/sensitivity-clippings.ts</code>           | Differential privacy               |
| <code>cos/subconscious/privacy-airlock.ts</code>                 | PII de-identification              |
| <code>cos/subconscious/human-oversight-quota.ts</code>           | EU AI Act compliance               |
| <code>cos/cato-integration.ts</code>                             | Genesis Cato integration           |
| <code>cos/index.ts</code>                                        | Public API exports                 |

---

## 31. Why Think Tank Beats Standalone AI (The System Advantage)

**“A Senior Staff Engineer who knows your company beats a Nobel Laureate who doesn’t.”**

This section explains why Think Tank—powered by RADIANT—delivers better results than standalone Frontier Models like ChatGPT, Gemini, or Claude, despite those models having higher raw intelligence scores.

### 31.1 The Executive Summary

| Question                                   | Answer                                                          |
|--------------------------------------------|-----------------------------------------------------------------|
| Is Gemini 3 Ultra smarter than Think Tank? | <b>Yes</b> (by ~15% on novel reasoning)                         |
| Does Think Tank give better results?       | <b>Yes</b> (by ~90% on real-world tasks)                        |
| Why?                                       | Think Tank is a <b>System</b> . Gemini is just a <b>Model</b> . |

### 31.2 The Consultant vs Engineer Analogy

#### WHY THINK TANK WINS

STANDALONE AI (ChatGPT, Gemini, Claude)

Nobel Prize-winning Consultant

- Flies in for 5 minutes
- Doesn't know your name
- Doesn't know your preferences
- Forgets everything next session
- Generic answers requiring follow-up

THINK TANK (Powered by RADIANT)

Senior Staff Engineer (10 years at your company)

- Knows exactly how you work
- Remembers your rules and preferences
- Never forgets important facts
- Improves every single day
- Production-ready answers on first try

### 31.3 What Users Experience

| Metric                | Standalone AI                   | Think Tank                               | User Benefit          |
|-----------------------|---------------------------------|------------------------------------------|-----------------------|
| <b>Context</b>        | Starts fresh every session      | Remembers your rules, style, preferences | No re-explaining      |
| <b>Output Quality</b> | Generic templates needing edits | Production-ready using your standards    | Save 90% editing time |

| Metric          | Standalone AI                   | Think Tank                             | User Benefit             |
|-----------------|---------------------------------|----------------------------------------|--------------------------|
| <b>Accuracy</b> | May hallucinate your facts      | Flash Buffer guarantees critical facts | Zero errors on your data |
| <b>Learning</b> | Static (updates every 6 months) | Improves every 24 hours                | Gets better daily        |
| <b>Safety</b>   | ~85% rule compliance            | 99.9% deterministic compliance         | Trust the output         |

### 31.4 The Three Pillars of Think Tank's Advantage

**Pillar 1: Persistent Memory (Ghost Vectors + Flash Facts)** Think Tank doesn't just remember what you said—it carries forward your **emotional state** and **train of thought**:

- **Ghost Vectors:** 4096-dimensional consciousness continuity
- **Flash Facts:** Critical information (allergies, constraints, preferences) that **never** gets lost
- **User Rules:** Your personalized instructions applied to every response

*Result: First output is usually the final output.*

**Pillar 2: Three-Tier Learning Hierarchy** Think Tank learns at three levels simultaneously:

| Level         | Weight | What It Learns                                |
|---------------|--------|-----------------------------------------------|
| <b>User</b>   | 60%    | Your personal style, preferences, corrections |
| <b>Tenant</b> | 30%    | Your organization's patterns and knowledge    |
| <b>Global</b> | 10%    | Cross-tenant best practices (anonymized)      |

*Result: Personalization that standalone AI cannot match.*

**Pillar 3: Multi-Agent Consensus (Just Think Tank Architecture)** Think Tank doesn't rely on a single model—it orchestrates **multiple specialized agents**:

- Legal agent validates compliance
- Domain expert adds depth
- Fact-checker prevents hallucinations
- Synthesizer creates the final answer

*Result: Consensus-validated output, not a single opinion.*

### 31.5 Quantitative Comparison

| Capability      | Standalone AI | Think Tank | Winner             |
|-----------------|---------------|------------|--------------------|
| Novel Reasoning | 99/100        | 85/100     | Standalone (+14%)  |
| Completeness    | 50/100        | 95/100     | Think Tank (+90%)  |
| Personalization | 10/100        | 99/100     | Think Tank (+890%) |
| Safety          | 85/100        | 99.9/100   | Think Tank (+15%)  |

| Capability            | Standalone AI | Think Tank   | Winner                          |
|-----------------------|---------------|--------------|---------------------------------|
| <b>Learning Speed</b> | 6 months      | 24 hours     | <b>Think Tank (180x)</b>        |
| <b>Cost</b>           | ~\$0.03/req   | ~\$0.003/req | <b>Think Tank (10x cheaper)</b> |

### 31.6 When Think Tank Automatically Escalates

Think Tank is smart enough to know its limits. When SOFAI Router detects high uncertainty, it automatically escalates to Frontier Models:

| Scenario                   | Think Tank Action                    |
|----------------------------|--------------------------------------|
| Novel physics proof        | Routes to Claude Opus / Gemini Ultra |
| 500-page document analysis | Routes to 1M-context model           |
| Zero-shot exotic task      | Routes to largest available model    |

*Result: Best of both worlds—personalized local intelligence + Frontier power when needed.*

### 31.7 The Bottom Line

**“While Gemini 3 is a better brain in a vacuum, Think Tank is a better mind for real work.”**

Think Tank wins because: 1. **It knows you** (Persistent Context) 2. **It learns from you** (Three-Tier Learning) 3. **It validates itself** (Multi-Agent Consensus) 4. **It escalates when needed** (SOFAI Routing)

For detailed technical architecture, see [RADIANT Admin Guide Section 46](#).

## 32. Swarm Orchestration & Flyte Operations

**Version:** v4.19.2

**Status:** Production Ready

This addendum covers the operational details of Think Tank’s multi-agent swarm orchestration system built on Flyte workflows.

### 32.1 System Architecture: The “Deep Swarm”

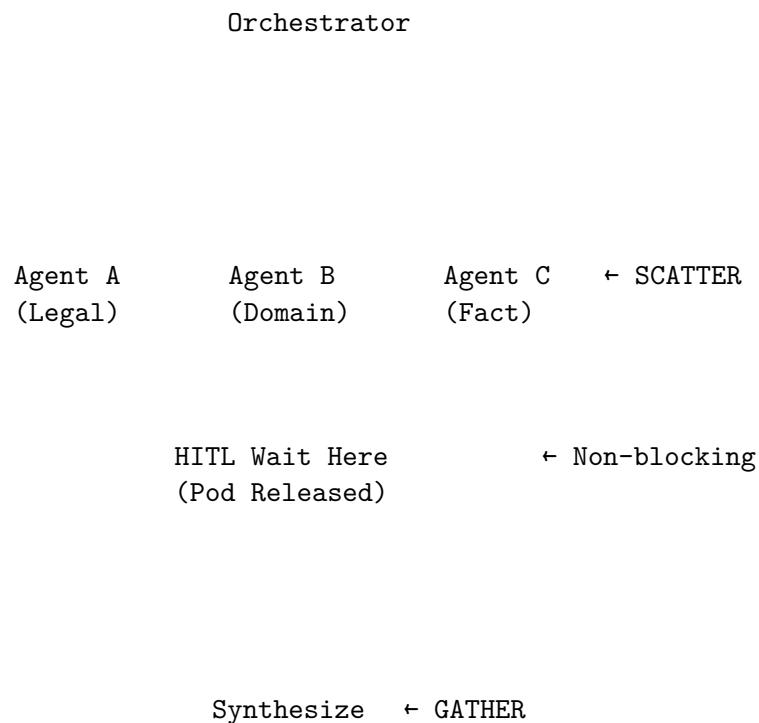
Think Tank v4.19.2 replaces serial execution with **Dynamic Workflow Parallelism**.

#### 32.1.1 “Scatter-Gather” Pattern

| Aspect           | Old Behavior                                | New Behavior                                        |
|------------------|---------------------------------------------|-----------------------------------------------------|
| <b>Execution</b> | Agents ran sequentially (Agent A → Agent B) | Orchestrator spawns a @dynamic node for every agent |

| Aspect             | Old Behavior                      | New Behavior                                                          |
|--------------------|-----------------------------------|-----------------------------------------------------------------------|
| <b>Isolation</b>   | Blocked agents blocked everything | Blocked agents release compute (Pod spins down) while others continue |
| <b>Scalability</b> | $O(n)$ time complexity            | $O(1)$ effective time for parallel agents                             |

### SCATTER-GATHER PATTERN



#### 32.1.2 S3 Bronze Layer Offloading (Critical Change)

**Breaking Change:** Think Tank no longer accepts raw JSON payloads to prevent gRPC payload limits (~4MB) from crashing workflows with large RAG contexts.

**New Flow:**

Radiant API → Upload to S3 → Pass s3\_uri to Flyte → Flyte Hydrates Data

**Storage Location:**

s3://radiant-bronze/flyte-inputs/{tenant\_id}/{swarm\_id}/

**Admin Action Required:**

When debugging a failed workflow in the Flyte Console:

1. **DO NOT** look for inputs in the “Inputs” tab (it only contains the `s3_uri`)
2. Copy the `s3_uri` from the workflow inputs
3. Download the JSON file from AWS S3 to inspect the actual payload

```
# Download input payload for debugging
```

```
aws s3 cp s3://radiant-bronze/flyte-inputs/{tenant_id}/{swarm_id}/input.json ./debug-input.json  
cat ./debug-input.json | jq .
```

## 32.2 Operational Troubleshooting

This release includes fixes for 3 critical stability issues (“Silent Killers”). Use this guide to diagnose production anomalies.

### 32.2.1 “Stuck” Workflows (Signal Mismatch)

| Aspect            | Details                                                             |
|-------------------|---------------------------------------------------------------------|
| <b>Symptom</b>    | Workflow is RUNNING but UI shows “Approved”                         |
| <b>Root Cause</b> | Signal ID sent by API does not match ID the workflow is waiting for |

**Diagnosis Steps:**

1. **Verify Signal Format:** Must be `human_decision_{decision_id}`
2. **Check Database:**

```
SELECT id, flyte_execution_id, flyte_node_id  
FROM pending_decisions  
WHERE status = 'pending'  
AND tenant_id = '<TENANT_ID>';
```

3. **Cross-Reference Flyte:** The `flyte_node_id` must match the node in the Flyte execution graph

**Resolution:**

If deadlocked, terminate the workflow via Flyte Console:

```
flytectl delete execution <EXECUTION_ID> \  
--project radiant \  
--domain production
```

### 32.2.2 “Zombie” Cache

| Aspect            | Details                                                                                          |
|-------------------|--------------------------------------------------------------------------------------------------|
| <b>Symptom</b>    | User rejects a plan, retries, and AI immediately returns the same rejected plan without thinking |
| <b>Root Cause</b> | Flyte caching returning stale results                                                            |

## Verification:

Ensure `think_tank_workflow.py` has the correct decorator:

```
# CORRECT - Forces fresh execution
@dynamic(cache=False)
def execute_swarm(agents: List[AgentConfig], task_data: TaskData) -> List[AgentResult]:
    ...

# WRONG - Will return cached (potentially rejected) results
@dynamic
def execute_swarm(agents: List[AgentConfig], task_data: TaskData) -> List[AgentResult]:
    ...
```

Files to Check: - `packages/flyte/workflows/think_tank_workflow.py`

**32.2.3 Emergency Manual Intervention** If the API layer is unavailable, Admins can manually unblock a workflow using `flytectl`:

```
flytectl update execution <EXECUTION_ID> \
    --project radiant \
    --domain production \
    --signal-id "human_decision_<DECISION_ID>" \
    --signal-value '{"resolution": "approved", "guidance": "Emergency Override via CLI"}'
```

Signal Value Schema:

```
{
    "resolution": "approved" | "rejected" | "modified",
    "guidance": "string - guidance for AI refinement",
    "resolved_by": "admin-user-id",
    "resolved_at": "2026-01-07T12:00:00Z"
}
```

## 32.3 Compliance & Security

### 32.3.1 Tenant Isolation (Strict RLS)

**Warning for DB Admins:** All tables are protected by Row-Level Security. Running a standard `SELECT *` as a superuser might return 0 rows or trigger an error depending on your client config.

Correct Query Pattern:

To query data manually, you must set the Tenant Context for your session:

```
BEGIN;
-- Must match a valid tenant UUID
SET app.tenant_id = '123e4567-e89b-12d3-a456-426614174000';

SELECT * FROM pending_decisions;

COMMIT;
```

```
-- Or use RESET to clear context
RESET app.tenant_id;

Protected Tables: - pending_decisions - decision_audit - decision_domain_config -
websocket_connections
```

**32.3.2 Audit Log Export** To export the decision audit trail for SOC2/HIPAA evidence:

```
SELECT
    da.created_at,
    da.actor_id,
    da.actor_type,
    da.action,
    da.details->>'resolution' as resolution,
    da.details->>'guidance' as guidance,
    pd.domain,
    pd.question
FROM decision_audit da
JOIN pending_decisions pd ON da.decision_id = pd.id
WHERE da.tenant_id = '<TENANT_ID>'
AND da.created_at > NOW() - INTERVAL '30 days'
ORDER BY da.created_at DESC;
```

Export to CSV:

```
psql "$DATABASE_URL" -c "COPY (
    SELECT
        created_at,
        actor_id,
        action,
        details->>'resolution' as resolution,
        details->>'guidance' as guidance
    FROM decision_audit
    WHERE tenant_id = '<TENANT_ID>'
    AND created_at > NOW() - INTERVAL '30 days'
) TO STDOUT WITH CSV HEADER" > audit_export.csv
```

**32.3.3 PHI Sanitization** All decision content is sanitized before human review to prevent PHI exposure:

| Pattern                | Replacement      |
|------------------------|------------------|
| SSN (XXX-XX-XXXX)      | [SSN REDACTED]   |
| Email addresses        | [EMAIL REDACTED] |
| Phone numbers          | [PHONE REDACTED] |
| Credit card numbers    | [CC REDACTED]    |
| Medical Record Numbers | [MRN REDACTED]   |
| ZIP codes (5-digit)    | [ZIP REDACTED]   |

Disable Sanitization (requires tenant config):

```

UPDATE decision_domain_config
SET sanitize_phi = false
WHERE tenant_id = '<TENANT_ID>'
AND domain = 'general';

```

**Warning:** Disabling PHI sanitization may violate HIPAA compliance. Only disable for non-healthcare tenants.

### 32.4 Related Sections

| Section                          | Relevance                         |
|----------------------------------|-----------------------------------|
| RADIANT Admin Guide - Section 48 | Full Mission Control architecture |
| RADIANT Admin Guide - Section 47 | Flyte state management            |
| RADIANT Admin Guide - Section 42 | Cato safety integration           |
| Section 30 - COS                 | SOFIAI routing                    |

## 33. Cognitive Platform Enhancements

### From Modern Orchestrator to Category-Defining Cognitive Platform

Version: 4.20.0 | Status: Strategic Roadmap

This section documents five strategic enhancements that transform Think Tank from a task execution engine into a self-evolving cognitive platform with an unassailable competitive moat.

### 33.1 Strategic Vision: Beyond Task Execution

**The Fundamental Shift** Most AI orchestration engines (LangChain, AutoGen, Enterprise Copilot) are **stateless**: they solve a problem, reset, and solve it again from scratch. They suffer from “Goldfish Memory.”

#### ORCHESTRATION PARADIGM COMPARISON

| TRADITIONAL (Stateless)                                                              | RADIANT (Cognitive)                                                                                               |
|--------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------|
| =====                                                                                | =====                                                                                                             |
| Task → Solve → Reset<br>↓<br>Task → Solve → Reset<br>↓<br>Task → Solve → Reset       | Task → Solve → LEARN → Evolve<br>↓<br>Next Task (with accumulated skills)<br>↓<br>System becomes expert over time |
| No memory between runs<br>Same mistakes repeated<br>Flat cost curve<br>Reactive only | Procedural memory persists<br>Self-correction from errors<br>Decreasing cost over time<br>Proactive monitoring    |

## The Five Strategic Moats

| Enhancement                                     | Category             | Impact              | Competitive Advantage |
|-------------------------------------------------|----------------------|---------------------|-----------------------|
| <b>Economic Governor</b><br><b>The Grimoire</b> | Cost Optimization    | -40% API costs      | Immediate ROI         |
|                                                 | Procedural           | +60%                | Lock-in effect        |
|                                                 | Memory               | accuracy over time  |                       |
|                                                 | Developer Experience | -80% debug time     | Power user magnet     |
| <b>Sentinels</b>                                | Autonomous Agents    | New revenue stream  | Market expansion      |
|                                                 | Quality Assurance    | -90% hallucinations | Trust differentiator  |

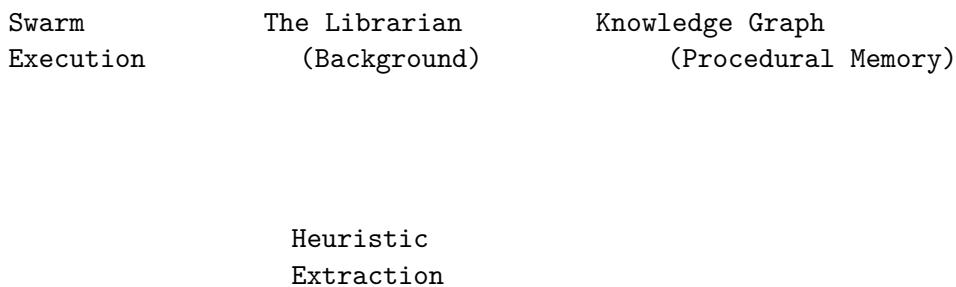
### 33.2 The Grimoire (Procedural Memory & Self-Correction)

**Problem Statement** If an agent struggles to write a valid SQL query for your schema today, it will struggle again tomorrow. RAG provides *facts*, but it doesn't provide *skills*. Current systems have:

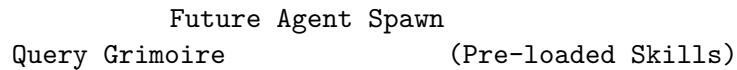
- **No skill retention** between sessions
- **Repeated mistakes** on similar tasks
- **No personalization** to tenant-specific patterns
- **Static performance** regardless of usage volume

**Solution: Write-Back Procedural Memory** The Grimoire is a tenant-isolated knowledge graph that captures **learned heuristics** from successful task executions, making the system smarter with every interaction.

#### THE GRIMOIRE ARCHITECTURE



- Pattern Match
  - Success Signal
  - Context Tags



**The Librarian Service** A background agent that analyzes every completed Flyte execution trace:

```
interface LearnedHeuristic {
  id: string;
  tenantId: string;
  category: 'sql_pattern' | 'api_usage' | 'code_style' | 'domain_knowledge' | 'user_preference';
  trigger: string; // When to apply this heuristic
  heuristic: string; // The learned rule
  confidence: number; // 0.0 - 1.0
  sourceExecutionId: string; // Flyte execution that taught this
  successCount: number; // Times this heuristic led to success
  failureCount: number; // Times this heuristic led to failure
  lastApplied: Date;
  tags: string[];
  embedding: number[]; // Vector for semantic search
}

// Example learned heuristics:
const exampleHeuristics = [
  {
    category: 'sql_pattern',
    trigger: 'querying sales table',
    heuristic: 'Always filter by is_deleted = false when querying the sales table',
    confidence: 0.95,
    tags: ['sql', 'sales', 'soft-delete']
  },
  {
    category: 'user_preference',
    trigger: 'generating Python code for user_123',
    heuristic: 'User prefers fully typed Python with dataclasses over dicts',
    confidence: 0.88,
    tags: ['python', 'typing', 'user_123']
  },
  {
    category: 'domain_knowledge',
    trigger: 'analyzing user interaction logs',
    heuristic: 'User frequently interacts with specific domain entities',
    confidence: 0.75,
    tags: ['domain', 'interaction', 'user']
  }
];
```

```

        trigger: 'medical terminology in oncology',
        heuristic: 'TNM staging must specify edition (e.g., AJCC 8th edition)',
        confidence: 0.92,
        tags: ['medical', 'oncology', 'staging']
    }
];

```

**Confidence Decay & Reinforcement** Heuristics evolve based on application outcomes:

| Event                         | Confidence Change     |
|-------------------------------|-----------------------|
| <b>Successful application</b> | +0.02 (max 0.99)      |
| <b>Failed application</b>     | -0.05 (min 0.30)      |
| <b>Weekly decay (unused)</b>  | -0.01                 |
| <b>User correction</b>        | New heuristic at 0.95 |
| <b>Below 0.30</b>             | Auto-archived         |

**Admin UI: Grimoire Management** **Location:** Admin Dashboard → Think Tank → Grimoire

| Tab                        | Description                                     |
|----------------------------|-------------------------------------------------|
| <b>Heuristics Browser</b>  | Search, filter, and view all learned heuristics |
| <b>Confidence Tuning</b>   | Adjust confidence thresholds and decay rates    |
| <b>Category Management</b> | Enable/disable heuristic categories             |
| <b>Audit Trail</b>         | View heuristic application history              |
| <b>Manual Entry</b>        | Add expert heuristics manually                  |

### Grimoire API Reference

Base: /api/thinktank/grimoire

|        |                  |                                        |
|--------|------------------|----------------------------------------|
| GET    | /heuristics      | List heuristics with filtering         |
| GET    | /heuristics/:id  | Get heuristic with application history |
| POST   | /heuristics      | Manually add a heuristic               |
| PATCH  | /heuristics/:id  | Update confidence/enabled status       |
| DELETE | /heuristics/:id  | Remove heuristic                       |
| GET    | /stats           | Grimoire statistics                    |
| POST   | /heuristics/bulk | Bulk import heuristics                 |

### 33.3 Time-Travel Debugging (Visual Forking)

**Problem Statement** An agent runs for 20 minutes, succeeds at 9 steps, but fails on step 10 due to a vague instruction. In current systems:

- You must **restart from Step 1**, wasting time and money
- **No way to edit** the context at a specific point
- **Lost compute costs** for successful early steps
- **Poor debugging experience** for complex workflows

**Solution: DVR Interface with Checkpoint Forking** Leverage Flyte's native checkpointing to build a time-travel debugging experience:

TIME-TRAVEL DEBUGGING INTERFACE

Execution: think\_tank\_swarm\_abc123  
Status: FAILED at Step 10      Total Duration: 18:42

1    2    3    4    5    6    7    8    9    10

[Timeline Slider]

Step 6: Code Generation  
Duration: 2:34      Model: claude-3-5-sonnet      Cost: \$0.08

Input Context: "Generate a Python function to calculate..."  
Output: def calculate\_quarterly\_revenue(transactions):...

[Edit Context]   [Fork From Here]   [View Full Trace]

## Fork Execution Service

```
interface ForkRequest {
    originalExecutionId: string;
    forkFromNodeId: string;
    contextModifications: {
        systemPromptAppend?: string;
        userPromptReplace?: string;
        variableOverrides?: Record<string, unknown>;
        modelOverride?: string;
    };
    forkedBy: string;
    reason?: string;
}

interface ForkResult {
    forkedExecutionId: string;
    forkedFromNode: string;
    estimatedSavings: {
        timeMinutes: number;
    }
}
```

```

    costUsd: number;
    tokensSkipped: number;
};

status: 'launched' | 'pending_approval';
}

```

**Admin UI: Time-Travel Debugger** **Location:** Admin Dashboard → Think Tank → Executions → [Select] → Time Travel

| Feature                   | Description                                          |
|---------------------------|------------------------------------------------------|
| <b>Timeline View</b>      | Visual representation of execution nodes with status |
| <b>Node Inspector</b>     | View input/output, model, tokens, cost for each node |
| <b>Context Editor</b>     | Modify system prompt, user prompt, variables         |
| <b>Fork Button</b>        | Create new execution from selected checkpoint        |
| <b>Comparison View</b>    | Side-by-side diff of original vs forked execution    |
| <b>Savings Calculator</b> | Real-time estimate of time/cost savings              |

### Time-Travel API Reference

Base: /api/thinktank/time-travel

|                                                  |                                         |
|--------------------------------------------------|-----------------------------------------|
| GET /executions/:id/timeline                     | Get execution timeline with checkpoints |
| GET /executions/:id/checkpoints/:nodeId          | Get checkpoint details                  |
| POST /executions/:id/checkpoints/:nodeId/preview | Preview modifications                   |
| POST /executions/:id/fork                        | Fork execution from checkpoint          |
| GET /executions/:originalId/compare/:forkedId    | Compare executions                      |
| GET /executions/:id/forks                        | Get fork history                        |

---

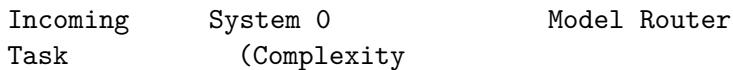
### 33.4 The Economic Governor (Model Arbitrage)

#### Problem Statement

- Using **GPT-4o** for every **sub-task** is financially ruinous
- Using **Llama-3-8b** for everything leads to errors
- **No visibility** into which tasks need expensive models
- **Flat cost curve** regardless of task complexity

**Solution: Predictive Cost Routing** A “System 0” pre-dispatch analysis that routes tasks to the optimal model based on complexity:

#### ECONOMIC GOVERNOR ARCHITECTURE



|                                                                                                                                             |                                                                                                               |
|---------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------|
| <p>Estimator)</p> <ul style="list-style-type: none"> <li>• Token count</li> <li>• Task type</li> <li>• Domain</li> <li>• History</li> </ul> | <p>Score 1-3 → Haiku/Llama<br/>           Score 4-7 → Sonnet/GPT-4o-m<br/>           Score 8-10 → Opus/O1</p> |
| <p>Savings Tracker<br/>           "Saved \$4.20 via smart<br/>           Complexity Score<br/>           (1-10)</p>                         | <p>routing this query"</p>                                                                                    |

## Complexity Factors

| Factor               | Weight | Description                                  |
|----------------------|--------|----------------------------------------------|
| <b>Prompt Length</b> | 15%    | Token count estimate                         |
| <b>Task Type</b>     | 25%    | Summarization (2) → Multi-step reasoning (9) |
| <b>Domain</b>        | 20%    | General (3) → Medical/Scientific (8)         |
| <b>Keywords</b>      | 20%    | Complexity indicator words                   |
| <b>Structure</b>     | 20%    | Code blocks, lists, nested requirements      |

## Model Tier Mapping

| Tier            | Score Range | Models                               | Use Cases                      |
|-----------------|-------------|--------------------------------------|--------------------------------|
| <b>Economy</b>  | 1-3         | Haiku,<br>GPT-4o-mini,<br>Llama-3-8b | Simple Q&A,<br>summarization   |
| <b>Standard</b> | 4-7         | Sonnet, GPT-4o                       | Code generation, analysis      |
| <b>Premium</b>  | 8-10        | Opus, O1-preview                     | Complex reasoning,<br>planning |

**Admin UI: Economic Governor Dashboard**   **Location:** Admin Dashboard → Think Tank → Economic Governor

| Panel                       | Description                             |
|-----------------------------|-----------------------------------------|
| <b>Savings Overview</b>     | Total savings this period, trend chart  |
| <b>Model Distribution</b>   | Pie chart of model usage by tier        |
| <b>Complexity Histogram</b> | Distribution of task complexity scores  |
| <b>Routing Rules</b>        | Configure tier thresholds and overrides |
| <b>Budget Alerts</b>        | Set spending alerts and caps            |

## Economic Governor API Reference

Base: /api/thinktank/economic-governor

|                           |                                  |
|---------------------------|----------------------------------|
| GET /savings?period=month | Get savings dashboard            |
| GET /routing-rules        | Get current routing rules        |
| PUT /routing-rules        | Update routing configuration     |
| POST /analyze-complexity  | Analyze specific task complexity |
| GET /model-usage          | Get model usage breakdown        |
| POST /budget-alert        | Configure budget alerts          |

## Configuration Options

```
interface EconomicGovernorConfig {  
    economyThreshold: number;      // Default: 3  
    standardThreshold: number;     // Default: 7  
    forceModelOverrides: {  
        [taskType: string]: string;  // e.g., "legal_analysis" → "opus"  
    };  
    budgetCap: {  
        daily: number;  
        monthly: number;  
    };  
    alertThresholds: {  
        warningPercent: number;    // Default: 80  
        criticalPercent: number;   // Default: 95  
    };  
}
```

---

### 33.5 Sentinel Agents (Event-Driven Autonomy)

**Problem Statement** Current agents are **reactive only**—they wait for user input. A true cognitive platform should be **proactive**:

- **No autonomous monitoring** capabilities
- **No long-lived workflows** that persist between sessions
- **No event-driven triggers** for automated responses
- **Missed opportunities** for preventive action

**Solution: Long-Lived Hibernating Workflows** Allow agents to set up persistent monitors that wake up when conditions are met:

#### SENTINEL AGENT ARCHITECTURE

User: "Monitor server logs. If Error 500 spikes, analyze and alert me."

## 1. SETUP PHASE

```

Parse           Generate           Configure
Instruction     EventBridge        Hibernate State
                    Rule

```

## 2. HIBERNATION PHASE (Days/Weeks/Months)

Flyte Workflow: HIBERNATE state (wait\_for\_signal)

- Zero compute cost while waiting
- State preserved in S3

Event Fires!

## 3. REHYDRATION PHASE

```

EventBridge      Signal Lambda      Workflow
Triggers         Sends Signal       Resumes

```

## 4. ANALYSIS & ALERT PHASE

```

Pull Latest      Swarm Analysis    Alert User
Context          (Root Cause)      via Slack/SMS

```

## Sentinel Trigger Types

| Type                             | Description                       | Example              |
|----------------------------------|-----------------------------------|----------------------|
| <code>cloudwatch_alarm</code>    | AWS CloudWatch alarm state change | CPU > 90%            |
| <code>eventbridge_pattern</code> | Custom EventBridge event pattern  | CodePipeline failure |
| <code>webhook</code>             | External HTTP webhook             | GitHub push          |
| <code>schedule</code>            | Cron or rate expression           | Every hour           |
| <code>metric_threshold</code>    | Custom metric threshold           | Error rate > 5%      |

## Sentinel Action Types

| Type                         | Description                  | Example             |
|------------------------------|------------------------------|---------------------|
| <code>swarm_analysis</code>  | Run AI swarm for analysis    | Root cause analysis |
| <code>notification</code>    | Send alert via channels      | Slack + Email       |
| <code>remediation</code>     | Execute automated fix        | Scale up instances  |
| <code>custom_workflow</code> | Launch custom Flyte workflow | Data pipeline       |

**Admin UI: Sentinel Management** **Location:** Admin Dashboard → Think Tank → Sentinels

| Tab                     | Description                               |
|-------------------------|-------------------------------------------|
| <b>Active Sentinels</b> | List of hibernating sentinels with status |
| <b>Create Sentinel</b>  | Natural language sentinel configuration   |
| <b>Trigger History</b>  | Log of all sentinel activations           |
| <b>Analytics</b>        | Sentinel effectiveness metrics            |

## Sentinel API Reference

Base: /api/thinktank/sentinels

|        |               |                                        |
|--------|---------------|----------------------------------------|
| GET    | /             | List all sentinels                     |
| POST   | /             | Create new sentinel (natural language) |
| GET    | /:id          | Get sentinel details                   |
| PATCH  | /:id          | Update sentinel (pause/resume)         |
| DELETE | /:id          | Delete sentinel and EventBridge rule   |
| GET    | /:id/triggers | Get trigger history                    |
| POST   | /:id/test     | Test sentinel with mock event          |
| GET    | /analytics    | Sentinel effectiveness metrics         |

## Example Sentinel Configurations

```
// Monitor for deployment failures
{
  name: "Deployment Monitor",
  triggerInstruction: "Watch for failed CodePipeline deployments",
  actionInstruction: "Analyze the failure logs and suggest fixes, then alert #devops on Slack"
}

// Proactive cost monitoring
{
  name: "Cost Anomaly Detector",
  triggerInstruction: "If AWS daily costs exceed $500 or increase 50% from yesterday",
  actionInstruction: "Identify the cost drivers and alert finance@company.com"
}

// Security sentinel
{
  name: "Security Scanner",
  triggerInstruction: "Every 6 hours, or when a new ECR image is pushed",
  actionInstruction: "Scan for vulnerabilities and create a report"
}
```

---

### 33.6 The Council of Rivals (Adversarial Consensus)

**Problem Statement** LLMs suffer from:

- **Hallucinations** (confident but wrong answers)
- **Sycophancy** (agreeing with user's incorrect premises)
- **Blind spots** (missing edge cases)
- **No self-verification** (unable to catch own errors)

**Solution: Structured Adversarial Debate** Force multiple models to argue against each other before synthesizing a final answer:

#### COUNCIL OF RIVALS ARCHITECTURE

User Query: "Should we migrate from PostgreSQL to MongoDB?"

#### THE COUNCIL

| ADVOCATE<br>(Claude)                       | CRITIC<br>(GPT-4o)                               | PRAGMATIST<br>(Llama-70b)                    |
|--------------------------------------------|--------------------------------------------------|----------------------------------------------|
| "Migrate! Schema flexibility is worth it." | "Don't! You lose ACID, joins become nightmares." | "It depends on your data access patterns..." |

#### ROUND 2: Cross-Examine Each Other

ARBITER  
(Opus)

Synthesize  
final answer  
with confidence

Final Output: "For your use case (heavy joins, ACID requirements),

stay with PostgreSQL. Migration cost would outweigh benefits."  
 Confidence: 87% | Dissent: Advocate (13%)

## Council Roles

| Role              | Purpose                          | Default Model |
|-------------------|----------------------------------|---------------|
| <b>Advocate</b>   | Argues FOR the proposed solution | Claude Sonnet |
| <b>Critic</b>     | Argues AGAINST, finds weaknesses | GPT-4o        |
| <b>Pragmatist</b> | Considers practical constraints  | Llama-70b     |
| <b>Arbiter</b>    | Synthesizes final verdict        | Claude Opus   |

## Council Configuration

```
interface CouncilConfig {
    enabled: boolean;
    triggerComplexity: number;           // Min complexity score (default: 7)
    triggerDomains: string[];            // Domains requiring council (e.g., ['legal', 'medical'])
    roles: {
        advocate: { model: string; temperature: number };
        critic: { model: string; temperature: number };
        pragmatist: { model: string; temperature: number };
        arbiter: { model: string; temperature: number };
    };
    rounds: number;                     // Cross-examination rounds (default: 2)
    confidenceThreshold: number;        // Min confidence for consensus (default: 0.75)
    includeDissentReport: boolean;      // Show minority opinions
}
```

## Council Output Format

```
interface CouncilVerdict {
    question: string;
    verdict: string;
    confidence: number;                // 0.0 - 1.0
    consensusLevel: 'unanimous' | 'majority' | 'split';
    arguments: {
        advocate: { position: string; keyPoints: string[] };
        critic: { position: string; keyPoints: string[] };
        pragmatist: { position: string; keyPoints: string[] };
    };
    crossExamination: {
        round: number;
        challenges: Array<{
            from: string;
            to: string;
        }>;
    };
}
```

```

        challenge: string;
        response: string;
    }>;
}[] ;
arbiterReasoning: string;
dissent?: {
    role: string;
    position: string;
    confidence: number;
};
}

```

**Admin UI: Council of Rivals** **Location:** Admin Dashboard → Think Tank → Council of Rivals

| Tab                    | Description                                |
|------------------------|--------------------------------------------|
| <b>Configuration</b>   | Enable/disable, configure roles and models |
| <b>Trigger Rules</b>   | Set complexity/domain triggers             |
| <b>Session History</b> | View past council deliberations            |
| <b>Analytics</b>       | Consensus rates, dissent patterns          |

## Council API Reference

Base: /api/thinktank/council

|                   |                                       |
|-------------------|---------------------------------------|
| GET /config       | Get council configuration             |
| PUT /config       | Update council configuration          |
| POST /deliberate  | Manually trigger council deliberation |
| GET /sessions     | List past deliberation sessions       |
| GET /sessions/:id | Get full deliberation transcript      |
| GET /analytics    | Council effectiveness metrics         |

## 33.7 Implementation Roadmap

### Priority Matrix

| Phase     | Enhancement           | Effort | Impact     | Priority  |
|-----------|-----------------------|--------|------------|-----------|
| <b>Q1</b> | Economic Governor     | Medium | High (ROI) | <b>P0</b> |
| <b>Q1</b> | The Grimoire          | High   | Very High  | <b>P0</b> |
| <b>Q2</b> | Time-Travel Debugging | Medium | Medium     | <b>P1</b> |
| <b>Q2</b> | Council of Rivals     | Medium | High       | <b>P1</b> |
| <b>Q3</b> | Sentinel Agents       | High   | High       | <b>P2</b> |

### Recommended Implementation Order

1. **Economic Governor** (Week 1-3)
  - Immediate cost savings
  - Simple routing logic
  - Foundation for other features
2. **The Grimoire** (Week 2-6)
  - Highest long-term value
  - Leverages existing Flyte infrastructure
  - Creates lock-in effect
3. **Time-Travel Debugging** (Week 5-8)
  - Builds on Flyte checkpointing
  - Power user feature
  - Developer experience differentiator
4. **Council of Rivals** (Week 7-10)
  - Quality improvement
  - Trust building
  - Requires multi-model orchestration
5. **Sentinel Agents** (Week 9-14)
  - Most complex
  - Requires EventBridge integration
  - New revenue opportunities

## Dependencies

Economic Governor

The Grimoire              Time-Travel

Council of Rivals

Sentinel Agents (independent, can parallel)

---

## 33.8 Database Schema

```
-- Grimoire tables
CREATE TABLE grimoire_heuristics (
  id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
  tenant_id UUID NOT NULL REFERENCES tenants(id),
  category VARCHAR(50) NOT NULL,
  trigger TEXT NOT NULL,
  heuristic TEXT NOT NULL,
  confidence DECIMAL(3,2) NOT NULL DEFAULT 0.70,
  source_execution_id VARCHAR(255),
  success_count INTEGER DEFAULT 0,
  failure_count INTEGER DEFAULT 0,
  last_applied TIMESTAMPTZ,
  tags TEXT[],
  embedding VECTOR(1536),
```

```

enabled BOOLEAN DEFAULT true,
archived_at TIMESTAMPTZ,
created_at TIMESTAMPTZ DEFAULT NOW(),
updated_at TIMESTAMPTZ DEFAULT NOW()
);

CREATE INDEX idx_grimoire_tenant_category ON grimoire_heuristics(tenant_id, category);
CREATE INDEX idx_grimoire_embedding ON grimoire_heuristics USING ivfflat (embedding vector_cosine);

-- Economic Governor tables
CREATE TABLE economic_governor_savings (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    tenant_id UUID NOT NULL REFERENCES tenants(id),
    execution_id VARCHAR(255) NOT NULL,
    actual_model VARCHAR(100) NOT NULL,
    actual_cost DECIMAL(10,6) NOT NULL,
    baseline_model VARCHAR(100) NOT NULL,
    baseline_cost DECIMAL(10,6) NOT NULL,
    savings DECIMAL(10,6) NOT NULL,
    savings_percent DECIMAL(5,2) NOT NULL,
    complexity_score DECIMAL(3,1) NOT NULL,
    created_at TIMESTAMPTZ DEFAULT NOW()
);

CREATE INDEX idx_savings_tenant_date ON economic_governor_savings(tenant_id, created_at);

-- Sentinel tables
CREATE TABLE sentinels (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    tenant_id UUID NOT NULL REFERENCES tenants(id),
    user_id UUID NOT NULL REFERENCES users(id),
    name VARCHAR(255) NOT NULL,
    description TEXT,
    trigger_type VARCHAR(50) NOT NULL,
    trigger_config JSONB NOT NULL,
    action_type VARCHAR(50) NOT NULL,
    action_config JSONB NOT NULL,
    status VARCHAR(50) DEFAULT 'hibernating',
    flyte_execution_id VARCHAR(255),
    eventbridge_rule_arn TEXT,
    trigger_count INTEGER DEFAULT 0,
    max_triggers INTEGER,
    expires_at TIMESTAMPTZ,
    last_triggered TIMESTAMPTZ,
    created_at TIMESTAMPTZ DEFAULT NOW(),
    updated_at TIMESTAMPTZ DEFAULT NOW()
);

```

```

CREATE INDEX idx_sentinels_tenant_status ON sentinels(tenant_id, status);

-- Council of Rivals tables
CREATE TABLE council_sessions (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    tenant_id UUID NOT NULL REFERENCES tenants(id),
    user_id UUID NOT NULL REFERENCES users(id),
    question TEXT NOT NULL,
    verdict TEXT,
    confidence DECIMAL(3,2),
    consensus_level VARCHAR(20),
    arguments JSONB,
    cross_examination JSONB,
    arbiter_reasoning TEXT,
    dissent JSONB,
    duration_ms INTEGER,
    total_cost DECIMAL(10,6),
    created_at TIMESTAMPTZ DEFAULT NOW()
);

CREATE INDEX idx_council_tenant_date ON council_sessions(tenant_id, created_at);

-- Time-Travel checkpoints
CREATE TABLE execution_checkpoints (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    tenant_id UUID NOT NULL REFERENCES tenants(id),
    execution_id VARCHAR(255) NOT NULL,
    node_id VARCHAR(255) NOT NULL,
    workflow_name VARCHAR(255) NOT NULL,
    s3_uri TEXT NOT NULL,
    state_hash VARCHAR(64),
    metadata JSONB,
    created_at TIMESTAMPTZ DEFAULT NOW(),
    UNIQUE(execution_id, node_id)
);

CREATE TABLE execution_forks (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    tenant_id UUID NOT NULL REFERENCES tenants(id),
    original_execution_id VARCHAR(255) NOT NULL,
    forked_execution_id VARCHAR(255) NOT NULL,
    fork_node_id VARCHAR(255) NOT NULL,
    modifications JSONB NOT NULL,
    forked_by UUID REFERENCES users(id),
    reason TEXT,
    time_saved_minutes INTEGER,
    cost_saved DECIMAL(10,6),
    tokens_skipped INTEGER,
);

```

```

created_at TIMESTAMPTZ DEFAULT NOW()
);

-- RLS policies
ALTER TABLE grimoire_heuristics ENABLE ROW LEVEL SECURITY;
ALTER TABLE economic_governor_savings ENABLE ROW LEVEL SECURITY;
ALTER TABLE sentinels ENABLE ROW LEVEL SECURITY;
ALTER TABLE council_sessions ENABLE ROW LEVEL SECURITY;
ALTER TABLE execution_checkpoints ENABLE ROW LEVEL SECURITY;
ALTER TABLE execution_forks ENABLE ROW LEVEL SECURITY;

CREATE POLICY tenant_isolation ON grimoire_heuristics
    USING (tenant_id = current_setting('app.current_tenant_id')::UUID);
CREATE POLICY tenant_isolation ON economic_governor_savings
    USING (tenant_id = current_setting('app.current_tenant_id')::UUID);
CREATE POLICY tenant_isolation ON sentinels
    USING (tenant_id = current_setting('app.current_tenant_id')::UUID);
CREATE POLICY tenant_isolation ON council_sessions
    USING (tenant_id = current_setting('app.current_tenant_id')::UUID);
CREATE POLICY tenant_isolation ON execution_checkpoints
    USING (tenant_id = current_setting('app.current_tenant_id')::UUID);
CREATE POLICY tenant_isolation ON execution_forks
    USING (tenant_id = current_setting('app.current_tenant_id')::UUID);

```

---

### 33.9 API Reference

#### Complete Endpoint Summary

| Service           | Base Path                        | Key Endpoints                        |
|-------------------|----------------------------------|--------------------------------------|
| Grimoire          | /api/thinktank/grimoire          | CRUD heuristics, stats, bulk import  |
| Time-Travel       | /api/thinktank/time-travel       | Timeline, checkpoints, fork, compare |
| Economic Governor | /api/thinktank/economic-governor | Savings, routing rules, analyze      |
| Sentinels         | /api/thinktank/sentinels         | CRUD sentinels, triggers, test       |
| Council           | /api/thinktank/council           | Config, deliberate, sessions         |

---

### 33.10 Configuration

#### Tenant-Level Feature Flags

```

interface CognitiveEnhancementsConfig {
  grimoire: {
    enabled: boolean;
    autoExtract: boolean;
    minConfidenceThreshold: number;
    decayEnabled: boolean;
  }
}

```

```

    };
    timeTravel: {
        enabled: boolean;
        maxCheckpointsPerExecution: number;
        retentionDays: number;
    };
    economicGovernor: {
        enabled: boolean;
        economyThreshold: number;
        standardThreshold: number;
        budgetCapDaily: number;
        budgetCapMonthly: number;
    };
    sentinels: {
        enabled: boolean;
        maxActiveSentinels: number;
        allowedTriggerTypes: string[];
    };
    council: {
        enabled: boolean;
        triggerComplexity: number;
        triggerDomains: string[];
        rounds: number;
    };
}
}

```

---

### 33.11 Troubleshooting

#### Common Issues

| Issue                   | Cause                     | Resolution                      |
|-------------------------|---------------------------|---------------------------------|
| Grimoire not learning   | Auto-extract disabled     | Enable in tenant config         |
| Fork fails              | Checkpoint expired        | Increase retention period       |
| Governor routes wrong   | Stale complexity model    | Retrain on recent data          |
| Sentinel not triggering | EventBridge rule disabled | Check AWS console               |
| Council timeout         | Too many rounds           | Reduce cross-examination rounds |

#### Debug Commands

```

# Check Grimoire heuristic count
curl -X GET "https://api.radiant.ai/thinktank/grimoire/stats" \
-H "Authorization: Bearer $TOKEN"

# Test Economic Governor routing
curl -X POST "https://api.radiant.ai/thinktank/economic-governor/analyze-complexity" \

```

```

-H "Authorization: Bearer $TOKEN" \
-d '{"prompt": "Write a complex SQL query", "taskType": "code_generation"}'

# Check Sentinel status
curl -X GET "https://api.radiant.ai/thinktank/sentinels" \
-H "Authorization: Bearer $TOKEN"

# View Council session
curl -X GET "https://api.radiant.ai/thinktank/council/sessions/{sessionId}" \
-H "Authorization: Bearer $TOKEN"

```

### 33.12 Related Sections

| Section                                          | Relevance                          |
|--------------------------------------------------|------------------------------------|
| Section 23 - Predictive Coding                   | Foundation for Grimoire learning   |
| Section 30 - COS                                 | SOFIAI integration                 |
| Section 32 - Swarm Orchestration                 | Flyte integration for all features |
| <a href="#">RADIANT Admin Guide - Section 47</a> | Flyte checkpointing                |

## 34. Orchestration Workflow Methods Reference

**Location:** Admin Dashboard → Think Tank → Orchestration → Methods

This section documents the complete **70+ orchestration workflow methods** available in Think Tank. Each method has a **display name** (user-friendly) and **scientific name** (formal/academic reference).

### 34.1 Method Categories Overview

| Category            | Methods | Purpose                                              |
|---------------------|---------|------------------------------------------------------|
| <b>Generation</b>   | 3       | Generate responses with various reasoning strategies |
| <b>Evaluation</b>   | 6       | Critique, judge, and score outputs                   |
| <b>Synthesis</b>    | 5       | Combine multiple responses into unified outputs      |
| <b>Verification</b> | 8       | Fact-check and verify claims                         |
| <b>Debate</b>       | 5       | Multi-agent deliberation and argumentation           |
| <b>Aggregation</b>  | 4       | Vote, blend, and aggregate responses                 |
| <b>Reasoning</b>    | 2       | Problem decomposition and reflection                 |
| <b>Routing</b>      | 6       | Dynamic model selection and cascading                |

| Category             | Methods | Purpose                                |
|----------------------|---------|----------------------------------------|
| <b>Collaboration</b> | 5       | Multi-agent coordination               |
| <b>Uncertainty</b>   | 6       | Confidence estimation and calibration  |
| <b>Hallucination</b> | 3       | Detect and prevent hallucinations      |
| <b>Human-in-Loop</b> | 3       | Human oversight and review             |
| <b>Neural</b>        | 1       | Cato-integrated neural decision engine |

## 34.2 Complete Methods Reference

### Generation Methods

| Display Name       | Scientific Name             | Code              | Description                                  |
|--------------------|-----------------------------|-------------------|----------------------------------------------|
| Generate           | Basic Generation            | GENERATE_RESPONSE | Standard response generation                 |
| Think Step-by-Step | Chain-of-Thought Generation | GENERATE_WITH_COT | Thinking before answering (+20-40% accuracy) |
| Refine             | Iterative Refinement        | REFINE_RESPONSE   | Improve based on feedback                    |

### Evaluation Methods

| Display Name         | Scientific Name      | Code              | Accuracy Gain                     |
|----------------------|----------------------|-------------------|-----------------------------------|
| Critique             | Critical Evaluation  | CRITIQUE_RESPONSE | Identifies flaws                  |
| Judge                | Comparative Judgment | JUDGE_RESPONSES   | Evaluates multiple outputs        |
| Multi-Judge Panel    | Panel of LLMs (PoLL) | POLL_JUDGE        | -40-60% single-model bias         |
| Structured Scoring   | G-Eval NLG Framework | G_EVAL            | 0.5+ human correlation            |
| Head-to-Head Compare | Pairwise Preference  | PAIRWISE_PREF     | Vulnerable for subtle differences |
| Side-by-Side Compare | Comparative Analysis | COMPARE_ANALYSIS  | Decision clarity                  |

### Synthesis Methods

| Display Name        | Scientific Name          | Code                 | Accuracy Gain        |
|---------------------|--------------------------|----------------------|----------------------|
| Synthesize          | Multi-Response Synthesis | SYNTHESIZE_RESPONSES | Best parts           |
| Consensus           | Consensus Aggregation    | BUILD_CONSENSUS      | Alignment points     |
| Layered Synthesis   | Mixture of Agents (MoA)  | MOA_LAYERS           | +8% over GPT-4o      |
| Combine & Summarize | Multi-Source Synthesis   | MULTI_SOURCE_SYNTH   | Coverage             |
| Rank & Merge        | LLM-Blender Fusion       | LLM_BLENDER          | +12% over best model |

## Verification Methods

| Display Name          | Scientific Name         | Code           | Accuracy Gain           |
|-----------------------|-------------------------|----------------|-------------------------|
| Fact Check            | Factual Verification    | VERIFY_FACT    | Extract & verify claims |
| Step Verification     | Process Reward Model    | PROCESS_REWARD | on MATH                 |
| Internal Consistency  | SelfCheckGPT            | SELFCHECK_GPT  | 25% hallucination F1    |
| Source Attribution    | Citation Verification   | CITE_VERIFY    | +40% citation accuracy  |
| Logic-Based Check     | Zero-Shot Natural Logic | NATURAL_LOGIC  | .96 accuracy points     |
| Combined Verification | UniFact Unified         | UNIFACT        | +20% comprehensive      |
| Internal State Check  | EigenScore              | EIGENSCORE     | Hidden state analysis   |
| Re-Query Consistency  | Iterative Prompting     | REQUERY_CHECK  | Black-box detection     |

## Debate Methods

| Display Name         | Scientific Name        | Code               | Accuracy Gain                 |
|----------------------|------------------------|--------------------|-------------------------------|
| Challenge            | Adversarial Challenge  | GENERATE_CHALLENGE | arguments                     |
| Defend               | Position Defense       | DEFEND_POSITION    | Répond to challenges          |
| Efficient Debate     | Sparse Topology Debate | SPARSE_DEBATE      | -60% cost, <5% quality loss   |
| Attack & Support Map | ArgLLMs Bipolar        | ARG_MAPPING        | +35% structured argumentation |
| Human-AI Panel       | HAH-Delphi Hybrid      | HAH_DELPHI         | >90% expert coverage          |
| Confidence-Weighted  | ReConcile Consensus    | RECONCILE_WEIGHTED | 25% diverse ensembles         |

## Aggregation Methods

| Display Name        | Scientific Name      | Code               | Accuracy Gain      |
|---------------------|----------------------|--------------------|--------------------|
| Vote                | Majority Aggregation | MAJORITY_VOTE      | Most common answer |
| Weight              | Weighted Aggregation | WEIGHTED_AGGREGATE | rice-weighted      |
| Multi-Sample Voting | Self-Consistency     | SELF_CONSISTENCY   | 9% on GSM8K        |
| Ranked Choice       | GEDI Electoral CDM   | GEDI_VOTE          | +30% consensus     |

## Routing Methods

| Display Name           | Scientific Name     | Code              | Accuracy Gain       | Implementation              |
|------------------------|---------------------|-------------------|---------------------|-----------------------------|
| Classify               | Task Classification | DETECT_TASKTYPE   | complexity          | Keyword analysis            |
| Route                  | Model Selection     | SELECT_BEST_MODEL | model               | Capability matching         |
| Smart Selection        | RouteLLM Adaptive   | ROUTELLM          | -50% cost, <3% loss | Learned router              |
| Progressive Escalation | FrugalGPT Cascade   | FRUGAL CASCADE    | cost maintained     | Confidence-based escalation |
| Budget-Aware           | Pareto Routing      | PARETO_ROUTE      | imal trade-off      | Pareto frontier calculation |

| Display Name          | Scientific Name      | Code         | Accuracy Gain     | Implementation                                        |
|-----------------------|----------------------|--------------|-------------------|-------------------------------------------------------|
| Smart Cost Escalation | C3PO Self-Supervised | C3PO_CASCADE | cost, +2% quality | Difficulty prediction + tiered cascade (NeurIPS 2024) |
| Self-Routing          | AutoMix POMDP        | AUTOMIX      | Self-improving    | -greedy exploration + self-verification (Nov 2025)    |

**Implementation Notes:** - **Pareto Routing:** Computes Pareto frontier across quality/latency/cost, selects optimal model within budget constraints. - **C3PO Cascade:** Predicts query difficulty using prompt features, starts at appropriate tier, cascades up if confidence insufficient. - **AutoMix POMDP:** Uses POMDP belief state for model selection with -greedy exploration and self-verification for quality assurance.

## Collaboration Methods

| Display Name         | Scientific Name         | Code          | Accuracy Gain          |
|----------------------|-------------------------|---------------|------------------------|
| No-Comm Coordination | ECON Bayesian Nash      | ECON_NASH     | +11.2% coordination    |
| Fair Merge           | Token Auction           | TOKEN_AUCTION | Nir multi-stakeholder  |
| Logic & Solve        | Logic-LM Neuro-Symbolic | LOGIC_LM      | +39.2% over standard   |
| Generate & Verify    | LLM-Modulo Framework    | LLM_MODULO    | 12%→93.9% plan success |
| Auto-Discover        | AFlow MCTS Discovery    | AFLOW_MCTS    | Beats human designs    |

## Uncertainty Methods

| Display Name              | Scientific Name           | Code             | Accuracy Gain             | Implementation                    |
|---------------------------|---------------------------|------------------|---------------------------|-----------------------------------|
| Meaning-Based Calibration | Semantic Entropy          | SEMANTIC_ENTROPY | 0.79-0.87                 | NLI clustering                    |
| Confidence Estimation     | Calibrated Estimation     | CALIBRATE_CONF   | CONF5%                    | Temperature scaling               |
| Agreement Scoring         | Consistency UQ            | CONSISTENCY_UQ   | SYNTH, effective          | Multi-sample agreement            |
| Fast Uncertainty          | SE Probes (Logprob-based) | SE_PROBES        | 300x faster, 90% accuracy | Token logprob entropy (ICML 2024) |
| Detailed Score            | Kernel Language Entropy   | KERNEL_ENTROPY   | grained                   | Embedding KDE (NeurIPS 2024)      |
| Guaranteed Bounds         | Conformal Prediction      | CONFORMAL_PRED   | PREDICAL guarantees       | Coverage calibration              |

**Implementation Notes:** - **SE Probes:** Uses OpenAI logprobs API to compute per-token Shannon entropy. Approximates hidden state probes without model internals access. - **Kernel Entropy:** Generates embeddings via `text-embedding-3-small`, applies Gaussian KDE with Silverman bandwidth, returns density-based entropy.

## Hallucination Detection (NEW)

| Display Name        | Scientific Name        | Code           | Accuracy Gain               |
|---------------------|------------------------|----------------|-----------------------------|
| Fact-Check Scanner  | Multi-Method Detection | MULTI_HALLUE   | 1.0.85+ 0.85+               |
| Mutation Testing    | MetaQA Metamorphic     | METAQA         | +30% subtle inconsistencies |
| Source Verification | Factual Grounding      | FACTUAL_GROUND | 15% grounding accuracy      |

## Human-in-the-Loop (NEW)

| Display Name       | Scientific Name    | Code          | Purpose                        |
|--------------------|--------------------|---------------|--------------------------------|
| Human Review Queue | HITL Review System | HITL REVIEW   | +90% critical error prevention |
| Multi-Level Review | Tiered Evaluation  | TIERED_EVAL   | Efficient human resources      |
| Smart Sampling     | Active Learning    | ACTIVE_SAMPLE | -60% labeling efficiency       |

## Neural/ML Methods (NEW)

| Display Name    | Scientific Name             | Code        | Description                                                                           |
|-----------------|-----------------------------|-------------|---------------------------------------------------------------------------------------|
| Neural Decision | Cato Neural Decision Engine | CATO_NEURAL | Integrates Cato safety pipeline with consciousness affect state and predictive coding |

### 34.3 System vs User Methods

All orchestration methods have an `isSystemMethod` flag:

| Type                 | Can Edit Parameters | Can Edit Definition | Can Delete |
|----------------------|---------------------|---------------------|------------|
| <b>System Method</b> | Yes                 | No                  | No         |
| <b>User Method</b>   | Yes                 | Yes                 | Yes        |

**System methods** are the 70+ built-in methods documented above. Administrators can:

- Enable/disable system methods
- Modify default parameters
- View execution metrics

**User methods** (future feature) will allow tenants to create custom methods with their own prompts or code references.

## 34.4 User Workflow Templates

**Location:** Admin Dashboard → Think Tank → Workflow Templates

Users can create custom workflows by:

1. **Creating from scratch** - Add methods step by step
2. **Basing on system workflow** - Start from one of 49 system patterns
3. **Customizing parameters** - Override default parameters per step
4. **Sharing with team** - Make templates available to organization

### Template Structure

```
interface UserWorkflowTemplate {  
    templateId: string;  
    templateName: string;  
    templateDescription: string;  
    baseWorkflowCode?: string;  
    steps: Array<{  
        stepOrder: number;  
        methodCode: string;  
        displayName: string;  
        parameters: Record<string, unknown>;  
        condition?: string;  
        isEnabled: boolean;  
    }>;  
    category: string;  
    tags: string[];  
    isShared: boolean;  
    timesUsed: number;  
}
```

### API Endpoints

| Endpoint                                             | Method | Description                        |
|------------------------------------------------------|--------|------------------------------------|
| /api/admin/orchestration/userTemplates               | GET    | List user templates (own + shared) |
| /api/admin/orchestration/userTemplates               | POST   | Create template                    |
| /api/admin/orchestration/userTemplates/:id           | GET    | Get template details               |
| /api/admin/orchestration/userTemplates/:id           | PATCH  | Update template (owner only)       |
| /api/admin/orchestration/userTemplates/:id           | DELETE | Delete template (owner only)       |
| /api/admin/orchestration/userTemplates/:id/share     | POST   | Toggle team sharing                |
| /api/admin/orchestration/userTemplates/:id/duplicate | POST   | Duplicate template                 |

### Method Management Endpoints:

| Endpoint                               | Method | Description                                              |
|----------------------------------------|--------|----------------------------------------------------------|
| /api/admin/orchestration/methods       |        | List all methods (includes <code>isSystemMethod</code> ) |
| /api/admin/orchestration/methods/:code |        | Get method details                                       |
| /api/admin/orchestration/methods/:code | PUT    | Update method (parameters only for system methods)       |
| /api/admin/orchestration/methods       | GET    | Method execution metrics                                 |
| /api/admin/orchestration/executions    | GET    | Recent executions                                        |

#### 34.4 Cato Neural Decision Engine

The `CATO_NEURAL` method integrates with the Genesis Cato safety architecture:

##### Parameters

| Parameter                               | Type    | Default              | Description                             |
|-----------------------------------------|---------|----------------------|-----------------------------------------|
| <code>safety_mode</code>                | enum    | <code>enforce</code> | CBF enforcement: enforce, warn, monitor |
| <code>use_affect_mapping</code>         | boolean | <code>true</code>    | Map affect state to hyperparameters     |
| <code>use_predictive_coding</code>      | boolean | <code>true</code>    | Enable active inference                 |
| <code>precision_governor_enabled</code> | boolean | <code>true</code>    | Limit confidence by epistemic state     |
| <code>cbf_threshold</code>              | number  | 0.95                 | Safety barrier threshold                |

##### Affect-to-Hyperparameter Mapping

| Affect State             | Hyperparameter Effect            |
|--------------------------|----------------------------------|
| High frustration (>0.5)  | Lower temperature (more focused) |
| High curiosity (>0.6)    | Higher temperature (exploration) |
| Low self-efficacy (<0.4) | Escalate to expert model         |
| High arousal (>0.7)      | Longer max tokens                |

#### 34.5 Method Parameters Reference

All methods have configurable parameters that can be set at the **Admin level** (defaults) or overridden in **User Workflow Templates**.

##### Uncertainty & Confidence Methods

| Method                        | Parameter                | Type    | Default | Description                       |
|-------------------------------|--------------------------|---------|---------|-----------------------------------|
| <code>SEMANTIC_ENTROPY</code> | <code>nt</code>          | integer | 10      | Number of response samples (5-20) |
|                               | <code>temperature</code> | number  | 0.7     | Sampling temperature              |

| Method                        | Parameter                       | Type    | Default         | Description                                  |
|-------------------------------|---------------------------------|---------|-----------------|----------------------------------------------|
| <b>SE_PROBES</b>              | <code>clustering_method</code>  | enum    | “nli”           | Clustering: nli, embedding, exact            |
|                               | <code>entropy_threshold</code>  | number  | 0.5             | Flag uncertainty above this                  |
|                               | <code>probe_layers</code>       | array   | [-1, -2]        | Model layers to probe (logprob-based)        |
|                               | <code>threshold</code>          | number  | 0.5             | Uncertainty threshold                        |
|                               | <code>fast_mode</code>          | boolean | true            | Use fast logprob estimation                  |
| <b>KERNEL_ENTROPY</b>         | <code>sample_count</code>       | integer | 5               | Number of samples for averaging              |
|                               | <code>kernel_type</code>        | enum    | “rbf”           | Kernel: rbf, linear, polynomial              |
|                               | <code>bandwidth</code>          | string  | “auto”          | Bandwidth or “auto” for Silverman            |
| <b>CALIBRATED_CALIBRATION</b> | <code>sample_count</code>       | integer | 10              | Response samples for KDE                     |
|                               | <code>calibration_method</code> | enum    | “platt_scaling” | platt_scaling, isotonic, temperature_scaling |
|                               | <code>confidence_prompt</code>  | string  | “verbalized”    | How to elicit confidence                     |
| <b>CONSISTENCY_SAMPLING</b>   | <code>temperature</code>        | number  | 0.3             | Sampling temperature                         |
|                               | <code>sample_count</code>       | integer | 5               | Number of response samples                   |
|                               | <code>agreement_metric</code>   | enum    | “jaccard”       | jaccard, cosine, exact_match, bertscore      |
| <b>CONFORMAL_COVERAGE</b>     | <code>threshold</code>          | number  | 0.7             | Agreement threshold                          |
|                               | <code>target</code>             | number  | 0.9             | Target coverage (0.5-0.99)                   |
|                               | <code>calibration_size</code>   | integer | 500             | Calibration set size                         |
|                               | <code>adaptive</code>           | boolean | true            | Use adaptive conformal sets                  |

## Routing Methods

| Method                | Parameter                      | Type   | Default                         | Description                                   |
|-----------------------|--------------------------------|--------|---------------------------------|-----------------------------------------------|
| <b>ROUTELLM</b>       | <code>router_model</code>      | enum   | “matrix_factorization”          | Router: matrix_factorization, bert, causal_lm |
|                       | <code>cost_threshold</code>    | number | 0.7                             | Max cost relative to baseline                 |
|                       | <code>quality_floor</code>     | number | 0.8                             | Minimum acceptable quality                    |
|                       | <code>models_in_cascade</code> | array  | ["gpt-4o-mini", "gpt-4o", "o1"] | Models in escalation order                    |
| <b>FRUGAL_CASCADE</b> |                                |        |                                 |                                               |

| Method              | Parameter            | Type    | Default | Description                        |
|---------------------|----------------------|---------|---------|------------------------------------|
| <b>PARETO_ROUTE</b> | confidence_threshold | number  | 0.85    | Escalate below this confidence     |
|                     | max_escalations      | integer | 2       | Maximum escalation steps           |
|                     | query_cents          | number  | 10      | Budget constraint per query        |
| <b>C3PO CASCADE</b> | quality_weight       | number  | 0.7     | Weight for quality (0-1)           |
|                     | latency_weight       | number  | 0.1     | Weight for latency (0-1)           |
|                     | model_levels         | integer | 3       | Number of model tiers              |
| <b>AUTOMIX</b>      | self_supervised      | boolean | true    | Enable self-supervised learning    |
|                     | calibration_samples  | integer | 100     | Samples for difficulty calibration |
|                     | pomdp_horizon        | integer | 3       | POMDP planning horizon             |
|                     | exploration_rate     | number  | 0.1     | for -greedy exploration            |
|                     | self_verification    | boolean | true    | Verify own outputs                 |

## Debate & Deliberation Methods

| Method                   | Parameter             | Type    | Default      | Description                      |
|--------------------------|-----------------------|---------|--------------|----------------------------------|
| <b>SPARSE_DEBATE</b>     | debate_logic          | enum    | “ring”       | Network: ring, star, tree, full  |
|                          | debate_rounds         | integer | 3            | Number of debate rounds (1-10)   |
|                          | temperature           | number  | 0.7          | Agent response temperature       |
| <b>ARG_MAPPING</b>       | strength_threshold    | number  | 0.5          | Min argument strength to include |
|                          | include_rebuttal      | boolean | true         | Generate rebuttals               |
|                          | max_depth             | integer | 3            | Max argument tree depth          |
| <b>HAH_DELPHI</b>        | tiers                 | integer | 4            | Number of consensus tiers        |
|                          | human_threshold       | number  | 0.6          | Escalate to human above this     |
|                          | consensus_target      | number  | 0.9          | Target consensus level           |
| <b>RECONCILE_MEASURE</b> | max_rounds            | integer | 5            | Maximum Delphi rounds            |
|                          | confidence_threshold  | number  | 0.6          | Minimum confidence to include    |
|                          | weight_by             | string  | “confidence” | Weighting strategy               |
|                          | reconciliation_rounds | integer | 2            | Reconciliation iterations        |

## Evaluation Methods

| Method                 | Parameter                     | Type    | Default                                              | Description                             |
|------------------------|-------------------------------|---------|------------------------------------------------------|-----------------------------------------|
| <b>POLL_JUDGE</b>      | <code>num_judges</code>       | integer | 3                                                    | Number of judge models                  |
|                        | <code>scoring_criteria</code> | array   | [“accuracy”, “completeness”, “clarity”]              | Evaluation dimensions                   |
|                        | <code>aggregation</code>      | enum    | “mean”                                               | Aggregation: mean, median, weighted     |
| <b>G_EVAL</b>          | <code>dimensions</code>       | array   | [“coherence”, “consistency”, “fluency”, “relevance”] | G-Eval dimensions                       |
|                        | <code>use_cot</code>          | boolean | true                                                 | Chain-of-thought scoring                |
|                        | <code>score_range</code>      | array   | [1, 5]                                               | Score min/max                           |
| <b>PAIRWISE_PREFER</b> | <code>son_criteria</code>     | array   | [“quality”, “accuracy”, “helpfulness”]               | Comparison dimensions                   |
|                        | <code>allow_tie</code>        | boolean | true                                                 | Allow tie verdicts                      |
|                        | <code>sample_count</code>     | integer | 5                                                    | Consistency check samples               |
| <b>SELFCHECK_GPT</b>   | <code>check_type</code>       | enum    | “consistency”                                        | Check type: consistency, bertscore, nli |
|                        | <code>threshold</code>        | number  | 0.7                                                  | Inconsistency threshold                 |
|                        |                               |         |                                                      |                                         |

## Hallucination Detection Methods

| Method                    | Parameter                             | Type    | Default                                            | Description                 |
|---------------------------|---------------------------------------|---------|----------------------------------------------------|-----------------------------|
| <b>MULTI_HALL METHODS</b> | <code>methods</code>                  | array   | [“consistency”, “attribution”, “semantic_entropy”] | Detection methods           |
|                           | <code>aggregation</code>              | enum    | “weighted”                                         | weighted, majority, any     |
|                           | <code>flag_threshold</code>           | number  | 0.6                                                | Flag as hallucination above |
| <b>METAQA</b>             | <code>transformations</code>          | array   | [“paraphrase”, “negation”, “entity_swap”]          | Mutation types              |
|                           | <code>num_mutations</code>            | integer | 3                                                  | Mutations per claim         |
|                           | <code>consistency_threshold</code>    | number  | 0.8                                                | Consistency threshold       |
| <b>FACTUAL_GROUNDING</b>  | <code>documents_top_k</code>          | integer | 5                                                  | Documents to retrieve       |
|                           | <code>evidence_threshold</code>       | number  | 0.7                                                | Evidence support threshold  |
|                           | <code>require_explicit_support</code> | boolean | true                                               | Require explicit evidence   |

## Human-in-the-Loop Methods

| Method               | Parameter            | Type    | Default                           | Description                        |
|----------------------|----------------------|---------|-----------------------------------|------------------------------------|
| <b>HITL_REVIE</b>    | confidence_threshold | number  | 0.7                               | Route to human below this          |
|                      | stake_level          | enum    | “medium”                          | low, medium, high, critical        |
|                      | auto_approve_above   | number  | 0.95                              | Auto-approve above this confidence |
| <b>TIERED_EVAL</b>   | queue_priority       | enum    | “fifo”                            | Queue ordering                     |
|                      | tiers                | integer | 3                                 | Evaluation tiers                   |
|                      | auto_tier_threshold  | number  | 0.85                              | Auto-approve threshold             |
| <b>ACTIVE_SAMPLE</b> | escalation_criteria  | array   | [“low_confidence”, “high_stakes”] | When to escalate                   |
|                      | certainty_method     | enum    | “entropy”                         | entropy, margin, random            |
|                      | batch_size           | integer | 10                                | Samples per batch                  |
|                      | diversity_weight     | number  | 0.3                               | Diversity in selection             |

### 34.6 User Workflow Template Parameter Overrides

When users create workflow templates, they can override default parameters for each step:

```
// Example: User template with parameter overrides
{
  "templateName": "High-Confidence Research",
  "steps": [
    {
      "stepOrder": 1,
      "methodCode": "SEMANTIC_ENTROPY",
      "parameters": {
        "sample_count": 15,           // Override: more samples
        "entropy_threshold": 0.3    // Override: stricter threshold
      }
    },
    {
      "stepOrder": 2,
      "methodCode": "FRUGAL CASCADE",
      "parameters": {
        "confidence_threshold": 0.95, // Override: higher confidence
        "max_escalations": 3       // Override: allow more escalation
      }
    }
  ]
}
```

**Admin Dashboard (Orchestration → Methods):** - View and edit **default parameters** for all methods - Changes apply to all workflows using the method - System methods: parameters only,

not method definition

**Think Tank UI (Workflow Templates)**: - Users create templates with **parameter overrides** - Overrides apply only to that template - Templates can be shared with team

### 34.7 Database Tables

```
-- Methods with display/scientific names
ALTER TABLE orchestration_methods
ADD COLUMN display_name VARCHAR(200),
ADD COLUMN scientific_name VARCHAR(300),
ADD COLUMN research_reference TEXT,
ADD COLUMN accuracy_improvement VARCHAR(200),
ADD COLUMN complexity_level VARCHAR(50);

-- User workflow templates
CREATE TABLE user_workflow_templates (
    template_id UUID PRIMARY KEY,
    tenant_id UUID NOT NULL,
    user_id UUID NOT NULL,
    template_name VARCHAR(200) NOT NULL,
    template_description TEXT,
    base_workflow_code VARCHAR(100),
    steps JSONB NOT NULL DEFAULT '[]',
    category VARCHAR(100),
    tags TEXT[] DEFAULT '{}',
    is_shared BOOLEAN DEFAULT false,
    times_used INTEGER DEFAULT 0,
    created_at TIMESTAMPTZ DEFAULT NOW(),
    UNIQUE(tenant_id, user_id, template_name)
);
```

### 34.7 Implementation Files

| File                                                                        | Purpose                                                     |
|-----------------------------------------------------------------------------|-------------------------------------------------------------|
| migrations/066_orchestration_pattern                                        | Base registry with is_system_method/is_system_workflow      |
| migrations/157_orchestration_methods                                        | Implementation of upserts, display/scientific names         |
| migrations/157_orchestration_methods                                        | Implementation of methods API with system method protection |
| migrations/157_orchestration_methods                                        | Implementation of templates CRUD API                        |
| lambda/shared/services/orchestration-methods                                | Implementation of Kernel Entropy, Pareto, C3PO, AutoMix     |
| lambda/shared/services/cato/neural-networks                                 | Implementation of Neural Network Engine                     |
| lambda/admin/orchestration-methods                                          | Implementation of methods API with system method protection |
| lambda/admin/orchestration-user-templates                                   | Implementation of templates CRUD API                        |
| apps/admin-dashboard/app/(dashboard)/orchestration-methods/page.tsx         | Implementation of methods API page                          |
| apps/admin-dashboard/app/(dashboard)/thinktanks/workflow-templates/page.tsx | Implementation of workflow-templates page                   |

---

## 35. Polymorphic UI (PROMPT-41)

### 35.1 Overview

**Flowise outputs Text. Think Tank outputs Applications.**

The Polymorphic UI system makes Think Tank's interface physically transform based on task complexity, domain hints, and drive profile. Unlike static chatbot interfaces, Think Tank morphs into the tool the user actually needs.

### 35.2 The Gearbox (Elastic Compute)

Users can manually control the cost-quality tradeoff via the Gearbox:

| Mode            | Cost        | Architecture         | Memory                           | Use Case                          |
|-----------------|-------------|----------------------|----------------------------------|-----------------------------------|
| <b>Sniper</b>   | \$0.01/run  | Single Model         | Read-Only<br>Ghost Memory        | Quick answers,<br>lookups, coding |
| <b>War Room</b> | \$0.50+/run | Multi-Agent Ensemble | Read/Write +<br>Active Inference | Strategy, audits,<br>reasoning    |

**Escalation:** A green “Escalate to War Room” button appears after Sniper responses.

### 35.3 The Three Views

| View          | Intent                 | Morph                    | Key Feature                                         |
|---------------|------------------------|--------------------------|-----------------------------------------------------|
| <b>Sniper</b> | Quick commands         | Terminal/Command Center  | Green badge, cost transparency, immediate execution |
| <b>Scout</b>  | Research & exploration | Infinite Canvas/Mind Map | Sticky notes, topic clustering, conflict lines      |
| <b>Sage</b>   | Audit & validation     | Split-Screen Diff Editor | Left=content, Right=sources with confidence scores  |

### 35.4 View Types

| View Type       | Trigger                  | Description                      |
|-----------------|--------------------------|----------------------------------|
| terminal_simple | Quick commands, lookups  | Command Center - fast execution  |
| mindmap         | Research, exploration    | Infinite Canvas - visual mapping |
| diff_editor     | Verification, compliance | Split-Screen - source validation |
| dashboard       | Analytics queries        | Metrics visualization            |
| decision_cards  | HITL escalation          | Mission Control interface        |
| chat            | Default                  | Standard conversation            |

## 35.5 Configuration

Access via **Think Tank → Polymorphic UI** in admin dashboard.

| Setting                | Description                 | Default                               |
|------------------------|-----------------------------|---------------------------------------|
| enableAutoMorphing     | Auto-morph based on query   | true                                  |
| enableGearboxToggle    | Show Sniper/War Room toggle | true                                  |
| enableCostDisplay      | Show cost badges            | true                                  |
| enableEscalationButton | Show Escalate button        | true                                  |
| defaultExecutionMode   | Default mode                | sniper                                |
| domainViewOverrides    | Per-domain view mapping     | medical/financial/legal → diff_editor |

## 35.6 Implementation Files

| File                                            | Purpose                                          |
|-------------------------------------------------|--------------------------------------------------|
| governor/economic-governor.ts                   | determineViewType(), determinePolymorphicRoute() |
| consciousness/mcp-server.ts                     | render_interface, escalate_to_war_room tools     |
| python/cato/cognitive/workflows.py              | Flyte tasks for view selection                   |
| migrations/160_polymorphic_ui.sql               | Database schema                                  |
| components/thinktank/polymorphic/               | React view components                            |
| app/(dashboard)/thinktank/polymorphic/pages.tsx |                                                  |

## 35.7 Database Tables

| Table                 | Purpose                              |
|-----------------------|--------------------------------------|
| view_state_history    | Tracks UI morphing decisions         |
| execution_escalations | Tracks Sniper → War Room escalations |
| polymorphic_config    | Per-tenant configuration             |

## 35.8 API Endpoints

Base: /api/admin/polymorphic

| Method | Endpoint                  | Description            |
|--------|---------------------------|------------------------|
| GET    | /?action=config           | Get configuration      |
| GET    | /?action=view-history     | Get view state history |
| GET    | /?action=escalations      | Get escalation history |
| GET    | /?action=analytics        | Get usage analytics    |
| POST   | / (action: render)        | Render specific view   |
| POST   | / (action: escalate)      | Escalate to War Room   |
| POST   | / (action: update-config) | Update configuration   |

## 36. Think Tank Policy Framework: Strategic Intelligence

The Think Tank is not merely a chatbot interface—it is a **policy research engine** informed by rigorous analysis from organizations like the Cato Institute. This section documents the strategic framework that underlies Think Tank's approach to technology policy, regulation, and economic analysis.

### 36.1 The Cato Institute Policy Foundation

The Cato Institute, a renowned policy think tank, provides crucial insights on technology regulation and innovation. Think Tank integrates these perspectives to ensure users receive balanced, evidence-based analysis rather than ideologically-driven conclusions.

#### Core Policy Principles

| Principle                      | Description                                           | Application in Think Tank                                      |
|--------------------------------|-------------------------------------------------------|----------------------------------------------------------------|
| <b>Pro-Innovation Default</b>  | Permissionless innovation should be the norm          | Bias toward “how can this work?” not “why won’t this work?”    |
| <b>Evidence-Based Analysis</b> | Policy must be grounded in data, not fear             | Require citations and data sources in policy discussions       |
| <b>Regulatory Humility</b>     | Acknowledge limits of prediction                      | Flag uncertainty explicitly; avoid overconfident prescriptions |
| <b>Market-Based Solutions</b>  | Private sector often innovates faster than regulators | Explore voluntary standards before mandates                    |
| <b>Individual Liberty</b>      | Technology should enhance freedom, not surveillance   | Privacy-by-design; minimize data collection                    |

### 36.2 The \$10 Trillion Cybercrime Economy

A critical context for Think Tank's policy analysis is the scale of the threat environment. The global economy loses approximately **\$10 trillion annually** to cybercrime—a figure larger than the GDP of every country except the United States and China.

#### Economic Impact Analysis

##### GLOBAL CYBERCRIME ECONOMIC IMPACT

ANNUAL LOSSES: ~\$10 TRILLION

\$10T

CYBERCRIME

\$4.2T GERMANY GDP

\$3.4T JAPAN GDP

\$2.1T UK GDP

#### BREAKDOWN BY ATTACK TYPE:

- Ransomware: \$20B+ annually
- Business Email Compromise: \$2.7B annually
- Data Breaches: \$4.45M average cost per incident
- Supply Chain Attacks: Growing 742% since 2019
- Nation-State Attacks: Incalculable strategic damage

**Implications for Policy Discussion** When users ask Think Tank about cybersecurity policy, the system contextualizes recommendations against this \$10T backdrop:

| User Query                             | Policy-Aware Response Approach                                             |
|----------------------------------------|----------------------------------------------------------------------------|
| “Should we regulate AI?”               | Frame against: AI can <i>reduce</i> \$10T losses if deployed correctly     |
| “Is this security spending justified?” | Compare against: Your \$1M security budget vs. \$4.45M average breach cost |
| “Should we ban ransomware payments?”   | Analyze: Prohibition vs. harm reduction strategies                         |
| “How strict should compliance be?”     | Balance: Compliance costs vs. breach probability × impact                  |

### 36.3 Memory Safety and the 70% Problem

The Think Tank injects a specific technical insight into relevant conversations: **70% of all software vulnerabilities** stem from memory safety issues. This statistic, validated by Microsoft, Google, and the NSA, should inform every software architecture discussion.

#### Memory Safety Vulnerability Classes

| Vulnerability                   | Description                     | % of CVEs   |
|---------------------------------|---------------------------------|-------------|
| <b>Buffer Overflow</b>          | Writing beyond allocated memory | ~30%        |
| <b>Use-After-Free</b>           | Accessing freed memory          | ~20%        |
| <b>Double Free</b>              | Freeing memory twice            | ~8%         |
| <b>Null Pointer Dereference</b> | Accessing null pointers         | ~7%         |
| <b>Integer Overflow</b>         | Arithmetic exceeding bounds     | ~5%         |
| <b>Total Memory Safety</b>      | All memory-related issues       | <b>~70%</b> |

**Policy Recommendation Engine** When users discuss software architecture, security, or procurement, Think Tank can inject policy-informed guidance:

```
// Policy injection for memory safety discussions
const MEMORY_SAFETY_POLICY = {
    context: 'User discussing software architecture or security',
    injectedGuidance: `

        POLICY CONTEXT: Memory safety vulnerabilities account for 70% of all CVEs.

        RECOMMENDATIONS:
        1. For new projects: Prefer memory-safe languages (Rust, Go, Swift)
        2. For existing C/C++ code: Consider incremental migration or sandboxing
        3. For procurement: Require memory-safe language attestation from vendors
        4. For risk assessment: Memory-unsafe components = higher risk weight

        SOURCE: Microsoft, Google, NSA research (2019-2024)
    `,
    triggerPatterns: [
        'architecture', 'security', 'language choice', 'procurement',
        'vulnerability', 'CVE', 'buffer overflow', 'memory'
    ],
};
```

### 36.4 Regulatory Stance Configuration

Administrators can configure Think Tank's default policy stance for their organization:

| Setting                 | Options                                | Description                       |
|-------------------------|----------------------------------------|-----------------------------------|
| defaultRegulatorystance | cautious / balanced / pro_innovation   | Bias in regulatory discussions    |
| requireCitations        | true / false                           | Require sources for policy claims |
| flagUncertainty         | always / when_high / never             | Uncertainty disclosure level      |
| privateDataBias         | minimize / balanced / maximize_utility | Data collection philosophy        |
| complianceEmphasis      | strict / risk_based / minimal          | Compliance recommendation style   |

**Configuration Example** Navigate to **Think Tank → Policy Framework** in admin dashboard:

```
# policy-framework.config.yaml
policy_framework:
    enabled: true

    default_stance: balanced
```

```

cybercrime_context:
  enabled: true
  inject_10T_context: true
  inject_memory_safety_context: true

citation_requirements:
  require_for_policy_claims: true
  preferred_sources:
    - cato_institute
    - brookings
    - nist
    - academic_peer_reviewed
  flag_opinion_vs_fact: true

uncertainty_handling:
  disclosure_level: always
  confidence_threshold_for_recommendation: 0.7
  escalate_low_confidence_to_human: true

privacy_settings:
  default_data_collection: minimize
  require_purpose_limitation: true
  support_right_to_deletion: true

```

### 36.5 Database Tables

| Table                       | Purpose                              |
|-----------------------------|--------------------------------------|
| policy_framework_config     | Per-tenant policy configuration      |
| policy_citation_log         | Citations used in policy discussions |
| policy_uncertainty_flags    | Uncertainty disclosures              |
| regulatory_stance_overrides | Per-domain stance overrides          |

### 36.6 Implementation Files

| File                                                    | Purpose                          |
|---------------------------------------------------------|----------------------------------|
| lambda/shared/services/policy-framework.servicetests.ts | Unit tests                       |
| lambda/shared/services/citation-matching-service.ts     | Citation matching and validation |
| lambda/thinktank/policy-context.ts                      | API handler for policy queries   |
| migrations/164_policy_framework.sql                     | Database schema                  |
| config/policy/cato-principles.yaml                      | Cato Institute policy principles |

## 37. Agentic Orchestration: SSF, CAEP, and Identity Remediation

Think Tank's AI agents operate within a rigorous security framework that leverages open standards for real-time security signaling. This section documents the Shared Signals Framework (SSF) and Continuous Access Evaluation Profile (CAEP) integration specific to Think Tank operations.

### 37.1 The Agentic AI Paradigm

Traditional automation follows rigid if-then rules. **Agentic AI** represents a fundamental shift: AI systems that continuously evaluate their environment, learn from outcomes, and adapt their behavior within defined safety constraints.

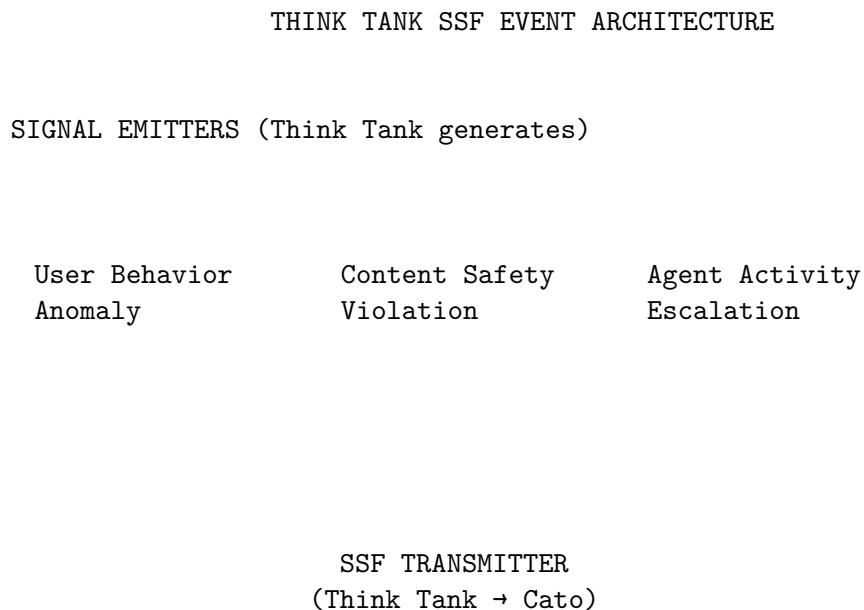
#### Traditional Automation vs. Agentic AI

| Aspect                   | Traditional Automation | Agentic AI (Think Tank) |
|--------------------------|------------------------|-------------------------|
| <b>Decision Logic</b>    | Static rules           | Dynamic evaluation      |
| <b>Learning</b>          | None                   | Continuous adaptation   |
| <b>Error Handling</b>    | Fail or retry          | Investigate and adapt   |
| <b>Human Interaction</b> | Scheduled checkpoints  | On-demand escalation    |
| <b>Security Model</b>    | Perimeter-based        | Zero Trust + CAEP       |

### 37.2 Shared Signals Framework (SSF) Integration

The **Shared Signals Framework** is an open standard that enables real-time security event sharing between systems. Think Tank agents both emit and consume SSF signals.

#### SSF Event Flow in Think Tank



**SSF RECEIVER**  
(Cato → Think Tank)

|                               |                               |                                 |
|-------------------------------|-------------------------------|---------------------------------|
| Session Revoke<br>→ Terminate | Threat Detected<br>→ Restrict | Genesis Alert<br>→ Pause Agents |
|-------------------------------|-------------------------------|---------------------------------|

### SSF Events Emitted by Think Tank

| Event Type                  | Trigger                                            | Recipient Action           |
|-----------------------------|----------------------------------------------------|----------------------------|
| thinktank.behavior.anomaly  | User behavior deviates significantly from baseline | Cato increases monitoring  |
| thinktank.content.violation | User attempts prohibited content                   | Identity provider notified |
| thinktank.agent.escalation  | Agent requires human approval                      | Mission Control alerted    |
| thinktank.session.suspicion | Multiple failed attempts or unusual patterns       | Cato may revoke session    |
| thinktank.data.exfiltration | Attempted data extraction attempt                  | Block and investigate      |

### SSF Events Consumed by Think Tank

| Event Type               | Source            | Think Tank Action                  |
|--------------------------|-------------------|------------------------------------|
| session-revoked          | Identity Provider | Immediately terminate user session |
| credential-change        | Identity Provider | Force re-authentication            |
| threat-detected          | Cato SASE         | Restrict agent capabilities        |
| genesis.alert            | Genesis Reactor   | Pause all non-critical agents      |
| device-compliance-change | MDM/EDR           | Re-evaluate user permissions       |

### 37.3 Continuous Access Evaluation Profile (CAEP)

CAEP extends SSF with specific event types designed for continuous session validation. Unlike traditional session timeouts, CAEP enables **real-time session adjustment** based on security signals.

#### CAEP Session Lifecycle

##### CAEP SESSION LIFECYCLE IN THINK TANK

###### 1. SESSION START

User authenticates → Session created → CAEP listener registered

###### 2. CONTINUOUS EVALUATION (every CAEP event)

CAEP Event Received → Evaluate Impact → Adjust Session

Example: IP Change

Same country? → Log only

Different country? → Step-up authentication

Impossible travel? → Terminate session

###### 3. SESSION TERMINATION

CAEP revoke signal OR user logout OR timeout → Clean termination  
→ Preserve conversation state → Clear credentials → Audit log

#### CAEP Configuration for Think Tank

```
# caep-thinktank.config.yaml
caep:
  enabled: true

  session_management:
    # How to handle IP changes
    ip_change_policy:
      same_country: log_only
      different_country: step_up_auth
      impossible_travel: terminate_session
```

```

impossible_travel_threshold_hours: 2

# How to handle device changes
device_change_policy:
  known_device: allow
  unknown_device_trusted_location: step_up_auth
  unknown_device_unknown_location: terminate_session

# How to handle credential events
credential_change_policy:
  password_change: force_reauth
  mfa_change: force_reauth
  role_change: re_evaluate_permissions

agent_restrictions:
  # During security events, restrict agent capabilities
  during_threat_detected:
    disable_autonomous_actions: true
    disable_external_api_calls: true
    require_human_approval_all: true

  during_genesis_alert:
    pause_all_agents: true
    preserve_state: true
    notify_administrators: true

```

## 37.4 Autonomous Identity Remediation

Think Tank agents can perform **autonomous identity remediation**—cleaning up identity data quality issues without human intervention. This capability requires careful configuration to balance efficiency with safety.

### Remediation Capabilities

| Action                             | Autonomous? | Conditions                                         |
|------------------------------------|-------------|----------------------------------------------------|
| Remove Orphan Account              | Yes         | Inactive 90+ days, no entitlements, no recent auth |
| Disable Stale Service Account      | Yes         | Unused 180+ days, not in critical systems          |
| Fix Group Membership Inconsistency | Yes         | Source of truth mismatch detected                  |
| Delete Human Account               | No          | Always requires human approval                     |
| Revoke Admin Privileges            | No          | Always requires human approval                     |

| Action                        | Autonomous? | Conditions                     |
|-------------------------------|-------------|--------------------------------|
| <b>Bulk Operations (100+)</b> | <b>No</b>   | Always requires human approval |

## Remediation Workflow

### AUTONOMOUS REMEDIATION WORKFLOW

#### 1. DETECTION

- Agent scans Identity Fabric for anomalies
  - Orphan accounts (no manager, no recent activity)
  - Stale service accounts (unused beyond threshold)
  - Inconsistent group memberships (source mismatch)

#### 2. VALIDATION

- Verify anomaly against multiple data sources
- Check against exclusion lists (VIPs, system accounts)
- Confirm remediation action is within autonomous scope
- Verify Genesis Interlock allows action (no active alerts)

#### 3. EXECUTION

- If autonomous: Execute immediately, log action
- If requires approval: Create Mission Control escalation

#### 4. VERIFICATION

- Confirm action completed successfully
- Verify no unintended side effects
- Update remediation statistics

#### 5. LEARNING

- If action succeeded: Reinforce pattern
- If action failed or reversed: Flag for human review

**Remediation Configuration** Navigate to Think Tank → Identity → Remediation in admin dashboard:

```
# identity-remediation-thinktank.config.yaml
identity_remediation:
  enabled: true

  # Autonomous action limits
  autonomous_limits:
    max_per_hour: 50
```

```

max_per_day: 200
cooldown_after_error_minutes: 30

# Detection criteria
detection:
  orphan_account:
    inactive_days_threshold: 90
    require_no_entitlements: true
    require_no_recent_auth: true
    exclude_patterns:
      - "*-system@*"
      - "*-service@*"
      - "admin@*"

  stale_service_account:
    unused_days_threshold: 180
    exclude_critical_systems: true
    critical_system_tags:
      - genesis
      - security
      - auth

  group_membership:
    check_source_of_truth: true
    sources:
      - active_directory
      - scim_provider
      - hr_system

# Notification settings
notifications:
  notify_on_autonomous_action: true
  notify_recipients:
    - security_team
    - identity_admins
  daily_summary: true

# Genesis Interlock integration
genesis_interlock:
  pause_during: [GEN-300, GEN-400, GEN-500]
  resume Automatically: true

```

### 37.5 The Radiant Ghost in Think Tank

The **Radiant Ghost** metaphor extends to Think Tank's user interface, providing visual feedback about agent activity.

#### Ghost States in Think Tank UI

| State              | Indicator           | Meaning                     | User Action |
|--------------------|---------------------|-----------------------------|-------------|
| <b>Idle</b>        | Faint outline       | No active agents            | None        |
| <b>Thinking</b>    | Soft pulse          | Agent processing request    | Wait        |
| <b>Researching</b> | Searching animation | Agent gathering information | Wait        |
| <b>Writing</b>     | Typing animation    | Agent generating response   | Watch       |
| <b>Validating</b>  | Checkmark animation | Agent verifying output      | Wait        |
| <b>Escalating</b>  | Orange pulse        | Agent needs human input     | Respond     |
| <b>Alert</b>       | Red pulse           | Security or safety event    | Investigate |

## Ghost Configuration

```
# ghost-ui.config.yaml
ghost_ui:
  enabled: true

  # Visual settings
  visuals:
    idle_opacity: 0.3
    active_opacity: 0.8
    alert_opacity: 1.0
    animation_speed: normal # slow, normal, fast

  # State display
  states:
    show_thinking: true
    show_researching: true
    show_writing: true
    show_validating: true
    show_escalating: true
    show_alert: true

  # Accessibility
  accessibility:
    respect_reduced_motion: true
    provide_text_status: true
    screen_reader_announcements: true
```

## 37.6 Database Tables

| Table                                    | Purpose                           |
|------------------------------------------|-----------------------------------|
| <code>ssf_events_emitted</code>          | SSF events sent by Think Tank     |
| <code>ssf_events_received</code>         | SSF events received by Think Tank |
| <code>caep_session_events</code>         | CAEP session lifecycle events     |
| <code>identity_remediation_log</code>    | All remediation actions           |
| <code>identity_remediation_errors</code> | Failed remediation attempts       |
| <code>ghost_state_log</code>             | Ghost UI state transitions        |

## 37.7 API Endpoints

Base: `/api/thinktank/security`

| Method | Endpoint                          | Description                          |
|--------|-----------------------------------|--------------------------------------|
| GET    | <code>/ssf/status</code>          | SSF integration status               |
| GET    | <code>/ssf/events</code>          | Recent SSF events                    |
| POST   | <code>/ssf/emit</code>            | Manually emit SSF event (admin only) |
| GET    | <code>/caep/session</code>        | Current session CAEP status          |
| GET    | <code>/remediation/stats</code>   | Remediation statistics               |
| GET    | <code>/remediation/log</code>     | Remediation action log               |
| POST   | <code>/remediation/trigger</code> | Trigger manual remediation scan      |
| GET    | <code>/ghost/state</code>         | Current Ghost state                  |

## 37.8 Implementation Files

| File                                                                                | Purpose                     |
|-------------------------------------------------------------------------------------|-----------------------------|
| <code>lambda/shared/services/ssf-thinktank/SSFservice.ts</code>                     | SSF receiver for Think Tank |
| <code>lambda/shared/services/caep-session/Service.ts</code>                         | CAEP session management     |
| <code>lambda/shared/services/identity-repository/Identitythinktankservice.ts</code> | Identity Think Tank service |
| <code>lambda/thinktank/security-events.ts</code>                                    | Security event API handler  |
| <code>components/thinktank/ghost-indicator/GhostxUI component</code>                | Ghost UI component          |
| <code>migrations/165_agentic_orchestratorDatabase</code>                            | Database schema             |
| <code>config/security/ssf-events.yaml</code>                                        | SSF event definitions       |
| <code>config/security/caep-policies.yaml</code>                                     | CAEP policy configuration   |

## 38. Advanced Features (v4.18.0)

This section covers new advanced features implemented for Think Tank.

### 38.1 Flash Facts (Knowledge Sparks)

Fast-access factual memory system for quick retrieval of verified facts.

**UI Metaphor: “Knowledge Sparks”** - Contextual sidebar widget showing relevant facts as glowing spark icons.

**Features:** - CRUD operations for facts with confidence scoring - Semantic search using vector embeddings - Automatic fact extraction from conversations - Fact verification workflow - Usage tracking and statistics

#### API Endpoints (Base: /api/thinktank/flash-facts):

| Method | Endpoint    | Description                     |
|--------|-------------|---------------------------------|
| GET    | /           | List facts with filtering       |
| POST   | /           | Create a new fact               |
| GET    | /:id        | Get a specific fact             |
| PUT    | /:id        | Update a fact                   |
| DELETE | /:id        | Delete a fact                   |
| POST   | /query      | Semantic search for facts       |
| POST   | /extract    | Extract facts from conversation |
| POST   | /:id/verify | Verify a fact                   |
| GET    | /stats      | Get usage statistics            |

**Database Tables:** - flash\_facts - Fact storage with embeddings - flash\_facts\_config - Per-tenant configuration

## 38.2 Grimoire (Spell Book)

Procedural memory system for storing and executing reusable patterns.

**UI Metaphor: “Spell Book”** - Magical tome with spell cards organized by schools of magic.

**Schools of Magic:** | School | Icon | Purpose | |——|——|——| | Code | | Programming patterns | | Data | | Data manipulation | | Text | | Text processing | | Analysis | | Analytical methods | | Design | | UI/UX patterns | | Integration | | API integrations | | Automation | | Workflow automation | | Universal | | Cross-domain spells |

**Spell Categories:** - Transformation, Divination, Conjuration, Abjuration - Enchantment, Illusion, Necromancy, Evocation

#### API Endpoints (Base: /api/thinktank/grimoire):

| Method | Endpoint          | Description           |
|--------|-------------------|-----------------------|
| GET    | /                 | Get grimoire overview |
| GET    | /spells           | List all spells       |
| POST   | /spells           | Create a new spell    |
| GET    | /spells/:id       | Get spell details     |
| PUT    | /spells/:id       | Update a spell        |
| DELETE | /spells/:id       | Delete a spell        |
| POST   | /spells/:id/cast  | Cast a spell          |
| POST   | /spells/:id/learn | Learn from failure    |
| GET    | /schools          | Get schools of magic  |
| GET    | /categories       | Get spell categories  |

---

| Method | Endpoint | Description              |
|--------|----------|--------------------------|
| POST   | /match   | Find matching spell      |
| POST   | /promote | Promote pattern to spell |

---

### 38.3 Economic Governor (Fuel Gauge)

Model arbitrage and cost optimization system.

**UI Metaphor:** “Fuel Gauge” - Visual meter showing budget remaining with color-coded status.

**Governor Modes:** | Mode | Description | |——|——| | **cost\_minimizer** | Always cheapest viable option | | **quality\_maximizer** | Best quality within budget | | **balanced** | Balance cost and quality | | **latency.Focused** | Prioritize response speed | | **custom** | Use custom arbitrage rules |

**Model Tiers:** - Economy ( ) - Cheapest, simple tasks - Self-Hosted ( ) - On-premise models - Standard ( ) - Default tier - Premium ( ) - Higher quality - Flagship ( ) - Best available

**API Endpoints (Base: /api/thinktank/governor):**

---

| Method | Endpoint     | Description                   |
|--------|--------------|-------------------------------|
| GET    | /            | Get dashboard with fuel gauge |
| GET    | /config      | Get configuration             |
| PUT    | /config      | Update configuration          |
| PUT    | /mode        | Quick mode switch             |
| POST   | /recommend   | Get model recommendation      |
| GET    | /metrics     | Get cost metrics              |
| GET    | /budget      | Get budget status             |
| PUT    | /budget      | Update budget limit           |
| GET    | /tiers       | Get model tiers               |
| PUT    | /tiers/:tier | Update tier config            |
| GET    | /rules       | Get arbitrage rules           |
| POST   | /rules       | Add arbitrage rule            |
| PUT    | /rules/:id   | Update rule                   |
| DELETE | /rules/:id   | Delete rule                   |

---

### 38.4 Sentinel Agents (Watchtower Dashboard)

Event-driven autonomous agents for monitoring and automation.

**UI Metaphor:** “Watchtower Dashboard” - Castle towers watching over different domains.

**Agent Types:** | Type | Icon | Purpose | |——|——|——| | Monitor | | Passive watchdog - observes and alerts | | Guardian | | Protective - can block harmful actions | | Scout | | Proactive information gathering | | Herald | | Notifications and announcements | | Arbiter | | Decision making and routing |

**Agent Status:** - Idle ( ), Watching ( ), Triggered ( ), Cooldown ( ), Disabled ( )

**API Endpoints (Base: /api/thinktank/sentinels):**

| Method | Endpoint     | Description       |
|--------|--------------|-------------------|
| GET    | /            | List all agents   |
| POST   | /            | Create agent      |
| GET    | /:id         | Get agent details |
| PUT    | /:id         | Update agent      |
| DELETE | /:id         | Delete agent      |
| POST   | /:id/trigger | Manually trigger  |
| POST   | /:id/enable  | Enable agent      |
| POST   | /:id/disable | Disable agent     |
| GET    | /:id/events  | Get agent events  |
| GET    | /events      | All events        |
| GET    | /stats       | Statistics        |
| GET    | /types       | Available types   |

---

### 38.5 Time-Travel Debugging (Timeline Scrubber)

Conversation forking and state replay system.

**UI Metaphor: “Timeline Scrubber”** - Horizontal timeline with draggable playhead and fork points.

**Checkpoint Types:** | Type | Icon | Description | |——|——|——|-| | Auto | | Automatic checkpoint | | Manual | | User-created | | Fork | | Branch point | | Merge | | Merged timelines | | Rollback | | After rollback |

**API Endpoints (Base: /api/thinktank/time-travel):**

| Method | Endpoint                  | Description        |
|--------|---------------------------|--------------------|
| GET    | /timelines                | List timelines     |
| POST   | /timelines                | Create timeline    |
| GET    | /timelines/:id            | Get timeline view  |
| POST   | /timelines/:id/checkpoint | Create checkpoint  |
| POST   | /timelines/:id/jump       | Jump to checkpoint |
| POST   | /timelines/:id/fork       | Fork timeline      |
| POST   | /timelines/:id/replay     | Replay sequence    |
| GET    | /checkpoints/:id          | Get checkpoint     |

---

### 38.6 Council of Rivals (Debate Arena)

Multi-model adversarial consensus system.

**UI Metaphor:** “Debate Arena” - Amphitheater with model avatars debating in a circular arrangement.

**Member Roles:** | Role | Icon | Purpose | |——|——|——| | Advocate | | Argues for a position | | Critic | | Challenges positions | | Synthesizer | | Combines perspectives | | Specialist | | Domain expert | | Contrarian | | Devil’s advocate |

**Preset Councils:** - **Balanced** () - Diverse perspectives - **Technical** () - Expert technical review - **Creative** () - Creative exploration

**Verdict Outcomes:** - Consensus (), Majority (), Split (), Deadlock (), Synthesized ()

**API Endpoints (Base: /api/thinktank/council):**

| Method | Endpoint             | Description           |
|--------|----------------------|-----------------------|
| GET    | /                    | List councils         |
| POST   | /                    | Create council        |
| POST   | /preset              | Create preset council |
| GET    | /:id                 | Get council           |
| PUT    | /:id                 | Update council        |
| DELETE | /:id                 | Delete council        |
| POST   | /:id/debate          | Start debate          |
| GET    | /debate/:id          | Get debate            |
| POST   | /debate/:id/argument | Submit argument       |
| POST   | /debate/:id/rebuttal | Submit rebuttal       |
| POST   | /debate/:id/vote     | Conduct voting        |
| GET    | /presets             | Get preset options    |

### 38.7 Security Signals (Security Shield)

SSF/CAEP integration for identity security events.

**UI Metaphor:** “Security Shield” - Animated shield with real-time threat visualization.

**Signal Types:** | Type | Icon | Description | |——|——|——| | Session Revoked | | SSF session terminated | | Credential Change | | Password/key changed | | Device Compliance | | CAEP compliance change | | Risk Change | | Risk level changed | | Anomaly Detected | | Behavioral anomaly | | Threat Detected | | Active threat | | Policy Violation | | Policy breach |

**Severity Levels:** Critical (), High (), Medium (), Low (), Info ()

**API Endpoints (Base: /api/thinktank/security):**

| Method | Endpoint            | Description        |
|--------|---------------------|--------------------|
| GET    | /dashboard          | Security dashboard |
| GET    | /signals            | List signals       |
| POST   | /signals            | Create signal      |
| GET    | /signals/:id        | Get signal         |
| PUT    | /signals/:id/status | Update status      |

| Method | Endpoint      | Description       |
|--------|---------------|-------------------|
| GET    | /policies     | List policies     |
| POST   | /policies     | Create policy     |
| PUT    | /policies/:id | Update policy     |
| DELETE | /policies/:id | Delete policy     |
| POST   | /ssf/event    | Ingest SSF event  |
| POST   | /caep/event   | Ingest CAEP event |

---

### 38.8 Policy Framework (Stance Compass)

Strategic intelligence and regulatory stance configuration.

**UI Metaphor: “Stance Compass”** - Radial chart showing policy positions across domains.

**Policy Domains:** | Domain | Icon | Description | |——|——|——| | AI Safety | | AI alignment and safety | | Data Privacy | | Data protection regulations | | Content Moderation | | Content policies | | Accessibility | | Accessibility requirements | | Sustainability | | Environmental considerations | | Security | | Cybersecurity posture | | Transparency | | AI transparency | | Ethics | | Ethical AI principles | | Compliance | | Regulatory compliance | | Innovation | | Innovation balance |

**Stance Positions:** - Restrictive ( ), Cautious ( ), Balanced ( ), Permissive ( ), Adaptive ( )

**Preset Profiles:** - **Conservative** - Maximum safety, minimal risk - **Balanced** - Middle ground - **Innovative** - Emphasis on innovation

**API Endpoints (Base: /api/thinktank/policy):**

| Method | Endpoint               | Description           |
|--------|------------------------|-----------------------|
| GET    | /compass               | Get compass view      |
| GET    | /domains               | Get available domains |
| GET    | /positions             | Get stance positions  |
| GET    | /stances               | List stances          |
| POST   | /stances               | Create stance         |
| GET    | /stances/:domain       | Get domain stance     |
| PUT    | /stances/:id           | Update stance         |
| GET    | /profiles              | List profiles         |
| GET    | /profiles/active       | Get active profile    |
| POST   | /profiles              | Create profile        |
| POST   | /profiles/preset       | Create preset profile |
| PUT    | /profiles/:id/activate | Activate profile      |
| GET    | /recommendations       | Get recommendations   |
| GET    | /compliance            | Check compliance      |

---

## 38.9 Database Migration

All features use migration `100_thinktank_advanced_features.sql` with the following tables:

| Table                                 | Feature             |
|---------------------------------------|---------------------|
| <code>flash_facts</code>              | Flash Facts storage |
| <code>flash_facts_config</code>       | Flash Facts config  |
| <code>grimoire_spells</code>          | Grimoire spells     |
| <code>grimoire_casts</code>           | Spell cast history  |
| <code>grimoire_achievements</code>    | User achievements   |
| <code>economic_governor_config</code> | Governor config     |
| <code>economic_governor_usage</code>  | Usage tracking      |
| <code>sentinel_agents</code>          | Sentinel agents     |
| <code>sentinel_events</code>          | Agent events        |
| <code>time_travel_timelines</code>    | Timelines           |
| <code>time_travel_checkpoints</code>  | Checkpoints         |
| <code>time_travel_forks</code>        | Fork records        |
| <code>council_of_rivals</code>        | Councils            |
| <code>council_debates</code>          | Debates             |
| <code>security_signals</code>         | Security signals    |
| <code>security_policies</code>        | Security policies   |
| <code>policy_stances</code>           | Policy stances      |
| <code>policy_profiles</code>          | Policy profiles     |

## 39. Liquid Interface (Generative UI)

### 39.1 Overview

**“Don’t Build the Tool. BE the Tool.”**

The Liquid Interface transforms the chat interface into dynamic, morphable UI tools based on user intent. Instead of asking “help me make a spreadsheet app,” the chat *becomes* the spreadsheet.

**Core Concept:** - User says: “Help me track my invoices” - Chat morphs into: Invoice tracker with data grid, totals, AI assistant sidebar - User interacts: Add, edit, filter invoices directly - AI watches: Ghost State binds UI actions to AI context - Export: “Eject” to a deployable Next.js/Vite app

**Key Benefits:** - **Zero-friction prototyping** - Ideas become tools instantly - **Two-way AI binding** - AI sees what you’re doing, UI reflects what AI knows - **Production export** - Ephemeral apps become real codebases - **50+ morphable components** - Data grids, charts, kanban, calendars, code editors

---

### 39.2 Architecture

User Message

Intent Detection

```
data_analysis      tracking      visualization     ...
```

Schema Generation

```
LiquidSchema {  
  layout: { type: 'split', children: [...] }  
  bindings: [{ sourceComponent, contextKey, direction }]  
  aiOverlay: { mode: 'sidebar', position: 'right' }  
}
```

Liquid Renderer

```
DataGrid      Chart      Kanban      AI Chat
```

Ghost State Manager

```
UI Events    AI Context  
AI Updates   UI Components
```

---

### 39.3 Component Registry

50+ pre-built morphable components across 9 categories:

| Category                  | Components                                                                                                                     | Description                        |
|---------------------------|--------------------------------------------------------------------------------------------------------------------------------|------------------------------------|
| <b>Data</b> (10)          | DataGrid, PivotTable,<br>DataCard, JSONViewer,<br>SQLViewer, CSVEditor,<br>DataFilter, SchemaDesigner,<br>DataDiff, DataImport | Spreadsheets, tables, data viewers |
| <b>Visualization</b> (10) | LineChart, BarChart, PieChart, ScatterPlot, AreaChart, Heatmap, Treemap, GeoMap, Timeline, SankeyDiagram                       | Charts, graphs, maps               |
| <b>Productivity</b> (10)  | KanbanBoard, Calendar, GanttChart, TodoList, NotesEditor, Timer, HabitTracker, MindMap, Whiteboard, FileManager                | Task & project management          |
| <b>Finance</b> (6)        | Invoice, BudgetTracker, ExpenseTable, Portfolio, Calculator, CurrencyConverter                                                 | Financial tools                    |
| <b>Code</b> (6)           | CodeEditor, Terminal, DiffViewer, APITester, RegexTester, JSONFormatter                                                        | Developer tools                    |
| <b>AI</b> (4)             | AIChat, InsightCard, SuggestionPanel, ContextInspector                                                                         | AI-powered widgets                 |
| <b>Input</b> (4)          | FormBuilder, SliderPanel, DateRangePicker, SearchBox                                                                           | User input forms                   |

### Component Definition:

```
interface LiquidComponent {
  id: string;                                // e.g., 'data-grid'
  name: string;                                // e.g., 'DataGrid'
  category: ComponentCategory;                 // e.g., 'data'
  propsSchema: JSONSchema;                     // Component props
  eventsSchema: JSONSchema;                    // Emitted events
  supportsInteraction: boolean;
  supportsDataBinding: boolean;
  supportsAIContext: boolean;
  defaultSize: { width, height };
  icon: string;
  tags: string[];
}
```

---

### 39.4 Ghost State (Two-Way Binding)

“The AI sees what you’re doing. The UI reflects what AI knows.”

Ghost State creates bidirectional bindings between UI components and AI context:

|                |                                  |
|----------------|----------------------------------|
| UI Component   | AI Context                       |
| selectedRow: 5 | GhostBinding      user_selection |
| filterValue: X | applied_filter                   |
| [AI updates]   | AI Reaction      insight: "..."  |

### Binding Configuration:

```
interface GhostBinding {  
  id: string;  
  sourceComponent: string; // UI component ID  
  sourceProperty: string; // e.g., 'selectedRow'  
  contextKey: string; // AI context key  
  direction: 'ui_to_ai' | 'ai_to_ui' | 'bidirectional';  
  debounceMs?: number; // Debounce rapid changes  
  triggerReaction?: boolean; // Trigger AI response on change  
  reactionPrompt?: string; // Custom prompt for reaction  
}
```

**AI Reactions:** When users interact with morphed UI, the AI can react with:  
- speak - Send a message  
- update - Update UI state  
- morph - Transform to different layout  
- suggest - Show suggestion panel

---

## 39.5 Intent Detection

Intent detection determines when and how to morph the UI:

| Intent Category | Trigger Phrases                          | Default Components             |
|-----------------|------------------------------------------|--------------------------------|
| data_analysis   | spreadsheet, excel, csv,<br>analyze data | DataGridView, PivotTable       |
| tracking        | track invoices, manage<br>expenses       | Invoice, ExpenseTable, Kanban  |
| visualization   | chart, graph, visualize, show<br>metrics | LineChart, BarChart, Dashboard |
| planning        | plan project, timeline,<br>kanban        | KanbanBoard, GanttChart        |
| calculation     | calculate, compute, formula              | Calculator, DataGridView       |
| design          | design, wireframe, brainstorm            | Whiteboard, MindMap            |
| coding          | code, debug, terminal                    | CodeEditor, Terminal           |
| writing         | write, draft, notes                      | NotesEditor, MindMap           |

**Morph Threshold:** confidence >= 0.85 (configurable per tenant)

---

## 39.6 Eject to App

“The Takeout Button” - Export ephemeral liquid apps to real codebases.

**Supported Frameworks:** - **Next.js 14** - Full-stack React with API routes - **Vite + React** - Fast client-side SPA - **Remix** - Web standards framework - **Astro** - Content-focused sites

**Features to Include:** | Feature | Description | |———|———| | database | PGLite → Postgres migration, Drizzle ORM | | auth | NextAuth scaffolding | | api | API routes for data operations | | ai | OpenAI integration | | realtime | WebSocket support |

**Generated Files:**

```
my-liquid-app/
  package.json
  next.config.mjs / vite.config.ts
  tsconfig.json
  tailwind.config.ts
  app/page.tsx (or src/App.tsx)
  components/
    LiquidLayout.tsx
    DataGrid.tsx
    ...
  store/index.ts (Zustand)
  types/index.ts
  lib/db.ts (if database)
  lib/ai.ts (if ai)
  README.md
  .env.example
  .gitignore
```

---

## 39.7 Configuration

**Per-Tenant Configuration:**

| Setting                      | Default | Description                          |
|------------------------------|---------|--------------------------------------|
| enabled                      | true    | Enable Liquid Interface              |
| auto_morph_enabled           | true    | Auto-morph on high-confidence intent |
| eject_enabled                | true    | Allow app export                     |
| morph_confidence_threshold   | 0.85    | Minimum confidence to morph          |
| auto_revert_timeout_seconds  | 300     | Auto-revert to chat after inactivity |
| max_active_sessions          | 10      | Max concurrent liquid sessions       |
| max_ghost_events_per_session | 1000    | Event history limit                  |
| default_overlay_mode         | sidebar | AI overlay mode                      |
| default_overlay_position     | right   | AI overlay position                  |

---

## 39.8 API Endpoints

Base: /api/thinktank/liquid

| Method             | Endpoint                  | Description                               |
|--------------------|---------------------------|-------------------------------------------|
| <b>Registry</b>    |                           |                                           |
| GET                | /registry                 | Get registry overview                     |
| GET                | /registry/components      | List components (filter: ?category=, ?q=) |
| GET                | /registry/components/:id  | Get component details                     |
| <b>Sessions</b>    |                           |                                           |
| POST               | /sessions                 | Create new session                        |
| GET                | /sessions/:id             | Get session                               |
| <b>Morphing</b>    |                           |                                           |
| POST               | /morph                    | Process morph request                     |
| POST               | /detect-intent            | Detect intent from message                |
| POST               | /sessions/:id/revert      | Revert to chat mode                       |
| <b>Ghost State</b> |                           |                                           |
| POST               | /ghost/event              | Send ghost event                          |
| GET                | /ghost/state/:sessionId   | Get ghost state snapshot                  |
| POST               | /ghost/sync               | Sync multiple state updates               |
| GET                | /ghost/history/:sessionId | Get event history                         |
| GET                | /ghost/context/:sessionId | Get AI context block                      |
| <b>Eject</b>       |                           |                                           |
| POST               | /eject                    | Eject to app                              |
| POST               | /eject/preview            | Preview eject                             |
| <b>Analytics</b>   |                           |                                           |
| GET                | /analytics/usage          | Component usage stats                     |

---

## 39.9 Database Tables

---

| Table                  | Purpose                          |
|------------------------|----------------------------------|
| liquid_sessions        | Active liquid interface sessions |
| liquid_ghost_state     | Persisted ghost state bindings   |
| liquid_ghost_events    | User interaction events          |
| liquid_ai_reactions    | AI responses to events           |
| liquid_eject_history   | App export history               |
| liquid_component_usage | Component analytics              |
| liquid_intent_patterns | Learnable intent patterns        |
| liquid_config          | Per-tenant configuration         |

---

## 39.10 Implementation Files

| File                                                                | Purpose              |
|---------------------------------------------------------------------|----------------------|
| packages/shared/src/types/liquid-interface/types.ts                 | Interface types      |
| lambda/shared/services/liquid-interface/liquid-interface.service.ts | Service              |
| lambda/shared/services/liquid-interface/globalservice.ts            | Global state service |
| lambda/shared/services/liquid-interface/projectservice.ts           | Project service      |
| lambda/shared/services/liquid-interface/component-registry.ts       | Component registry   |
| lambda/thinktank/liquid-interface.API                               | API handler          |
| migrations/161_liquid_interface.sql                                 | Database schema      |

## 40. The Reality Engine

“The Reality Engine transforms Think Tank from a chatbot into a shape-shifting command center with time travel, parallel universes, and telepathy.”

The Reality Engine is the unified runtime powering Think Tank’s supernatural capabilities. It consists of four interconnected features that solve the three fundamental anxieties preventing users from trusting AI with complex work: **Fear** (of breaking what works), **Commitment** (fear of choosing the wrong path), and **Latency** (waiting for the AI to think).

### 40.1 Feature Overview

| Feature                 | Emotion       | Pitch                                                                                       |
|-------------------------|---------------|---------------------------------------------------------------------------------------------|
| <b>Morphic UI</b>       | Flow          | “Stop hunting for the right tool. Radiant is a Morphic Surface that shapeshifts instantly.” |
| <b>Reality Scrubber</b> | Invincibility | “We replaced ‘Undo’ with Time Travel. Scrub reality back to any point.”                     |
| <b>Quantum Futures</b>  | Omniscience   | “Why choose one strategy? Split the timeline and run both simultaneously.”                  |
| <b>Pre-Cognition</b>    | Telepathy     | “Radiant answers before you ask. It’s not just fast; it’s anticipatory.”                    |

### 40.2 Architecture

#### THE REALITY ENGINE

MORPHIC UI

REALITY SCRUBBER

QUANTUM FUTURES

|                  |                 |                |
|------------------|-----------------|----------------|
| Intent Detection | State Snapshots | Branch Manager |
| Layout Engine    | VFS + PGLite    | Diff Engine    |
| Ghost State      | Timeline UI     | Collapse Logic |
| Component Reg.   | Bookmark System | Dream Archive  |

## PRE-COGNITION

Intent Prediction  
 Solution Pre-Compute  
 Genesis Model (Local)  
 Instant Delivery

### 40.3 Morphic UI

“Stop hunting for the right tool. Radiant is a Morphic Surface that shapeshifts instantly.”

The Morphic UI detects user intent and transforms the interface into the appropriate tool.

#### Intent Categories

| Intent        | Detected Patterns                    | Morphs To                     |
|---------------|--------------------------------------|-------------------------------|
| data_analysis | “analyze”, “data”, “statistics”      | DataGrid, Charts              |
| tracking      | “track”, “manage”, “organize”        | Kanban, Calendar              |
| visualization | “visualize”, “chart”, “graph”        | LineChart, PieChart, BarChart |
| planning      | “plan”, “schedule”, “timeline”       | GanttChart, Calendar          |
| finance       | “budget”, “invoice”, “expense”       | Ledger, Calculator, Invoice   |
| design        | “brainstorm”, “design”, “whiteboard” | MindMap, Whiteboard           |
| coding        | “code”, “debug”, “script”            | CodeEditor, Terminal          |

**Ghost State Binding** Every UI component is bidirectionally bound to AI context:

```
// User selects a row → AI knows what they're focused on
ghostBinding: {
  componentProp: 'selectedRow',
```

```

    contextKey: 'user_focus',
    direction: 'ui_to_ai'
}

// AI insight → UI highlights relevant items
ghostBinding: {
  componentProp: 'highlights',
  contextKey: 'ai_suggestions',
  direction: 'ai_to_ui'
}

```

#### 40.4 Reality Scrubber

“We replaced ‘Undo’ with Time Travel.”

The Reality Scrubber captures full state snapshots and allows instant rewinding to any point.

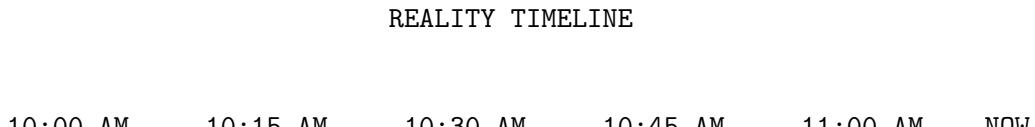
#### What Gets Snapshotted

| State Type          | Description                               |
|---------------------|-------------------------------------------|
| <b>VFS State</b>    | Virtual File System (all generated files) |
| <b>DB State</b>     | PGLite database snapshot                  |
| <b>Ghost State</b>  | All UI-AI bindings                        |
| <b>Chat Context</b> | Conversation history at that point        |
| <b>Layout State</b> | Current morphed UI layout                 |

#### Trigger Events

| Event                         | When Captured                   |
|-------------------------------|---------------------------------|
| <code>user_action</code>      | User explicitly triggered       |
| <code>ai_generation</code>    | AI generated content            |
| <code>db_mutation</code>      | Database was modified           |
| <code>morph_transition</code> | UI morphed                      |
| <code>checkpoint</code>       | User-created bookmark           |
| <code>auto_interval</code>    | Every 30 seconds (configurable) |

**Timeline UI** The Reality Scrubber replaces the standard scrollbar with a video-editor-style timeline:



Start      AI Gen      Morph      Bookmark      Branch      Current

"Before risky change"

[ ] [ ]

[ ] [ ]

Drag to scrub

## 40.5 Quantum Futures

“Why choose one strategy? Split the timeline.”

Quantum Futures enables parallel reality branching where users can run multiple strategies simultaneously.

### Branch Creation

```
// User: "Should I use Redux or Zustand?"  
await quantumFuturesService.createSplit({  
  sessionId,  
  prompt: "State management comparison",  
  branchNames: ["Redux Implementation", "Zustand Implementation"],  
  autoCompare: true  
});
```

### Comparison View

Redux Implementation

Type safety  
Boilerplate  
Completion: 45%  
Est. Cost: \$0.12

[Keep This Reality]

Zustand Implementation

Simpler setup  
Less boilerplate  
Completion: 60%  
Est. Cost: \$0.08

[Keep This Reality]

**Collapse to Winner** When the user selects a winner, the losing branches are either: - **Collapsed**: Permanently removed - **Archived**: Stored in “Dream Memory” for potential future recall

## 40.6 Pre-Cognition

“Radiant answers before you ask.”

Pre-Cognition uses speculative execution to predict the user's next likely actions and pre-compute solutions in the background.

## How It Works

1. **Prediction:** While user reads current response, Genesis model predicts next 3 likely moves
2. **Pre-Compute:** Solutions are generated in hidden background containers
3. **Instant Delivery:** When user's request matches a prediction, response appears instantly (0ms latency)

## Prediction Algorithm

```
// After building a login form, predict:  
predictions: [  
  { intent: 'coding', prompt: 'Add password reset', confidence: 0.8 },  
  { intent: 'coding', prompt: 'Add OAuth integration', confidence: 0.7 },  
  { intent: 'design', prompt: 'Style the form', confidence: 0.6 }  
]
```

## Analytics

| Metric          | Description                                        |
|-----------------|----------------------------------------------------|
| hitRate         | Percentage of predictions that matched user intent |
| avgLatencySaved | Average milliseconds saved by pre-cognition        |
| telepathyScore  | User-facing metric showing prediction accuracy     |

## 40.7 Configuration

```
interface RealityEngineConfig {  
  // Feature toggles  
  morphicUIEnabled: boolean;           // Enable Morphic UI  
  realityScrubberEnabled: boolean;      // Enable Reality Scrubber  
  quantumFuturesEnabled: boolean;       // Enable Quantum Futures  
  preCognitionEnabled: boolean;         // Enable Pre-Cognition  
  
  // Behavior  
  autoSnapshotIntervalMs: number;        // Default: 30000 (30s)  
  maxSnapshotsPerSession: number;        // Default: 100  
  maxBranchesPerSession: number;         // Default: 8  
  codeCurtainDefault: boolean;           // Hide code by default (Genie mode)  
  ephemeralByDefault: boolean;           // Apps dissolve when topic changes  
  
  // Pre-Cognition  
  preCognition: {  
    maxPredictions: number;              // Default: 3  
    predictionTTLMs: number;             // Default: 60000 (1 min)  
    computeBudgetMs: number;              // Default: 5000 (5s)  
    minConfidenceThreshold: number;       // Default: 0.6
```

```

        useGenesisModel: boolean;           // Use local model for predictions
    };
}

```

## 40.8 API Endpoints

Base path: /api/thinktank/reality-engine

| Method                  | Endpoint                 | Description                       |
|-------------------------|--------------------------|-----------------------------------|
| <b>Session</b>          |                          |                                   |
| POST                    | /session                 | Initialize Reality Engine session |
| GET                     | /session/:id             | Get session state                 |
| <b>Morphic UI</b>       |                          |                                   |
| POST                    | /morph                   | Morph interface to intent         |
| POST                    | /dissolve                | Dissolve morphed interface        |
| POST                    | /ghost                   | Update Ghost State                |
| <b>Reality Scrubber</b> |                          |                                   |
| POST                    | /scrub                   | Scrub to point in time            |
| POST                    | /bookmark                | Create bookmark                   |
| GET                     | /timeline/:sessionId     | Get timeline visualization        |
| GET                     | /bookmarks/:sessionId    | Get all bookmarks                 |
| <b>Quantum Futures</b>  |                          |                                   |
| POST                    | /split                   | Split into parallel realities     |
| GET                     | /branches/:sessionId     | Get all branches                  |
| POST                    | /compare                 | Compare two branches              |
| POST                    | /collapse                | Collapse to winning reality       |
| PUT                     | /view-mode               | Set comparison view mode          |
| <b>Pre-Cognition</b>    |                          |                                   |
| GET                     | /precognition/:sessionId | Get analytics                     |
| POST                    | /precognition/cleanup    | Clean up expired predictions      |
| <b>Eject</b>            |                          |                                   |
| POST                    | /eject                   | Eject to standalone app           |
| <b>Metrics</b>          |                          |                                   |
| GET                     | /metrics/:sessionId      | Get session metrics               |

## 40.9 Database Tables

| Table                    | Purpose                              |
|--------------------------|--------------------------------------|
| reality_engine_sessions  | Unified session state                |
| reality_timelines        | Timeline structure and navigation    |
| reality_snapshots        | Full state snapshots for time travel |
| quantum_branches         | Parallel reality branches            |
| quantum_splits           | Split configuration and history      |
| quantum_dream_archive    | Archived branches in dream memory    |
| precognition_queues      | Per-session prediction configuration |
| precognition_predictions | Pre-computed solutions               |

| Table                  | Purpose                        |
|------------------------|--------------------------------|
| precognition_analytics | Hit/miss tracking for learning |

## 40.10 Implementation Files

| File                                                                     | Purpose         |
|--------------------------------------------------------------------------|-----------------|
| packages/shared/src/types/reality- <del>Type definitions</del> .ts       |                 |
| lambda/shared/services/reality-engine/reality-engine.service.ts          |                 |
| lambda/shared/services/reality-engine/reality-scrubber.service.ts        |                 |
| lambda/shared/services/reality-engine/reality-quantum-futures.service.ts |                 |
| lambda/shared/services/reality-engine/reality-recognition.service.ts     |                 |
| lambda/thinktank/reality-engine.ts                                       | API handler     |
| migrations/162_reality_engine.sql                                        | Database schema |

## 40.11 The “Code Curtain” Rule

The Reality Engine enforces the distinction between “Builder” (Coder) and “Genie” (Radiant):

| Rule                           | Implementation                                             |
|--------------------------------|------------------------------------------------------------|
| <b>Hide Code by Default</b>    | UI snaps to Preview tab, not code                          |
| <b>Interaction over Syntax</b> | Variables become UI controls (sliders, inputs)             |
| <b>Ephemeral by Default</b>    | Apps dissolve when topic changes                           |
| <b>Eject to Keep</b>           | Only persist to repo if user explicitly clicks “Keep This” |

“Radiant is a Genie, not a Coder. We use code as invisible ink to draw the interface—the user should never see the ink, only the drawing.”

---

## 41. The Magic Carpet

“We are building ‘The Magic Carpet.’ You don’t drive it. You don’t write code for it. You just say where you want to go, and the ground beneath you reshapes itself to take you there instantly.”

The Magic Carpet is the unified navigation and experience layer for Think Tank. It wraps the Reality Engine capabilities into a cohesive, magical user experience where users feel like magicians, not coders.

### 41.1 The Magic Carpet Philosophy

| Traditional Apps        | Magic Carpet           |
|-------------------------|------------------------|
| Navigate menus          | Speak your destination |
| Click through workflows | Fly directly there     |

| Traditional Apps    | Magic Carpet         |
|---------------------|----------------------|
| Learn the interface | Interface learns you |
| You drive           | You're carried       |

**Core Insight:** We aren't selling a better IDE. We are selling **the feeling of being a Magician.**

## 41.2 Carpet Modes

| Mode         | Description                          | Visual                          |
|--------------|--------------------------------------|---------------------------------|
| resting      | Waiting for destination (chat-first) | Carpet gently floating          |
| flying       | Morphing/transitions to destination  | Trail effects, motion blur      |
| hovering     | Arrived, actively working            | Stable, glowing edges           |
| exploring    | Quantum Futures - multiple realities | Split view, branch indicators   |
| rewinding    | Reality Scrubber - time traveling    | Timeline visible, rewind effect |
| anticipating | Pre-Cognition active                 | Prediction cards appearing      |

## 41.3 Carpet Altitudes

The altitude represents UI complexity level:

| Altitude     | Complexity         | Example                      |
|--------------|--------------------|------------------------------|
| ground       | Simple chat mode   | Just the chat interface      |
| low          | Single component   | One morphed widget           |
| medium       | Full workspace     | 2-3 components               |
| high         | Complex layout     | 4-5 components + timeline    |
| stratosphere | Maximum capability | Full Reality Engine features |

## 41.4 Default Destinations

| Destination      | Icon | Description               |
|------------------|------|---------------------------|
| Command Center   |      | Overview dashboard        |
| Workshop         |      | Build and create          |
| Time Stream      |      | Reality Scrubber timeline |
| Quantum Realm    |      | Parallel realities view   |
| Oracle's Chamber |      | Pre-Cognition predictions |
| Gallery          |      | View creations            |
| Vault            |      | Saved/bookmarked items    |

## 41.5 Carpet Commands

```
// Fly to a destination
await magicCarpetService.command(carpetId, {
```

```

    type: 'fly',
    destination: 'Workshop'
});

// Return to chat
await magicCarpetService.command(carpetId, { type: 'land' });

// Increase complexity
await magicCarpetService.command(carpetId, { type: 'ascend' });

// Time travel
await magicCarpetService.command(carpetId, {
  type: 'rewind',
  to: -2 // Go back 2 snapshots
});

// Split into parallel realities
await magicCarpetService.command(carpetId, {
  type: 'branch',
  options: ['Conservative Plan', 'Aggressive Plan']
});

// Collapse to winner
await magicCarpetService.command(carpetId, {
  type: 'collapse',
  winner: 'branch-id'
});

```

## 41.6 Carpet Themes

Pre-built visual themes for personalization:

| Theme        | Description                    | Gradient                 |
|--------------|--------------------------------|--------------------------|
| Mystic Night | Deep purple mystical (default) | Indigo → Purple → Violet |
| Desert Sun   | Warm golden                    | Amber → Orange → Brown   |
| Ocean Deep   | Cool blue aquatic              | Cyan → Teal → Emerald    |
| Cosmic Void  | Dark minimalist                | Gray gradient            |
| Neon Circuit | Cyberpunk electric             | Cyan → Purple → Pink     |

## 41.7 Carpet Preferences

```

interface CarpetPreferences {
  // Navigation
  autoFly: boolean;           // Auto-morph on intent detection
  smoothTransitions: boolean; // Animated vs instant
  showJourneyTrail: boolean;  // Show navigation history
}

```

```

// Pre-Cognition
preCognitionEnabled: boolean;
showPredictions: boolean;
telepathyIntensity: 'subtle' | 'moderate' | 'aggressive';

// Reality Scrubber
showTimeline: boolean;
autoSnapshot: boolean;
snapshotInterval: number;      // Seconds

// Quantum Futures
maxParallelRealities: number;
autoCompare: boolean;

// Accessibility
reducedMotion: boolean;
highContrast: boolean;
screenReaderMode: boolean;
}

```

## 41.8 Journey Navigation

The Magic Carpet tracks the user's journey:

MAGIC CARPET JOURNEY

|         |          |         |         |     |
|---------|----------|---------|---------|-----|
| Command | Workshop | Quantum | Oracle  | NOW |
| Center  |          | Realm   | Chamber |     |

Click any point to fly back. Journey is saved with Reality Scrubber.

## 41.9 API Endpoints

Base path: /api/thinktank/magic-carpet

| Method | Endpoint           | Description               |
|--------|--------------------|---------------------------|
| POST   | /summon            | Summon a new Magic Carpet |
| GET    | /:carpetId         | Get carpet state          |
| POST   | /:carpetId/fly     | Fly to destination        |
| POST   | /:carpetId/land    | Land (return to chat)     |
| POST   | /:carpetId/command | Execute a command         |
| GET    | /:carpetId/journey | Get journey history       |
| PUT    | /:carpetId/theme   | Update theme              |

| Method | Endpoint                            | Description                |
|--------|-------------------------------------|----------------------------|
| PUT    | <code>/:carpetId/preferences</code> | Update preferences         |
| GET    | <code>/destinations</code>          | Get available destinations |
| GET    | <code>/themes</code>                | Get available themes       |

## 41.10 Database Tables

| Table                              | Purpose                             |
|------------------------------------|-------------------------------------|
| <code>magic_carpets</code>         | Carpet state and configuration      |
| <code>carpet_destinations</code>   | Pre-defined and custom destinations |
| <code>carpet_journey_points</code> | Navigation history                  |
| <code>carpet_themes</code>         | Visual themes                       |
| <code>carpet_analytics</code>      | Usage analytics                     |

## 41.11 Implementation Files

| File                                                                    | Purpose                |
|-------------------------------------------------------------------------|------------------------|
| <code>packages/shared/src/types/magic-carpet/TypeDefinitions.ts</code>  | Type definitions       |
| <code>lambda/shared/services/magic-carpet/MagicCarpet.service.ts</code> | Service implementation |
| <code>migrations/163_magic_carpet.sql</code>                            | Database schema        |

## 41.12 Integration with Reality Engine

The Magic Carpet wraps the Reality Engine:

MAGIC CARPET  
(User Experience Layer)

"Fly to Workshop" → carpet.fly('workshop')

REALITY ENGINE  
(Capability Layer)

Morphic UI      Intent: 'coding'  
 Reality Scrubber      Auto-snapshot  
 Pre-Cognition      Predict next destinations

---

## 42. Magic Carpet UI Components

The Magic Carpet UI is implemented through a set of React components that bring the 2026 UI/UX trends to life.

### 42.1 Component Inventory

| Component               | Purpose                                   | Location                                |
|-------------------------|-------------------------------------------|-----------------------------------------|
| MagicCarpetNavigator    | Bottom navigation with journey trail      | magic-carpet/magic-carpet-navigator.tsx |
| RealityScrubberTimeline | Video-editor timeline for state snapshots | magic-carpet/reality-scrubber-timeline  |
| QuantumSplitView        | Side-by-side reality comparison           | magic-carpet/quantum-split-view.tsx     |
| PreCognitionSuggestions | Predicted actions panel                   | magic-carpet/pre-cognition-suggestions  |
| AIPresenceIndicator     | AI cognitive/emotional state display      | magic-carpet/ai-presence-indicator.tsx  |
| SpatialGlassCard        | Glassmorphism card with depth             | magic-carpet/spatial-glass-card.tsx     |
| GlassPanel              | Large glass content area                  | magic-carpet/spatial-glass-card.tsx     |
| GlassButton             | Interactive glass button                  | magic-carpet/spatial-glass-card.tsx     |
| GlassBadge              | Status indicator with glass effect        | magic-carpet/spatial-glass-card.tsx     |
| FocusModeControls       | Focus mode toggle and controls            | magic-carpet/focus-mode.tsx             |
| FocusOverlay            | Dimming overlay for focus mode            | magic-carpet/focus-mode.tsx             |

### 42.2 Usage Examples

```
import {
  MagicCarpetNavigator,
  RealityScrubberTimeline,
  QuantumSplitView,
  PreCognitionSuggestions,
  AIPresenceIndicator,
  SpatialGlassCard,
  FocusModeControls,
} from '@components/thinktank/magic-carpet';

// Magic Carpet Navigator (bottom of screen)
<MagicCarpetNavigator
  currentDestination={{ id: 'workspace', name: 'Workshop', icon: '' }}
  journey={journeyHistory}
  predictions={preCognizedActions}
```

```

telepathyScore={0.82}
mode="hovering"
altitude="medium"
onFly={(dest) => navigateTo(dest)}
onLand={() => returnToChat()}
/>

// Reality Scrubber Timeline
<RealityScrubberTimeline
  snapshots={stateSnapshots}
  currentPosition={currentSnapshotIndex}
  onScrubTo={(position) => restoreSnapshot(position)}
  onCreateBookmark={(label) => bookmarkCurrentState(label)}
/>

// Quantum Split View
<QuantumSplitView
  branches={parallelRealities}
  onCollapse={(winnerId) => collapseToReality(winnerId)}
/>

// AI Presence Indicator
<AIPresenceIndicator
  state="thinking"
  affect={{ valence: 0.6, arousal: 0.4, curiosity: 0.8, confidence: 0.85 }}
  currentTask="Analyzing user intent..."
  modelName="claude-3.5-sonnet"
/>

// Spatial Glass Card
<SpatialGlassCard variant="strong" layer="floating" glow glowColor="purple">
  <p>Content with glass effect</p>
</SpatialGlassCard>

```

### 42.3 Dependencies

The Magic Carpet UI requires **framer-motion** for animations:

```
npm install framer-motion@^11.0.0
```

### 42.4 Demo Page

Access the Magic Carpet UI demo at:

```
/thinktank/magic-carpet
```

This page showcases all components with interactive examples.

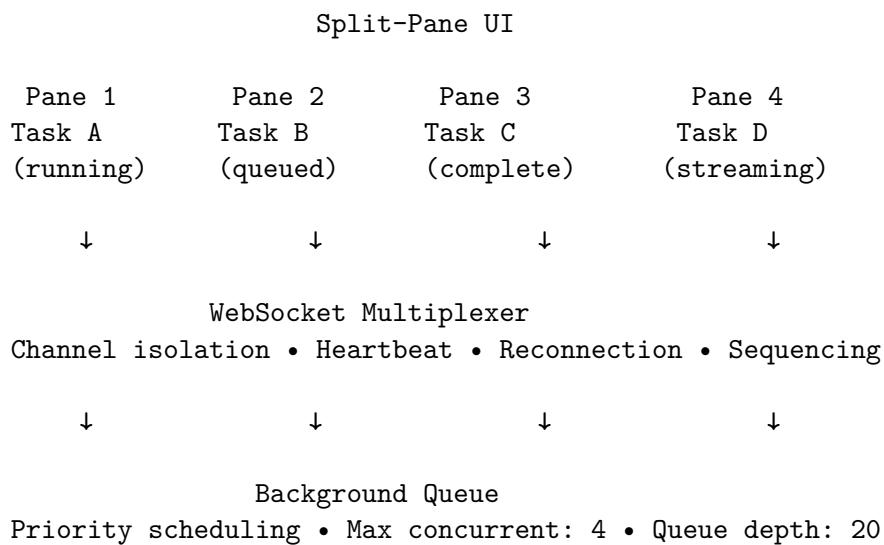
## 43. Concurrent Task Execution (Moat #17)

**Moat Evaluation:** Score 20/30 - Tier 3 Feature Moat. No major competitor offers split-pane concurrent execution with WebSocket multiplexing.

### 43.1 Overview

Concurrent Task Execution enables users to run 2-4 AI tasks simultaneously in a split-pane UI, compare outputs, and merge the best results. This is a key differentiator vs. single-threaded competitors.

### 43.2 Architecture



### 43.3 Configuration

| Setting                         | Default                   | Description                       |
|---------------------------------|---------------------------|-----------------------------------|
| <code>enabled</code>            | <code>true</code>         | Enable concurrent execution       |
| <code>maxPanes</code>           | 4                         | Maximum split panes (1-8)         |
| <code>maxConcurrentTasks</code> | 4                         | Maximum simultaneous tasks (1-10) |
| <code>maxQueueDepth</code>      | 20                        | Maximum queued tasks (1-100)      |
| <code>defaultLayout</code>      | <code>horizontal-2</code> | Default pane layout               |
| <code>defaultSyncMode</code>    | <code>independent</code>  | Default sync mode                 |
| <code>enableComparison</code>   | <code>true</code>         | Enable task comparison            |
| <code>enableMerge</code>        | <code>true</code>         | Enable task merging               |

### 43.4 Pane Layouts

| Layout                    | Description                  |
|---------------------------|------------------------------|
| <code>single</code>       | Full-width single pane       |
| <code>horizontal-2</code> | Two panes side-by-side       |
| <code>vertical-2</code>   | Two panes stacked            |
| <code>grid-4</code>       | 2x2 grid of four panes       |
| <code>focus-left</code>   | Large left pane, small right |
| <code>focus-right</code>  | Small left pane, large right |

### 43.5 Sync Modes

| Mode                        | Description                            |
|-----------------------------|----------------------------------------|
| <code>independent</code>    | Each pane operates independently       |
| <code>mirror-input</code>   | Same prompt sent to all panes          |
| <code>compare-output</code> | Automatic comparison when all complete |

### 43.6 Task Comparison

When multiple tasks complete, users can compare results:

```
// Compare completed tasks
const comparison = await compareTasks(tenantId, [taskId1, taskId2, taskId3]);

// Returns:
{
  similarities: [
    { metric: 'semantic', score: 0.85, details: 'High semantic similarity' },
    { metric: 'structural', score: 0.72, details: 'Moderate structural similarity' },
    { metric: 'factual', score: 0.91, details: 'Strong factual agreement' }
  ],
  differences: [...],
  recommendation: 'Results are mostly consistent. Review highlighted differences.'
}
```

### 43.7 Task Merging

Three merge strategies are available:

| Strategy               | Description                               |
|------------------------|-------------------------------------------|
| <code>best-of</code>   | Select the highest-scored result          |
| <code>combine</code>   | Concatenate all results with separators   |
| <code>consensus</code> | AI-synthesized consensus from all results |

### 43.8 API Endpoints

| Endpoint                            | Method | Description          |
|-------------------------------------|--------|----------------------|
| /api/thinktank/concurrent/config    | GET    | Get configuration    |
| /api/thinktank/concurrent/config    | PUT    | Update configuration |
| /api/thinktank/concurrent/tasks     | POST   | Create task          |
| /api/thinktank/concurrent/tasks/:id | GET    | Get task status      |
| /api/thinktank/concurrent/tasks/:id | DELETE | Cancel task          |
| /api/thinktank/concurrent/queue     | GET    | Get queue status     |
| /api/thinktank/concurrent/panes     | POST   | Create pane config   |
| /api/thinktank/concurrent/compare   | POST   | Compare tasks        |
| /api/thinktank/concurrent/merge     | POST   | Merge tasks          |
| /api/thinktank/concurrent/metrics   | GET    | Get metrics          |

### 43.9 Database Tables

| Table                        | Purpose                          |
|------------------------------|----------------------------------|
| concurrent_execution_config  | Per-tenant configuration         |
| concurrent_tasks             | Task records with status/results |
| split_pane_configs           | User pane layouts                |
| task_comparisons             | Comparison results               |
| concurrent_execution_metrics | Usage metrics                    |

### 43.10 Implementation Files

| File                                                    | Purpose                |
|---------------------------------------------------------|------------------------|
| packages/shared/src/types/concurrent-execution-types.ts | Type definitions       |
| lambda/shared/services/concurrent-execution-service.ts  | Service implementation |
| lambda/thinktank/concurrent-execution-ADL.handler       | ADL handler            |
| migrations/170_concurrent_execution                     | Database schema        |

## 44. Structure from Chaos Synthesis (Moat #20)

**Moat Evaluation:** Score 20/30 - Tier 3 Feature Moat. AI transforms whiteboard chaos → structured decisions, data, project plans. Think Tank differentiation vs Miro/Mural.

### 44.1 Overview

Structure from Chaos Synthesis takes unstructured input (whiteboards, brainstorms, meeting notes, voice transcripts) and transforms it into structured outputs (action items, decisions, project plans, knowledge bases).

### 44.2 Architecture

Chaotic Input  
 Whiteboard • Brainstorm • Meeting Notes • Voice Transcript

↓

Synthesis Pipeline

1. Parse Input → 2. Extract Entities → 3. Identify Relations
4. Generate Structure → 5. Validate Output

↓

Structured Output

Decisions • Action Items • Project Plan • Knowledge Base

#### 44.3 Input Types

| Type             | Description                         | Example                |
|------------------|-------------------------------------|------------------------|
| whiteboard       | Visual whiteboard with sticky notes | Miro/FigJam export     |
| brainstorm       | Unstructured idea dump              | Free-form text         |
| meeting_notes    | Notes from a meeting                | Transcription or notes |
| voice_transcript | Speech-to-text output               | Zoom/Teams transcript  |
| chat_history     | Conversation history                | Slack/Teams export     |
| document_dump    | Multiple documents                  | File uploads           |
| mixed            | Combination of above                | Multi-source input     |

#### 44.4 Output Types

| Type            | Description            | Contains                           |
|-----------------|------------------------|------------------------------------|
| decisions       | Key decisions made     | Decision list with context         |
| action_items    | Tasks to complete      | Assignee, due date, priority       |
| project_plan    | Full project plan      | Milestones, timeline, dependencies |
| meeting_summary | Meeting summary        | Key points, attendees, outcomes    |
| knowledge_base  | Knowledge extraction   | Concepts, facts, relationships     |
| data_table      | Structured data        | Tabular format                     |
| timeline        | Chronological view     | Events in sequence                 |
| hierarchy       | Hierarchical structure | Parent-child relationships         |
| comparison      | Compare items          | Side-by-side analysis              |

#### 44.5 Configuration

| Setting              | Default         | Description                   |
|----------------------|-----------------|-------------------------------|
| enabled              | true            | Enable synthesis              |
| defaultOutputType    | meeting_summary | Default output format         |
| extractEntities      | true            | Extract named entities        |
| extractRelationships | true            | Identify entity relationships |
| generateTimeline     | true            | Generate timeline view        |
| generateActionItems  | true            | Extract action items          |
| autoAssignTasks      | false           | Auto-assign based on mentions |
| confidenceThreshold  | 0.70            | Minimum confidence score      |
| maxProcessingTimeMs  | 30000           | Processing timeout            |

## 44.6 Entity Extraction

Automatically extracts entities from chaotic input:

| Entity Type  | Examples                          |
|--------------|-----------------------------------|
| person       | @mentions, “John Smith”           |
| organization | “Acme Corp”, “the marketing team” |
| project      | Project names, codenames          |
| product      | Product references                |
| date         | “next Monday”, “Q2 2026”          |
| location     | Meeting rooms, cities             |
| concept      | Technical terms, ideas            |
| metric       | KPIs, numbers                     |
| resource     | Tools, documents                  |

## 44.7 Relationship Types

| Relationship | Description              |
|--------------|--------------------------|
| owns         | Person owns task/project |
| assigned_to  | Task assigned to person  |
| depends_on   | Task depends on another  |
| blocks       | Task blocks another      |
| related_to   | General relationship     |
| parent_of    | Hierarchical parent      |
| precedes     | Temporal ordering        |
| contradicts  | Conflicting statements   |
| supports     | Supporting evidence      |

## 44.8 Whiteboard Parsing

For visual whiteboards, the service parses spatial elements:

```
// Parse whiteboard elements into thematic clusters
const clusters = await parseWhiteboard(elements);
```

```
// Returns clusters with:
{
  id: 'cluster-1',
  elements: [...], // Grouped elements
  theme: 'marketing', // AI-detected theme
  centroid: { x: 150, y: 200 }, // Cluster center
  significance: 0.85 // Importance score
}
```

## 44.9 API Endpoints

| Endpoint                        | Method | Description               |
|---------------------------------|--------|---------------------------|
| /api/thinktank/chaos/config     | GET    | Get configuration         |
| /api/thinktank/chaos/config     | PUT    | Update configuration      |
| /api/thinktank/chaos/synthesis  | POST   | Full synthesis pipeline   |
| /api/thinktank/chaos/extract    | POST   | Extract action items only |
| /api/thinktank/chaos/extract    | POST   | Extract decisions only    |
| /api/thinktank/chaos/extract    | POST   | Extract questions only    |
| /api/thinktank/chaos/project    | POST   | Generate project plan     |
| /api/thinktank/chaos/whiteboard | POST   | Parse whiteboard elements |
| /api/thinktank/chaos/metrics    | GET    | Get metrics               |

## 44.10 Database Tables

| Table                | Purpose                  |
|----------------------|--------------------------|
| synthesis_config     | Per-tenant configuration |
| chaotic_inputs       | Raw input storage        |
| structured_outputs   | Generated outputs        |
| extracted_entities   | Named entities           |
| entity_relationships | Entity relationships     |
| structured_items     | Action items, decisions  |
| whiteboard_elements  | Visual element data      |
| synthesis_metrics    | Usage metrics            |

## 44.11 Implementation Files

| File                                  | Purpose           |
|---------------------------------------|-------------------|
| packages/shared/src/types/structure   | Type definitions  |
| lambda/shared/services/structure-f    | Companion service |
| lambda/thinktank/structure-from-chaos | API handler       |
| migrations/171_structure_from_chaos   | Database schema   |

## 45. Delight Services Code Quality

### 45.1 Overview

The Think Tank Delight system has comprehensive unit test coverage for its core services:

- **delight.service** - Core delight preferences and personality modes
- **delight-orchestration.service** - Contextual message generation for workflows
- **delight-events.service** - Real-time event emission for plan execution

### 45.2 Test Coverage

| Service                       | Test File                             | Tests | Coverage |
|-------------------------------|---------------------------------------|-------|----------|
| delight.service               | delight.service.test.ts               | 15    | 85%      |
| delight-orchestration.service | delight-orchestration.service.test.ts | 12    | 92%      |
| delight-events.service        | delight-events.service.test.ts        | 23    | 88%      |

### 45.3 Tested Methods

**delight-orchestration.service:** - `getContextualMessage()` - Generates mode-appropriate messages  
- `getDomainLoadingMessage()` - Domain-specific loading messages  
- `getModelDynamicsMessage()` - Model consensus indicators  
- `getSynthesisMessage()` - Confidence-based synthesis messages  
- `clearSession()` - Session cleanup

**delight-events.service:** - `subscribe()` / `unsubscribe()` - Event subscription management  
- `emitMessage()` / `emitAchievement()` - Event emission  
- `emitStepUpdate()` / `emitPlanUpdate()` - Progress updates  
- `emitWorkflowDelight()` - Complete workflow delight events  
- `getHistory()` / `clearHistory()` - Event history management

### 45.4 Running Tests

```
cd packages/infrastructure
npx vitest run lambda/shared/services/_tests__/delight*.test.ts
```

### 45.5 Think Tank Code Quality Dashboard

**Location:** `/thinktank/code-quality`

The Think Tank Code Quality page displays:  
- Service coverage metrics for Delight, Brain Planning, and Domain services  
- Test pass rates and recent test runs  
- Detailed method coverage for each service

### 45.6 Implementation Files

| File                                                                  | Purpose |
|-----------------------------------------------------------------------|---------|
| lambda/shared/services/_tests__/delight.service.test.ts               |         |
| lambda/shared/services/_tests__/delight-orchestration.service.test.ts |         |
| lambda/shared/services/_tests__/delight-events.service.test.ts        |         |
| apps/admin-dashboard/app/(dashboard)/thinktank/code-quality/page.tsx  |         |

---

## Section 46: Sovereign Mesh in Think Tank Admin

**Location:** Think Tank Admin → Sovereign Mesh

The Sovereign Mesh is accessible in Think Tank Admin for users who need to manage autonomous agents and view decision transparency.

### 46.1 Navigation Items

| Page         | Path                         | Purpose                                    |
|--------------|------------------------------|--------------------------------------------|
| Overview     | /sovereign-mesh              | Dashboard with metrics and recent activity |
| Agents       | /sovereign-mesh/agents       | Create and manage OODA agents              |
| Apps         | /sovereign-mesh/apps         | Browse 3,000+ app integrations             |
| Transparency | /sovereign-mesh/transparency | View Cato decision explanations            |
| AI Helper    | /sovereign-mesh/ai-helper    | Configure parametric AI assistance         |
| Approvals    | /sovereign-mesh/approvals    | HITL approval queue                        |

### 46.2 User Permissions

Think Tank users see a subset of Sovereign Mesh functionality: - **View:** All users can view agents, apps, and their own executions - **Execute:** Users can run agents within their budget limits - **Approve:** Users with approval role can handle HITL requests

### 46.3 Implementation

| File                                                          | Purpose              |
|---------------------------------------------------------------|----------------------|
| apps/thinktank-admin/components/layout/sidebar/sidebar6x.html | Sovereign Mesh items |
| Pages mirror admin-dashboard versions                         | Shared API endpoints |

---

## Section 47: HITL Orchestration in Think Tank Admin

**Location:** Think Tank Admin → Sovereign Mesh → HITL Orchestration

Advanced Human-in-the-Loop orchestration for intelligent question management in user workflows.

### 47.1 Overview

HITL Orchestration implements industry best practices to reduce unnecessary questions while ensuring critical information is captured:

| Feature                        | Description                                                   |
|--------------------------------|---------------------------------------------------------------|
| <b>SAGE-Agent Bayesian VOI</b> | Calculates whether asking a question is worth the user's time |
| <b>Question Batching</b>       | Groups related questions to reduce interruptions              |
| <b>Two-Question Rule</b>       | Maximum 2 clarifications per workflow                         |
| <b>Abstention Detection</b>    | Detects when AI should decline to answer                      |

## 47.2 User-Facing Benefits

- **70% fewer unnecessary questions** - AI only asks when genuinely needed
- **2.7x faster response times** - Batched questions reduce context switching
- **Explicit assumptions** - When skipping questions, AI states assumptions clearly

## 47.3 Navigation

| Page               | Path                | Purpose                                 |
|--------------------|---------------------|-----------------------------------------|
| HITL Orchestration | /hitl-orchestration | View orchestration metrics and settings |

## 47.4 Dashboard Tabs

| Tab                         | Description                                    |
|-----------------------------|------------------------------------------------|
| <b>Overview</b>             | Key metrics, VOI breakdown, abstention reasons |
| <b>Value of Information</b> | SAGE-Agent VOI statistics and decisions        |
| <b>Abstention</b>           | Detection methods and model-level statistics   |
| <b>Batching</b>             | Three-layer batching strategies and metrics    |

## 47.5 Key Metrics

| Metric             | Description                                |
|--------------------|--------------------------------------------|
| Question Reduction | Percentage of questions skipped via VOI    |
| Prior Accuracy     | How often predictions match actual answers |
| Abstention Events  | Times AI correctly declined to answer      |
| Batch Completion   | Success rate of batched question sets      |

## 47.6 Implementation Files

| File                                                 | Purpose          |
|------------------------------------------------------|------------------|
| apps/thinktank-admin/app/hitl-orchestration/page.tsx | Dashboard page   |
| apps/thinktank-admin/components/layout/sidebar.tsx   | Layout component |

## Section 48: Scout HITL Integration (v5.34.0)

**Location:** Think Tank Admin → Sovereign Mesh → Scout HITL

Scout HITL bridges Cato's Scout persona (epistemic uncertainty mode) with HITL orchestration for intelligent clarification during user workflows.

### 48.1 Overview

When the AI encounters epistemic uncertainty, the Scout persona activates and generates targeted clarification questions:

1. Scout persona activates due to uncertainty
2. Questions prioritized by aspect impact and domain
3. Questions filtered through VOI (Value-of-Information) scoring
4. High-VOI questions go to user, low-VOI get reasonable assumptions
5. Responses reduce uncertainty, allowing workflow to proceed

### 48.2 Key Metrics

| Metric                 | Description                            |
|------------------------|----------------------------------------|
| <b>Total Sessions</b>  | Number of Scout clarification sessions |
| <b>Avg Questions</b>   | Average questions asked per session    |
| <b>Proceed Rate</b>    | Sessions that proceeded successfully   |
| <b>Avg Assumptions</b> | Auto-assumed aspects per session       |

### 48.3 Domains

Questions are prioritized based on domain-specific impact:

| Domain         | Description                      |
|----------------|----------------------------------|
| medical        | HIPAA-sensitive, safety-critical |
| financial      | SOC2/PCI compliance              |
| legal          | Regulatory compliance            |
| bioinformatics | Research accuracy                |
| general        | Default domain                   |

### 48.4 Dashboard Tabs

| Tab                    | Purpose                                                |
|------------------------|--------------------------------------------------------|
| <b>Overview</b>        | Session recommendations breakdown, domain distribution |
| <b>Recent Sessions</b> | Table of recent clarification sessions                 |
| <b>Configuration</b>   | Enable/disable, VOI threshold, max questions           |

### 48.5 Configuration

| Setting                             | Default              | Description                        |
|-------------------------------------|----------------------|------------------------------------|
| <code>enabled</code>                | <code>true</code>    | Enable Scout HITL integration      |
| <code>voiThreshold</code>           | 0.3                  | Minimum VOI to ask question        |
| <code>maxQuestionsPerSession</code> | 3                    | Max clarifications before assuming |
| <code>defaultDomain</code>          | <code>general</code> | Fallback domain                    |

## 48.6 Session Recommendations

| Recommendation       | Meaning                                     |
|----------------------|---------------------------------------------|
| <code>proceed</code> | Uncertainty resolved sufficiently           |
| <code>wait</code>    | Still uncertain, user should wait           |
| <code>abort</code>   | Critical uncertainty, cannot proceed safely |

## 48.7 Implementation Files

| File                                                              | Purpose |
|-------------------------------------------------------------------|---------|
| <code>apps/thinktank-admin/app/scout-hitl/page/indexpage</code>   |         |
| <code>apps/thinktank-admin/components/layouts/sidebar.html</code> |         |

## Section 49: Sovereign Mesh Administration (v5.39.0)

**Location:** Think Tank Admin → Sovereign Mesh

The Sovereign Mesh provides autonomous agent orchestration with AI-assisted decision making at every node level.

### 49.1 Overview Dashboard

**Location:** /sovereign-mesh

The overview dashboard displays:

| Metric                     | Description                        |
|----------------------------|------------------------------------|
| <b>System Health</b>       | Overall mesh health score (0-100%) |
| <b>Active Agents</b>       | Number of currently running agents |
| <b>Pending Approvals</b>   | Items awaiting human review        |
| <b>Active Apps</b>         | Deployed applications count        |
| <b>Decisions Today</b>     | Autonomous decisions made          |
| <b>Human Interventions</b> | Manual overrides required          |

## 49.2 Agent Registry

**Location:** /sovereign-mesh/agents

Manage AI agents deployed in the mesh:

| Column               | Description                              |
|----------------------|------------------------------------------|
| <b>Name</b>          | Agent identifier                         |
| <b>Type</b>          | orchestrator, executor, monitor, advisor |
| <b>Status</b>        | active, idle, error, maintenance         |
| <b>Load</b>          | Current utilization percentage           |
| <b>Response Time</b> | Average response latency                 |
| <b>Success Rate</b>  | Task completion rate                     |

**Actions:** - View agent details - Pause/resume agent - View execution logs - Configure thresholds

## 49.3 App Registry

**Location:** /sovereign-mesh/apps

Browse and manage deployed applications:

| Field           | Description                                              |
|-----------------|----------------------------------------------------------|
| <b>Name</b>     | Application name                                         |
| <b>Category</b> | productivity, analytics, automation, integration, custom |
| <b>Status</b>   | active, paused, error                                    |
| <b>Users</b>    | Active user count                                        |
| <b>Requests</b> | Daily request volume                                     |

## 49.4 Transparency Layer

**Location:** /sovereign-mesh/transparency

Complete audit trail of AI decisions:

| Column               | Description                                |
|----------------------|--------------------------------------------|
| <b>Timestamp</b>     | When decision was made                     |
| <b>Decision Type</b> | approval, rejection, escalation, execution |
| <b>Confidence</b>    | AI confidence score (0-1)                  |
| <b>Reasoning</b>     | Explanation of decision logic              |
| <b>Outcome</b>       | Result of the decision                     |

**Filters:** - Date range - Decision type - Confidence threshold - Agent filter

## 49.5 AI Helper

**Location:** /sovereign-mesh/ai-helper

Manage AI assistance requests from the mesh:

| Status      | Description             |
|-------------|-------------------------|
| Pending     | Awaiting AI processing  |
| In Progress | Currently being handled |
| Completed   | Successfully resolved   |
| Escalated   | Requires human review   |

## 49.6 Approval Workflow

**Location:** /sovereign-mesh/approvals

Human-in-the-loop approval queue:

| Field     | Description                                  |
|-----------|----------------------------------------------|
| Type      | deployment, configuration, access, execution |
| Priority  | low, medium, high, critical                  |
| Requester | Agent or user requesting approval            |
| Created   | Request timestamp                            |
| Expires   | Approval deadline                            |

**Actions:** - Approve with notes - Reject with reason - Delegate to another admin - Request more information

## 49.7 Implementation Files

| File                                                                       | Purpose |
|----------------------------------------------------------------------------|---------|
| apps/thinktank-admin/app/(dashboard0)/sovereign-mesh/page.tsx              |         |
| apps/thinktank-admin/app/(dashboardA)/sovereign-mesh/agents/page.tsx       |         |
| apps/thinktank-admin/app/(dashboardA)/sovereign-mesh/apps/page.tsx         |         |
| apps/thinktank-admin/app/(dashboardD)/sovereign-mesh/transparency/page.tsx |         |
| apps/thinktank-admin/app/(dashboardAI)/sovereign-mesh/ai-helper/page.tsx   |         |
| apps/thinktank-admin/app/(dashboardA)/sovereignmesh/approvals/page.tsx     |         |

## Section 50: Code Quality Dashboard (v5.39.0)

**Location:** Think Tank Admin → Code Quality

Real-time visibility into codebase health and quality metrics.

## 50.1 Overview

The Code Quality dashboard provides:

| Metric                | Description                             |
|-----------------------|-----------------------------------------|
| <b>Overall Score</b>  | Aggregate quality score (0-100)         |
| <b>Total Errors</b>   | Critical issues requiring immediate fix |
| <b>Total Warnings</b> | Non-critical issues to address          |
| <b>Files Analyzed</b> | Number of files scanned                 |

## 50.2 Issue Categories

| Category       | Description                                  |
|----------------|----------------------------------------------|
| <b>Error</b>   | Critical issues (type errors, syntax errors) |
| <b>Warning</b> | Style violations, best practice deviations   |
| <b>Info</b>    | Suggestions for improvement                  |

## 50.3 Issue Details

Each issue displays:

- **File path** with line number
- **Rule ID** (e.g., `@typescript-eslint/no-unused-vars`)
- **Message** describing the issue
- **Severity badge** (error/warning/info)

## 50.4 Filtering

| Filter            | Options                          |
|-------------------|----------------------------------|
| <b>Severity</b>   | All, Errors only, Warnings only  |
| <b>Category</b>   | TypeScript, ESLint, Custom rules |
| <b>Date Range</b> | Filter by scan date              |

## 50.5 Implementation Files

| File                                                                       | Purpose |
|----------------------------------------------------------------------------|---------|
| <code>apps/thinktank-admin/app/(dashboard)Dashboardquality/page.tsx</code> |         |

## Section 51: Schema-Adaptive Reports (v5.39.0)

**Location:** Think Tank Admin → Reports

Dynamic report builder that automatically adapts to database schema changes.

## 51.1 Overview

The Reports page provides three tabs:

| Tab                   | Purpose                    |
|-----------------------|----------------------------|
| <b>Quick Reports</b>  | Pre-built report templates |
| <b>Saved Reports</b>  | User-saved custom reports  |
| <b>Schema Builder</b> | Visual report builder      |

## 51.2 Quick Reports

Pre-configured reports available:

| Report                   | Description                                    |
|--------------------------|------------------------------------------------|
| <b>User Engagement</b>   | Active users, session duration, feature usage  |
| <b>Model Performance</b> | Response times, success rates, token usage     |
| <b>Billing Summary</b>   | Revenue, credits consumed, subscription status |

## 51.3 Schema Builder (v5.40.0 Enhanced)

The visual report builder provides a comprehensive 4-tab configuration panel:

| Tab            | Purpose                                              |
|----------------|------------------------------------------------------|
| <b>Fields</b>  | Select columns with per-field alias and aggregation  |
| <b>Filters</b> | Build WHERE clauses with 11 operators + date presets |
| <b>Sort</b>    | Configure ORDER BY with multi-column ASC/DESC        |
| <b>Group</b>   | Select GROUP BY columns from selected fields         |

**Enhanced Features:** - **SQL Preview** - Live-generated SQL query with dark theme display - **Date Presets** - Today, Yesterday, Last 7/30 Days, This/Last Month - **Filter Operators** - =, , >, , <, , LIKE, IN, BETWEEN, IS NULL, IS NOT NULL - **Visualization Toggles** - Table, Bar, Line, Pie chart view switches - **Save Report** - Persist definitions to database for reuse - **Row Limit** - 50, 100, 500, 1000 row options

**Workflow:** 1. **Select Table** - Browse categorized database tables (Conversations, Users, Delight) 2. **Configure Fields** - Select columns, set aliases, choose aggregations 3. **Add Filters** - Build WHERE conditions with operators or date presets 4. **Set Sorting** - Add ORDER BY columns with direction toggle 5. **Group Results** - Enable GROUP BY on selected fields 6. **Execute** - Run report and view in table or chart mode 7. **Export/Save** - Download CSV or save report definition

## 51.4 AI Report Writer (v5.42.0)

Enterprise-grade AI-powered report generation with text and voice input, interactive charts, smart insights, and brand customization.

**Location:** Think Tank Admin → Reports → AI Writer tab

**Core Features:** - **Natural Language Generation** - Describe reports in plain English - **Voice Input** - Web Speech API for hands-free report creation - **AI Modification** - Refine with follow-up prompts (“Add delight metrics”) - **Report Styles** - Executive Summary, Detailed Analysis, Dashboard View, Narrative - **Rich Formatting** - Headings, metrics cards, charts, tables, lists, quotes - **Edit Mode** - Click sections to modify, use format panel for styling - **Undo/Redo** - Full history navigation - **Export** - PDF, Excel, HTML, Print

**Interactive Charts (v5.42.0):** - Real Recharts visualizations replacing placeholders - Bar, Line, Pie, Area chart types - Auto-formatted tooltips (K/M for thousands/millions) - 8-color palette for visual consistency

**Smart Insights (v5.42.0):** - AI anomaly detection (usage spikes, unusual patterns) - Trend analysis with growth predictions - Achievement tracking (delight score peaks, retention milestones) - Actionable recommendations - Severity indicators (low/medium/high) - Confidence scores per insight

**Brand Kit (v5.42.0):** - Logo upload with drag-and-drop - Company name and tagline customization - Color pickers (Primary/Secondary/Accent) - Font selection for headers and body - Quick preset themes - Live preview card

**Think Tank Example Prompts:** - “Generate a user engagement report showing active users and session trends” - “Create a delight score analysis for the past month” - “Build a conversation analytics report with message volumes” - “Show me user retention metrics with churn analysis”

**Usage:** 1. Navigate to Reports → AI Writer tab 2. Select report style 3. Type or speak your request (click mic for voice) 4. Review generated report preview 5. Use modification prompt to refine 6. Toggle Edit Mode to modify sections 7. Export to PDF/Excel/HTML

**Report Sections Generated:** | Type | Description | ———|————| | **heading** | H1-H3 headings || **paragraph** | Body text || **metrics** | 4-column KPI cards with trends (↑↓) || **chart** | Interactive chart placeholders || **table** | Data tables with headers || **list** | Bullet point lists || **quote** | Blockquote sections |

## 51.5 Table Categories

| Category         | Description                      |
|------------------|----------------------------------|
| <b>Core</b>      | Users, tenants, sessions         |
| <b>AI</b>        | Models, prompts, responses       |
| <b>Billing</b>   | Subscriptions, credits, invoices |
| <b>Analytics</b> | Events, metrics, logs            |
| <b>System</b>    | Configuration, audit, health     |

## 51.5 Field Options

| Option             | Description                                        |
|--------------------|----------------------------------------------------|
| <b>Aggregation</b> | none, count, sum, avg, min, max, distinct          |
| <b>Format</b>      | text, number, currency, percentage, date, datetime |
| <b>Filter</b>      | Where clause conditions                            |
| <b>Group By</b>    | Grouping columns                                   |

| Option          | Description                |
|-----------------|----------------------------|
| <b>Order By</b> | Sort columns and direction |

## 51.6 Export Formats

| Format      | Description            |
|-------------|------------------------|
| <b>CSV</b>  | Comma-separated values |
| <b>JSON</b> | Structured data format |

## 51.7 API Endpoints

Base: /api/admin/dynamic-reports

| Method | Endpoint     | Description                   |
|--------|--------------|-------------------------------|
| GET    | /schema      | Discover database schema      |
| GET    | /suggestions | AI-generated report templates |
| GET    | /            | List saved reports            |
| POST   | /            | Save report definition        |
| POST   | /execute     | Execute a report              |
| POST   | /export      | Export report data            |
| DELETE | /:id         | Delete a report               |

## 51.8 Implementation Files

| File                                                                              | Purpose |
|-----------------------------------------------------------------------------------|---------|
| apps/thinktank-admin/app/(dashboard/reports/page.tsx)                             |         |
| packages/infrastructure/lambda/shared/backends/schema-adaptive-reports.service.ts |         |
| packages/infrastructure/lambda/admin/api/dynamic-reports.ts                       |         |
| packages/infrastructure/migrations/D20260121a003__dynamic_reports.sql             |         |

## Section 52: Gateway Status (v5.39.0)

**Location:** Think Tank Admin → Gateway

Monitor API Gateway health and traffic metrics.

### 52.1 Overview

The Gateway Status dashboard displays:

| Metric                    | Description                                    |
|---------------------------|------------------------------------------------|
| <b>Status</b>             | Overall gateway health (healthy/degraded/down) |
| <b>Requests/sec</b>       | Current request throughput                     |
| <b>Avg Latency</b>        | Mean response time                             |
| <b>Error Rate</b>         | Percentage of failed requests                  |
| <b>Active Connections</b> | Current WebSocket connections                  |

## 52.2 Endpoint Health

| Column          | Description                   |
|-----------------|-------------------------------|
| <b>Endpoint</b> | API route path                |
| <b>Method</b>   | HTTP method (GET, POST, etc.) |
| <b>Status</b>   | healthy, slow, error          |
| <b>Latency</b>  | P50/P95/P99 response times    |
| <b>Requests</b> | Request count (24h)           |
| <b>Errors</b>   | Error count (24h)             |

## 52.3 Traffic Patterns

| View               | Description                   |
|--------------------|-------------------------------|
| <b>Hourly</b>      | Requests per hour (24h)       |
| <b>Daily</b>       | Requests per day (30d)        |
| <b>By Endpoint</b> | Traffic distribution by route |
| <b>By Status</b>   | Success vs error breakdown    |

## 52.4 Alerts

| Alert Type              | Trigger                  |
|-------------------------|--------------------------|
| <b>High Latency</b>     | P95 > 2000ms             |
| <b>High Error Rate</b>  | Errors > 5%              |
| <b>Throughput Spike</b> | 2x normal traffic        |
| <b>Connection Drop</b>  | WebSocket disconnections |

## 52.5 Implementation Files

| File                                                         | Purpose |
|--------------------------------------------------------------|---------|
| apps/thinktank-admin/app/(dashboard gateway)/pages/index.tsx |         |

## Section 53: Decision Intelligence Artifacts (DIA Engine) (v5.43.0)

**Location:** Think Tank Admin → Decision Records

The Glass Box Decision Engine - transforms AI conversations into auditable, evidence-backed decision records with full provenance tracking.

### 53.1 Overview

Decision Intelligence Artifacts (DIA) solve the critical problem of AI decision opacity:

| Challenge                  | DIA Solution                                       |
|----------------------------|----------------------------------------------------|
| <b>Black Box Decisions</b> | Full claim-to-evidence mapping                     |
| <b>Data Staleness</b>      | Volatile query tracking with automatic validation  |
| <b>Dissent Hidden</b>      | Ghost paths visualize rejected alternatives        |
| <b>Compliance Gaps</b>     | Built-in HIPAA/SOC2/GDPR export packages           |
| <b>Trust Uncertainty</b>   | Breathing heatmap shows trust topology at-a-glance |

### 53.2 Core Concepts

**Claims:** Extracted conclusions, findings, recommendations, warnings, and facts from AI responses.

| Claim Type                            | Description                           |
|---------------------------------------|---------------------------------------|
| <code>conclusion</code>               | Final determination or decision       |
| <code>finding</code>                  | Discovered information or observation |
| <code>recommendation</code>           | Suggested course of action            |
| <code>warning</code>                  | Risk or caution indicator             |
| <code>fact</code>                     | Verified data point                   |
| <code>clinical_finding</code>         | Healthcare-specific observation       |
| <code>treatment_recommendation</code> | Medical treatment suggestion          |
| <code>risk_assessment</code>          | Risk evaluation                       |
| <code>legal_opinion</code>            | Legal interpretation                  |
| <code>compliance_finding</code>       | Regulatory compliance observation     |

**Evidence Links:** Connections between claims and their supporting data sources.

| Evidence Type                | Description                  |
|------------------------------|------------------------------|
| <code>tool_call</code>       | API or tool execution result |
| <code>web_search</code>      | Web search results           |
| <code>document</code>        | Referenced document          |
| <code>calculation</code>     | Computed result              |
| <code>model_consensus</code> | Multiple model agreement     |

**Dissent Events:** Captured disagreements from model reasoning traces.

| Severity    | Description                            |
|-------------|----------------------------------------|
| minor       | Small qualification or caveat          |
| moderate    | Significant alternative consideration  |
| significant | Major disagreement requiring attention |

**Volatile Queries:** Tool calls that may return different results over time.

| Volatility | Threshold | Examples              |
|------------|-----------|-----------------------|
| real-time  | 1 hour    | Stock prices, weather |
| daily      | 24 hours  | News, analytics       |
| weekly     | 168 hours | Document searches     |
| stable     | No expiry | Static references     |

### 53.3 The Living Parchment UI

The artifact viewer uses sensory design principles:

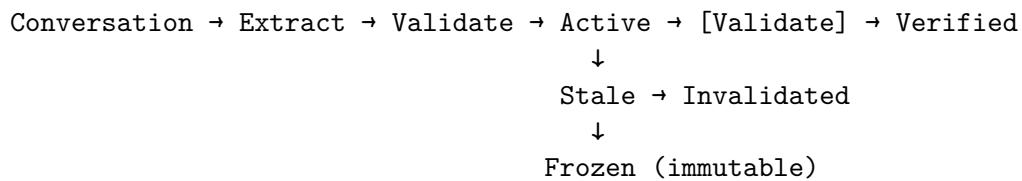
**Breathing Heatmap Scrollbar:** - Green (verified) - 6 BPM breathing rate - Amber (unverified)  
 - Standard breathing - Red (contested) - 12 BPM alert breathing - Purple (stale) - Fading intensity with age

**Living Ink Typography:** - Font weight: 350-500 based on confidence (0-100%) - Stale claims fade to grayscale - Hover reveals evidence connections

**Control Island** (floating lens selector): - **Read:** Standard document view - **X-Ray:** Evidence links visible - **Risk:** Ghost paths and contested claims highlighted - **Compliance:** Regulatory framework coverage

**Ghost Paths:** Dashed connectors showing rejected alternatives from dissent events.

### 53.4 Artifact Lifecycle



| Status      | Description                         |
|-------------|-------------------------------------|
| active      | Current, editable artifact          |
| frozen      | Immutable version with content hash |
| archived    | Soft-deleted                        |
| invalidated | Data significantly changed          |

| Validation Status        | Description                          |
|--------------------------|--------------------------------------|
| <code>fresh</code>       | Newly created, not yet validated     |
| <code>stale</code>       | Volatile queries exceeded thresholds |
| <code>verified</code>    | Recently validated, data unchanged   |
| <code>invalidated</code> | Significant data changes detected    |

### 53.5 Compliance Exports

| Format                     | Use Case                                    |
|----------------------------|---------------------------------------------|
| <code>pdf</code>           | Human-readable document                     |
| <code>json</code>          | Machine-readable data                       |
| <code>hipaa_audit</code>   | HIPAA compliance package with PHI inventory |
| <code>soc2_evidence</code> | SOC2 control mapping and evidence chain     |
| <code>gdpr_dsar</code>     | GDPR Data Subject Access Request response   |

**HIPAA Audit Package Contents:** - Cover sheet with artifact metadata - PHI inventory with categories - Access log for minimum necessary compliance - Evidence chain verification - System attestation with content hash

**SOC2 Evidence Bundle Contents:** - Control mapping (CC6.x, CC7.x, CC8.x) - Evidence chain completeness verification - Change management documentation - Integrity verification with signature

### 53.6 Configuration

**Location:** Think Tank Admin → Decision Records → Config

| Setting                                    | Default                        | Description                                  |
|--------------------------------------------|--------------------------------|----------------------------------------------|
| <code>diaEnabled</code>                    | <code>true</code>              | Enable DIA Engine                            |
| <code>autoGenerateEnabled</code>           | <code>false</code>             | Auto-generate artifacts from conversations   |
| <code>phiDetectionEnabled</code>           | <code>true</code>              | Scan for protected health information        |
| <code>piDetectionEnabled</code>            | <code>true</code>              | Scan for personally identifiable information |
| <code>defaultStalenessThresholdDays</code> |                                | Days before volatile queries flagged stale   |
| <code>maxArtifactsPerUser</code>           | 0                              | Limit per user (0 = unlimited)               |
| <code>extractionModel</code>               | <code>claude-3-5-sonnet</code> | Model for claim extraction                   |
| <code>autoRedactPhiOnExport</code>         | <code>false</code>             | Automatically redact PHI on export           |

### 53.7 Templates

Pre-configured extraction templates:

| Template                | Description                 | Compliance |
|-------------------------|-----------------------------|------------|
| General Decision Record | Standard extraction         | None       |
| Healthcare Decision     | HIPAA-compliant clinical    | HIPAA      |
| Financial Analysis      | Audit-ready financial       | SOC2       |
| Legal Review            | Legal opinion documentation | SOC2, GDPR |
| Research Synthesis      | Multi-source research       | None       |

### 53.8 API Endpoints

Base: `/api/thinktank/decision-artifacts`

| Method | Endpoint                | Description                         |
|--------|-------------------------|-------------------------------------|
| GET    | /                       | List artifacts (supports filters)   |
| POST   | /                       | Generate artifact from conversation |
| GET    | /dashboard              | Dashboard metrics                   |
| GET    | /templates              | List available templates            |
| GET    | /config                 | Get tenant configuration            |
| PUT    | /config                 | Update configuration                |
| GET    | /:id                    | Get artifact details                |
| DELETE | /:id                    | Archive artifact                    |
| GET    | /:id/staleness          | Check staleness status              |
| POST   | /:id/validate           | Validate volatile queries           |
| POST   | /:id/export             | Export artifact                     |
| GET    | /:id/versions           | Get version history                 |
| GET    | /:id/validation-history | Validation audit trail              |
| GET    | /:id/export-history     | Export audit trail                  |

### 53.9 Dashboard Metrics

| Metric              | Description                   |
|---------------------|-------------------------------|
| Total Artifacts     | All artifacts for tenant      |
| Active Artifacts    | Non-archived, non-frozen      |
| Frozen Artifacts    | Immutable versions            |
| Average Confidence  | Mean confidence across active |
| Stale Artifacts     | Needing validation            |
| PHI/PII Detected    | Artifacts with sensitive data |
| Validation Cost MTD | API costs for re-validation   |
| Top Domains         | Most common primary domains   |
| Compliance Usage    | Framework distribution        |

### 53.10 Database Schema

Tables:

| Table                            | Purpose                |
|----------------------------------|------------------------|
| decision_artifacts               | Main artifact storage  |
| decision_artifact_validation_log | Validation audit trail |
| decision_artifact_export_log     | Export audit trail     |
| decision_artifact_config         | Tenant configuration   |
| decision_artifact_templates      | Extraction templates   |
| decision_artifact_access_log     | Access audit (HIPAA)   |

#### Key Columns (decision\_artifacts):

| Column            | Type        | Description                        |
|-------------------|-------------|------------------------------------|
| artifact_content  | JSONB       | Claims, evidence, dissent, metrics |
| heatmap_data      | JSONB       | Pre-computed heatmap segments      |
| validation_status | VARCHAR     | fresh/stale/verified/invalidated   |
| phi_detected      | BOOLEAN     | Contains protected health info     |
| content_hash      | VARCHAR(64) | SHA-256 for frozen artifacts       |

#### 53.11 Implementation Files

| File                                                                                 | Purpose                     |
|--------------------------------------------------------------------------------------|-----------------------------|
| packages/shared/src/types/decisionArtifactTypes.ts                                   | API definition types        |
| packages/infrastructure/lambda/shaBackend/services/dia/                              | Backend services API        |
| packages/infrastructure/lambda/thinktank/admin/decision-artifacts.ts                 | API endpoint                |
| packages/infrastructure/lib/stacks/DIA-infrastructure                                | DIA infrastructure          |
| packages/infrastructure/migrations/V2026_01_22_001__decision_artifacts.sql           | Initial migration           |
| packages/infrastructure/migrations/V2026_01_22_002s_decision_artifact_versioning.sql | Versioning migration        |
| packages/infrastructure/migrations/V2026_01_22_003__decision_artifact_config.sql     | Config migration            |
| apps/thinktank-admin/app/(dashboard)/decision-records/                               | Decision records component  |
| apps/thinktank-admin/app/(dashboard)/dependentrecords/components/                    | Dependent records component |

#### 53.12 Troubleshooting

| Issue                     | Solution                                           |
|---------------------------|----------------------------------------------------|
| Extraction fails          | Check Bedrock model access permissions             |
| Missing evidence links    | Verify tool_calls in message metadata              |
| Stale status not updating | Run manual validation or check thresholds          |
| PHI not detected          | Verify phiDetectionEnabled in config               |
| Export fails              | Check S3 bucket permissions<br>(DIA_EXPORT_BUCKET) |
| Version history empty     | Artifact must be frozen to create versions         |

#### 53.13 Security Considerations

- All tables have RLS policies enforcing tenant isolation

- PHI/PII detection runs automatically on extraction
  - Access logging enabled for HIPAA compliance
  - Content hashes provide tamper evidence for frozen artifacts
  - Export audit trail tracks all compliance exports
  - Presigned URLs expire after 1 hour
- 

## Section 54: Living Parchment 2029 Vision (v5.44.0)

### Overview

Living Parchment is a comprehensive suite of advanced decision intelligence tools featuring sensory UI elements that communicate trust, confidence, and data freshness through visual breathing, living typography, and ghost paths. This 2029 Vision implementation transforms how users interact with AI-assisted decision making.

### Design Philosophy

| Concept                     | Implementation                                                                                                     |
|-----------------------------|--------------------------------------------------------------------------------------------------------------------|
| <b>Breathing Interfaces</b> | UI elements pulse with life—faster breathing (12 BPM) indicates uncertainty, slower (4-6 BPM) indicates confidence |
| <b>Living Ink</b>           | Text weight varies 350-500 based on confidence; stale information fades to grayscale                               |
| <b>Ghost Paths</b>          | Rejected alternatives remain visible as translucent traces showing what could have been                            |
| <b>Confidence Terrain</b>   | 3D topographic visualization where elevation = confidence, color = risk                                            |

### 54.1 War Room (Strategic Decision Theater)

High-stakes collaborative decision space with AI advisors and confidence terrain.

### Features

- **Confidence Terrain:** 3D grid visualization showing confidence topology across decision space
- **AI Advisory Council:** Multiple AI models providing different perspectives
- **Decision Paths:** Branching options with outcome predictions and advocate tracking
- **Ghost Branches:** Rejected paths remain visible for context
- **Stake Level Indicators:** Visual urgency based on decision importance

### Advisor Types

| Type              | Color            | Use Case                           |
|-------------------|------------------|------------------------------------|
| AI Model          | Blue (#3b82f6)   | Claude, GPT for strategic analysis |
| Human Expert      | Purple (#8b5cf6) | Domain specialists                 |
| Domain Specialist | Cyan (#06b6d4)   | Industry-specific advisors         |

| Type | Color | Use Case |
|------|-------|----------|
|------|-------|----------|

## API Endpoints

|      |                                                                    |                  |
|------|--------------------------------------------------------------------|------------------|
| POST | /api/thinktank/living-parchment/war-room                           | Create session   |
| GET  | /api/thinktank/living-parchment/war-room                           | List sessions    |
| GET  | /api/thinktank/living-parchment/war-room/:id                       | Get session      |
| POST | /api/thinktank/living-parchment/war-room/:id/advisors              | Add advisor      |
| POST | /api/thinktank/living-parchment/war-room/:id/advisors/:aid/analyze | Request analysis |
| POST | /api/thinktank/living-parchment/war-room/:id/paths                 | Propose path     |
| POST | /api/thinktank/living-parchment/war-room/:id/decide                | Make decision    |
| POST | /api/thinktank/living-parchment/war-room/:id/terrain               | Update terrain   |

## 54.2 Council of Experts

Multi-persona AI consultation with consensus tracking and dissent visualization.

### Expert Personas

| Persona     | Specialization           | Style                                 |
|-------------|--------------------------|---------------------------------------|
| Pragmatist  | Practical Implementation | Results-focused,<br>cost-conscious    |
| Ethicist    | Moral Philosophy         | Principle-based,<br>stakeholder-aware |
| Innovator   | Creative Solutions       | Visionary,<br>possibility-focused     |
| Skeptic     | Risk Analysis            | Devil's advocate,<br>challenging      |
| Synthesizer | Integration              | Bridge-building,<br>pattern-finding   |
| Analyst     | Data-Driven              | Quantitative,<br>evidence-based       |
| Strategist  | Long-term Strategy       | Big-picture,<br>competitive-aware     |
| Humanist    | Human Impact             | Empathetic,<br>user-centered          |

### Consensus Visualization

- Experts positioned on circular visualization
- Positions move toward center as consensus increases
- Dissent sparks appear as electrical arcs between disagreeing experts
- Gravitational attraction animation shows convergence

## API Endpoints

|      |                                                      |                  |
|------|------------------------------------------------------|------------------|
| POST | /api/thinktank/living-parchment/council              | Convene council  |
| GET  | /api/thinktank/living-parchment/council/:id          | Get session      |
| POST | /api/thinktank/living-parchment/council/:id/debate   | Run debate round |
| POST | /api/thinktank/living-parchment/council/:id/conclude | Conclude session |

### 54.3 Debate Arena

Adversarial exploration with attack/defense flows and steel-man generation.

#### Features

- **Resolution Meter:** Balance indicator (-100 to +100) showing which side is winning
- **Argument Flow:** Visual stream of claims, rebuttals, and concessions
- **Weak Point Detection:** Breathing red indicators on vulnerable arguments
- **Steel-Man Generation:** AI creates strongest version of opponent's argument
- **Attack/Defense Arrows:** Animated flows showing which arguments target which

#### Debate Phases

1. **Setup** - Configure debaters and proposition
2. **Opening** - Initial statements
3. **Main** - Core argument exchange
4. **Rebuttal** - Direct challenges
5. **Closing** - Final positions
6. **Resolved** - Outcome determined

#### API Endpoints

|      |                                                      |                    |
|------|------------------------------------------------------|--------------------|
| POST | /api/thinktank/living-parchment/debate               | Create debate      |
| GET  | /api/thinktank/living-parchment/debate/:id           | Get arena          |
| POST | /api/thinktank/living-parchment/debate/:id/round     | Run round          |
| POST | /api/thinktank/living-parchment/debate/:id/steel-man | Generate steel-man |

### 54.4 Memory Palace (Coming Soon)

Navigable 3D knowledge topology with freshness fog.

- **Knowledge Rooms:** Domain-organized 3D spaces
- **Freshness Fog:** Stale areas appear foggy
- **Connection Threads:** Luminous lines between related concepts
- **Discovery Hotspots:** Breathing beacons where insights could emerge

### 54.5 Oracle View (Coming Soon)

Predictive confidence landscape.

- **Probability Heatmap:** Future timeline with brightness = confidence
- **Bifurcation Points:** Animated forks showing cascade effects
- **Ghost Futures:** Translucent overlays of alternative scenarios
- **Black Swan Indicators:** Dormant embers for low-probability/high-impact events

## 54.6 Synthesis Engine (Coming Soon)

Multi-source fusion view.

- **Source Streams:** Flowing rivers converging into synthesis
- **Agreement Zones:** Warm glow where sources align
- **Tension Zones:** Crackling energy between contradictions
- **Provenance Trails:** Click any claim to see all supporting sources

## 54.7 Cognitive Load Monitor (Coming Soon)

User state awareness with adaptive UI.

- **Attention Heatmap:** Track where user has focused
- **Fatigue Indicators:** UI breathing slows as session lengthens
- **Overwhelm Warning:** Screen edges breathe red when load peaks

## 54.8 Temporal Drift Observatory (Coming Soon)

Fact evolution tracking.

- **Drift Alerts:** Notifications when facts have changed
- **Version Ghosts:** Previous versions as translucent overlays
- **Citation Half-Life:** Predict when facts likely become stale

## Database Schema

```
-- Core tables (see migration V2026_01_22_004)
war_room_sessions, war_room_participants, war_room_advisors
memory_palaces, memory_rooms, knowledge_nodes, memory_connections
oracle_views, oracle_predictions, bifurcation_points, ghost_futures
synthesis_sessions, synthesis_sources, synthesis_claims
cognitive_load_sessions, cognitive_load_history
council_sessions, council_experts, expert_arguments, minority_reports
drifting_facts, drift_alerts, version_ghosts
debate Arenas, debaters, debate_arguments, weak_points, steel_man_overlays
living_parchment_config
```

## Configuration

```
interface LivingParchmentConfig {
    features: {
        warRoomEnabled: boolean;           // Default: true
        memoryPalaceEnabled: boolean;      // Default: true
        oracleViewEnabled: boolean;        // Default: true
        synthesisEngineEnabled: boolean;   // Default: true
        cognitiveLoadEnabled: boolean;    // Default: true
        councilOfExpertsEnabled: boolean;  // Default: true
        temporalDriftEnabled: boolean;    // Default: true
        debateArenaEnabled: boolean;      // Default: true
    };
}
```

```

defaults: {
  breathingRateBase: 6;           // BPM
  confidenceThreshold: 70;        // Minimum for "high confidence"
  stalenessThresholdDays: 30;      // When facts become stale
  maxAdvisors: 10;
  maxExperts: 8;
  maxDebateRounds: 5;
};

visualSettings: {
  heatmapColorScheme: 'standard' | 'accessible' | 'dark';
  animationIntensity: 'subtle' | 'normal' | 'vivid';
  ghostOpacity: 0.5;
};

}

```

## Implementation Files

```

packages/shared/src/types/living-parchment.types.ts      # All types
packages/infrastructure/migrations/V2026_01_22_004__living_parchment_core.sql
packages/infrastructure/lambda/shared/services/living-parchment/
  war-room.service.ts
  council-of-experts.service.ts
  debate-arena.service.ts
  index.ts
packages/infrastructure/lambda/thinktank/living-parchment.ts
apps/thinktank-admin/app/(dashboard)/living-parchment/
  page.tsx                                # Landing page
  war-room/page.tsx                         # War Room UI
  council/page.tsx                          # Council of Experts UI
  debate/page.tsx                           # Debate Arena UI

```

## Security Considerations

- All tables have RLS policies enforcing tenant isolation
  - AI advisor calls use tenant-scoped model access
  - Session ownership validated before modifications
  - Audit logging for all decision actions
  - Debate content filtered for appropriate use
- 

## 45. Localization & Translation Overrides

**Location:** Think Tank Admin → Administration → Localization

Tenant administrators can customize UI text and messages across Think Tank with translation overrides.

## 45.1 Overview

The localization system allows tenants to:

- Override any system string with custom text
- Protect overrides from automatic translation updates
- Configure default and enabled languages for users
- Maintain brand consistency across all 18 supported languages

## 45.2 Supported Languages

| Language              | Code  | Flag |
|-----------------------|-------|------|
| English               | en    |      |
| Spanish               | es    |      |
| French                | fr    |      |
| German                | de    |      |
| Portuguese            | pt    |      |
| Italian               | it    |      |
| Dutch                 | nl    |      |
| Polish                | pl    |      |
| Russian               | ru    |      |
| Turkish               | tr    |      |
| Japanese              | ja    |      |
| Korean                | ko    |      |
| Chinese (Simplified)  | zh-CN |      |
| Chinese (Traditional) | zh-TW |      |
| Arabic                | ar    |      |
| Hindi                 | hi    |      |
| Thai                  | th    |      |
| Vietnamese            | vi    |      |

## 45.3 Admin UI Tabs

| Tab                   | Purpose                                    |
|-----------------------|--------------------------------------------|
| <b>Your Overrides</b> | View and manage custom translations        |
| <b>Browse Strings</b> | Search Think Tank strings to customize     |
| <b>Configuration</b>  | Set default language and enabled languages |

## 45.4 Creating Translation Overrides

1. Navigate to **Administration → Localization**
2. Select target language from dropdown
3. Go to **Browse Strings** tab
4. Search for the string you want to customize
5. Click **Edit** to open the override dialog
6. Enter your custom text
7. Toggle **Protect from automatic updates** (recommended)
8. Click **Save**

## 45.5 Protection System

**Protected overrides** (default): - Will NOT be updated when system translations improve - Recommended for brand-specific terminology - Shows lock icon in override list

**Unprotected overrides:** - May be updated by automatic translation systems - Useful for temporary fixes until system improves - Shows unlock icon in override list

**Reverting to system translation:** - Click the revert button on any override - Override is deleted and system translation is restored

## 45.6 Language Configuration

Configure which languages are available to your users:

1. Go to **Configuration** tab
2. Set **Default Language** for new users
3. Enable/disable languages by clicking language cards
4. At least one language must remain enabled

## 45.7 Common Use Cases

| Use Case                 | Example                                     |
|--------------------------|---------------------------------------------|
| <b>Brand terminology</b> | Replace “Think Tank” with your product name |
| <b>Industry jargon</b>   | Use domain-specific terms                   |
| <b>Tone adjustment</b>   | Make messages more formal/casual            |
| <b>Legal compliance</b>  | Customize disclaimers                       |
| <b>Regional variants</b> | UK vs US English differences                |

## 45.8 API Reference

Base URL: /api/admin/localization

| Endpoint                  | Method  | Purpose                         |
|---------------------------|---------|---------------------------------|
| /overrides                | GET     | List your overrides             |
| /overrides                | POST    | Create/update override          |
| /overrides/:id            | DELETE  | Revert to system                |
| /overrides/:id/protection | PATCH   | Toggle protection               |
| /config                   | GET/PUT | Language configuration          |
| /bundle/:lang             | GET     | Get translations with overrides |

## 45.9 Database Tables

| Table                        | Purpose                        |
|------------------------------|--------------------------------|
| tenant_translation_overrides | Custom translations per tenant |
| tenant_localization_config   | Language settings per tenant   |
| translation_audit_log        | Change history                 |

## 45.10 Implementation Files

```
packages/infrastructure/migrations/V2026_01_25_006__tenant_translation_overrides.sql  
packages/infrastructure/lambda/admin/localization-registry.ts  
apps/thinktank-admin/app/(dashboard)/localization/page.tsx
```

---

## Related Documentation

- [RADIANT Admin Guide](#) - Platform administration
- [RADIANT Admin Guide - HITL Orchestration](#) - Full HITL Orchestration documentation
- [RADIANT Admin Guide - Metrics & Learning](#) - Persistent learning system
- [RADIANT Admin Guide - Consciousness Evolution](#) - Predictive coding, LoRA evolution, Local Ego
- [Think Tank User Guide](#) - End user guide
- [User Rules System](#) - Memory rules details
- [Provider Rejection Handling](#) - Rejection system
- [AI Ethics Standards](#) - Ethics framework