

Contents

Think Tank User Guide	2
Table of Contents	2
1. Welcome to Think Tank	3
What Makes Think Tank Different	3
How Think Tank Works	3
2. Getting Started	4
First Time Setup	4
Authentication & Security	4
Language Settings	5
The Main Interface	5
3. The Dashboard	5
Key Metrics	5
Quick Actions	6
Your Profile	6
4. Conversations	7
Starting a New Conversation	7
Understanding Responses	7
Conversation Actions	7
Conversation Search	7
5. My Rules - Personalizing AI Responses	7
Creating a Custom Rule	7
Rule Types	8
Using Preset Rules	8
Managing Rules	9
Best Practices	9
6. Domain Modes	9
Available Domains	9
How Domain Detection Works	9
Domain Indicators	10
Overriding Domain Detection	10
7. Delight System - AI Personality	10
Personality Modes	10
Personality Elements	10
Adjusting Personality	11
8. Collaboration Features	11
Real-Time Collaboration	11
Sharing Conversations	12
9. Advanced Features	12
Polymorphic UI - Adaptive Views	12
The Grimoire - AI Learning	13
Magic Carpet Navigator	13
Artifacts - Generated Code	14
Liquid Interface - Chat Morphs Into Tools (v5.52.8)	14
10. How Think Tank's Memory Works	15
Cato - The AI's Personality & Memory	15
Cortex - The Enterprise Knowledge Library	15

How They Work Together	16
What This Means for You	16
Privacy Note	16
11. Understanding AI Decisions	17
Brain Plans	17
Confidence Levels	17
When Think Tank Asks for Clarification	18
Epistemic Humility	18
14. Safety & Governance	18
Five-Layer Safety Stack	18
Governance Presets	19
Human-in-the-Loop (HITL)	19
What Think Tank Will Never Do	19
Reporting Issues	19
12. Decision Records	19
Accessing Decision Records	19
The Living Parchment View	20
Verifying Claims	20
Exporting for Compliance	20
Exporting Conversations Directly (v5.52.16)	20
13. Living Parchment	21
War Room (Strategic Decision Theater)	21
Council of Experts	21
Debate Arena	21
Understanding Living Parchment UI	22
14. Safety & Governance	22
15. Keyboard Shortcuts	22
16. Troubleshooting	22
Common Issues	22
Getting Help	23
17. Glossary	23
Document History	24

Think Tank User Guide

Version: 5.52.29

Last Updated: January 25, 2026

Audience: End Users of Think Tank

Table of Contents

1. Welcome to Think Tank
2. Getting Started
3. The Dashboard
4. Conversations
5. My Rules - Personalizing AI Responses

6. Domain Modes
 7. Delight System - AI Personality
 8. Collaboration Features
 9. Advanced Features
 10. How Think Tank's Memory Works
 11. Understanding AI Decisions
 12. Decision Records
 13. Living Parchment
 14. Safety & Governance
 15. Keyboard Shortcuts
 16. Troubleshooting
 17. Glossary
-

1. Welcome to Think Tank

Think Tank is an advanced AI assistant platform that adapts to your needs, learns your preferences, and provides intelligent responses across a wide range of domains. Unlike simple chatbots, Think Tank:

- **Adapts to Your Expertise** - Automatically detects your knowledge domain and adjusts responses
- **Remembers Your Preferences** - Your rules and settings persist across conversations
- **Shows Its Thinking** - Transparent about its reasoning and confidence levels
- **Keeps You Safe** - Built-in safety guardrails protect against harmful outputs
- **Collaborates** - Work together with colleagues in real-time sessions

What Makes Think Tank Different

Traditional Chatbots	Think Tank
One-size-fits-all responses	Adapts to your domain expertise
Forgets your preferences	Persistent user rules and context
Black-box decisions	Transparent reasoning with Brain Plans
Single interaction mode	Multiple view modes (Sniper, Scout, Sage)
No safety guarantees	Five-layer safety architecture

How Think Tank Works

YOUR QUESTION

BRAIN PLANNER

Domain

Your Rules

Context

Detection Applied Analysis

MODEL SELECTION & ROUTING

Sniper	Scout	Sage	War Room
Fast & Cheap	Research & Explore	Analysis & Compare	Multi-Agent Debate

SAFETY VALIDATION (Cato)

Content safety Governance check Cost approval

YOUR ANSWER

+ Confidence indicator + Sources + Suggestions

2. Getting Started

First Time Setup

1. **Log In** - Access Think Tank through your organization's portal
2. **Set Your Preferences** - Visit Settings to configure your experience
3. **Create Rules** - Add personal rules to customize AI responses
4. **Start a Conversation** - Type your first message or use voice input

Authentication & Security

Think Tank supports multiple authentication methods for secure access:

- **Email/Password** - Traditional sign-in with optional MFA
- **Social Sign-In** - Google, Microsoft, Apple, GitHub
- **Enterprise SSO** - SAML 2.0 / OIDC via your organization
- **Passkeys** - Passwordless authentication using biometrics

Multi-Factor Authentication (MFA) may be required by your organization. When enabled, you'll need an authenticator app (Google Authenticator, Authy, etc.) to generate verification codes.

Detailed Guides: See [Authentication User Guide](#) and [MFA Guide](#)

Language Settings

Think Tank supports **18 languages** including:

Western	Asian	RTL
English, Spanish, French, German, Portuguese, Italian, Dutch, Polish, Russian, Turkish	Japanese, Korean, Chinese (Simplified/Traditional), Hindi, Thai, Vietnamese	Arabic

To change your language: 1. Click your profile icon → **Settings** 2. Select **Language & Region** 3. Choose your preferred language 4. The interface updates immediately

Search works in all languages, with special **CJK (Chinese/Japanese/Korean) bi-gram search** for accurate results.

Detailed Guide: See [Internationalization Guide](#)

The Main Interface

[Logo] Think Tank [User] [Settings]

SIDE BAR MAIN CONTENT AREA

Dashboard Your conversations and AI responses
Users appear here
Messages
My Rules
Delight
...
...

[Message input...] [Send]

3. The Dashboard

The Dashboard provides an overview of your Think Tank activity.

Key Metrics

Metric	Description
Active Users	How many people in your organization are using Think Tank
Conversations	Total number of conversations you've had

Metric	Description
User Rules	How many personal rules you've created
API Requests	Volume of AI requests (for awareness)

Quick Actions

From the dashboard, you can quickly:

- **Manage Users** - View and manage team members (if you have permissions)
- **Configure Delight** - Customize the AI's personality
- **Domain Modes** - Adjust how Think Tank handles different topics

Your Profile

Access your profile by clicking your avatar in the top-right corner. Your profile includes:

Activity Heatmap A GitHub-style visualization of your conversation activity over the past year:

Feature	Description
Breathing Animation	Cells pulse based on activity intensity - more active days “breathe” faster
AI Insights	Automatic pattern detection with natural language explanations
Streak Tracking	Current and longest streaks highlighted with badges
Sound Feedback	Optional audio cues when hovering over active days
Accessibility Mode	Full narrative summary for screen readers

AI Insights Examples: - “You’re a weekday warrior! Most activity happens Monday-Friday” (92% confidence) - “Amazing! Your longest streak is 14 days. That’s dedication! ” - “Activity has slowed recently. A quick session could reignite momentum!”

Color Legend: - Empty (dark) → No activity - Light purple → Low activity - Bright purple → High activity - Dashed border → Predicted future activity

Profile Stats

Stat	Description
Conversations	Total conversations you’ve had
Tokens Used	AI tokens consumed (for awareness)
Messages	Total messages exchanged
Achievements	Unlocked gamification badges

4. Conversations

Starting a New Conversation

1. Type your message in the input field at the bottom
2. Press **Enter** or click **Send**
3. Wait for the AI response (a typing indicator shows Think Tank is working)

Understanding Responses

Think Tank responses may include:

- **Main Answer** - The AI's response to your question
- **Confidence Indicator** - How certain the AI is about its answer
- **Sources** - References or citations when available
- **Suggestions** - Related questions you might want to ask

Conversation Actions

Action	How To
Share conversation	Click the share icon to create a shareable link
Export	Download conversation as text or markdown
Delete	Remove a conversation from your history
Branch	Create an alternative thread from any point

Conversation Search

Use the search bar to find past conversations by: - Keywords in messages - Date range - Conversation status (active, archived)

5. My Rules - Personalizing AI Responses

My Rules lets you set **persistent preferences** for how Think Tank responds to you. These rules are applied to every conversation.

Creating a Custom Rule

ADD CUSTOM RULE

Rule Summary *

Prefer concise bullet points

Rule Text *

When responding, use bullet points instead of long paragraphs. Keep responses under 200 words unless I specifically ask for more detail.

Rule Type

Format

[Cancel] [Create Rule]

1. Navigate to **My Rules** from the sidebar
2. Click **Add Custom Rule**
3. Fill in:
 - **Rule Summary** - Brief name (e.g., “Prefer concise answers”)
 - **Rule Text** - Detailed instruction for the AI
 - **Rule Type** - Category of the rule
4. Click **Create Rule**

Rule Types

Type	Use For	Example
Preference	General response style	“I prefer bullet points over paragraphs”
Restriction	Things to avoid	“Never use jargon without explaining it”
Format	Response structure	“Always include a summary at the end”
Sources	Citation preferences	“Cite academic sources when available”
Tone	Communication style	“Use a professional but friendly tone”
Topic	Subject-specific rules	“For medical topics, always recommend consulting a doctor”
Privacy	Data handling	“Don’t reference my previous conversations”

Using Preset Rules

Think Tank provides pre-made rules you can add with one click:

1. Go to **My Rules** → **Add from Presets**
2. Browse categories (e.g., “Response Style”, “Privacy”, “Formatting”)
3. Click **Add** next to any rule you want
4. Rules marked **Popular** are used by many users

Managing Rules

- **Toggle On/Off** - Temporarily disable a rule without deleting it
- **Times Applied** - See how often each rule has been used
- **Delete** - Remove a rule permanently

Best Practices

- Start with 3-5 core rules
 - Be specific in your rule text
 - Review rules periodically - remove ones that don't help
 - Use preset rules as starting points, then customize
-

6. Domain Modes

Think Tank automatically detects what domain your question relates to and adjusts its behavior accordingly.

Available Domains

Domain	Icon	Optimized For
General		Everyday questions and tasks
Medical		Healthcare and medical topics (with appropriate disclaimers)
Legal		Legal research and analysis
Code		Programming and development
Academic		Research and educational content
Creative		Writing, content creation, brainstorming
Scientific		Scientific research and analysis

How Domain Detection Works

"What are the side effects of ibuprofen?"

DOMAIN DETECTION

Keywords: "side effects", "ibuprofen" → MEDICAL
Intent: Information seeking
Confidence: 94%

AUTOMATIC ADJUSTMENTS

Model: Claude Sonnet (high accuracy for medical)
Temperature: 0.3 (factual, precise)
Style: Include disclaimers, cite sources
Guardrails: Add "consult healthcare provider" notice

1. You submit a question
2. Think Tank analyzes keywords, context, and intent
3. The appropriate domain mode is automatically selected
4. The AI adjusts its:
 - **Model selection** - Best AI model for that domain
 - **Temperature** - Creativity vs. precision balance
 - **Response style** - Technical depth, tone, formatting

Domain Indicators

You'll see a small badge indicating the detected domain:

[Medical] Analyzing your health question...

Overriding Domain Detection

If Think Tank picks the wrong domain: 1. Click the domain badge 2. Select the correct domain 3. Your choice is remembered for similar questions

7. Delight System - AI Personality

The Delight System controls Think Tank's personality, making interactions more engaging and human-like.

Personality Modes

Mode	Description	Best For
Auto (Recommended)	Adapts based on context	Most users
Professional	Formal and direct	Business use
Friendly	Warm and approachable	Casual conversations
Playful	Fun and expressive	Creative work
Minimal	Just the facts	Quick lookups

Personality Elements

Think Tank's personality includes:

Messages Contextual messages that appear during interactions: - **Pre-execution** - “Let me think about that...” - **During execution** - “Analyzing your data...” - **Post-execution** - “Here’s what I found!”

Achievements Unlock achievements as you use Think Tank: - **First Conversation** - Start your journey - **Power User** - 100 conversations - **Rule Master** - Create 10 custom rules - **Domain Expert** - Use 5 different domains

Achievement rarities: Common, Uncommon, Rare, Epic, Legendary

Easter Eggs Hidden surprises triggered by special phrases or patterns. Discover them yourself! (Hint: Try asking about the meaning of life...)

Sounds (Optional) Audio feedback for actions (can be disabled in Settings): - Notification sounds - Achievement unlocks - Transition effects

Adjusting Personality

1. Go to **Settings** → **Personality**
 2. Select your preferred mode
 3. Changes apply immediately
-

8. Collaboration Features

Real-Time Collaboration

Work together with colleagues on the same conversation.

Starting a Collaborative Session

1. Navigate to **Collaborate** from the sidebar
2. Click **Create Session** or join an existing one
3. Share the session link with colleagues
4. Everyone sees messages in real-time

Enhanced Collaboration Features The enhanced collaboration mode includes:

Feature	Description
Chat	Real-time messaging with all participants
Branches	Create alternative discussion threads
AI Roundtable	Multiple AI perspectives on a topic
Knowledge Graph	Visual map of discussed concepts
Playback	Review the conversation timeline

AI Facilitator Enable the AI Facilitator to: - Summarize long discussions - Suggest next topics - Identify areas of agreement/disagreement - Keep the conversation productive

Participant Roles

- **Owner** - Full control, can delete session
 - **Participant** - Can send messages and interact
 - **Guest** - View-only access (via guest link)

Sharing Conversations

Share a conversation without real-time collaboration:

1. Open any conversation
 2. Click the **Share** button
 3. Choose sharing options:
 - **Public link** - Anyone with link can view
 - **Copy allowed** - Viewers can copy content
 4. Click **Create Share Link**
 5. Copy and send the link

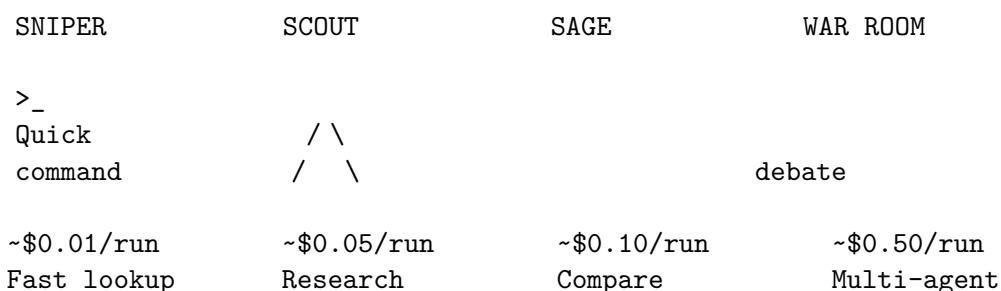
9. Advanced Features

Polymorphic UI - Adaptive Views

Think Tank's interface automatically adapts based on your query type.

View Types

POLYMORPHIC VIEW SELECTION



View	When It Appears	Best For
Sniper (Terminal)	Quick commands, lookups	Fast answers, low cost
Scout (Mind Map)	Research, exploration	Complex research
Sage (Diff Editor)	Validation, comparison	Reviewing changes

View	When It Appears	Best For
Dashboard	Data queries	Analytics
Decision Cards	Choices needed	Human approval required
Chat	General conversation	Default view

Cost Indicators Different views have different costs (reflected in credits used): - **Sniper Mode** - ~\$0.01/run (fast, single model) - **War Room Mode** - ~\$0.50+/run (multi-agent, thorough)

The Economic Governor automatically routes your query to the most cost-effective mode that can handle it.

Manual Escalation If Sniper mode isn't giving good results: 1. Click **Escalate to War Room** 2. The AI will use more models and deeper analysis 3. Results are more thorough but take longer

The Grimoire - AI Learning

The Grimoire is Think Tank's procedural memory - rules it has learned from successful interactions.

What You'll See

- **Heuristics** - Learned rules like "When asked about X, always consider Y"
- **Confidence Scores** - How sure the AI is about each rule
- **Domain Tags** - Which domains the rule applies to

Reinforcing Learning You can help the AI learn: - **Thumbs Up** - Increases confidence in a heuristic - **Thumbs Down** - Decreases confidence - This feedback improves future responses

Magic Carpet Navigator

An advanced navigation system with intent-based routing.

Opening the Navigator Press **K** (Mac) or **Ctrl+K** (Windows) to open the destination selector.

Destinations

Destination	Icon	Purpose
Command Center		Overview dashboard
Workshop		Build and create
Time Stream		Reality Scrubber (history)
Quantum Realm		Parallel realities (branches)
Oracle's Chamber		Pre-Cognition (predictions)
Gallery		View creations
Vault		Saved items

Journey Breadcrumbs The navigator shows your path through the application, making it easy to retrace steps.

Artifacts - Generated Code

When Think Tank generates code or components, they appear as Artifacts.

Artifact Features

- **Live Preview** - See generated UI components
- **Code View** - Inspect the source code
- **Validation** - Safety checks ensure code is secure
- **Reflexion** - If generation fails, AI retries with improvements

Safety Validation All generated code passes through Cato safety validation:
- No dangerous operations
- Only allowed dependencies
- Follows security best practices

Liquid Interface - Chat Morphs Into Tools (v5.52.8)

In **Advanced Mode**, the chat interface can transform (“morph”) into specialized tools when you need them. This is called the Liquid Interface - “Don’t Build the Tool. BE the Tool.”

Enabling Advanced Mode

1. Toggle **Advanced Mode** in the header (lightning bolt icon)
2. Tool trigger buttons appear in the toolbar
3. Click any tool icon to morph the chat into that tool

Available Tools

Tool	Icon	What It Does
Data Grid		Interactive spreadsheet for data manipulation
Chart		Visualize data as bar, line, pie, or area charts
Kanban		Task board with multiple frameworks (see below)
Calculator		Full calculator with memory and operations
Code Editor		Write and run code with output panel
Document		Rich text editor for writing

Kanban Board Variants The Kanban tool supports 5 different productivity frameworks:

Variant	Best For	Key Features
Standard	General task tracking	Traditional columns, drag-and-drop
Scrumban	Agile teams	Sprint goals, velocity, story points
Enterprise	Portfolio management	Multi-lane boards (Strategic/Ops/Support)
Personal	Individual productivity	Simple 3-column, WIP limit of 3
Pomodoro	Focus sessions	Built-in 25-min timer, break tracking

Using Pomodoro Kanban: 1. Select “Pomodoro Kanban” from the variant dropdown 2. Add tasks with estimated pomodoros () 3. Click **Start** on a task to begin a 25-minute focus session 4. Timer shows in header - take a 5-minute break when it ends 5. Track completed pomodoros per task

Analytics Panel: Click **Analytics** to see: - Total tasks and completed count - Average cycle time (how long tasks take) - Throughput (tasks completed per week)

Returning to Chat Click the **X** button in the tool header to close the morphed view and return to chat.

10. How Think Tank’s Memory Works

Think Tank uses two interconnected systems to remember things and access knowledge.

Cato - The AI’s Personality & Memory

Cato is the AI’s “self” - its personality, emotional state, and personal memory of you.

What Cato Remembers	Example
Your Preferences	“This user prefers detailed explanations”
Past Conversations	Topics you’ve discussed, corrections you’ve made
Current Mood	Confidence level, engagement, curiosity
Communication Style	Formal vs casual, concise vs detailed

How it helps you: The AI adapts its responses based on what it knows about you. If you’ve told it you prefer bullet points, it remembers. If you corrected it before, it learns.

Cortex - The Enterprise Knowledge Library

Cortex is your organization’s knowledge graph - facts, documents, and relationships extracted from enterprise data.

Knowledge Tier	What’s There	Speed
Hot	Current session context	Instant
Warm	Knowledge graph, verified facts	Fast
Cold	Archives, compliance data	Slower

How it helps you: When you ask a question, the AI can pull relevant facts from your organization’s knowledge base - not just generic internet knowledge.

How They Work Together

YOUR MESSAGE: "What's the status of Project Alpha?"

CATO checks:

- Your role (Project Manager)
- Your preferences (prefers executive summaries)
- Your mood context (busy, needs quick answers)

CORTEX retrieves:

- Project Alpha timeline (from knowledge graph)
- Recent status updates (from documents)
- Related milestones (from relationships)

AI RESPONSE:

Personalized (knows you want summaries)
+ Informed (has actual project data)
+ Contextual (understands your role)

What This Means for You

Without Memory Integration	With Memory Integration
AI gives generic answers	AI gives personalized + informed answers
You re-explain context every time	AI remembers your preferences
No access to company data	Enterprise facts in every response
Each session starts fresh	Learning persists across sessions

Privacy Note

- **Personal memories** (Cato) are tied to your user account
- **Enterprise knowledge** (Cortex) follows your organization's access controls
- You can ask "What do you remember about me?" to see stored context
- Admins can configure retention periods and what gets remembered

11. Understanding AI Decisions

Think Tank is designed to be transparent about how it makes decisions.

Brain Plans

When Think Tank processes your request, it creates a Brain Plan showing:

BRAIN PLAN

Orchestration: research
Domain: Scientific
Confidence: 87%

EXECUTION STEPS

Step 1: Analyze query context	[Done]
Step 2: Select relevant models	[Done]
Step 3: Gather information	[Running]
Step 4: Synthesize response	[Pending]
Step 5: Validate and format	[Pending]

Model: Claude Sonnet 4.0
Est. Cost: \$0.03
Est. Time: ~8 seconds

1. **Orchestration Mode** - How the AI will approach your question
 - thinking - Standard reasoning
 - extended_thinking - Deep multi-step analysis
 - coding - Code generation
 - creative - Creative writing
 - research - Research synthesis
 - multi_model - Multiple AI perspectives
2. **Domain Detection** - The identified knowledge area
3. **Model Selection** - Which AI model will be used and why
4. **Steps** - The planned execution steps
5. **Cost Estimate** - Expected credits to be used

Confidence Levels

Think Tank indicates how confident it is in responses:

Indicator	Meaning
High Confidence	AI is very sure about this answer
Medium Confidence	Reasonably sure, but verify important details
Low Confidence	Uncertain - treat as a starting point

When Think Tank Asks for Clarification

If the AI is uncertain about your intent, it will ask clarifying questions rather than guess. This is intentional - it's better to ask than give a wrong answer.

Epistemic Humility

Think Tank acknowledges when it doesn't know something: - "I'm not certain, but..." - "Based on my training data (which may be outdated)..." - "I don't have enough information to answer this definitively"

14. Safety & Governance

Think Tank includes multiple safety layers to protect you and ensure responsible AI use.

Five-Layer Safety Stack

L4	COGNITIVE LAYER	
	Active Inference • Precision Governor • Planning	
L3	CONTROL LAYER	ALWAYS ENFORCED
	Control Barrier Functions • CANNOT be bypassed	
L2	PERCEPTION LAYER	
	Uncertainty Detection • Fracture Prevention	
L1	SENSORY LAYER	
	Immediate Veto • Dangerous Request Blocking	
L0	RECOVERY LAYER	
	Safe State Return • Error Recovery	

Layer	Name	What It Does
L4	Cognitive	Active inference, precision control
L3	Control	Barrier functions - always enforced
L2	Perception	Uncertainty detection, fracture prevention
L1	Sensory	Immediate veto for dangerous requests
L0	Recovery	Returns to safe state if issues occur

Governance Presets

Your organization may use different governance levels:

Preset	Icon	Behavior
Paranoid		Every action requires approval
Balanced		Auto-approve low-risk, checkpoint medium/high
Cowboy		Full autonomy with notifications

You'll see a governance badge indicating the current level.

Human-in-the-Loop (HITL)

For high-stakes decisions, Think Tank may pause and ask for your approval:

Approval Required

This action will modify your database.
Cost: \$2.50 estimated

[Approve] [Modify] [Reject]

What Think Tank Will Never Do

These are hardcoded safety limits that cannot be changed: - Generate harmful content - Execute destructive actions without confirmation - Bypass safety barriers - Delete audit logs - Expose sensitive data

Reporting Issues

If Think Tank produces concerning output: 1. Click the **Report** button on the response 2. Select the issue type 3. Add any additional context 4. Submit for review

12. Decision Records

Decision Records capture the AI's reasoning, evidence, and conclusions in an auditable format. This feature helps you understand and verify AI-assisted decisions.

Accessing Decision Records

After significant conversations, Think Tank automatically extracts: - **Claims** - Key conclusions and recommendations - **Evidence** - Supporting data and sources - **Dissent** - Alternative viewpoints considered but rejected - **Compliance** - Regulatory implications if applicable

The Living Parchment View

Decision Records use a special “Living Parchment” interface where: - **Breathing colors** indicate trust levels (green = verified, amber = unverified, red = contested) - **Font weight** reflects confidence (bolder = more confident) - **Ghost paths** show rejected alternatives as faded traces

Verifying Claims

Click any claim to see: 1. The evidence supporting it 2. The AI’s reasoning chain 3. Any dissenting opinions 4. Data freshness indicators

Exporting for Compliance

Export decision records in various formats: - **HIPAA Audit Package** - For healthcare compliance
- **SOC2 Evidence Bundle** - For security audits - **GDPR DSAR Response** - For data requests

Exporting Conversations Directly (v5.52.16)

You can export any conversation directly from the sidebar:

1. **Hover** over any conversation in the sidebar
2. **Click** the (more options) button that appears
3. **Select** an export format:
 - **Generate Decision Record** - Creates a Decision Intelligence Artifact with claims, evidence, and dissent
 - **Export HIPAA Audit Package** - PHI-redacted export for healthcare compliance
 - **Export SOC2 Evidence** - Audit trail for security compliance
 - **Export GDPR DSAR** - Data subject access request format
 - **Export as PDF** - Standard PDF export

SIDE BAR – Conversation Actions

Today

"Drug interaction analysis" [] []

Generate Decision
Record

Export HIPAA
Export SOC2
Export GDPR DSAR

Export as PDF

Note: PHI (Protected Health Information) is automatically redacted in compliance exports by default.

13. Living Parchment

Living Parchment is Think Tank's advanced decision intelligence suite with sensory UI that communicates trust through visual breathing, living typography, and confidence terrain.

War Room (Strategic Decision Theater)

For high-stakes decisions, enter the War Room:

1. **Navigate to** Living Parchment → War Room
2. **Create a session** with your decision question
3. **Add AI advisors** - multiple perspectives on your problem
4. **View the Confidence Terrain** - 3D visualization where height = confidence
5. **Explore Decision Paths** - branching options with predicted outcomes
6. **Make your decision** with full documentation

Understanding the Terrain: - Green peaks = High confidence areas - Amber slopes = Moderate uncertainty

- Red valleys = Risk zones requiring attention

Council of Experts

Summon diverse AI perspectives that debate and converge:

1. **Convene a Council** with your question
2. **Watch 8 expert personas** discuss:
 - Pragmatist (practical focus)
 - Ethicist (moral considerations)
 - Innovator (creative solutions)
 - Skeptic (devil's advocate)
 - Synthesizer (finding common ground)
 - Analyst (data-driven insights)
 - Strategist (long-term thinking)
 - Humanist (human impact)
3. **Observe consensus forming** as experts move toward center
4. **Review minority reports** - valid dissenting views preserved

Debate Arena

Test any idea through adversarial exploration:

1. **Create a debate** with your proposition
2. **Watch AI debaters** argue both sides
3. **Track the Resolution Meter** showing which side is winning

4. **Identify weak points** (breathing red indicators)
5. **Generate Steel-Man** - AI creates the strongest version of the opposing argument

Understanding Living Parchment UI

Visual Element	Meaning
Fast breathing (12 BPM)	High uncertainty, needs attention
Slow breathing (4-6 BPM)	Confident, stable information
Bold text	High confidence claim
Light text	Lower confidence, verify before acting
Faded/gray text	Stale information, may need refresh
Ghost overlays	Rejected alternatives (what could have been)

14. Safety & Governance

Think Tank includes multiple safety layers to protect you and your organization. See Section 14: Safety & Governance in the main guide for details on: - Five-layer Cato safety architecture - Control Barrier Functions (CBFs) - Human-in-the-Loop approvals - Governance presets

15. Keyboard Shortcuts

Shortcut	Action
K / Ctrl+K	Open Magic Carpet Navigator
Enter / Ctrl+Enter	Send message
Escape	Close dialogs/modals
/ / Ctrl+ /	Open keyboard shortcuts help
N / Ctrl+N	New conversation
S / Ctrl+S	Save/Export conversation

16. Troubleshooting

Common Issues

“Response is taking too long”

- Complex queries in War Room mode take longer
- Check your network connection
- Try simplifying your question

“AI gave an incorrect answer”

1. Check if you’re in the right domain mode

2. Provide more context in your question
3. Use the thumbs down button to provide feedback
4. Consider adding a rule to prevent this in the future

“I can’t access a feature”

- Some features require specific permissions
- Contact your administrator for access
- Check if the feature is enabled for your organization

“My rules aren’t being applied”

1. Verify the rule is toggled ON
2. Check the rule isn’t too vague
3. Make sure the rule doesn’t conflict with other rules
4. Try making the rule more specific

Getting Help

- **In-App Help** - Click the ? icon in the header
 - **Documentation** - Access guides from Settings
 - **Support** - Contact your organization’s IT team
-

17. Glossary

Term	Definition
Brain Plan	Think Tank’s execution plan showing how it will answer your question
Breathing UI	Visual elements that pulse to communicate confidence and data freshness
CBF	Control Barrier Function - safety guardrails that cannot be bypassed
Confidence Terrain	3D visualization where elevation = confidence, color = risk
Council of Experts	Multi-persona AI consultation with 8 distinct viewpoints
Debate Arena	Adversarial exploration tool for stress-testing ideas
Decision Record	Auditable capture of AI reasoning, evidence, and conclusions
Delight	The personality and engagement system
Domain Mode	Specialized configuration for different knowledge areas
Ego	Think Tank’s persistent identity and emotional state
Ghost Path	Translucent overlay showing rejected alternatives
Governance Preset	Organization-wide safety/autonomy settings
Grimoire	The AI’s learned procedural memory
Heuristic	A learned rule or shortcut the AI uses
HTL	Human-in-the-Loop - requiring human approval
Living Ink	Typography that varies weight based on confidence (350-500)
Living Parchment	Advanced decision intelligence suite with sensory UI

Term	Definition
Magic Carpet	Intent-based navigation system
My Rules	Personal preferences that customize AI responses
Polymorphic UI	Interface that adapts based on query type
Sniper Mode	Fast, low-cost single-model execution
Steel-Man	AI-generated strongest version of an opposing argument
War Room	Strategic Decision Theater for high-stakes collaborative decisions
War Room Mode	Thorough multi-agent execution

Document History

Version	Date	Changes
5.35.0	Jan 2026	Initial comprehensive user guide
5.35.0	Jan 2026	Added visual diagrams and flowcharts
5.43.0	Jan 22, 2026	Added Decision Records section (DIA Engine)
5.44.0	Jan 22, 2026	Added Living Parchment section (War Room, Council, Debate Arena)
5.52.0	Jan 23, 2026	Simulator now uses real API data with graceful fallbacks

Think Tank is designed to be your intelligent partner. The more you use it and customize it to your needs, the more valuable it becomes. Happy thinking!