

# Contents

<b>RADIANT v4.17.0 - Troubleshooting Guide</b>	<b>1</b>
Common Issues and Solutions . . . . .	1
CDK Deployment Failures . . . . .	1
Aurora Database Issues . . . . .	1
Lambda Function Errors . . . . .	2
LiteLLM / ECS Issues . . . . .	2
SageMaker Issues (Tier 3+) . . . . .	3
Cognito Authentication Issues . . . . .	3
Admin Dashboard Issues . . . . .	3
Log Locations . . . . .	4
Viewing Logs . . . . .	4
Health Check Endpoints . . . . .	4
Emergency Procedures . . . . .	5
Database Restore from Snapshot . . . . .	5
Point-in-Time Recovery . . . . .	5
Rollback CDK Deployment . . . . .	5
Disable Problematic Model . . . . .	5
Force Scale Down SageMaker . . . . .	5
Performance Benchmarks . . . . .	6
Support Checklist . . . . .	6
Useful AWS CLI Commands . . . . .	6

## RADIANT v4.17.0 - Troubleshooting Guide

### Common Issues and Solutions

#### CDK Deployment Failures

Issue	Likely Cause	Solution
Bootstrap failed	Wrong account/region	Verify <code>aws sts get-caller-identity</code>
Stack timeout	Slow resource creation	Check CloudFormation events in AWS Console
Resource limit	Service quota exceeded	Request quota increase via AWS Service Quotas
IAM permission denied	Insufficient permissions	Ensure IAM user has AdministratorAccess
Circular dependency	Stack references	Check stack dependencies in CDK code
Asset upload failed	S3 bucket permissions	Verify CDK bootstrap bucket exists

#### Aurora Database Issues

Issue	Likely Cause	Solution
Connection refused	Security group rules	Verify Lambda SG can reach Aurora SG on port 5432
Authentication failed	Wrong credentials	Check Secrets Manager for correct credentials
Connection timeout	Missing VPC endpoints	Add RDS VPC endpoint to private subnets
Too many connections	Connection exhaustion	Use RDS Proxy or increase max_connections
Slow queries	Missing indexes	Run EXPLAIN ANALYZE and add appropriate indexes
RLS blocking access	Tenant ID not set	Ensure <code>app.current_tenant_id</code> is set in session

## Lambda Function Errors

Issue	Likely Cause	Solution
Cold start > 10s	VPC attachment	Use provisioned concurrency for critical functions
Timeout	Slow downstream services	Increase timeout, check DB/API latency
Out of memory	Large payloads/responses	Increase memory allocation (also increases CPU)
Permission denied	IAM role misconfigured	Check Lambda execution role policies
Module not found	Missing dependency	Verify all dependencies in package.json
Handler not found	Incorrect handler path	Check function configuration in CDK

## LiteLLM / ECS Issues

Issue	Likely Cause	Solution
503 Service Unavailable	ECS task unhealthy	Check ECS service events and task logs
Provider timeout	Invalid API key	Verify provider secrets in Secrets Manager
Rate limited	Too many requests	Implement exponential backoff retry
Wrong model response	Model misconfigured	Check config.yaml model mappings
Container crashes	Memory exhaustion	Increase task memory in CDK

Issue	Likely Cause	Solution
No healthy targets	Health check failing	Verify health check endpoint returns 200

### SageMaker Issues (Tier 3+)

Issue	Likely Cause	Solution
Endpoint failed to create	Insufficient capacity	Try different instance type or region
InvocationError	Model loading failed	Check CloudWatch logs for model errors
Slow cold start	Large model size	Use warm pools or smaller model variant
Capacity error	Instance quota reached	Request SageMaker quota increase
Timeout	Long inference time	Increase endpoint timeout or optimize model

### Cognito Authentication Issues

Issue	Likely Cause	Solution
Invalid grant	Expired refresh token	Re-authenticate user
User not confirmed	Email not verified	Check email or manually confirm user
MFA required	MFA not set up	Complete MFA setup flow
Invalid client	Wrong client ID	Verify app client ID in configuration
Callback URL mismatch	URL not whitelisted	Add URL to allowed callbacks in Cognito

### Admin Dashboard Issues

Issue	Likely Cause	Solution
403 Forbidden	CloudFront OAC issue	Verify S3 bucket policy allows CloudFront
API calls fail	CORS configuration	Check API Gateway CORS settings
Login redirect loop	Cookie domain mismatch	Verify cookie domain matches site domain
Blank page	Build error	Check <code>next build</code> output for errors
Slow load	Large bundle size	Enable code splitting and lazy loading

## Log Locations

Component	CloudWatch Log Group
API Gateway	/aws/api-gateway/radiant-{env}-api
Lambda Functions	/aws/lambda/Radiant-{env}-*
LiteLLM (ECS)	/ecs/radiant-{env}-litellm
SageMaker Endpoints	/aws/sagemaker/Endpoints/radiant-*
Aurora PostgreSQL	/aws/rds/cluster/radiant-{env}/postgresql
CloudFront	Standard CloudFront logs in S3

## Viewing Logs

```
# Tail Lambda logs in real-time
aws logs tail /aws/lambda/Radiant-dev-router --follow

# View ECS logs
aws logs tail /ecs/radiant-dev-litellm --follow

# Search logs for errors
aws logs filter-log-events \
--log-group-name /aws/lambda/Radiant-dev-router \
--filter-pattern "ERROR" \
--start-time $(date -d '1 hour ago' +%s)000
```

---

## Health Check Endpoints

```
# Platform API health
curl https://api.YOUR_DOMAIN/health
# Expected: {"status": "healthy", "version": "4.17.0"}

# LiteLLM health
curl https://api.YOUR_DOMAIN/v2/litellm/health
# Expected: {"status": "healthy"}

# Admin API health
curl https://admin-api.YOUR_DOMAIN/health

# Model registry status
curl https://api.YOUR_DOMAIN/v2/models/status
```

---

## Emergency Procedures

### Database Restore from Snapshot

```
# List available snapshots
aws rds describe-db-cluster-snapshots \
  --db-cluster-identifier radiant-prod-cluster \
  --query 'DBClusterSnapshots[*].[DBClusterSnapshotIdentifier,SnapshotCreateTime]' \
  --output table

# Restore from snapshot
aws rds restore-db-cluster-from-snapshot \
  --db-cluster-identifier radiant-prod-restored \
  --snapshot-identifier your-snapshot-id \
  --engine aurora-postgresql \
  --vpc-security-group-ids sg-xxx \
  --db-subnet-group-name radiant-prod-db-subnet
```

### Point-in-Time Recovery

```
aws rds restore-db-cluster-to-point-in-time \
  --source-db-cluster-identifier radiant-prod-cluster \
  --db-cluster-identifier radiant-prod-recovered \
  --restore-to-time "2024-12-20T10:00:00Z" \
  --vpc-security-group-ids sg-xxx \
  --db-subnet-group-name radiant-prod-db-subnet
```

### Rollback CDK Deployment

```
# Rollback to previous deployment
cd packages/infrastructure
npx cdk deploy Radiant-prod-API \
  --context environment=prod \
  --context tier=3 \
  --rollback
```

### Disable Problematic Model

```
-- Connect to Aurora and run:
UPDATE models
SET status = 'disabled',
    disabled_reason = 'Emergency disable due to errors',
    updated_at = NOW()
WHERE model_id = 'problematic-model-id';
```

### Force Scale Down SageMaker

```
# Scale endpoint to 0 instances
aws sagemaker update-endpoint-weights-and-capacities \
```

```
--endpoint-name radiant-prod-model-endpoint \
--desired-weights-and-capacities '[{"VariantName": "AllTraffic", "DesiredInstanceCount": 0}]'
```

---

## Performance Benchmarks

Metric	Target	Acceptable	Action if Exceeded
API Gateway p50 latency	< 50ms	< 100ms	Check Lambda cold starts
API Gateway p99 latency	< 200ms	< 500ms	Enable provisioned concurrency
Chat streaming start	< 500ms	< 1s	Check LiteLLM/provider latency
Admin dashboard load	< 2s	< 3s	Optimize bundle, enable CDN caching
Model warm-up time	< 3 min	< 5 min	Use larger instance or warm pools
Aurora query latency	< 10ms	< 50ms	Add indexes, optimize queries

---

## Support Checklist

When reporting issues, include:

1. **Environment:** dev/staging/prod
  2. **Tier:** 1-5
  3. **Error message:** Full error text
  4. **Request ID:** From response headers
  5. **Timestamp:** When the error occurred
  6. **Steps to reproduce:** What actions led to the error
  7. **Relevant logs:** CloudWatch log excerpts
- 

## Useful AWS CLI Commands

```
# Check stack status
aws cloudformation describe-stacks --stack-name Radiant-dev-API \
--query 'Stacks[0].StackStatus'

# List recent CloudFormation events
aws cloudformation describe-stack-events --stack-name Radiant-dev-API \
--query 'StackEvents[0:10].[Timestamp,ResourceStatus,ResourceType,LogicalResourceId]' \
--output table

# Check Lambda function configuration
```

```
aws lambda get-function-configuration --function-name Radiant-dev-router

# List ECS services
aws ecs list-services --cluster radiant-dev-cluster

# Describe ECS service
aws ecs describe-services --cluster radiant-dev-cluster \
--services radiant-dev-litellm

# Check Secrets Manager secret
aws secretsmanager get-secret-value --secret-id radiant/dev/db-credentials \
--query 'SecretString' --output text | jq .
```