

RADIANT API Reference

RADIANT Team

2025-12-25

Contents

RADIANT API Reference	1
Chat Completions	2
Create Chat Completion	2
Models	3
List Models	3
Get Model	4
Embeddings	4
Create Embeddings	4
Billing	5
Get Credit Balance	5
Get Usage	5
Webhooks	5
List Webhooks	5
Create Webhook	5
Test Webhook	6
Batch Processing	6
Create Batch Job	6
Get Batch Job	6
List Batch Jobs	7
Error Codes	7
Error Categories	7
Common Error Codes	7
Rate Limits	8
SDKs	8
TypeScript/JavaScript	8
Python	8
CLI	9
Changelog	9
Support	9

RADIANT API Reference

Complete API documentation for the RADIANT AI Platform.

Base URL: <https://api.radiant.example.com/v2>

Authentication: Bearer token (API key or JWT)

Authorization: Bearer rad_your_api_key

Chat Completions

Create Chat Completion

POST /v2/chat/completions

Create a chat completion with any supported model.

Request Body:

Field	Type	Required	Description
model	string		Model ID (e.g., gpt-4o, claude-3-sonnet)
messages	array		Array of message objects
max_tokens	integer		Maximum tokens to generate
temperature	number		Sampling temperature (0-2)
top_p	number		Nucleus sampling (0-1)
stream	boolean		Stream response tokens
stop	string/array		Stop sequences
functions	array		Function definitions
function_call	string/object		Function calling mode

Message Object:

Field	Type	Required	Description
role	string		system, user, assistant, function
content	string		Message content
name	string		Function name (for function messages)

Example Request:

```
{  
  "model": "gpt-4o",  
  "messages": [  
    {"role": "system", "content": "You are a helpful assistant."},  
    {"role": "user", "content": "Hello!"}  
  ],  
  "max_tokens": 1000,  
  "temperature": 0.7  
}
```

Response:

```
{
  "id": "chatcmpl_abc123",
  "object": "chat.completion",
  "created": 1703980800,
  "model": "gpt-4o",
  "choices": [
    {
      "index": 0,
      "message": {
        "role": "assistant",
        "content": "Hello! How can I help you today?"
      },
      "finish_reason": "stop"
    }
  ],
  "usage": {
    "prompt_tokens": 25,
    "completion_tokens": 10,
    "total_tokens": 35
  }
}
```

Models

List Models

GET /v2/models

List all available models.

Query Parameters:

Parameter	Type	Description
category	string	Filter by category (chat, embedding, image)
provider	string	Filter by provider (openai, anthropic, google)

Response:

```
{
  "object": "list",
  "data": [
    {
      "id": "gpt-4o",
      "object": "model",
      "created": 1703980800,
```

```

    "owned_by": "openai",
    "display_name": "GPT-4o",
    "category": "chat",
    "context_window": 128000,
    "input_cost_per_1k": 0.005,
    "output_cost_per_1k": 0.015,
    "capabilities": ["chat", "vision", "function_calling"]
  }
]
}

```

Get Model

GET /v2/models/{model_id}

Get details for a specific model.

Embeddings

Create Embeddings

POST /v2/embeddings

Generate embeddings for text.

Request Body:

Field	Type	Required	Description
model	string		Embedding model ID
input	string/array		Text to embed
encoding_format	string		float or base64

Response:

```
{
  "object": "list",
  "data": [
    {
      "object": "embedding",
      "index": 0,
      "embedding": [0.0023, -0.0091, ...]
    }
  ],
  "model": "text-embedding-3-small",
  "usage": {
    "prompt_tokens": 8,
    "total_tokens": 8
  }
}
```

```
    }  
}
```

Billing

Get Credit Balance

```
GET /v2/billing/credits
```

Get current credit balance.

Response:

```
{  
  "data": {  
    "available": 150.50,  
    "reserved": 10.00,  
    "currency": "USD",  
    "updated_at": "2024-01-15T10:30:00Z"  
  }  
}
```

Get Usage

```
GET /v2/billing/usage
```

Get usage data for a period.

Query Parameters:

Parameter	Type	Description
<code>start_date</code>	string	Start date (YYYY-MM-DD)
<code>end_date</code>	string	End date (YYYY-MM-DD)
<code>group_by</code>	string	day, model, endpoint

Webhooks

List Webhooks

```
GET /v2/webhooks
```

List configured webhooks.

Create Webhook

```
POST /v2/webhooks
```

Request Body:

```
{  
  "url": "https://your-server.com/webhook",  
  "event_types": ["billing.low_balance", "usage.quota_reached"],  
  "description": "Billing alerts"  
}
```

Response:

```
{  
  "data": {  
    "id": "wh_abc123",  
    "url": "https://your-server.com/webhook",  
    "secret": "whsec_xyz789...",  
    "event_types": ["billing.low_balance", "usage.quota_reached"],  
    "is_active": true,  
    "created_at": "2024-01-15T10:30:00Z"  
  }  
}
```

Test Webhook

POST /v2/webhooks/{webhook_id}/test

Send a test event to the webhook.

Batch Processing

Create Batch Job

POST /v2/batch/jobs

Create a batch processing job.

Request Body:

```
{  
  "type": "completions",  
  "model": "gpt-4o",  
  "input_file": "batch-input.jsonl",  
  "options": {  
    "system_prompt": "You are a helpful assistant.",  
    "max_tokens": 500  
  }  
}
```

Get Batch Job

GET /v2/batch/jobs/{job_id}

Get batch job status and results.

List Batch Jobs

GET /v2/batch/jobs

List all batch jobs.

Error Codes

RADIANT uses standardized error codes across all endpoints. See [Error Codes Reference](#) for the complete list.

Error Categories

Category	Code Range	Description
Authentication	RADIANT_AUTH_1xxx	Token, API key, session errors
Authorization	RADIANT_AUTHZ_2xxx	Permission, role, tenant errors
Validation	RADIANT_VAL_3xxx	Input validation errors
Resource	RADIANT_RES_4xxx	Not found, conflict, quota errors
Rate Limiting	RADIANT_RATE_5xxx	Throttling and rate limit errors
AI/Model	RADIANT_AI_6xxx	Model, provider, inference errors
Billing	RADIANT_BILL_7xxx	Credits, subscription errors
Storage	RADIANT_STOR_8xxx	File upload, storage errors
Internal	RADIANT_INT_9xxx	Server, database, timeout errors

Common Error Codes

Code	HTTP	Retryable	Description
RADIANT_AUTH_1001	401		Invalid authentication token
RADIANT_AUTH_1004	401		Invalid API key
RADIANT_VAL_3001	400		Required field missing
RADIANT_RES_4001	404		Resource not found
RADIANT_RATE_5001	429		Rate limit exceeded
RADIANT_BILL_7001	402		Insufficient credits
RADIANT_AI_6004	502		AI provider error
RADIANT_INT_9001	500		Internal server error

Error Response Format:

```
{  
  "error": {  
    "code": "RADIANT_RATE_5001",  
    "message": "Too many requests. Please slow down.",  
    "category": "rate_limit",  
    "retryable": true,  
    "timestamp": "2024-12-25T10:30:00.000Z"  
  }  
}
```

```
    }  
}
```

Retry-After Header: Retryable errors include `Retry-After` header with seconds to wait.

Rate Limits

Tier	Requests/min	Tokens/min
Free	10	10,000
Starter	50	50,000
Professional	100	200,000
Business	500	1,000,000
Enterprise	2,000	Unlimited

Rate limit headers:

```
X-RateLimit-Limit: 100  
X-RateLimit-Remaining: 95  
X-RateLimit-Reset: 1703980860
```

SDKs

TypeScript/JavaScript

```
npm install @radiant/sdk  
  
import { RadiantClient } from '@radiant/sdk';  
  
const client = new RadiantClient({ apiKey: 'your-key' });  
const response = await client.chat.create({  
  model: 'gpt-4o',  
  messages: [{ role: 'user', content: 'Hello!' }],  
});
```

Python

```
pip install radiant-sdk  
  
from radiant import RadiantClient  
  
client = RadiantClient(api_key="your-key")  
response = client.chat.create(  
  model="gpt-4o",  
  messages=[{"role": "user", "content": "Hello!"}],  
)
```

CLI

```
npm install -g @radiant/cli
radiant auth login
radiant chat send "Hello!"
```

Changelog

See [CHANGELOG.md](#) for version history.

Support

- **Email:** support@radiant.example.com
- **Documentation:** <https://docs.radiant.example.com>
- **Status:** <https://status.radiant.example.com>