

# RADIANT Platform v4.18.2 - Complete Documentation

RADIANT Team

2025-12-28

# Contents

<b>RADIANT v4.17.0 - Deployed System Guide</b>	<b>15</b>
Table of Contents . . . . .	15
1. System Overview . . . . .	15
1.1 Platform Summary . . . . .	15
1.2 Environment Tiers . . . . .	16
1.3 Region Configuration . . . . .	16
2. Infrastructure Components . . . . .	16
2.1 VPC Architecture . . . . .	16
2.2 Security Groups . . . . .	17
2.3 IAM Roles . . . . .	17
3. Deployment Architecture . . . . .	17
3.1 CDK Stacks . . . . .	17
3.2 Deployment Commands . . . . .	17
3.3 Environment Variables . . . . .	17
4. Service Endpoints . . . . .	18
4.1 API Gateway Endpoints . . . . .	18
4.2 Health Check Endpoints . . . . .	18
4.3 Internal Endpoints . . . . .	18
5. Database Operations . . . . .	18
5.1 Aurora Configuration . . . . .	18
5.2 Connection Management . . . . .	19
5.3 Row-Level Security . . . . .	19
5.4 Database Migrations . . . . .	19
6. Lambda Functions . . . . .	19
6.1 Function Inventory . . . . .	19
6.2 Monitoring Lambda . . . . .	20
6.3 Cold Start Optimization . . . . .	20
7. AI Provider Integration . . . . .	20
7.1 External Providers . . . . .	20
7.2 Self-Hosted Models (SageMaker) . . . . .	20
7.3 Model Thermal States . . . . .	21
8. Monitoring & Observability . . . . .	21
8.1 CloudWatch Dashboards . . . . .	21
8.2 Key Metrics . . . . .	21
8.3 Alerting . . . . .	21
8.4 Log Groups . . . . .	22
9. Security Configuration . . . . .	22
9.1 Encryption . . . . .	22
9.2 Secrets Management . . . . .	22
9.3 WAF Rules . . . . .	22
10. Backup & Recovery . . . . .	22

10.1 Automated Backups . . . . .	22
10.2 Recovery Procedures . . . . .	23
10.3 Disaster Recovery . . . . .	23
11. Maintenance Procedures . . . . .	23
11.1 Routine Maintenance . . . . .	23
11.2 Scaling Operations . . . . .	23
11.3 Deployment Checklist . . . . .	24
12. Troubleshooting . . . . .	24
12.1 Common Issues . . . . .	24
12.2 Emergency Procedures . . . . .	24
Version History . . . . .	25
<b>RADIANT Architecture v4.17.0</b>	<b>26</b>
Overview . . . . .	26
Technology Stack . . . . .	26
Monorepo Structure . . . . .	26
Phase 1 Architecture . . . . .	26
@radian/shared Package . . . . .	26
Swift Deployer App . . . . .	27
CDK Infrastructure Stacks . . . . .	27
Infrastructure Tiers . . . . .	27
Deployment Flow . . . . .	27
Phase 1 Deliverables . . . . .	28
Future Phases . . . . .	28
<b>RADIANT v4.18.0 - Deployment Guide</b>	<b>29</b>
Overview . . . . .	29
What's New in v4.18.0 . . . . .	29
Prerequisites . . . . .	29
AWS Account Setup . . . . .	29
Development Environment . . . . .	29
Quick Start . . . . .	30
Automated Deployment . . . . .	30
Verify Deployment . . . . .	30
Manual Deployment . . . . .	30
1. Install Dependencies . . . . .	30
2. Build Shared Package . . . . .	30
3. Build Lambda Functions . . . . .	30
4. Build Admin Dashboard . . . . .	30
CDK Deployment . . . . .	31
Bootstrap CDK (One-time per account/region) . . . . .	31
Deploy All Stacks . . . . .	31
Deploy Individual Stacks . . . . .	31
Database Migrations . . . . .	31
Post-Deployment Configuration . . . . .	31
Create First Super Admin . . . . .	31
Configure AI Providers . . . . .	32
Environment Configuration . . . . .	32
Tiers . . . . .	32
Environment Variables . . . . .	32
Verification . . . . .	32
Health Checks . . . . .	32
Smoke Tests . . . . .	33
Troubleshooting . . . . .	33

Common Issues . . . . .	33
Logs . . . . .	33
Cleanup . . . . .	33
CI/CD Pipeline . . . . .	33
Workflow Stages . . . . .	33
Pre-commit Hooks . . . . .	34
Manual Deployment from CI . . . . .	34
Testing Before Deployment . . . . .	34
Support . . . . .	34
<b>RADIANT Platform - Administrator Guide</b>	<b>35</b>
Table of Contents . . . . .	35
1. Introduction . . . . .	35
1.1 What is RADIANT? . . . . .	35
1.2 Administrator Roles . . . . .	36
1.3 Key Concepts . . . . .	36
2. Accessing the Admin Dashboard . . . . .	37
2.1 URL and Login . . . . .	37
2.2 First Login . . . . .	37
2.3 Session Management . . . . .	37
2.4 Password Requirements . . . . .	38
3. Dashboard Overview . . . . .	38
3.1 Main Dashboard . . . . .	38
3.2 Navigation Menu . . . . .	38
4. Tenant Management . . . . .	39
4.1 Viewing Tenants . . . . .	39
4.2 Creating a Tenant . . . . .	39
4.3 Tenant Details . . . . .	39
4.4 Tenant Actions . . . . .	40
4.5 Data Isolation . . . . .	40
5. User & Administrator Management . . . . .	40
5.1 Administrator Roles . . . . .	40
5.2 Managing Administrators . . . . .	40
5.3 Viewing Tenant Users . . . . .	41
5.4 User Actions . . . . .	41
6. AI Model Configuration . . . . .	41
6.1 Model Registry . . . . .	41
6.2 Model Categories . . . . .	42
6.3 Model Configuration . . . . .	42
6.4 Self-Hosted Models . . . . .	43
6.5 Thermal States (Self-Hosted) . . . . .	43
7. Provider Management . . . . .	43
7.1 External Providers . . . . .	43
7.2 Adding Provider Credentials . . . . .	44
7.3 Provider Health Monitoring . . . . .	44
7.4 Fallback Configuration . . . . .	44
8. Billing & Subscriptions . . . . .	44
8.1 Subscription Tiers . . . . .	44
8.2 Credit System . . . . .	45
8.3 Managing Subscriptions . . . . .	45
8.4 Usage Reports . . . . .	45
8.5 Billing Alerts . . . . .	45
9. Storage Management . . . . .	46
9.1 Storage Overview . . . . .	46

9.2 Storage Tiers . . . . .	46
9.3 File Management . . . . .	46
10. Orchestration & Neural Engine . . . . .	46
10.1 Brain Router . . . . .	46
10.2 Neural Patterns . . . . .	47
10.3 Workflow Templates . . . . .	47
11. Localization . . . . .	47
11.1 Translation Management . . . . .	47
11.2 Supported Languages . . . . .	47
11.3 AI Translation . . . . .	47
12. Configuration Management . . . . .	48
12.1 System Configuration . . . . .	48
12.2 Tenant Overrides . . . . .	48
12.3 SSM Parameters . . . . .	48
13. Security & Compliance . . . . .	48
13.1 Security Dashboard . . . . .	48
13.2 Anomaly Detection . . . . .	49
13.3 Compliance Reports . . . . .	49
13.4 Generating Reports . . . . .	49
14. Cost Analytics . . . . .	49
14.1 Cost Dashboard . . . . .	49
14.2 Cost Alerts . . . . .	50
14.3 Cost Optimization . . . . .	50
15. A/B Testing & Experiments . . . . .	50
15.1 Experiment Dashboard . . . . .	50
15.2 Creating an Experiment . . . . .	50
15.3 Statistical Analysis . . . . .	51
16. Audit & Monitoring . . . . .	51
16.1 Audit Logs . . . . .	51
16.2 Log Filtering . . . . .	51
16.3 Log Export . . . . .	51
16.4 Real-Time Monitoring . . . . .	51
17. Database Migrations . . . . .	52
17.1 Migration Workflow . . . . .	52
17.2 Pending Migrations . . . . .	52
17.3 Approving Migrations . . . . .	52
18. API Management . . . . .	52
18.1 API Keys . . . . .	52
18.2 Rate Limiting . . . . .	52
18.3 Webhooks . . . . .	53
19. Troubleshooting . . . . .	53
19.1 Common Issues . . . . .	53
19.2 Support Resources . . . . .	53
19.3 Log Locations . . . . .	54
20. Delight System Administration . . . . .	54
20.1 Accessing Delight Admin . . . . .	54
20.2 Dashboard Overview . . . . .	54
20.3 Managing Categories . . . . .	54
20.4 Managing Messages . . . . .	54
20.5 Statistics Dashboard . . . . .	55
20.6 Managing Achievements . . . . .	55
20.7 Managing Easter Eggs . . . . .	55
20.8 API Endpoints . . . . .	55
Appendix: Quick Reference . . . . .	56

Keyboard Shortcuts . . . . .	56
Status Indicators . . . . .	56
<b>RADIANT Deployer - Administrator Guide</b>	<b>57</b>
Table of Contents . . . . .	57
1. Introduction . . . . .	57
1.1 What is RADIANT Deployer? . . . . .	57
1.2 Architecture . . . . .	58
1.3 Key Features . . . . .	58
2. System Requirements . . . . .	59
2.1 Hardware Requirements . . . . .	59
2.2 Software Requirements . . . . .	60
2.3 AWS Requirements . . . . .	60
3. Installation . . . . .	61
3.1 Download and Install . . . . .	61
3.2 Initial Permissions . . . . .	62
3.3 Verify Installation . . . . .	62
4. First-Time Setup . . . . .	62
4.1 Setup Wizard . . . . .	62
4.2 AWS Credentials Setup . . . . .	62
4.3 Environment Configuration . . . . .	63
4.4 AI Assistant Setup (Optional) . . . . .	64
5. AWS Credentials Management . . . . .	64
5.1 Credential Sets . . . . .	64
5.2 Adding Credentials . . . . .	64
5.3 Credential Validation . . . . .	64
5.4 Security Best Practices . . . . .	64
5.5 Importing from AWS CLI . . . . .	65
6. Deployment Operations . . . . .	65
6.1 Deployment Dashboard . . . . .	65
6.2 Starting a Deployment . . . . .	65
6.3 Deployment Phases . . . . .	65
6.4 Deployment Progress . . . . .	66
6.5 Deployment Settings . . . . .	67
6.6 Operation Timeouts . . . . .	67
7. Multi-Region Deployments . . . . .	67
7.1 Overview . . . . .	67
7.2 Supported Regions . . . . .	68
7.3 Adding a Region . . . . .	68
7.4 Region Consistency . . . . .	68
8. Snapshots & Rollbacks . . . . .	68
8.1 Snapshot Types . . . . .	68
8.2 Creating a Snapshot . . . . .	68
8.3 Snapshot Contents . . . . .	69
8.4 Restoring from Snapshot . . . . .	69
8.5 Automatic Rollback . . . . .	69
9. AI Assistant . . . . .	69
9.1 Overview . . . . .	69
9.2 Enabling AI Assistant . . . . .	70
9.3 Using the Assistant . . . . .	70
9.4 Sample Interactions . . . . .	70
9.5 Offline Mode . . . . .	70
10. Package Management . . . . .	71
10.1 Package System . . . . .	71

10.2 Viewing Packages . . . . .	71
10.3 Version Management . . . . .	71
10.4 Lock-Step Mode . . . . .	71
11. Monitoring & Health Checks . . . . .	71
11.1 Health Dashboard . . . . .	71
11.2 Health Check Types . . . . .	72
11.3 Alerts . . . . .	72
12. Security Features . . . . .	72
12.1 Credential Security . . . . .	72
12.2 Deployment Locks . . . . .	72
12.3 Audit Logging . . . . .	73
12.4 Secret Detection . . . . .	73
13. Troubleshooting . . . . .	73
13.1 Common Issues . . . . .	73
13.2 Log Locations . . . . .	73
13.3 Getting Help . . . . .	74
14. Reference . . . . .	74
14.1 Keyboard Shortcuts . . . . .	74
14.2 CLI Commands . . . . .	74
14.3 Environment Variables . . . . .	74
14.4 File Locations . . . . .	74
Appendix A: IAM Policy Requirements . . . . .	75
Appendix B: Glossary . . . . .	75
<b>RADIANT Deployer Architecture &amp; Deployment Packages</b>	<b>76</b>
Overview . . . . .	76
Deployment Modes . . . . .	76
Mode Definitions . . . . .	76
INSTALL Mode (Fresh Installation) . . . . .	77
UPDATE Mode (Upgrade Existing) . . . . .	77
ROLLBACK Mode (Revert to Previous) . . . . .	77
Deployment Package Structure . . . . .	77
Package Manifest . . . . .	78
Package Storage Locations . . . . .	79
Key Implementation Files . . . . .	79
Swift Deployer . . . . .	79
Build Tools . . . . .	79
Data Flow Diagrams . . . . .	80
Install Flow . . . . .	80
Update Flow . . . . .	80
Rollback Flow . . . . .	81
Verification Checklist . . . . .	81
Deployment Modes . . . . .	81
AI Registry Seeding . . . . .	81
Deployment Packages . . . . .	81
Parameter Rules . . . . .	81
Related Documentation . . . . .	81
<b>Think Tank AI - User Guide</b>	<b>82</b>
Welcome to Think Tank . . . . .	82
Table of Contents . . . . .	82
1. Getting Started . . . . .	82
1.1 Creating Your Account . . . . .	82
1.2 Choosing a Plan . . . . .	83

1.3 Your First Chat . . . . .	83
2. Your Dashboard . . . . .	83
2.1 Interface Overview . . . . .	83
2.2 Sidebar Navigation . . . . .	83
2.3 Quick Actions . . . . .	84
3. Chatting with AI . . . . .	84
3.1 Sending Messages . . . . .	84
3.2 Message Actions . . . . .	84
3.3 Streaming Responses . . . . .	84
3.4 Multi-Turn Conversations . . . . .	84
3.5 Code in Chats . . . . .	85
4. Choosing Models . . . . .	85
4.1 Model Selection . . . . .	85
4.2 Model Categories . . . . .	85
4.3 Auto Mode . . . . .	86
4.4 Model Comparison . . . . .	86
5. Focus Modes & Personas . . . . .	87
5.1 Focus Modes . . . . .	87
5.2 Custom Personas . . . . .	87
5.3 Sharing Personas . . . . .	88
6. Canvas & Artifacts . . . . .	88
6.1 What is Canvas? . . . . .	88
6.2 Creating Artifacts . . . . .	88
6.3 Artifact Actions . . . . .	88
6.4 Full Canvas Mode . . . . .	88
7. Collaboration Features . . . . .	89
7.1 Sharing Chats . . . . .	89
7.2 Real-Time Collaboration . . . . .	89
7.3 Team Workspaces . . . . .	89
7.4 Comments & Annotations . . . . .	89
8. Managing Your Account . . . . .	89
8.1 Profile Settings . . . . .	89
8.2 Preferences . . . . .	90
8.3 Data & Privacy . . . . .	90
8.4 Connected Apps . . . . .	90
9. Credits & Billing . . . . .	90
9.1 Understanding Credits . . . . .	90
9.2 Credit Usage . . . . .	91
9.3 Viewing Usage . . . . .	91
9.4 Purchasing Credits . . . . .	91
9.5 Subscription Management . . . . .	91
10. Tips & Best Practices . . . . .	92
10.1 Writing Better Prompts . . . . .	92
10.2 Getting Better Results . . . . .	92
10.3 Saving Credits . . . . .	92
10.4 Organizing Chats . . . . .	92
11. Keyboard Shortcuts . . . . .	92
11.1 General . . . . .	92
11.2 Chat . . . . .	93
11.3 Navigation . . . . .	93
12. FAQ . . . . .	93
Getting Started . . . . .	93
Credits & Billing . . . . .	93
Privacy & Security . . . . .	94

Troubleshooting . . . . .	94
Need Help? . . . . .	94
Version History . . . . .	94
13. Delight System . . . . .	94
13.1 What is Delight? . . . . .	94
13.2 Personality Modes . . . . .	95
13.3 Achievements . . . . .	95
13.4 Easter Eggs . . . . .	95
13.5 Sound Effects . . . . .	95
13.6 Customizing Delight . . . . .	95
<b>RADIANT API Reference . . . . .</b>	<b>96</b>
Chat Completions . . . . .	96
Create Chat Completion . . . . .	96
Models . . . . .	97
List Models . . . . .	97
Get Model . . . . .	98
Embeddings . . . . .	98
Create Embeddings . . . . .	98
Billing . . . . .	99
Get Credit Balance . . . . .	99
Get Usage . . . . .	99
Webhooks . . . . .	99
List Webhooks . . . . .	99
Create Webhook . . . . .	99
Test Webhook . . . . .	100
Batch Processing . . . . .	100
Create Batch Job . . . . .	100
Get Batch Job . . . . .	100
List Batch Jobs . . . . .	100
Error Codes . . . . .	100
Error Categories . . . . .	101
Common Error Codes . . . . .	101
Rate Limits . . . . .	101
SDKs . . . . .	102
TypeScript/JavaScript . . . . .	102
Python . . . . .	102
CLI . . . . .	102
Changelog . . . . .	102
Support . . . . .	102
<b>RADIANT API Versioning Guide . . . . .</b>	<b>103</b>
Overview . . . . .	103
Versioning Strategy . . . . .	103
URL Path Versioning . . . . .	103
Version Lifecycle . . . . .	103
Support Policy . . . . .	103
Breaking vs Non-Breaking Changes . . . . .	103
Non-Breaking Changes (No Version Bump) . . . . .	103
Breaking Changes (Require New Version) . . . . .	104
Version Header Support . . . . .	104
Request Headers . . . . .	104
Response Headers . . . . .	104
Deprecation Process . . . . .	104

Timeline . . . . .	104
Deprecation Headers . . . . .	104
Deprecation Warnings . . . . .	105
Migration Guide Template . . . . .	105
v1 to v2 Migration . . . . .	105
Feature Flags . . . . .	105
Beta Features . . . . .	105
Graduated Features . . . . .	106
SDK Versioning . . . . .	106
SDK Version Matrix . . . . .	106
SDK Version Headers . . . . .	106
OpenAPI Specification . . . . .	106
Versioned Specs . . . . .	106
Schema Versioning . . . . .	107
Testing Versions . . . . .	107
Version Compatibility Tests . . . . .	107
Client Communication . . . . .	107
Changelog . . . . .	107
Email Notifications . . . . .	107
Best Practices . . . . .	108
For API Developers . . . . .	108
For API Consumers . . . . .	108
Contact . . . . .	108
<b>RADIANT Error Codes Reference</b>	<b>109</b>
Overview . . . . .	109
Error Code Format . . . . .	109
Authentication Errors (1xxx) . . . . .	109
Authorization Errors (2xxx) . . . . .	110
Validation Errors (3xxx) . . . . .	110
Resource Errors (4xxx) . . . . .	110
Rate Limiting Errors (5xxx) . . . . .	111
AI/Model Errors (6xxx) . . . . .	111
Billing Errors (7xxx) . . . . .	111
Storage Errors (8xxx) . . . . .	112
Internal Errors (9xxx) . . . . .	112
Usage in Code . . . . .	112
TypeScript/JavaScript . . . . .	112
Response Format . . . . .	113
Client Handling . . . . .	113
Retry Logic . . . . .	113
Error Display . . . . .	113
Adding New Error Codes . . . . .	114
See Also . . . . .	114
<b>RADIANT Testing Guide</b>	<b>115</b>
Overview . . . . .	115
Quick Start . . . . .	115
Unit Testing . . . . .	115
Lambda Handler Tests . . . . .	115
Shared Module Tests . . . . .	117
E2E Testing (Admin Dashboard) . . . . .	117
Setup . . . . .	117
Test Structure . . . . .	118

Writing E2E Tests . . . . .	118
Swift Testing . . . . .	118
Test Structure . . . . .	118
Running Swift Tests . . . . .	119
Writing Swift Tests . . . . .	119
Test Utilities . . . . .	120
Mock Factories . . . . .	120
Assertion Helpers . . . . .	120
Mocking Guidelines . . . . .	120
Database Mocking . . . . .	120
AWS SDK Mocking . . . . .	121
External API Mocking . . . . .	121
CI/CD Integration . . . . .	121
Coverage Requirements . . . . .	121
Best Practices . . . . .	122
Do's . . . . .	122
Don'ts . . . . .	122
Test Naming Convention . . . . .	122
Debugging Tests . . . . .	122
Vitest . . . . .	122
Playwright . . . . .	122
Swift . . . . .	123
See Also . . . . .	123
<b>RADIANT Compliance Guide</b>	<b>124</b>
Overview . . . . .	124
Compliance Matrix . . . . .	124
SOC 2 Controls . . . . .	124
Trust Service Criteria . . . . .	124
Evidence Collection . . . . .	125
Annual Audit Checklist . . . . .	125
HIPAA Compliance . . . . .	125
Applicability . . . . .	125
Technical Safeguards . . . . .	125
Administrative Safeguards . . . . .	126
Physical Safeguards . . . . .	126
PHI Data Handling . . . . .	126
BAA Requirements . . . . .	126
GDPR Compliance . . . . .	126
Data Subject Rights . . . . .	126
Data Export (Right to Access) . . . . .	126
Data Deletion (Right to Erasure) . . . . .	127
Data Processing Agreement . . . . .	127
Data Residency . . . . .	127
Consent Management . . . . .	128
Data Classification . . . . .	128
Classification Levels . . . . .	128
PII Fields . . . . .	128
Encryption . . . . .	129
At Rest . . . . .	129
In Transit . . . . .	129
Key Rotation . . . . .	129
Audit Logging . . . . .	129
What We Log . . . . .	129

Log Format . . . . .	129
Log Protection . . . . .	130
Incident Response . . . . .	130
Classification . . . . .	130
Breach Notification . . . . .	130
Vendor Management . . . . .	130
Approved Sub-Processors . . . . .	130
Vendor Assessment . . . . .	130
Contact . . . . .	131
<b>RADIANT Cost Optimization Guide</b>	<b>132</b>
Overview . . . . .	132
Current Architecture Costs . . . . .	132
Estimated Monthly Costs by Tier . . . . .	132
Cost Breakdown by Service . . . . .	132
Optimization Strategies . . . . .	132
1. Database Optimization . . . . .	132
2. Lambda Optimization . . . . .	133
3. S3 Optimization . . . . .	134
4. API Gateway Optimization . . . . .	134
5. CloudWatch Optimization . . . . .	135
6. ElastiCache Optimization . . . . .	135
7. Data Transfer Optimization . . . . .	136
Cost Monitoring . . . . .	136
AWS Cost Explorer . . . . .	136
CloudWatch Billing Alerts . . . . .	136
Cost Allocation Tags . . . . .	137
Environment-Specific Recommendations . . . . .	137
Development . . . . .	137
Staging . . . . .	137
Production . . . . .	137
Monthly Cost Review Checklist . . . . .	137
Tools . . . . .	137
<b>RADIANT Data Retention Policy</b>	<b>139</b>
Overview . . . . .	139
Retention Schedule . . . . .	139
User Data . . . . .	139
System Data . . . . .	139
Billing Data . . . . .	139
Implementation . . . . .	140
Database Retention . . . . .	140
S3 Lifecycle Policies . . . . .	140
CloudWatch Log Retention . . . . .	141
Data Deletion . . . . .	141
User-Initiated Deletion . . . . .	141
Administrative Deletion . . . . .	142
Tenant Offboarding . . . . .	143
Legal Holds . . . . .	144
Implementing a Legal Hold . . . . .	144
Suspending Retention Policies . . . . .	144
Audit Trail . . . . .	145
Retention Actions Log . . . . .	145
Compliance Reporting . . . . .	145

Verification . . . . .	145
Monthly Retention Audit . . . . .	145
Compliance Queries . . . . .	146
Contact . . . . .	146
<b>RADIANT Disaster Recovery Guide</b>	<b>147</b>
Overview . . . . .	147
Recovery Objectives . . . . .	147
Backup Strategy . . . . .	147
Database Backups . . . . .	147
S3 Backups . . . . .	148
Secrets Backup . . . . .	148
Failure Scenarios . . . . .	148
Scenario 1: Single AZ Failure . . . . .	148
Scenario 2: Database Corruption . . . . .	148
Scenario 3: Region Failure . . . . .	149
Scenario 4: Accidental Deletion . . . . .	149
Scenario 5: Security Breach . . . . .	150
Recovery Procedures . . . . .	150
Database Recovery Runbook . . . . .	150
Full Service Recovery Runbook . . . . .	151
Testing DR Procedures . . . . .	152
Quarterly DR Drill . . . . .	152
DR Test Checklist . . . . .	152
Communication Plan . . . . .	152
Escalation Matrix . . . . .	152
Status Page Updates . . . . .	152
Infrastructure as Code . . . . .	153
Contacts . . . . .	153
<b>RADIANT Performance Guide</b>	<b>154</b>
Overview . . . . .	154
Architecture Performance . . . . .	154
Request Flow . . . . .	154
Latency Targets . . . . .	154
Caching Strategy . . . . .	154
Multi-Layer Caching . . . . .	154
Cache Keys . . . . .	155
Cache Invalidation . . . . .	155
Database Optimization . . . . .	155
Connection Pooling . . . . .	155
Query Optimization . . . . .	156
RLS Performance . . . . .	156
Lambda Optimization . . . . .	156
Cold Start Reduction . . . . .	156
Memory Optimization . . . . .	156
Bundling . . . . .	156
Rate Limiting . . . . .	157
Tier Limits . . . . .	157
Implementation . . . . .	157
Load Testing . . . . .	157
Running Tests . . . . .	157
Performance Baselines . . . . .	157
Scaling . . . . .	158

Horizontal Scaling . . . . .	158
Vertical Scaling . . . . .	158
Monitoring . . . . .	158
Key Metrics . . . . .	158
Alerting Thresholds . . . . .	159
Best Practices . . . . .	159
Do's . . . . .	159
Don'ts . . . . .	159
Troubleshooting . . . . .	159
High Latency . . . . .	159
High Error Rate . . . . .	159
Scaling Issues . . . . .	159
<b>RADIANT AI Registry Seed Data System</b>	<b>160</b>
Overview . . . . .	160
Architecture . . . . .	160
Critical Rules . . . . .	161
Rule 1: NO HARDCODING IN DEPLOYER APP . . . . .	161
Rule 2: INSTALLER SEEDS, UPDATER PRESERVES . . . . .	161
Rule 3: ADMIN CONTROLS ALL . . . . .	161
Seed Data Structure . . . . .	161
manifest.json . . . . .	161
providers.json . . . . .	162
external-models.json . . . . .	162
self-hosted-models.json . . . . .	162
Building Packages with Seed Data . . . . .	163
List Available Seed Versions . . . . .	163
Build with Specific Seed Version . . . . .	163
Package Manifest with Seed Data . . . . .	163
Seed Data Categories . . . . .	163
External Providers (21) . . . . .	163
External Models (50+) . . . . .	164
Self-Hosted Models (38) . . . . .	164
Pricing Structure . . . . .	164
External Providers . . . . .	164
Self-Hosted Models . . . . .	164
Creating New Seed Versions . . . . .	165
1. Create Version Directory . . . . .	165
2. Create Required Files . . . . .	165
3. Update Registry . . . . .	165
4. Test Build . . . . .	165
Database Seeding . . . . .	165
Swift Service API . . . . .	166
SeedDataService . . . . .	166
Usage in DeploymentService . . . . .	166
Related Documentation . . . . .	166
<b>RADIANT v4.17.0 - Troubleshooting Guide</b>	<b>167</b>
Common Issues and Solutions . . . . .	167
CDK Deployment Failures . . . . .	167
Aurora Database Issues . . . . .	167
Lambda Function Errors . . . . .	167
LitellM / ECS Issues . . . . .	168
SageMaker Issues (Tier 3+) . . . . .	168

Cognito Authentication Issues . . . . .	168
Admin Dashboard Issues . . . . .	169
Log Locations . . . . .	169
Viewing Logs . . . . .	169
Health Check Endpoints . . . . .	170
Emergency Procedures . . . . .	170
Database Restore from Snapshot . . . . .	170
Point-in-Time Recovery . . . . .	170
Rollback CDK Deployment . . . . .	170
Disable Problematic Model . . . . .	171
Force Scale Down SageMaker . . . . .	171
Performance Benchmarks . . . . .	171
Support Checklist . . . . .	171
Useful AWS CLI Commands . . . . .	171
<b>RADIANT CDK Stack Dependencies</b>	<b>173</b>
Stack Dependency Graph . . . . .	173
Stack Definitions . . . . .	174
Layer 1: Foundation . . . . .	174
Layer 2: Core Infrastructure . . . . .	174
Layer 3: Authentication . . . . .	174
Layer 4: Application Services . . . . .	174
Layer 5: AI Services . . . . .	175
Layer 6: Specialized Services . . . . .	175
CDK Implementation . . . . .	175
Explicit Dependencies . . . . .	175
Cross-Stack References . . . . .	176
Deployment Order . . . . .	176
Fresh Install . . . . .	176
Update (with dependencies) . . . . .	177
Selective Deployment . . . . .	177
Rollback Considerations . . . . .	177
Validation Script . . . . .	177
Related Documentation . . . . .	178

# RADIANT v4.17.0 - Deployed System Guide

**Operations and infrastructure documentation for the deployed RADIANT platform**

Last Updated: {{BUILD\_DATE}} Version: {{RADIANT\_VERSION}}

---

## Table of Contents

1. System Overview
  2. Infrastructure Components
  3. Deployment Architecture
  4. Service Endpoints
  5. Database Operations
  6. Lambda Functions
  7. AI Provider Integration
  8. Monitoring & Observability
  9. Security Configuration
  10. Backup & Recovery
  11. Maintenance Procedures
  12. Troubleshooting
- 

## 1. System Overview

### 1.1 Platform Summary

RADIANT is a multi-tenant AWS SaaS platform deployed across multiple AWS services:

Component	AWS Service	Purpose
Compute	Lambda	Serverless API handlers
Database	Aurora PostgreSQL	Primary data store
Cache	ElasticCache Redis	Session & response caching
Storage	S3	File storage, artifacts
Auth	Cognito	User authentication
API	API Gateway	REST & WebSocket APIs
AI	SageMaker, Bedrock	Model hosting
CDN	CloudFront	Static asset delivery
DNS	Route 53	Domain management

## 1.2 Environment Tiers

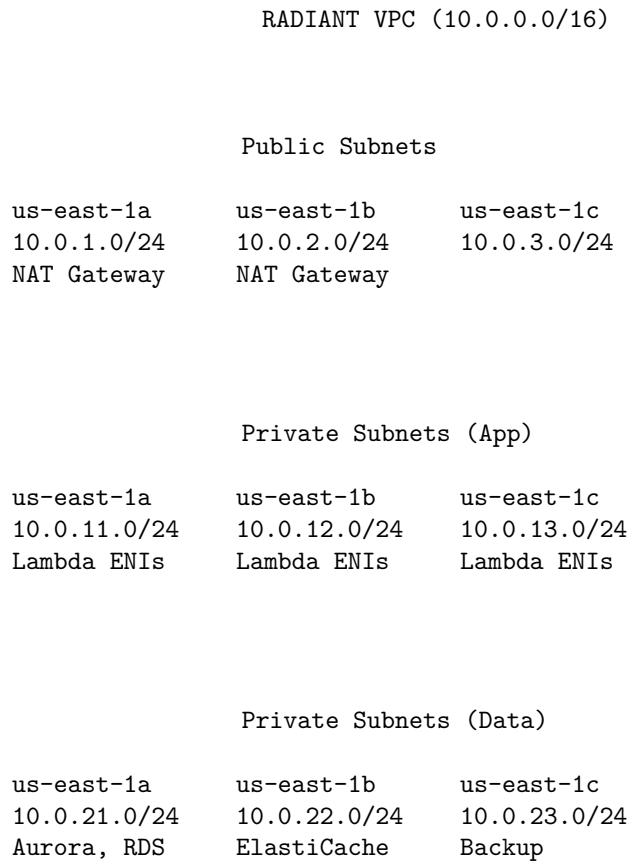
Environment	Purpose	Domain
<b>Production</b>	Live customer traffic	api.{RADIANT_DOMAIN}
<b>Staging</b>	Pre-release testing	staging-api.{RADIANT_DOMAIN}
<b>Development</b>	Development & testing	dev-api.{RADIANT_DOMAIN}

## 1.3 Region Configuration

Region	Role	Services
us-east-1	Primary	All services
us-west-2	DR Failover	Database replica, S3 replication
eu-west-1	EU Data Residency	Optional for GDPR compliance

## 2. Infrastructure Components

### 2.1 VPC Architecture



## 2.2 Security Groups

Security Group	Inbound Rules	Purpose
sg-lambda	None (outbound only)	Lambda functions
sg-aurora	5432 from sg-lambda	Database access
sg-redis	6379 from sg-lambda	Cache access
sg-alb	443 from 0.0.0.0/0	Load balancer

## 2.3 IAM Roles

Role	Attached Policies	Used By
radiant-lambda-role	AWSLambdaVPCAccess, SecretsManager, S3, SageMaker	Lambda functions
radiant-api-gw-role	CloudWatchLogs, Lambda invoke	API Gateway
radiant-sagemaker-role	SageMakerFullAccess, S3	SageMaker endpoints

## 3. Deployment Architecture

### 3.1 CDK Stacks

Stack	Resources	Dependencies
NetworkingStack	VPC, Subnets, NAT	None
SecurityStack	IAM, KMS, Security Groups	Networking
DataStack	Aurora, ElastiCache, S3	Security
AuthStack	Cognito User Pool	Data
ApiStack	API Gateway, Lambda	Auth, Data
AiStack	SageMaker, Bedrock config	Api
MonitoringStack	CloudWatch, Alarms	All
AdminStack	Admin Dashboard hosting	Api

### 3.2 Deployment Commands

```
# Deploy all stacks to development
npm run deploy:dev

# Deploy to staging with approval
npm run deploy:staging

# Deploy to production (requires approval)
npm run deploy:prod

# Deploy specific stack
cd packages/infrastructure
cdk deploy RadiantApiStack --context environment=prod
```

### 3.3 Environment Variables

Variable	Description	Source
AURORA_CLUSTER_ARN	Database cluster ARN	CDK output
AURORA_SECRET_ARN	Database credentials	Secrets Manager
COGNITO_USER_POOL_ID	Auth pool ID	CDK output
REDIS_URL	Cache connection string	CDK output
S3_BUCKET	Storage bucket name	CDK output

## 4. Service Endpoints

### 4.1 API Gateway Endpoints

Endpoint	Method	Description
/v1/chat	POST	Send chat message
/v1/chat/stream	WebSocket	Streaming responses
/v1/models	GET	List available models
/v1/usage	GET	Usage statistics
/admin/v1/*	Various	Admin API

### 4.2 Health Check Endpoints

Endpoint	Expected Response	Checks
/health	{"status": "healthy"}	Lambda running
/health/db	{"status": "connected"}	Aurora connection
/health/cache	{"status": "connected"}	Redis connection
/health/providers	Provider status array	AI provider APIs

### 4.3 Internal Endpoints

Service	Endpoint	Port
Aurora PostgreSQL	radiant-cluster.xxx.us-east-1.rds.amazonaws.com	5432
ElastiCache Redis	radiant-cache.xxx.cache.amazonaws.com	6379
SageMaker Runtime	Via AWS SDK	-

## 5. Database Operations

### 5.1 Aurora Configuration

Setting	Value
Engine	PostgreSQL 15.4
Instance Type	Serverless v2
Min ACUs	2 (dev), 8 (prod)
Max ACUs	16 (dev), 128 (prod)

Setting	Value
Storage	Auto-scaling to 128 TB
Encryption	AES-256 (KMS)

## 5.2 Connection Management

```
-- Check active connections
SELECT count(*) FROM pg_stat_activity;

-- View connection by application
SELECT application_name, count(*)
FROM pg_stat_activity
GROUP BY application_name;

-- Terminate idle connections
SELECT pg_terminate_backend(pid)
FROM pg_stat_activity
WHERE state = 'idle'
AND query_start < now() - interval '10 minutes';
```

## 5.3 Row-Level Security

All tenant data is protected by RLS:

```
-- Verify RLS is enabled
SELECT tablename, rowsecurity
FROM pg_tables
WHERE schemaname = 'public' AND rowsecurity = true;

-- Set tenant context (done automatically by Lambda)
SET app.current_tenant_id = 'tenant-uuid';

-- Queries automatically filter by tenant
SELECT * FROM users; -- Only returns current tenant's users
```

## 5.4 Database Migrations

```
# Run pending migrations
cd packages/infrastructure
npm run migrate:up

# Rollback last migration
npm run migrate:down

# View migration status
npm run migrate:status
```

---

# 6. Lambda Functions

## 6.1 Function Inventory

Function	Memory	Timeout	Trigger
radiant-api	1024 MB	30s	API Gateway
radiant-brain	2048 MB	60s	API Gateway
radiant-webhooks	512 MB	30s	EventBridge
radiant-scheduler	512 MB	300s	EventBridge (cron)
radiant-batch	3008 MB	900s	SQS

## 6.2 Monitoring Lambda

```
# View recent invocations
aws lambda get-function \
--function-name radiant-api \
--query 'Configuration.[FunctionName,MemorySize,Timeout,LastModified]'

# Check concurrent executions
aws cloudwatch get-metric-statistics \
--namespace AWS/Lambda \
--metric-name ConcurrentExecutions \
--dimensions Name=FunctionName,Value=radiant-api \
--start-time $(date -u -d '1 hour ago' +%Y-%m-%dT%H:%M:%SZ) \
--end-time $(date -u +%Y-%m-%dT%H:%M:%SZ) \
--period 300 \
--statistics Maximum
```

## 6.3 Cold Start Optimization

Strategy	Implementation
Provisioned Concurrency	10 instances for <code>radiant-api</code>
Keep-Warm	CloudWatch scheduled ping every 5 min
Connection Pooling	RDS Proxy for database
Lazy Loading	Defer non-critical imports

# 7. AI Provider Integration

## 7.1 External Providers

Provider	Models	Authentication
OpenAI	GPT-4, GPT-3.5	API Key
Anthropic	Claude 3, Claude 2	API Key
Google	Gemini Pro, PaLM 2	Service Account
Cohere	Command, Embed	API Key
Mistral	Mistral Large, Medium	API Key

## 7.2 Self-Hosted Models (SageMaker)

Model	Instance Type	Thermal Default
LLaMA-2-70B	ml.g5.48xlarge	COLD
Stable Diffusion XL	ml.g5.2xlarge	COLD
Whisper Large	ml.g4dn.xlarge	COLD
CodeLlama-34B	ml.g5.12xlarge	COLD

### 7.3 Model Thermal States

```
# Check endpoint status
aws sagemaker describe-endpoint \
--endpoint-name radiant-llama-70b

# Scale endpoint to warm
aws sagemaker update-endpoint-weights-and-capacities \
--endpoint-name radiant-llama-70b \
--desired-weights-and-capacities '[{"VariantName": "AllTraffic", "DesiredInstanceCount": 1}]'

# Scale to zero (cold)
aws sagemaker update-endpoint-weights-and-capacities \
--endpoint-name radiant-llama-70b \
--desired-weights-and-capacities '[{"VariantName": "AllTraffic", "DesiredInstanceCount": 0}]'
```

---

## 8. Monitoring & Observability

### 8.1 CloudWatch Dashboards

Dashboard	Metrics
Platform Overview	Request rate, latency, errors
Database Performance	Connections, ACUs, query time
AI Provider Health	Success rate, latency by provider
Billing & Usage	Credits consumed, revenue

### 8.2 Key Metrics

Metric	Warning	Critical
API Error Rate	> 1%	> 5%
P99 Latency	> 5s	> 15s
Aurora ACU Usage	> 80%	> 95%
Lambda Errors	> 10/min	> 50/min

### 8.3 Alerting

```
# List configured alarms
aws cloudwatch describe-alarms \
--alarm-name-prefix radiant

# View alarm state
aws cloudwatch describe-alarm-history \
--alarm-name radiant-high-error-rate
```

## 8.4 Log Groups

Log Group	Retention	Content
/aws/lambda/radiant-api	30 days	API request logs
/aws/lambda/radiant-brain	30 days	Model routing logs
/radiant/audit	365 days	Security audit logs
/radiant/billing	2 years	Billing events

---

## 9. Security Configuration

### 9.1 Encryption

Data	Encryption
Database at rest	AES-256 (KMS CMK)
S3 objects	AES-256 (SSE-S3)
Secrets	KMS envelope encryption
API traffic	TLS 1.3

### 9.2 Secrets Management

```
# List RADIANT secrets
aws secretsmanager list-secrets \
  --filters Key=name,Values=radiant

# Rotate database credentials
aws secretsmanager rotate-secret \
  --secret-id radiant/aurora-credentials

# View secret metadata
aws secretsmanager describe-secret \
  --secret-id radiant/openai-api-key
```

### 9.3 WAF Rules

Rule	Action	Purpose
Rate Limiting	Block	> 1000 req/min per IP
SQL Injection	Block	SQLi pattern detection
XSS	Block	Cross-site scripting
Geo Blocking	Block	Sanctioned countries

---

## 10. Backup & Recovery

### 10.1 Automated Backups

Resource	Frequency	Retention
Aurora Snapshots	Daily	35 days
S3 Objects	Continuous	Versioned
DynamoDB	Point-in-time	35 days
Secrets	Auto-versioned	30 versions

## 10.2 Recovery Procedures

### Database Point-in-Time Recovery:

```
aws rds restore-db-cluster-to-point-in-time \
--source-db-cluster-identifier radiant-cluster \
--db-cluster-identifier radiant-cluster-restored \
--restore-to-time 2024-01-15T10:00:00Z
```

### S3 Object Recovery:

```
# List object versions
aws s3api list-object-versions \
--bucket radiant-storage \
--prefix tenant/123

# Restore specific version
aws s3api copy-object \
--bucket radiant-storage \
--copy-source radiant-storage/file.txt?versionId=xxx \
--key file.txt
```

## 10.3 Disaster Recovery

RPO	RTO	Strategy
1 hour	4 hours	Cross-region Aurora replica, S3 CRR

# 11. Maintenance Procedures

## 11.1 Routine Maintenance

Task	Frequency	Procedure
Security patches	Weekly	Automated via CDK
Log rotation	Daily	Automated
Database vacuum	Weekly	Automated
Certificate renewal	60 days	ACM auto-renewal

## 11.2 Scaling Operations

```
# Scale Aurora
aws rds modify-db-cluster \
--db-cluster-identifier radiant-cluster \
--serverless-v2-scaling-configuration MinCapacity=16,MaxCapacity=64
```

```
# Update Lambda memory
aws lambda update-function-configuration \
--function-name radiant-api \
--memory-size 2048
```

## 11.3 Deployment Checklist

- Run tests in staging
  - Review CloudWatch for anomalies
  - Announce maintenance window
  - Deploy with --require-approval
  - Verify health checks
  - Monitor error rates for 30 min
  - Update status page
- 

# 12. Troubleshooting

## 12.1 Common Issues

High Latency:

```
# Check database performance
aws cloudwatch get-metric-statistics \
--namespace AWS/RDS \
--metric-name DatabaseConnections \
--dimensions Name=DBClusterIdentifier,Value=radiant-cluster

# Check Lambda duration
aws logs filter-log-events \
--log-group-name /aws/lambda/radiant-api \
--filter-pattern "REPORT Duration"
```

Connection Errors:

```
# Verify security group rules
aws ec2 describe-security-groups --group-ids sg-xxx

# Check VPC endpoints
aws ec2 describe-vpc-endpoints \
--filters Name=vpcl-id,Values=vpcl-xxx
```

## 12.2 Emergency Procedures

Rollback Deployment:

```
# List recent deployments
aws cloudformation list-stacks \
--stack-status-filter UPDATE_COMPLETE

# Rollback to previous version
cdk deploy --rollback
```

Disable Problematic Feature:

```
# Update feature flag
aws ssm put-parameter \
--name /radiant/features/new-feature \
--value "false" \
--overwrite
```

---

## Version History

Version	Date	Changes
4.17.0	<code>{{BUILD_DATE}}</code>	Initial system guide

---

*This documentation is automatically generated as part of the RADIANT build process.*

# RADIANT Architecture v4.17.0

## Overview

RADIANT is a multi-tenant AWS SaaS platform for AI model access and orchestration. It consists of:

1. **Swift Deployer App** - macOS GUI for infrastructure deployment
2. **AWS Infrastructure** - CDK-based cloud infrastructure
3. **Admin Dashboard** - Next.js admin UI (Phase 3)

## Technology Stack

Layer	Technology
Swift App	SwiftUI, macOS 13.0+, Swift 5.9+, GRDB
Infrastructure	AWS CDK (TypeScript), Aurora PostgreSQL, Lambda
Dashboard	Next.js 14, TypeScript, Tailwind CSS
AI Integration	106+ models, LiteLLM, SageMaker

## Monorepo Structure

```
radiant/
  packages/
    shared/          # @radiant/shared - Types & constants
    infrastructure/ # @radiant/infrastructure - CDK stacks
  apps/
    swift-deployer/ # RadiantDeployer macOS app
    functions/       # Lambda functions (Phase 2)
    migrations/     # Database migrations (Phase 2)
    docs/           # Specifications
```

## Phase 1 Architecture

### @radiant/shared Package

Single source of truth for all types and constants:

```
// Types
- app.types.ts      // ManagedApp, DeploymentStatus, etc.
- environment.types.ts // Environment, TierConfig, etc.
- ai.types.ts       // AIPublisher, AIModel, etc.
- admin.types.ts    // Administrator, Permissions
```

```

- billing.types.ts      // UsageEvent, Invoice, etc.
- compliance.types.ts // PHI, AuditLog, etc.

// Constants
- version.ts           // RADIANT_VERSION = "4.17.0"
- tiers.ts             // Tier 1-5 configurations
- regions.ts           // AWS region configs
- providers.ts         // AI provider definitions

```

## Swift Deployer App

```

RadiantDeployer/
    RadiantDeployerApp.swift # @main entry point
    AppState.swift          # Global state management
    Models/
        ManagedApp.swift    # App configuration model
        Credentials.swift   # AWS credentials model
        Deployment.swift    # Deployment state/progress
    Services/
        CredentialService.swift # Keychain credential storage
        CDKService.swift       # CDK command execution
        AWSService.swift       # AWS API interactions
    Views/
        MainView.swift        # NavigationSplitView container
        AppsView.swift        # Application grid
        DeployView.swift      # Deployment wizard
        ProvidersView.swift   # AI provider list
        ModelsView.swift      # AI model catalog
        SettingsView.swift    # Preferences

```

## CDK Infrastructure Stacks

```

packages/infrastructure/
    bin/radiant.ts          # CDK app entry point
    lib/stacks/
        foundation-stack.ts  # KMS, SSM parameters
        networking-stack.ts  # VPC, subnets, endpoints
        security-stack.ts    # Security groups, WAF

```

## Infrastructure Tiers

Tier	Name	VPC CIDR	AZs	NAT	Aurora ACU	Features
1	SEED	/20	2	1	0.5-2	Dev only
2	STARTUP	/18	2	1	1-8	WAF, GuardDuty
3	GROWTH	/17	3	2	2-16	Self-hosted models, HIPAA
4	SCALE	/16	3	3	4-64	Multi-region, Global DB
5	ENTERPRISE	/14	3	3	8-128	Full enterprise

## Deployment Flow

Swift Deployer

(macOS)

1. Configure credentials
2. Select app/tier/env
3. Initiate deployment

CDK Deploy  
(via Process)

4. Bootstrap AWS
5. Synth stacks
6. Deploy in order

AWS Account  
(VPC, RDS,  
Lambda, etc)

## Phase 1 Deliverables

- Monorepo structure (pnpm workspace)
- @radiant/shared package with all types
- @radiant/infrastructure base CDK stacks
- RadiantDeployer Swift app with full UI
- Credential management via Keychain
- CDK service for deployment execution
- Documentation

## Future Phases

Phase	Sections	Description
2	3-7	AI stacks, Lambdas, Database
3	8-9	Admin Dashboard, Deployment Guide
4	10-17	Visual AI, Brain, Analytics
5	18-28	Think Tank, Collaboration
6	29-35	Registry, Time Machine
7	36-39	Neural Engine, Workflows
8	40-42	Isolation, i18n, Config
9	43-46	Billing System

# RADIANT v4.18.0 - Deployment Guide

## Overview

This guide covers deploying the RADIANT platform from development to production. The platform consists of:

1. **AWS Infrastructure** - CDK stacks for all cloud resources
2. **Admin Dashboard** - Next.js admin interface
3. **Swift Deployer App** - macOS deployment tool with AI assistant

## What's New in v4.18.0

- **Unified Package System** - Deploy with atomic component versioning
- **AI Assistant** - Claude-powered deployment guidance in Swift app
- **Configurable Timeouts** - SSM-synced operation timeouts
- **Cost Management** - Real-time cost tracking and alerts
- **Compliance Reports** - SOC2, HIPAA, GDPR, ISO27001 reporting

## Prerequisites

### AWS Account Setup

Requirement	Action	Verification
AWS Account	Create or use existing	Account ID available
IAM User	Create with AdministratorAccess	Access keys configured
AWS CLI	Install and configure	<code>aws sts get-caller-identity</code>
Route 53 Domain (optional)	Register domain	Domain in hosted zone
ACM Certificate (optional)	Request in us-east-1	Certificate validated

### Development Environment

Requirement	Version	Verification
Node.js	20.x LTS	<code>node --version</code>
pnpm	8.x+	<code>pnpm --version</code>
AWS CDK CLI	2.x	<code>cdk --version</code>
Xcode	15.x+	<code>xcode-select -p</code>

Requirement	Version	Verification
Swift	5.9+	<code>swift --version</code>

## Quick Start

### Automated Deployment

Use the deployment script for a streamlined deployment:

```
# Deploy to dev environment
./scripts/deploy.sh --environment dev

# Deploy to staging
./scripts/deploy.sh --environment staging

# Deploy to production
./scripts/deploy.sh --environment prod
```

**Note:** RADIANT uses a unified deployment model. All features are available in every deployment. Licensing restrictions are handled at the application level, not infrastructure level.

### Verify Deployment

After deployment, verify all resources:

```
./scripts/verify-deployment.sh --environment dev
```

---

## Manual Deployment

### 1. Install Dependencies

```
cd radiant
npx pnpm install
```

### 2. Build Shared Package

```
cd packages/shared
npm run build
```

### 3. Build Lambda Functions

```
cd packages/infrastructure/lambda
npm install --legacy-peer-deps
npm run build
```

### 4. Build Admin Dashboard

```
cd apps/admin-dashboard
npm install
npm run build
```

## CDK Deployment

### Bootstrap CDK (One-time per account/region)

```
cd packages/infrastructure
npx cdk bootstrap aws://ACCOUNT_ID/us-east-1 --qualifier radiant
```

### Deploy All Stacks

```
# Deploy in order with dependencies
npx cdk deploy --all \
  --context environment=dev \
  --context tier=1 \
  --require-approval never
```

### Deploy Individual Stacks

```
# Phase 1: Foundation
npx cdk deploy Radiant-dev-Foundation --context environment=dev --context tier=1
npx cdk deploy Radiant-dev-Networking --context environment=dev --context tier=1

# Phase 2: Security & Data
npx cdk deploy Radiant-dev-Security --context environment=dev --context tier=1
npx cdk deploy Radiant-dev-Data --context environment=dev --context tier=1
npx cdk deploy Radiant-dev-Storage --context environment=dev --context tier=1

# Phase 3: Auth & AI
npx cdk deploy Radiant-dev-Auth --context environment=dev --context tier=1
npx cdk deploy Radiant-dev-AI --context environment=dev --context tier=1

# Phase 4: API & Admin
npx cdk deploy Radiant-dev-API --context environment=dev --context tier=1
npx cdk deploy Radiant-dev-Admin --context environment=dev --context tier=1
```

## Database Migrations

After infrastructure is deployed, run database migrations:

```
# Connect to Aurora and run migrations
cd packages/infrastructure/migrations
./run-migrations.sh --environment dev
```

Migration files are applied in order (44 total): 1. 001\_initial\_schema.sql - Base tables 2. 002\_tenant\_isolation.sql - RLS policies 3. 003\_ai\_models.sql - Providers and models 4. 004\_usage\_billing.sql - Usage tracking 5. 005\_admin\_approval.sql - Audit logs 6. 006\_self\_hosted\_models.sql - SageMaker config 7. 007\_external\_providers.sql - Provider settings 8. ... (see migrations/ for full list) 44. 044\_cost\_experiments\_security.sql - Cost tracking, A/B testing, security

## Post-Deployment Configuration

### Create First Super Admin

```
aws cognito-idp admin-create-user \
  --user-pool-id YOUR_ADMIN_POOL_ID \
  --username admin@example.com \
```

```
--user-attributes Name=email,Value=admin@example.com \
--temporary-password TempPass123! \
--message-action SUPPRESS

aws cognito-oidc admin-add-user-to-group \
--user-pool-id YOUR_ADMIN_POOL_ID \
--username admin@example.com \
--group-name super_admin
```

## Configure AI Providers

1. Navigate to Admin Dashboard → Providers
2. Add API keys for external providers:
  - OpenAI
  - Anthropic
  - Google AI
  - xAI (Grok)
  - DeepSeek
3. Verify connectivity with test requests

## Environment Configuration

### Tiers

Tier	Name	Use Case
1	SEED	Development, testing
2	STARTUP	Small production
3	GROWTH	Medium production
4	SCALE	Large production
5	ENTERPRISE	Enterprise with compliance

### Environment Variables

Required environment variables for deployment:

```
export AWS_REGION=us-east-1
export AWS_ACCOUNT_ID=123456789012
export RADIANT_ENVIRONMENT=dev # dev, staging, prod
export RADIANT_TIER=1          # 1-5
export RADIANT_DOMAIN@example.com
```

## Verification

### Health Checks

```
# API Health
curl https://YOUR_API_ENDPOINT/health

# Expected response:
# {"status": "healthy", "version": "4.18.0"}
```

## Smoke Tests

```
# Test chat completions
curl -X POST https://YOUR_API_ENDPOINT/v1/chat/completions \
-H "Authorization: Bearer YOUR_TOKEN" \
-H "Content-Type: application/json" \
-d '{"model": "gpt-4o-mini", "messages": [{"role": "user", "content": "Hello"}]}'
```

## Troubleshooting

### Common Issues

Issue	Cause	Solution
CDK bootstrap fails	Missing permissions	Ensure IAM user has AdministratorAccess
Aurora connection timeout	Security group	Check VPC endpoint and security group rules
Lambda cold starts	Function size	Enable provisioned concurrency for critical functions
Cognito auth fails	Pool configuration	Verify callback URLs and client settings

## Logs

```
# View Lambda logs
aws logs tail /aws/lambda/Radiant-dev-router --follow
```

```
# View ECS logs (LiteLLM)
aws logs tail /ecs/radiant-dev-litellm --follow
```

## Cleanup

To destroy all resources:

```
cd packages/infrastructure
npx cdk destroy --all --context environment=dev --context tier=1
```

**Warning:** This will delete all data including databases. Export data before destroying.

## CI/CD Pipeline

RADIANT includes a GitHub Actions CI/CD pipeline:

### Workflow Stages

Stage	Trigger	Actions
Lint	All PRs	ESLint, TypeScript check
Build	All PRs	Build shared, infrastructure, dashboard
Test	All PRs	Unit tests, coverage report
CDK Synth	All PRs	Validate CDK templates
Deploy Dev	Merge to develop	Auto-deploy to dev
Deploy Prod	Merge to main	Deploy with approval

## Pre-commit Hooks

Pre-commit hooks run automatically via Husky:

```
# Hooks run on every commit:  
- lint-staged (ESLint, Prettier)  
- Secret detection  
- Version bump enforcement  
- Discrete validation
```

To bypass (not recommended):

```
git commit --no-verify -m "message"
```

## Manual Deployment from CI

```
# Trigger deployment workflow  
gh workflow run deploy.yml -f environment=staging -f tier=2
```

## Testing Before Deployment

Always run tests before deploying:

```
# Run all tests  
pnpm test  
  
# Run E2E tests  
cd apps/admin-dashboard && pnpm test:e2e  
  
# Run CDK synthesis (validates templates)  
cd packages/infrastructure && npx cdk synth
```

See [Testing Guide](#) for comprehensive testing information.

## Support

For issues, check: 1. CloudWatch Logs for error details 2. CDK diff to verify expected changes 3. AWS Console for resource status 4. [Troubleshooting Guide](#) 5. [Error Codes Reference](#)

# RADIANT Platform - Administrator Guide

Complete guide for managing the RADIANT AI Platform via the Admin Dashboard

Version: 4.18.1 | Last Updated: December 2024

---

## Table of Contents

1. Introduction
  2. Accessing the Admin Dashboard
  3. Dashboard Overview
  4. Tenant Management
  5. User & Administrator Management
  6. AI Model Configuration
  7. Provider Management
  8. Billing & Subscriptions
  9. Storage Management
  10. Orchestration & Neural Engine
  11. Localization
  12. Configuration Management
  13. Security & Compliance
  14. Cost Analytics
  15. A/B Testing & Experiments
  16. Audit & Monitoring
  17. Database Migrations
  18. API Management
  19. Troubleshooting
  20. Delight System Administration
- 

## 1. Introduction

### 1.1 What is RADIANT?

RADIANT is a multi-tenant AWS SaaS platform providing unified access to 106+ AI models through:

- **50 External Provider Models:** OpenAI, Anthropic, Google, xAI, DeepSeek, and more
- **56 Self-Hosted Models:** Running on AWS SageMaker for cost control and privacy
- **Intelligent Routing:** Brain router for optimal model selection
- **Neural Engine:** Personalization learning from user interactions

## 1.2 Administrator Roles

Role	Permissions	Use Case
<b>Super Admin</b>	Full access to all features	Platform owner
<b>Admin</b>	Tenant management, billing, models	Operations team
<b>Operator</b>	Read access, limited actions	Support team
<b>Auditor</b>	Read-only access to logs	Compliance team

### Role Details

**Super Admin** - The highest privilege level with unrestricted access:  
- Create and delete tenants  
- Manage all administrators  
- Access all billing and financial data  
- Modify system-wide configuration  
- Approve production database migrations  
- Impersonate any tenant for debugging  
- Access compliance and audit reports  
- Typically limited to 1-3 people (CTO, lead engineer)

**Admin** - Day-to-day operations management:  
- Create and modify tenants (cannot delete)  
- Manage users within tenants  
- Configure AI models and providers  
- View billing data (cannot modify pricing)  
- Monitor system health  
- Cannot access other admin accounts  
- Typically assigned to operations team members

**Operator** - Limited support and monitoring:  
- View tenant information (read-only)  
- View user issues and support tickets  
- Monitor system health dashboards  
- Cannot modify any configuration  
- Cannot access billing or sensitive data  
- Typically assigned to support staff

**Auditor** - Compliance and security review:  
- Full read access to audit logs  
- Access to compliance reports  
- Cannot modify anything  
- Cannot view sensitive data (API keys, passwords)  
- Access is logged for compliance  
- Typically assigned to compliance officers or external auditors

## 1.3 Key Concepts

Concept	Description
<b>Tenant</b>	Organization with isolated data
<b>User</b>	End-user within a tenant
<b>Subscription</b>	Billing tier (1-7)
<b>Credits</b>	Currency for AI usage
<b>API Key</b>	Authentication for API access
<b>App</b>	Consumer application (Think Tank, etc.)

### Tenant Architecture Explained

A **Tenant** represents a complete organization using RADIANT. Each tenant has:

- **Complete Data Isolation:** All data is stored with tenant IDs and protected by PostgreSQL Row-Level Security (RLS). One tenant can never access another tenant's data, even if there's a bug in application code.
- **Separate Billing:** Each tenant has its own subscription, credit balance, and usage tracking. Costs are attributed to the correct tenant automatically.
- **Custom Configuration:** Tenants can customize model access, rate limits, and feature flags without affecting other tenants.
- **User Management:** Each tenant manages their own users, roles, and permissions independently.

### User vs Administrator

**Users** are end-users who interact with RADIANT-powered applications like Think Tank. They:  
- Sign up and log in via Cognito  
- Use AI models through the API or applications  
- Have credits deducted for usage  
- Cannot access the Admin Dashboard

**Administrators** manage the RADIANT platform itself. They:

- Access the Admin Dashboard - Manage tenants, users, and billing
- Configure AI models and providers
- Have no credits (administrative access is separate)

### Credit System Explained

Credits are RADIANT's universal currency for AI usage:

- **1 credit = \$0.01 USD** (configurable per deployment)
- Different models cost different amounts based on their API pricing
- Credits are deducted in real-time as requests complete
- Tenants can purchase credits or receive them through subscriptions
- Credits can be tracked, audited, and reported on

**Example Credit Costs:** | Model | Cost per 1K tokens | |-----|-----| | GPT-4o | 5 credits input, 15 credits output | | GPT-4o-mini | 0.5 credits input, 1.5 credits output | | Claude 3.5 Sonnet | 3 credits input, 15 credits output | | Self-hosted Llama | 0.2 credits (all) |

### API Key Types

RADIANT supports multiple API key types:

- **User API Keys:** Tied to a specific user, inherit user's permissions
  - **Service API Keys:** For server-to-server communication, not tied to a user
  - **Admin API Keys:** For administrative operations, require elevated permissions
  - **Scoped Keys:** Limited to specific models, endpoints, or rate limits
- 

## 2. Accessing the Admin Dashboard

### 2.1 URL and Login

1. Navigate to: <https://admin.your-domain.com>
2. Enter your email address
3. Enter your password
4. Complete MFA verification (required)

### 2.2 First Login

On first login:

1. You'll receive a temporary password via email
2. Enter the temporary password
3. Set a new password (12+ characters, mixed case, numbers, symbols)
4. Set up MFA using an authenticator app
5. You'll be redirected to the dashboard

### 2.3 Session Management

Setting	Value
<b>Session Duration</b>	8 hours
<b>Idle Timeout</b>	30 minutes
<b>Concurrent Sessions</b>	3 maximum
<b>Remember Device</b>	30 days

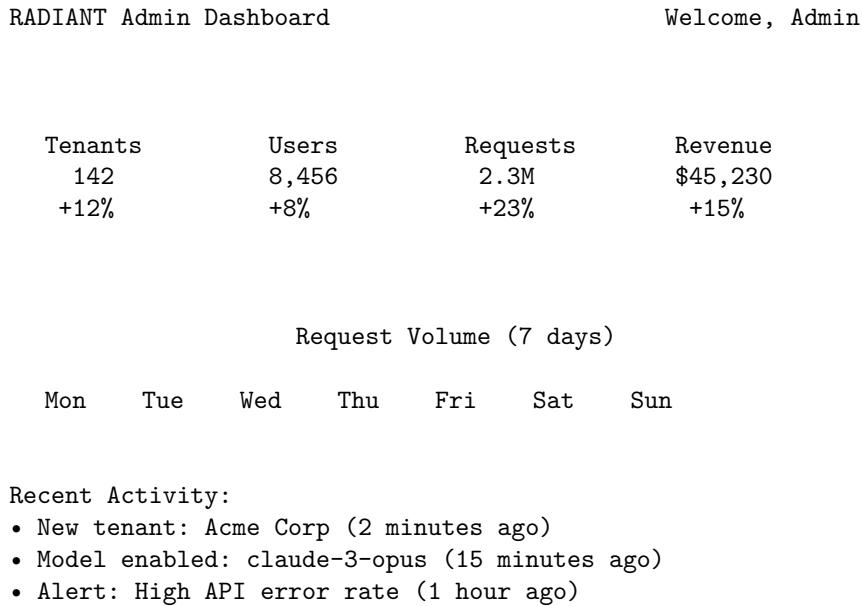
#### **2.4 Password Requirements**

- Minimum 12 characters
  - At least one uppercase letter
  - At least one lowercase letter
  - At least one number
  - At least one special character
  - Cannot reuse last 10 passwords

### 3. Dashboard Overview

### 3.1 Main Dashboard

The dashboard displays key metrics at a glance:



### 3.2 Navigation Menu

Section	Description
<b>Dashboard</b>	Overview and metrics
<b>Tenants</b>	Tenant management
<b>Users</b>	User management
<b>Models</b>	AI model configuration
<b>Providers</b>	Provider management
<b>Billing</b>	Subscriptions and credits
<b>Storage</b>	Storage usage
<b>Orchestration</b>	Neural engine settings
<b>Localization</b>	Translation management
<b>Configuration</b>	System settings
<b>Security</b>	Security monitoring

Section	Description
<b>Compliance</b>	Compliance reports
<b>Experiments</b>	A/B testing
<b>Cost</b>	Cost analytics
<b>Audit</b>	Audit logs
<b>Migrations</b>	Database migrations
<b>Notifications</b>	System alerts
<b>Settings</b>	Personal settings

---

## 4. Tenant Management

### 4.1 Viewing Tenants

Navigate to **Tenants** to see all organizations:

Column	Description
<b>Name</b>	Organization name
<b>Plan</b>	Subscription tier
<b>Users</b>	User count
<b>Status</b>	Active/Suspended/Trial
<b>Created</b>	Creation date
<b>Last Active</b>	Last API call

### 4.2 Creating a Tenant

1. Click “+ New Tenant”
2. Fill in required fields:
  - **Name:** Organization name
  - **Slug:** URL-friendly identifier
  - **Plan:** Initial subscription tier
  - **Admin Email:** Primary admin email
3. Configure optional settings:
  - Custom domain
  - Branding settings
  - Feature flags
4. Click “Create Tenant”

### 4.3 Tenant Details

View comprehensive tenant information:

Tenant: Acme Corporation

Overview      Users      Billing      Settings

Tenant ID: tn\_abc123xyz  
Status: Active

Plan: Professional (Tier 4)  
Created: 2024-01-15  
Last Active: 2 minutes ago

Usage This Month:

API Requests: 145,234  
Tokens Used: 12.5M  
Storage: 2.3 GB  
Credits Used: \$1,234.56

[Edit] [Suspend] [Delete] [Impersonate]

#### 4.4 Tenant Actions

Action	Description	Permission
Edit	Modify tenant settings	Admin
Suspend	Temporarily disable	Admin
Delete	Permanently remove	Super Admin
Impersonate	Login as tenant admin	Super Admin
Export	Export tenant data	Admin

#### 4.5 Data Isolation

Each tenant has complete data isolation:

- Separate database rows with RLS
- Unique API keys
- Isolated storage buckets
- Independent usage tracking

---

## 5. User & Administrator Management

### 5.1 Administrator Roles

Role	Dashboard Access	API Access	Billing	Audit
Super Admin	Full	Full	Full	Full
Admin	Full	Full	Read	Read
Operator	Read	Read	None	Read
Auditor	Logs only	None	None	Full

### 5.2 Managing Administrators

Navigate to **Administrators** to:

1. **Invite New Admin:**
  - Click “+ Invite Administrator”
  - Enter email address
  - Select role

- Click “Send Invitation”
2. **Modify Admin:**
    - Click on administrator row
    - Edit role or permissions
    - Click “Save Changes”
  3. **Remove Admin:**
    - Click “Remove” button
    - Confirm removal
    - Admin’s sessions are invalidated immediately

### 5.3 Viewing Tenant Users

Navigate to **Tenants** → [Tenant] → **Users** to see:

Field	Description
<b>Email</b>	User email
<b>Name</b>	Display name
<b>Role</b>	Tenant role
<b>Status</b>	Active/Invited/Disabled
<b>Last Login</b>	Last authentication
<b>API Keys</b>	Number of active keys

### 5.4 User Actions

Action	Description
<b>Reset Password</b>	Send password reset email
<b>Disable</b>	Prevent login
<b>Enable</b>	Restore access
<b>Delete</b>	Remove user data
<b>View Sessions</b>	See active sessions

## 6. AI Model Configuration

### 6.1 Model Registry

Navigate to **Models** to see all available models:

AI Models 106 Total

Filter: [All ] Category: [All ] Status: [Enabled ]

Model	Provider	Category	Tier	Status
gpt-4o	OpenAI	Chat	1	Enabled
gpt-4o-mini	OpenAI	Chat	1	Enabled
claude-3-opus	Anthropic	Chat	2	Enabled
claude-3-sonnet	Anthropic	Chat	1	Enabled
gemini-pro	Google	Chat	1	Enabled

llama-3.1-70b	Self-Host	Chat	3	Enabled
whisper-large	Self-Host	Audio	3	Disabled

[+ Add Model] [Import Models] [Export Config]

## 6.2 Model Categories

Category	Description	Example Models
<b>Chat/LLM</b>	Text generation	GPT-4o, Claude 3, Gemini
<b>Embedding</b>	Vector embeddings	text-embedding-3-large
<b>Vision</b>	Image understanding	GPT-4V, Claude Vision
<b>Audio</b>	Speech-to-text	Whisper, Deepgram
<b>Image</b>	Image generation	DALL-E 3, Stable Diffusion
<b>Code</b>	Code generation	Codestral, DeepSeek Coder
<b>Scientific</b>	Research models	BioGPT, ChemLLM

### Category Details

**Chat/LLM (Large Language Models):** The core of RADIANT. These models handle conversational AI, content generation, summarization, and general-purpose text tasks. They’re the most commonly used and include flagship models from OpenAI, Anthropic, Google, and open-source alternatives.

**Embedding Models:** Convert text into numerical vectors for semantic search, similarity matching, and retrieval-augmented generation (RAG). Essential for building knowledge bases and search functionality. Vectors are typically 1536-3072 dimensions.

**Vision Models:** Analyze images, extract text (OCR), describe visual content, and answer questions about images. Increasingly important for document processing, accessibility, and multimodal applications.

**Audio Models:** Transcribe speech to text, translate audio, and identify speakers. Whisper is the most popular, offering excellent accuracy across 99 languages. Used for meeting transcription, accessibility, and voice interfaces.

**Image Generation:** Create images from text descriptions. DALL-E 3 offers the best prompt following, while Stable Diffusion provides more customization options. Consider content policies when enabling these.

**Code Models:** Specialized for programming tasks including code generation, explanation, debugging, and refactoring. Some are fine-tuned on specific languages or frameworks.

**Scientific Models:** Domain-specific models trained on scientific literature. Useful for research applications but require careful evaluation for accuracy.

## 6.3 Model Configuration

Click on a model to configure:

Setting	Description
<b>Enabled</b>	Available for use
<b>Min Tier</b>	Minimum subscription tier
<b>Rate Limits</b>	Requests per minute
<b>Max Tokens</b>	Maximum context/output
<b>Temperature Range</b>	Allowed temperature values
<b>Price Override</b>	Custom pricing

## Configuration Settings Explained

**Enabled:** When disabled, the model is hidden from users and API requests return “model not found”. Use this to temporarily remove models during maintenance or to restrict access to specific models.

**Min Tier:** Sets the minimum subscription tier required to access this model. For example, setting GPT-4 to Tier 2 means Free tier users cannot use it. This helps control costs and create upgrade incentives.

**Rate Limits:** Controls requests per minute per user for this model. Prevents abuse and ensures fair access. Set based on the provider’s rate limits and your capacity: - Conservative: 10-20 requests/minute - Standard: 50-100 requests/minute

- High: 200+ requests/minute (requires provider rate limit increases)

**Max Tokens:** Limits context window and output length. Useful for controlling costs since longer contexts cost more. Set based on use case: - Short tasks (Q&A): 4,096 tokens - Medium tasks (writing): 16,384 tokens - Long tasks (analysis): 32,768+ tokens

**Temperature Range:** Restricts the temperature parameter users can set. Temperature controls randomness: - 0.0: Deterministic, consistent outputs - 0.7: Balanced creativity and consistency - 1.0+: More creative, less predictable

Restricting range (e.g., 0.0-1.0) prevents users from setting extreme values that produce poor results.

**Price Override:** Allows custom pricing different from the default. Useful for: - Offering discounts on specific models - Increasing prices for premium models - Matching competitor pricing - A/B testing pricing strategies

## 6.4 Self-Hosted Models

For Tier 3+ deployments:

1. Navigate to **Models → Self-Hosted**
2. Click “**+ Add Self-Hosted Model**”
3. Configure:
  - **Model ID:** Unique identifier
  - **SageMaker Endpoint:** Endpoint name
  - **Instance Type:** ml.g5.xlarge, etc.
  - **Auto-Scaling:** Min/max instances
4. Deploy model to SageMaker

## 6.5 Thermal States (Self-Hosted)

State	Description	Response Time
<b>HOT</b>	Always running	<100ms
<b>WARM</b>	Scaled down	<5s
<b>COLD</b>	Stopped	30-60s
<b>OFF</b>	Disabled	N/A

---

## 7. Provider Management

### 7.1 External Providers

Navigate to **Providers** to manage API integrations:

Provider	Models	Status	Health
<b>OpenAI</b>	12	Configured	99.9%
<b>Anthropic</b>	6	Configured	99.8%
<b>Google AI</b>	8	Configured	99.7%
<b>xAI</b>	2	Configured	99.5%
<b>DeepSeek</b>	4	Not configured	-

## 7.2 Adding Provider Credentials

1. Click on provider name
2. Click “Configure”
3. Enter API credentials:
  - **API Key:** Provider API key
  - **Organization ID:** (if applicable)
  - **Base URL:** (for custom endpoints)
4. Click “Test Connection”
5. Click “Save”

## 7.3 Provider Health Monitoring

View real-time provider health:

Provider Health: OpenAI

Status:	Healthy
Uptime (30d):	99.94%
Avg Latency:	245ms
P95 Latency:	520ms
Error Rate:	0.02%

Last 24 Hours:

12am      6am      12pm      6pm      12am

## 7.4 Fallback Configuration

Configure provider fallbacks:

1. Navigate to **Providers → Fallbacks**
2. Set priority order for each model category
3. Configure automatic failover rules
4. Set retry policies

---

# 8. Billing & Subscriptions

## 8.1 Subscription Tiers

Tier	Name	Monthly	Features
1	Free	\$0	Basic models, 1K requests
2	Starter	\$29	More models, 10K requests
3	Professional	\$99	All external models, 100K requests
4	Business	\$299	Priority support, 500K requests
5	Enterprise	\$999	Self-hosted, unlimited
6	Enterprise+	Custom	Custom SLAs, dedicated support
7	Ultimate	Custom	On-premise options

## 8.2 Credit System

Credits are the universal currency for AI usage:

Model Type	Cost per 1M Tokens
GPT-4o	500 credits
GPT-4o-mini	50 credits
Claude 3 Opus	600 credits
Claude 3 Sonnet	150 credits
Self-hosted	20 credits

## 8.3 Managing Subscriptions

Navigate to **Billing → Subscriptions**:

1. View current subscription
2. Upgrade/downgrade tier
3. Add credit packages
4. View invoices
5. Update payment method

## 8.4 Usage Reports

Generate usage reports:

1. Navigate to **Billing → Reports**
2. Select date range
3. Choose grouping (by tenant/model/user)
4. Export as CSV/PDF

## 8.5 Billing Alerts

Configure alerts for:

- Credit balance low
- Usage spike
- Approaching quota
- Failed payments

## 9. Storage Management

### 9.1 Storage Overview

Navigate to **Storage** to monitor:

#### Storage Overview

Total Used: 234.5 GB of 500 GB (47%)

##### By Type:

Documents:	120.3 GB (51%)
Images:	45.2 GB (19%)
Audio:	38.7 GB (17%)
Video:	22.1 GB (9%)
Other:	8.2 GB (4%)

##### Top Tenants:

1. Acme Corp	45.2 GB
2. TechStart	32.1 GB
3. DataCo	28.4 GB

### 9.2 Storage Tiers

Tier	Included	Additional
Free	1 GB	N/A
Starter	10 GB	\$0.10/GB
Professional	100 GB	\$0.08/GB
Business	500 GB	\$0.05/GB
Enterprise	2 TB	\$0.03/GB

### 9.3 File Management

Manage uploaded files:

- View file metadata
  - Download files
  - Delete files
  - Set retention policies
- 

## 10. Orchestration & Neural Engine

### 10.1 Brain Router

The Brain Router automatically selects optimal models:

Factor	Weight	Description
<b>Cost</b>	30%	Price optimization
<b>Quality</b>	30%	Output quality
<b>Speed</b>	20%	Response latency
<b>Availability</b>	20%	Provider health

## 10.2 Neural Patterns

Configure orchestration patterns:

Pattern	Description	Use Case
<b>Single</b>	One model	Simple requests
<b>Fallback</b>	Primary + backup	High availability
<b>Parallel</b>	Multiple simultaneous	Consensus
<b>Chain</b>	Sequential models	Complex tasks

## 10.3 Workflow Templates

Create reusable workflows:

1. Navigate to **Orchestration → Workflows**
  2. Click “+ New Workflow”
  3. Define steps and conditions
  4. Set triggers and parameters
  5. Save and activate
- 

# 11. Localization

## 11.1 Translation Management

Navigate to **Localization** to manage:

- Supported languages
- Translation strings
- AI translation settings

## 11.2 Supported Languages

Language	Code	Status
English	en	Default
Spanish	es	Enabled
French	fr	Enabled
German	de	Enabled
Japanese	ja	Enabled
Chinese	zh	Enabled

## 11.3 AI Translation

Enable AI-powered translation:

1. Navigate to **Localization** → **Settings**
  2. Enable “**AI Translation**”
  3. Select translation model
  4. Configure quality settings
- 

## 12. Configuration Management

### 12.1 System Configuration

Navigate to **Configuration** to manage:

Category	Settings
<b>General</b>	Platform name, domain, timezone
<b>Email</b>	SMTP settings, templates
<b>Security</b>	Password policy, MFA settings
<b>API</b>	Rate limits, CORS settings
<b>Features</b>	Feature flags

### 12.2 Tenant Overrides

Allow tenant-specific configuration:

1. Navigate to **Configuration** → **Tenant Overrides**
2. Select tenant
3. Override specific settings
4. Save changes

### 12.3 SSM Parameters

System configuration is stored in AWS SSM:

Parameter	Description
/radiant/prod/database/url	Database connection
/radiant/prod/api/rate-limit	API rate limits
/radiant/prod/features/*	Feature flags

---

## 13. Security & Compliance

### 13.1 Security Dashboard

Navigate to **Security** to monitor:

Security Dashboard      Threat Level: Low

Active Threats: 0  
Failed Logins: 23 (last 24h)  
Suspicious IPs: 2 blocked

MFA Adoption: 94%

Recent Alerts:

Unusual login location - user@acme.com (2h ago)  
Resolved: Brute force attempt blocked (5h ago)  
Resolved: API key rotated - tenant xyz (1d ago)

## 13.2 Anomaly Detection

Automatic detection of:

- Impossible travel (geographic anomalies)
- Session hijacking attempts
- Brute force attacks
- Unusual API patterns

## 13.3 Compliance Reports

Navigate to **Compliance** to generate:

Framework	Description
<b>SOC 2</b>	Service organization controls
<b>HIPAA</b>	Healthcare data protection
<b>GDPR</b>	EU data protection
<b>ISO 27001</b>	Information security

## 13.4 Generating Reports

1. Click “**Generate Report**”
  2. Select framework
  3. Choose date range
  4. Select metrics to include
  5. Generate PDF/CSV
- 

# 14. Cost Analytics

## 14.1 Cost Dashboard

Navigate to **Cost** to view:

Cost Analytics

Period: Last 30 Days

Total Spend: \$12,456.78 (+12% vs last month)  
Projected: \$14,200.00 (this month)

By Provider:

OpenAI: \$6,234.56 (50%)  
Anthropic: \$3,456.78 (28%)  
Self-hosted: \$1,234.56 (10%)  
Other: \$1,530.88 (12%)

#### AI Recommendations:

Switch 23% of GPT-4 calls to GPT-4-mini (save \$890/mo)  
Enable caching for repeated queries (save \$340/mo)

## 14.2 Cost Alerts

Configure alerts:

- Daily budget exceeded
- Weekly spend spike
- Per-tenant limits
- Per-model thresholds

## 14.3 Cost Optimization

Review AI-powered recommendations:

1. Navigate to **Cost → Insights**
  2. Review suggestions
  3. Click “**Apply**” to implement (requires approval)
  4. Track savings over time
- 

# 15. A/B Testing & Experiments

## 15.1 Experiment Dashboard

Navigate to **Experiments** to manage:

Experiment	Status	Variants	Sample Size
Model routing v2	Running	3	45,234
Prompt optimization	Running	2	12,456
Temperature test	Completed	4	89,123

## 15.2 Creating an Experiment

1. Click “**+ New Experiment**”
2. Configure:
  - **Name:** Descriptive name
  - **Hypothesis:** What you’re testing
  - **Variants:** Control + treatments
  - **Traffic Split:** Percentage per variant
  - **Success Metric:** What to measure
3. Set targeting rules
4. Start experiment

## 15.3 Statistical Analysis

View results with:

- Conversion rates per variant
  - Statistical significance (p-value)
  - Confidence intervals
  - Sample size recommendations
- 

# 16. Audit & Monitoring

## 16.1 Audit Logs

Navigate to **Audit** to view all actions:

Column	Description
<b>Timestamp</b>	When action occurred
<b>Actor</b>	Who performed action
<b>Action</b>	What was done
<b>Resource</b>	What was affected
<b>IP Address</b>	Source IP
<b>Details</b>	Additional context

## 16.2 Log Filtering

Filter by:

- Date range
- Actor (user/admin)
- Action type
- Resource type
- Severity level

## 16.3 Log Export

Export logs for compliance:

1. Set filter criteria
2. Click “**Export**”
3. Choose format (CSV/JSON)
4. Download file

## 16.4 Real-Time Monitoring

Navigate to **Monitoring** for:

- Live request stream
  - Error rate graphs
  - Latency percentiles
  - Active users count
-

## 17. Database Migrations

### 17.1 Migration Workflow

RADIANT uses dual-admin approval for production migrations:

1. **Submit:** Admin submits migration
2. **Review:** Second admin reviews
3. **Approve:** Second admin approves
4. **Execute:** Migration runs
5. **Verify:** Automatic verification

### 17.2 Pending Migrations

Navigate to **Migrations** to see:

Database Migrations

Pending Approval:

```
#045 - Add user preferences table
Submitted by: alice@company.com (2 hours ago)
[View SQL] [Approve] [Reject]
```

Recent Migrations:

```
#044 - Cost tracking tables (applied 2024-12-24)
#043 - Experiment framework (applied 2024-12-20)
#042 - Security anomalies (applied 2024-12-15)
```

### 17.3 Approving Migrations

1. Review the SQL in “**View SQL**”
  2. Check for potential issues
  3. Click “**Approve**” or “**Reject**”
  4. Add comment explaining decision
- 

## 18. API Management

### 18.1 API Keys

Manage platform API keys:

1. Navigate to **Settings** → **API Keys**
2. View existing keys
3. Create new keys with scopes
4. Revoke compromised keys

### 18.2 Rate Limiting

Configure rate limits:

Level	Default	Configurable
<b>Global</b>	10,000/min	Yes
<b>Per-Tenant</b>	1,000/min	Yes
<b>Per-User</b>	100/min	Yes
<b>Per-Key</b>	60/min	Yes

## 18.3 Webhooks

Configure outgoing webhooks:

1. Navigate to **Settings → Webhooks**
  2. Add webhook URL
  3. Select events to send
  4. Test webhook
  5. Enable webhook
- 

# 19. Troubleshooting

## 19.1 Common Issues

### High Error Rate

1. Check **Providers** for unhealthy providers
2. Review **Audit** logs for patterns
3. Check **Monitoring** for load spikes
4. Verify API key validity

### Slow Response Times

1. Check provider latency in **Providers**
2. Review model selection in **Orchestration**
3. Check for cold-start issues (self-hosted)
4. Verify database performance

### Authentication Failures

1. Check user status in **Users**
2. Verify MFA configuration
3. Review **Audit** logs for login attempts
4. Check for IP blocks in **Security**

## 19.2 Support Resources

Resource	Description
<b>Documentation</b>	This guide + online docs
<b>Status Page</b>	status.radiant.example.com
<b>Support Email</b>	support@radiant.example.com
<b>Emergency</b>	+1-555-RADIANT

### 19.3 Log Locations

Service	Log Group
API Gateway	/aws/apigateway/radiant
Lambda	/aws/lambda/radiant-*
Admin Dashboard	/aws/cloudfront/admin
Database	/aws/rds/cluster/radiant

## 20. Delight System Administration

The Delight System provides personality, achievements, and engagement features for Think Tank users.

### 20.1 Accessing Delight Admin

Navigate to **Think Tank → Delight** in the admin sidebar.

### 20.2 Dashboard Overview

The Delight dashboard shows:

Metric	Description
<b>Messages Shown</b>	Total delight messages displayed
<b>Achievements Unlocked</b>	Total achievements earned by users
<b>Easter Eggs Found</b>	Hidden features discovered
<b>Active Users</b>	Users with Delight enabled

### 20.3 Managing Categories

Toggle entire categories on/off:

Category	Purpose
Domain Loading	Messages while loading domain expertise
Domain Transition	Messages when switching topics
Time Awareness	Time-of-day contextual messages
Model Dynamics	Messages about AI consensus/disagreement
Complexity Signals	Feedback on query complexity
Synthesis Quality	Post-response quality indicators
Achievements	Milestone celebrations
Wellbeing	Break/health reminders
Easter Eggs	Hidden features
Sounds	Audio feedback

### 20.4 Managing Messages

- **Create:** Add new delight messages with targeting options
- **Edit:** Modify text, triggers, and display settings
- **Delete:** Remove messages (soft delete)
- **Toggle:** Enable/disable individual messages

## Message Targeting Options

Option	Values
Injection Point	pre_execution, during_execution, post_execution
Trigger Type	domain_loading, time_aware, model_dynamics, etc.
Domain Families	science, humanities, creative, technical, etc.
Time Contexts	morning, afternoon, evening, night, weekend
Display Style	subtle, moderate, expressive

## 20.5 Statistics Dashboard

Access detailed usage statistics at **Delight → View Statistics**:

- **Weekly Trends:** 12-week activity history
- **Top Messages:** Most-shown messages with engagement data
- **Achievement Stats:** Unlock rates, time-to-unlock averages
- **Easter Egg Stats:** Discovery rates by egg
- **User Engagement:** Leaderboard by achievement points

## 20.6 Managing Achievements

Configure achievement unlock criteria:

Setting	Description
Threshold	Number required to unlock
Rarity	common, uncommon, rare, epic, legendary
Points	Score value for leaderboards
Hidden	Only visible after unlock

## 20.7 Managing Easter Eggs

Configure hidden features:

Setting	Description
Trigger Type	key_sequence, text_input, time_based, random
Trigger Value	The activation pattern
Effect Type	mode_change, visual_transform, sound_play
Duration	How long the effect lasts

## 20.8 API Endpoints

Endpoint	Method	Description
/api/admin/delight/dashboard	GET	Dashboard data
/api/admin/delight/statistics	GET	Detailed statistics
/api/admin/delight/categories/:id	PATCH	Toggle category
/api/admin/delight/messages	POST	Create message
/api/admin/delight/messages/:id	PUT/DELETE	Update/delete
/api/admin/delight/user-engagement	GET	User leaderboard

## Appendix: Quick Reference

### Keyboard Shortcuts

Shortcut	Action
G + D	Go to Dashboard
G + T	Go to Tenants
G + M	Go to Models
G + B	Go to Billing
G + A	Go to Audit
?	Show shortcuts

### Status Indicators

Icon	Meaning
Green circle	Healthy/Success
Yellow triangle	Warning
Red circle	Error/Failed
Blue circle	In Progress
Grey circle	Disabled/Pending

---

*Document Version: 4.18.1 Last Updated: December 2024*

# RADIANT Deployer - Administrator Guide

Complete guide for deploying and managing RADIANT infrastructure using the Swift Deployer App

Version: 4.18.1 | Last Updated: December 2024

---

## Table of Contents

1. Introduction
  2. System Requirements
  3. Installation
  4. First-Time Setup
  5. AWS Credentials Management
  6. Deployment Operations
  7. Multi-Region Deployments
  8. Snapshots & Rollbacks
  9. AI Assistant
  10. Package Management
  11. Monitoring & Health Checks
  12. Security Features
  13. Troubleshooting
  14. Reference
- 

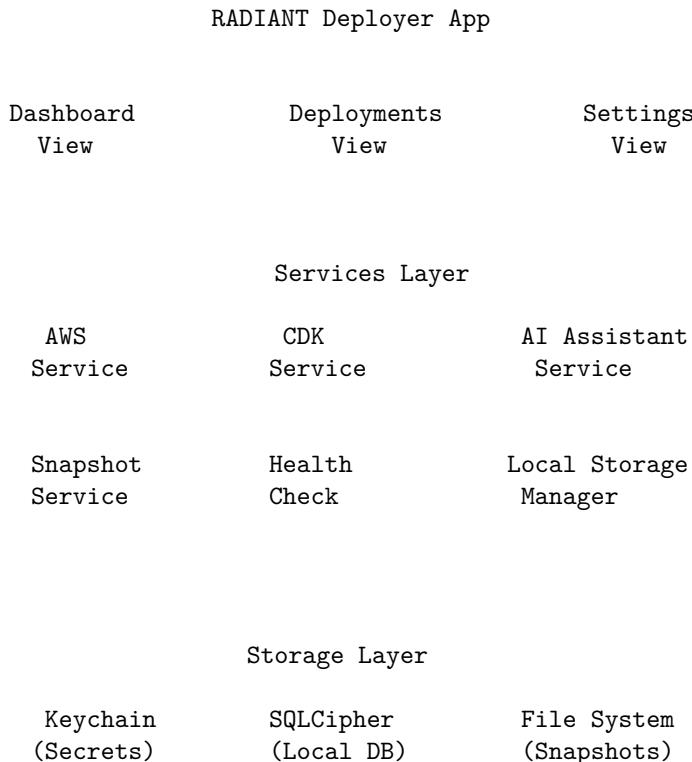
## 1. Introduction

### 1.1 What is RADIANT Deployer?

RADIANT Deployer is a native macOS application that provides a complete deployment management solution for the RADIANT platform. It offers:

- **One-Click Deployments:** Deploy entire infrastructure stacks with a single click
- **AI-Powered Assistance:** Claude-powered assistant for deployment guidance
- **Snapshot Management:** Create and restore deployment snapshots
- **Multi-Region Support:** Deploy across multiple AWS regions
- **Health Monitoring:** Real-time health checks and status monitoring
- **Secure Credential Storage:** Keychain-integrated credential management

## 1.2 Architecture



## 1.3 Key Features

Feature	Description
<b>Deployment Wizard</b>	Step-by-step guided deployment process
<b>Lock-Step Mode</b>	Ensures component version consistency
<b>Automatic Rollback</b>	Reverts failed deployments automatically
<b>Offline Mode</b>	Core functionality works without internet
<b>Audit Logging</b>	Complete deployment history tracking

### Deployment Wizard Explained

The Deployment Wizard breaks down the complex AWS infrastructure deployment into manageable steps. Instead of manually running CDK commands and configuring services, the wizard:

- Validates Prerequisites:** Checks that all required software (Node.js, AWS CLI, CDK) is installed and properly configured before starting
- Guides Configuration:** Walks you through selecting environment (dev/staging/prod), tier (1-5), and region settings with explanations for each option
- Shows Progress:** Displays real-time progress as each CloudFormation stack deploys, with estimated time remaining
- Handles Errors:** If any step fails, provides clear error messages and suggested remediation steps

## Lock-Step Mode Explained

Lock-Step Mode prevents version mismatches between RADIANT components that could cause compatibility issues:

- **What it does:** Ensures the Admin Dashboard, Lambda functions, and database schema are all on compatible versions
- **Why it matters:** A v4.18 Lambda function might expect database columns that don't exist in a v4.17 schema, causing runtime errors
- **How it works:** Before deployment, the system checks version numbers across all components and blocks deployment if drift exceeds the configured maximum (default: 1 minor version)
- **When to disable:** Only disable during development when testing individual component changes

## Automatic Rollback Explained

Automatic Rollback protects your production environment by reverting failed deployments:

- **Trigger conditions:** Activates when any deployment phase fails (CDK deploy error, health check failure, migration error)
- **Rollback process:** Restores the pre-deployment snapshot, which includes database state, Lambda code, and configuration
- **Recovery time:** Typically 5-10 minutes depending on the size of changes being reverted
- **Notification:** Sends alerts via configured channels (email, Slack) when rollback occurs

## Offline Mode Explained

Offline Mode allows essential operations when internet connectivity is unavailable:

- **Available offline:** Viewing deployment history, browsing local snapshots, reviewing configuration, accessing cached documentation
- **Requires internet:** Deploying to AWS, health checks, AI assistant queries, credential validation
- **Data sync:** When connectivity returns, local changes sync automatically with the cloud state

## Audit Logging Explained

Every action in the Deployer is logged for compliance and troubleshooting:

- **What's logged:** User identity, timestamp, action type, parameters, outcome, duration
- **Storage:** Logs stored locally in SQLCipher database and optionally synced to CloudWatch
- **Retention:** Local logs kept for 90 days by default (configurable)
- **Export:** Logs can be exported to CSV/JSON for compliance audits

---

## 2. System Requirements

### 2.1 Hardware Requirements

Component	Minimum	Recommended
macOS Version	13.0 (Ventura)	14.0+ (Sonoma)
Processor	Apple Silicon or Intel	Apple Silicon M1+
Memory	8 GB RAM	16 GB RAM
Storage	2 GB free	10 GB free
Display	1280x800	1440x900+

## Why These Requirements?

**macOS 13.0+:** Required for Swift 5.9 runtime and modern SwiftUI features. Older versions lack required system APIs for secure keychain access and modern networking.

**Apple Silicon Recommended:** While Intel Macs are supported, Apple Silicon provides 2-3x faster CDK synthesis and compilation times. The app is built as a universal binary supporting both architectures.

**8 GB RAM Minimum:** CDK synthesis loads the entire infrastructure definition into memory. Complex deployments (Tier 3+) with many resources may require more memory. With 8 GB, you may experience slowdowns during synthesis.

**16 GB RAM Recommended:** Allows comfortable multitasking while deployments run in the background. Essential if you're also running Docker, IDEs, or other development tools.

**2 GB Storage Minimum:** Covers the application itself (~200 MB), local database (~50 MB), and several snapshots. However, snapshots can grow large over time.

**10 GB Storage Recommended:** Provides room for multiple deployment snapshots, comprehensive logs, and CDK cache. Each full snapshot can be 100-500 MB depending on your deployment size.

## 2.2 Software Requirements

Software	Version	Purpose	Installation
Xcode	15.0+	Swift runtime (Command Line Tools sufficient)	<code>xcode-select --install</code>
AWS CLI	2.x	AWS operations	<code>brew install awscli</code>
Node.js	20.x LTS	CDK operations	<code>brew install node@20</code>
AWS CDK	2.120+	Infrastructure deployment	<code>npm install -g aws-cdk</code>
pnpm	8.x+	Package management	<code>npm install -g pnpm</code>

### Software Explained

**Xcode Command Line Tools:** Provides the Swift runtime and compiler. You don't need the full Xcode IDE - just the command line tools (4 GB vs 12 GB download). The Deployer uses Swift for its native performance and seamless macOS integration.

**AWS CLI v2:** The official AWS command-line interface. Used internally by the Deployer to execute AWS operations, assume IAM roles, and query service status. Version 2 is required for SSO support and improved credential handling.

**Node.js 20 LTS:** Required to run the AWS CDK, which is written in TypeScript. LTS (Long Term Support) versions receive security updates for 30 months. The Deployer manages Node.js processes internally during CDK operations.

**AWS CDK 2.120+:** The Cloud Development Kit that defines RADIANT's infrastructure as TypeScript code. Version 2.120+ includes critical bug fixes for Aurora PostgreSQL and Lambda layer handling. The CDK synthesizes your infrastructure into CloudFormation templates.

**pnpm 8.x:** A fast, disk-efficient package manager used to install CDK dependencies. Chosen over npm for its superior handling of monorepo workspaces and 2-3x faster installation times.

## 2.3 AWS Requirements

Requirement	Details	How to Verify
<b>AWS Account</b>	Active account with billing enabled	Check AWS Console billing page
<b>IAM User</b>	AdministratorAccess or equivalent	<code>aws sts get-caller-identity</code>
<b>Regions</b>	Access to us-east-1 (required) + additional regions	<code>aws ec2 describe-regions</code>
<b>Service Quotas</b>	Default quotas sufficient for Tier 1-2	AWS Service Quotas console

### AWS Requirements Explained

**Active AWS Account with Billing:** RADIANT deploys real AWS resources that incur costs. Billing must be enabled and a valid payment method on file. New accounts have a \$5 pending charge verification. Free tier covers some resources for 12 months but won't cover all RADIANT components.

**IAM User with AdministratorAccess:** The Deployer needs broad permissions to create and manage resources across many AWS services. For production, you can use a scoped-down policy (see Appendix A), but AdministratorAccess is recommended for initial setup to avoid permission errors.

**Required Permissions Include:** - CloudFormation (stack operations) - Lambda (function management) - API Gateway (REST API setup) - RDS/Aurora (database provisioning) - Cognito (user pool management) - S3 (storage buckets) - IAM (role creation) - SSM (parameter storage) - Secrets Manager (credential storage) - CloudWatch (logging and monitoring)

**us-east-1 Required:** This region is required even if you deploy to other regions because: - ACM certificates for CloudFront must be in us-east-1 - Some global services (IAM, Route 53) operate from us-east-1 - CDK bootstrap resources are region-specific

**Service Quotas:** Default AWS quotas are sufficient for Tier 1-2 deployments. Tier 3+ may require quota increases for: - Lambda concurrent executions (default: 1,000) - RDS instances (default: 40) - VPC Elastic IPs (default: 5)

To request quota increases: AWS Console > Service Quotas > Select service > Request increase

## 3. Installation

### 3.1 Download and Install

#### Option A: Pre-built Application (Recommended)

1. Download the latest release from GitHub Releases
2. Drag `RadiantDeployer.app` to `/Applications`
3. Right-click and select “Open” (first launch only)
4. Grant necessary permissions when prompted

#### Option B: Build from Source

```
# Clone the repository
git clone https://github.com/your-org/radiant.git
cd radiant/apps/swift-deployer
```

```
# Build the application
```

```

swift build -c release

# Run the application
swift run RadiantDeployer

```

## 3.2 Initial Permissions

The app requires the following permissions:

Permission	Purpose	How to Grant
<b>Keychain Access</b>	Store AWS credentials securely	Approve on first credential save
<b>Network Access</b>	Connect to AWS and AI services	Approve in System Settings
<b>File Access</b>	Save snapshots and logs	Approve when prompted

## 3.3 Verify Installation

1. Launch RadiantDeployer
  2. Navigate to **Settings** → **About**
  3. Verify version shows 4.18.1
  4. Check all services show green status
- 

# 4. First-Time Setup

## 4.1 Setup Wizard

On first launch, the Setup Wizard guides you through:

### Setup Wizard

Step 1: Welcome	Complete
Step 2: AWS Credentials	In Progress
Step 3: Environment Configuration	Pending
Step 4: AI Assistant Setup	Pending
Step 5: Verification	Pending

## 4.2 AWS Credentials Setup

1. Click “Add AWS Credentials”
2. Enter your credentials:
  - **Name:** Descriptive name (e.g., “Production Account”)
  - **Access Key ID:** AKIA... (20 characters)
  - **Secret Access Key:** Your secret key (40 characters)
  - **Region:** Primary region (e.g., us-east-1)
3. Click “Validate” to test connectivity
4. Click “Save” to store securely in Keychain

## 4.3 Environment Configuration

Configure your deployment environment:

Setting	Description	Default
<b>Environment</b>	dev, staging, or prod	dev
<b>Tier</b>	Infrastructure tier (1-5)	1
<b>Domain</b>	Your domain name	Required for Tier 2+
<b>Stack Prefix</b>	CDK stack name prefix	radiant

### Environment Types Explained

**Development (dev):** For active development and testing. Features relaxed security settings, smaller instance sizes, and no deletion protection. Data can be reset freely. Cost: ~\$50-150/month for Tier 1.

**Staging (staging):** Pre-production environment that mirrors production configuration. Use this to test deployments before going live. Same security as production but can use smaller instances. Cost: ~60% of production.

**Production (prod):** Live environment serving real users. Includes deletion protection, multi-AZ deployments, automated backups, and enhanced monitoring. Never deploy untested changes directly to production.

### Infrastructure Tiers Explained

Tier	Name	Monthly Cost	Use Case	Resources
1	SEED	\$50-150	Development, POC	Single-AZ, t3.small instances, 20GB storage
2	STARTUP	\$200-400	Small production	Multi-AZ database, WAF, basic monitoring
3	GROWTH	\$1,000-2,500	Medium production	Self-hosted models, HIPAA compliance, enhanced security
4	SCALE	\$4,000-8,000	Large production	Multi-region, global database, dedicated instances
5	ENTERPRISE	\$15,000-35,000	Enterprise	Custom SLAs, dedicated support, on-premise options

**Choosing the Right Tier:** Start with Tier 1 for development. Move to Tier 2 when you have paying customers. Tier 3+ is for organizations with compliance requirements or high traffic.

### Domain Configuration

For Tier 2+, you need a custom domain:

1. Register a domain in Route 53 or transfer an existing domain
2. Create a hosted zone in Route 53 for your domain
3. Request an ACM certificate in us-east-1 for \*.yourdomain.com
4. Validate the certificate via DNS (automatic if using Route 53)

The Deployer will configure: - api.yourdomain.com - API Gateway endpoint - admin.yourdomain.com - Admin Dashboard - app.yourdomain.com - Think Tank application

## 4.4 AI Assistant Setup (Optional)

Enable the Claude-powered AI assistant:

1. Navigate to **Settings** → **AI Assistant**
  2. Enter your Anthropic API key
  3. Select response style:
    - **Concise**: Brief, action-focused responses
    - **Detailed**: In-depth explanations
    - **Tutorial**: Step-by-step guidance
  4. Test the connection with a sample query
- 

## 5. AWS Credentials Management

### 5.1 Credential Sets

Manage multiple AWS accounts:

Field	Description	Example
<b>Name</b>	Friendly identifier	“Production”
<b>Access Key ID</b>	AWS access key	AKIAIOSFODNN7EXAMPLE
<b>Secret Access Key</b>	AWS secret	(stored encrypted)
<b>Region</b>	Default region	us-east-1
<b>Role ARN</b>	Optional assume role	arn:aws:iam::123:role/deploy

### 5.2 Adding Credentials

1. Navigate to **Credentials** tab
2. Click “+ Add Credential Set”
3. Fill in the required fields
4. Click “Validate” to test
5. Click “Save”

### 5.3 Credential Validation

The app validates:

- Access key format (AKIA prefix, 20 chars)
- Secret key length (40+ chars)
- Region validity
- AWS connectivity (STS GetCallerIdentity)
- Required permissions

### 5.4 Security Best Practices

Practice	Recommendation
<b>Rotate Keys</b>	Every 90 days
<b>Least Privilege</b>	Use scoped IAM policies
<b>MFA</b>	Enable on AWS account
<b>Audit</b>	Review access logs regularly
<b>Backup</b>	Export credentials securely

## 5.5 Importing from AWS CLI

```
# The app can import from ~/.aws/credentials  
# Navigate to Credentials + Import from AWS CLI
```

---

# 6. Deployment Operations

## 6.1 Deployment Dashboard

The main dashboard shows:

Environment: dev	Status: Healthy
Version: 4.18.1	Last Deploy: 2024-12-25 10:30:00

Deploy [Button]	Rollback [Button]	Settings [Button]
-----------------	-------------------	-------------------

Recent Deployments:

2024-12-25 10:30	v4.18.1	prod	Success	4m 32s
2024-12-24 15:45	v4.18.0	prod	Success	5m 12s
2024-12-24 09:00	v4.17.0	dev	Success	3m 45s

## 6.2 Starting a Deployment

1. Select target **Environment** (dev/staging/prod)
2. Select **Tier** (1-5)
3. Review deployment plan
4. Click “**Start Deployment**”
5. Monitor progress in real-time

## 6.3 Deployment Phases

Phase	Duration	Description
1. Validation	~30s	Credential and configuration check
2. Snapshot	~1m	Create pre-deployment backup
3. CDK Synth	~1m	Generate CloudFormation templates
4. CDK Deploy	~10-20m	Deploy infrastructure
5. Migration	~2m	Run database migrations
6. Health Check	~1m	Verify all services
7. Cleanup	~30s	Remove temporary resources

### Phase Details

#### Phase 1 - Validation (30 seconds)

Before any changes are made, the system validates:

- AWS credentials are valid and not expired
- IAM permissions are sufficient for all required operations
- Target environment exists or can be created
- No conflicting deployments are in progress (deployment lock check)
- Required software versions are installed (Node.js, CDK, etc.)
- Network connectivity to AWS services

If validation fails, you'll see specific error messages explaining what needs to be fixed.

### **Phase 2 - Snapshot (1 minute)**

A pre-deployment snapshot captures the current state so you can rollback if needed:

- Database schema and critical data (not full data backup)
- Current Lambda function code versions
- SSM Parameter Store values
- Current CloudFormation stack states
- Configuration files

Snapshots are stored locally and can be managed in the Snapshots tab.

### **Phase 3 - CDK Synth (1 minute)**

The CDK synthesizes TypeScript infrastructure code into CloudFormation templates:

- Reads infrastructure definitions from `packages/infrastructure/`
- Resolves all construct dependencies
- Generates CloudFormation JSON/YAML templates
- Calculates resource changes (what will be created/updated/deleted)
- Validates templates against AWS CloudFormation rules

You can review the generated templates before proceeding.

### **Phase 4 - CDK Deploy (10-20 minutes)**

The actual AWS resource deployment happens in this phase:

- CloudFormation stacks are created or updated in dependency order
- Resources are provisioned (databases, Lambda functions, API Gateway, etc.)
- IAM roles and policies are configured
- Networking (VPC, subnets, security groups) is set up
- DNS records are created/updated

This is the longest phase. Progress shows which stack is currently deploying.

### **Phase 5 - Migration (2 minutes)**

Database migrations ensure your schema matches the deployed code:

- Connects to Aurora PostgreSQL using Data API
- Runs pending migration files in sequence
- Creates new tables, columns, indexes as needed
- Updates RLS (Row-Level Security) policies
- Verifies migration success with integrity checks

Migrations are idempotent - running them twice won't cause issues.

### **Phase 6 - Health Check (1 minute)**

Verifies all deployed services are functioning:

- API Gateway responds to health endpoint
- Lambda functions can be invoked
- Database connections succeed
- Cognito user pools are accessible
- S3 buckets are reachable
- CloudFront distributions are deployed

If health checks fail, automatic rollback is triggered (if enabled).

### **Phase 7 - Cleanup (30 seconds)**

Final cleanup tasks:

- Removes temporary files created during deployment
- Clears CDK staging buckets of old assets
- Updates deployment history in local database
- Releases deployment lock
- Sends completion notification

## **6.4 Deployment Progress**

Deploying to Production

[ ] 65%

```
Current Phase: CDK Deploy
Stack: Radiant-prod-API (5 of 9)
Elapsed: 8m 23s | Estimated: 4m remaining
```

```
Validation complete
Snapshot created: snap-20241225-103000
CDK synthesis complete
Deploying stacks...
  Radiant-prod-Foundation
  Radiant-prod-Networking
  Radiant-prod-Security
  Radiant-prod-Data
  Radiant-prod-API
  Radiant-prod-Auth
  Radiant-prod-AI
  Radiant-prod-Admin
  Radiant-prod-Monitoring
```

```
[Cancel Deployment]
```

## 6.5 Deployment Settings

Configure deployment behavior:

Setting	Description	Default
<b>Auto-Rollback</b>	Rollback on failure	Enabled
<b>Lock-Step Mode</b>	Require version consistency	Enabled
<b>Max Version Drift</b>	Maximum version difference	1
<b>Approval Required</b>	Require confirmation for prod	Enabled
<b>Notification</b>	Send completion notifications	Enabled

## 6.6 Operation Timeouts

Operation	Default Timeout	Configurable
CDK Deploy	30 minutes	Yes
Health Check	5 minutes	Yes
Migration	10 minutes	Yes
Snapshot	5 minutes	Yes
Rollback	15 minutes	Yes

---

# 7. Multi-Region Deployments

## 7.1 Overview

Deploy RADIANT across multiple AWS regions for:

- **High Availability:** Survive regional outages
- **Low Latency:** Serve users from nearest region
- **Compliance:** Data residency requirements

## 7.2 Supported Regions

Region	Code	Primary Use
US East (N. Virginia)	us-east-1	Primary (required)
US West (Oregon)	us-west-2	West coast users
EU (Ireland)	eu-west-1	European users
EU (Frankfurt)	eu-central-1	GDPR compliance
Asia Pacific (Singapore)	ap-southeast-1	APAC users
Asia Pacific (Tokyo)	ap-northeast-1	Japanese users

## 7.3 Adding a Region

1. Navigate to **Multi-Region** tab
2. Click “**Add Region**”
3. Configure:
  - **Region:** Select from available regions
  - **Is Primary:** Set primary region flag
  - **Stack Prefix:** Region-specific prefix
  - **Endpoint:** Custom domain for region
4. Click “**Deploy to Region**”

## 7.4 Region Consistency

Monitor version consistency across regions:

Multi-Region Status		Consistency: 100%		
Region	Version	Status	Last Deploy	
us-east-1 (P)	4.18.1	Healthy	2024-12-25 10:30	
eu-west-1	4.18.1	Healthy	2024-12-25 10:35	
ap-southeast-1	4.18.1	Healthy	2024-12-25 10:40	

[Deploy All] [Check Consistency] [Sync Versions]

---

## 8. Snapshots & Rollbacks

### 8.1 Snapshot Types

Type	Description	Retention
Pre-Deploy	Automatic before each deployment	30 days
Manual	User-initiated backup	Until deleted
Scheduled	Periodic backups	Configurable

### 8.2 Creating a Snapshot

1. Navigate to **Snapshots** tab

2. Click “Create Snapshot”
3. Enter description (optional)
4. Select components to include:
  - Database state
  - Configuration
  - Lambda code
  - Infrastructure state
5. Click “Create”

### 8.3 Snapshot Contents

```
Snapshot: snap-20241225-103000
  metadata.json          # Snapshot info
  database/
    schema.sql           # Database schema
    data.sql              # Critical data
  config/
    ssm-parameters.json # SSM parameters
    secrets.json         # Secret references
  lambda/
    functions.zip        # Lambda code packages
  infrastructure/
    state.json           # CDK state
```

### 8.4 Restoring from Snapshot

1. Navigate to **Snapshots** tab
2. Select snapshot to restore
3. Click “**Restore**”
4. Confirm restoration scope:
  - Full restoration
  - Database only
  - Configuration only
5. Click “**Confirm Restore**”

### 8.5 Automatic Rollback

When enabled, failed deployments automatically:

1. Stop deployment process
  2. Identify last known good state
  3. Restore from pre-deploy snapshot
  4. Verify system health
  5. Send notification
- 

## 9. AI Assistant

### 9.1 Overview

The Claude-powered AI Assistant provides:

- **Deployment Guidance:** Step-by-step help
- **Error Translation:** User-friendly error explanations
- **Troubleshooting:** Problem diagnosis

- **Best Practices:** Recommendations

## 9.2 Enabling AI Assistant

1. Navigate to **Settings → AI Assistant**
2. Enter Anthropic API key
3. Toggle “**Enable AI Assistant**”
4. Select response style

## 9.3 Using the Assistant

Access the assistant via:

- **Chat Panel:** Click AI icon in toolbar
- **Contextual Help:** Click “?” on any screen
- **Error Dialogs:** Click “Explain” on errors

## 9.4 Sample Interactions

AI Assistant

You: What does error "AccessDenied: User is not authorized to perform dynamodb:CreateTable" mean?

Assistant: This error indicates your AWS credentials don't have permission to create DynamoDB tables. To fix this:

1. Check your IAM user/role has the required permissions
2. Add the "AmazonDynamoDBFullAccess" managed policy, or
3. Add specific permissions:
  - dynamodb:CreateTable
  - dynamodb:DescribeTable
  - dynamodb:UpdateTable

Would you like me to show the exact IAM policy needed?

[Yes, show policy] [No, I'll figure it out]

## 9.5 Offline Mode

When offline, the assistant provides:

- Pre-cached common error explanations
- Local troubleshooting guides
- Fallback recommendations

## 10. Package Management

### 10.1 Package System

RADIANT uses atomic packages for deployment:

```
radiant-4.18.1.pkg
  manifest.json          # Package manifest
  checksums.sha256        # Component checksums
  radiant/
    infrastructure/
    lambda/
    dashboard/
  thinktank/              # Think Tank components
    api/
    frontend/
```

### 10.2 Viewing Packages

Navigate to **Packages** tab to see:

- Installed packages
- Available updates
- Package history
- Component versions

### 10.3 Version Management

Version Type	Format	Example
Radiant	Major.Minor.Patch	4.18.1
Think Tank	Major.Minor.Patch	3.2.0
Package	Combined	4.18.1+3.2.0

### 10.4 Lock-Step Mode

When enabled:

- All components must have same minor version
- Maximum version drift configurable
- Automatic sync available

---

## 11. Monitoring & Health Checks

### 11.1 Health Dashboard

System Health   Overall: Healthy

Service	Status	Latency	Last Check
API Gateway	Up	45ms	10s ago
Lambda (Router)	Up	120ms	10s ago

Aurora PostgreSQL	Up	12ms	10s ago
DynamoDB	Up	8ms	10s ago
Cognito	Up	85ms	10s ago
S3 Storage	Up	35ms	10s ago
CloudFront	Up	22ms	10s ago
SageMaker (if T3+)	Up	250ms	10s ago

[Refresh] [Run Full Check] [View History]

## 11.2 Health Check Types

Check	Frequency	Timeout
<b>Quick</b>	Every 60s	5s
<b>Standard</b>	Every 5m	30s
<b>Deep</b>	Manual/Deploy	2m

## 11.3 Alerts

Configure alerts for:

- Service degradation
- High latency
- Error rate spikes
- Failed deployments

# 12. Security Features

## 12.1 Credential Security

Feature	Implementation
<b>Storage</b>	macOS Keychain (encrypted)
<b>Memory</b>	Cleared after use
<b>Transport</b>	TLS 1.3 only
<b>Validation</b>	Format + connectivity check

## 12.2 Deployment Locks

Prevent concurrent deployments:

Deployment Lock: Active  
 Acquired: 2024-12-25 10:30:00  
 Owner: deployer@example.com  
 Environment: production  
 Expires: 2024-12-25 11:30:00

## 12.3 Audit Logging

All operations are logged:

```
{  
  "timestamp": "2024-12-25T10:30:00Z",  
  "operation": "deployment.start",  
  "user": "admin@example.com",  
  "environment": "production",  
  "version": "4.18.1",  
  "status": "success",  
  "duration_ms": 272000  
}
```

## 12.4 Secret Detection

Pre-commit checks scan for:

- AWS access keys
  - API keys
  - Passwords
  - Private keys
- 

# 13. Troubleshooting

## 13.1 Common Issues

### Deployment Fails at CDK Synth

**Symptoms:** Deployment stops at synthesis phase

**Solutions:** 1. Check Node.js version: `node --version` (need 20.x) 2. Clear CDK cache: `rm -rf cdk.out` 3. Update CDK: `npm update -g aws-cdk` 4. Check TypeScript errors in `packages/infrastructure`

### AWS Credentials Invalid

**Symptoms:** “Invalid credentials” error

**Solutions:** 1. Verify access key format (starts with AKIA) 2. Check secret key hasn’t expired 3. Verify IAM user is active 4. Test with AWS CLI: `aws sts get-caller-identity`

### Health Check Timeout

**Symptoms:** Services show unhealthy after deployment

**Solutions:** 1. Wait 2-3 minutes for cold start 2. Check CloudWatch logs for errors 3. Verify security group rules 4. Check VPC endpoint configuration

## 13.2 Log Locations

Log Type	Location
<b>App Logs</b>	<code>~/Library/Logs/RadiantDeployer/</code>
<b>Deployment Logs</b>	<code>~/Library/Application Support/RadiantDeployer/deployments/</code>
<b>AWS Logs</b>	<code>CloudWatch Log Groups</code>

### 13.3 Getting Help

1. **AI Assistant:** Built-in help
  2. **Documentation:** This guide + online docs
  3. **Support:** support@radianit.example.com
- 

## 14. Reference

### 14.1 Keyboard Shortcuts

Shortcut	Action
+ D	Start deployment
+ R	Refresh status
+ S	Create snapshot
+ ,	Open settings
+ ?	Open AI assistant
+ L	View logs

### 14.2 CLI Commands

```
# Build and run from source
cd apps/swift-deployer
swift build -c release
swift run RadiantDeployer

# Run with specific config
swift run RadiantDeployer --environment prod --tier 3

# Headless deployment
swift run RadiantDeployer deploy --non-interactive
```

### 14.3 Environment Variables

Variable	Description
RADIANT_ENV	Override environment
RADIANT_TIER	Override tier
RADIANT_DEBUG	Enable debug logging
RADIANT_AI_KEY	Anthropic API key

### 14.4 File Locations

File	Location
Configuration	~/Library/Application Support/RadiantDeployer/config.json
Snapshots	~/Library/Application Support/RadiantDeployer/snapshots/
Logs	~/Library/Logs/RadiantDeployer/
Database	~/Library/Application Support/RadiantDeployer/local.db

---

## Appendix A: IAM Policy Requirements

Minimum IAM permissions for deployment:

```
{  
  "Version": "2012-10-17",  
  "Statement": [  
    {  
      "Effect": "Allow",  
      "Action": [  
        "cloudformation:*",  
        "s3:*",  
        "lambda:*",  
        "apigateway:*",  
        "cognito-idp:*",  
        "rds:*",  
        "dynamodb:*",  
        "sns:*",  
        "sns:*",  
        "events:*",  
        "logs:*",  
        "iam:PassRole",  
        "iam:CreateRole",  
        "iam:AttachRolePolicy",  
        "ssm:*",  
        "secretsmanager:*",  
        "ecr:*",  
        "ecs:*"  
      ],  
      "Resource": "*"  
    }  
  ]  
}
```

---

## Appendix B: Glossary

Term	Definition
<b>CDK</b>	AWS Cloud Development Kit
<b>Stack</b>	CloudFormation stack deployed by CDK
<b>Snapshot</b>	Point-in-time backup of deployment
<b>Lock-Step</b>	Version consistency enforcement
<b>Tier</b>	Infrastructure sizing level (1-5)

---

*Document Version: 4.18.1 Last Updated: December 2024*

# RADIANT Deployer Architecture & Deployment Packages

## Technical Architecture Document

Version: 4.18.1 | Last Updated: December 2024

---

## Overview

The RadiantDeployer Swift app operates in three distinct modes, each with different behaviors for parameter handling, package selection, and database operations.

## Deployment Modes

### Mode Definitions

#### DEPLOYER OPERATIONAL MODES

INSTALL (Fresh)	UPDATE (Upgrade)	ROLLBACK (Revert)
Use Default Parameters	Read Current From Instance	Read Target Snapshot
Seed AI Registry	Merge User Changes	Restore Previous
Create New Instance	Apply Delta Changes	Apply Snapshot

## INSTALL Mode (Fresh Installation)

**Trigger:** No existing deployment detected for app/environment combination

**Key Behaviors:** 1. Uses DEFAULT parameters from tier configuration 2. Runs ALL database migrations (fresh) 3. SEEDS the AI Registry with providers and models 4. Creates initial admin user 5. Stores deployment metadata

**Parameter Source:** InstallationParameters.defaults()

```
// Parameters are initialized with tier-appropriate defaults
let parameters = InstallationParameters.defaults(
  appId: app.id,
  environment: environment,
  tier: .growth // Based on selected tier
)
```

## UPDATE Mode (Upgrade Existing)

**Trigger:** Existing deployment detected AND target version >= current version

**Key Behaviors:** 1. Fetches current parameters FROM the running instance 2. Creates pre-update snapshot for rollback 3. MERGES user changes with current parameters 4. Validates parameter changes are safe 5. Runs INCREMENTAL migrations only 6. **DOES NOT** seed AI Registry (preserves admin customizations)

**Parameter Source:** Running instance API + user modifications

```
// Parameters fetched from instance, then merged with user changes
let currentParameters = await fetchCurrentParameters(app, environment, credentials)
let updatedParameters = mergeParameters(current: currentParameters, changes: userChanges)
```

## ROLLBACK Mode (Revert to Previous)

**Trigger:** User explicitly requests rollback OR target version < current version

**Key Behaviors:** 1. Loads target snapshot from S3 2. Creates safety snapshot of current state 3. Deploys with SNAPSHOT parameters (not current, not defaults) 4. Optionally restores database from RDS snapshot 5. Does not modify AI Registry

**Parameter Source:** Selected snapshot

---

## Deployment Package Structure

Deployment packages are self-contained, versioned bundles containing everything needed to deploy a specific version of RADIANT.

```
radiant-4.18.0-abc123.radpkg
  manifest.json           # Package metadata & verification
  checksums.sha256         # File integrity verification

  infrastructure/          # CDK Stacks (compiled)
    cdk.out/                # Synthesized CloudFormation
    lib/                     # CDK TypeScript (compiled)
```

```

cdk.json          # CDK configuration

migrations/      # Database migrations
  radiant/        # Core schema migrations
  thinktank/      # Think Tank specific
  seeds/          # Seed data (AI Registry, etc.)

functions/       # Lambda function code
  api/            # API handlers
  admin/          # Admin handlers
  billing/        # Billing handlers
  thermal/        # Thermal management

admin-dashboard/ # Next.js admin dashboard
  .next/          # Compiled Next.js

config/          # Default configurations
  defaults.json   # Default parameters per tier
  providers.json  # AI provider seed data
  models.json     # AI model seed data

```

## Package Manifest

```
{
  "packageFormat": "radpkg-v1",
  "version": "4.18.0",
  "buildId": "abc123def456",
  "buildTimestamp": "2024-12-24T10:30:00Z",

  "components": {
    "radiantPlatform": {
      "version": "4.18.0",
      "minUpgradeFrom": "4.15.0"
    },
    "thinkTank": {
      "version": "3.2.0",
      "minUpgradeFrom": "3.0.0"
    }
  },

  "compatibility": {
    "minimumDeployerVersion": "4.16.0",
    "supportedTiers": ["SEED", "STARTER", "GROWTH", "SCALE", "ENTERPRISE"],
    "supportedRegions": ["us-east-1", "us-west-2", "eu-west-1", "ap-southeast-1"]
  },

  "installBehavior": {
    "seedAIRegistry": true,
    "createInitialAdmin": true,
    "runFullMigrations": true
  },

  "updateBehavior": {
    "seedAIRegistry": false,

```

```

    "preserveAdminCustomizations": true,
    "runIncrementalMigrations": true,
    "createPreUpdateSnapshot": true
}
}

```

---

## Package Storage Locations

1. DEPLOYER APP CACHE (Local)  
`~/Library/Application Support/RadiantDeployer/packages/  
 radiant-4.18.0-abc123.radpkg  
 radiant-4.17.0-def456.radpkg  
 index.json`
  2. S3 RELEASE BUCKET (Cloud - Official Releases)  
`s3://radiant-releases-{region}/  
 stable/  
 radiant-4.18.0-abc123.radpkg  
 latest.json  
 beta/  
 radiant-4.19.0-beta1-xyz789.radpkg  
 archive/  
 radiant-4.17.0-def456.radpkg`
  3. DEPLOYED INSTANCE (Cloud - Per Instance)  
`s3://radiant-{appId}-{env}-deployments/  
 current/  
 radiant-4.18.0-abc123.radpkg  
 snapshots/  
 snapshot-2024-12-24T10-30-00Z/  
 package.radpkg  
 parameters.json  
 db-snapshot-id.txt  
 ...`
- 

## Key Implementation Files

### Swift Deployer

File	Purpose
Models/InstallationParameters.swift	DeploymentMode enum, TierLevel, InstallationParameters, InstanceParameters, ParameterChanges, DeploymentSnapshot
Services/DeploymentService.swift	Mode detection, executeInstall, executeUpdate, executeRollback, parameter fetching/merging
Services/PackageService.swift	Package discovery, download, verification, caching
Views/Deployment/ParameterEditorView.swift	Editing parameters based on mode

## Build Tools

File	Purpose
tools/scripts/build-package.sh	Build deployment packages from source
tools/version-manager.ts	Version bumping and synchronization

## Data Flow Diagrams

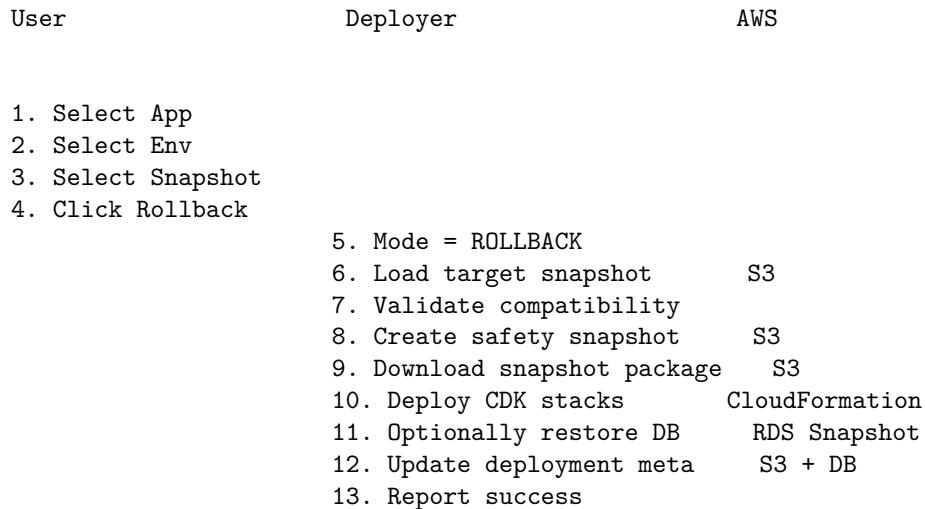
### Install Flow

User	Deployer	AWS
1. Select App		
2. Select Env		
3. Select Tier		
4. Click Deploy		
	5. Check instance exists	(None found)
	6. Mode = INSTALL	
	7. Load DEFAULT params	
	8. Download latest package	S3
	9. Verify package integrity	
	10. Deploy CDK stacks	CloudFormation
	11. Run ALL migrations	Aurora
	12. SEED AI Registry	Aurora
	13. Create initial admin	Cognito
	14. Store deployment meta	S3 + DB
	15. Report success	

### Update Flow

User	Deployer	AWS + Instance
1. Select App		
2. Select Env		
3. Change Params		
4. Click Update		
	5. Check instance exists	(Found!)
	6. Mode = UPDATE	
	7. Fetch CURRENT params	Radiant API
	8. Create snapshot	S3
	9. MERGE user changes	
	10. Validate changes	
	11. Download target package	S3
	12. Deploy CDK stacks	CloudFormation
	13. Run INCREMENTAL migrations	Aurora
	14. SKIP AI Registry seeding	
	15. Update deployment meta	S3 + DB
	16. Report success	

## Rollback Flow



## Verification Checklist

### Deployment Modes

- **On INSTALL:** Parameters come from defaults
- **On UPDATE:** Parameters come from running instance + user changes
- **On ROLLBACK:** Parameters come from selected snapshot

### AI Registry Seeding

- AI Registry is seeded ONLY on fresh install
- On UPDATE, AI Registry is preserved (not touched)
- Admins can add/remove providers via Admin Dashboard

### Deployment Packages

- Packages are created by build-package.sh script
- Package creation is triggered by code changes or version bumps
- Packages are stored in local cache, S3 release bucket, and instance bucket

### Parameter Rules

- Region CANNOT be changed after install
  - Tier CAN be changed on update (with feature validation)
  - All parameter changes are tracked via snapshots
- 

## Related Documentation

- [Deployer Admin Guide](#) - User-facing documentation
- [Deployment Guide](#) - Deployment procedures
- [API Reference](#) - API documentation

# Think Tank AI - User Guide

**Your gateway to 100+ AI models in one place**

Version: 3.2.0 (Platform: RADIANT 4.18.1) Last Updated: December 2024

---

## Welcome to Think Tank

Think Tank is your all-in-one AI assistant platform. Access the world's best AI models—GPT-4, Claude, Gemini, and 100+ more—from a single, beautiful interface.

---

## Table of Contents

1. Getting Started
  2. Your Dashboard
  3. Chatting with AI
  4. Choosing Models
  5. Focus Modes & Personas
  6. Canvas & Artifacts
  7. Collaboration Features
  8. Managing Your Account
  9. Credits & Billing
  10. Tips & Best Practices
  11. Keyboard Shortcuts
  12. FAQ
  13. Delight System
- 

## 1. Getting Started

### 1.1 Creating Your Account

1. Visit [thinktank.ai](https://thinktank.ai)
2. Click **Get Started Free**
3. Sign up with:
  - Email and password
  - Google account
  - Microsoft account
  - Apple ID

4. Verify your email
5. Complete your profile

## 1.2 Choosing a Plan

Plan	Price	Best For
<b>Free</b>	\$0/month	Trying out Think Tank
<b>Starter</b>	\$29/month	Individual creators
<b>Pro</b>	\$99/month	Power users
<b>Team</b>	\$49/user/month	Small teams
<b>Business</b>	\$199/user/month	Organizations

## 1.3 Your First Chat

1. Click **New Chat** or press **Ctrl+N**
2. Type your question or request
3. Press **Enter** or click **Send**
4. Watch as AI responds in real-time

**Try these starter prompts:** - “Explain quantum computing like I’m 10 years old” - “Write a professional email declining a meeting” - “Help me debug this Python code: [paste code]” - “Create a meal plan for the week”

---

## 2. Your Dashboard

### 2.1 Interface Overview

The screenshot shows the Think Tank dashboard. At the top left is the brand name "Think Tank". To its right are two buttons: "[Search]" and "[Credits: 450]". Below the header, there's a sidebar on the left with sections for "Chats" and "Personas". Under "Chats", it lists "Starred", "Today", "Chat 1", "Chat 2", "Yesterday", and "Chat 3". Under "Personas", it lists "Canvas" and "Settings". The main area is titled "Chat Area" and contains a message from the AI: "Welcome! How can I help you today?". Below the message input field, it says "Type your message...". To the right of the input field is a "[Send]" button. At the bottom of the main area, it shows "Model: GPT-4 Turbo | Focus: General |".

## 2.2 Sidebar Navigation

Icon	Section	Description
	Chats	All your conversations
	Starred	Important chats
	Personas	Custom AI personalities
	Canvas	Visual workspace
	Usage	Credit usage stats
	Settings	Account settings

## 2.3 Quick Actions

Action	Shortcut
New Chat	<b>Ctrl+N</b>
Search	<b>Ctrl+K</b>
Toggle Sidebar	<b>Ctrl+B</b>
Settings	<b>Ctrl+,</b>

## 3. Chatting with AI

### 3.1 Sending Messages

**Text Messages:** - Type in the input box - Press **Enter** to send - Use **Shift+Enter** for new lines

**Attachments:** - Click to attach files - Drag & drop images, PDFs, code files - Paste images directly (**Ctrl+V**)

**Voice Input:** - Click microphone icon - Speak your message - Click again to stop

### 3.2 Message Actions

Hover over any message to see actions:

Icon	Action	Description
	Copy	Copy message text
/	Regenerate	Get a new response
	Edit	Modify your message
	Rate	Help improve AI
	Pin	Keep message visible
	Delete	Remove message

### 3.3 Streaming Responses

AI responses stream in real-time. You can: - **Stop:** Click to stop generation - **Continue:** Ask “continue” if response was cut off - **Regenerate:** Get a different response

### 3.4 Multi-Turn Conversations

Think Tank remembers your conversation context:

You: I'm planning a trip to Japan

AI: That's exciting! When are you planning to visit...

You: What about the food?

AI: Japanese cuisine is incredible! Based on your trip...

[AI remembers you're going to Japan]

### 3.5 Code in Chats

Code is automatically syntax-highlighted:

```
def hello_world():
    print("Hello, Think Tank!")
```

Click **Copy** to copy code blocks, or **Run** for supported languages.

---

## 4. Choosing Models

### 4.1 Model Selection

Click the model selector at the bottom of the chat:

Select Model

Favorites

GPT-4 Turbo	\$0.02/msg
Claude 3 Opus	\$0.03/msg

Recommended

GPT-4o	\$0.01/msg
Claude 3.5 Sonnet	\$0.01/msg
Gemini 1.5 Pro	\$0.01/msg

Writing

Coding

Analysis

Creative

[View All 100+ Models]

### 4.2 Model Categories

Category	Best For	Top Models
<b>General</b>	Everyday tasks	GPT-4o, Claude 3.5
<b>Writing</b>	Content creation	Claude 3 Opus, GPT-4
<b>Coding</b>	Programming help	GPT-4 Turbo, CodeLlama
<b>Analysis</b>	Data & research	Gemini 1.5, Claude 3
<b>Creative</b>	Art & ideas	GPT-4, Mistral Large
<b>Vision</b>	Image understanding	GPT-4V, LLaVA
<b>Fast</b>	Quick responses	GPT-3.5, Claude Instant

## Choosing the Right Model

**For everyday questions and tasks:** Start with GPT-4o or Claude 3.5 Sonnet. These models offer the best balance of quality, speed, and cost. They handle most tasks excellently including writing, answering questions, brainstorming, and light coding.

**For professional writing:** Claude 3 Opus excels at long-form content, maintaining consistent tone, and nuanced writing. GPT-4 is also excellent for business documents and creative writing.

**For coding and technical work:** GPT-4 Turbo has strong coding abilities across many languages. For specialized tasks, consider CodeLlama (open source, good for common languages) or specialized models like DeepSeek Coder.

**For data analysis:** Gemini 1.5 Pro handles very long documents (up to 1 million tokens) making it ideal for analyzing large datasets or documents. Claude 3 is excellent for nuanced analytical reasoning.

**For creative projects:** GPT-4 and Mistral Large are both creative and can help with brainstorming, storytelling, and idea generation. They're less constrained in creative contexts.

**For image understanding:** GPT-4V (Vision) and Claude 3 Vision can analyze images, read text from photos, describe scenes, and answer questions about visual content.

**For quick, simple tasks:** GPT-3.5 Turbo and Claude Instant are much faster and cheaper. Use them for simple questions, formatting, or when you need instant responses.

## 4.3 Auto Mode

Let Think Tank choose the best model:

1. Enable **Auto Mode** in settings
2. Our Brain Router analyzes your request
3. Automatically selects optimal model
4. Balances quality, speed, and cost

### How Auto Mode Works

When you enable Auto Mode, Think Tank's Brain Router analyzes each message you send and selects the best model based on:

- **Task complexity:** Simple questions go to fast models; complex tasks go to powerful models
- **Content type:** Coding questions route to code-specialized models; creative requests to creative models
- **Your history:** Learns your preferences over time and adjusts recommendations
- **Cost efficiency:** Avoids using expensive models when cheaper ones would work equally well
- **Current availability:** Routes around any models experiencing slowdowns

**When to use Auto Mode:** - You're not sure which model to use - You want to optimize cost without sacrificing quality - You have varied tasks throughout the day - You're new to Think Tank

**When to choose manually:** - You need a specific model's unique capabilities - You're doing specialized work (e.g., always want Claude for writing) - You're comparing models intentionally - You have strong preferences for certain models

## 4.4 Model Comparison

Split-screen to compare models:

1. Click **Compare** button
2. Select 2-4 models
3. Send message to all simultaneously
4. See responses side-by-side

---

## 5. Focus Modes & Personas

### 5.1 Focus Modes

Pre-configured modes for specific tasks:

Mode	Optimized For
<b>Professional</b>	Business writing, emails
<b>Developer</b>	Code, debugging, architecture
<b>Research</b>	Analysis, citations, accuracy
<b>Creative</b>	Stories, brainstorming
<b>Learning</b>	Explanations, tutoring
<b>Concise</b>	Brief, direct answers

**To switch modes:** 1. Click the Focus selector 2. Choose your mode 3. AI adapts its style

#### Focus Mode Details

**Professional Mode:** The AI adopts a business-appropriate tone. Responses are polished, formal, and suitable for workplace communication. Great for drafting emails, reports, presentations, and client communications. Avoids casual language and ensures professional formatting.

**Developer Mode:** Optimized for technical work. The AI provides code with proper syntax highlighting, explains technical concepts clearly, suggests best practices, and can help debug issues. Responses include code comments and consider edge cases.

**Research Mode:** Emphasizes accuracy and thoroughness. The AI cites sources when possible, acknowledges uncertainty, presents multiple perspectives, and structures information logically. Ideal for academic work, fact-checking, and deep analysis.

**Creative Mode:** Removes constraints on creativity. The AI is more willing to explore unusual ideas, use vivid language, and think outside the box. Perfect for brainstorming, creative writing, storytelling, and generating innovative solutions.

**Learning Mode:** The AI becomes a patient tutor. Explanations start from basics and build up, concepts are broken into digestible pieces, and the AI checks understanding before moving on. Great for studying new topics.

**Concise Mode:** Responses are brief and to the point. The AI avoids lengthy explanations and gets straight to the answer. Useful when you need quick facts or are in a hurry.

### 5.2 Custom Personas

Create your own AI personalities:

1. Go to **Personas** → **Create New**
2. Configure:
  - **Name:** “Marketing Expert”
  - **Personality:** Professional, enthusiastic
  - **Expertise:** Digital marketing, SEO
  - **Style:** Uses bullet points, data-driven
3. Click **Save**

**Example Persona:**

Name: Code Reviewer  
Personality: Thorough, constructive  
Instructions: Review code for bugs, security issues, and best practices. Always suggest improvements with examples.

### 5.3 Sharing Personas

- **Public:** Share with all Think Tank users
  - **Team:** Share within your organization
  - **Private:** Only you can use
- 

## 6. Canvas & Artifacts

### 6.1 What is Canvas?

Canvas is your visual workspace for complex outputs: - Code files with syntax highlighting - Diagrams and flowcharts - Documents and reports - Data tables - Mind maps

### 6.2 Creating Artifacts

When AI generates complex content, it appears as an artifact:

```
business_plan.md

# Business Plan

## Executive Summary
...

[Copy] [Download] [Edit] [Version History]
```

### 6.3 Artifact Actions

Action	Description
<b>Copy</b>	Copy content to clipboard
<b>Download</b>	Save as file
<b>Edit</b>	Modify directly
<b>Versions</b>	View previous versions
<b>Share</b>	Generate share link
<b>To Canvas</b>	Open in full Canvas view

### 6.4 Full Canvas Mode

For larger projects:

1. Click **Canvas** in sidebar
2. Create new canvas or open existing
3. Add multiple artifacts

4. Arrange spatially
  5. Connect related items
- 

## 7. Collaboration Features

### 7.1 Sharing Chats

Share any conversation:

1. Click **Share ( )** on a chat
2. Choose visibility:
  - **Link:** Anyone with link
  - **Team:** Your organization
  - **Private:** Specific people
3. Copy and share the link

### 7.2 Real-Time Collaboration

Work together on the same chat:

1. Share chat with **Edit** access
2. Multiple users can:
  - Send messages simultaneously
  - See each other's cursors
  - React to messages
3. Changes sync in real-time

### 7.3 Team Workspaces

For Team and Business plans:

- **Shared Chats:** Team-visible conversations
- **Shared Personas:** Team AI configurations
- **Shared Canvas:** Collaborative workspaces
- **Usage Dashboard:** Team analytics

### 7.4 Comments & Annotations

Add notes to any message:

1. Hover over message
  2. Click **Comment ( )**
  3. Add your note
  4. Tag teammates with @mention
- 

## 8. Managing Your Account

### 8.1 Profile Settings

Access via **Settings → Profile**:

Setting	Description
Display Name	Your visible name
Email	Login email
Avatar	Profile picture
Language	Interface language
Timezone	For scheduled features

## 8.2 Preferences

Customize your experience:

Preference	Options
Theme	Light, Dark, System
Font Size	Small, Medium, Large
Default Model	Your preferred model
Auto Mode	Enable/disable
Sound Effects	On/Off
Notifications	Email, Push, None

## 8.3 Data & Privacy

Control your data:

- **Export Data:** Download all your chats
- **Delete History:** Remove chat history
- **Training Opt-Out:** Exclude from AI training
- **Data Retention:** Set auto-delete period

## 8.4 Connected Apps

Manage integrations:

- Google Drive
- Dropbox
- Notion
- Slack
- GitHub

---

## 9. Credits & Billing

### 9.1 Understanding Credits

Credits are Think Tank's universal currency:

Credit Value	Equivalent
1 credit	\$0.01
100 credits	\$1.00

## 9.2 Credit Usage

Different models cost different amounts:

Model	~Cost per Message
GPT-3.5 Turbo	0.5 credits
GPT-4o	1-2 credits
GPT-4 Turbo	2-3 credits
Claude 3.5 Sonnet	1-2 credits
Claude 3 Opus	3-5 credits

*Actual cost depends on message length*

## 9.3 Viewing Usage

Check your usage in **Settings → Usage**:

Credit Usage - December 2024

Balance: 450 credits  
Used this month: 1,550 credits  
Included: 2,000 credits

[ ] 77% used

By Model:  
GPT-4 Turbo 800 credits (52%)  
Claude 3 500 credits (32%)  
Other 250 credits (16%)

## 9.4 Purchasing Credits

Need more credits?

1. Go to **Settings → Billing**
2. Click **Buy Credits**
3. Select amount:
  - 500 credits - \$5
  - 1,000 credits - \$9 (10% bonus)
  - 5,000 credits - \$40 (20% bonus)
4. Complete payment

## 9.5 Subscription Management

Manage your plan:

- **Upgrade:** Get more features and credits
- **Downgrade:** Switch to lower tier (end of period)
- **Cancel:** Cancel subscription (keep access until end)
- **Invoices:** Download billing history

## 10. Tips & Best Practices

### 10.1 Writing Better Prompts

Be Specific:

```
"Write about dogs"  
"Write a 200-word blog post about training golden  
    retriever puppies, focusing on positive reinforcement"
```

Provide Context:

```
"Fix this code"  
"Fix this Python code that should sort a list but  
    throws an IndexError: [paste code]"
```

Set the Format:

```
"Give me ideas"  
"Give me 5 blog post ideas about sustainable living,  
    formatted as bullet points with a brief description"
```

### 10.2 Getting Better Results

Technique	Example
<b>Chain of thought</b>	"Think step by step..."
<b>Role assignment</b>	"Act as a senior developer..."
<b>Examples</b>	"Here's an example of what I want..."
<b>Constraints</b>	"In 100 words or less..."
<b>Iteration</b>	"Good, but make it more formal"

### 10.3 Saving Credits

- Use **Auto Mode** for optimal model selection
- Use **GPT-3.5** for simple tasks
- Be concise in your prompts
- Avoid regenerating unnecessarily
- Use **Focus Modes** for specialized tasks

### 10.4 Organizing Chats

- Star important conversations
- **Folders:** Group related chats
- **Tags:** Add searchable labels
- **Search:** Find any past conversation

---

## 11. Keyboard Shortcuts

### 11.1 General

Shortcut	Action
Ctrl+N	New chat
Ctrl+K	Search

Shortcut	Action
<b>Ctrl+B</b>	Toggle sidebar
<b>Ctrl+,</b>	Settings
<b>Ctrl+/-</b>	Show shortcuts
<b>Escape</b>	Close modal

## 11.2 Chat

Shortcut	Action
<b>Enter</b>	Send message
<b>Shift+Enter</b>	New line
<b>Ctrl+↑</b>	Edit last message
<b>Ctrl+Shift+C</b>	Copy last response
<b>Ctrl+Shift+R</b>	Regenerate
<b>Ctrl+.</b>	Stop generation

## 11.3 Navigation

Shortcut	Action
<b>Alt+↑/↓</b>	Previous/next chat
<b>Ctrl+1-9</b>	Switch to chat 1-9
<b>Ctrl+Tab</b>	Cycle tabs

---

## 12. FAQ

### Getting Started

**Q: Is Think Tank free?** A: Yes! The Free plan includes 50 credits/month. Upgrade for more credits and features.

**Q: Which AI model should I use?** A: Enable Auto Mode and let us choose, or: - General tasks → GPT-4o or Claude 3.5 Sonnet - Complex analysis → GPT-4 Turbo or Claude 3 Opus - Quick answers → GPT-3.5 Turbo

**Q: Can I use Think Tank on mobile?** A: Yes! Visit [thinktank.ai](https://thinktank.ai) on any mobile browser, or download our iOS/Android apps.

### Credits & Billing

**Q: What happens when I run out of credits?** A: You can purchase more credits or wait for your monthly refresh (paid plans).

**Q: Do unused credits roll over?** A: Monthly included credits expire. Purchased credits never expire.

**Q: Can I get a refund?** A: Contact support within 14 days for subscription refunds.

## Privacy & Security

**Q: Is my data used to train AI?** A: By default, no. You can verify in Settings → Privacy.

**Q: Who can see my chats?** A: Only you, unless you explicitly share them.

**Q: Is my data encrypted?** A: Yes, with AES-256 encryption at rest and TLS 1.3 in transit.

## Troubleshooting

**Q: Why is the AI response slow?** A: Complex queries or busy times may cause delays. Try a faster model.

**Q: Why did my response get cut off?** A: Models have output limits. Type “continue” to get the rest.

**Q: I found a bug. How do I report it?** A: Click **Help** → **Report Issue** or email support@thinktank.ai.

---

## Need Help?

- **Help Center:** [help.thinktank.ai](http://help.thinktank.ai)
  - **Live Chat:** Click the chat bubble
  - **Email:** [support@thinktank.ai](mailto:support@thinktank.ai)
  - **Twitter:** @ThinkTankAI
  - **Discord:** [discord.gg/thinktank](https://discord.gg/thinktank)
- 

## Version History

Version	Date	Changes
3.2.0	December 2024	Time Machine, enhanced collaboration, A/B experiments, <b>Delight System</b>
3.1.0	November 2024	Canvas improvements, new models
3.0.0	October 2024	Initial release

---

## 13. Delight System

Think Tank includes a personality system called “Delight” that makes your AI experience more engaging.

### 13.1 What is Delight?

Delight adds contextual, friendly messages during your conversations:

- **Domain Loading:** “Consulting the fundamental forces...” when you ask physics questions
- **Time Awareness:** “Burning the midnight tokens” during late-night sessions
- **Model Dynamics:** “Consensus forming...” when multiple models agree
- **Wellbeing Nudges:** “You’ve been thinking hard. Time for a break?”

## 13.2 Personality Modes

Choose your preferred personality style in **Settings → Delight**:

Mode	Description
<b>Professional</b>	Minimal, business-appropriate feedback
<b>Subtle</b>	Light touches of personality
<b>Expressive</b>	Full personality with humor
<b>Playful</b>	Maximum fun, includes easter eggs

## 13.3 Achievements

Earn achievements as you use Think Tank:

Achievement	How to Unlock
Domain Explorer	Explore 10+ knowledge domains
Week Warrior	Use Think Tank 7 days in a row
Renaissance Mind	Explore 50+ domains
Monthly Mind	30-day streak

View your achievements in **Settings → Achievements**.

## 13.4 Easter Eggs

Think Tank has hidden surprises! Try: - Typing special phrases - Using keyboard shortcuts - Exploring during special times

Discover them yourself—that's half the fun!

## 13.5 Sound Effects

Enable audio feedback in **Settings → Delight → Sounds**:

Theme	Style
Default	Pleasant chimes
Mission Control	NASA-inspired beeps
Library	Page turns, book sounds
Workshop	Tool clicks
Emissions	Tesla-style... fun

## 13.6 Customizing Delight

Toggle individual features: - Domain messages - Model personality - Time awareness - Achievements - Wellbeing nudges - Easter eggs - Sound effects

Set intensity level (1-10) to control how often messages appear.

---

*Thank you for using Think Tank! We're constantly improving based on your feedback.*

© 2024 Think Tank AI. All rights reserved.

# RADIANT API Reference

Complete API documentation for the RADIANT AI Platform.

**Base URL:** `https://api.radiant.example.com/v2`

**Authentication:** Bearer token (API key or JWT)

**Authorization:** `Bearer rad_your_api_key`

---

## Chat Completions

### Create Chat Completion

`POST /v2/chat/completions`

Create a chat completion with any supported model.

**Request Body:**

Field	Type	Required	Description
<code>model</code>	string		Model ID (e.g., <code>gpt-4o</code> , <code>claude-3-sonnet</code> )
<code>messages</code>	array		Array of message objects
<code>max_tokens</code>	integer		Maximum tokens to generate
<code>temperature</code>	number		Sampling temperature (0-2)
<code>top_p</code>	number		Nucleus sampling (0-1)
<code>stream</code>	boolean		Stream response tokens
<code>stop</code>	string/array		Stop sequences
<code>functions</code>	array		Function definitions
<code>function_call</code>	string/object		Function calling mode

**Message Object:**

Field	Type	Required	Description
<code>role</code>	string		<code>system</code> , <code>user</code> , <code>assistant</code> , <code>function</code>
<code>content</code>	string		Message content
<code>name</code>	string		Function name (for function messages)

**Example Request:**

```
{  
  "model": "gpt-4o",  
  "messages": [  
    {"role": "user", "content": "Hello!"},  
    {"role": "assistant", "content": "Hello!"}  
  ]}
```

```

    "messages": [
        {"role": "system", "content": "You are a helpful assistant."},
        {"role": "user", "content": "Hello!"}
    ],
    "max_tokens": 1000,
    "temperature": 0.7
}

```

**Response:**

```

{
    "id": "chatcmpl_abc123",
    "object": "chat.completion",
    "created": 1703980800,
    "model": "gpt-4o",
    "choices": [
        {
            "index": 0,
            "message": {
                "role": "assistant",
                "content": "Hello! How can I help you today?"
            },
            "finish_reason": "stop"
        }
    ],
    "usage": {
        "prompt_tokens": 25,
        "completion_tokens": 10,
        "total_tokens": 35
    }
}

```

---

## Models

### List Models

GET /v2/models

List all available models.

**Query Parameters:**

Parameter	Type	Description
category	string	Filter by category (chat, embedding, image)
provider	string	Filter by provider (openai, anthropic, google)

**Response:**

```
{
    "object": "list",
    "data": [

```

```

    {
      "id": "gpt-4o",
      "object": "model",
      "created": 1703980800,
      "owned_by": "openai",
      "display_name": "GPT-4o",
      "category": "chat",
      "context_window": 128000,
      "input_cost_per_1k": 0.005,
      "output_cost_per_1k": 0.015,
      "capabilities": ["chat", "vision", "function_calling"]
    }
  ]
}

```

## Get Model

GET /v2/models/{model\_id}

Get details for a specific model.

---

## Embeddings

### Create Embeddings

POST /v2/embeddings

Generate embeddings for text.

**Request Body:**

Field	Type	Required	Description
model	string		Embedding model ID
input	string/array		Text to embed
encoding_format	string		float or base64

**Response:**

```

{
  "object": "list",
  "data": [
    {
      "object": "embedding",
      "index": 0,
      "embedding": [0.0023, -0.0091, ...]
    }
  ],
  "model": "text-embedding-3-small",
  "usage": {
    "prompt_tokens": 8,
    "total_tokens": 8
  }
}

```

---

## Billing

### Get Credit Balance

```
GET /v2/billing/credits
```

Get current credit balance.

#### Response:

```
{  
  "data": {  
    "available": 150.50,  
    "reserved": 10.00,  
    "currency": "USD",  
    "updated_at": "2024-01-15T10:30:00Z"  
  }  
}
```

### Get Usage

```
GET /v2/billing/usage
```

Get usage data for a period.

#### Query Parameters:

Parameter	Type	Description
start_date	string	Start date (YYYY-MM-DD)
end_date	string	End date (YYYY-MM-DD)
group_by	string	day, model, endpoint

---

## Webhooks

### List Webhooks

```
GET /v2/webhooks
```

List configured webhooks.

### Create Webhook

```
POST /v2/webhooks
```

#### Request Body:

```
{  
  "url": "https://your-server.com/webhook",  
  "event_types": ["billing.low_balance", "usage.quota_reached"],  
  "description": "Billing alerts"  
}
```

#### Response:

```
{  
  "data": {  
    "id": "wh_abc123",  
    "url": "https://your-server.com/webhook",  
    "secret": "whsec_xyz789...",  
    "event_types": ["billing.low_balance", "usage.quota_reached"],  
    "is_active": true,  
    "created_at": "2024-01-15T10:30:00Z"  
  }  
}
```

## Test Webhook

POST /v2/webhooks/{webhook\_id}/test

Send a test event to the webhook.

---

## Batch Processing

### Create Batch Job

POST /v2/batch/jobs

Create a batch processing job.

#### Request Body:

```
{  
  "type": "completions",  
  "model": "gpt-4o",  
  "input_file": "batch-input.jsonl",  
  "options": {  
    "system_prompt": "You are a helpful assistant.",  
    "max_tokens": 500  
  }  
}
```

### Get Batch Job

GET /v2/batch/jobs/{job\_id}

Get batch job status and results.

### List Batch Jobs

GET /v2/batch/jobs

List all batch jobs.

---

## Error Codes

RADIANT uses standardized error codes across all endpoints. See [Error Codes Reference](#) for the complete list.

## Error Categories

Category	Code Range	Description
Authentication	RADIANT_AUTH_1xxx	Token, API key, session errors
Authorization	RADIANT_AUTHZ_2xxx	Permission, role, tenant errors
Validation	RADIANT_VAL_3xxx	Input validation errors
Resource	RADIANT_RES_4xxx	Not found, conflict, quota errors
Rate Limiting	RADIANT_RATE_5xxx	Throttling and rate limit errors
AI/Model	RADIANT_AI_6xxx	Model, provider, inference errors
Billing	RADIANT_BILL_7xxx	Credits, subscription errors
Storage	RADIANT_STOR_8xxx	File upload, storage errors
Internal	RADIANT_INT_9xxx	Server, database, timeout errors

## Common Error Codes

Code	HTTP	Retryable	Description
RADIANT_AUTH_1001	401		Invalid authentication token
RADIANT_AUTH_1004	401		Invalid API key
RADIANT_VAL_3001	400		Required field missing
RADIANT_RES_4001	404		Resource not found
RADIANT_RATE_5001	429		Rate limit exceeded
RADIANT_BILL_7001	402		Insufficient credits
RADIANT_AI_6004	502		AI provider error
RADIANT_INT_9001	500		Internal server error

### Error Response Format:

```
{  
  "error": {  
    "code": "RADIANT_RATE_5001",  
    "message": "Too many requests. Please slow down.",  
    "category": "rate_limit",  
    "retryable": true,  
    "timestamp": "2024-12-25T10:30:00.000Z"  
  }  
}
```

**Retry-After Header:** Retryable errors include `Retry-After` header with seconds to wait.

## Rate Limits

Tier	Requests/min	Tokens/min
Free	10	10,000
Starter	50	50,000
Professional	100	200,000
Business	500	1,000,000
Enterprise	2,000	Unlimited

Rate limit headers:

```
X-RateLimit-Limit: 100
X-RateLimit-Remaining: 95
X-RateLimit-Reset: 1703980860
```

---

## SDKs

### TypeScript/JavaScript

```
npm install @radiant/sdk

import { RadiantClient } from '@radiant/sdk';

const client = new RadiantClient({ apiKey: 'your-key' });
const response = await client.chat.create({
  model: 'gpt-4o',
  messages: [{ role: 'user', content: 'Hello!' }],
});
```

### Python

```
pip install radiant-sdk

from radiant import RadiantClient

client = RadiantClient(api_key="your-key")
response = client.chat.create(
    model="gpt-4o",
    messages=[{"role": "user", "content": "Hello!"}],
)
```

### CLI

```
npm install -g @radiant/cli
radiant auth login
radiant chat send "Hello!"
```

---

## Changelog

See [CHANGELOG.md](#) for version history.

## Support

- **Email:** support@radiant.example.com
- **Documentation:** <https://docs.radiant.example.com>
- **Status:** <https://status.radiant.example.com>

# RADIANT API Versioning Guide

## Overview

This document describes the API versioning strategy for the RADIANT platform.

## Versioning Strategy

### URL Path Versioning

RADIANT uses URL path versioning as the primary versioning mechanism:

`https://api.radiant.example.com/v2/models`  
`https://api.radiant.example.com/v2/chat/completions`

### Version Lifecycle

Version	Status	Support End	Deprecation
v1	Deprecated	2024-06-01	Sunset
v2	Current	-	-
v3	Planned	-	Q3 2025

### Support Policy

- **Current version:** Full support, all new features
- **Previous version:** Security fixes only, 12 months after new version
- **Deprecated:** No fixes, 6-month sunset warning

## Breaking vs Non-Breaking Changes

### Non-Breaking Changes (No Version Bump)

These changes can be made to the current version:

- Adding new endpoints
- Adding new optional request parameters
- Adding new response fields
- Adding new enum values (with graceful handling)
- Performance improvements
- Bug fixes that don't change behavior

## Breaking Changes (Require New Version)

These changes require a new API version:

- Removing endpoints
- Removing request/response fields
- Changing field types
- Changing validation rules
- Changing authentication methods
- Changing error response formats
- Removing enum values
- Changing default values

## Version Header Support

### Request Headers

```
# Specify API version via header (optional override)
X-API-Version: 2024-12-01
```

```
# Request specific features
X-API-Features: beta-orchestration,streaming-v2
```

### Response Headers

```
# Current API version
X-API-Version: 2
X-API-Version-Date: 2024-12-01
```

```
# Deprecation warning
Deprecation: true
Sunset: Sat, 01 Jun 2024 00:00:00 GMT
Link: <https://api.radiant.example.com/v2>; rel="successor-version"
```

## Deprecation Process

### Timeline

Day 0:  
Announce deprecation  
Add Deprecation header  
Update documentation

Month 3:  
Send reminder emails  
Log deprecation warnings

Month 6:  
Begin returning 299 status for deprecated endpoints  
Increase warning frequency

Month 12:  
Sunset - Return 410 Gone  
Redirect to new version docs

### Deprecation Headers

```
// Add deprecation headers to old endpoints
function addDeprecationHeaders(res: Response, sunset: Date): void {
```

```

    res.setHeader('Deprecation', 'true');
    res.setHeader('Sunset', sunset.toUTCString());
    res.setHeader('Link', '<https://api.radiant.example.com/v3>; rel="successor-version"');
}

```

## Deprecation Warnings

```

// Log deprecation usage for migration tracking
async function logDeprecatedUsage(req: Request): Promise<void> {
  await analytics.track({
    event: 'deprecated_api_usage',
    properties: {
      endpoint: req.path,
      version: extractVersion(req),
      apiKey: extractKeyId(req),
      tenant: extractTenantId(req),
    },
  });
}

```

## Migration Guide Template

### v1 to v2 Migration

```

# Migrating from API v1 to v2

## Breaking Changes

### 1. Authentication
- v1: API key in query string (`?api_key=xxx`)
- v2: API key in header (`Authorization: Bearer xxx`)

### 2. Response Format
- v1: Flat response (`{ models: [...] }`)
- v2: Wrapped response (`{ data: [...], meta: [...] }`)

### 3. Error Format
- v1: `{ error: "message" }`
- v2: `{ error: { code: "...", message: "...", details: [...] } }`
```

### ## Migration Steps

1. Update authentication headers
2. Update response parsing
3. Update error handling
4. Test all endpoints
5. Switch base URL from /v1 to /v2

## Feature Flags

### Beta Features

```

// Enable beta features via header
const betaFeatures = {
```

```

'beta-orchestration': true,
'beta-streaming-v2': true,
'beta-function-calling': true,
};

function isBetaEnabled(req: Request, feature: string): boolean {
  const features = req.headers['x-api-features']?.split(',') || [];
  return features.includes(feature) && betaFeatures[feature];
}

```

## Graduated Features

```

// Track feature graduation
const featureGraduation = {
  'function-calling': {
    beta: '2024-06-01',
    stable: '2024-09-01',
    version: 'v2',
  },
  'streaming-v2': {
    beta: '2024-09-01',
    stable: null, // Still in beta
    version: 'v2',
  },
};

```

## SDK Versioning

### SDK Version Matrix

SDK	Latest	Min API Version	Max API Version
JavaScript	2.5.0	v2	v2
Python	2.3.0	v2	v2
Go	1.2.0	v2	v2
Ruby	1.1.0	v2	v2

### SDK Version Headers

```

# SDKs include version info
User-Agent: radiant-js/2.5.0 node/20.10.0
X-Radiant-SDK: js
X-Radiant-SDK-Version: 2.5.0

```

## OpenAPI Specification

### Versioned Specs

```

/docs/openapi/v2.yaml      # Current version
/docs/openapi/v2-beta.yaml # With beta features
/docs/openapi/v1.yaml      # Deprecated version

```

## Schema Versioning

```
# openapi.yaml
openapi: 3.1.0
info:
  title: RADIANT API
  version: 2.0.0
  x-api-version: v2
  x-version-date: '2024-12-01'
  x-deprecation-date: null
```

## Testing Versions

### Version Compatibility Tests

```
describe('API Version Compatibility', () => {
  it('should support v2 endpoints', async () => {
    const res = await fetch('/v2/health');
    expect(res.status).toBe(200);
  });

  it('should return 410 for sunset v1 endpoints', async () => {
    const res = await fetch('/v1/health');
    expect(res.status).toBe(410);
    expect(res.headers.get('Link')).toContain('/v2');
  });

  it('should include deprecation headers for deprecated endpoints', async () => {
    const res = await fetch('/v2/deprecated-endpoint');
    expect(res.headers.get('Deprecation')).toBe('true');
    expect(res.headers.get('Sunset')).toBeDefined();
  });
});
```

## Client Communication

### Changelog

Maintain a public changelog:

```
# API Changelog

## 2024-12-24 (v2.5)
- Added: Orchestration patterns endpoint
- Added: Workflow proposals endpoint
- Changed: Increased rate limits for Professional tier

## 2024-12-01 (v2.4)
- Added: AI translation for localization
- Deprecated: Legacy /translate endpoint (sunset 2025-06-01)
```

### Email Notifications

```
// Notify developers of breaking changes
async function notifyApiChanges(change: ApiChange): Promise<void> {
```

```

const affectedKeys = await getApiKeysUsingEndpoint(change.endpoint);

for (const key of affectedKeys) {
  await sendEmail({
    to: key.ownerEmail,
    subject: `RADIANT API: ${change.type} - ${change.endpoint}`,
    template: 'api-change-notification',
    data: {
      change,
      migrationGuide: change.migrationGuideUrl,
      deadline: change.sunsetDate,
    },
  });
}

```

## Best Practices

### For API Developers

- Plan for change:** Design APIs to be extensible
- Use optional fields:** Make new fields optional with defaults
- Version from day one:** Include version in all endpoints
- Document everything:** Keep OpenAPI specs updated
- Communicate early:** 12-month deprecation notice minimum

### For API Consumers

- Pin versions:** Don't use unversioned endpoints
- Handle unknown fields:** Ignore unexpected response fields
- Monitor deprecation headers:** Set up alerts
- Test regularly:** Run integration tests against current version
- Subscribe to updates:** Follow changelog and email updates

## Contact

Role	Contact	Purpose
API Support	api-support@radiant.example.com	Usage questions
Developer Relations	devrel@radiant.example.com	SDKs, docs
Engineering	engineering@radiant.example.com	Bug reports

# RADIANT Error Codes Reference

Standardized error codes for consistent API responses across all RADIANT services.

## Overview

All RADIANT errors follow a consistent format:

```
{  
  "error": {  
    "code": "RADIANT_AUTH_1001",  
    "message": "Invalid authentication token. Please sign in again.",  
    "category": "authentication",  
    "retryable": false,  
    "timestamp": "2024-12-25T10:30:00.000Z"  
  }  
}
```

## Error Code Format

RADIANT\_<CATEGORY>\_<NUMBER>

- **RADIANT** - Prefix for all error codes
  - **CATEGORY** - Short category identifier (AUTH, VAL, RES, etc.)
  - **NUMBER** - Unique 4-digit number within category
- 

## Authentication Errors (1xxx)

Errors related to authentication and identity.

Code	HTTP	Retryable	Description
RADIANT_AUTH_100101			Invalid authentication token
RADIANT_AUTH_100201			Token has expired
RADIANT_AUTH_100301			Missing authentication token
RADIANT_AUTH_100401			Invalid API key
RADIANT_AUTH_100501			API key has expired
RADIANT_AUTH_100601			API key has been revoked
RADIANT_AUTH_100703			Insufficient API key scope
RADIANT_AUTH_100801			Multi-factor authentication required
RADIANT_AUTH_100901			Session has expired

---

## Authorization Errors (2xxx)

Errors related to permissions and access control.

Code	HTTP	Retryable	Description
RADIANT_AUTHZ_2001	403		Forbidden - access denied
RADIANT_AUTHZ_2002	403		Tenant ID mismatch
RADIANT_AUTHZ_2003	403		Required role not assigned
RADIANT_AUTHZ_2004	403		Permission denied
RADIANT_AUTHZ_2005	403		Resource access denied
RADIANT_AUTHZ_2006	403		Subscription tier insufficient

---

## Validation Errors (3xxx)

Errors related to input validation.

Code	HTTP	Retryable	Description
RADIANT_VAL_3001	400		Required field is missing
RADIANT_VAL_3002	400		Invalid field format
RADIANT_VAL_3003	400		Value out of allowed range
RADIANT_VAL_3004	400		Invalid data type
RADIANT_VAL_3005	400		Constraint violation
RADIANT_VAL_3006	400		Schema mismatch
RADIANT_VAL_3007	400		Invalid JSON in request body
RADIANT_VAL_3008	400		Maximum length exceeded
RADIANT_VAL_3009	400		Minimum length required

---

## Resource Errors (4xxx)

Errors related to resources and entities.

Code	HTTP	Retryable	Description
RADIANT_RES_4001	404		Resource not found
RADIANT_RES_4002	409		Resource already exists
RADIANT_RES_4003	410		Resource has been deleted
RADIANT_RES_4004	423		Resource is locked
RADIANT_RES_4005	409		Resource conflict
RADIANT_RES_4006	429		Resource quota exceeded

---

## Rate Limiting Errors (5xxx)

Errors related to rate limiting and throttling.

Code	HTTP	Retryable	Description
RADIANT_RATE_5001	429		Rate limit exceeded
RADIANT_RATE_5002	429		Tenant rate limit exceeded
RADIANT_RATE_5003	429		User rate limit exceeded
RADIANT_RATE_5004	429		API key rate limit exceeded
RADIANT_RATE_5005	429		Model rate limit exceeded
RADIANT_RATE_5006	429		Burst limit exceeded

**Retry-After Header:** Rate limit errors include `Retry-After` header with seconds to wait.

---

## AI/Model Errors (6xxx)

Errors related to AI models and inference.

Code	HTTP	Retryable	Description
RADIANT_AI_6001	404		Model not found
RADIANT_AI_6002	503		Model temporarily unavailable
RADIANT_AI_6003	503		Model overloaded
RADIANT_AI_6004	502		AI provider error
RADIANT_AI_6005	400		Context length exceeded
RADIANT_AI_6006	400		Content filtered by safety
RADIANT_AI_6007	400		Invalid AI request
RADIANT_AI_6008	500		Streaming error
RADIANT_AI_6009	504		AI request timeout
RADIANT_AI_6010	503		Model is cold (warming up)

---

## Billing Errors (7xxx)

Errors related to billing, credits, and subscriptions.

Code	HTTP	Retryable	Description
RADIANT_BILL_7001	402		Insufficient credits
RADIANT_BILL_7002	402		Payment required
RADIANT_BILL_7003	402		Payment failed
RADIANT_BILL_7004	402		Subscription expired
RADIANT_BILL_7005	402		Subscription cancelled
RADIANT_BILL_7006	400		Invalid coupon code
RADIANT_BILL_7007	429		Usage quota exceeded

## Storage Errors (8xxx)

Errors related to file storage.

Code	HTTP	Retryable	Description
RADIANT_STOR_8001	413		Storage quota exceeded
RADIANT_STOR_8002	413		File too large
RADIANT_STOR_8003	415		Invalid file type
RADIANT_STOR_8004	500		Upload failed
RADIANT_STOR_8005	404		File not found

---

## Internal Errors (9xxx)

Internal server errors and system failures.

Code	HTTP	Retryable	Description
RADIANT_INT_9001	500		Internal server error
RADIANT_INT_9002	500		Database error
RADIANT_INT_9003	500		Cache error
RADIANT_INT_9004	500		Queue processing error
RADIANT_INT_9005	503		Service unavailable
RADIANT_INT_9006	502		Dependency failure
RADIANT_INT_9007	500		Configuration error
RADIANT_INT_9008	504		Request timeout

---

## Usage in Code

### TypeScript/JavaScript

```
import {
  ErrorCode,
  RadianError,
  createNotFoundError,
  createValidationError,
  isRetryableError
} from '@radiant/shared';

// Using factory functions (recommended)
throw createNotFoundError('User', userId);
throw createValidationError('Email is required', 'email');

// Direct construction
throw new RadianError(ErrorCode.AUTH_INVALID_TOKEN, 'Custom message', {
  details: { tokenPrefix: 'rad...' },
  requestId: context.awsRequestId,
});
```

```
// Check if retryable
if (isRetryableError(error.code)) {
    // Implement retry logic
}
```

## Response Format

```
// RadiantError automatically formats responses
const error = new RadiantError(ErrorCodes.RESOURCE_NOT_FOUND);
return error.toResponse();

// Returns:
// {
//   statusCode: 404,
//   headers: { 'Content-Type': 'application/json' },
//   body: '{"error":{"code":"RADIANT_RES_4001","message":"Resource not found.","category":"resource",'
// }
```

---

## Client Handling

### Retry Logic

```
async function callWithRetry(fn: () => Promise<Response>, maxRetries = 3) {
    for (let i = 0; i < maxRetries; i++) {
        try {
            const response = await fn();
            if (response.ok) return response;

            const error = await response.json();
            if (!error.error.retryable) throw error;

            const retryAfter = response.headers.get('Retry-After') || '5';
            await sleep(parseInt(retryAfter) * 1000);
        } catch (e) {
            if (i === maxRetries - 1) throw e;
        }
    }
}
```

### Error Display

```
function getUserMessage(error: RadiantError): string {
    // Error codes include user-friendly messages
    return error.message;
}

function shouldShowRetryButton(error: RadiantError): boolean {
    return error.retryable;
}
```

---

## Adding New Error Codes

1. Add the code to packages/shared/src/errors/codes.ts:

```
export const ErrorCodes = {
  // ... existing codes
  MY_NEW_ERROR: 'RADIANT_CAT_NNNN',
} as const;
```

2. Add metadata:

```
export const ErrorCodeMetadata: Record<ErrorCode, {...}> = {
  // ... existing metadata
  [ErrorCodes.MY_NEW_ERROR]: {
    httpStatus: 400,
    category: 'category',
    retryable: false,
    userMessage: 'User-friendly error message.',
  },
};
```

3. Update this documentation.
- 

## See Also

- [API Reference](#)
- [Contributing Guide](#)
- [Troubleshooting](#)

# RADIANT Testing Guide

Comprehensive guide for testing RADIANT components.

## Overview

RADIANT uses a multi-layered testing strategy:

Layer	Tool	Location	Purpose
Unit Tests	Vitest	**/__tests__/*.test.ts	Test individual functions/components
Integration Tests	Vitest	**/__tests__/*.integration.ts	Test service interactions
E2E Tests	Playwright	apps/admin-dashboard/e2e/	Test user workflows
Swift Tests	XCTest	apps/swift-deployer/Tests/	Test Swift services

## Quick Start

```
# Run all tests
pnpm test

# Run with coverage
pnpm test:coverage

# Run E2E tests
cd apps/admin-dashboard && pnpm test:e2e

# Run Swift tests
cd apps/swift-deployer && swift test
```

## Unit Testing

### Lambda Handler Tests

Tests for Lambda handlers are located in \_\_tests\_\_/ directories:

```
packages/infrastructure/lambda/
  admin/
    __tests__/
```

```
    handler.test.ts
billing/
  __tests__/
    handler.test.ts
shared/
  __tests__/
    auth.test.ts
    errors.test.ts
    services.test.ts
```

## Running Lambda Tests

```
cd packages/infrastructure
```

```
# Run all Lambda tests
pnpm test
```

```
# Run specific handler
pnpm test -- admin
pnpm test -- billing
```

```
# Run with coverage
pnpm test:coverage
```

```
# Watch mode
pnpm test:watch
```

## Writing Lambda Tests

```
import { describe, it, expect, vi, beforeEach } from 'vitest';
import type { APIGatewayProxyEvent, Context } from 'aws-lambda';

// Mock dependencies
vi.mock('../shared/db', () => ({
  listTenants: vi.fn(),
  getTenantById: vi.fn(),
}));

import { handler } from '../handler';
import { listTenants, getTenantById } from '../shared/db';

// Create mock context
const mockContext = {
  awsRequestId: 'test-request-id',
  functionName: 'test-handler',
  // ... other required fields
} as Context;

// Create mock event helper
function createMockEvent(overrides = {}): APIGatewayProxyEvent {
  return {
    httpMethod: 'GET',
    path: '/test',
    headers: { Authorization: 'Bearer test-token' },
    // ... default values
  };
}
```

```

    ...overrides,
  );
}

describe('Handler', () => {
  beforeEach(() => {
    vi.clearAllMocks();
  });

  it('should return 200 for valid request', async () => {
    (listTenants as ReturnType<typeof vi.fn>).mockResolvedValue([]);

    const event = createMockEvent({ path: '/admin/tenants' });
    const result = await handler(event, mockContext);

    expect(result.statusCode).toBe(200);
  });
});

```

## Shared Module Tests

Test shared utilities and services:

```

// packages/infrastructure/lambda/shared/_tests_/errors.test.ts
import { describe, it, expect } from 'vitest';
import {
  ValidationError,
  NotFoundError,
  isOperationalError,
  toAppError,
} from '../errors';

describe('Error Classes', () => {
  it('should create ValidationError with 400 status', () => {
    const error = new ValidationError('Invalid input');

    expect(error.statusCode).toBe(400);
    expect(error.code).toBe('VALIDATION_ERROR');
  });
});

```

---

## E2E Testing (Admin Dashboard)

### Setup

```

cd apps/admin-dashboard

# Install Playwright browsers
npx playwright install

# Run E2E tests
pnpm test:e2e

```

```

# Run with UI
pnpm test:e2e:ui

# Run specific test file
pnpm test:e2e -- dashboard.spec.ts

```

## Test Structure

```

apps/admin-dashboard/e2e/
  dashboard.spec.ts      # Dashboard navigation tests
  deployment.spec.ts    # Deployment workflow tests
  fixtures/
    test-data.json       # Test fixtures

```

## Writing E2E Tests

```

// apps/admin-dashboard/e2e/dashboard.spec.ts
import { test, expect } from '@playwright/test';

test.describe('Dashboard', () => {
  test.beforeEach(async ({ page }) => {
    // Mock authentication
    await page.addInitScript(() => {
      localStorage.setItem('auth_token', 'test-token');
      localStorage.setItem('user', JSON.stringify({
        id: 'test-user',
        email: 'test@example.com',
        role: 'admin',
      }));
    });
  });

  test('should display dashboard home', async ({ page }) => {
    await page.goto('/');
    await expect(page.getByRole('heading', { level: 1 }))
      .toContainText('Dashboard');
  });

  test('should navigate to models page', async ({ page }) => {
    await page.goto('/');
    await page.getByRole('link', { name: 'Models' }).click();
    await expect(page).toHaveURL('/models');
  });
});

```

---

## Swift Testing

### Test Structure

```

apps/swift-deployer/Tests/
  RadianDeployerTests.swift          # Basic unit tests
  RadianDeployerTests/

```

```

E2ETests/
    DeploymentE2ETests.swift      # E2E workflow tests
ServiceTests/
    LocalStorageManagerTests.swift
    CredentialServiceTests.swift

```

## Running Swift Tests

```

cd apps/swift-deployer

# Run all tests
swift test

# Run specific test class
swift test --filter LocalStorageManagerTests

# Run with verbose output
swift test -v

# Generate coverage (requires llvm-cov)
swift test --enable-code-coverage

```

## Writing Swift Tests

```

import XCTest
@testable import RadiantDeployer

final class LocalStorageManagerTests: XCTestCase {
    var storageManager: LocalStorageManager!

    override func setUp() {
        super.setUp()
        storageManager = LocalStorageManager.shared
    }

    override func tearDown() {
        // Cleanup
        super.tearDown()
    }

    func testSaveAndLoadConfiguration() async throws {
        // Given
        let key = "test_config"
        let config = TestConfiguration(name: "Test", value: 42)

        // When
        try await storageManager.save(config, forKey: key)
        let loaded: TestConfiguration? = try await storageManager.load(forKey: key)

        // Then
        XCTAssertNotNil(loaded)
        XCTAssertEqual(loaded?.name, "Test")
        XCTAssertEqual(loaded?.value, 42)
    }
}

```

```
}
```

---

## Test Utilities

### Mock Factories

Use the shared testing utilities:

```
import {
  createMockTenant,
  createMockUser,
  createMockApiKey,
  createMockChatRequest,
  createMockChatResponse,
  createMockApiGatewayEvent,
  createMockLambdaContext,
} from '@radiant/shared/testing';

// Create mock data
const tenant = createMockTenant({ name: 'Test Corp' });
const user = createMockUser({ tenantId: tenant.id });
const event = createMockApiGatewayEvent({
  httpMethod: 'POST',
  body: JSON.stringify({ model: 'gpt-4o' }),
});
```

### Assertion Helpers

```
import {
  assertDefined,
  assertEquals,
  assertMatch,
  assertContains,
  assertThrows,
  waitFor,
  sleep,
} from '@radiant/shared/testing';

// Custom assertions
assertDefined(result, 'Result should not be null');
assertMatch(response.id, /^chatcmpl_/, 'Invalid response ID format');

// Wait for async condition
await waitFor(() => service.isReady(), { timeout: 5000 });
```

---

## Mocking Guidelines

### Database Mocking

```
vi.mock('../db/client', () => ({
  executeStatement: vi.fn(),
```

```

});;

import { executeStatement } from '../db/client';

// Mock return value
(executeStatement as ReturnType<typeof vi.fn>).mockResolvedValue({
  rows: [{ id: '123', name: 'Test' }],
});

```

## AWS SDK Mocking

```

vi.mock('@aws-sdk/client-s3', () => ({
  S3Client: vi.fn().mockImplementation(() => ({
    send: vi.fn(),
  })),
  PutObjectCommand: vi.fn(),
});

```

## External API Mocking

```

vi.mock('node-fetch', () => ({
  default: vi.fn(),
}));

import fetch from 'node-fetch';

(fetch as ReturnType<typeof vi.fn>).mockResolvedValue({
  ok: true,
  json: () => Promise.resolve({ data: 'mocked' }),
});

```

---

## CI/CD Integration

Tests run automatically in GitHub Actions:

```

# .github/workflows/ci.yml
jobs:
  test:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v4
      - name: Setup Node.js
        uses: actions/setup-node@v4
      - name: Install dependencies
        run: pnpm install
      - name: Run tests
        run: pnpm test
        env:
          DATABASE_URL: postgres://test@localhost:5432/test

```

## Coverage Requirements

- Minimum 80% coverage for new code

- Critical paths require 90%+ coverage
  - All error handling paths must be tested
- 

## Best Practices

### Do's

- Test behavior, not implementation
- Use descriptive test names
- Mock external dependencies
- Test error cases and edge cases
- Keep tests fast and isolated
- Use factory functions for test data

### Don'ts

- Don't test private methods directly
- Don't share state between tests
- Don't test framework code
- Don't ignore flaky tests
- Don't hardcode test data

## Test Naming Convention

```
describe('ServiceName', () => {
  describe('methodName', () => {
    it('should return X when given Y', () => {});
    it('should throw error when invalid input', () => {});
    it('should handle empty array gracefully', () => {});
  });
});
```

---

## Debugging Tests

### Vitest

```
# Run with verbose output
pnpm test -- --reporter=verbose

# Run single test
pnpm test -- -t "should return 200"

# Debug mode
node --inspect-brk node_modules/.bin/vitest run
```

### Playwright

```
# Debug mode with browser
pnpm test:e2e -- --debug

# Generate trace on failure
```

```
pnpm test:e2e -- --trace on  
  
# View trace  
npx playwright show-trace trace.zip
```

## Swift

```
# Run with verbose output  
swift test -v  
  
# Run single test  
swift test --filter "testSaveAndLoadConfiguration"
```

---

## See Also

- Contributing Guide
- Error Codes Reference
- API Reference

# RADIANT Compliance Guide

## Overview

This document outlines RADIANT's compliance posture for SOC 2, HIPAA, and GDPR requirements.

## Compliance Matrix

Framework	Tier Required	Status
SOC 2 Type II	All tiers	Controls implemented
HIPAA	Tier 3+ (GROWTH)	BAA available
GDPR	All tiers (EU data)	DPA available
PCI DSS	N/A	Not applicable (no card data)

## SOC 2 Controls

### Trust Service Criteria

#### Security (Common Criteria)

Control	Implementation
CC1.1 - Board oversight	Documented security policies
CC2.1 - Communication	Security awareness training
CC3.1 - Risk assessment	Annual risk assessments
CC4.1 - Monitoring	CloudWatch, GuardDuty
CC5.1 - Logical access	IAM, Cognito, RLS
CC6.1 - System operations	Runbooks, on-call
CC7.1 - Change management	CI/CD, PR reviews
CC8.1 - Risk mitigation	WAF, rate limiting
CC9.1 - Entity risk	Vendor assessments

### Availability

Control	Implementation
A1.1 - Capacity planning	Auto-scaling, monitoring
A1.2 - Environmental protection	Multi-AZ, DR procedures
A1.3 - Recovery	Backups, PITR, runbooks

## Confidentiality

Control	Implementation
C1.1 - Data classification	PII tagging, encryption
C1.2 - Data disposal	Lifecycle policies

## Evidence Collection

```
// Automated evidence collection
const auditLogs = {
  // All admin actions logged
  source: 'audit_logs table',
  retention: '7 years',

  // Access logs
  accessLogs: 'CloudWatch Logs',

  // Configuration changes
  configChanges: 'AWS Config',

  // Security events
  securityEvents: 'GuardDuty findings',
};


```

## Annual Audit Checklist

- Access review completed
- Penetration test completed
- Vulnerability scan completed
- Security training completed
- Incident response test completed
- DR test completed
- Vendor assessments updated
- Policies reviewed and updated

## HIPAA Compliance

### Applicability

HIPAA compliance is available for Tier 3 (GROWTH) and above, which includes:

- Encryption at rest (AES-256)
- Encryption in transit (TLS 1.3)
- Audit logging
- Access controls
- BAA with AWS

### Technical Safeguards

Requirement	Implementation
Access Control (§164.312(a))	Cognito MFA, RLS, RBAC
Audit Controls (§164.312(b))	CloudTrail, audit_logs table
Integrity Controls (§164.312(c))	Checksums, versioning
Transmission Security (§164.312(e))	TLS 1.3, VPC endpoints

## Administrative Safeguards

Requirement	Implementation
Security Officer	Designated in org
Workforce Training	Annual security training
Access Management	Quarterly access reviews
Incident Response	Documented procedures

## Physical Safeguards

Handled by AWS: - Data center security - Device controls - Facility access

## PHI Data Handling

```
-- PHI fields are encrypted at column level
CREATE TABLE patient_data (
    id UUID PRIMARY KEY,
    tenant_id UUID NOT NULL,
    -- PHI fields use additional encryption
    encrypted_data BYTEA NOT NULL,
    encryption_key_id VARCHAR(255) NOT NULL,
    created_at TIMESTAMPTZ DEFAULT NOW()
);

-- Enable RLS for tenant isolation
ALTER TABLE patient_data ENABLE ROW LEVEL SECURITY;
```

## BAA Requirements

Before processing PHI: 1. Sign BAA with RADIANT 2. Enable HIPAA-eligible services only 3. Configure CloudTrail logging 4. Enable AWS Config 5. Review shared responsibility model

## GDPR Compliance

### Data Subject Rights

Right	Implementation
Right to Access	Data export API
Right to Rectification	Self-service + API
Right to Erasure	Deletion API + cascade
Right to Restrict	Processing flags
Right to Portability	JSON/CSV export
Right to Object	Consent management

### Data Export (Right to Access)

```
// API endpoint for data export
// GET /api/v2/gdpr/export
async function exportUserData(userId: string): Promise<UserDataExport> {
    return {
        personalData: await getPersonalData(userId),
```

```

        activityLogs: await getActivityLogs(userId),
        preferences: await getPreferences(userId),
        exportedAt: new Date().toISOString(),
        format: 'JSON',
    };
}

```

## Data Deletion (Right to Erasure)

```

// API endpoint for data deletion
// DELETE /api/v2/gdpr/delete
async function deleteUserData(userId: string): Promise<DeletionResult> {
    // Cascade delete all user data
    await deletePersonalData(userId);
    await deleteActivityLogs(userId);
    await deletePreferences(userId);
    await deleteApiKeys(userId);

    // Anonymize audit logs (retain for compliance)
    await anonymizeAuditLogs(userId);

    return {
        deletedAt: new Date().toISOString(),
        confirmation: generateDeletionCertificate(userId),
    };
}

```

## Data Processing Agreement

DPA includes:

- Nature and purpose of processing
- Types of personal data
- Categories of data subjects
- Sub-processor list
- Technical measures
- Audit rights

## Data Residency

Region	Data Location	Backup Location
EU	eu-west-1 (Ireland)	eu-central-1 (Frankfurt)
US	us-east-1 (Virginia)	us-west-2 (Oregon)
APAC	ap-northeast-1 (Tokyo)	ap-southeast-1 (Singapore)

```

// Enforce data residency
const dataResidency = {
    EU: ['eu-west-1', 'eu-central-1'],
    US: ['us-east-1', 'us-west-2'],
    APAC: ['ap-northeast-1', 'ap-southeast-1'],
};

// Route requests to appropriate region
function routeByResidency(tenantRegion: string): string {
    return dataResidency[tenantRegion][0];
}

```

## Consent Management

```
-- Consent tracking table
CREATE TABLE consent_records (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    user_id UUID NOT NULL REFERENCES users(id),
    consent_type VARCHAR(50) NOT NULL,
    granted BOOLEAN NOT NULL,
    granted_at TIMESTAMPTZ,
    withdrawn_at TIMESTAMPTZ,
    ip_address INET,
    user_agent TEXT,
    created_at TIMESTAMPTZ DEFAULT NOW()
);

-- Consent types: marketing, analytics, essential, third_party
```

## Data Classification

### Classification Levels

Level	Description	Examples	Controls
Public	No restrictions	Marketing content	None
Internal	Business use	Metrics, configs	Access control
Confidential	Sensitive business	API keys, billing	Encryption, audit
Restricted	Highly sensitive	PHI, PII, credentials	Full controls

### PII Fields

```
// Fields classified as PII
const piiFields = [
    'email',
    'display_name',
    'phone_number',
    'ip_address',
    'user_agent',
    'billing_address',
    'payment_method',
];

// Automatic PII detection and tagging
function tagPiiFields(data: Record<string, unknown>): void {
    for (const field of piiFields) {
        if (data[field]) {
            // Tag for audit and retention policies
            data[`[${field}]_pii`] = true;
        }
    }
}
```

## Encryption

### At Rest

Data Type	Encryption	Key Management
Database	AES-256	AWS KMS
S3	AES-256	AWS KMS
Secrets	AES-256	Secrets Manager
Backups	AES-256	AWS KMS

### In Transit

Connection	Protocol	Minimum Version
API	TLS	1.2 (1.3 preferred)
Database	TLS	1.2
Internal	TLS	1.2

### Key Rotation

```
// Automatic key rotation
const kmsKey = new kms.Key(this, 'Key', {
  enableKeyRotation: true, // Annual rotation
  rotationPeriod: cdk.Duration.days(365),
});
```

## Audit Logging

### What We Log

Event Type	Retention	Purpose
Authentication	2 years	Security
Authorization	2 years	Security
Data access	7 years	Compliance
Admin actions	7 years	Compliance
Configuration changes	7 years	Compliance
API requests	90 days	Operations

### Log Format

```
{
  "timestamp": "2024-12-24T10:30:00Z",
  "event_type": "data_access",
  "actor": {
    "id": "user-123",
    "type": "admin",
    "ip": "192.168.1.100"
  },
  "resource": {
    "type": "model",
  }
}
```

```

    "id": "model-456"
},
"action": "read",
"outcome": "success",
"metadata": {}
}

```

## Log Protection

- Logs are immutable (write-once)
- Logs are encrypted at rest
- Access requires special IAM role
- Log deletion requires dual approval

## Incident Response

### Classification

Severity	Response Time	Examples
Critical	1 hour	Data breach, service down
High	4 hours	Attempted breach, partial outage
Medium	24 hours	Policy violation
Low	72 hours	Minor security event

### Breach Notification

Jurisdiction	Requirement	Timeline
GDPR	DPA + affected users	72 hours
HIPAA	HHS + affected individuals	60 days
State laws	Varies by state	Varies

## Vendor Management

### Approved Sub-Processors

Vendor	Purpose	Location	DPA
AWS	Infrastructure	Global	Yes
OpenAI	AI provider	US	Yes
Anthropic	AI provider	US	Yes
Google Cloud	AI provider	Global	Yes

### Vendor Assessment

Annual assessment includes:

- Security questionnaire
- SOC 2 report review
- Penetration test results
- Insurance verification

## Contact

Role	Contact
Data Protection Officer	dpo@radiant.example.com
Security Team	security@radiant.example.com
Compliance Team	compliance@radiant.example.com

# RADIANT Cost Optimization Guide

## Overview

This guide provides strategies for optimizing AWS costs for the RADIANT platform while maintaining performance and reliability.

## Current Architecture Costs

### Estimated Monthly Costs by Tier

Tier	Infrastructure	Est. Monthly Cost
SEED (Dev)	Minimal	\$50-150
STARTUP	Small production	\$200-400
GROWTH	Self-hosted models	\$1,000-2,500
SCALE	Multi-region	\$4,000-8,000
ENTERPRISE	Global, full HA	\$15,000-35,000

### Cost Breakdown by Service

Service	% of Total	Optimization Potential
Aurora	30-40%	High
Lambda	15-25%	Medium
API Gateway	5-10%	Low
S3	5-10%	Medium
CloudFront	5-10%	Low
ElastiCache	10-15%	Medium
Other	10-15%	Varies

## Optimization Strategies

### 1. Database Optimization

#### Aurora Serverless v2

```
// Use Serverless v2 for variable workloads
const cluster = new rds.DatabaseCluster(this, 'Database', {
  serverlessV2MinCapacity: 0.5,    // Scale to near-zero
  serverlessV2MaxCapacity: 16,     // Scale up when needed
});
```

**Savings:** 40-60% vs. provisioned instances for variable workloads

## Reserved Instances (Steady Workloads)

```
# Purchase reserved capacity for predictable workloads
aws rds purchase-reserved-db-instances-offering \
  --reserved-db-instances-offering-id xxx \
  --db-instance-count 1
```

**Savings:** 30-60% for 1-3 year terms

## Read Replicas Strategy

```
// Use read replicas only when needed
// Scale readers with traffic
readers: [
  rds.ClusterInstance.serverlessV2('reader', {
    scaleWithWriter: true, // Auto-scale with primary
  }),
],
```

## 2. Lambda Optimization

### Right-Size Memory

```
// Test different memory sizes to find optimal cost/performance
const memoryOptions = [256, 512, 1024, 2048];

// Use AWS Lambda Power Tuning tool
// https://github.com/alexcasalboni/aws-lambda-power-tuning
```

Function Type	Recommended Memory	Reason
Simple CRUD	256-512 MB	Light compute
API Router	512-1024 MB	Balanced
AI Processing	1024-2048 MB	Heavy compute

### Provisioned Concurrency (Strategic)

```
// Only use for latency-critical functions
new lambda.Alias(this, 'LiveAlias', {
  aliasName: 'live',
  version: fn.currentVersion,
  provisionedConcurrentExecutions: 5, // Keep 5 warm
});
```

**Cost:** ~\$0.015/hour per provisioned instance **Use when:** P99 latency requirements < 200ms

### ARM64 (Graviton2)

```
// 20% cheaper, often faster
const fn = new lambda.Function(this, 'Function', {
  architecture: lambda.Architecture.ARM_64,
  runtime: lambda.Runtime.NODEJS_20_X,
});
```

**Savings:** 20% on compute costs

### 3. S3 Optimization

#### Intelligent Tiering

```
const bucket = new s3.Bucket(this, 'Storage', {
  intelligentTieringConfigurations: [
    {
      name: 'auto-tier',
      archiveAccessTierTime: cdk.Duration.days(90),
      deepArchiveAccessTierTime: cdk.Duration.days(180),
    },
  ],
});
```

**Savings:** Up to 95% for infrequently accessed data

#### Lifecycle Rules

```
const bucket = new s3.Bucket(this, 'Storage', {
  lifecycleRules: [
    // Move old versions to cheaper storage
    {
      noncurrentVersionTransitions: [
        {
          storageClass: s3.StorageClass.INFREQUENT_ACCESS,
          transitionAfter: cdk.Duration.days(30),
        },
        {
          storageClass: s3.StorageClass.GLACIER,
          transitionAfter: cdk.Duration.days(90),
        },
      ],
    },
    // Delete old logs
    {
      prefix: 'logs/',
      expiration: cdk.Duration.days(90),
    },
  ],
});
```

### 4. API Gateway Optimization

#### HTTP API vs REST API

```
// HTTP API is 70% cheaper than REST API
// Use when you don't need REST API features

// HTTP API: $1.00/million requests
// REST API: $3.50/million requests
```

Feature	REST API	HTTP API
Cost	\$3.50/M	\$1.00/M
Lambda integration	Yes	Yes
Request validation	Yes	No
API keys/usage plans	Yes	No
Caching	Yes	No

## Caching

```
// Enable caching for GET endpoints
const method = resource.addMethod('GET', integration, {
  cacheKeyParameters: ['method.request.querystring.id'],
});

// Cache stage setting
stage.cacheClusterEnabled = true;
stage.cacheClusterSize = '0.5'; // 0.5 GB minimum
```

**Note:** Cache costs \$0.02/hour (0.5 GB). Calculate break-even point.

## 5. CloudWatch Optimization

### Log Retention

```
// Don't keep logs forever
new logs.LogGroup(this, 'LogGroup', {
  retention: logs.RetentionDays.ONE_MONTH, // Adjust per environment
});
```

Environment	Retention	Reason
Development	7 days	Quick debugging
Staging	14 days	Testing cycles
Production	90 days	Compliance needs

### Metric Filters vs. Logs Insights

```
// Use metric filters for known patterns
// Cheaper than running Logs Insights queries repeatedly

new logs.MetricFilter(this, 'ErrorMetric', {
  logGroup,
  metricNamespace: 'Radiant',
  metricName: 'Errors',
  filterPattern: logs.FilterPattern.literal('ERROR'),
});
```

## 6. ElastiCache Optimization

### Reserved Nodes

```
# Purchase reserved nodes for production
aws elasticache purchase-reserved-cache-nodes-offering \
  --reserved-cache-nodes-offering-id xxx
```

**Savings:** 30-55% for 1-3 year terms

### Right-Size Nodes

Use Case	Recommended	Memory
Development	cache.t3.micro	0.5 GB
Small Prod	cache.t3.small	1.4 GB
Medium Prod	cache.r6g.large	13 GB

Use Case	Recommended	Memory
Large Prod	cache.r6g.xlarge	26 GB

## 7. Data Transfer Optimization

### Use VPC Endpoints

```
// Avoid NAT Gateway costs for AWS services
vpc.addInterfaceEndpoint('S3Endpoint', {
    service: ec2.InterfaceVpcEndpointAwsService.S3,
});

vpc.addInterfaceEndpoint('SecretsManagerEndpoint', {
    service: ec2.InterfaceVpcEndpointAwsService.SECRETS_MANAGER,
});
```

**Savings:** \$0.045/GB saved vs. NAT Gateway

### CloudFront for S3

```
// Serve S3 content through CloudFront
// Cheaper data transfer + better performance
const distribution = new cloudfront.Distribution(this, 'CDN', {
    defaultBehavior: {
        origin: new origins.S3Origin(bucket),
    },
});
```

## Cost Monitoring

### AWS Cost Explorer

```
# Get cost breakdown by service
aws ce get-cost-and-usage \
--time-period Start=2024-12-01,End=2024-12-31 \
--granularity MONTHLY \
--metrics BlendedCost \
--group-by Type=DIMENSION,Key=SERVICE
```

### CloudWatch Billing Alerts

```
// Alert before surprise bills
new cloudwatch.Alarm(this, 'BillingAlarm', {
    metric: new cloudwatch.Metric({
        namespace: 'AWS/Billing',
        metricName: 'EstimatedCharges',
        dimensionsMap: { Currency: 'USD' },
        statistic: 'Maximum',
        period: cdk.Duration.hours(6),
    }),
    threshold: 1000, // $1000 threshold
    evaluationPeriods: 1,
});
```

## Cost Allocation Tags

```
// Tag all resources for cost tracking
cdk.Tags.of(this).add('Project', 'radianit');
cdk.Tags.of(this).add('Environment', environment);
cdk.Tags.of(this).add('CostCenter', 'platform');
```

## Environment-Specific Recommendations

### Development

- Use Aurora Serverless v2 (scales to zero)
- Minimal Lambda memory
- No provisioned concurrency
- Short log retention
- Single-AZ deployments

**Target:** < \$100/month

### Staging

- Aurora Serverless v2
- Moderate Lambda memory
- No provisioned concurrency
- 14-day log retention
- Single-AZ acceptable

**Target:** < \$300/month

### Production

- Aurora Serverless v2 or Reserved (if predictable)
- Right-sized Lambda memory
- Provisioned concurrency for critical paths
- 90-day log retention
- Multi-AZ required

**Target:** Optimize for reliability, then cost

## Monthly Cost Review Checklist

- Review AWS Cost Explorer for anomalies
- Check for unused resources (idle RDS, orphan EBS)
- Review Lambda right-sizing opportunities
- Check S3 storage class distribution
- Review data transfer costs
- Validate reserved capacity utilization
- Update cost allocation tags
- Project next month's costs

## Tools

- AWS Cost Explorer
- AWS Trusted Advisor
- AWS Compute Optimizer

- Lambda Power Tuning
- Infracost - Cost estimation for IaC

# RADIANT Data Retention Policy

## Overview

This document defines data retention periods and deletion procedures for all data stored in the RADIANT platform.

## Retention Schedule

### User Data

Data Type	Active Retention	Archive	Total Retention	Deletion
Account info	Active + 30 days	N/A	Account lifetime + 30 days	Automatic
Usage history	2 years	5 years	7 years	Automatic
Chat history	90 days	1 year	1 year	Automatic
Uploaded files	Active	30 days post-delete	Active + 30 days	On request
API keys	Active	N/A	Revoked + 90 days	Automatic

### System Data

Data Type	Retention	Purpose	Deletion
Audit logs	7 years	Compliance	Automatic
Access logs	2 years	Security	Automatic
Error logs	90 days	Debugging	Automatic
Metrics	15 months	CloudWatch default	Automatic
Backups	35 days	Recovery	Automatic

### Billing Data

Data Type	Retention	Purpose	Legal Basis
Invoices	7 years	Tax compliance	Legal requirement
Transactions	7 years	Financial audit	Legal requirement
Payment methods	Active	Processing	Contract
Receipts	7 years	Tax compliance	Legal requirement

# Implementation

## Database Retention

```
-- Automatic data cleanup job (runs daily)
CREATE OR REPLACE FUNCTION cleanup_expired_data()
RETURNS void AS $$

BEGIN
    -- Delete expired chat messages (90 days)
    DELETE FROM chat_messages
    WHERE created_at < NOW() - INTERVAL '90 days'
    AND archived = false;

    -- Archive chat messages older than 90 days
    UPDATE chat_messages
    SET archived = true, archived_at = NOW()
    WHERE created_at < NOW() - INTERVAL '90 days'
    AND archived = false;

    -- Delete archived messages older than 1 year
    DELETE FROM chat_messages
    WHERE archived = true
    AND archived_at < NOW() - INTERVAL '1 year';

    -- Delete revoked API keys (90 days after revocation)
    DELETE FROM api_keys
    WHERE revoked_at < NOW() - INTERVAL '90 days';

    -- Delete expired sessions
    DELETE FROM user_sessions
    WHERE expires_at < NOW();

    -- Log cleanup
    INSERT INTO system_jobs (job_name, completed_at, records_affected)
    VALUES ('cleanup_expired_data', NOW(),
        (SELECT count(*) FROM pg_stat_user_tables WHERE relname IN
        ('chat_messages', 'api_keys', 'user_sessions')));

END;
$$ LANGUAGE plpgsql;

-- Schedule daily at 3 AM UTC
SELECT cron.schedule('data-cleanup', '0 3 * * *', 'SELECT cleanup_expired_data()');
```

## S3 Lifecycle Policies

```
const bucket = new s3.Bucket(this, 'Storage', {
    lifecycleRules: [
        // User uploads - delete 30 days after object deletion marker
        {
            id: 'delete-old-versions',
            noncurrentVersionExpiration: cdk.Duration.days(30),
        },
        // Temp files - delete after 7 days
        {
            id: 'delete-temp-files',
            noncurrentVersionExpiration: cdk.Duration.days(7),
        }
    ]
});
```

```

        id: 'cleanup-temp',
        prefix: 'temp/',
        expiration: cdk.Duration.days(7),
    },

    // Logs - transition to Glacier after 90 days, delete after 2 years
    {
        id: 'archive-logs',
        prefix: 'logs/',
        transitions: [
            {
                storageClass: s3.StorageClass.GLACIER,
                transitionAfter: cdk.Duration.days(90),
            },
        ],
        expiration: cdk.Duration.days(730), // 2 years
    },

    // Backups - delete after 35 days
    {
        id: 'cleanup-backups',
        prefix: 'backups/',
        expiration: cdk.Duration.days(35),
    },
],
);

```

## CloudWatch Log Retention

```

// Set retention for all log groups
const logRetention: Record<string, logs.RetentionDays> = {
    // Application logs
    '/aws/lambda/radiant-*': logs.RetentionDays.THREE_MONTHS,

    // API Gateway logs
    '/aws/apigateway/radiant-*': logs.RetentionDays.THREE_MONTHS,

    // Database logs (longer for compliance)
    '/aws/rds/cluster/radiant-*': logs.RetentionDays.TWO_YEARS,

    // Audit logs (longest retention)
    '/radiant/audit/*': logs.RetentionDays.TEN_YEARS,
};

```

## Data Deletion

### User-Initiated Deletion

#### Account Deletion Flow

```

async function deleteUserAccount(userId: string): Promise<void> {
    // 1. Verify identity (MFA required)
    await verifyIdentity(userId);
}

```

```

// 2. Cancel active subscriptions
await cancelSubscriptions(userId);

// 3. Export data (optional, user-requested)
const exportUrl = await exportUserData(userId);

// 4. Mark account for deletion (30-day grace period)
await markForDeletion(userId, {
  scheduledAt: addDays(new Date(), 30),
  reason: 'user_requested',
});

// 5. Send confirmation email
await sendDeletionConfirmation(userId, exportUrl);
}

// Actual deletion after grace period
async function executeAccountDeletion(userId: string): Promise<void> {
  // Delete in order (respect foreign keys)
  await deleteApiKeys(userId);
  await deleteChatHistory(userId);
  await deleteFiles(userId);
  await deletePreferences(userId);
  await deleteBillingHistory(userId); // Anonymize, don't delete
  await deleteAccount(userId);

  // Anonymize audit logs
  await anonymizeAuditLogs(userId);

  // Log deletion for compliance
  await logAccountDeletion(userId);
}

```

## Data Categories Deleted

Category	Action	Timing
Profile	Delete	Immediate
Preferences	Delete	Immediate
Chat history	Delete	Immediate
Files	Delete	Immediate
API keys	Revoke + Delete	Immediate
Billing history	Anonymize	Immediate
Audit logs	Anonymize	Immediate
Backups	Excluded	Expires naturally

## Administrative Deletion

```

// Bulk deletion for compliance (e.g., GDPR request)
async function adminBulkDelete(
  tenantId: string,
  options: {
    dataTypes: string[];
    olderThan: Date;
  }
)

```

```

        reason: string;
        approvedBy: string[];
    }
): Promise<DeletionReport> {
    // Require dual admin approval
    if (options.approvedBy.length < 2) {
        throw new Error('Dual admin approval required');
    }

    // Log the deletion request
    await logAdminAction({
        action: 'bulk_delete',
        tenantId,
        options,
    });

    // Execute deletion
    const results = await Promise.all(
        options.dataTypes.map(type =>
            deleteDataByType(tenantId, type, options.olderThan)
        )
    );
}

return {
    requestId: generateRequestId(),
    deletedAt: new Date(),
    recordsDeleted: results.reduce((a, b) => a + b, 0),
    dataTypes: options.dataTypes,
};
}
}

```

## Tenant Offboarding

```

async function offboardTenant(tenantId: string): Promise<void> {
    // 1. Export all data (required for compliance)
    const exportUrl = await exportTenantData(tenantId);

    // 2. Notify all users
    await notifyTenantUsers(tenantId, 'account_closing');

    // 3. Wait for grace period (30 days default)
    await scheduleTenantDeletion(tenantId, {
        gracePeriod: 30,
        exportUrl,
    });

    // 4. After grace period, delete all data
    // (Handled by scheduled job)
}

```

## Legal Holds

### Implementing a Legal Hold

```
-- Legal hold table
CREATE TABLE legal_holds (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    tenant_id UUID REFERENCES tenants(id),
    user_id UUID REFERENCES users(id),
    hold_type VARCHAR(50) NOT NULL, -- 'litigation', 'investigation', 'regulatory'
    description TEXT,
    started_at TIMESTAMPTZ DEFAULT NOW(),
    expires_at TIMESTAMPTZ,
    created_by UUID REFERENCES administrators(id),
    CONSTRAINT legal_holds_target CHECK (tenant_id IS NOT NULL OR user_id IS NOT NULL)
);

-- Prevent deletion of held data
CREATE OR REPLACE FUNCTION check_legal_hold()
RETURNS TRIGGER AS $$$
BEGIN
    IF EXISTS (
        SELECT 1 FROM legal_holds
        WHERE (tenant_id = OLD.tenant_id OR user_id = OLD.user_id)
        AND (expires_at IS NULL OR expires_at > NOW())
    ) THEN
        RAISE EXCEPTION 'Cannot delete data under legal hold';
    END IF;
    RETURN OLD;
END;
$$ LANGUAGE plpgsql;
```

### Suspending Retention Policies

```
// Suspend automatic deletion during legal hold
async function applyLegalHold(params: {
    holdId: string;
    scope: 'tenant' | 'user';
    targetId: string;
}): Promise<void> {
    // Update retention flags
    await updateRetentionPolicy(params.targetId, {
        suspended: true,
        holdId: params.holdId,
    });

    // Exclude from cleanup jobs
    await excludeFromCleanup(params.targetId);

    // Notify compliance team
    await notifyCompliance('legal_hold_applied', params);
}
```

# Audit Trail

## Retention Actions Log

```
-- Log all retention-related actions
CREATE TABLE retention_actions (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    action_type VARCHAR(50) NOT NULL, -- 'delete', 'archive', 'export', 'hold'
    target_type VARCHAR(50) NOT NULL, -- 'user', 'tenant', 'data_type'
    target_id VARCHAR(255) NOT NULL,
    records_affected INTEGER,
    performed_by UUID,
    reason TEXT,
    metadata JSONB,
    created_at TIMESTAMPTZ DEFAULT NOW()
);

-- Index for compliance queries
CREATE INDEX idx_retention_actions_date ON retention_actions(created_at);
CREATE INDEX idx_retention_actions_target ON retention_actions(target_type, target_id);
```

## Compliance Reporting

```
// Generate retention compliance report
async function generateRetentionReport(
    startDate: Date,
    endDate: Date
): Promise<RetentionReport> {
    return {
        period: { start: startDate, end: endDate },

        // Data deleted by type
        deletions: await getRetentionActions('delete', startDate, endDate),

        // Data archived
        archives: await getRetentionActions('archive', startDate, endDate),

        // Active legal holds
        legalHolds: await getActiveLegalHolds(),

        // Policy violations (data past retention not deleted)
        violations: await getRetentionViolations(),

        // User deletion requests
        userRequests: await getUserDeletionRequests(startDate, endDate),
    };
}
```

# Verification

## Monthly Retention Audit

- Verify cleanup jobs running successfully
- Check for retention policy violations

- Review legal holds status
- Verify S3 lifecycle policies active
- Confirm CloudWatch log retention settings
- Review user deletion requests processed
- Update retention schedule if needed

## Compliance Queries

```
-- Find data past retention period
SELECT
    'chat_messages' as table_name,
    COUNT(*) as records,
    MIN(created_at) as oldest_record
FROM chat_messages
WHERE created_at < NOW() - INTERVAL '90 days'
AND archived = false

UNION ALL

SELECT
    'api_keys' as table_name,
    COUNT(*) as records,
    MIN(revoked_at) as oldest_record
FROM api_keys
WHERE revoked_at < NOW() - INTERVAL '90 days';
```

## Contact

Role	Contact	Purpose
Data Protection Officer	dpo@radiant.example.com	GDPR requests
Legal	legal@radiant.example.com	Legal holds
Compliance	compliance@radiant.example.com	Audit questions

# RADIANT Disaster Recovery Guide

## Overview

This document outlines disaster recovery (DR) procedures for the RADIANT platform, including backup strategies, recovery procedures, and business continuity plans.

## Recovery Objectives

Metric	Target	Maximum
<b>RTO</b> (Recovery Time Objective)	1 hour	4 hours
<b>RPO</b> (Recovery Point Objective)	5 minutes	1 hour

## Backup Strategy

### Database Backups

#### Automated Backups (Aurora)

```
// CDK Configuration
const database = new rds.DatabaseCluster(this, 'Database', {
  backup: {
    retention: cdk.Duration.days(35),           // 35 days retention
    preferredWindow: '03:00-04:00',             // 3-4 AM UTC
  },
  deletionProtection: true,
  storageEncrypted: true,
});
```

#### Point-in-Time Recovery

Aurora supports point-in-time recovery (PITR) to any second within the retention period.

```
# Restore to specific point in time
aws rds restore-db-cluster-to-point-in-time \
--source-db-cluster-identifier radiant-production \
--db-cluster-identifier radiant-production-restored \
--restore-to-time "2024-12-24T10:30:00Z" \
--vpc-security-group-ids sg-xxx \
--db-subnet-group-name radiant-production
```

## Manual Snapshots

```
# Create manual snapshot before major changes
aws rds create-db-cluster-snapshot \
  --db-cluster-identifier radiant-production \
  --db-cluster-snapshot-identifier radiant-production-pre-migration-$(date +%Y%m%d)
```

## S3 Backups

### Versioning

```
// All S3 buckets have versioning enabled
const bucket = new s3.Bucket(this, 'Storage', {
  versioned: true,
  lifecycleRules: [
    {
      noncurrentVersionExpiration: cdk.Duration.days(90),
    },
  ],
});
```

### Cross-Region Replication

```
// Production buckets replicate to DR region
const replicationRule = {
  destination: {
    bucket: drBucket.bucketArn,
    storageClass: s3.StorageClass.STANDARD_IA,
  },
  status: 'Enabled',
};
```

### Secrets Backup

```
# Export secrets for DR (store securely!)
aws secretsmanager get-secret-value \
  --secret-id radiant-production-db \
  --query SecretString \
  --output text > /secure/path/db-credentials.json
```

## Failure Scenarios

### Scenario 1: Single AZ Failure

**Impact:** Partial service degradation **Recovery:** Automatic (Multi-AZ)

Aurora automatically fails over to a read replica in another AZ.

```
# Monitor failover
aws rds describe-events \
  --source-type db-cluster \
  --source-identifier radiant-production \
  --duration 60
```

### Scenario 2: Database Corruption

**Impact:** Data integrity issues **Recovery:** Point-in-time restore

1. Identify corruption time
2. Restore to point before corruption
3. Validate data integrity
4. Switch traffic to restored database

```
# Step 1: Identify issue time from logs
aws logs filter-log-events \
--log-group-name /aws/rds/cluster/radiant-production/error \
--start-time $(date -d '24 hours ago' +%s000)

# Step 2: Restore
aws rds restore-db-cluster-to-point-in-time \
--source-db-cluster-identifier radiant-production \
--db-cluster-identifier radiant-dr-$((date +%Y%m%d%H%M)) \
--restore-to-time "2024-12-24T09:00:00Z"

# Step 3: Update Lambda environment to use new cluster
aws lambda update-function-configuration \
--function-name radiant-production-router \
--environment "Variables={DB_CLUSTER_ARN=arn:aws:rds:...}"
```

### Scenario 3: Region Failure

**Impact:** Complete service outage **Recovery:** Failover to DR region

1. Activate DR region infrastructure
2. Promote Aurora Global Database secondary
3. Update Route 53 to point to DR region
4. Verify service health

```
# Step 1: Promote DR database
aws rds failover-global-cluster \
--global-cluster-identifier radiant-global \
--target-db-cluster-identifier radiant-dr-cluster

# Step 2: Update DNS
aws route53 change-resource-record-sets \
--hosted-zone-id Z123456 \
--change-batch file://dr-dns-failover.json
```

### Scenario 4: Accidental Deletion

**Impact:** Data loss **Recovery:** Restore from backup

```
# Restore deleted S3 objects
aws s3api list-object-versions \
--bucket radiant-storage-production \
--prefix "deleted/path/" \
--query 'DeleteMarkers[?IsLatest==`true`]'

# Restore specific version
aws s3api delete-object \
--bucket radiant-storage-production \
--key "path/to/file" \
--version-id "delete-marker-version-id"
```

## Scenario 5: Security Breach

**Impact:** Potential data exposure **Recovery:** Isolation and investigation

1. Isolate affected systems
2. Rotate all credentials
3. Investigate scope
4. Restore from known-good backup
5. Notify affected parties

```
# Step 1: Disable API access
aws apigateway update-stage \
  --rest-api-id abc123 \
  --stage-name v2 \
  --patch-operations op=replace,path=/throttling/rateLimit,value=0

# Step 2: Rotate database credentials
aws secretsmanager rotate-secret \
  --secret-id radiant-production-db

# Step 3: Invalidate all sessions
aws cognito-oidc admin-user-global-sign-out \
  --user-pool-id us-east-1_xxx \
  --username "*"
```

## Recovery Procedures

### Database Recovery Runbook

```
#!/bin/bash
# database-recovery.sh

set -e

CLUSTER_ID="radiant-production"
RESTORE_TIME="${1:-$(date -d '1 hour ago' -Iseconds)}"
NEW_CLUSTER_ID="radiant-dr-$(date +%Y%m%d%H%M)"

echo " Starting database recovery..."
echo "   Source: $CLUSTER_ID"
echo "   Restore time: $RESTORE_TIME"
echo "   New cluster: $NEW_CLUSTER_ID"

# Create restored cluster
aws rds restore-db-cluster-to-point-in-time \
  --source-db-cluster-identifier "$CLUSTER_ID" \
  --db-cluster-identifier "$NEW_CLUSTER_ID" \
  --restore-to-time "$RESTORE_TIME" \
  --db-subnet-group-name radiant-production \
  --vpc-security-group-ids sg-xxx

echo " Waiting for cluster to be available..."
aws rds wait db-cluster-available \
  --db-cluster-identifier "$NEW_CLUSTER_ID"
```

```

# Create instance
aws rds create-db-instance \
--db-instance-identifier "${NEW_CLUSTER_ID}-instance-1" \
--db-cluster-identifier "$NEW_CLUSTER_ID" \
--db-instance-class db.r6g.large \
--engine aurora-postgresql

echo " Waiting for instance to be available..."
aws rds wait db-instance-available \
--db-instance-identifier "${NEW_CLUSTER_ID}-instance-1"

echo " Database restored successfully!"
echo " Endpoint: $(aws rds describe-db-clusters \
--db-cluster-identifier "$NEW_CLUSTER_ID" \
--query 'DBClusters[0].Endpoint' --output text)"

```

## Full Service Recovery Runbook

```

#!/bin/bash
# full-recovery.sh

set -e

echo " RADIANT Full Service Recovery"
echo "====="

# Step 1: Database
echo "Step 1: Recovering database..." 
./scripts/dr/database-recovery.sh

# Step 2: Update Lambda configurations
echo "Step 2: Updating Lambda configurations..." 
for fn in router admin billing localization configuration; do
    aws lambda update-function-configuration \
    --function-name "radian-production-$fn" \
    --environment "Variables={DB_CLUSTER_ARN=$NEW_DB_ARN}"
done

# Step 3: Clear caches
echo "Step 3: Clearing caches..." 
redis-cli -h radiant-cache.xxx.cache.amazonaws.com FLUSHALL

# Step 4: Verify health
echo "Step 4: Verifying service health..." 
curl -f https://api.radiant.example.com/v2/health || exit 1

# Step 5: Run smoke tests
echo "Step 5: Running smoke tests..." 
k6 run --env BASE_URL=https://api.radiant.example.com tests/load/k6-config.js

echo " Recovery complete!"

```

# Testing DR Procedures

## Quarterly DR Drill

1. Preparation
  - Schedule maintenance window
  - Notify stakeholders
  - Prepare rollback plan
2. Execution
  - Simulate failure scenario
  - Execute recovery procedures
  - Measure RTO/RPO
3. Validation
  - Verify data integrity
  - Run integration tests
  - Check all services
4. Documentation
  - Record actual RTO/RPO
  - Document issues encountered
  - Update procedures

## DR Test Checklist

- Database point-in-time recovery tested
- S3 object recovery tested
- Secret rotation tested
- Lambda rollback tested
- DNS failover tested
- Communication plan executed
- Recovery time recorded
- Post-mortem completed

# Communication Plan

## Escalation Matrix

Severity	Response Time	Notify
SEV1	15 min	Eng Lead, CTO, Status Page
SEV2	30 min	Eng Lead, Status Page
SEV3	2 hours	On-call team

## Status Page Updates

```
# Update status page (example with Statuspage.io)
curl -X POST https://api.statuspage.io/v1/pages/xxx/incidents \
-H "Authorization: OAuth $STATUSPAGE_API_KEY" \
-d '{
  "incident": {
    "name": "Service Degradation",
    "status": "investigating",
    "body": "We are investigating reports of API errors."
  }
}'
```

## Infrastructure as Code

All DR infrastructure is defined in CDK:

```
// lib/stacks/dr-stack.ts
export class DRStack extends cdk.Stack {
  constructor(scope: Construct, id: string, props: DRStackProps) {
    super(scope, id, props);

    // Global Database for cross-region replication
    const globalCluster = new rds.CfnGlobalCluster(this, 'GlobalCluster', {
      globalClusterIdentifier: 'radian-global',
      sourceDbClusterIdentifier: props.primaryClusterArn,
    });

    // S3 Cross-Region Replication
    const drBucket = new s3.Bucket(this, 'DRBucket', {
      bucketName: `radian-storage-dr-${props.drRegion}`,
    });
  }
}
```

## Contacts

Role	Contact	Backup
DR Coordinator	dr@radian-global.com	cto@radian-global.com
Database Admin	dba@radian-global.com	platform@radian-global.com
Security	security@radian-global.com	cto@radian-global.com

# RADIANT Performance Guide

## Overview

This guide covers performance optimization, caching strategies, and scalability considerations for the RADIANT platform.

## Architecture Performance

### Request Flow

```
Client → CloudFront → WAF → API Gateway → Lambda → Aurora  
                                ↓  
                                Redis Cache
```

### Latency Targets

Component	Target	Max Acceptable
CloudFront edge	< 50ms	100ms
WAF processing	< 5ms	20ms
API Gateway	< 20ms	50ms
Lambda cold start	< 500ms	1000ms
Lambda execution	< 200ms	500ms
Database query	< 50ms	200ms
<b>Total P95</b>	<b>&lt; 500ms</b>	<b>2000ms</b>

## Caching Strategy

### Multi-Layer Caching

```
CloudFront CDN  
TTL: 5m for static, 1m for API (with stale-while-revalidate)
```

↓

```
API Gateway Cache  
TTL: 60s for GET endpoints
```

↓

Redis Cache

```
Session: 24h, Config: 5m, Translations: 1h
```

```
↓
```

```
Aurora Database  
Query cache, Connection pooling
```

## Cache Keys

```
// Session cache  
`session:${tenantId}:${userId}` → TTL: 24h  
  
// Configuration cache  
`config:${tenantId}:${key}` → TTL: 5m  
`config:global:${key}` → TTL: 5m  
  
// Translation cache  
`i18n:${language}:bundle` → TTL: 1h  
`i18n:${language}:${key}` → TTL: 1h  
  
// Model cache  
`models:${tenantId}:list` → TTL: 5m  
`models:${tenantId}:${modelId}` → TTL: 5m  
  
// Rate limit cache  
`ratelimit:${tenantId}:${endpoint}` → TTL: 1m  
`ratelimit:ip:${ip}` → TTL: 5m
```

## Cache Invalidation

```
// Pattern-based invalidation  
await redis.del(`config:${tenantId}:*`);  
  
// Event-driven invalidation  
eventBridge.putEvents({  
  Entries: [  
    Source: 'radiant.config',  
    DetailType: 'ConfigUpdated',  
    Detail: JSON.stringify({ tenantId, key }),  
  ],  
});
```

## Database Optimization

### Connection Pooling

```
// RDS Proxy configuration  
const pool = {  
  min: 2,  
  max: 10,  
  idleTimeoutMillis: 30000,  
  connectionTimeoutMillis: 5000,  
};
```

## Query Optimization

```
-- Always use indexes
CREATE INDEX idx_models_tenant_status ON ai_models(tenant_id, status);
CREATE INDEX idx_transactions_tenant_date ON credit_transactions(tenant_id, created_at DESC);

-- Use covering indexes for common queries
CREATE INDEX idx_models_list ON ai_models(tenant_id, status, is_enabled)
    INCLUDE (display_name, category, input_cost_per_1k);

-- Partition large tables by date
CREATE TABLE audit_logs_2024_01 PARTITION OF audit_logs
    FOR VALUES FROM ('2024-01-01') TO ('2024-02-01');
```

## RLS Performance

```
-- Set tenant context once per request
SET app.current_tenant_id = 'tenant-123';

-- All subsequent queries automatically filtered
SELECT * FROM models; -- Implicitly filtered by RLS
```

## Lambda Optimization

### Cold Start Reduction

```
// 1. Minimize dependencies
// 2. Use Lambda layers for shared code
// 3. Enable provisioned concurrency for critical functions

// Provisioned concurrency config
new lambda.Function(this, 'Router', {
    // ... config
    provisionedConcurrentExecutions: 10, // Keep 10 warm
});
```

### Memory Optimization

```
// Memory vs CPU tradeoff
// More memory = more CPU = faster execution

// Recommended settings by function type:
const memoryConfig = {
    router: 1024,      // Main API - balanced
    billing: 512,     // Light compute
    aiProxy: 2048,    // Heavy compute for AI
    migration: 256,   // Infrequent, light
};
```

### Bundling

```
// esbuild configuration for minimal bundle size
{
    bundle: true,
    minify: true,
```

```

treeShaking: true,
external: ['aws-sdk'], // Use Lambda runtime SDK
target: 'node20',
}

```

## Rate Limiting

### Tier Limits

Tier	RPS	Burst	Daily
Free	10	20	1,000
Starter	50	100	10,000
Professional	100	200	50,000
Business	500	1,000	250,000
Enterprise	2,000	5,000	Unlimited

### Implementation

```

// Token bucket algorithm in Redis
const rateLimiter = {
  async checkLimit(tenantId: string, limit: number): Promise<boolean> {
    const key = `ratelimit:${tenantId}`;
    const current = await redis.incr(key);

    if (current === 1) {
      await redis.expire(key, 1); // 1 second window
    }

    return current <= limit;
  }
};

```

## Load Testing

### Running Tests

```

# Install k6
brew install k6

# Run smoke test
k6 run --env BASE_URL=https://api-dev.radiant.example.com tests/load/k6-config.js

# Run with specific scenario
k6 run --env BASE_URL=https://api-dev.radiant.example.com \
-e SCENARIO=load tests/load/k6-config.js

```

### Performance Baselines

Metric	Baseline	Target
Throughput	500 RPS	2000 RPS
P50 Latency	100ms	50ms

Metric	Baseline	Target
P95 Latency	500ms	200ms
P99 Latency	1000ms	500ms
Error Rate	< 1%	< 0.1%

## Scaling

### Horizontal Scaling

Component	Scaling Method
Lambda	Automatic (up to account limit)
Aurora	Read replicas, Serverless v2
Redis	ElastiCache cluster mode
API Gateway	Automatic

### Vertical Scaling

```
// Aurora Serverless v2 scaling
const database = new rds.DatabaseCluster(this, 'Database', {
  serverlessV2MinCapacity: 0.5, // Minimum ACUs
  serverlessV2MaxCapacity: 16, // Maximum ACUs
});

// Lambda memory scaling
const lambda = new lambda.Function(this, 'Function', {
  memorySize: 2048, // More memory = more CPU
});
```

## Monitoring

### Key Metrics

```
// CloudWatch metrics to monitor
const metrics = {
  // Latency
  'AWS/ApiGateway/Latency': 'p95 < 500ms',
  'AWS/Lambda/Duration': 'p95 < 200ms',
  'AWS/RDS/ReadLatency': 'avg < 50ms',

  // Throughput
  'AWS/ApiGateway/Count': 'track trends',
  'AWS/Lambda/Invocations': 'track trends',

  // Errors
  'AWS/ApiGateway/5XXError': 'rate < 1%',
  'AWS/Lambda/Errors': 'rate < 1%',

  // Resources
  'AWS/Lambda/ConcurrentExecutions': '< 80% of limit',
  'AWS/RDS/CPUUtilization': '< 80%',
```

```
'AWS/RDS/DatabaseConnections': '< 80% of max',  
};
```

## Alerting Thresholds

Metric	Warning	Critical
API P95 Latency	> 1s	> 3s
Error Rate	> 1%	> 5%
Lambda Concurrent	> 500	> 800
DB CPU	> 70%	> 85%
DB Connections	> 60%	> 80%

## Best Practices

### Do's

- Cache aggressively with proper invalidation
- Use connection pooling
- Minimize cold starts with provisioned concurrency
- Use read replicas for read-heavy workloads
- Implement circuit breakers for external services
- Use async processing for non-critical paths

### Don'ts

- Don't make synchronous calls to external APIs in hot paths
- Don't use Lambda for long-running tasks (> 15 min)
- Don't store large objects in Redis
- Don't rely on API Gateway caching for dynamic data
- Don't skip database indexes

## Troubleshooting

### High Latency

1. Check Lambda cold starts (enable provisioned concurrency)
2. Check database query times (add indexes)
3. Check external API latency (add caching/circuit breaker)
4. Check connection pool exhaustion

### High Error Rate

1. Check Lambda errors in CloudWatch Logs
2. Check database connection errors
3. Check rate limiting (429 errors)
4. Check WAF blocked requests

### Scaling Issues

1. Check Lambda concurrent execution limit
2. Check database connection limit
3. Check API Gateway throttling
4. Check Redis memory usage

# RADIANT AI Registry Seed Data System

## Technical Documentation

Version: 4.18.1 | Last Updated: December 2024

---

## Overview

The RADIANT Seed Data System manages versioned AI provider and model configurations that are used to populate the AI Registry during fresh installations. Seed data is stored separately from packages, can be versioned independently, and is selectable when building deployment packages.

## Architecture

### SEED DATA ARCHITECTURE

```
config/seeds/
    registry.json      # Index of all seed versions
    v1/
        manifest.json # Seed data version 1.0.0
        providers.json # Version metadata and stats
        external-models.json # 21 external providers
        self-hosted-models.json # 50+ external models
        services.json     # 38 self-hosted models
    v2/                  # Future seed versions...
```

Build Time:

```
build-package.sh      Select seed version      Include in package
--seed-version 1
```

Deploy Time (INSTALL only):

```
DeploymentService    Read seeds from package    INSERT to database
.executeInstall()
```

## Critical Rules

### Rule 1: NO HARDCODING IN DEPLOYER APP

The Swift Deployer app **MUST NOT** contain hardcoded lists of providers or models:

```
// WRONG - Never do this
let providers = ["openai", "anthropic", "google", ...]

// CORRECT - Fetch from Radiant API after deployment
let providers = try await radiantAPI.fetchProviders()
```

### Rule 2: INSTALLER SEEDS, UPDATER PRESERVES

Mode	Seed Behavior
INSTALL	Seeds database with complete provider/model list
UPDATE	NEVER touches AI Registry - preserves admin customizations
ROLLBACK	Restores from snapshot - does not re-seed

### Rule 3: ADMIN CONTROLS ALL

Everything in seed data is **editable by the administrator** post-deployment: - Enable/disable providers and models - Change pricing markup - Add new providers/models - Delete providers/models

## Seed Data Structure

### manifest.json

```
{
  "version": "1.0.0",
  "name": "RADIANT AI Registry Seed Data",
  "description": "Complete provider and model seed data for fresh installations",
  "createdAt": "2024-12-25T00:00:00Z",
  "updatedAt": "2024-12-25T00:00:00Z",
  "compatibility": {
    "minRadiantVersion": "4.16.0",
    "maxRadiantVersion": "5.0.0"
  },
  "files": {
    "providers": "providers.json",
    "externalModels": "external-models.json",
    "selfHostedModels": "self-hosted-models.json",
    "services": "services.json"
  },
  "stats": {
    "externalProviders": 21,
    "externalModels": 50,
    "selfHostedModels": 38,
```

```

    "services": 5
  },
  "pricing": {
    "externalMarkup": 1.40,
    "selfHostedMarkup": 1.75
  }
}

```

## providers.json

Each provider includes:

Field	Description
<code>id</code>	Unique identifier
<code>name</code>	Internal name
<code>displayName</code>	Human-readable name
<code>category</code>	Provider category (text_generation, image_generation, etc.)
<code>apiBaseUrl</code>	API endpoint
<code>authType</code>	Authentication type (bearer, api_key, iam)
<code>secretName</code>	AWS Secrets Manager path for API key
<code>features</code>	Supported features (streaming, vision, etc.)
<code>compliance</code>	Compliance certifications (SOC2, GDPR, HIPAA)
<code>rateLimit</code>	Rate limiting configuration

## external-models.json

Each model includes:

Field	Description
<code>id</code>	Unique model identifier
<code>providerId</code>	Reference to provider
<code>modelId</code>	Provider's model ID
<code>litellmId</code>	LiteLLM routing ID
<code>category</code>	Model category
<code>capabilities</code>	Model capabilities
<code>contextWindow</code>	Max input tokens
<code>maxOutput</code>	Max output tokens
<code>pricing</code>	Cost per 1K tokens + markup
<code>minTier</code>	Minimum tier required

## self-hosted-models.json

Each self-hosted model includes:

Field	Description
<code>id</code>	Unique model identifier
<code>instanceType</code>	SageMaker instance type
<code>thermal</code>	Thermal management config (COLD/WARM/HOT)
<code>license</code>	Open-source license
<code>pricing</code>	Hourly rate + per-unit pricing
<code>minTier</code>	Minimum tier required (typically 3+)

## Building Packages with Seed Data

### List Available Seed Versions

```
./tools/scripts/build-package.sh --list-seeds
```

Output:

```
Available Seed Data Versions:  
v1.0.0 - 21 providers, 50 external models, 38 self-hosted models
```

### Build with Specific Seed Version

```
# Use default (latest) seed version  
./tools/scripts/build-package.sh  
  
# Use specific seed version  
./tools/scripts/build-package.sh --seed-version 1
```

### Package Manifest with Seed Data

The generated package manifest includes seed data information:

```
{  
  "schemaVersion": "2.1",  
  "package": {  
    "version": "4.18.1"  
  },  
  "seedData": {  
    "version": "1.0.0",  
    "hash": "abc123...",  
    "externalProviders": 21,  
    "externalModels": 50,  
    "selfHostedModels": 38,  
    "services": 5  
  },  
  "installBehavior": {  
    "seedAIRegistry": true  
  },  
  "updateBehavior": {  
    "seedAIRegistry": false  
  }  
}
```

---

## Seed Data Categories

### External Providers (21)

Category	Providers
Text Generation	OpenAI, Anthropic, Google, xAI, DeepSeek, Mistral, Cohere
Image Generation	OpenAI Images, Stability AI, FLUX
Video Generation	Runway, Luma AI

Category	Providers
Audio	ElevenLabs, OpenAI Audio
Embeddings	OpenAI Embeddings, Voyage AI
Search	Perplexity
3D Generation	Meshy
Self-Hosted	SageMaker (internal)

## External Models (50+)

Category	Example Models
Text	GPT-4o, Claude Sonnet 4, Gemini 2.0, Grok 3, DeepSeek R1
Reasoning	O1, O3 Mini, DeepSeek Reasoner
Code	Codestral
Image	DALL-E 3, Stable Diffusion 3, FLUX Pro
Video	Gen-3 Alpha, Ray 2
Audio	Whisper, TTS-1, Multilingual V2

## Self-Hosted Models (38)

Category	Models
Vision Classification	EfficientNet, Swin Transformer, CLIP
Object Detection	YOLOv8 (Nano/Small/Medium/XLarge), Grounding DINO
Segmentation	SAM, SAM 2, MobileSAM
Speech	Whisper Large V3, Parakeet TDT
Scientific	AlphaFold 2, ESM-2
Medical	nnU-Net, MedSAM
Geospatial	Prithvi 100M/600M
3D	Nerfstudio
LLM	Mistral 7B, Llama 3 70B

---

## Pricing Structure

### External Providers

Default markup: **40% (1.40x)**

Example: GPT-4o - Provider cost: \$0.0025/1K input, \$0.01/1K output - Tenant cost: \$0.0035/1K input, \$0.014/1K output

### Self-Hosted Models

Default markup: **75% (1.75x)**

Example: YOLOv8 Medium - Infrastructure cost: ~\$2.47/hour + \$0.005/image - Tenant cost: ~\$4.32/hour + \$0.00875/image

---

## Creating New Seed Versions

### 1. Create Version Directory

```
mkdir config/seeds/v2
```

### 2. Create Required Files

- manifest.json - Version metadata
- providers.json - Provider definitions
- external-models.json - External model definitions
- self-hosted-models.json - Self-hosted model definitions
- services.json - Service definitions

### 3. Update Registry

Add new version to config/seeds/registry.json:

```
{
  "versions": [
    {
      "version": "2.0.0",
      "directory": "v2",
      "releaseDate": "2025-01-15",
      "status": "stable",
      "changelog": "Added new providers and models..."
    },
    // ... existing versions
  ]
}
```

### 4. Test Build

```
./tools/scripts/build-package.sh --seed-version 2
```

---

## Database Seeding

During fresh installation, the DeploymentService generates SQL migrations from seed data:

```
-- Only runs if providers table is empty
DO $$
BEGIN
  IF NOT EXISTS (SELECT 1 FROM providers LIMIT 1) THEN
    -- Insert providers
    INSERT INTO providers (...) VALUES (...);

    -- Insert external models
    INSERT INTO models (...) VALUES (...);

    -- Insert self-hosted models
    INSERT INTO self_hosted_models (...) VALUES (...);
  END IF;
END $$;
```

Key behaviors:

- Uses `ON CONFLICT DO NOTHING` to preserve admin changes
- Only runs on fresh install (empty database)
- Logs completion with model counts

---

## Swift Service API

### SeedDataService

```
actor SeedDataService {
    /// List available seed versions
    func listAvailableSeedVersions() async throws -> [SeedDataInfo]

    /// Load complete seed data for a version
    func loadSeedData(version: String) async throws -> SeedData

    /// Generate SQL migration from seed data
    func generateSeedMigration(seedData: SeedData) -> String
}
```

### Usage in DeploymentService

```
func executeInstall(...) async throws -> DeploymentExecutionResult {
    // Load seed data from package
    let seedData = try await seedDataService.loadSeedData(
        version: package.manifest.seedData?.version ?? "1.0.0"
    )

    // Generate and run seed migration
    let seedSQL = seedDataService.generateSeedMigration(seedData: seedData)
    try await runMigration(sql: seedSQL)
}
```

---

## Related Documentation

- Deployer Architecture - Deployment modes and package management
- Deployer Admin Guide - User-facing deployment documentation
- API Reference - Provider and model API endpoints

# RADIANT v4.17.0 - Troubleshooting Guide

## Common Issues and Solutions

### CDK Deployment Failures

Issue	Likely Cause	Solution
Bootstrap failed	Wrong account/region	Verify <code>aws sts get-caller-identity</code>
Stack timeout	Slow resource creation	Check CloudFormation events in AWS Console
Resource limit	Service quota exceeded	Request quota increase via AWS Service Quotas
IAM permission denied	Insufficient permissions	Ensure IAM user has AdministratorAccess
Circular dependency	Stack references	Check stack dependencies in CDK code
Asset upload failed	S3 bucket permissions	Verify CDK bootstrap bucket exists

### Aurora Database Issues

Issue	Likely Cause	Solution
Connection refused	Security group rules	Verify Lambda SG can reach Aurora SG on port 5432
Authentication failed	Wrong credentials	Check Secrets Manager for correct credentials
Connection timeout	Missing VPC endpoints	Add RDS VPC endpoint to private subnets
Too many connections	Connection exhaustion	Use RDS Proxy or increase <code>max_connections</code>
Slow queries	Missing indexes	Run EXPLAIN ANALYZE and add appropriate indexes
RLS blocking access	Tenant ID not set	Ensure <code>app.current_tenant_id</code> is set in session

### Lambda Function Errors

Issue	Likely Cause	Solution
Cold start > 10s	VPC attachment	Use provisioned concurrency for critical functions
Timeout	Slow downstream services	Increase timeout, check DB/API latency
Out of memory	Large payloads/responses	Increase memory allocation (also increases CPU)
Permission denied	IAM role misconfigured	Check Lambda execution role policies
Module not found	Missing dependency	Verify all dependencies in package.json
Handler not found	Incorrect handler path	Check function configuration in CDK

## LiteLLM / ECS Issues

Issue	Likely Cause	Solution
503 Service Unavailable	ECS task unhealthy	Check ECS service events and task logs
Provider timeout	Invalid API key	Verify provider secrets in Secrets Manager
Rate limited	Too many requests	Implement exponential backoff retry
Wrong model response	Model misconfigured	Check config.yaml model mappings
Container crashes	Memory exhaustion	Increase task memory in CDK
No healthy targets	Health check failing	Verify health check endpoint returns 200

## SageMaker Issues (Tier 3+)

Issue	Likely Cause	Solution
Endpoint failed to create	Insufficient capacity	Try different instance type or region
InvocationError	Model loading failed	Check CloudWatch logs for model errors
Slow cold start	Large model size	Use warm pools or smaller model variant
Capacity error	Instance quota reached	Request SageMaker quota increase
Timeout	Long inference time	Increase endpoint timeout or optimize model

## Cognito Authentication Issues

Issue	Likely Cause	Solution
Invalid grant	Expired refresh token	Re-authenticate user

Issue	Likely Cause	Solution
User not confirmed	Email not verified	Check email or manually confirm user
MFA required	MFA not set up	Complete MFA setup flow
Invalid client	Wrong client ID	Verify app client ID in configuration
Callback URL mismatch	URL not whitelisted	Add URL to allowed callbacks in Cognito

## Admin Dashboard Issues

Issue	Likely Cause	Solution
403 Forbidden	CloudFront OAC issue	Verify S3 bucket policy allows CloudFront
API calls fail	CORS configuration	Check API Gateway CORS settings
Login redirect loop	Cookie domain mismatch	Verify cookie domain matches site domain
Blank page	Build error	Check <code>next build</code> output for errors
Slow load	Large bundle size	Enable code splitting and lazy loading

## Log Locations

Component	CloudWatch Log Group
API Gateway	/aws/api-gateway/radiant-{env}-api
Lambda Functions	/aws/lambda/Radiant-{env}-*
LiteLLM (ECS)	/ecs/radiant-{env}-litellm
SageMaker Endpoints	/aws/sagemaker/Endpoints/radiant-*
Aurora PostgreSQL	/aws/rds/cluster/radiant-{env}/postgresql
CloudFront	Standard CloudFront logs in S3

## Viewing Logs

```
# Tail Lambda logs in real-time
aws logs tail /aws/lambda/Radiant-dev-router --follow

# View ECS logs
aws logs tail /ecs/radiant-dev-litellm --follow

# Search logs for errors
aws logs filter-log-events \
--log-group-name /aws/lambda/Radiant-dev-router \
--filter-pattern "ERROR" \
--start-time $(date -d '1 hour ago' +%s)000
```

## Health Check Endpoints

```
# Platform API health
curl https://api.YOUR_DOMAIN/health
# Expected: {"status": "healthy", "version": "4.17.0"}

# LiteLLM health
curl https://api.YOUR_DOMAIN/v2/litellm/health
# Expected: {"status": "healthy"}

# Admin API health
curl https://admin-api.YOUR_DOMAIN/health

# Model registry status
curl https://api.YOUR_DOMAIN/v2/models/status
```

---

## Emergency Procedures

### Database Restore from Snapshot

```
# List available snapshots
aws rds describe-db-cluster-snapshots \
  --db-cluster-identifier radiant-prod-cluster \
  --query 'DBClusterSnapshots[*].[DBClusterSnapshotIdentifier,SnapshotCreateTime]' \
  --output table

# Restore from snapshot
aws rds restore-db-cluster-from-snapshot \
  --db-cluster-identifier radiant-prod-restored \
  --snapshot-identifier your-snapshot-id \
  --engine aurora-postgresql \
  --vpc-security-group-ids sg-xxx \
  --db-subnet-group-name radiant-prod-db-subnet
```

### Point-in-Time Recovery

```
aws rds restore-db-cluster-to-point-in-time \
  --source-db-cluster-identifier radiant-prod-cluster \
  --db-cluster-identifier radiant-prod-recovered \
  --restore-to-time "2024-12-20T10:00:00Z" \
  --vpc-security-group-ids sg-xxx \
  --db-subnet-group-name radiant-prod-db-subnet
```

### Rollback CDK Deployment

```
# Rollback to previous deployment
cd packages/infrastructure
npx cdk deploy Radiant-prod-API \
  --context environment=prod \
  --context tier=3 \
  --rollback
```

## Disable Problematic Model

```
-- Connect to Aurora and run:  
UPDATE models  
SET status = 'disabled',  
    disabled_reason = 'Emergency disable due to errors',  
    updated_at = NOW()  
WHERE model_id = 'problematic-model-id';
```

## Force Scale Down SageMaker

```
# Scale endpoint to 0 instances  
aws sagemaker update-endpoint-weights-and-capacities \  
  --endpoint-name radiant-prod-model-endpoint \  
  --desired-weights-and-capacities '[{"VariantName": "AllTraffic", "DesiredInstanceCount": 0}]'
```

---

## Performance Benchmarks

Metric	Target	Acceptable	Action if Exceeded
API Gateway p50 latency	< 50ms	< 100ms	Check Lambda cold starts
API Gateway p99 latency	< 200ms	< 500ms	Enable provisioned concurrency
Chat streaming start	< 500ms	< 1s	Check LiteLLM/provider latency
Admin dashboard load	< 2s	< 3s	Optimize bundle, enable CDN caching
Model warm-up time	< 3 min	< 5 min	Use larger instance or warm pools
Aurora query latency	< 10ms	< 50ms	Add indexes, optimize queries

---

## Support Checklist

When reporting issues, include:

1. **Environment:** dev/staging/prod
  2. **Tier:** 1-5
  3. **Error message:** Full error text
  4. **Request ID:** From response headers
  5. **Timestamp:** When the error occurred
  6. **Steps to reproduce:** What actions led to the error
  7. **Relevant logs:** CloudWatch log excerpts
- 

## Useful AWS CLI Commands

```
# Check stack status  
aws cloudformation describe-stacks --stack-name Radiant-dev-API \  
  --region us-west-2
```

```
--query 'Stacks[0].StackStatus'

# List recent CloudFormation events
aws cloudformation describe-stack-events --stack-name Radiant-dev-API \
--query 'StackEvents[0:10].[Timestamp,ResourceStatus,ResourceType,LogicalResourceId]' \
--output table

# Check Lambda function configuration
aws lambda get-function-configuration --function-name Radiant-dev-router

# List ECS services
aws ecs list-services --cluster radiant-dev-cluster

# Describe ECS service
aws ecs describe-services --cluster radiant-dev-cluster \
--services radiant-dev-litellm

# Check Secrets Manager secret
aws secretsmanager get-secret-value --secret-id radiant/dev/db-credentials \
--query 'SecretString' --output text | jq .
```

# RADIANT CDK Stack Dependencies

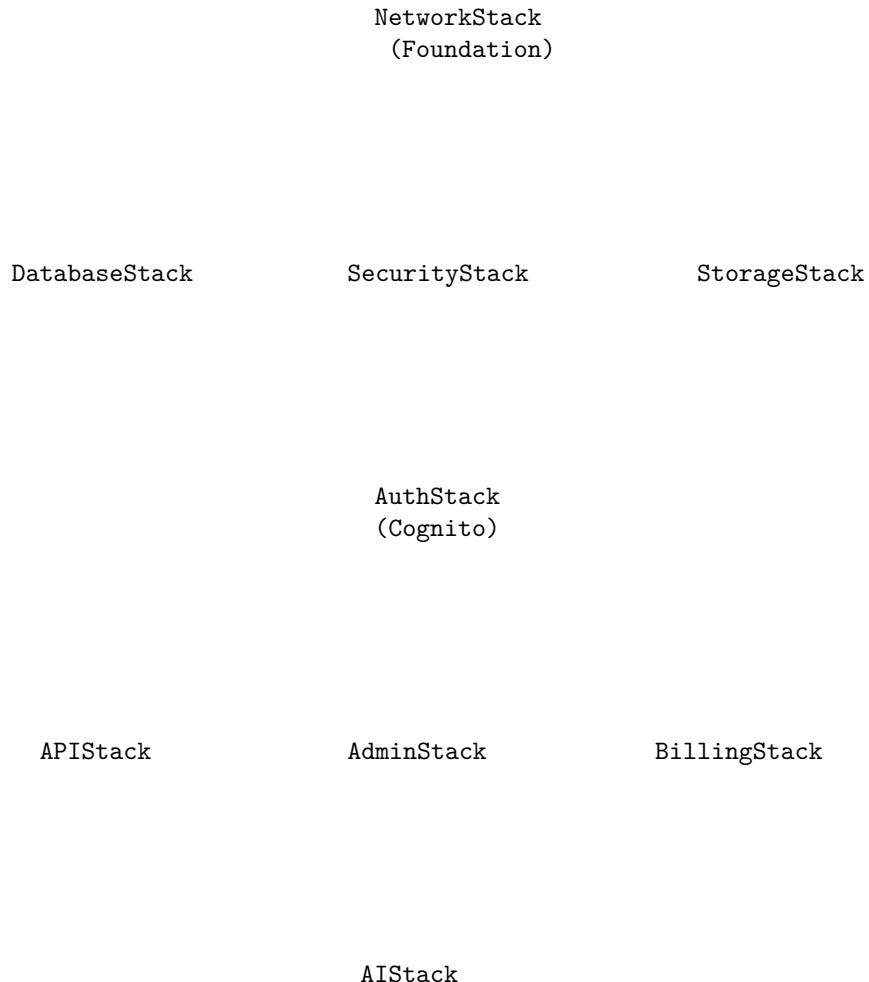
## Technical Reference

Version: 4.18.1 | Last Updated: December 2024

This document defines the explicit dependency graph for RADIANT CDK stacks to ensure correct deployment ordering.

---

## Stack Dependency Graph



(Models/LLM)

ThermalStack      PerceptionSvc      SageMaker  
(Scaling)            (Vision)            (Self-Host)

---

## Stack Definitions

### Layer 1: Foundation

Stack	Purpose	Exports
<b>NetworkStack</b>	VPC, Subnets, NAT Gateways	VPC ID, Subnet IDs, Security Group IDs

### Layer 2: Core Infrastructure

Stack	Purpose	Dependencies	Exports
<b>DatabaseStack</b>	Aurora PostgreSQL, RDS Proxy	NetworkStack	Cluster ARN, Secret ARN, Proxy Endpoint
<b>SecurityStack</b>	KMS Keys, WAF, GuardDuty	NetworkStack	Key ARNs, WAF ACL ARN
<b>StorageStack</b>	S3 Buckets, CloudFront	NetworkStack	Bucket ARNs, Distribution ID

### Layer 3: Authentication

Stack	Purpose	Dependencies	Exports
<b>AuthStack</b>	Cognito User/Identity Pools	Database, Security, Storage	Pool IDs, Client IDs

### Layer 4: Application Services

Stack	Purpose	Dependencies	Exports
<b>APIStack</b>	API Gateway, Lambda handlers	Auth, Database, Security	API Endpoint, Lambda ARNs
<b>AdminStack</b>	Admin API, Dashboard hosting	Auth, Database, Security	Admin API Endpoint
<b>BillingStack</b>	Stripe integration, Usage tracking	Auth, Database	Billing API Endpoint

## Layer 5: AI Services

Stack	Purpose	Dependencies	Exports
<b>AISStack</b>	LiteLLM, Model routing	API, Database, Security	AI Gateway Endpoint

## Layer 6: Specialized Services

Stack	Purpose	Dependencies	Exports
<b>ThermalStack</b>	Model scaling, State management	AI, Database	Thermal API Endpoint
<b>PerceptionStack</b>	Computer vision pipeline	AI, Storage	Perception API Endpoint
<b>SageMakerStack</b>	Self-hosted model endpoints	AI, Network	Endpoint ARNs

## CDK Implementation

### Explicit Dependencies

```
// packages/infrastructure/lib/main.ts

import { App } from 'aws-cdk-lib';

const app = new App();

// Layer 1
const networkStack = new NetworkStack(app, 'Network', { env });

// Layer 2
const databaseStack = new DatabaseStack(app, 'Database', {
  env,
  vpc: networkStack.vpc,
});
databaseStack.addDependency(networkStack);

const securityStack = new SecurityStack(app, 'Security', {
  env,
  vpc: networkStack.vpc,
});
securityStack.addDependency(networkStack);

const storageStack = new StorageStack(app, 'Storage', {
  env,
  vpc: networkStack.vpc,
});
storageStack.addDependency(networkStack);

// Layer 3
```

```

const authStack = new AuthStack(app, 'Auth', {
  env,
  database: databaseStack,
  security: securityStack,
  storage: storageStack,
});
authStack.addDependency(databaseStack);
authStack.addDependency(securityStack);
authStack.addDependency(storageStack);

// Layer 4
const apiStack = new APIStack(app, 'API', {
  env,
  auth: authStack,
  database: databaseStack,
  security: securityStack,
});
apiStack.addDependency(authStack);

// ... continue for remaining stacks

```

## Cross-Stack References

```

// Example: APIStack referencing DatabaseStack exports

export class APIStack extends Stack {
  constructor(scope: Construct, id: string, props: APIStackProps) {
    super(scope, id, props);

    // Use exports from DatabaseStack
    const clusterArn = props.database.clusterArn;
    const secretArn = props.database.secretArn;

    // Create Lambda with database access
    const handler = new Function(this, 'ApiHandler', {
      environment: {
        AURORA_CLUSTER_ARN: clusterArn,
        AURORA_SECRET_ARN: secretArn,
      },
    });
  }
}

```

---

## Deployment Order

### Fresh Install

```

# Deploy in dependency order
cdk deploy NetworkStack
cdk deploy DatabaseStack SecurityStack StorageStack --parallel
cdk deploy AuthStack
cdk deploy APIStack AdminStack BillingStack --parallel

```

```
cdk deploy AIStack  
cdk deploy ThermalStack PerceptionStack SageMakerStack --parallel
```

## Update (with dependencies)

```
# CDK handles ordering automatically when using addDependency  
cdk deploy --all
```

## Selective Deployment

```
# Deploy specific stack and its dependencies  
cdk deploy AIStack --require-approval never
```

---

## Rollback Considerations

Stack	Rollback Safe	Notes
NetworkStack	Caution	May affect all dependent stacks
DatabaseStack	Caution	Requires DB snapshot for data preservation
SecurityStack	Safe	KMS keys have deletion protection
StorageStack	Caution	S3 buckets may have data
AuthStack	Safe	Cognito pools preserved
APIStack	Safe	Stateless Lambda functions
AdminStack	Safe	Stateless
BillingStack	Caution	May have pending transactions
AIStack	Safe	Stateless routing
ThermalStack	Safe	State in database
SageMakerStack	Caution	May have running endpoints

---

## Validation Script

```
#!/bin/bash  
# tools/scripts/validate-stack-deps.sh  
  
echo "Validating CDK stack dependencies..."  
  
# Check for circular dependencies  
cdk synth --quiet 2>&1 | grep -i "circular" && {  
    echo " Circular dependency detected!"  
    exit 1  
}
```

```
# Verify deployment order
cdk diff --all 2>&1 | head -50

echo " Stack dependencies validated"
```

---

## Related Documentation

- [Deployment Guide](#) - Full deployment procedures
- [Deployer Architecture](#) - Package and deployment flow