

RADIANT Platform - Administrator Guide

RADIANT Team

2025-12-25

Contents

RADIANT Platform - Administrator Guide	3
Table of Contents	3
1. Introduction	3
1.1 What is RADIANT?	3
1.2 Administrator Roles	4
1.3 Key Concepts	4
2. Accessing the Admin Dashboard	5
2.1 URL and Login	5
2.2 First Login	5
2.3 Session Management	6
2.4 Password Requirements	6
3. Dashboard Overview	6
3.1 Main Dashboard	6
3.2 Navigation Menu	7
4. Tenant Management	7
4.1 Viewing Tenants	7
4.2 Creating a Tenant	7
4.3 Tenant Details	8
4.4 Tenant Actions	8
4.5 Data Isolation	8
5. User & Administrator Management	9
5.1 Administrator Roles	9
5.2 Managing Administrators	9
5.3 Viewing Tenant Users	9
5.4 User Actions	10
6. AI Model Configuration	10
6.1 Model Registry	10
6.2 Model Categories	10
6.3 Model Configuration	11
6.4 Self-Hosted Models	12
6.5 Thermal States (Self-Hosted)	12
7. Provider Management	12
7.1 External Providers	12
7.2 Adding Provider Credentials	13

7.3 Provider Health Monitoring	13
7.4 Fallback Configuration	13
8. Billing & Subscriptions	13
8.1 Subscription Tiers	13
8.2 Credit System	14
8.3 Managing Subscriptions	14
8.4 Usage Reports	14
8.5 Billing Alerts	14
9. Storage Management	15
9.1 Storage Overview	15
9.2 Storage Tiers	15
9.3 File Management	15
10. Orchestration & Neural Engine	16
10.1 Brain Router	16
10.2 Neural Patterns	16
10.3 Workflow Templates	16
11. Localization	16
11.1 Translation Management	16
11.2 Supported Languages	16
11.3 AI Translation	17
12. Configuration Management	17
12.1 System Configuration	17
12.2 Tenant Overrides	17
12.3 SSM Parameters	17
13. Security & Compliance	18
13.1 Security Dashboard	18
13.2 Anomaly Detection	18
13.3 Compliance Reports	18
13.4 Generating Reports	18
14. Cost Analytics	19
14.1 Cost Dashboard	19
14.2 Cost Alerts	19
14.3 Cost Optimization	19
15. A/B Testing & Experiments	20
15.1 Experiment Dashboard	20
15.2 Creating an Experiment	20
15.3 Statistical Analysis	20
16. Audit & Monitoring	20
16.1 Audit Logs	20
16.2 Log Filtering	21
16.3 Log Export	21
16.4 Real-Time Monitoring	21
17. Database Migrations	21
17.1 Migration Workflow	21
17.2 Pending Migrations	21
17.3 Approving Migrations	22
18. API Management	22
18.1 API Keys	22

18.2 Rate Limiting	22
18.3 Webhooks	22
19. Troubleshooting	23
19.1 Common Issues	23
19.2 Support Resources	23
19.3 Log Locations	23
Appendix: Quick Reference	24
Keyboard Shortcuts	24
Status Indicators	24

RADIANT Platform - Administrator Guide

Complete guide for managing the RADIANT AI Platform via the Admin Dashboard

Version: 4.18.1 | Last Updated: December 2024

Table of Contents

1. Introduction
 2. Accessing the Admin Dashboard
 3. Dashboard Overview
 4. Tenant Management
 5. User & Administrator Management
 6. AI Model Configuration
 7. Provider Management
 8. Billing & Subscriptions
 9. Storage Management
 10. Orchestration & Neural Engine
 11. Localization
 12. Configuration Management
 13. Security & Compliance
 14. Cost Analytics
 15. A/B Testing & Experiments
 16. Audit & Monitoring
 17. Database Migrations
 18. API Management
 19. Troubleshooting
-

1. Introduction

1.1 What is RADIANT?

RADIANT is a multi-tenant AWS SaaS platform providing unified access to 106+ AI models through:

- **50 External Provider Models:** OpenAI, Anthropic, Google, xAI, DeepSeek, and more
- **56 Self-Hosted Models:** Running on AWS SageMaker for cost control and privacy
- **Intelligent Routing:** Brain router for optimal model selection
- **Neural Engine:** Personalization learning from user interactions

1.2 Administrator Roles

Role	Permissions	Use Case
Super Admin	Full access to all features	Platform owner
Admin	Tenant management, billing, models	Operations team
Operator	Read access, limited actions	Support team
Auditor	Read-only access to logs	Compliance team

Role Details **Super Admin** - The highest privilege level with unrestricted access: - Create and delete tenants - Manage all administrators - Access all billing and financial data - Modify system-wide configuration - Approve production database migrations - Impersonate any tenant for debugging - Access compliance and audit reports - Typically limited to 1-3 people (CTO, lead engineer)

Admin - Day-to-day operations management: - Create and modify tenants (cannot delete) - Manage users within tenants - Configure AI models and providers - View billing data (cannot modify pricing) - Monitor system health - Cannot access other admin accounts - Typically assigned to operations team members

Operator - Limited support and monitoring: - View tenant information (read-only) - View user issues and support tickets - Monitor system health dashboards - Cannot modify any configuration - Cannot access billing or sensitive data - Typically assigned to support staff

Auditor - Compliance and security review: - Full read access to audit logs - Access to compliance reports - Cannot modify anything - Cannot view sensitive data (API keys, passwords) - Access is logged for compliance - Typically assigned to compliance officers or external auditors

1.3 Key Concepts

Concept	Description
Tenant	Organization with isolated data
User	End-user within a tenant
Subscription	Billing tier (1-7)
Credits	Currency for AI usage
API Key	Authentication for API access
App	Consumer application (Think Tank, etc.)

Tenant Architecture Explained A **Tenant** represents a complete organization using RADIAN. Each tenant has:

- **Complete Data Isolation:** All data is stored with tenant IDs and protected by PostgreSQL Row-Level Security (RLS). One tenant can never access another tenant's data, even if there's a bug in application code.

- **Separate Billing:** Each tenant has its own subscription, credit balance, and usage tracking. Costs are attributed to the correct tenant automatically.
- **Custom Configuration:** Tenants can customize model access, rate limits, and feature flags without affecting other tenants.
- **User Management:** Each tenant manages their own users, roles, and permissions independently.

User vs Administrator **Users** are end-users who interact with RADIANT-powered applications like Think Tank. They:

- Sign up and log in via Cognito
- Use AI models through the API or applications
- Have credits deducted for usage
- Cannot access the Admin Dashboard

Administrators manage the RADIANT platform itself. They:

- Access the Admin Dashboard
- Manage tenants, users, and billing
- Configure AI models and providers
- Have no credits (administrative access is separate)

Credit System Explained Credits are RADIANT's universal currency for AI usage:

- **1 credit = \$0.01 USD** (configurable per deployment)
- Different models cost different amounts based on their API pricing
- Credits are deducted in real-time as requests complete
- Tenants can purchase credits or receive them through subscriptions
- Credits can be tracked, audited, and reported on

Example Credit Costs:

Model	Cost per 1K tokens
GPT-4o	5 credits input, 15 credits output
GPT-4o-mini	0.5 credits input, 1.5 credits output
Claude 3.5 Sonnet	3 credits input, 15 credits output
Self-hosted Llama	0.2 credits (all)

API Key Types RADIANT supports multiple API key types:

- **User API Keys:** Tied to a specific user, inherit user's permissions
 - **Service API Keys:** For server-to-server communication, not tied to a user
 - **Admin API Keys:** For administrative operations, require elevated permissions
 - **Scoped Keys:** Limited to specific models, endpoints, or rate limits
-

2. Accessing the Admin Dashboard

2.1 URL and Login

1. Navigate to: <https://admin.your-domain.com>
2. Enter your email address
3. Enter your password
4. Complete MFA verification (required)

2.2 First Login

On first login:

1. You'll receive a temporary password via email
2. Enter the temporary password

3. Set a new password (12+ characters, mixed case, numbers, symbols)
 4. Set up MFA using an authenticator app
 5. You'll be redirected to the dashboard

2.3 Session Management

Setting	Value
Session Duration	8 hours
Idle Timeout	30 minutes
Concurrent Sessions	3 maximum
Remember Device	30 days

2.4 Password Requirements

- Minimum 12 characters
 - At least one uppercase letter
 - At least one lowercase letter
 - At least one number
 - At least one special character
 - Cannot reuse last 10 passwords

3. Dashboard Overview

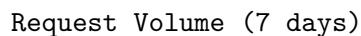
3.1 Main Dashboard

The dashboard displays key metrics at a glance:

RADIANT Admin Dashboard

Welcome, Admin

Tenants	Users	Requests	Revenue
142	8,456	2.3M	\$45,230
+12%	+8%	+23%	+15%



Mon Tue Wed Thu Fri Sat Sun

Recent Activity:

- New tenant: Acme Corp (2 minutes ago)
 - Model enabled: claude-3-opus (15 minutes ago)
 - Alert: High API error rate (1 hour ago)

3.2 Navigation Menu

Section	Description
Dashboard	Overview and metrics
Tenants	Tenant management
Users	User management
Models	AI model configuration
Providers	Provider management
Billing	Subscriptions and credits
Storage	Storage usage
Orchestration	Neural engine settings
Localization	Translation management
Configuration	System settings
Security	Security monitoring
Compliance	Compliance reports
Experiments	A/B testing
Cost	Cost analytics
Audit	Audit logs
Migrations	Database migrations
Notifications	System alerts
Settings	Personal settings

4. Tenant Management

4.1 Viewing Tenants

Navigate to **Tenants** to see all organizations:

Column	Description
Name	Organization name
Plan	Subscription tier
Users	User count
Status	Active/Suspended/Trial
Created	Creation date
Last Active	Last API call

4.2 Creating a Tenant

1. Click “+ New Tenant”
2. Fill in required fields:
 - **Name:** Organization name
 - **Slug:** URL-friendly identifier

- **Plan:** Initial subscription tier
 - **Admin Email:** Primary admin email
3. Configure optional settings:
 - Custom domain
 - Branding settings
 - Feature flags
 4. Click “**Create Tenant**”

4.3 Tenant Details

View comprehensive tenant information:

Tenant: Acme Corporation

[Overview](#) [Users](#) [Billing](#) [Settings](#)

Tenant ID: tn_abc123xyz
 Status: Active
 Plan: Professional (Tier 4)
 Created: 2024-01-15
 Last Active: 2 minutes ago

Usage This Month:

API Requests: 145,234
 Tokens Used: 12.5M
 Storage: 2.3 GB
 Credits Used: \$1,234.56

[Edit] [Suspend] [Delete] [Impersonate]

4.4 Tenant Actions

Action	Description	Permission
Edit	Modify tenant settings	Admin
Suspend	Temporarily disable	Admin
Delete	Permanently remove	Super Admin
Impersonate	Login as tenant admin	Super Admin
Export	Export tenant data	Admin

4.5 Data Isolation

Each tenant has complete data isolation:

- Separate database rows with RLS
 - Unique API keys
 - Isolated storage buckets
 - Independent usage tracking
-

5. User & Administrator Management

5.1 Administrator Roles

Role	Dashboard Access	API Access	Billing	Audit
Super Admin	Full	Full	Full	Full
Admin	Full	Full	Read	Read
Operator	Read	Read	None	Read
Auditor	Logs only	None	None	Full

5.2 Managing Administrators

Navigate to **Administrators** to:

1. **Invite New Admin:**
 - Click “+ Invite Administrator”
 - Enter email address
 - Select role
 - Click “Send Invitation”
2. **Modify Admin:**
 - Click on administrator row
 - Edit role or permissions
 - Click “Save Changes”
3. **Remove Admin:**
 - Click “Remove” button
 - Confirm removal
 - Admin’s sessions are invalidated immediately

5.3 Viewing Tenant Users

Navigate to **Tenants** → [Tenant] → **Users** to see:

Field	Description
Email	User email
Name	Display name
Role	Tenant role
Status	Active/Invited/Disabled
Last Login	Last authentication
API Keys	Number of active keys

5.4 User Actions

Action	Description
Reset Password	Send password reset email
Disable	Prevent login
Enable	Restore access
Delete	Remove user data
View Sessions	See active sessions

6. AI Model Configuration

6.1 Model Registry

Navigate to **Models** to see all available models:

AI Models				
106 Total				
Filter: [All] Category: [All] Status: [Enabled]				
Model	Provider	Category	Tier	Status
gpt-4o	OpenAI	Chat	1	Enabled
gpt-4o-mini	OpenAI	Chat	1	Enabled
claude-3-opus	Anthropic	Chat	2	Enabled
claude-3-sonnet	Anthropic	Chat	1	Enabled
gemini-pro	Google	Chat	1	Enabled
llama-3.1-70b	Self-Host	Chat	3	Enabled
whisper-large	Self-Host	Audio	3	Disabled

[+ Add Model] [Import Models] [Export Config]

6.2 Model Categories

Category	Description	Example Models
Chat/LLM	Text generation	GPT-4o, Claude 3, Gemini
Embedding	Vector embeddings	text-embedding-3-large
Vision	Image understanding	GPT-4V, Claude Vision
Audio	Speech-to-text	Whisper, Deepgram
Image	Image generation	DALL-E 3, Stable Diffusion
Code	Code generation	Codestral, DeepSeek Coder
Scientific	Research models	BioGPT, ChemLLM

Category Details Chat/LLM (Large Language Models): The core of RADIANT. These models handle conversational AI, content generation, summarization, and general-purpose text tasks. They're the most commonly used and include flagship models from OpenAI, Anthropic, Google, and open-source alternatives.

Embedding Models: Convert text into numerical vectors for semantic search, similarity matching, and retrieval-augmented generation (RAG). Essential for building knowledge bases and search functionality. Vectors are typically 1536-3072 dimensions.

Vision Models: Analyze images, extract text (OCR), describe visual content, and answer questions about images. Increasingly important for document processing, accessibility, and multimodal applications.

Audio Models: Transcribe speech to text, translate audio, and identify speakers. Whisper is the most popular, offering excellent accuracy across 99 languages. Used for meeting transcription, accessibility, and voice interfaces.

Image Generation: Create images from text descriptions. DALL-E 3 offers the best prompt following, while Stable Diffusion provides more customization options. Consider content policies when enabling these.

Code Models: Specialized for programming tasks including code generation, explanation, debugging, and refactoring. Some are fine-tuned on specific languages or frameworks.

Scientific Models: Domain-specific models trained on scientific literature. Useful for research applications but require careful evaluation for accuracy.

6.3 Model Configuration

Click on a model to configure:

Setting	Description
Enabled	Available for use
Min Tier	Minimum subscription tier
Rate Limits	Requests per minute
Max Tokens	Maximum context/output
Temperature Range	Allowed temperature values
Price Override	Custom pricing

Configuration Settings Explained **Enabled:** When disabled, the model is hidden from users and API requests return “model not found”. Use this to temporarily remove models during maintenance or to restrict access to specific models.

Min Tier: Sets the minimum subscription tier required to access this model. For example, setting GPT-4 to Tier 2 means Free tier users cannot use it. This helps control costs and create upgrade incentives.

Rate Limits: Controls requests per minute per user for this model. Prevents abuse and ensures fair access. Set based on the provider’s rate limits and your capacity:

- Conservative: 10-20 requests/minute
- Standard: 50-100 requests/minute

- High: 200+ requests/minute (requires provider rate limit increases)

Max Tokens: Limits context window and output length. Useful for controlling costs since longer contexts cost more. Set based on use case: - Short tasks (Q&A): 4,096 tokens - Medium tasks (writing): 16,384 tokens - Long tasks (analysis): 32,768+ tokens

Temperature Range: Restricts the temperature parameter users can set. Temperature controls randomness: - 0.0: Deterministic, consistent outputs - 0.7: Balanced creativity and consistency - 1.0+: More creative, less predictable

Restricting range (e.g., 0.0-1.0) prevents users from setting extreme values that produce poor results.

Price Override: Allows custom pricing different from the default. Useful for: - Offering discounts on specific models - Increasing prices for premium models - Matching competitor pricing - A/B testing pricing strategies

6.4 Self-Hosted Models

For Tier 3+ deployments:

1. Navigate to **Models → Self-Hosted**
2. Click “**+ Add Self-Hosted Model**”
3. Configure:
 - **Model ID:** Unique identifier
 - **SageMaker Endpoint:** Endpoint name
 - **Instance Type:** ml.g5.xlarge, etc.
 - **Auto-Scaling:** Min/max instances
4. Deploy model to SageMaker

6.5 Thermal States (Self-Hosted)

State	Description	Response Time
HOT	Always running	<100ms
WARM	Scaled down	<5s
COLD	Stopped	30-60s
OFF	Disabled	N/A

7. Provider Management

7.1 External Providers

Navigate to **Providers** to manage API integrations:

Provider	Models	Status	Health
OpenAI	12	Configured	99.9%
Anthropic	6	Configured	99.8%
Google AI	8	Configured	99.7%
xAI	2	Configured	99.5%

Provider	Models	Status	Health
DeepSeek	4	Not configured	-

7.2 Adding Provider Credentials

1. Click on provider name
2. Click “Configure”
3. Enter API credentials:
 - **API Key:** Provider API key
 - **Organization ID:** (if applicable)
 - **Base URL:** (for custom endpoints)
4. Click “Test Connection”
5. Click “Save”

7.3 Provider Health Monitoring

View real-time provider health:

Provider Health: OpenAI

Status:	Healthy
Uptime (30d):	99.94%
Avg Latency:	245ms
P95 Latency:	520ms
Error Rate:	0.02%

Last 24 Hours:

12am 6am 12pm 6pm 12am

7.4 Fallback Configuration

Configure provider fallbacks:

1. Navigate to **Providers → Fallbacks**
2. Set priority order for each model category
3. Configure automatic failover rules
4. Set retry policies

8. Billing & Subscriptions

8.1 Subscription Tiers

Tier	Name	Monthly	Features
1	Free	\$0	Basic models, 1K requests
2	Starter	\$29	More models, 10K requests
3	Professional	\$99	All external models, 100K requests
4	Business	\$299	Priority support, 500K requests
5	Enterprise	\$999	Self-hosted, unlimited
6	Enterprise+	Custom	Custom SLAs, dedicated support
7	Ultimate	Custom	On-premise options

8.2 Credit System

Credits are the universal currency for AI usage:

Model Type	Cost per 1M Tokens
GPT-4o	500 credits
GPT-4o-mini	50 credits
Claude 3 Opus	600 credits
Claude 3 Sonnet	150 credits
Self-hosted	20 credits

8.3 Managing Subscriptions

Navigate to **Billing → Subscriptions**:

1. View current subscription
2. Upgrade/downgrade tier
3. Add credit packages
4. View invoices
5. Update payment method

8.4 Usage Reports

Generate usage reports:

1. Navigate to **Billing → Reports**
2. Select date range
3. Choose grouping (by tenant/model/user)
4. Export as CSV/PDF

8.5 Billing Alerts

Configure alerts for:

- Credit balance low
- Usage spike
- Approaching quota
- Failed payments

9. Storage Management

9.1 Storage Overview

Navigate to **Storage** to monitor:

Storage Overview

Total Used: 234.5 GB of 500 GB (47%)

By Type:

Documents:	120.3 GB (51%)
Images:	45.2 GB (19%)
Audio:	38.7 GB (17%)
Video:	22.1 GB (9%)
Other:	8.2 GB (4%)

Top Tenants:

1. Acme Corp	45.2 GB
2. TechStart	32.1 GB
3. DataCo	28.4 GB

9.2 Storage Tiers

Tier	Included	Additional
Free	1 GB	N/A
Starter	10 GB	\$0.10/GB
Professional	100 GB	\$0.08/GB
Business	500 GB	\$0.05/GB
Enterprise	2 TB	\$0.03/GB

9.3 File Management

Manage uploaded files:

- View file metadata
- Download files
- Delete files
- Set retention policies

10. Orchestration & Neural Engine

10.1 Brain Router

The Brain Router automatically selects optimal models:

Factor	Weight	Description
Cost	30%	Price optimization
Quality	30%	Output quality
Speed	20%	Response latency
Availability	20%	Provider health

10.2 Neural Patterns

Configure orchestration patterns:

Pattern	Description	Use Case
Single	One model	Simple requests
Fallback	Primary + backup	High availability
Parallel	Multiple simultaneous	Consensus
Chain	Sequential models	Complex tasks

10.3 Workflow Templates

Create reusable workflows:

1. Navigate to **Orchestration → Workflows**
 2. Click “+ New Workflow”
 3. Define steps and conditions
 4. Set triggers and parameters
 5. Save and activate
-

11. Localization

11.1 Translation Management

Navigate to **Localization** to manage:

- Supported languages
- Translation strings
- AI translation settings

11.2 Supported Languages

Language	Code	Status
English	en	Default

Language	Code	Status
Spanish	es	Enabled
French	fr	Enabled
German	de	Enabled
Japanese	ja	Enabled
Chinese	zh	Enabled

11.3 AI Translation

Enable AI-powered translation:

1. Navigate to **Localization → Settings**
 2. Enable “**AI Translation**”
 3. Select translation model
 4. Configure quality settings
-

12. Configuration Management

12.1 System Configuration

Navigate to **Configuration** to manage:

Category	Settings
General	Platform name, domain, timezone
Email	SMTP settings, templates
Security	Password policy, MFA settings
API	Rate limits, CORS settings
Features	Feature flags

12.2 Tenant Overrides

Allow tenant-specific configuration:

1. Navigate to **Configuration → Tenant Overrides**
2. Select tenant
3. Override specific settings
4. Save changes

12.3 SSM Parameters

System configuration is stored in AWS SSM:

Parameter	Description
/radiant/prod/database/url	Database connection
/radiant/prod/api/rate-limit	API rate limits
/radiant/prod/features/*	Feature flags

13. Security & Compliance

13.1 Security Dashboard

Navigate to **Security** to monitor:

Security Dashboard

Threat Level: Low

Active Threats: 0
Failed Logins: 23 (last 24h)
Suspicious IPs: 2 blocked
MFA Adoption: 94%

Recent Alerts:

Unusual login location - user@acme.com (2h ago)
Resolved: Brute force attempt blocked (5h ago)
Resolved: API key rotated - tenant xyz (1d ago)

13.2 Anomaly Detection

Automatic detection of:

- Impossible travel (geographic anomalies)
- Session hijacking attempts
- Brute force attacks
- Unusual API patterns

13.3 Compliance Reports

Navigate to **Compliance** to generate:

Framework	Description
SOC 2	Service organization controls
HIPAA	Healthcare data protection
GDPR	EU data protection
ISO 27001	Information security

13.4 Generating Reports

1. Click “**Generate Report**”
2. Select framework

3. Choose date range
 4. Select metrics to include
 5. Generate PDF/CSV
-

14. Cost Analytics

14.1 Cost Dashboard

Navigate to **Cost** to view:

Cost Analytics

Period: Last 30 Days

Total Spend:	\$12,456.78	(+12% vs last month)
Projected:	\$14,200.00	(this month)

By Provider:

OpenAI:	\$6,234.56 (50%)
Anthropic:	\$3,456.78 (28%)
Self-hosted:	\$1,234.56 (10%)
Other:	\$1,530.88 (12%)

AI Recommendations:

Switch 23% of GPT-4 calls to GPT-4-mini (save \$890/mo)
Enable caching for repeated queries (save \$340/mo)

14.2 Cost Alerts

Configure alerts:

- Daily budget exceeded
- Weekly spend spike
- Per-tenant limits
- Per-model thresholds

14.3 Cost Optimization

Review AI-powered recommendations:

1. Navigate to **Cost → Insights**
 2. Review suggestions
 3. Click “**Apply**” to implement (requires approval)
 4. Track savings over time
-

15. A/B Testing & Experiments

15.1 Experiment Dashboard

Navigate to **Experiments** to manage:

Experiment	Status	Variants	Sample Size
Model routing v2	Running	3	45,234
Prompt optimization	Running	2	12,456
Temperature test	Completed	4	89,123

15.2 Creating an Experiment

1. Click “+ New Experiment”
2. Configure:
 - **Name:** Descriptive name
 - **Hypothesis:** What you’re testing
 - **Variants:** Control + treatments
 - **Traffic Split:** Percentage per variant
 - **Success Metric:** What to measure
3. Set targeting rules
4. Start experiment

15.3 Statistical Analysis

View results with:

- Conversion rates per variant
- Statistical significance (p-value)
- Confidence intervals
- Sample size recommendations

16. Audit & Monitoring

16.1 Audit Logs

Navigate to **Audit** to view all actions:

Column	Description
Timestamp	When action occurred
Actor	Who performed action
Action	What was done
Resource	What was affected
IP Address	Source IP
Details	Additional context

16.2 Log Filtering

Filter by:

- Date range
- Actor (user/admin)
- Action type
- Resource type
- Severity level

16.3 Log Export

Export logs for compliance:

1. Set filter criteria
2. Click “**Export**”
3. Choose format (CSV/JSON)
4. Download file

16.4 Real-Time Monitoring

Navigate to **Monitoring** for:

- Live request stream
 - Error rate graphs
 - Latency percentiles
 - Active users count
-

17. Database Migrations

17.1 Migration Workflow

RADIANT uses dual-admin approval for production migrations:

1. **Submit**: Admin submits migration
2. **Review**: Second admin reviews
3. **Approve**: Second admin approves
4. **Execute**: Migration runs
5. **Verify**: Automatic verification

17.2 Pending Migrations

Navigate to **Migrations** to see:

Database Migrations

Pending Approval:

#045 - Add user preferences table

Submitted by: alice@company.com (2 hours ago)
[View SQL] [Approve] [Reject]

Recent Migrations:

- #044 - Cost tracking tables (applied 2024-12-24)
- #043 - Experiment framework (applied 2024-12-20)
- #042 - Security anomalies (applied 2024-12-15)

17.3 Approving Migrations

1. Review the SQL in “**View SQL**”
 2. Check for potential issues
 3. Click “**Approve**” or “**Reject**”
 4. Add comment explaining decision
-

18. API Management

18.1 API Keys

Manage platform API keys:

1. Navigate to **Settings** → **API Keys**
2. View existing keys
3. Create new keys with scopes
4. Revoke compromised keys

18.2 Rate Limiting

Configure rate limits:

Level	Default	Configurable
Global	10,000/min	Yes
Per-Tenant	1,000/min	Yes
Per-User	100/min	Yes
Per-Key	60/min	Yes

18.3 Webhooks

Configure outgoing webhooks:

1. Navigate to **Settings** → **Webhooks**
2. Add webhook URL
3. Select events to send
4. Test webhook
5. Enable webhook

19. Troubleshooting

19.1 Common Issues

High Error Rate

1. Check **Providers** for unhealthy providers
2. Review **Audit** logs for patterns
3. Check **Monitoring** for load spikes
4. Verify API key validity

Slow Response Times

1. Check provider latency in **Providers**
2. Review model selection in **Orchestration**
3. Check for cold-start issues (self-hosted)
4. Verify database performance

Authentication Failures

1. Check user status in **Users**
2. Verify MFA configuration
3. Review **Audit** logs for login attempts
4. Check for IP blocks in **Security**

19.2 Support Resources

Resource	Description
Documentation	This guide + online docs
Status Page	status.radiant.example.com
Support Email	support@radiant.example.com
Emergency	+1-555-RADIANT

19.3 Log Locations

Service	Log Group
API Gateway	/aws/apigateway/radiant
Lambda	/aws/lambda/radiant-*
Admin Dashboard	/aws/cloudfront/admin
Database	/aws/rds/cluster/radiant

Appendix: Quick Reference

Keyboard Shortcuts

Shortcut	Action
G + D	Go to Dashboard
G + T	Go to Tenants
G + M	Go to Models
G + B	Go to Billing
G + A	Go to Audit
?	Show shortcuts

Status Indicators

Icon	Meaning
Green circle	Healthy/Success
Yellow triangle	Warning
Red circle	Error/Failed
Blue circle	In Progress
Grey circle	Disabled/Pending

Document Version: 4.18.1 Last Updated: December 2024